
Julia Language Documentation

Release 0.6.0-dev

Jeff Bezanson, Stefan Karpinski, Viral Shah, Alan Edelman, et al.

December 07, 2016

I	The Julia Manual	1
1	Introduction	3
2	Getting Started	5
3	Variables	9
4	Integers and Floating-Point Numbers	13
5	Mathematical Operations and Elementary Functions	25
6	Complex and Rational Numbers	35
7	Strings	41
8	Functions	55
9	Control Flow	65
10	Scope of Variables	81
11	Types	89
12	Methods	109
13	Constructors	117
14	Conversion and Promotion	127
15	Interfaces	133
16	Modules	139
17	Documentation	147
18	Metaprogramming	161
19	Multi-dimensional Arrays	177
20	Linear algebra	189

21	Networking and Streams	193
22	Parallel Computing	199
23	Date and DateTime	221
24	Running External Programs	231
25	Calling C and Fortran Code	237
26	Handling Operating System Variation	255
27	Interacting With Julia	257
28	Embedding Julia	263
29	Packages	271
30	Package Development	279
31	Profiling	291
32	Memory allocation analysis	297
33	Stack Traces	299
34	Performance Tips	305
35	Workflow Tips	325
36	Style Guide	327
37	Frequently Asked Questions	335
38	Noteworthy Differences from other Languages	347
39	Unicode Input	355
II	The Julia Standard Library	357
40	Essentials	359
41	Collections and Data Structures	381
42	Mathematics	405
43	Numbers	433
44	Strings	443
45	Arrays	451
46	Tasks and Parallel Computing	485
47	Linear Algebra	501
48	Constants	553

49 Filesystem	555
50 I/O and Network	561
51 Punctuation	575
52 Sorting and Related Functions	577
53 Package Manager Functions	583
54 Dates and Time	587
55 Iteration utilities	597
56 Unit Testing	601
57 C Interface	607
58 LLVM Interface	611
59 C Standard Library	613
60 Dynamic Linker	615
61 Profiling	617
62 SIMD Support	619
Bibliography	621

Part I

The Julia Manual

Introduction

Scientific computing has traditionally required the highest performance, yet domain experts have largely moved to slower dynamic languages for daily work. We believe there are many good reasons to prefer dynamic languages for these applications, and we do not expect their use to diminish. Fortunately, modern language design and compiler techniques make it possible to mostly eliminate the performance trade-off and provide a single environment productive enough for prototyping and efficient enough for deploying performance-intensive applications. The Julia programming language fills this role: it is a flexible dynamic language, appropriate for scientific and numerical computing, with performance comparable to traditional statically-typed languages.

Because Julia’s compiler is different from the interpreters used for languages like Python or R, you may find that Julia’s performance is unintuitive at first. If you find that something is slow, we highly recommend reading through the *Performance Tips* (page 305) section before trying anything else. Once you understand how Julia works, it’s easy to write code that’s nearly as fast as C.

Julia features optional typing, multiple dispatch, and good performance, achieved using type inference and *just-in-time (JIT) compilation*, implemented using *LLVM*. It is multi-paradigm, combining features of imperative, functional, and object-oriented programming. Julia provides ease and expressiveness for high-level numerical computing, in the same way as languages such as R, MATLAB, and Python, but also supports general programming. To achieve this, Julia builds upon the lineage of mathematical programming languages, but also borrows much from popular dynamic languages, including *Lisp*, *Perl*, *Python*, *Lua*, and *Ruby*.

The most significant departures of Julia from typical dynamic languages are:

- The core language imposes very little; the standard library is written in Julia itself, including primitive operations like integer arithmetic
- A rich language of types for constructing and describing objects, that can also optionally be used to make type declarations
- The ability to define function behavior across many combinations of argument types via *multiple dispatch*
- Automatic generation of efficient, specialized code for different argument types
- Good performance, approaching that of statically-compiled languages like C

Although one sometimes speaks of dynamic languages as being “typeless”, they are definitely not: every object, whether primitive or user-defined, has a type. The lack of type declarations in most dynamic languages, however, means that one cannot instruct the compiler about the types of values, and often cannot explicitly talk about types at all. In static languages, on the other hand, while one can — and usually must — annotate types for the compiler, types exist only at compile time and cannot be manipulated or expressed at run time. In Julia, types are themselves run-time objects, and can also be used to convey information to the compiler.

While the casual programmer need not explicitly use types or multiple dispatch, they are the core unifying features of Julia: functions are defined on different combinations of argument types, and applied by dispatching to the most specific matching definition. This model is a good fit for mathematical programming, where it is unnatural for the first

argument to “own” an operation as in traditional object-oriented dispatch. Operators are just functions with special notation — to extend addition to new user-defined data types, you define new methods for the `+` function. Existing code then seamlessly applies to the new data types.

Partly because of run-time type inference (augmented by optional type annotations), and partly because of a strong focus on performance from the inception of the project, Julia’s computational efficiency exceeds that of other dynamic languages, and even rivals that of statically-compiled languages. For large scale numerical problems, speed always has been, continues to be, and probably always will be crucial: the amount of data being processed has easily kept pace with Moore’s Law over the past decades.

Julia aims to create an unprecedented combination of ease-of-use, power, and efficiency in a single language. In addition to the above, some advantages of Julia over comparable systems include:

- Free and open source ([MIT licensed](#))
- User-defined types are as fast and compact as built-ins
- No need to vectorize code for performance; devectorized code is fast
- Designed for parallelism and distributed computation
- Lightweight “green” threading ([coroutines](#))
- Unobtrusive yet powerful type system
- Elegant and extensible conversions and promotions for numeric and other types
- Efficient support for [Unicode](#), including but not limited to [UTF-8](#)
- Call C functions directly (no wrappers or special APIs needed)
- Powerful shell-like capabilities for managing other processes
- Lisp-like macros and other metaprogramming facilities

Getting Started

Julia installation is straightforward, whether using precompiled binaries or compiling from source. Download and install Julia by following the instructions at <http://julialang.org/downloads/>.

The easiest way to learn and experiment with Julia is by starting an interactive session (also known as a read-eval-print loop or “repl”) by double-clicking the Julia executable or running `julia` from the command line:

```
$ julia

      _       _ _(_)_      | A fresh approach to technical computing
  (_)_      | (_)_(_)      | Documentation: http://docs.julialang.org
    _ _    _|_|_ _ _ _    | Type "?help" for help.
  | | | | | | | | / _ ` | |
  | | | | | | | | (| | | | | Version 0.5.0-dev+2440 (2016-02-01 02:22 UTC)
 _/ | \_ ' _|_|_ | \_ ' _| | Commit 2bb94d6 (11 days old master)
|__/_/                          | x86_64-apple-darwin13.1.0

julia> 1 + 2
3

julia> ans
3
```

To exit the interactive session, type `^D` — the control key together with the `d` key or type `quit()`. When run in interactive mode, `julia` displays a banner and prompts the user for input. Once the user has entered a complete expression, such as `1 + 2`, and hits enter, the interactive session evaluates the expression and shows its value. If an expression is entered into an interactive session with a trailing semicolon, its value is not shown. The variable `ans` is bound to the value of the last evaluated expression whether it is shown or not. The `ans` variable is only bound in interactive sessions, not when Julia code is run in other ways.

To evaluate expressions written in a source file `file.jl`, write `include("file.jl")`.

To run code in a file non-interactively, you can give it as the first argument to the `julia` command:

```
$ julia script.jl arg1 arg2...
```

As the example implies, the following command-line arguments to `julia` are taken as command-line arguments to the program `script.jl`, passed in the global constant `ARGS`. The name of the script itself is passed in as the global `PROGRAM_FILE`. Note that `ARGS` is also set when script code is given using the `-e` option on the command line (see the `julia help` output below) but `PROGRAM_FILE` will be empty. For example, to just print the arguments given to a script, you could do this:

```
$ julia -e 'println(PROGRAM_FILE); for x in ARGS; println(x); end' foo bar
```

```
foo
bar
```

Or you could put that code into a script and run it:

```
$ echo 'println(PROGRAM_FILE); for x in ARGS; println(x); end' > script.jl
$ julia script.jl foo bar
script.jl
foo
bar
```

The `--` delimiter can be used to separate command-line args to the scriptfile from args to Julia:

```
$ julia --color=yes -O -- foo.jl arg1 arg2..
```

Julia can be started in parallel mode with either the `-p` or the `--machinefile` options. `-p n` will launch an additional `n` worker processes, while `--machinefile file` will launch a worker for each line in file `file`. The machines defined in `file` must be accessible via a passwordless `ssh` login, with Julia installed at the same location as the current host. Each machine definition takes the form `[count*][user@]host[:port][bind_addr[:port]]`. `user` defaults to current user, `port` to the standard `ssh` port. `count` is the number of workers to spawn on the node, and defaults to 1. The optional `bind-to` `bind_addr[:port]` specifies the ip-address and port that other workers should use to connect to this worker.

If you have code that you want executed whenever Julia is run, you can put it in `~/.juliarc.jl`:

```
$ echo 'println("Greetings! 好! 안녕하세요?")' > ~/.juliarc.jl
$ julia
Greetings! 好! 안녕하세요?
...

```

There are various ways to run Julia code and provide options, similar to those available for the `perl` and `ruby` programs:

```
julia [switches] -- [programfile] [args...]
-v, --version          Display version information
-h, --help             Print this message

-J, --sysimage <file>  Start up with the given system image file
--precompiled={yes|no} Use precompiled code from system image if available
--compilecache={yes|no} Enable/disable incremental precompilation of modules
-H, --home <dir>       Set location of `julia` executable
--startup-file={yes|no} Load ~/.juliarc.jl
--handle-signals={yes|no} Enable or disable Julia's default signal handlers

-e, --eval <expr>      Evaluate <expr>
-E, --print <expr>     Evaluate and show <expr>
-L, --load <file>      Load <file> immediately on all processors

-p, --procs {N|auto}   Integer value N launches N additional local worker processes
                        "auto" launches as many workers as the number of local cores
--machinefile <file>   Run processes on hosts listed in <file>

-i                     Interactive mode; REPL runs and isinteractive() is true
-q, --quiet            Quiet startup (no banner)
--color={yes|no}       Enable or disable color text
--history-file={yes|no} Load or save history

--compile={yes|no|all|min} Enable or disable JIT compiler, or request exhaustive compilation
```

```

-C, --cpu-target <target> Limit usage of cpu features up to <target>
-O, --optimize={0,1,2,3} Set the optimization level (default is 2 if unspecified or 3 if specified)
-g, -g <level> Enable / Set the level of debug info generation (default is 1 if unspecified)
--inline={yes|no} Control whether inlining is permitted (overrides functions declared as @inline)
--check-bounds={yes|no} Emit bounds checks always or never (ignoring declarations)
--math-mode={ieee,fast} Disallow or enable unsafe floating point optimizations (overrides @fastmath)

--depwarn={yes|no|error} Enable or disable syntax and method deprecation warnings ("error" turns warnings into errors)

--output-o name Generate an object file (including system image data)
--output-ji name Generate a system image data file (.ji)
--output-bc name Generate LLVM bitcode (.bc)
--output-incremental=no Generate an incremental output file (rather than complete)

--code-coverage={none|user|all}, --code-coverage Count executions of source lines (omitting setting is equivalent to "user")
--track-allocation={none|user|all}, --track-allocation Count bytes allocated by each source line

```

2.1 Resources

In addition to this manual, there are various other resources that may help new users get started with Julia:

- [Julia and IJulia cheatsheet](#)
- [Learn Julia in a few minutes](#)
- [Learn Julia the Hard Way](#)
- [Julia by Example](#)
- [Hands-on Julia](#)
- [Tutorial for Homer Reid's numerical analysis class](#)
- [An introductory presentation](#)
- [Videos from the Julia tutorial at MIT](#)
- [Forio Julia Tutorials](#)
- [YouTube videos from the JuliaCons](#)

Variables

A variable, in Julia, is a name associated (or bound) to a value. It's useful when you want to store a value (that you obtained after some math, for example) for later use. For example:

```
# Assign the value 10 to the variable x
julia> x = 10
10

# Doing math with x's value
julia> x + 1
11

# Reassign x's value
julia> x = 1 + 1
2

# You can assign values of other types, like strings of text
julia> x = "Hello World!"
"Hello World!"
```

Julia provides an extremely flexible system for naming variables. Variable names are case-sensitive, and have no semantic meaning (that is, the language will not treat variables differently based on their names).

```
julia> x = 1.0
1.0

julia> y = -3
-3

julia> Z = "My string"
"My string"

julia> customary_phrase = "Hello world!"
"Hello world!"

julia> UniversalDeclarationOfHumanRightsStart = "人人生而自由，在尊严和权利上一律平等。"
"人人生而自由，在尊严和权利上一律平等。"
```

Unicode names (in UTF-8 encoding) are allowed:

```
julia> δ = 0.00001
1.0e-5

julia> 안녕하세요 = "Hello"
"Hello"
```

In the Julia REPL and several other Julia editing environments, you can type many Unicode math symbols by typing the backslashed LaTeX symbol name followed by tab. For example, the variable name δ can be entered by typing `\delta-tab`, or even α^2 by `\alpha-tab-\hat-tab-_2-tab`. Julia will even let you redefine built-in constants and functions if needed:

```
julia> pi
π = 3.1415926535897...

julia> pi = 3
WARNING: imported binding for pi overwritten in module Main
3

julia> pi
3

julia> sqrt(100)
10.0

julia> sqrt = 4
WARNING: imported binding for sqrt overwritten in module Main
4
```

However, this is obviously not recommended to avoid potential confusion.

3.1 Allowed Variable Names

Variable names must begin with a letter (A-Z or a-z), underscore, or a subset of Unicode code points greater than 00A0; in particular, [Unicode character categories](#) Lu/Ll/Lt/Lm/Lo/Nl (letters), Sc/So (currency and other symbols), and a few other letter-like characters (e.g. a subset of the Sm math symbols) are allowed. Subsequent characters may also include ! and digits (0-9 and other characters in categories Nd/No), as well as other Unicode code points: diacritics and other modifying marks (categories Mn/Mc/Me/Sk), some punctuation connectors (category Pc), primes, and a few other characters.

Operators like `+` are also valid identifiers, but are parsed specially. In some contexts, operators can be used just like variables; for example `(+)` refers to the addition function, and `(+) = f` will reassign it. Most of the Unicode infix operators (in category Sm), such as \oplus , are parsed as infix operators and are available for user-defined methods (e.g. you can use `const \otimes = kron` to define \otimes as an infix Kronecker product).

The only explicitly disallowed names for variables are the names of built-in statements:

```
julia> else = false
ERROR: syntax: unexpected "else"
...

julia> try = "No"
ERROR: syntax: unexpected "="
...
```

3.2 Stylistic Conventions

While Julia imposes few restrictions on valid names, it has become useful to adopt the following conventions:

- Names of variables are in lower case.

- Word separation can be indicated by underscores (`'_'`), but use of underscores is discouraged unless the name would be hard to read otherwise.
- Names of `Types` and `Modules` begin with a capital letter and word separation is shown with upper camel case instead of underscores.
- Names of `functions` and `macros` are in lower case, without underscores.
- Functions that write to their arguments have names that end in `!`. These are sometimes called “mutating” or “in-place” functions because they are intended to produce changes in their arguments after the function is called, not just return a value.

Integers and Floating-Point Numbers

Integers and floating-point values are the basic building blocks of arithmetic and computation. Built-in representations of such values are called numeric primitives, while representations of integers and floating-point numbers as immediate values in code are known as numeric literals. For example, `1` is an integer literal, while `1.0` is a floating-point literal; their binary in-memory representations as objects are numeric primitives.

Julia provides a broad range of primitive numeric types, and a full complement of arithmetic and bitwise operators as well as standard mathematical functions are defined over them. These map directly onto numeric types and operations that are natively supported on modern computers, thus allowing Julia to take full advantage of computational resources. Additionally, Julia provides software support for *Arbitrary Precision Arithmetic* (page 21), which can handle operations on numeric values that cannot be represented effectively in native hardware representations, but at the cost of relatively slower performance.

The following are Julia's primitive numeric types:

- **Integer types:**

Type	Signed?	Number of bits	Smallest value	Largest value
<code>Int8</code>	x	8	-2^7	$2^7 - 1$
<code>UInt8</code>		8	0	$2^8 - 1$
<code>Int16</code>	x	16	-2^{15}	$2^{15} - 1$
<code>UInt16</code>		16	0	$2^{16} - 1$
<code>Int32</code>	x	32	-2^{31}	$2^{31} - 1$
<code>UInt32</code>		32	0	$2^{32} - 1$
<code>Int64</code>	x	64	-2^{63}	$2^{63} - 1$
<code>UInt64</code>		64	0	$2^{64} - 1$
<code>Int128</code>	x	128	-2^{127}	$2^{127} - 1$
<code>UInt128</code>		128	0	$2^{128} - 1$
<code>Bool</code>	N/A	8	false (0)	true (1)

- **Floating-point types:**

Type	Precision	Number of bits
<code>Float16</code>	half	16
<code>Float32</code> (page 437)	single	32
<code>Float64</code> (page 437)	double	64

Additionally, full support for *Complex and Rational Numbers* (page 35) is built on top of these primitive numeric types. All numeric types interoperate naturally without explicit casting, thanks to a flexible, user-extensible *type promotion system* (page 127).

4.1 Integers

Literal integers are represented in the standard manner:

```
julia> 1
1

julia> 1234
1234
```

The default type for an integer literal depends on whether the target system has a 32-bit architecture or a 64-bit architecture:

```
# 32-bit system:
julia> typeof(1)
Int32

# 64-bit system:
julia> typeof(1)
Int64
```

The Julia internal variable `Sys.WORD_SIZE` (page 553) indicates whether the target system is 32-bit or 64-bit.:

```
# 32-bit system:
julia> Sys.WORD_SIZE
32

# 64-bit system:
julia> Sys.WORD_SIZE
64
```

Julia also defines the types `Int` and `UInt`, which are aliases for the system’s signed and unsigned native integer types respectively.:

```
# 32-bit system:
julia> Int
Int32
julia> UInt
UInt32

# 64-bit system:
julia> Int
Int64
julia> UInt
UInt64
```

Larger integer literals that cannot be represented using only 32 bits but can be represented in 64 bits always create 64-bit integers, regardless of the system type:

```
# 32-bit or 64-bit system:
julia> typeof(3000000000)
Int64
```

Unsigned integers are input and output using the `0x` prefix and hexadecimal (base 16) digits `0–9a–f` (the capitalized digits `A–F` also work for input). The size of the unsigned value is determined by the number of hex digits used:

```
julia> 0x1
0x01
```

```
julia> typeof(ans)
UInt8

julia> 0x123
0x0123

julia> typeof(ans)
UInt16

julia> 0x1234567
0x01234567

julia> typeof(ans)
UInt32

julia> 0x123456789abcdef
0x0123456789abcdef

julia> typeof(ans)
UInt64
```

This behavior is based on the observation that when one uses unsigned hex literals for integer values, one typically is using them to represent a fixed numeric byte sequence, rather than just an integer value.

Recall that the variable `ans` (page 362) is set to the value of the last expression evaluated in an interactive session. This does not occur when Julia code is run in other ways.

Binary and octal literals are also supported:

```
julia> 0b10
0x02

julia> typeof(ans)
UInt8

julia> 0o10
0x08

julia> typeof(ans)
UInt8
```

The minimum and maximum representable values of primitive numeric types such as integers are given by the `typemin()` (page 365) and `typemax()` (page 365) functions:

```
julia> (typemin{Int32}, typemax{Int32})
(-2147483648, 2147483647)

julia> for T in [Int8, Int16, Int32, Int64, Int128, UInt8, UInt16, UInt32, UInt64, UInt128]
    println("$($lpad(T, 7)): [$($typemin(T)), $($typemax(T))]")
end
  Int8: [-128, 127]
  Int16: [-32768, 32767]
  Int32: [-2147483648, 2147483647]
  Int64: [-9223372036854775808, 9223372036854775807]
Int128: [-170141183460469231731687303715884105728, 170141183460469231731687303715884105727]
  UInt8: [0, 255]
  UInt16: [0, 65535]
  UInt32: [0, 4294967295]
  UInt64: [0, 18446744073709551615]
Int128: [0, 340282366920938463463374607431768211455]
```

The values returned by `typemin()` (page 365) and `typemax()` (page 365) are always of the given argument type. (The above expression uses several features we have yet to introduce, including *for loops* (page 71), *Strings* (page 41), and *Interpolation* (page 46), but should be easy enough to understand for users with some existing programming experience.)

4.1.1 Overflow behavior

In Julia, exceeding the maximum representable value of a given type results in a wraparound behavior:

```
julia> x = typemax{Int64}
9223372036854775807

julia> x + 1
-9223372036854775808

julia> x + 1 == typemin{Int64}
true
```

Thus, arithmetic with Julia integers is actually a form of *modular arithmetic*. This reflects the characteristics of the underlying arithmetic of integers as implemented on modern computers. In applications where overflow is possible, explicit checking for wraparound produced by overflow is essential; otherwise, the `BigInt` type in *Arbitrary Precision Arithmetic* (page 21) is recommended instead.

4.1.2 Division errors

Integer division (the `div` function) has two exceptional cases: dividing by zero, and dividing the lowest negative number (`typemin()` (page 365)) by -1. Both of these cases throw a *DivideError* (page 374). The remainder and modulus functions (`rem` and `mod`) throw a *DivideError* (page 374) when their second argument is zero.

4.2 Floating-Point Numbers

Literal floating-point numbers are represented in the standard formats:

```
julia> 1.0
1.0

julia> 1.
1.0

julia> 0.5
0.5

julia> .5
0.5

julia> -1.23
-1.23

julia> 1e10
1.0e10
```

```
julia> 2.5e-4
0.00025
```

The above results are all `Float64` values. Literal `Float32` values can be entered by writing an `f` in place of `e`:

```
julia> 0.5f0
0.5f0

julia> typeof(ans)
Float32

julia> 2.5f-4
0.00025f0
```

Values can be converted to `Float32` easily:

```
julia> Float32(-1.5)
-1.5f0

julia> typeof(ans)
Float32
```

Hexadecimal floating-point literals are also valid, but only as `Float64` values:

```
julia> 0x1p0
1.0

julia> 0x1.8p3
12.0

julia> 0x.4p-1
0.125

julia> typeof(ans)
Float64
```

Half-precision floating-point numbers are also supported (`Float16`), but only as a storage format. In calculations they'll be converted to `Float32`:

```
julia> sizeof(Float16(4.))
2

julia> 2*Float16(4.)
Float16(8.0)
```

The underscore `_` can be used as digit separator:

```
julia> 10_000, 0.000_000_005, 0xdead_beef, 0b1011_0010
(10000, 5.0e-9, 0xdeadbeef, 0xb2)
```

4.2.1 Floating-point zero

Floating-point numbers have [two zeros](#), positive zero and negative zero. They are equal to each other but have different binary representations, as can be seen using the `bits` function: :

```
julia> 0.0 == -0.0
true

julia> bits(0.0)
```

```
"0000000000000000000000000000000000000000000000000000000000000000"

julia> bits(-0.0)
"1000000000000000000000000000000000000000000000000000000000000000"
```

4.2.2 Special floating-point values

There are three specified standard floating-point values that do not correspond to any point on the real number line:

Special value			Name	Description
Float16	Float32	Float64		
Inf16	Inf32	Inf	positive infinity	a value greater than all finite floating-point values
-Inf16	-Inf32	-Inf	negative infinity	a value less than all finite floating-point values
NaN16	NaN32	NaN	not a number	a value not == to any floating-point value (including itself)

For further discussion of how these non-finite floating-point values are ordered with respect to each other and other floats, see *Numeric Comparisons* (page 27). By the [IEEE 754 standard](#), these floating-point values are the results of certain arithmetic operations:

```
julia> 1/Inf
0.0

julia> 1/0
Inf

julia> -5/0
-Inf

julia> 0.000001/0
Inf

julia> 0/0
NaN

julia> 500 + Inf
Inf

julia> 500 - Inf
-Inf

julia> Inf + Inf
Inf

julia> Inf - Inf
NaN

julia> Inf * Inf
Inf

julia> Inf / Inf
NaN

julia> 0 * Inf
NaN
```


The `typemin()` (page 365) and `typemax()` (page 365) functions also apply to floating-point types:

```
julia> (typemin(Float16), typemax(Float16))
(-Inf16, Inf16)

julia> (typemin(Float32), typemax(Float32))
(-Inf32, Inf32)

julia> (typemin(Float64), typemax(Float64))
(-Inf, Inf)
```

4.2.3 Machine epsilon

Most real numbers cannot be represented exactly with floating-point numbers, and so for many purposes it is important to know the distance between two adjacent representable floating-point numbers, which is often known as [machine epsilon](#).

Julia provides `eps()` (page 366), which gives the distance between 1.0 and the next larger representable floating-point value:

```
julia> eps(Float32)
1.1920929f-7

julia> eps(Float64)
2.220446049250313e-16

julia> eps() # same as eps(Float64)
2.220446049250313e-16
```

These values are 2.0^{-23} and 2.0^{-52} as `Float32` and `Float64` values, respectively. The `eps()` (page 366) function can also take a floating-point value as an argument, and gives the absolute difference between that value and the next representable floating point value. That is, `eps(x)` yields a value of the same type as `x` such that `x + eps(x)` is the next representable floating-point value larger than `x`:

```
julia> eps(1.0)
2.220446049250313e-16

julia> eps(1000.)
1.1368683772161603e-13

julia> eps(1e-27)
1.793662034335766e-43

julia> eps(0.0)
5.0e-324
```

The distance between two adjacent representable floating-point numbers is not constant, but is smaller for smaller values and larger for larger values. In other words, the representable floating-point numbers are densest in the real number line near zero, and grow sparser exponentially as one moves farther away from zero. By definition, `eps(1.0)` is the same as `eps(Float64)` since 1.0 is a 64-bit floating-point value.

Julia also provides the `nextfloat()` (page 436) and `prevfloat()` (page 436) functions which return the next largest or smallest representable floating-point number to the argument respectively: :

```
julia> x = 1.25f0
1.25f0

julia> nextfloat(x)
```

```
1.2500001f0

julia> prevfloat(x)
1.2499999f0

julia> bits(prevfloat(x))
"00111111100111111111111111111111"

julia> bits(x)
"00111111101000000000000000000000"

julia> bits(nextfloat(x))
"00111111101000000000000000000001"
```

This example highlights the general principle that the adjacent representable floating-point numbers also have adjacent binary integer representations.

4.2.4 Rounding modes

If a number doesn't have an exact floating-point representation, it must be rounded to an appropriate representable value, however, if wanted, the manner in which this rounding is done can be changed according to the rounding modes presented in the [IEEE 754 standard](#).

```
julia> x = 1.1; y = 0.1;

julia> x + y
1.2000000000000002

julia> setrounding(Float64, RoundDown) do
           x + y
       end
1.2
```

The default mode used is always *RoundNearest* (page 417), which rounds to the nearest representable value, with ties rounded towards the nearest value with an even least significant bit.

Warning: Rounding is generally only correct for basic arithmetic functions (*+* (page 405), *-* (page 405), *** (page 443), */* (page 405) and *sqrt* (page 420)) and type conversion operations. Many other functions assume the default *RoundNearest* (page 417) mode is set, and can give erroneous results when operating under other rounding modes.

4.2.5 Background and References

Floating-point arithmetic entails many subtleties which can be surprising to users who are unfamiliar with the low-level implementation details. However, these subtleties are described in detail in most books on scientific computation, and also in the following references:

- The definitive guide to floating point arithmetic is the [IEEE 754-2008 Standard](#); however, it is not available for free online.
- For a brief but lucid presentation of how floating-point numbers are represented, see John D. Cook's [article](#) on the subject as well as his [introduction](#) to some of the issues arising from how this representation differs in behavior from the idealized abstraction of real numbers.
- Also recommended is Bruce Dawson's [series of blog posts on floating-point numbers](#).

- For an excellent, in-depth discussion of floating-point numbers and issues of numerical accuracy encountered when computing with them, see David Goldberg’s paper [What Every Computer Scientist Should Know About Floating-Point Arithmetic](#).
- For even more extensive documentation of the history of, rationale for, and issues with floating-point numbers, as well as discussion of many other topics in numerical computing, see the [collected writings](#) of William Kahan, commonly known as the “Father of Floating-Point”. Of particular interest may be [An Interview with the Old Man of Floating-Point](#).

4.3 Arbitrary Precision Arithmetic

To allow computations with arbitrary-precision integers and floating point numbers, Julia wraps the [GNU Multiple Precision Arithmetic Library \(GMP\)](#) and the [GNU MPFR Library](#), respectively. The [BigInt](#) (page 437) and [BigFloat](#) (page 437) types are available in Julia for arbitrary precision integer and floating point numbers respectively.

Constructors exist to create these types from primitive numerical types, and `parse()` (page 369) can be used to construct them from `AbstractStrings`. Once created, they participate in arithmetic with all other numeric types thanks to Julia's *type promotion and conversion mechanism* (page 127):

[illegible]

However, type promotion between the primitive types above and *BigInt* (page 437)/*BigFloat* (page 437) is not automatic and must be explicitly stated.

```
julia> x = typemin(Int64)
-9223372036854775808

julia> x = x - 1
9223372036854775807

julia> typeof(x)
Int64

julia> y = BigInt(typemin(Int64))
-9223372036854775808

julia> y = y - 1
-9223372036854775809

julia> typeof(y)
BigInt
```

The default precision (in number of bits of the significand) and rounding mode of *BigFloat* (page 437) operations can be changed globally by calling *setprecision()* (page 440) and *setrounding()* (page 437), and all further

calculations will take these changes in account. Alternatively, the precision or the rounding can be changed only within the execution of a particular block of code by using the same functions with a `do` block:

[illegible]

4.4 Numeric Literal Coefficients

To make common numeric formulas and expressions clearer, Julia allows variables to be immediately preceded by a numeric literal, implying multiplication. This makes writing polynomial expressions much cleaner:

```
julia> x = 3
3
julia> 2x^2 - 3x + 1
10
julia> 1.5x^2 - .5x + 1
13.0
```

It also makes writing exponential functions more elegant:

```
julia> 2^2x
64
```

The precedence of numeric literal coefficients is the same as that of unary operators such as negation. So $2^3 \times$ is parsed as $2^3 \times$, and 2×3 is parsed as $2 \times (3)$.

Numeric literals also work as coefficients to parenthesized expressions:

```
julia> 2(x-1)^2 - 3(x-1) + 1
3
```

Additionally, parenthesized expressions can be used as coefficients to variables, implying multiplication of the expression by the variable:

```
julia> (x-1)x
6
```

Neither juxtaposition of two parenthesized expressions, nor placing a variable before a parenthesized expression, however, can be used to imply multiplication:

```
julia> (x-1)(x+1)
ERROR: MethodError: objects of type Int64 are not callable
...

```

```
julia> x(x+1)
ERROR: MethodError: objects of type Int64 are not callable
...

```

Both expressions are interpreted as function application: any expression that is not a numeric literal, when immediately followed by a parenthetical, is interpreted as a function applied to the values in parentheses (see [Functions](#) (page 55) for more about functions). Thus, in both of these cases, an error occurs since the left-hand value is not a function.

The above syntactic enhancements significantly reduce the visual noise incurred when writing common mathematical formulae. Note that no whitespace may come between a numeric literal coefficient and the identifier or parenthesized expression which it multiplies.

4.4.1 Syntax Conflicts

Juxtaposed literal coefficient syntax may conflict with two numeric literal syntaxes: hexadecimal integer literals and engineering notation for floating-point literals. Here are some situations where syntactic conflicts arise:

- The hexadecimal integer literal expression `0xff` could be interpreted as the numeric literal `0` multiplied by the variable `xff`.
- The floating-point literal expression `1e10` could be interpreted as the numeric literal `1` multiplied by the variable `e10`, and similarly with the equivalent `E` form.

In both cases, we resolve the ambiguity in favor of interpretation as a numeric literals:

- Expressions starting with 0x are always hexadecimal literals.
- Expressions starting with a numeric literal followed by e or E are always floating-point literals.

4.5 Literal zero and one

Julia provides functions which return literal 0 and 1 corresponding to a specified type or the type of a given variable.

Function	Description
<i>zero</i> (<i>x</i>) (page 435)	Literal zero of type <i>x</i> or type of variable <i>x</i>
<i>one</i> (<i>x</i>) (page 435)	Literal one of type <i>x</i> or type of variable <i>x</i>

These functions are useful in *Numeric Comparisons* (page 27) to avoid overhead from unnecessary *type conversion* (page 127).

Examples:

[illegible]

Mathematical Operations and Elementary Functions

Julia provides a complete collection of basic arithmetic and bitwise operators across all of its numeric primitive types, as well as providing portable, efficient implementations of a comprehensive collection of standard mathematical functions.

5.1 Arithmetic Operators

The following [arithmetic operators](#) are supported on all primitive numeric types:

Expression	Name	Description
<code>+x</code>	unary plus	the identity operation
<code>-x</code>	unary minus	maps values to their additive inverses
<code>x + y</code>	binary plus	performs addition
<code>x - y</code>	binary minus	performs subtraction
<code>x * y</code>	times	performs multiplication
<code>x / y</code>	divide	performs division
<code>x \ y</code>	inverse divide	equivalent to <code>y / x</code>
<code>x ^ y</code>	power	raises <code>x</code> to the <code>y</code> th power
<code>x % y</code>	remainder	equivalent to <code>rem(x, y)</code>

as well as the negation on `Bool` types:

Expression	Name	Description
<code>!x</code>	negation	changes <code>true</code> to <code>false</code> and vice versa

Julia’s promotion system makes arithmetic operations on mixtures of argument types “just work” naturally and automatically. See [Conversion and Promotion](#) (page 127) for details of the promotion system.

Here are some simple examples using arithmetic operators:

```
julia> 1 + 2 + 3
6

julia> 1 - 2
-1

julia> 3*2/12
0.5
```

(By convention, we tend to space operators more tightly if they get applied before other nearby operators. For instance, we would generally write `-x + 2` to reflect that first `x` gets negated, and then 2 is added to that result.)

5.2 Bitwise Operators

The following [bitwise operators](#) are supported on all primitive integer types:

Expression	Name
<code>~x</code>	bitwise not
<code>x & y</code>	bitwise and
<code>x y</code>	bitwise or
<code>x ^ y</code>	bitwise xor (exclusive or)
<code>x >>> y</code>	logical shift right
<code>x >> y</code>	arithmetic shift right
<code>x << y</code>	logical/arithmetic shift left

Here are some examples with bitwise operators:

```
julia> ~123
-124

julia> 123 & 234
106

julia> 123 | 234
251

julia> 123 ^ 234
145

julia> xor(123, 234)
145

julia> ~UInt32(123)
0xffffffff84

julia> ~UInt8(123)
0x84
```

5.3 Updating operators

Every binary arithmetic and bitwise operator also has an updating version that assigns the result of the operation back into its left operand. The updating version of the binary operator is formed by placing a `=` immediately after the operator. For example, writing `x += 3` is equivalent to writing `x = x + 3`:

```
julia> x = 1
1

julia> x += 3
4

julia> x
4
```

The updating versions of all the binary arithmetic and bitwise operators are:

```
+= -= *= /= \= ÷= %= ^= &= |= = >>= >>= <<=
```


Note: An updating operator rebinds the variable on the left-hand side. As a result, the type of the variable may change.

```
julia> x = 0x01; typeof(x)
UInt8

julia> x *= 2 #Same as x = x * 2
2

julia> isa(x, Int)
true
```

5.4 Numeric Comparisons

Standard comparison operations are defined for all the primitive numeric types:

Operator	Name
<code>==</code> (page 409)	equality
<code>!=</code> (page 410) \neq (page 410)	inequality
<code><</code> (page 411)	less than
<code><=</code> (page 410) \leq (page 410)	less than or equal to
<code>></code> (page 411)	greater than
<code>>=</code> (page 410) \geq (page 410)	greater than or equal to

Here are some simple examples:

```
julia> 1 == 1
true

julia> 1 == 2
false

julia> 1 != 2
true

julia> 1 == 1.0
true

julia> 1 < 2
true

julia> 1.0 > 3
false

julia> 1 >= 1.0
true

julia> -1 <= 1
true

julia> -1 <= -1
true

julia> -1 <= -2
```

```
false

julia> 3 < -0.5
false
```

Integers are compared in the standard manner — by comparison of bits. Floating-point numbers are compared according to the [IEEE 754 standard](#):

- Finite numbers are ordered in the usual manner.
- Positive zero is equal but not greater than negative zero.
- `Inf` is equal to itself and greater than everything else except `NaN`.
- `-Inf` is equal to itself and less than everything else except `NaN`.
- `NaN` is not equal to, not less than, and not greater than anything, including itself.

The last point is potentially surprising and thus worth noting:

```
julia> NaN == NaN
false

julia> NaN != NaN
true

julia> NaN < NaN
false

julia> NaN > NaN
false
```

and can cause especial headaches with [Arrays](#) (page 177):

```
julia> [1 NaN] == [1 NaN]
false
```

Julia provides additional functions to test numbers for special values, which can be useful in situations like hash key comparisons:

Function	Tests if
isequal(x, y) (page 362)	<code>x</code> and <code>y</code> are identical
isfinite(x) (page 436)	<code>x</code> is a finite number
isinf(x) (page 436)	<code>x</code> is infinite
isnan(x) (page 436)	<code>x</code> is not a number

[isequal\(\)](#) (page 362) considers `NaN`s equal to each other:

```
julia> isequal(NaN, NaN)
true

julia> isequal([1 NaN], [1 NaN])
true

julia> isequal(NaN, NaN32)
true
```

[isequal\(\)](#) (page 362) can also be used to distinguish signed zeros:

```
julia> -0.0 == 0.0
true
```

```
julia> isequal(-0.0, 0.0)
false
```

Mixed-type comparisons between signed integers, unsigned integers, and floats can be tricky. A great deal of care has been taken to ensure that Julia does them correctly.

For other types, `isequal()` (page 362) defaults to calling `==()` (page 409), so if you want to define equality for your own types then you only need to add a `==()` (page 409) method. If you define your own equality function, you should probably define a corresponding `hash()` (page 363) method to ensure that `isequal(x, y)` implies `hash(x) == hash(y)`.

5.4.1 Chaining comparisons

Unlike most languages, with the notable exception of Python, comparisons can be arbitrarily chained:

```
julia> 1 < 2 <= 2 < 3 == 3 > 2 >= 1 == 1 < 3 != 5
true
```

Chaining comparisons is often quite convenient in numerical code. Chained comparisons use the `&&` operator for scalar comparisons, and the `&` (page 411) operator for elementwise comparisons, which allows them to work on arrays. For example, `0 .< A .< 1` gives a boolean array whose entries are true where the corresponding elements of `A` are between 0 and 1.

The operator `.<` (page 410) is intended for array objects; the operation `A .< B` is valid only if `A` and `B` have the same dimensions. The operator returns an array with boolean entries and with the same dimensions as `A` and `B`. Such operators are called *elementwise*; Julia offers a suite of elementwise operators: `.*` (page 406), `.+` (page 405), etc. Some of the elementwise operators can take a scalar operand such as the example `0 .< A .< 1` in the preceding paragraph. This notation means that the scalar operand should be replicated for each entry of the array.

Note the evaluation behavior of chained comparisons:

```
julia> v(x) = (println(x); x)
v (generic function with 1 method)

julia> v(1) < v(2) <= v(3)
2
1
3
true

julia> v(1) > v(2) <= v(3)
2
1
false
```

The middle expression is only evaluated once, rather than twice as it would be if the expression were written as `v(1) < v(2) && v(2) <= v(3)`. However, the order of evaluations in a chained comparison is undefined. It is strongly recommended not to use expressions with side effects (such as printing) in chained comparisons. If side effects are required, the short-circuit `&&` operator should be used explicitly (see *Short-Circuit Evaluation* (page 69)).

5.4.2 Operator Precedence

Julia applies the following order of operations, from highest precedence to lowest:

Category	Operators
Syntax	<code>.</code> followed by <code>:</code>
Exponentiation	<code>^</code> and its elementwise equivalent <code>.^</code>
Fractions	<code>//</code> and <code>./</code>
Multiplication	<code>*</code> <code>/</code> <code>%</code> <code>&</code> <code>\</code> and <code>.*</code> <code>./</code> <code>.*</code> <code>.\</code>
Bitshifts	<code><<</code> <code>>></code> <code>>>></code> and <code>.<<</code> <code>.>></code> <code>.>>></code>
Addition	<code>+</code> <code>-</code> <code> </code> and <code>.+</code> <code>.-</code>
Syntax	<code>:</code> <code>..</code> followed by <code> ></code>
Comparisons	<code>></code> <code><</code> <code>>=</code> <code><=</code> <code>==</code> <code>===</code> <code>!=</code> <code>!==</code> <code><:</code> and <code>.></code> <code>.<</code> <code>.>=</code> <code>.<=</code> <code>==</code> <code>!:=</code>
Control flow	<code>&&</code> followed by <code> </code> followed by <code>?</code>
Assignments	<code>=</code> <code>+=</code> <code>-=</code> <code>*=</code> <code>/=</code> <code>//=</code> <code>\=</code> <code>^=</code> <code>÷=</code> <code>%=</code> <code> =</code> <code>&=</code> <code>=</code> <code><<=</code> <code>>>=</code> <code>>>>=</code> and <code>.*=</code> <code>./=</code> <code>./.=</code> <code>.\=</code> <code>.^=</code> <code>÷=</code> <code>.*=</code>

5.4.3 Elementary Functions

Julia provides a comprehensive collection of mathematical functions and operators. These mathematical operations are defined over as broad a class of numerical values as permit sensible definitions, including integers, floating-point numbers, rationals, and complexes, wherever such definitions make sense.

Moreover, these functions (like any Julia function) can be applied in “vectorized” fashion to arrays and other collections with the syntax `f.(A)`, e.g. `sin.(A)` will compute the elementwise sine of each element of an array `A`. See [Dot Syntax for Vectorizing Functions](#) (page 63).

5.5 Numerical Conversions

Julia supports three forms of numerical conversion, which differ in their handling of inexact conversions.

- The notation `T(x)` or `convert(T, x)` converts `x` to a value of type `T`.
 - If `T` is a floating-point type, the result is the nearest representable value, which could be positive or negative infinity.
 - If `T` is an integer type, an `InexactError` is raised if `x` is not representable by `T`.
- `x % T` converts an integer `x` to a value of integer type `T` congruent to `x` modulo 2^n , where `n` is the number of bits in `T`. In other words, the binary representation is truncated to fit.
- The [Rounding functions](#) (page 31) take a type `T` as an optional argument. For example, `round{Int}(x)` is a shorthand for `Int(round(x))`.

The following examples show the different forms.

```
julia> Int8(127)
127

julia> Int8(128)
ERROR: InexactError()
in Int8{::Int64} at ./sysimg.jl:66
...

julia> Int8(127.0)
127
```

```

julia> Int8(3.14)
ERROR: InexactError()
  in convert(::Type{Int8}, ::Float64) at ./float.jl:635
  in Int8(::Float64) at ./sysimg.jl:66
...

julia> Int8(128.0)
ERROR: InexactError()
  in convert(::Type{Int8}, ::Float64) at ./float.jl:635
  in Int8(::Float64) at ./sysimg.jl:66
...

julia> 127 % Int8
127

julia> 128 % Int8
-128

julia> round(Int8, 127.4)
127

julia> round(Int8, 127.6)
ERROR: InexactError()
  in trunc(::Type{Int8}, ::Float64) at ./float.jl:628
  in round(::Type{Int8}, ::Float64) at ./float.jl:332
...

```

See *Conversion and Promotion* (page 127) for how to define your own conversions and promotions.

5.5.1 Rounding functions

Function	Description	Return type
<code>round(x)</code> (page 416)	round x to the nearest integer	<code>typeof(x)</code>
<code>round(T, x)</code> (page 416)	round x to the nearest integer	<code>T</code>
<code>floor(x)</code> (page 417)	round x towards $-\text{Inf}$	<code>typeof(x)</code>
<code>floor(T, x)</code> (page 417)	round x towards $-\text{Inf}$	<code>T</code>
<code>ceil(x)</code> (page 417)	round x towards $+\text{Inf}$	<code>typeof(x)</code>
<code>ceil(T, x)</code> (page 417)	round x towards $+\text{Inf}$	<code>T</code>
<code>trunc(x)</code> (page 417)	round x towards zero	<code>typeof(x)</code>
<code>trunc(T, x)</code> (page 417)	round x towards zero	<code>T</code>

5.5.2 Division functions

Function	Description
<code>div(x, y)</code> (page 407)	truncated division; quotient rounded towards zero
<code>fld(x, y)</code> (page 407)	floored division; quotient rounded towards $-\text{Inf}$
<code>cld(x, y)</code> (page 407)	ceiling division; quotient rounded towards $+\text{Inf}$
<code>rem(x, y)</code> (page 407)	remainder; satisfies $x == \text{div}(x, y) * y + \text{rem}(x, y)$; sign matches x
<code>mod(x, y)</code> (page 407)	modulus; satisfies $x == \text{fld}(x, y) * y + \text{mod}(x, y)$; sign matches y
<code>mod1(x, y)</code> (page 408)	<code>mod()</code> with offset 1; returns $r \in (0, y]$ for $y > 0$ or $r \in [y, 0)$ for $y < 0$, where $\text{mod}(r, y) == \text{mod}(x, y)$
<code>mod2pi(x)</code> (page 407)	modulus with respect to 2π ; $0 \leq \text{mod2pi}(x) < 2\pi$
<code>divrem(x, y)</code> (page 407)	returns $(\text{div}(x, y), \text{rem}(x, y))$
<code>fldmod(x, y)</code> (page 408)	returns $(\text{fld}(x, y), \text{mod}(x, y))$
<code>gcd(x, y, ...)</code> (page 421)	greatest positive common divisor of x, y, \dots
<code>lcm(x, y, ...)</code> (page 421)	least positive common multiple of x, y, \dots

5.5.3 Sign and absolute value functions

Function	Description
<code>abs(x)</code> (page 418)	a positive value with the magnitude of x
<code>abs2(x)</code> (page 419)	the squared magnitude of x
<code>sign(x)</code> (page 419)	indicates the sign of x , returning -1, 0, or +1
<code>signbit(x)</code> (page 419)	indicates whether the sign bit is on (true) or off (false)
<code>copysign(x, y)</code> (page 419)	a value with the magnitude of x and the sign of y
<code>flipsign(x, y)</code> (page 419)	a value with the magnitude of x and the sign of $x * y$

5.5.4 Powers, logs and roots

Function	Description
<code>sqrt(x)</code> (page 420) \sqrt{x}	square root of x
<code>cbrt(x)</code> (page 420) $\sqrt[3]{x}$	cube root of x
<code>hypot(x, y)</code> (page 414)	hypotenuse of right-angled triangle with other sides of length x and y
<code>exp(x)</code> (page 415)	natural exponential function at x
<code>expm1(x)</code> (page 415)	accurate $\exp(x) - 1$ for x near zero
<code>ldexp(x, n)</code> (page 415)	$x \cdot 2^n$ computed efficiently for integer values of n
<code>log(x)</code> (page 414)	natural logarithm of x
<code>log(b, x)</code> (page 414)	base b logarithm of x
<code>log2(x)</code> (page 415)	base 2 logarithm of x
<code>log10(x)</code> (page 415)	base 10 logarithm of x
<code>log1p(x)</code> (page 415)	accurate $\log(1+x)$ for x near zero
<code>exponent(x)</code> (page 434)	binary exponent of x
<code>significand(x)</code> (page 434)	binary significand (a.k.a. mantissa) of a floating-point number x

For an overview of why functions like `hypot()` (page 414), `expm1()` (page 415), and `log1p()` (page 415) are necessary and useful, see John D. Cook’s excellent pair of blog posts on the subject: `expm1`, `log1p`, `erfc`, and `hypot`.

5.5.5 Trigonometric and hyperbolic functions

All the standard trigonometric and hyperbolic functions are also defined:

<code>sin</code>	<code>cos</code>	<code>tan</code>	<code>cot</code>	<code>sec</code>	<code>csc</code>
<code>sinh</code>	<code>cosh</code>	<code>tanh</code>	<code>coth</code>	<code>sech</code>	<code>csch</code>
<code>asin</code>	<code>acos</code>	<code>atan</code>	<code>acot</code>	<code>asec</code>	<code>acsc</code>
<code>asinh</code>	<code>acosh</code>	<code>atanh</code>	<code>acoth</code>	<code>asech</code>	<code>acsch</code>
<code>sinc</code>	<code>cosc</code>	<code>atan2</code>			

These are all single-argument functions, with the exception of `atan2`, which gives the angle in radians between the x -axis and the point specified by its arguments, interpreted as x and y coordinates.

Additionally, `sinpi(x)` (page 412) and `cospi(x)` (page 412) are provided for more accurate computations of `sin(pi*x)` (page 412) and `cos(pi*x)` (page 412) respectively.

In order to compute trigonometric functions with degrees instead of radians, suffix the function with `d`. For example, `sind(x)` (page 412) computes the sine of x where x is specified in degrees. The complete list of trigonometric functions with degree variants is:

<code>sind</code>	<code>cosd</code>	<code>tand</code>	<code>cotd</code>	<code>secd</code>	<code>cscd</code>
<code>asind</code>	<code>acosd</code>	<code>atand</code>	<code>acotd</code>	<code>asecd</code>	<code>acscd</code>

5.5.6 Special functions

Function	Description
<code>erf(x)</code> (page 420)	error function at x
<code>erfc(x)</code> (page 420)	complementary error function, i.e. the accurate version of <code>1 - erf(x)</code>
<code>erfinv(x)</code> (page 420)	inverse function to <code>erf()</code> (page 420)
<code>erfcinv(x)</code>	inverse function to <code>erfc()</code> (page 420)
<code>erfi(x)</code> (page 420)	imaginary error function defined as $-im * erf(x)$
<code>erfcx(x)</code> (page 420)	scaled complementary error function, i.e. accurate <code>exp(x^2) * erfc(x)</code>

Table 5.1 – continued from previous page

Function	Description
<code>dawson(x)</code> (page 420)	scaled imaginary error function, a.k.a. Dawson function
<code>gamma(x)</code> (page 422)	gamma function at x
<code>lgamma(x)</code> (page 422)	accurate $\log(\text{gamma}(x))$ for large x
<code>lfact(x)</code> (page 422)	accurate $\log(\text{factorial}(x))$ for large x ; same as <code>lgamma(x)</code>
<code>digamma(x)</code> (page 422)	digamma function (i.e. the derivative of <code>lgamma()</code>)
<code>beta(x, y)</code> (page 424)	beta function at x, y
<code>lbeta(x, y)</code> (page 424)	accurate $\log(\text{beta}(x, y))$ for large x or y
<code>eta(x)</code> (page 424)	Dirichlet eta function at x
<code>zeta(x)</code> (page 424)	Riemann zeta function at x
<code>airy(z)</code> (page 422), <code>airyai(z)</code> (page 422), <code>airy(0, z)</code>	Airy Ai function at z
<code>airyprime(z)</code> (page 423), <code>airyaiprime(z)</code> (page 423), <code>airy(1, z)</code>	derivative of the Airy Ai function at z
<code>airybi(z)</code> (page 423), <code>airy(2, z)</code>	Airy Bi function at z
<code>airybiprime(z)</code> (page 423), <code>airy(3, z)</code>	derivative of the Airy Bi function at z
<code>airyx(z)</code> (page 423), <code>airyx(k, z)</code>	scaled Airy AI function and k th derivatives at z
<code>besselj(nu, z)</code> (page 423)	Bessel function of the first kind of order nu at z
<code>besselj0(z)</code> (page 423)	<code>besselj(0, z)</code>
<code>besselj1(z)</code> (page 423)	<code>besselj(1, z)</code>
<code>besseljx(nu, z)</code> (page 423)	scaled Bessel function of the first kind of order nu at z
<code>bessely(nu, z)</code> (page 423)	Bessel function of the second kind of order nu at z
<code>bessely0(z)</code> (page 423)	<code>bessely(0, z)</code>
<code>bessely1(z)</code> (page 423)	<code>bessely(1, z)</code>
<code>besselyx(nu, z)</code> (page 423)	scaled Bessel function of the second kind of order nu at z
<code>besselh(nu, k, z)</code> (page 423)	Bessel function of the third kind (a.k.a. Hankel function)
<code>hankelh1(nu, z)</code> (page 423)	<code>besselh(nu, 1, z)</code>
<code>hankelh1x(nu, z)</code> (page 423)	scaled <code>besselh(nu, 1, z)</code>
<code>hankelh2(nu, z)</code> (page 423)	<code>besselh(nu, 2, z)</code>
<code>hankelh2x(nu, z)</code> (page 423)	scaled <code>besselh(nu, 2, z)</code>
<code>besseli(nu, z)</code> (page 424)	modified Bessel function of the first kind of order nu at z
<code>besselix(nu, z)</code> (page 424)	scaled modified Bessel function of the first kind of order nu at z
<code>besselk(nu, z)</code> (page 424)	modified Bessel function of the second kind of order nu at z
<code>besselkx(nu, z)</code> (page 424)	scaled modified Bessel function of the second kind of order nu at z

Complex and Rational Numbers

Julia ships with predefined types representing both complex and rational numbers, and supports all *standard mathematical operations* (page 25) on them. *Conversion and Promotion* (page 127) are defined so that operations on any combination of predefined numeric types, whether primitive or composite, behave as expected.

6.1 Complex Numbers

The global constant *im* (page 435) is bound to the complex number *i*, representing the principal square root of -1. It was deemed harmful to co-opt the name *i* for a global constant, since it is such a popular index variable name. Since Julia allows numeric literals to be *juxtaposed with identifiers as coefficients* (page 22), this binding suffices to provide convenient syntax for complex numbers, similar to the traditional mathematical notation:

```
julia> 1 + 2im
1 + 2im
```

You can perform all the standard arithmetic operations with complex numbers:

```
julia> (1 + 2im) * (2 - 3im)
8 + 1im

julia> (1 + 2im) / (1 - 2im)
-0.6 + 0.8im

julia> (1 + 2im) + (1 - 2im)
2 + 0im

julia> (-3 + 2im) - (5 - 1im)
-8 + 3im

julia> (-1 + 2im)^2
-3 - 4im

julia> (-1 + 2im)^2.5
2.7296244647840084 - 6.960664459571898im

julia> (-1 + 2im)^(1 + 1im)
-0.27910381075826657 + 0.08708053414102428im

julia> 3(2 - 5im)
6 - 15im

julia> 3(2 - 5im)^2
```

```
-63 - 60im  
  
julia> 3(2 - 5im)^-1.0  
0.20689655172413796 + 0.5172413793103449im
```

The promotion mechanism ensures that combinations of operands of different types just work:

```
julia> 2(1 - 1im)  
2 - 2im  
  
julia> (2 + 3im) - 1  
1 + 3im  
  
julia> (1 + 2im) + 0.5  
1.5 + 2.0im  
  
julia> (2 + 3im) - 0.5im  
2.0 + 2.5im  
  
julia> 0.75(1 + 2im)  
0.75 + 1.5im  
  
julia> (2 + 3im) / 2  
1.0 + 1.5im  
  
julia> (1 - 3im) / (2 + 2im)  
-0.5 - 1.0im  
  
julia> 2im^2  
-2 + 0im  
  
julia> 1 + 3/4im  
1.0 - 0.75im
```

Note that `3/4im == 3/(4*im) == -(3/4*im)`, since a literal coefficient binds more tightly than division.

Standard functions to manipulate complex values are provided:

```
julia> real(1 + 2im)  
1  
  
julia> imag(1 + 2im)  
2  
  
julia> conj(1 + 2im)  
1 - 2im  
  
julia> abs(1 + 2im)  
2.23606797749979  
  
julia> abs2(1 + 2im)  
5  
  
julia> angle(1 + 2im)  
1.1071487177940904
```

As usual, the absolute value (*abs()* (page 418)) of a complex number is its distance from zero. *abs2()* (page 419) gives the square of the absolute value, and is of particular use for complex numbers where it avoids taking a square root. *angle()* (page 420) returns the phase angle in radians (also known as the *argument* or *arg* function). The full gamut of other *Elementary Functions* (page 30) is also defined for complex numbers:

```
julia> sqrt(1im)
0.7071067811865476 + 0.7071067811865475im

julia> sqrt(1 + 2im)
1.272019649514069 + 0.7861513777574233im

julia> cos(1 + 2im)
2.0327230070196656 - 3.0518977991518im

julia> exp(1 + 2im)
-1.1312043837568135 + 2.4717266720048188im

julia> sinh(1 + 2im)
-0.4890562590412937 + 1.4031192506220405im
```

Note that mathematical functions typically return real values when applied to real numbers and complex values when applied to complex numbers. For example, `sqrt()` (page 420) behaves differently when applied to `-1` versus `-1 + 0im` even though `-1 == -1 + 0im`:

```
julia> sqrt(-1)
ERROR: DomainError:
sqrt will only return a complex result if called with a complex argument. Try sqrt(complex(x)).
in sqrt(::Int64) at ./math.jl:278
...

julia> sqrt(-1 + 0im)
0.0 + 1.0im
```

The *literal numeric coefficient notation* (page 22) does not work when constructing complex number from variables. Instead, the multiplication must be explicitly written out:

```
julia> a = 1; b = 2; a + b*im
1 + 2im
```

However, this is *not* recommended; Use the `complex()` (page 434) function instead to construct a complex value directly from its real and imaginary parts.:

```
julia> complex(a,b)
1 + 2im
```

This construction avoids the multiplication and addition operations.

`Inf` (page 436) and `NaN` (page 436) propagate through complex numbers in the real and imaginary parts of a complex number as described in the *Special floating-point values* (page 18) section:

```
julia> 1 + Inf*im
1.0 + Inf*im

julia> 1 + NaN*im
1.0 + NaN*im
```

6.2 Rational Numbers

Julia has a rational number type to represent exact ratios of integers. Rationals are constructed using the `//` (page 408) operator:

```
julia> 2//3
2//3
```

If the numerator and denominator of a rational have common factors, they are reduced to lowest terms such that the denominator is non-negative:

```
julia> 6//9
2//3

julia> -4//8
-1//2

julia> 5//-15
-1//3

julia> -4// -12
1//3
```

This normalized form for a ratio of integers is unique, so equality of rational values can be tested by checking for equality of the numerator and denominator. The standardized numerator and denominator of a rational value can be extracted using the `numerator()` (page 408) and `denominator()` (page 408) functions:

```
julia> numerator(2//3)
2

julia> denominator(2//3)
3
```

Direct comparison of the numerator and denominator is generally not necessary, since the standard arithmetic and comparison operations are defined for rational values:

```
julia> 2//3 == 6//9
true

julia> 2//3 == 9//27
false

julia> 3//7 < 1//2
true

julia> 3//4 > 2//3
true

julia> 2//4 + 1//6
2//3

julia> 5//12 - 1//4
1//6

julia> 5//8 * 3//12
5//32

julia> 6//5 / 10//7
21//25
```

Rationals can be easily converted to floating-point numbers:

```
julia> float(3//4)
0.75
```

Conversion from rational to floating-point respects the following identity for any integral values of a and b , with the exception of the case $a == 0$ and $b == 0$:

```
julia> isequal(float(a//b), a/b)
true
```

Constructing infinite rational values is acceptable:

```
julia> 5//0
1//0

julia> -3//0
-1//0

julia> typeof(ans)
Rational{Int64}
```

Trying to construct a *NaN* (page 436) rational value, however, is not:

```
julia> 0//0
ERROR: ArgumentError: invalid rational: zero(Int64)//zero(Int64)
  in Rational{Int64}(::Int64, ::Int64) at ./rational.jl:8
  in //(::Int64, ::Int64) at ./rational.jl:22
...
```

As usual, the promotion system makes interactions with other numeric types effortless:

```
julia> 3//5 + 1
8//5

julia> 3//5 - 0.5
0.09999999999999998

julia> 2//7 * (1 + 2im)
2//7 + 4//7*im

julia> 2//7 * (1.5 + 2im)
0.42857142857142855 + 0.5714285714285714im

julia> 3//2 / (1 + 2im)
3//10 - 3//5*im

julia> 1//2 + 2im
1//2 + 2//1*im

julia> 1 + 2//3im
1//1 - 2//3*im

julia> 0.5 == 1//2
true

julia> 0.33 == 1//3
false

julia> 0.33 < 1//3
true

julia> 1//3 - 0.33
0.00333333333333332993
```

Strings

Strings are finite sequences of characters. Of course, the real trouble comes when one asks what a character is. The characters that English speakers are familiar with are the letters A, B, C, etc., together with numerals and common punctuation symbols. These characters are standardized together with a mapping to integer values between 0 and 127 by the [ASCII](#) standard. There are, of course, many other characters used in non-English languages, including variants of the ASCII characters with accents and other modifications, related scripts such as Cyrillic and Greek, and scripts completely unrelated to ASCII and English, including Arabic, Chinese, Hebrew, Hindi, Japanese, and Korean. The [Unicode](#) standard tackles the complexities of what exactly a character is, and is generally accepted as the definitive standard addressing this problem. Depending on your needs, you can either ignore these complexities entirely and just pretend that only ASCII characters exist, or you can write code that can handle any of the characters or encodings that one may encounter when handling non-ASCII text. Julia makes dealing with plain ASCII text simple and efficient, and handling Unicode is as simple and efficient as possible. In particular, you can write C-style string code to process ASCII strings, and they will work as expected, both in terms of performance and semantics. If such code encounters non-ASCII text, it will gracefully fail with a clear error message, rather than silently introducing corrupt results. When this happens, modifying the code to handle non-ASCII data is straightforward.

There are a few noteworthy high-level features about Julia's strings:

- The built-in concrete type used for strings (and string literals) in Julia is `String` (page 443). This supports the full range of [Unicode](#) characters via the [UTF-8](#) encoding. (A `transcode()` (page 443) function is provided to convert to/from other Unicode encodings.)
- All string types are subtypes of the abstract type `AbstractString`, and external packages define additional `AbstractString` subtypes (e.g. for other encodings). If you define a function expecting a string argument, you should declare the type as `AbstractString` in order to accept any string type.
- Like C and Java, but unlike most dynamic languages, Julia has a first-class type representing a single character, called `Char`. This is just a special kind of 32-bit bitstype whose numeric value represents a Unicode code point.
- As in Java, strings are immutable: the value of an `AbstractString` object cannot be changed. To construct a different string value, you construct a new string from parts of other strings.
- Conceptually, a string is a *partial function* from indices to characters: for some index values, no character value is returned, and instead an exception is thrown. This allows for efficient indexing into strings by the byte index of an encoded representation rather than by a character index, which cannot be implemented both efficiently and simply for variable-width encodings of Unicode strings.

7.1 Characters

A `Char` value represents a single character: it is just a 32-bit bitstype with a special literal representation and appropriate arithmetic behaviors, whose numeric value is interpreted as a [Unicode code point](#). Here is how `Char` values are input and shown:

```
julia> 'x'
'x'

julia> typeof(ans)
Char
```

You can convert a `Char` to its integer value, i.e. code point, easily:

```
julia> Int('x')
120

julia> typeof(ans)
Int64
```

On 32-bit architectures, `typeof(ans)` (page 363) will be `Int32`. You can convert an integer value back to a `Char` just as easily:

```
julia> Char(120)
'x'
```

Not all integer values are valid Unicode code points, but for performance, the `Char()` conversion does not check that every character value is valid. If you want to check that each converted value is a valid code point, use the `isvalid()` (page 445) function:

```
julia> Char(0x110000)
'\U110000'

julia> isvalid(Char, 0x110000)
false
```

As of this writing, the valid Unicode code points are `U+00` through `U+d7ff` and `U+e000` through `U+10ffff`. These have not all been assigned intelligible meanings yet, nor are they necessarily interpretable by applications, but all of these values are considered to be valid Unicode characters.

You can input any Unicode character in single quotes using `\u` followed by up to four hexadecimal digits or `\U` followed by up to eight hexadecimal digits (the longest valid value only requires six):

```
julia> '\u0'
'\0'

julia> '\u78'
'x'

julia> '\u2200'
'∀'

julia> '\U10ffff'
'\U10ffff'
```

Julia uses your system's locale and language settings to determine which characters can be printed as-is and which must be output using the generic, escaped `\u` or `\U` input forms. In addition to these Unicode escape forms, all of C's traditional escaped input forms can also be used:

```
julia> Int('\0')
0

julia> Int('\t')
9

julia> Int('\n')
```



```

10
julia> Int('\e')
27
julia> Int('\x7f')
127
julia> Int('\177')
127
julia> Int('\xff')
255

```

You can do comparisons and a limited amount of arithmetic with Char values:

```

julia> 'A' < 'a'
true
julia> 'A' <= 'a' <= 'Z'
false
julia> 'A' <= 'X' <= 'Z'
true
julia> 'x' - 'a'
23
julia> 'A' + 1
'B'

```

7.2 String Basics

String literals are delimited by double quotes or triple double quotes:

```

julia> str = "Hello, world.\n"
"Hello, world.\n"
julia> """Contains "quote" characters"""
"Contains \"quote\" characters"

```

If you want to extract a character from a string, you index into it:

```

julia> str[1]
'H'
julia> str[6]
','
julia> str[end]
'\n'

```

All indexing in Julia is 1-based: the first element of any integer-indexed object is found at index 1, and the last element is found at index `n`, when the string has a length of `n`.

In any indexing expression, the keyword `end` can be used as a shorthand for the last index (computed by `endof(str)` (page 383)). You can perform arithmetic and other operations with `end`, just like a normal value:

```
julia> str[end-1]
','

julia> str[end÷2]
','
```

Using an index less than 1 or greater than `end` raises an error:

```
julia> str[0]
ERROR: BoundsError: attempt to access 14-element Array{UInt8,1} at index [0]

julia> str[end+1]
ERROR: BoundsError: attempt to access 14-element Array{UInt8,1} at index [15]
```

You can also extract a substring using range indexing:

```
julia> str[4:9]
"lo, wo"
```

Notice that the expressions `str[k]` and `str[k:k]` do not give the same result:

```
julia> str[6]
','

julia> str[6:6]
","
```

The former is a single character value of type `Char`, while the latter is a string value that happens to contain only a single character. In Julia these are very different things.

7.3 Unicode and UTF-8

Julia fully supports Unicode characters and strings. As *discussed above* (page 41), in character literals, Unicode code points can be represented using Unicode `\u` and `\U` escape sequences, as well as all the standard C escape sequences. These can likewise be used to write string literals:

```
julia> s = "\u2200 x \u2203 y"
"∀ x ∃ y"
```

Whether these Unicode characters are displayed as escapes or shown as special characters depends on your terminal's locale settings and its support for Unicode. String literals are encoded using the UTF-8 encoding. UTF-8 is a variable-width encoding, meaning that not all characters are encoded in the same number of bytes. In UTF-8, ASCII characters — i.e. those with code points less than 0x80 (128) — are encoded as they are in ASCII, using a single byte, while code points 0x80 and above are encoded using multiple bytes — up to four per character. This means that not every byte index into a UTF-8 string is necessarily a valid index for a character. If you index into a string at such an invalid byte index, an error is thrown:

```
julia> s[1]
'∀'

julia> s[2]
ERROR: UnicodeError: invalid character index
  in slow_utf8_next(::Array{UInt8,1}, ::UInt8, ::Int64) at ./strings/string.jl:67
  in next at ./strings/string.jl:92 [inlined]
  in getindex(::String, ::Int64) at ./strings/basic.jl:71
  ...
```

```
julia> s[3]
ERROR: UnicodeError: invalid character index
  in slow_utf8_next(::Array{UInt8,1}, ::UInt8, ::Int64) at ./strings/string.jl:67
  in next at ./strings/string.jl:92 [inlined]
  in getindex(::String, ::Int64) at ./strings/basic.jl:71
  ...

julia> s[4]
' '
```

In this case, the character √ is a three-byte character, so the indices 2 and 3 are invalid and the next character's index is 4; this next valid index can be computed by `nextind(s, 1)` (page 449), and the next index after that by `nextind(s, 4)` and so on.

Because of variable-length encodings, the number of characters in a string (given by `length(s)` (page 443)) is not always the same as the last index. If you iterate through the indices 1 through `endof(s)` (page 383) and index into `s`, the sequence of characters returned when errors aren't thrown is the sequence of characters comprising the string `s`. Thus we have the identity that `length(s) <= endof(s)`, since each character in a string must have its own index. The following is an inefficient and verbose way to iterate through the characters of `s`:

```
julia> for i = 1:endof(s)
    try
        println(s[i])
    catch
        # ignore the index error
    end
end

√
x
∃
y
```

The blank lines actually have spaces on them. Fortunately, the above awkward idiom is unnecessary for iterating through the characters in a string, since you can just use the string as an iterable object, no exception handling required:

```
julia> for c in s
    println(c)
end

√
x
∃
y
```

Julia uses the UTF-8 encoding by default, and support for new encodings can be added by packages. For example, the `LegacyStrings.jl` package implements `UTF16String` and `UTF32String` types. Additional discussion of other encodings and how to implement support for them is beyond the scope of this document for the time being. For further discussion of UTF-8 encoding issues, see the section below on *byte array literals* (page ??). The `transcode()` (page 443) function is provided to convert data between the various UTF-xx encodings, primarily for working with external data and libraries.

7.4 Interpolation

One of the most common and useful string operations is concatenation:

```
julia> greet = "Hello"
"Hello"

julia> whom = "world"
"world"

julia> string(greet, " ", whom, ".\n")
"Hello, world.\n"
```

Constructing strings like this can become a bit cumbersome, however. To reduce the need for these verbose calls to `string()` (page 443), Julia allows interpolation into string literals using `$`, as in Perl:

```
julia> "$greet, $whom.\n"
"Hello, world.\n"
```

This is more readable and convenient and equivalent to the above string concatenation — the system rewrites this apparent single string literal into a concatenation of string literals with variables.

The shortest complete expression after the `$` is taken as the expression whose value is to be interpolated into the string. Thus, you can interpolate any expression into a string using parentheses:

```
julia> "1 + 2 = $(1 + 2)"
"1 + 2 = 3"
```

Both concatenation and string interpolation call `string()` (page 443) to convert objects into string form. Most non-`AbstractString` objects are converted to strings closely corresponding to how they are entered as literal expressions:

```
julia> v = [1,2,3]
3-element Array{Int64,1}:
 1
 2
 3

julia> "v: $v"
"v: [1,2,3]"
```

`string()` (page 443) is the identity for `AbstractString` and `Char` values, so these are interpolated into strings as themselves, unquoted and unescaped:

```
julia> c = 'x'
'x'

julia> "hi, $c"
"hi, x"
```

To include a literal `$` in a string literal, escape it with a backslash:

```
julia> print("I have \$100 in my account.\n")
I have $100 in my account.
```

7.5 Triple-Quoted String Literals

When strings are created using triple-quotes (`"""..."""`) they have some special behavior that can be useful for creating longer blocks of text. First, if the opening `"""` is followed by a newline, the newline is stripped from the resulting string.

```
"""hello"""
```

is equivalent to

```
"""
hello"""
```

but

```
"""
hello"""
```

will contain a literal newline at the beginning. Trailing whitespace is left unaltered. They can contain `"` symbols without escaping. Triple-quoted strings are also dedented to the level of the least-indented line. This is useful for defining strings within code that is indented. For example:

```
julia> str = """
        Hello,
        world.
        """
" Hello,\n world.\n"
```

In this case the final (empty) line before the closing `"""` sets the indentation level.

Note that line breaks in literal strings, whether single- or triple-quoted, result in a newline (LF) character `\n` in the string, even if your editor uses a carriage return `\r` (CR) or CRLF combination to end lines. To include a CR in a string, use an explicit escape `\r`; for example, you can enter the literal string `"a CRLF line ending\r\n"`.

7.6 Common Operations

You can lexicographically compare strings using the standard comparison operators:

```
julia> "abracadabra" < "xylophone"
true

julia> "abracadabra" == "xylophone"
false

julia> "Hello, world." != "Goodbye, world."
true

julia> "1 + 2 = 3" == "1 + 2 = $(1 + 2)"
true
```

You can search for the index of a particular character using the `search()` (page 446) function:

```
julia> search("xylophone", 'x')
1

julia> search("xylophone", 'p')
5
```

```
julia> search("xylophone", 'z')
0
```

You can start the search for a character at a given offset by providing a third argument:

```
julia> search("xylophone", 'o')
4

julia> search("xylophone", 'o', 5)
7

julia> search("xylophone", 'o', 8)
0
```

You can use the `contains()` (page 446) function to check if a substring is contained in a string:

```
julia> contains("Hello, world.", "world")
true

julia> contains("Xylophon", "o")
true

julia> contains("Xylophon", "a")
false

julia> contains("Xylophon", 'o')
ERROR: MethodError: no method matching contains(::String, ::Char)
Closest candidates are:
  contains(!Matched::Function, ::Any, !Matched::Any) at reduce.jl:640
  contains(::AbstractString, !Matched::AbstractString) at strings/search.jl:366
...
```

The last error is because `'o'` is a character literal, and `contains()` (page 446) is a generic function that looks for subsequences. To look for an element in a sequence, you must use `in()` (page 383) instead.

Two other handy string functions are `repeat()` (page 530) and `join()` (page 448):

```
julia> repeat(".:Z:.", 10)
".:Z:.:Z:.:Z:.:Z:.:Z:.:Z:.:Z:.:Z:.:Z:."

julia> join(["apples", "bananas", "pineapples"], ", ", " and ")
"apples, bananas and pineapples"
```

Some other useful functions include:

- `endof(str)` (page 383) gives the maximal (byte) index that can be used to index into `str`.
- `length(str)` (page 443) the number of characters in `str`.
- `i = start(str)` (page 381) gives the first valid index at which a character can be found in `str` (typically 1).
- `c, j = next(str, i)` (page 381) returns next character at or after the index `i` and the next valid character index following that. With `start()` (page 381) and `endof()` (page 383), can be used to iterate through the characters in `str`.
- `ind2chr(str, i)` (page 449) gives the number of characters in `str` up to and including any at index `i`.
- `chr2ind(str, j)` (page 449) gives the index at which the `j`th character in `str` occurs.

7.7 Non-Standard String Literals

There are situations when you want to construct a string or use string semantics, but the behavior of the standard string construct is not quite what is needed. For these kinds of situations, Julia provides *non-standard string literals* (page 171). A non-standard string literal looks like a regular double-quoted string literal, but is immediately prefixed by an identifier, and doesn't behave quite like a normal string literal. Regular expressions, byte array literals and version number literals, as described below, are some examples of non-standard string literals. Other examples are given in the *metaprogramming* (page 171) section.

7.8 Regular Expressions

Julia has Perl-compatible regular expressions (regexes), as provided by the [PCRE](#) library. Regular expressions are related to strings in two ways: the obvious connection is that regular expressions are used to find regular patterns in strings; the other connection is that regular expressions are themselves input as strings, which are parsed into a state machine that can be used to efficiently search for patterns in strings. In Julia, regular expressions are input using non-standard string literals prefixed with various identifiers beginning with `r`. The most basic regular expression literal without any options turned on just uses `r"..."`:

```
julia> r"^s*(?:#|$)"
r"^s*(?:#|$)"

julia> typeof(ans)
Regex
```

To check if a regex matches a string, use `ismatch()` (page 445):

```
julia> ismatch(r"^s*(?:#|$)", "not a comment")
false

julia> ismatch(r"^s*(?:#|$)", "# a comment")
true
```

As one can see here, `ismatch()` (page 445) simply returns true or false, indicating whether the given regex matches the string or not. Commonly, however, one wants to know not just whether a string matched, but also *how* it matched. To capture this information about a match, use the `match()` (page 445) function instead:

```
julia> match(r"^s*(?:#|$)", "not a comment")

julia> match(r"^s*(?:#|$)", "# a comment")
RegexMatch("#")
```

If the regular expression does not match the given string, `match()` (page 445) returns `nothing` — a special value that does not print anything at the interactive prompt. Other than not printing, it is a completely normal value and you can test for it programmatically:

```
m = match(r"^s*(?:#|$)", line)
if m === nothing
    println("not a comment")
else
    println("blank or comment")
end
```

If a regular expression does match, the value returned by `match()` (page 445) is a `RegexMatch` object. These objects record how the expression matches, including the substring that the pattern matches and any captured substrings,

if there are any. This example only captures the portion of the substring that matches, but perhaps we want to capture any non-blank text after the comment character. We could do the following:

```
julia> m = match(r"^\\s*(?:#\\s*(.*)\\s*$|\\s*$)", "# a comment ")
RegexMatch("# a comment ", 1="a comment")
```

When calling `match()` (page 445), you have the option to specify an index at which to start the search. For example:

```
julia> m = match(r"[0-9]", "aaaa1aaaa2aaaa3", 1)
RegexMatch("1")

julia> m = match(r"[0-9]", "aaaa1aaaa2aaaa3", 6)
RegexMatch("2")

julia> m = match(r"[0-9]", "aaaa1aaaa2aaaa3", 11)
RegexMatch("3")
```

You can extract the following info from a `RegexMatch` object:

- the entire substring matched: `m.match`
- the captured substrings as an array of strings: `m.captures`
- the offset at which the whole match begins: `m.offset`
- the offsets of the captured substrings as a vector: `m.offsets`

For when a capture doesn't match, instead of a substring, `m.captures` contains `nothing` in that position, and `m.offsets` has a zero offset (recall that indices in Julia are 1-based, so a zero offset into a string is invalid). Here is a pair of somewhat contrived examples:

```
julia> m = match(r"(a|b)(c)?(d)", "acd")
RegexMatch("acd", 1="a", 2="c", 3="d")

julia> m.match
"acd"

julia> m.captures
3-element Array{Union{SubString{String},Void},1}:
 "a"
 "c"
 "d"

julia> m.offset
1

julia> m.offsets
3-element Array{Int64,1}:
 1
 2
 3

julia> m = match(r"(a|b)(c)?(d)", "ad")
RegexMatch("ad", 1="a", 2=nothing, 3="d")

julia> m.match
"ad"

julia> m.captures
3-element Array{Union{SubString{String},Void},1}:
 "a"
```



```
nothing
"d"

julia> m.offset
1

julia> m.offsets
3-element Array{Int64,1}:
 1
 0
 2
```

It is convenient to have captures returned as an array so that one can use destructuring syntax to bind them to local variables:

```
julia> first, second, third = m.captures; first
"a"
```

Captures can also be accessed by indexing the `RegexMatch` object with the number or name of the capture group:

```
julia> m=match(r"(?<hour>\d+):(?<minute>\d+)", "12:45")
RegexMatch("12:45", hour="12", minute="45")
julia> m[:minute]
"45"
julia> m[2]
"45"
```

Captures can be referenced in a substitution string when using `replace()` (page 446) by using `\n` to refer to the *n*th capture group and prefixing the substitution string with `s`. Capture group 0 refers to the entire match object. Named capture groups can be referenced in the substitution with `g<groupname>`. For example:

```
julia> replace("first second", r"(\w+) (?<agroup>\w+)", s"\g<agroup> \1")
"second first"
```

Numbered capture groups can also be referenced as `\g<n>` for disambiguation, as in:

```
julia> replace("a", r".", s"\g<0>1")
"a1"
```

You can modify the behavior of regular expressions by some combination of the flags `i`, `m`, `s`, and `x` after the closing double quote mark. These flags have the same meaning as they do in Perl, as explained in this excerpt from the [perlre manpage](#):

```
i  Do case-insensitive pattern matching.

    If locale matching rules are in effect, the case map is taken
    from the current locale for code points less than 255, and
    from Unicode rules for larger code points. However, matches
    that would cross the Unicode rules/non-Unicode rules boundary
    (ords 255/256) will not succeed.

m  Treat string as multiple lines. That is, change "^" and "$"
    from matching the start or end of the string to matching the
    start or end of any line anywhere within the string.

s  Treat string as single line. That is, change "." to match any
    character whatsoever, even a newline, which normally it would
    not match.

    Used together, as r""ms, they let the "." match any character
```

whatsoever, **while** still allowing `"^"` and `"$"` to match, respectively, just after and just before newlines within the string.

- x Tells the regular expression parser to ignore most whitespace that **is** neither backslashed nor within a character class. You can use this to **break** up your regular expression into (slightly) more readable parts. The `'#'` character **is** also treated as a metacharacter introducing a comment, just as **in** ordinary code.

For example, the following regex has all three flags turned on:

```
julia> r"a+.*b+.*?d$"ism
r"a+.*b+.*?d$"ims

julia> match(r"a+.*b+.*?d$"ism, "Goodbye,\nOh, angry,\nBad world\n")
RegexMatch("angry,\nBad world")
```

Triple-quoted regex strings, of the form `r"""..."""`, are also supported (and may be convenient for regular expressions containing quotation marks or newlines).

7.9 Byte Array Literals

Another useful non-standard string literal is the byte-array string literal: `b"..."`. This form lets you use string notation to express literal byte arrays — i.e. arrays of `UInt8` values. The rules for byte array literals are the following:

- ASCII characters and ASCII escapes produce a single byte.
- `\x` and octal escape sequences produce the *byte* corresponding to the escape value.
- Unicode escape sequences produce a sequence of bytes encoding that code point in UTF-8.

There is some overlap between these rules since the behavior of `\x` and octal escapes less than 0x80 (128) are covered by both of the first two rules, but here these rules agree. Together, these rules allow one to easily use ASCII characters, arbitrary byte values, and UTF-8 sequences to produce arrays of bytes. Here is an example using all three:

```
julia> b"DATA\xff\u2200"
8-element Array{UInt8,1}:
 0x44
 0x41
 0x54
 0x41
 0xff
 0xe2
 0x88
 0x80
```

The ASCII string “DATA” corresponds to the bytes 68, 65, 84, 65. `\xff` produces the single byte 255. The Unicode escape `\u2200` is encoded in UTF-8 as the three bytes 226, 136, 128. Note that the resulting byte array does not correspond to a valid UTF-8 string — if you try to use this as a regular string literal, you will get a syntax error:

```
julia> "DATA\xff\u2200"
ERROR: syntax: invalid UTF-8 sequence
...
```

Also observe the significant distinction between `\xff` and `\uff`: the former escape sequence encodes the *byte* 255, whereas the latter escape sequence represents the *code point* 255, which is encoded as two bytes in UTF-8:

```
julia> b"\xff"
1-element Array{UInt8,1}:
 0xff

julia> b"\uff"
2-element Array{UInt8,1}:
 0xc3
 0xbf
```

In character literals, this distinction is glossed over and `\xff` is allowed to represent the code point 255, because characters *always* represent code points. In strings, however, `\x` escapes always represent bytes, not code points, whereas `\u` and `\U` escapes always represent code points, which are encoded in one or more bytes. For code points less than `\u80`, it happens that the UTF-8 encoding of each code point is just the single byte produced by the corresponding `\x` escape, so the distinction can safely be ignored. For the escapes `\x80` through `\xff` as compared to `\u80` through `\uff`, however, there is a major difference: the former escapes all encode single bytes, which — unless followed by very specific continuation bytes — do not form valid UTF-8 data, whereas the latter escapes all represent Unicode code points with two-byte encodings.

If this is all extremely confusing, try reading “[The Absolute Minimum Every Software Developer Absolutely, Positively Must Know About Unicode and Character Sets](#)”. It’s an excellent introduction to Unicode and UTF-8, and may help alleviate some confusion regarding the matter.

7.10 Version Number Literals

Version numbers can easily be expressed with non-standard string literals of the form `v"..."`. Version number literals create `VersionNumber` objects which follow the specifications of [semantic versioning](#), and therefore are composed of major, minor and patch numeric values, followed by pre-release and build alpha-numeric annotations. For example, `v"0.2.1-rc1+win64"` is broken into major version 0, minor version 2, patch version 1, pre-release `rc1` and build `win64`. When entering a version literal, everything except the major version number is optional, therefore e.g. `v"0.2"` is equivalent to `v"0.2.0"` (with empty pre-release/build annotations), `v"2"` is equivalent to `v"2.0.0"`, and so on.

`VersionNumber` objects are mostly useful to easily and correctly compare two (or more) versions. For example, the constant `VERSION` holds Julia version number as a `VersionNumber` object, and therefore one can define some version-specific behavior using simple statements as:

```
if v"0.2" <= VERSION < v"0.3-"
    # do something specific to 0.2 release series
end
```

Note that in the above example the non-standard version number `v"0.3-"` is used, with a trailing `-`: this notation is a Julia extension of the standard, and it’s used to indicate a version which is lower than any `0.3` release, including all of its pre-releases. So in the above example the code would only run with stable `0.2` versions, and exclude such versions as `v"0.3.0-rc1"`. In order to also allow for unstable (i.e. pre-release) `0.2` versions, the lower bound check should be modified like this: `v"0.2-" <= VERSION`.

Another non-standard version specification extension allows one to use a trailing `+` to express an upper limit on build versions, e.g. `VERSION > v"0.2-rc1+"` can be used to mean any version above `0.2-rc1` and any of its builds: it will return `false` for version `v"0.2-rc1+win64"` and `true` for `v"0.2-rc2"`.

It is good practice to use such special versions in comparisons (particularly, the trailing `-` should always be used on upper bounds unless there’s a good reason not to), but they must not be used as the actual version number of anything, as they are invalid in the semantic versioning scheme.

Besides being used for the `VERSION` (page 553) constant, `VersionNumber` objects are widely used in the `Pkg` (page 581) module, to specify packages versions and their dependencies.

Functions

In Julia, a function is an object that maps a tuple of argument values to a return value. Julia functions are not pure mathematical functions, in the sense that functions can alter and be affected by the global state of the program. The basic syntax for defining functions in Julia is:

```
function f(x,y)
    x + y
end
```

There is a second, more terse syntax for defining a function in Julia. The traditional function declaration syntax demonstrated above is equivalent to the following compact “assignment form”:

```
f(x,y) = x + y
```

In the assignment form, the body of the function must be a single expression, although it can be a compound expression (see *Compound Expressions* (page 65)). Short, simple function definitions are common in Julia. The short function syntax is accordingly quite idiomatic, considerably reducing both typing and visual noise.

A function is called using the traditional parenthesis syntax:

```
julia> f(2,3)
5
```

Without parentheses, the expression `f` refers to the function object, and can be passed around like any value:

```
julia> g = f;

julia> g(2,3)
5
```

As with variables, Unicode can also be used for function names:

```
julia> ∑(x,y) = x + y
∑ (generic function with 1 method)
```

8.1 Argument Passing Behavior

Julia function arguments follow a convention sometimes called “pass-by-sharing”, which means that values are not copied when they are passed to functions. Function arguments themselves act as new variable *bindings* (new locations that can refer to values), but the values they refer to are identical to the passed values. Modifications to mutable values (such as Arrays) made within a function will be visible to the caller. This is the same behavior found in Scheme, most Lisps, Python, Ruby and Perl, among other dynamic languages.

8.2 The `return` Keyword

The value returned by a function is the value of the last expression evaluated, which, by default, is the last expression in the body of the function definition. In the example function, `f`, from the previous section this is the value of the expression `x + y`. As in C and most other imperative or functional languages, the `return` keyword causes a function to return immediately, providing an expression whose value is returned:

```
function g(x,y)
    return x * y
    x + y
end
```

Since function definitions can be entered into interactive sessions, it is easy to compare these definitions:

```
f(x,y) = x + y

function g(x,y)
    return x * y
    x + y
end

julia> f(2,3)
5

julia> g(2,3)
6
```

Of course, in a purely linear function body like `g`, the usage of `return` is pointless since the expression `x + y` is never evaluated and we could simply make `x * y` the last expression in the function and omit the `return`. In conjunction with other control flow, however, `return` is of real use. Here, for example, is a function that computes the hypotenuse length of a right triangle with sides of length `x` and `y`, avoiding overflow:

```
function hypot(x,y)
    x = abs(x)
    y = abs(y)
    if x > y
        r = y/x
        return x*sqrt(1+r*r)
    end
    if y == 0
        return zero(x)
    end
    r = x/y
    return y*sqrt(1+r*r)
end
```

There are three possible points of return from this function, returning the values of three different expressions, depending on the values of `x` and `y`. The `return` on the last line could be omitted since it is the last expression.

8.3 Operators Are Functions

In Julia, most operators are just functions with support for special syntax. (The exceptions are operators with special evaluation semantics like `&&` and `||`. These operators cannot be functions since *short-circuit evaluation* (page 69) requires that their operands are not evaluated before evaluation of the operator.) Accordingly, you can also apply them using parenthesized argument lists, just as you would any other function:

```
julia> 1 + 2 + 3
6

julia> +(1,2,3)
6
```

The infix form is exactly equivalent to the function application form — in fact the former is parsed to produce the function call internally. This also means that you can assign and pass around operators such as `+` (page 405) and `*` (page 443) just like you would with other function values:

```
julia> f = +;

julia> f(1,2,3)
6
```

Under the name `f`, the function does not support infix notation, however.

8.4 Operators With Special Names

A few special expressions correspond to calls to functions with non-obvious names. These are:

Expression	Calls
<code>[A B C ...]</code>	<code>hcat()</code> (page 462)
<code>[A, B, C, ...]</code>	<code>vcate()</code> (page 462)
<code>[A B; C D; ...]</code>	<code>hvcate()</code> (page 463)
<code>A'</code>	<code>ctranspose()</code> (page 534)
<code>A.'</code>	<code>transpose()</code> (page 534)
<code>1:n</code>	<code>colon()</code> (page 409)
<code>A[i]</code>	<code>getindex()</code> (page 454)
<code>A[i]=x</code>	<code>setindex!()</code> (page 461)

These functions are included in the `Base.Operators` module even though they do not have operator-like names.

8.5 Anonymous Functions

Functions in Julia are *first-class objects*: they can be assigned to variables, and called using the standard function call syntax from the variable they have been assigned to. They can be used as arguments, and they can be returned as values. They can also be created anonymously, without being given a name, using either of these syntaxes:

```
julia> x -> x^2 + 2x - 1
(::#1) (generic function with 1 method)

julia> function (x)
    x^2 + 2x - 1
end
(::#3) (generic function with 1 method)
```

This creates a function taking one argument x and returning the value of the polynomial $x^2 + 2x - 1$ at that value. Notice that the result is a generic function, but with a compiler-generated name based on consecutive numbering.

The primary use for anonymous functions is passing them to functions which take other functions as arguments. A classic example is `map()` (page 390), which applies a function to each value of an array and returns a new array containing the resulting values:

```
julia> map(round, [1.2, 3.5, 1.7])
3-element Array{Float64,1}:
 1.0
 4.0
 2.0
```

This is fine if a named function effecting the transform one wants already exists to pass as the first argument to `map()` (page 390). Often, however, a ready-to-use, named function does not exist. In these situations, the anonymous function construct allows easy creation of a single-use function object without needing a name:

```
julia> map(x -> x^2 + 2x - 1, [1, 3, -1])
3-element Array{Int64,1}:
 2
14
-2
```

An anonymous function accepting multiple arguments can be written using the syntax $(x, y, z) \rightarrow 2x + y - z$. A zero-argument anonymous function is written as $() \rightarrow 3$. The idea of a function with no arguments may seem strange, but is useful for “delaying” a computation. In this usage, a block of code is wrapped in a zero-argument function, which is later invoked by calling it as `f()`.

8.6 Multiple Return Values

In Julia, one returns a tuple of values to simulate returning multiple values. However, tuples can be created and destructured without needing parentheses, thereby providing an illusion that multiple values are being returned, rather than a single tuple value. For example, the following function returns a pair of values:

```
julia> function foo(a,b)
    a+b, a*b
end;
```

If you call it in an interactive session without assigning the return value anywhere, you will see the tuple returned:

```
julia> foo(2,3)
(5, 6)
```

A typical usage of such a pair of return values, however, extracts each value into a variable. Julia supports simple tuple “destructuring” that facilitates this:

```
julia> x, y = foo(2,3);

julia> x
5

julia> y
6
```

You can also return multiple values via an explicit usage of the `return` keyword:

```
function foo(a,b)
    return a+b, a*b
end
```

This has the exact same effect as the previous definition of `foo`.

8.7 Varargs Functions

It is often convenient to be able to write functions taking an arbitrary number of arguments. Such functions are traditionally known as “varargs” functions, which is short for “variable number of arguments”. You can define a varargs function by following the last argument with an ellipsis:

```
julia> bar(a,b,x...) = (a,b,x)
bar (generic function with 1 method)
```

The variables `a` and `b` are bound to the first two argument values as usual, and the variable `x` is bound to an iterable collection of the zero or more values passed to `bar` after its first two arguments:

```
julia> bar(1,2)
(1,2,())

julia> bar(1,2,3)
(1,2,(3,))

julia> bar(1,2,3,4)
(1,2,(3,4))

julia> bar(1,2,3,4,5,6)
(1,2,(3,4,5,6))
```

In all these cases, `x` is bound to a tuple of the trailing values passed to `bar`.

It is possible to constrain the number of values passed as a variable argument; this will be discussed later in [Parametrically-constrained Varargs methods](#) (page 115).

On the flip side, it is often handy to “splice” the values contained in an iterable collection into a function call as individual arguments. To do this, one also uses `...` but in the function call instead:

```
julia> x = (3,4)
(3,4)

julia> bar(1,2,x...)
(1,2,(3,4))
```

In this case a tuple of values is spliced into a varargs call precisely where the variable number of arguments go. This need not be the case, however:

```
julia> x = (2,3,4)
(2,3,4)

julia> bar(1,x...)
(1,2,(3,4))

julia> x = (1,2,3,4)
(1,2,3,4)

julia> bar(x...)
(1,2,(3,4))
```

Furthermore, the iterable object spliced into a function call need not be a tuple:

```
julia> x = [3,4]
2-element Array{Int64,1}:
 3
 4
```

```
julia> bar(1,2,x...)
(1,2,(3,4))

julia> x = [1,2,3,4]
4-element Array{Int64,1}:
 1
 2
 3
 4

julia> bar(x...)
(1,2,(3,4))
```

Also, the function that arguments are spliced into need not be a varargs function (although it often is):

```
julia> baz(a,b) = a + b;

julia> args = [1,2]
2-element Array{Int64,1}:
 1
 2

julia> baz(args...)
3

julia> args = [1,2,3]
3-element Array{Int64,1}:
 1
 2
 3

julia> baz(args...)
ERROR: MethodError: no method matching baz(::Int64, ::Int64, ::Int64)
Closest candidates are:
  baz(::Any, ::Any) at none:1
...
```

As you can see, if the wrong number of elements are in the spliced container, then the function call will fail, just as it would if too many arguments were given explicitly.

8.8 Optional Arguments

In many cases, function arguments have sensible default values and therefore might not need to be passed explicitly in every call. For example, the library function `parse(type,num,base)` (page 369) interprets a string as a number in some base. The `base` argument defaults to 10. This behavior can be expressed concisely as:

```
function parse(type, num, base=10)
    ###
end
```

With this definition, the function can be called with either two or three arguments, and 10 is automatically passed when a third argument is not specified:

```
julia> parse(Int, "12", 10)
12
```

```
julia> parse{Int, "12", 3}
5

julia> parse{Int, "12"}
12
```

Optional arguments are actually just a convenient syntax for writing multiple method definitions with different numbers of arguments (see *Note on Optional and keyword Arguments* (page 115)).

8.9 Keyword Arguments

Some functions need a large number of arguments, or have a large number of behaviors. Remembering how to call such functions can be difficult. Keyword arguments can make these complex interfaces easier to use and extend by allowing arguments to be identified by name instead of only by position.

For example, consider a function `plot` that plots a line. This function might have many options, for controlling line style, width, color, and so on. If it accepts keyword arguments, a possible call might look like `plot(x, y, width=2)`, where we have chosen to specify only line width. Notice that this serves two purposes. The call is easier to read, since we can label an argument with its meaning. It also becomes possible to pass any subset of a large number of arguments, in any order.

Functions with keyword arguments are defined using a semicolon in the signature:

```
function plot(x, y; style="solid", width=1, color="black")
    ###
end
```

When the function is called, the semicolon is optional: one can either call `plot(x, y, width=2)` or `plot(x, y; width=2)`, but the former style is more common. An explicit semicolon is required only for passing varargs or computed keywords as described below.

Keyword argument default values are evaluated only when necessary (when a corresponding keyword argument is not passed), and in left-to-right order. Therefore default expressions may refer to prior keyword arguments.

The types of keyword arguments can be made explicit as follows:

```
function f(;x::Int64=1)
    ###
end
```

Extra keyword arguments can be collected using `...`, as in varargs functions:

```
function f(x; y=0, kwargs...)
    ###
end
```

Inside `f`, `kwargs` will be a collection of `(key, value)` tuples, where each `key` is a symbol. Such collections can be passed as keyword arguments using a semicolon in a call, e.g. `f(x, z=1; kwargs...)`. Dictionaries can also be used for this purpose.

One can also pass `(key, value)` tuples, or any iterable expression (such as a `=>` pair) that can be assigned to such a tuple, explicitly after a semicolon. For example, `plot(x, y; (:width, 2))` and `plot(x, y; :width => 2)` are equivalent to `plot(x, y, width=2)`. This is useful in situations where the keyword name is computed at runtime.

The nature of keyword arguments makes it possible to specify the same argument more than once. For example, in the call `plot(x, y; options..., width=2)` it is possible that the `options` structure also contains a value

for `width`. In such a case the rightmost occurrence takes precedence; in this example, `width` is certain to have the value 2.

8.10 Evaluation Scope of Default Values

Optional and keyword arguments differ slightly in how their default values are evaluated. When optional argument default expressions are evaluated, only *previous* arguments are in scope. In contrast, *all* the arguments are in scope when keyword arguments default expressions are evaluated. For example, given this definition:

```
function f(x, a=b, b=1)
    ###
end
```

the `b` in `a=b` refers to a `b` in an outer scope, not the subsequent argument `b`. However, if `a` and `b` were keyword arguments instead, then both would be created in the same scope and the `b` in `a=b` would refer to the subsequent argument `b` (shadowing any `b` in an outer scope), which would result in an undefined variable error (since the default expressions are evaluated left-to-right, and `b` has not been assigned yet).

8.11 Do-Block Syntax for Function Arguments

Passing functions as arguments to other functions is a powerful technique, but the syntax for it is not always convenient. Such calls are especially awkward to write when the function argument requires multiple lines. As an example, consider calling `map()` (page 390) on a function with several cases:

```
map(x->begin
    if x < 0 && iseven(x)
        return 0
    elseif x == 0
        return 1
    else
        return x
    end
end,
[A, B, C])
```

Julia provides a reserved word `do` for rewriting this code more clearly:

```
map([A, B, C]) do x
    if x < 0 && iseven(x)
        return 0
    elseif x == 0
        return 1
    else
        return x
    end
end
```

The `do x` syntax creates an anonymous function with argument `x` and passes it as the first argument to `map()` (page 390). Similarly, `do a, b` would create a two-argument anonymous function, and a plain `do` would declare that what follows is an anonymous function of the form `() -> ...`.

How these arguments are initialized depends on the “outer” function; here, `map()` (page 390) will sequentially set `x` to `A`, `B`, `C`, calling the anonymous function on each, just as would happen in the syntax `map(func, [A, B, C])`.

This syntax makes it easier to use functions to effectively extend the language, since calls look like normal code blocks. There are many possible uses quite different from `map()` (page 390), such as managing system state. For example, there is a version of `open()` (page 561) that runs code ensuring that the opened file is eventually closed:

```
open("outfile", "w") do io
    write(io, data)
end
```

This is accomplished by the following definition:

```
function open(f::Function, args...)
    io = open(args...)
    try
        f(io)
    finally
        close(io)
    end
end
```

Here, `open()` (page 561) first opens the file for writing and then passes the resulting output stream to the anonymous function you defined in the `do ... end` block. After your function exits, `open()` (page 561) will make sure that the stream is properly closed, regardless of whether your function exited normally or threw an exception. (The `try/finally` construct will be described in [Control Flow](#) (page 65).)

With the `do` block syntax, it helps to check the documentation or implementation to know how the arguments of the user function are initialized.

8.12 Dot Syntax for Vectorizing Functions

In technical-computing languages, it is common to have “vectorized” versions of functions, which simply apply a given function $f(x)$ to each element of an array A to yield a new array via $f(A)$. This kind of syntax is convenient for data processing, but in other languages vectorization is also often required for performance: if loops are slow, the “vectorized” version of a function can call fast library code written in a low-level language. In Julia, vectorized functions are *not* required for performance, and indeed it is often beneficial to write your own loops (see [Performance Tips](#) (page 305)), but they can still be convenient. Therefore, *any* Julia function f can be applied elementwise to any array (or other collection) with the syntax $f.(A)$.

Of course, you can omit the dot if you write a specialized “vector” method of f , e.g. via $f(A::AbstractArray) = \text{map}(f, A)$, and this is just as efficient as $f.(A)$. But that approach requires you to decide in advance which functions you want to vectorize.

More generally, $f.(args...)$ is actually equivalent to `broadcast(f, args...)`, which allows you to operate on multiple arrays (even of different shapes), or a mix of arrays and scalars (see [Broadcasting](#) (page 184)). For example, if you have $f(x, y) = 3x + 4y$, then $f.(pi, A)$ will return a new array consisting of $f(pi, a)$ for each a in A , and $f.(vector1, vector2)$ will return a new vector consisting of $f(vector1[i], vector2[i])$ for each index i (throwing an exception if the vectors have different length).

Moreover, *nested* $f.(args...)$ calls are *fused* into a single broadcast loop. For example, $\sin.(\cos.(X))$ is equivalent to `broadcast(x -> sin(cos(x)), X)`, similar to `[sin(cos(x)) for x in X]`: there is only a single loop over X , and a single array is allocated for the result. [In contrast, $\sin(\cos(X))$ in a typical “vectorized” language would first allocate one temporary array for $\text{tmp} = \cos(X)$, and then compute $\sin(\text{tmp})$ in a separate loop, allocating a second array.] This loop fusion is not a compiler optimization that may or may not occur, it is a *syntactic guarantee* whenever nested $f.(args...)$ calls are encountered. Technically, the fusion stops as soon as a “non-dot” function is encountered; for example, in $\sin.(\text{sort}(\cos.(X)))$ the \sin and \cos loops cannot be merged because of the intervening `sort` function.

Finally, the maximum efficiency is typically achieved when the output array of a vectorized operation is *pre-allocated*, so that repeated calls do not allocate new arrays over and over again for the results (*Pre-allocating outputs* (page 317):). A convenient syntax for this is `X .= ...`, which is equivalent to `broadcast!(identity, X, ...)` except that, as above, the `broadcast!` loop is fused with any nested “dot” calls. For example, `X .= sin.(Y)` is equivalent to `broadcast!(sin, X, Y)`, overwriting `X` with `sin.(Y)` in-place. If the left-hand side is an array-indexing expression, e.g. `X[2:end] .= sin.(Y)`, then it translates to `broadcast!` on a view, e.g. `broadcast!(sin, view(X, 2:endof(X)), Y)`, so that the left-hand side is updated in-place.

(In future versions of Julia, operators like `.*` will also be handled with the same mechanism: they will be equivalent to `broadcast` calls and will be fused with other nested “dot” calls. `X .+= Y` is equivalent to `X .= X .+ Y` and will eventually result in a fused in-place assignment. Similarly for `.*=` etcetera.)

8.13 Further Reading

We should mention here that this is far from a complete picture of defining functions. Julia has a sophisticated type system and allows multiple dispatch on argument types. None of the examples given here provide any type annotations on their arguments, meaning that they are applicable to all types of arguments. The type system is described in *Types* (page 89) and defining a function in terms of methods chosen by multiple dispatch on run-time argument types is described in *Methods* (page 109).

Control Flow

Julia provides a variety of control flow constructs:

- *Compound Expressions* (page 65): `begin` and `(;)`.
- *Conditional Evaluation* (page 66): `if-elseif-else` and `?:` (ternary operator).
- *Short-Circuit Evaluation* (page 69): `&&`, `||` and chained comparisons.
- *Repeated Evaluation: Loops* (page 71): `while` and `for`.
- *Exception Handling* (page 73): `try-catch`, `error()` (page 374) and `throw()` (page 374).
- *Tasks (aka Coroutines)* (page 78): `yieldto()` (page 485).

The first five control flow mechanisms are standard to high-level programming languages. *Task* (page 485)s are not so standard: they provide non-local control flow, making it possible to switch between temporarily-suspended computations. This is a powerful construct: both exception handling and cooperative multitasking are implemented in Julia using tasks. Everyday programming requires no direct usage of tasks, but certain problems can be solved much more easily by using tasks.

9.1 Compound Expressions

Sometimes it is convenient to have a single expression which evaluates several subexpressions in order, returning the value of the last subexpression as its value. There are two Julia constructs that accomplish this: `begin` blocks and `(;)` chains. The value of both compound expression constructs is that of the last subexpression. Here's an example of a `begin` block:

```
julia> z = begin
           x = 1
           y = 2
           x + y
       end
3
```

Since these are fairly small, simple expressions, they could easily be placed onto a single line, which is where the `(;)` chain syntax comes in handy:

```
julia> z = (x = 1; y = 2; x + y)
3
```

This syntax is particularly useful with the terse single-line function definition form introduced in *Functions* (page 55). Although it is typical, there is no requirement that `begin` blocks be multiline or that `(;)` chains be single-line:

```
julia> begin x = 1; y = 2; x + y end
3

julia> (x = 1;
        y = 2;
        x + y)
3
```

9.2 Conditional Evaluation

Conditional evaluation allows portions of code to be evaluated or not evaluated depending on the value of a boolean expression. Here is the anatomy of the `if-elseif-else` conditional syntax:

```
if x < y
    println("x is less than y")
elseif x > y
    println("x is greater than y")
else
    println("x is equal to y")
end
```

If the condition expression `x < y` is `true`, then the corresponding block is evaluated; otherwise the condition expression `x > y` is evaluated, and if it is `true`, the corresponding block is evaluated; if neither expression is `true`, the `else` block is evaluated. Here it is in action:

```
julia> function test(x, y)
    if x < y
        println("x is less than y")
    elseif x > y
        println("x is greater than y")
    else
        println("x is equal to y")
    end
end
test (generic function with 1 method)

julia> test(1, 2)
x is less than y

julia> test(2, 1)
x is greater than y

julia> test(1, 1)
x is equal to y
```

The `elseif` and `else` blocks are optional, and as many `elseif` blocks as desired can be used. The condition expressions in the `if-elseif-else` construct are evaluated until the first one evaluates to `true`, after which the associated block is evaluated, and no further condition expressions or blocks are evaluated.

`if` blocks are “leaky”, i.e. they do not introduce a local scope. This means that new variables defined inside the `if` clauses can be used after the `if` block, even if they weren’t defined before. So, we could have defined the `test` function above as

```
julia> function test(x,y)
    if x < y
        relation = "less than"
```



```

elseif x == y
    relation = "equal to"
else
    relation = "greater than"
end
println("x is ", relation, " y.")
end
test (generic function with 1 method)

```

The variable `relation` is declared inside the `if` block, but used outside. However, when depending on this behavior, make sure all possible code paths define a value for the variable. The following change to the above function results in a runtime error

```

julia> function test(x,y)
    if x < y
        relation = "less than"
    elseif x == y
        relation = "equal to"
    end
    println("x is ", relation, " y.")
end
test (generic function with 1 method)

julia> test(1,2)
x is less than y.

julia> test(2,1)
ERROR: UndefVarError: relation not defined
  in test(::Int64, ::Int64) at ./none:7
...

```

`if` blocks also return a value, which may seem unintuitive to users coming from many other languages. This value is simply the return value of the last executed statement in the branch that was chosen, so

```

julia> x = 3
3

julia> if x > 0
    "positive!"
else
    "negative..."
end
"positive!"

```

Note that very short conditional statements (one-liners) are frequently expressed using Short-Circuit Evaluation in Julia, as outlined in the next section.

Unlike C, MATLAB, Perl, Python, and Ruby — but like Java, and a few other stricter, typed languages — it is an error if the value of a conditional expression is anything but `true` or `false`:

```

julia> if 1
    println("true")
end
ERROR: TypeError: non-boolean (Int64) used in boolean context
...

```

This error indicates that the conditional was of the wrong type: `Int64` rather than the required `Bool`.

The so-called “ternary operator”, `?:`, is closely related to the `if-elseif-else` syntax, but is used where a conditional choice between single expression values is required, as opposed to conditional execution of longer blocks of

code. It gets its name from being the only operator in most languages taking three operands:

```
a ? b : c
```

The expression `a`, before the `?`, is a condition expression, and the ternary operation evaluates the expression `b`, before the `:`, if the condition `a` is `true` or the expression `c`, after the `:`, if it is `false`.

The easiest way to understand this behavior is to see an example. In the previous example, the `println` call is shared by all three branches: the only real choice is which literal string to print. This could be written more concisely using the ternary operator. For the sake of clarity, let's try a two-way version first:

```
julia> x = 1; y = 2;

julia> println(x < y ? "less than" : "not less than")
less than

julia> x = 1; y = 0;

julia> println(x < y ? "less than" : "not less than")
not less than
```

If the expression `x < y` is true, the entire ternary operator expression evaluates to the string `"less than"` and otherwise it evaluates to the string `"not less than"`. The original three-way example requires chaining multiple uses of the ternary operator together:

```
julia> test(x, y) = println(x < y ? "x is less than y"      :
                           x > y ? "x is greater than y" : "x is equal to y")
test (generic function with 1 method)

julia> test(1, 2)
x is less than y

julia> test(2, 1)
x is greater than y

julia> test(1, 1)
x is equal to y
```

To facilitate chaining, the operator associates from right to left.

It is significant that like `if-elseif-else`, the expressions before and after the `:` are only evaluated if the condition expression evaluates to `true` or `false`, respectively:

```
julia> v(x) = (println(x); x)
v (generic function with 1 method)

julia> 1 < 2 ? v("yes") : v("no")
yes
"yes"

julia> 1 > 2 ? v("yes") : v("no")
no
"no"
```

9.3 Short-Circuit Evaluation

Short-circuit evaluation is quite similar to conditional evaluation. The behavior is found in most imperative programming languages having the `&&` and `||` boolean operators: in a series of boolean expressions connected by these operators, only the minimum number of expressions are evaluated as are necessary to determine the final boolean value of the entire chain. Explicitly, this means that:

- In the expression `a && b`, the subexpression `b` is only evaluated if `a` evaluates to `true`.
- In the expression `a || b`, the subexpression `b` is only evaluated if `a` evaluates to `false`.

The reasoning is that `a && b` must be `false` if `a` is `false`, regardless of the value of `b`, and likewise, the value of `a || b` must be `true` if `a` is `true`, regardless of the value of `b`. Both `&&` and `||` associate to the right, but `&&` has higher precedence than `||` does. It's easy to experiment with this behavior:

```
julia> t(x) = (println(x); true)
t (generic function with 1 method)

julia> f(x) = (println(x); false)
f (generic function with 1 method)

julia> t(1) && t(2)
1
2
true

julia> t(1) && f(2)
1
2
false

julia> f(1) && t(2)
1
false

julia> f(1) && f(2)
1
false

julia> t(1) || t(2)
1
true

julia> t(1) || f(2)
1
true

julia> f(1) || t(2)
1
2
true

julia> f(1) || f(2)
1
2
false
```

You can easily experiment in the same way with the associativity and precedence of various combinations of `&&` and `||` operators.

This behavior is frequently used in Julia to form an alternative to very short `if` statements. Instead of `if <cond> <statement> end`, one can write `<cond> && <statement>` (which could be read as: `<cond> and then <statement>`). Similarly, instead of `if ! <cond> <statement> end`, one can write `<cond> || <statement>` (which could be read as: `<cond> or else <statement>`).

For example, a recursive factorial routine could be defined like this:

```
julia> function fact(n::Int)
    n >= 0 || error("n must be non-negative")
    n == 0 && return 1
    n * fact(n-1)
end
fact (generic function with 1 method)

julia> fact(5)
120

julia> fact(0)
1

julia> fact(-1)
ERROR: n must be non-negative
in fact(::Int64) at ./none:2
...
```

Boolean operations *without* short-circuit evaluation can be done with the bitwise boolean operators introduced in [Mathematical Operations and Elementary Functions](#) (page 25): `&` and `|`. These are normal functions, which happen to support infix operator syntax, but always evaluate their arguments:

```
julia> f(1) & t(2)
1
2
false

julia> t(1) | t(2)
1
2
true
```

Just like condition expressions used in `if`, `elseif` or the ternary operator, the operands of `&&` or `||` must be boolean values (`true` or `false`). Using a non-boolean value anywhere except for the last entry in a conditional chain is an error:

```
julia> 1 && true
ERROR: TypeError: non-boolean (Int64) used in boolean context
...
```

On the other hand, any type of expression can be used at the end of a conditional chain. It will be evaluated and returned depending on the preceding conditionals:

```
julia> true && (x = rand(2,2))
2×2 Array{Float64,2}:
 0.768448  0.673959
 0.940515  0.395453

julia> false && (x = rand(2,2))
false
```

9.4 Repeated Evaluation: Loops

There are two constructs for repeated evaluation of expressions: the `while` loop and the `for` loop. Here is an example of a `while` loop:

```
julia> i = 1;

julia> while i <= 5
    println(i)
    i += 1
end

1
2
3
4
5
```

The `while` loop evaluates the condition expression (`i <= 5` in this case), and as long it remains `true`, keeps also evaluating the body of the `while` loop. If the condition expression is `false` when the `while` loop is first reached, the body is never evaluated.

The `for` loop makes common repeated evaluation idioms easier to write. Since counting up and down like the above `while` loop does is so common, it can be expressed more concisely with a `for` loop:

```
julia> for i = 1:5
    println(i)
end

1
2
3
4
5
```

Here the `1:5` is a `Range` object, representing the sequence of numbers 1, 2, 3, 4, 5. The `for` loop iterates through these values, assigning each one in turn to the variable `i`. One rather important distinction between the previous `while` loop form and the `for` loop form is the scope during which the variable is visible. If the variable `i` has not been introduced in another scope, in the `for` loop form, it is visible only inside of the `for` loop, and not afterwards. You'll either need a new interactive session instance or a different variable name to test this:

```
julia> for j = 1:5
    println(j)
end

1
2
3
4
5

julia> j
ERROR: UndefVarError: j not defined
...
```

See *Scope of Variables* (page 81) for a detailed explanation of variable scope and how it works in Julia.

In general, the `for` loop construct can iterate over any container. In these cases, the alternative (but fully equivalent) keyword `in` or `∈` is typically used instead of `=`, since it makes the code read more clearly:

```
julia> for i in [1,4,0]
    println(i)
end
```

```
end
1
4
0

julia> for s ∈ ["foo", "bar", "baz"]
    println(s)
end
foo
bar
baz
```

Various types of iterable containers will be introduced and discussed in later sections of the manual (see, e.g., *Multi-dimensional Arrays* (page 177)).

It is sometimes convenient to terminate the repetition of a `while` before the test condition is falsified or stop iterating in a `for` loop before the end of the iterable object is reached. This can be accomplished with the `break` keyword:

```
julia> i = 1;

julia> while true
    println(i)
    if i >= 5
        break
    end
    i += 1
end
1
2
3
4
5

julia> for i = 1:1000
    println(i)
    if i >= 5
        break
    end
end
1
2
3
4
5
```

The above `while` loop would never terminate on its own, and the `for` loop would iterate up to 1000. These loops are both exited early by using the `break` keyword.

In other circumstances, it is handy to be able to stop an iteration and move on to the next one immediately. The `continue` keyword accomplishes this:

```
julia> for i = 1:10
    if i % 3 != 0
        continue
    end
    println(i)
end
3
6
```

9

This is a somewhat contrived example since we could produce the same behavior more clearly by negating the condition and placing the `println` call inside the `if` block. In realistic usage there is more code to be evaluated after the `continue`, and often there are multiple points from which one calls `continue`.

Multiple nested `for` loops can be combined into a single outer loop, forming the cartesian product of its iterables:

```
julia> for i = 1:2, j = 3:4
        println((i, j))
    end
(1, 3)
(1, 4)
(2, 3)
(2, 4)
```

A `break` statement inside such a loop exits the entire nest of loops, not just the inner one.

9.5 Exception Handling

When an unexpected condition occurs, a function may be unable to return a reasonable value to its caller. In such cases, it may be best for the exceptional condition to either terminate the program, printing a diagnostic error message, or if the programmer has provided code to handle such exceptional circumstances, allow that code to take the appropriate action.

9.5.1 Built-in Exceptions

Exceptions are thrown when an unexpected condition has occurred. The built-in Exceptions listed below all interrupt the normal flow of control.

Exception
ArgumentError (page 374)
BoundsError (page 374)
CompositeException
DivideError (page 374)
DomainError (page 375)
EOFError (page 375)
ErrorException (page 375)
InexactError (page 375)
InitError (page 375)
InterruptException (page 375)
InvalidStateException
KeyError (page 375)
LoadError (page 375)
OutOfMemoryError (page 375)
ReadOnlyMemoryError (page 375)
RemoteException (page 490)
MethodError (page 375)
OverflowError (page 375)
ParseError (page 375)
SystemError (page 375)
TypeError (page 375)
UndefRefError (page 375)
UndefVarError (page 375)
UnicodeError

For example, the `sqrt()` (page 420) function throws a `DomainError` (page 375) if applied to a negative real value:

```
julia> sqrt(-1)
ERROR: DomainError:
sqrt will only return a complex result if called with a complex argument. Try sqrt(complex(x)).
in sqrt(::Int64) at ./math.jl:278
...
```

You may define your own exceptions in the following way:

```
julia> type MyCustomException <: Exception end
```

9.5.2 The `throw()` function

Exceptions can be created explicitly with `throw()` (page 374). For example, a function defined only for nonnegative numbers could be written to `throw()` (page 374) a `DomainError` (page 375) if the argument is negative:

```
julia> f(x) = x>=0 ? exp(-x) : throw(DomainError())
f (generic function with 1 method)

julia> f(1)
0.36787944117144233

julia> f(-1)
ERROR: DomainError:
in f(::Int64) at ./none:1
...
```

Note that `DomainError` (page 375) without parentheses is not an exception, but a type of exception. It needs to be called to obtain an `Exception` object:


```
julia> typeof(DomainError()) <: Exception
true

julia> typeof(DomainError) <: Exception
false
```

Additionally, some exception types take one or more arguments that are used for error reporting:

```
julia> throw(UndefVarError(:x))
ERROR: UndefVarError: x not defined
...
```

This mechanism can be implemented easily by custom exception types following the way *UndefVarError* (page 375) is written:

```
julia> type MyUndefVarError <: Exception
    var::Symbol
end
julia> Base.showerror(io::IO, e::MyUndefVarError) = print(io, e.var, " not defined");
```

Note: When writing an error message, it is preferred to make the first word lowercase. For example, `size(A) == size(B) || throw(DimensionMismatch("size of A not equal to size of B"))`

is preferred over

```
size(A) == size(B) || throw(DimensionMismatch("Size of A not equal to size of B")).
```

However, sometimes it makes sense to keep the uppercase first letter, for instance if an argument to a function is a capital letter: `size(A,1) == size(B,2) || throw(DimensionMismatch("A has first dimension..."))`.

9.5.3 Errors

The *error()* (page 374) function is used to produce an *ErrorException* (page 375) that interrupts the normal flow of control.

Suppose we want to stop execution immediately if the square root of a negative number is taken. To do this, we can define a fussy version of the *sqrt()* (page 420) function that raises an error if its argument is negative:

```
julia> fussy_sqrt(x) = x >= 0 ? sqrt(x) : error("negative x not allowed")
fussy_sqrt (generic function with 1 method)

julia> fussy_sqrt(2)
1.4142135623730951

julia> fussy_sqrt(-1)
ERROR: negative x not allowed
in fussy_sqrt(::Int64) at ./none:1
...
```

If *fussy_sqrt* is called with a negative value from another function, instead of trying to continue execution of the calling function, it returns immediately, displaying the error message in the interactive session:

```
julia> function verbose_fussy_sqrt(x)
    println("before fussy_sqrt")
    r = fussy_sqrt(x)
end
```

```
        println("after fussy_sqrt")
        return r
    end
verbose_fussy_sqrt (generic function with 1 method)

julia> verbose_fussy_sqrt(2)
before fussy_sqrt
after fussy_sqrt
1.4142135623730951

julia> verbose_fussy_sqrt(-1)
before fussy_sqrt
ERROR: negative x not allowed
  in fussy_sqrt at ./none:1 [inlined]
  in verbose_fussy_sqrt(::Int64) at ./none:3
...
```

9.5.4 Warnings and informational messages

Julia also provides other functions that write messages to the standard error I/O, but do not throw any `Exceptions` and hence do not interrupt execution:

```
julia> info("Hi"); 1+1
INFO: Hi
2

julia> warn("Hi"); 1+1
WARNING: Hi
2

julia> error("Hi"); 1+1
ERROR: Hi
  in error(::String) at ./error.jl:21
...
```

9.5.5 The `try/catch` statement

The `try/catch` statement allows for `Exceptions` to be tested for. For example, a customized square root function can be written to automatically call either the real or complex square root method on demand using `Exceptions`:

```
julia> f(x) = try
           sqrt(x)
         catch
           sqrt(complex(x, 0))
         end
f (generic function with 1 method)

julia> f(1)
1.0

julia> f(-1)
0.0 + 1.0im
```

It is important to note that in real code computing this function, one would compare `x` to zero instead of catching an exception. The exception is much slower than simply comparing and branching.

`try/catch` statements also allow the `Exception` to be saved in a variable. In this contrived example, the following example calculates the square root of the second element of `x` if `x` is indexable, otherwise assumes `x` is a real number and returns its square root:

```
julia> sqrt_second(x) = try
    sqrt(x[2])
catch y
    if isa(y, DomainError)
        sqrt(complex(x[2], 0))
    elseif isa(y, BoundsError)
        sqrt(x)
    end
end
sqrt_second (generic function with 1 method)

julia> sqrt_second([1 4])
2.0

julia> sqrt_second([1 -4])
0.0 + 2.0im

julia> sqrt_second(9)
3.0

julia> sqrt_second(-9)
ERROR: DomainError:
 in sqrt_second(::Int64) at ./none:7
...
```

Note that the symbol following `catch` will always be interpreted as a name for the exception, so care is needed when writing `try/catch` expressions on a single line. The following code will *not* work to return the value of `x` in case of an error:

```
try bad() catch x end
```

Instead, use a semicolon or insert a line break after `catch`:

```
try bad() catch; x end

try bad()
catch
    x
end
```

The `catch` clause is not strictly necessary; when omitted, the default return value is `nothing`.

```
julia> try error() end #Returns nothing
```

The power of the `try/catch` construct lies in the ability to unwind a deeply nested computation immediately to a much higher level in the stack of calling functions. There are situations where no error has occurred, but the ability to unwind the stack and pass a value to a higher level is desirable. Julia provides the `rethrow()` (page 374), `backtrace()` (page 374) and `catch_backtrace()` (page 374) functions for more advanced error handling.

9.5.6 finally Clauses

In code that performs state changes or uses resources like files, there is typically clean-up work (such as closing files) that needs to be done when the code is finished. Exceptions potentially complicate this task, since they can cause a

block of code to exit before reaching its normal end. The `finally` keyword provides a way to run some code when a given block of code exits, regardless of how it exits.

For example, here is how we can guarantee that an opened file is closed:

```
f = open("file")
try
    # operate on file f
finally
    close(f)
end
```

When control leaves the `try` block (for example due to a `return`, or just finishing normally), `close(f)` will be executed. If the `try` block exits due to an exception, the exception will continue propagating. A `catch` block may be combined with `try` and `finally` as well. In this case the `finally` block will run after `catch` has handled the error.

9.6 Tasks (aka Coroutines)

Tasks are a control flow feature that allows computations to be suspended and resumed in a flexible manner. This feature is sometimes called by other names, such as symmetric coroutines, lightweight threads, cooperative multitasking, or one-shot continuations.

When a piece of computing work (in practice, executing a particular function) is designated as a *Task* (page 485), it becomes possible to interrupt it by switching to another *Task* (page 485). The original *Task* (page 485) can later be resumed, at which point it will pick up right where it left off. At first, this may seem similar to a function call. However there are two key differences. First, switching tasks does not use any space, so any number of task switches can occur without consuming the call stack. Second, switching among tasks can occur in any order, unlike function calls, where the called function must finish executing before control returns to the calling function.

This kind of control flow can make it much easier to solve certain problems. In some problems, the various pieces of required work are not naturally related by function calls; there is no obvious “caller” or “callee” among the jobs that need to be done. An example is the producer-consumer problem, where one complex procedure is generating values and another complex procedure is consuming them. The consumer cannot simply call a producer function to get a value, because the producer may have more values to generate and so might not yet be ready to return. With tasks, the producer and consumer can both run as long as they need to, passing values back and forth as necessary.

Julia provides the functions `produce()` (page 486) and `consume()` (page 486) for solving this problem. A producer is a function that calls `produce()` (page 486) on each value it needs to produce:

```
julia> function producer()
    produce("start")
    for n=1:4
        produce(2n)
    end
    produce("stop")
end;
```

To consume values, first the producer is wrapped in a *Task* (page 485), then `consume()` (page 486) is called repeatedly on that object:

```
julia> p = Task(producer);

julia> consume(p)
"start"

julia> consume(p)
```

```

2

julia> consume(p)
4

julia> consume(p)
6

julia> consume(p)
8

julia> consume(p)
"stop"

```

One way to think of this behavior is that `producer` was able to return multiple times. Between calls to `produce()` (page 486), the producer’s execution is suspended and the consumer has control.

A `Task` can be used as an iterable object in a `for` loop, in which case the loop variable takes on all the produced values:

```

julia> for x in Task(producer)
        println(x)
    end
start
2
4
6
8
stop

```

Note that the `Task()` (page 485) constructor expects a 0-argument function. A common pattern is for the producer to be parameterized, in which case a partial function application is needed to create a 0-argument *anonymous function* (page 57). This can be done either directly or by use of a convenience macro:

```

function mytask(myarg)
    ...
end

taskHdl = Task{() -> mytask(7)}
# or, equivalently
taskHdl = @task mytask(7)

```

`produce()` (page 486) and `consume()` (page 486) do not launch threads that can run on separate CPUs. True kernel threads are discussed under the topic of *Parallel Computing* (page 199).

9.6.1 Core task operations

While `produce()` (page 486) and `consume()` (page 486) illustrate the essential nature of tasks, they are actually implemented as library functions using a more primitive function, `yieldto()` (page 485). `yieldto(task, value)` suspends the current task, switches to the specified `task`, and causes that task’s last `yieldto()` (page 485) call to return the specified `value`. Notice that `yieldto()` (page 485) is the only operation required to use task-style control flow; instead of calling and returning we are always just switching to a different task. This is why this feature is also called “symmetric coroutines”; each task is switched to and from using the same mechanism.

`yieldto()` (page 485) is powerful, but most uses of tasks do not invoke it directly. Consider why this might be. If you switch away from the current task, you will probably want to switch back to it at some point, but knowing when to

switch back, and knowing which task has the responsibility of switching back, can require considerable coordination. For example, `produce()` (page 486) needs to maintain some state to remember who the consumer is. Not needing to manually keep track of the consuming task is what makes `produce()` (page 486) easier to use than `yieldto()` (page 485).

In addition to `yieldto()` (page 485), a few other basic functions are needed to use tasks effectively.

- `current_task()` (page 485) gets a reference to the currently-running task.
- `istaskdone()` (page 485) queries whether a task has exited.
- `istaskstarted()` (page 485) queries whether a task has run yet.
- `task_local_storage()` (page 486) manipulates a key-value store specific to the current task.

9.6.2 Tasks and events

Most task switches occur as a result of waiting for events such as I/O requests, and are performed by a scheduler included in the standard library. The scheduler maintains a queue of runnable tasks, and executes an event loop that restarts tasks based on external events such as message arrival.

The basic function for waiting for an event is `wait()` (page 490). Several objects implement `wait()` (page 490); for example, given a `Process` object, `wait()` (page 490) will wait for it to exit. `wait()` (page 490) is often implicit; for example, a `wait()` (page 490) can happen inside a call to `read()` (page 562) to wait for data to be available.

In all of these cases, `wait()` (page 490) ultimately operates on a `Condition` (page 486) object, which is in charge of queueing and restarting tasks. When a task calls `wait()` (page 490) on a `Condition` (page 486), the task is marked as non-runnable, added to the condition's queue, and switches to the scheduler. The scheduler will then pick another task to run, or block waiting for external events. If all goes well, eventually an event handler will call `notify()` (page 486) on the condition, which causes tasks waiting for that condition to become runnable again.

A task created explicitly by calling `Task` (page 485) is initially not known to the scheduler. This allows you to manage tasks manually using `yieldto()` (page 485) if you wish. However, when such a task waits for an event, it still gets restarted automatically when the event happens, as you would expect. It is also possible to make the scheduler run a task whenever it can, without necessarily waiting for any events. This is done by calling `schedule()` (page 486), or using the `@schedule` (page 486) or `@async` (page 493) macros (see *Parallel Computing* (page 199) for more details).

9.6.3 Task states

Tasks have a `state` field that describes their execution status. A task state is one of the following symbols:

Symbol	Meaning
<code>:runnable</code>	Currently running, or available to be switched to
<code>:waiting</code>	Blocked waiting for a specific event
<code>:queued</code>	In the scheduler's run queue about to be restarted
<code>:done</code>	Successfully finished executing
<code>:failed</code>	Finished with an uncaught exception

Scope of Variables

The *scope* of a variable is the region of code within which a variable is visible. Variable scoping helps avoid variable naming conflicts. The concept is intuitive: two functions can both have arguments called `x` without the two `x`'s referring to the same thing. Similarly there are many other cases where different blocks of code can use the same name without referring to the same thing. The rules for when the same variable name does or doesn't refer to the same thing are called scope rules; this section spells them out in detail.

Certain constructs in the language introduce *scope blocks*, which are regions of code that are eligible to be the scope of some set of variables. The scope of a variable cannot be an arbitrary set of source lines; instead, it will always line up with one of these blocks. There are two main types of scopes in Julia, *global scope* and *local scope*, the latter can be nested. The constructs introducing scope blocks are:

Scope name	block/construct introducing this kind of scope	
global (page 82)	module, baremodule, at interactive prompt (REPL)	
local (page 82)	soft (page 83)	for, while, comprehensions, try-catch-finally, let
	hard (page 84)	functions (either syntax, anonymous & do-blocks), type, immutable, macro

Notably missing from this table are *begin blocks* (page 65) and *if blocks* (page 66), which do *not* introduce new scope blocks. All three types of scopes follow somewhat different rules which will be explained below as well as some extra rules for certain blocks.

Julia uses *lexical scoping*, meaning that a function's scope does not inherit from its caller's scope, but from the scope in which the function was defined. For example, in the following code the `x` inside `foo` refers to the `x` in the global scope of its module `Bar`:

```
julia> module Bar
    x = 1
    foo() = x
end;
```

and not a `x` in the scope where `foo` is used:

```
julia> import Bar

julia> x = -1;

julia> Bar.foo()
1
```

Thus *lexical scope* means that the scope of variables can be inferred from the source code alone.

10.1 Global Scope

Each module introduces a new *global scope*, separate from the global scope of all other modules; there is no all-encompassing global scope. Modules can introduce variables of other modules into their scope through the *using* or *import* (page 139) statements or through qualified access using the dot-notation, i.e. each module is a so-called *namespace*. Note that variable bindings can only be changed within their global scope and not from an outside module.

```
julia> module A
    a = 1 # a global in A's scope
end;

julia> module B
    module C
        c = 2
    end
    b = C.c # can access the namespace of a nested global scope
            # through a qualified access
    import A # makes module A available
    d = A.a
end;

julia> module D
    b = a # errors as D's global scope is separate from A's
end;
ERROR: UndefVarError: a not defined
...

julia> module E
    import A # make module A available
    A.a = 2 # throws below error
end;
ERROR: cannot assign variables in other modules
...
```

Note that the interactive prompt (aka REPL) is in the global scope of the module `Main`.

10.2 Local Scope

A new local scope is introduced by most code-blocks, see above *table* (page 81) for a complete list. A local scope *usually* inherits all the variables from its parent scope, both for reading and writing. There are two subtypes of local scopes, hard and soft, with slightly different rules concerning what variables are inherited. Unlike global scopes, local scopes are not namespaces, thus variables in an inner scope cannot be retrieved from the parent scope through some sort of qualified access.

The following rules and examples pertain to both hard and soft local scopes. A newly introduced variable in a local scope does not back-propagate to its parent scope. For example, here the `z` is not introduced into the top-level scope:

```
julia> for i = 1:10
    z = i
end

julia> z
ERROR: UndefVarError: z not defined
...
```


(Note, in this and all following examples it is assumed that their top-level is a global scope with a clean workspace, for instance a newly started REPL.)

Inside a local scope a variable can be forced to be a local variable using the `local` keyword:

```
julia> x = 0;

julia> for i = 1:10
    local x
    x = i + 1
end

julia> x
0
```

Inside a local scope a new global variable can be defined using the keyword `global`:

```
julia> for i = 1:10
    global z
    z = i
end

julia> z
10
```

The location of both the `local` and `global` keywords within the scope block is irrelevant. The following is equivalent to the last example (although stylistically worse):

```
julia> for i = 1:10
    z = i
    global z
end

julia> z
10
```

Multiple global or local definitions can be on one line and can also be paired with assignments:

```
julia> for i = 1:10
    global x = i, y, z
    local a = 4, b, c = 1
end
```

10.2.1 Soft Local Scope

In a soft local scope, all variables are inherited from its parent scope unless a variable is specifically marked with the keyword `local`.

Soft local scopes are introduced by for-loops, while-loops, comprehensions, try-catch-finally-blocks, and let-blocks. There are some extra rules for *let-blocks* (page 86) and for *for-loops and comprehensions* (page 87).

In the following example the `x` and `y` refer always to the same variables as the soft local scope inherits both read and write variables:

```
julia> x, y = 0, 1;

julia> for i = 1:10
    x = i + y + 1
end
```

```
julia> x
12
```

Within soft scopes, the *global* keyword is never necessary, although allowed. The only case when it would change the semantics is (currently) a syntax error:

```
julia> let
    local j = 2
    let
        global j = 3
    end
end
ERROR: syntax: `global j`: j is local variable in the enclosing scope
...
```

10.2.2 Hard Local Scope

Hard local scopes are introduced by function definitions (in all their forms), type & immutable-blocks, and macro-definitions.

In a hard local scope, all variables are inherited from its parent scope unless:

- an assignment would result in a modified *global* variable, or
- a variable is specifically marked with the keyword `local`.

Thus global variables are only inherited for reading but not for writing:

```
julia> x, y = 1, 2;

julia> function foo()
    x = 2 # assignment introduces a new local
    return x + y # y refers to the global
end;

julia> foo()
4

julia> x
1
```

An explicit `global` is needed to assign to a global variable:

```
julia> x = 1;

julia> function foobar()
    global x = 2
end;

julia> foobar();

julia> x
2
```

Note that *nested functions* can behave differently to functions defined in the global scope as they can modify their parent scope's *local* variables:

```
julia> x, y = 1, 2;
```

```
julia> function baz()
    x = 2 # introduces a new local
    function bar()
        x = 10 # modifies the parent's x
        return x + y # y is global
    end
    return bar() + x # 12 + 10 (x is modified in call of bar())
end;

julia> baz()
22

julia> x, y
(1, 2)
```

The distinction between inheriting global and local variables for assignment can lead to some slight differences between functions defined in local vs. global scopes. Consider the modification of the last example by moving `bar` to the global scope:

```
julia> x, y = 1, 2;

julia> function bar()
    x = 10 # local
    return x + y
end;

julia> function quz()
    x = 2 # local
    return bar() + x # 12 + 2 (x is not modified)
end;

julia> quz()
14

julia> x, y
(1, 2)
```

Note that above subtlety does not pertain to type and macro definitions as they can only appear at the global scope. There are special scoping rules concerning the evaluation of default and keyword function arguments which are described in the [Function section](#) (page 62).

An assignment introducing a variable used inside a function, type or macro definition need not come before its inner usage:

```
julia> f = y -> y + a
(::#2) (generic function with 1 method)

julia> f(3)
ERROR: UndefVarError: a not defined
  in (::#2#3)(::Int64) at ./none:1
...

julia> a = 1
1

julia> f(3)
4
```

This behavior may seem slightly odd for a normal variable, but allows for named functions — which are just normal variables holding function objects — to be used before they are defined. This allows functions to be defined in whatever order is intuitive and convenient, rather than forcing bottom up ordering or requiring forward declarations, as long as they are defined by the time they are actually called. As an example, here is an inefficient, mutually recursive way to test if positive integers are even or odd:

```
julia> even(n) = n == 0 ? true : odd(n-1);

julia> odd(n) = n == 0 ? false : even(n-1);

julia> even(3)
false

julia> odd(3)
true
```

Julia provides built-in, efficient functions to test for oddness and evenness called `iseven()` (page 439) and `isodd()` (page 439) so the above definitions should only be taken as examples.

10.2.3 Hard vs. Soft Local Scope

Blocks which introduce a soft local scope, such as loops, are generally used to manipulate the variables in their parent scope. Thus their default is to fully access all variables in their parent scope.

Conversely, the code inside blocks which introduce a hard local scope (function, type, and macro definitions) can be executed at any place in a program. Remotely changing the state of global variables in other modules should be done with care and thus this is an opt-in feature requiring the `global` keyword.

The reason to allow *modifying local* variables of parent scopes in nested functions is to allow constructing *closures* which have a private state, for instance the `state` variable in the following example:

```
julia> let
    state = 0
    global counter
    counter() = state += 1
end;

julia> counter()
1

julia> counter()
2
```

See also the closures in the examples in the next two sections.

10.2.4 Let Blocks

Unlike assignments to local variables, `let` statements allocate new variable bindings each time they run. An assignment modifies an existing value location, and `let` creates new locations. This difference is usually not important, and is only detectable in the case of variables that outlive their scope via closures. The `let` syntax accepts a comma-separated series of assignments and variable names:

```
julia> x, y, z = -1, -1, -1;

julia> let x = 1, z
    println("x: $x, y: $y") # x is local variable, y the global
    println("z: $z") # errors as z has not been assigned yet but is local
end
```

```

    end
x: 1, y: -1
ERROR: UndefVarError: z not defined
...

```

The assignments are evaluated in order, with each right-hand side evaluated in the scope before the new variable on the left-hand side has been introduced. Therefore it makes sense to write something like `let x = x` since the two `x` variables are distinct and have separate storage. Here is an example where the behavior of `let` is needed:

```

julia> Fs = Array{Any}(2); i = 1;

julia> while i <= 2
           Fs[i] = ()->i
           i += 1
       end

julia> Fs[1]()
3

julia> Fs[2]()
3

```

Here we create and store two closures that return variable `i`. However, it is always the same variable `i`, so the two closures behave identically. We can use `let` to create a new binding for `i`:

```

julia> Fs = Array{Any}(2); i = 1;

julia> while i <= 2
           let i = i
               Fs[i] = ()->i
           end
           i += 1
       end

julia> Fs[1]()
1

julia> Fs[2]()
2

```

Since the `begin` construct does not introduce a new scope, it can be useful to use a zero-argument `let` to just introduce a new scope block without creating any new bindings:

```

julia> let
           local x = 1
           let
               local x = 2
           end
           x
       end

1

```

Since `let` introduces a new scope block, the inner `local x` is a different variable than the outer `local x`.

10.2.5 For Loops and Comprehensions

`for` loops and [comprehensions](#) (page 179) have the following behavior: any new variables introduced in their body scopes are freshly allocated for each loop iteration. This is in contrast to `while` loops which reuse the variables for

all iterations. Therefore these constructs are similar to `while` loops with `let` blocks inside:

```
julia> Fs = Array{Any}(2);

julia> for j = 1:2
           Fs[j] = ()->j
       end

julia> Fs[1]()
1

julia> Fs[2]()
2
```

`for` loops will reuse existing variables for its iteration variable:

```
julia> i = 0;

julia> for i = 1:3
       end

julia> i
3
```

However, comprehensions do not do this, and always freshly allocate their iteration variables:

```
julia> x = 0;

julia> [ x for x = 1:3 ];

julia> x
0
```

10.3 Constants

A common use of variables is giving names to specific, unchanging values. Such variables are only assigned once. This intent can be conveyed to the compiler using the `const` keyword:

```
julia> const e = 2.71828182845904523536;

julia> const pi = 3.14159265358979323846;
```

The `const` declaration is allowed on both global and local variables, but is especially useful for globals. It is difficult for the compiler to optimize code involving global variables, since their values (or even their types) might change at almost any time. If a global variable will not change, adding a `const` declaration solves this performance problem.

Local constants are quite different. The compiler is able to determine automatically when a local variable is constant, so local constant declarations are not necessary for performance purposes.

Special top-level assignments, such as those performed by the `function` and `type` keywords, are constant by default.

Note that `const` only affects the variable binding; the variable may be bound to a mutable object (such as an array), and that object may still be modified.

Types

Type systems have traditionally fallen into two quite different camps: static type systems, where every program expression must have a type computable before the execution of the program, and dynamic type systems, where nothing is known about types until run time, when the actual values manipulated by the program are available. Object orientation allows some flexibility in statically typed languages by letting code be written without the precise types of values being known at compile time. The ability to write code that can operate on different types is called polymorphism. All code in classic dynamically typed languages is polymorphic: only by explicitly checking types, or when objects fail to support operations at run-time, are the types of any values ever restricted.

Julia’s type system is dynamic, but gains some of the advantages of static type systems by making it possible to indicate that certain values are of specific types. This can be of great assistance in generating efficient code, but even more significantly, it allows method dispatch on the types of function arguments to be deeply integrated with the language. Method dispatch is explored in detail in *Methods* (page 109), but is rooted in the type system presented here.

The default behavior in Julia when types are omitted is to allow values to be of any type. Thus, one can write many useful Julia programs without ever explicitly using types. When additional expressiveness is needed, however, it is easy to gradually introduce explicit type annotations into previously “untyped” code. Doing so will typically increase both the performance and robustness of these systems, and perhaps somewhat counterintuitively, often significantly simplify them.

Describing Julia in the lingo of *type systems*, it is: dynamic, nominative and parametric. Generic types can be parameterized, and the hierarchical relationships between types are *explicitly declared*, rather than *implied by compatible structure*. One particularly distinctive feature of Julia’s type system is that concrete types may not subtype each other: all concrete types are final and may only have abstract types as their supertypes. While this might at first seem unduly restrictive, it has many beneficial consequences with surprisingly few drawbacks. It turns out that being able to inherit behavior is much more important than being able to inherit structure, and inheriting both causes significant difficulties in traditional object-oriented languages. Other high-level aspects of Julia’s type system that should be mentioned up front are:

- There is no division between object and non-object values: all values in Julia are true objects having a type that belongs to a single, fully connected type graph, all nodes of which are equally first-class as types.
- There is no meaningful concept of a “compile-time type”: the only type a value has is its actual type when the program is running. This is called a “run-time type” in object-oriented languages where the combination of static compilation with polymorphism makes this distinction significant.
- Only values, not variables, have types — variables are simply names bound to values.
- Both abstract and concrete types can be parameterized by other types. They can also be parameterized by symbols, by values of any type for which `isbits()` (page 367) returns true (essentially, things like numbers and bools that are stored like C types or structs with no pointers to other objects), and also by tuples thereof. Type parameters may be omitted when they do not need to be referenced or restricted.

Julia’s type system is designed to be powerful and expressive, yet clear, intuitive and unobtrusive. Many Julia programmers may never feel the need to write code that explicitly uses types. Some kinds of programming, however, become clearer, simpler, faster and more robust with declared types.

11.1 Type Declarations

The `::` operator can be used to attach type annotations to expressions and variables in programs. There are two primary reasons to do this:

1. As an assertion to help confirm that your program works the way you expect,
2. To provide extra type information to the compiler, which can then improve performance in some cases

When appended to an expression computing a value, the `::` operator is read as “is an instance of”. It can be used anywhere to assert that the value of the expression on the left is an instance of the type on the right. When the type on the right is concrete, the value on the left must have that type as its implementation — recall that all concrete types are final, so no implementation is a subtype of any other. When the type is abstract, it suffices for the value to be implemented by a concrete type that is a subtype of the abstract type. If the type assertion is not true, an exception is thrown, otherwise, the left-hand value is returned:

```
julia> (1+2)::AbstractFloat
ERROR: TypeError: typeassert: expected AbstractFloat, got Int64
...

julia> (1+2)::Int
3
```

This allows a type assertion to be attached to any expression in-place.

When appended to a variable on the left-hand side of an assignment, or as part of a `local` declaration, the `::` operator means something a bit different: it declares the variable to always have the specified type, like a type declaration in a statically-typed language such as C. Every value assigned to the variable will be converted to the declared type using `convert()` (page 364):

```
julia> function foo()
    x::Int8 = 100
    x
end
foo (generic function with 1 method)

julia> foo()
100

julia> typeof(ans)
Int8
```

This feature is useful for avoiding performance “gotchas” that could occur if one of the assignments to a variable changed its type unexpectedly.

This “declaration” behavior only occurs in specific contexts:

```
local x::Int8 # in a local declaration
x::Int8 = 10 # as the left-hand side of an assignment
```

and applies to the whole current scope, even before the declaration. Currently, type declarations cannot be used in global scope, e.g. in the REPL, since Julia does not yet have constant-type globals.

Declarations can also be attached to function definitions:


```
function sinc(x)::Float64
    if x == 0
        return 1
    end
    return sin(pi*x)/(pi*x)
end
```

Returning from this function behaves just like an assignment to a variable with a declared type: the value is always converted to `Float64`.

11.2 Abstract Types

Abstract types cannot be instantiated, and serve only as nodes in the type graph, thereby describing sets of related concrete types: those concrete types which are their descendants. We begin with abstract types even though they have no instantiation because they are the backbone of the type system: they form the conceptual hierarchy which makes Julia’s type system more than just a collection of object implementations.

Recall that in *Integers and Floating-Point Numbers* (page 13), we introduced a variety of concrete types of numeric values: `Int8`, `UInt8`, `Int16`, `UInt16`, `Int32`, `UInt32`, `Int64`, `UInt64`, `Int128`, `UInt128`, `Float16`, `Float32` (page 437), and `Float64` (page 437). Although they have different representation sizes, `Int8`, `Int16`, `Int32`, `Int64` and `Int128` all have in common that they are signed integer types. Likewise `UInt8`, `UInt16`, `UInt32`, `UInt64` and `UInt128` are all unsigned integer types, while `Float16`, `Float32` (page 437) and `Float64` (page 437) are distinct in being floating-point types rather than integers. It is common for a piece of code to make sense, for example, only if its arguments are some kind of integer, but not really depend on what particular *kind* of integer. For example, the greatest common denominator algorithm works for all kinds of integers, but will not work for floating-point numbers. Abstract types allow the construction of a hierarchy of types, providing a context into which concrete types can fit. This allows you, for example, to easily program to any type that is an integer, without restricting an algorithm to a specific type of integer.

Abstract types are declared using the `abstract` keyword. The general syntaxes for declaring an abstract type are:

```
abstract <name>
abstract <name> <: <supertype>
```

The `abstract` keyword introduces a new abstract type, whose name is given by `<name>`. This name can be optionally followed by `<:` and an already-existing type, indicating that the newly declared abstract type is a subtype of this “parent” type.

When no supertype is given, the default supertype is `Any` — a predefined abstract type that all objects are instances of and all types are subtypes of. In type theory, `Any` is commonly called “top” because it is at the apex of the type graph. Julia also has a predefined abstract “bottom” type, at the nadir of the type graph, which is written as `Union{}`. It is the exact opposite of `Any`: no object is an instance of `Union{}` and all types are supertypes of `Union{}`.

Let’s consider some of the abstract types that make up Julia’s numerical hierarchy:

```
abstract Number
abstract Real <: Number
abstract AbstractFloat <: Real
abstract Integer <: Real
abstract Signed <: Integer
abstract Unsigned <: Integer
```

The `Number` type is a direct child type of `Any`, and `Real` is its child. In turn, `Real` has two children (it has more, but only two are shown here; we’ll get to the others later): `Integer` and `AbstractFloat`, separating the world into representations of integers and representations of real numbers. Representations of real numbers include, of course, floating-point types, but also include other types, such as rationals. Hence, `AbstractFloat` is a proper subtype

of `Real`, including only floating-point representations of real numbers. Integers are further subdivided into `Signed` and `Unsigned` varieties.

The `<:` operator in general means “is a subtype of”, and, used in declarations like this, declares the right-hand type to be an immediate supertype of the newly declared type. It can also be used in expressions as a subtype operator which returns `true` when its left operand is a subtype of its right operand:

```
julia> Integer <: Number
true

julia> Integer <: AbstractFloat
false
```

An important use of abstract types is to provide default implementations for concrete types. To give a simple example, consider:

```
function myplus(x, y)
    x+y
end
```

The first thing to note is that the above argument declarations are equivalent to `x::Any` and `y::Any`. When this function is invoked, say as `myplus(2, 5)`, the dispatcher chooses the most specific method named `myplus` that matches the given arguments. (See [Methods](#) (page 109) for more information on multiple dispatch.)

Assuming no method more specific than the above is found, Julia next internally defines and compiles a method called `myplus` specifically for two `Int` arguments based on the generic function given above, i.e., it implicitly defines and compiles:

```
function myplus(x::Int, y::Int)
    x+y
end
```

and finally, it invokes this specific method.

Thus, abstract types allow programmers to write generic functions that can later be used as the default method by many combinations of concrete types. Thanks to multiple dispatch, the programmer has full control over whether the default or more specific method is used.

An important point to note is that there is no loss in performance if the programmer relies on a function whose arguments are abstract types, because it is recompiled for each tuple of argument concrete types with which it is invoked. (There may be a performance issue, however, in the case of function arguments that are containers of abstract types; see [Performance Tips](#) (page 305).)

11.3 Bits Types

A bits type is a concrete type whose data consists of plain old bits. Classic examples of bits types are integers and floating-point values. Unlike most languages, Julia lets you declare your own bits types, rather than providing only a fixed set of built-in bits types. In fact, the standard bits types are all defined in the language itself:

```
bitstype 16 Float16 <: AbstractFloat
bitstype 32 Float32 <: AbstractFloat
bitstype 64 Float64 <: AbstractFloat

bitstype 8 Bool <: Integer
bitstype 32 Char

bitstype 8 Int8 <: Signed
bitstype 8 UInt8 <: Unsigned
```

```

bitstype 16 Int16    <: Signed
bitstype 16 UInt16   <: Unsigned
bitstype 32 Int32    <: Signed
bitstype 32 UInt32   <: Unsigned
bitstype 64 Int64    <: Signed
bitstype 64 UInt64   <: Unsigned
bitstype 128 Int128  <: Signed
bitstype 128 UInt128 <: Unsigned

```

The general syntaxes for declaration of a `bitstype` are:

```

bitstype <bits> <name>
bitstype <bits> <name> <: <supertype>

```

The number of bits indicates how much storage the type requires and the name gives the new type a name. A bits type can optionally be declared to be a subtype of some supertype. If a supertype is omitted, then the type defaults to having `Any` as its immediate supertype. The declaration of `Bool` above therefore means that a boolean value takes eight bits to store, and has `Integer` as its immediate supertype. Currently, only sizes that are multiples of 8 bits are supported. Therefore, boolean values, although they really need just a single bit, cannot be declared to be any smaller than eight bits.

The types `Bool`, `Int8` and `UInt8` all have identical representations: they are eight-bit chunks of memory. Since Julia’s type system is nominative, however, they are not interchangeable despite having identical structure. A fundamental difference between them is that they have different supertypes: `Bool`’s direct supertype is `Integer`, `Int8`’s is `Signed`, and `UInt8`’s is `Unsigned`. All other differences between `Bool`, `Int8`, and `UInt8` are matters of behavior — the way functions are defined to act when given objects of these types as arguments. This is why a nominative type system is necessary: if structure determined type, which in turn dictates behavior, then it would be impossible to make `Bool` behave any differently than `Int8` or `UInt8`.

11.4 Composite Types

Composite types are called records, structures (`structs` in C), or objects in various languages. A composite type is a collection of named fields, an instance of which can be treated as a single value. In many languages, composite types are the only kind of user-definable type, and they are by far the most commonly used user-defined type in Julia as well.

In mainstream object oriented languages, such as C++, Java, Python and Ruby, composite types also have named functions associated with them, and the combination is called an “object”. In purer object-oriented languages, such as Ruby or Smalltalk, all values are objects whether they are composites or not. In less pure object oriented languages, including C++ and Java, some values, such as integers and floating-point values, are not objects, while instances of user-defined composite types are true objects with associated methods. In Julia, all values are objects, but functions are not bundled with the objects they operate on. This is necessary since Julia chooses which method of a function to use by multiple dispatch, meaning that the types of *all* of a function’s arguments are considered when selecting a method, rather than just the first one (see [Methods](#) (page 109) for more information on methods and dispatch). Thus, it would be inappropriate for functions to “belong” to only their first argument. Organizing methods into function objects rather than having named bags of methods “inside” each object ends up being a highly beneficial aspect of the language design.

Since composite types are the most common form of user-defined concrete type, they are simply introduced with the `type` keyword followed by a block of field names, optionally annotated with types using the `::` operator:

```

julia> type Foo
    bar
    baz::Int
    qux::Float64
end

```

Fields with no type annotation default to `Any`, and can accordingly hold any type of value.

New objects of composite type `Foo` are created by applying the `Foo` type object like a function to values for its fields:

```
julia> foo = Foo("Hello, world.", 23, 1.5)
Foo("Hello, world.", 23, 1.5)

julia> typeof(foo)
Foo
```

When a type is applied like a function it is called a *constructor*. Two constructors are generated automatically (these are called *default constructors*). One accepts any arguments and calls `convert()` (page 364) to convert them to the types of the fields, and the other accepts arguments that match the field types exactly. The reason both of these are generated is that this makes it easier to add new definitions without inadvertently replacing a default constructor.

Since the `bar` field is unconstrained in type, any value will do. However, the value for `baz` must be convertible to `Int`:

```
julia> Foo(), 23.5, 1)
ERROR: InexactError()
  in convert(::Type{Int64}, ::Float64) at ./float.jl:656
  in Foo(::Tuple{}, ::Float64, ::Int64) at ./none:2
  ...
```

You may find a list of field names using the `fieldnames` function.

```
julia> fieldnames(foo)
3-element Array{Symbol,1}:
 :bar
 :baz
 :qux
```

You can access the field values of a composite object using the traditional `foo.bar` notation:

```
julia> foo.bar
"Hello, world."

julia> foo.baz
23

julia> foo.qux
1.5
```

You can also change the values as one would expect:

```
julia> foo.qux = 2
2

julia> foo.bar = 1//2
1//2
```

Composite types with no fields are singletons; there can be only one instance of such types:

```
type NoFields
end

julia> NoFields() === NoFields()
true
```

The `===` function confirms that the “two” constructed instances of `NoFields` are actually one and the same. Singleton types are described in further detail [below](#) (page 102).

There is much more to say about how instances of composite types are created, but that discussion depends on both *Parametric Types* (page 96) and on *Methods* (page 109), and is sufficiently important to be addressed in its own section: *Constructors* (page 117).

11.5 Immutable Composite Types

It is also possible to define *immutable* composite types by using the keyword `immutable` instead of `type`:

```
immutable Complex
    real::Float64
    imag::Float64
end
```

Such types behave much like other composite types, except that instances of them cannot be modified. Immutable types have several advantages:

- They are more efficient in some cases. Types like the `Complex` example above can be packed efficiently into arrays, and in some cases the compiler is able to avoid allocating immutable objects entirely.
- It is not possible to violate the invariants provided by the type's constructors.
- Code using immutable objects can be easier to reason about.

An immutable object might contain mutable objects, such as arrays, as fields. Those contained objects will remain mutable; only the fields of the immutable object itself cannot be changed to point to different objects.

A useful way to think about immutable composites is that each instance is associated with specific field values — the field values alone tell you everything about the object. In contrast, a mutable object is like a little container that might hold different values over time, and so is not identified with specific field values. In deciding whether to make a type immutable, ask whether two instances with the same field values would be considered identical, or if they might need to change independently over time. If they would be considered identical, the type should probably be immutable.

To recap, two essential properties define immutability in Julia:

- An object with an immutable type is passed around (both in assignment statements and in function calls) by copying, whereas a mutable type is passed around by reference.
- It is not permitted to modify the fields of a composite immutable type.

It is instructive, particularly for readers whose background is C/C++, to consider why these two properties go hand in hand. If they were separated, i.e., if the fields of objects passed around by copying could be modified, then it would become more difficult to reason about certain instances of generic code. For example, suppose `x` is a function argument of an abstract type, and suppose that the function changes a field: `x.isprocessed = true`. Depending on whether `x` is passed by copying or by reference, this statement may or may not alter the actual argument in the calling routine. Julia sidesteps the possibility of creating functions with unknown effects in this scenario by forbidding modification of fields of objects passed around by copying.

11.6 Declared Types

The three kinds of types discussed in the previous three sections are actually all closely related. They share the same key properties:

- They are explicitly declared.
- They have names.
- They have explicitly declared supertypes.

- They may have parameters.

Because of these shared properties, these types are internally represented as instances of the same concept, `DataType`, which is the type of any of these types:

```
julia> typeof(Real)
DataType

julia> typeof(Int)
DataType
```

A `DataType` may be abstract or concrete. If it is concrete, it has a specified size, storage layout, and (optionally) field names. Thus a bits type is a `DataType` with nonzero size, but no field names. A composite type is a `DataType` that has field names or is empty (zero size).

Every concrete value in the system is an instance of some `DataType`.

11.7 Type Unions

A type union is a special abstract type which includes as objects all instances of any of its argument types, constructed using the special `Union` function:

```
julia> IntOrString = Union{Int, AbstractString}
Union{AbstractString, Int64}

julia> 1 :: IntOrString
1

julia> "Hello!" :: IntOrString
"Hello!"

julia> 1.0 :: IntOrString
ERROR: type: typeassert: expected Union{AbstractString, Int64}, got Float64
```

The compilers for many languages have an internal union construct for reasoning about types; Julia simply exposes it to the programmer.

11.8 Parametric Types

An important and powerful feature of Julia's type system is that it is parametric: types can take parameters, so that type declarations actually introduce a whole family of new types — one for each possible combination of parameter values. There are many languages that support some version of [generic programming](#), wherein data structures and algorithms to manipulate them may be specified without specifying the exact types involved. For example, some form of generic programming exists in ML, Haskell, Ada, Eiffel, C++, Java, C#, F#, and Scala, just to name a few. Some of these languages support true parametric polymorphism (e.g. ML, Haskell, Scala), while others support ad-hoc, template-based styles of generic programming (e.g. C++, Java). With so many different varieties of generic programming and parametric types in various languages, we won't even attempt to compare Julia's parametric types to other languages, but will instead focus on explaining Julia's system in its own right. We will note, however, that because Julia is a dynamically typed language and doesn't need to make all type decisions at compile time, many traditional difficulties encountered in static parametric type systems can be relatively easily handled.

All declared types (the `DataType` variety) can be parameterized, with the same syntax in each case. We will discuss them in the following order: first, parametric composite types, then parametric abstract types, and finally parametric bits types.

11.8.1 Parametric Composite Types

Type parameters are introduced immediately after the type name, surrounded by curly braces:

```
type Point{T}
    x::T
    y::T
end
```

This declaration defines a new parametric type, `Point{T}`, holding two “coordinates” of type `T`. What, one may ask, is `T`? Well, that’s precisely the point of parametric types: it can be any type at all (or a value of any bits type, actually, although here it’s clearly used as a type). `Point{Float64}` is a concrete type equivalent to the type defined by replacing `T` in the definition of `Point` with `Float64` (page 437). Thus, this single declaration actually declares an unlimited number of types: `Point{Float64}`, `Point{AbstractString}`, `Point{Int64}`, etc. Each of these is now a usable concrete type:

```
julia> Point{Float64}
Point{Float64}

julia> Point{AbstractString}
Point{AbstractString}
```

The type `Point{Float64}` is a point whose coordinates are 64-bit floating-point values, while the type `Point{AbstractString}` is a “point” whose “coordinates” are string objects (see [Strings](#) (page 41)). However, `Point` itself is also a valid type object:

```
julia> Point
Point{T}
```

Here the `T` is the dummy type symbol used in the original declaration of `Point`. What does `Point` by itself mean? It is a type that contains all the specific instances `Point{Float64}`, `Point{AbstractString}`, etc.:

```
julia> Point{Float64} <: Point
true

julia> Point{AbstractString} <: Point
true
```

Other types, of course, are not subtypes of it:

```
julia> Float64 <: Point
false

julia> AbstractString <: Point
false
```

Concrete `Point` types with different values of `T` are never subtypes of each other:

```
julia> Point{Float64} <: Point{Int64}
false

julia> Point{Float64} <: Point{Real}
false
```

This last point is very important:

- **Even though** `Float64 <: Real` **we DO NOT have** `Point{Float64} <: Point{Real}`.

In other words, in the parlance of type theory, Julia’s type parameters are *invariant*, rather than being covariant (or even contravariant). This is for practical reasons: while any instance of `Point{Float64}` may conceptually be like an instance of `Point{Real}` as well, the two types have different representations in memory:

- An instance of `Point{Float64}` can be represented compactly and efficiently as an immediate pair of 64-bit values;
- An instance of `Point{Real}` must be able to hold any pair of instances of `Real`. Since objects that are instances of `Real` can be of arbitrary size and structure, in practice an instance of `Point{Real}` must be represented as a pair of pointers to individually allocated `Real` objects.

The efficiency gained by being able to store `Point{Float64}` objects with immediate values is magnified enormously in the case of arrays: an `Array{Float64}` can be stored as a contiguous memory block of 64-bit floating-point values, whereas an `Array{Real}` must be an array of pointers to individually allocated `Real` objects — which may well be boxed 64-bit floating-point values, but also might be arbitrarily large, complex objects, which are declared to be implementations of the `Real` abstract type.

Since `Point{Float64}` is not a subtype of `Point{Real}`, the following method can't be applied to arguments of type `Point{Float64}`:

```
function norm(p::Point{Real})
    sqrt(p.x^2 + p.y^2)
end
```

The correct way to define a method that accepts all arguments of type `Point{T}` where `T` is a subtype of `Real` is:

```
function norm{T<:Real}(p::Point{T})
    sqrt(p.x^2 + p.y^2)
end
```

More examples will be discussed later in [Methods](#) (page 109).

How does one construct a `Point` object? It is possible to define custom constructors for composite types, which will be discussed in detail in [Constructors](#) (page 117), but in the absence of any special constructor declarations, there are two default ways of creating new composite objects, one in which the type parameters are explicitly given and the other in which they are implied by the arguments to the object constructor.

Since the type `Point{Float64}` is a concrete type equivalent to `Point` declared with `Float64` (page 437) in place of `T`, it can be applied as a constructor accordingly:

```
julia> Point{Float64}(1.0,2.0)
Point{Float64}(1.0,2.0)

julia> typeof(ans)
Point{Float64}
```

For the default constructor, exactly one argument must be supplied for each field:

```
julia> Point{Float64}(1.0)
ERROR: MethodError: Cannot `convert` an object of type Float64 to an object of type Point{Float64}
This may have arisen from a call to the constructor Point{Float64}(...),
since type constructors fall back to convert methods.
in Point{Float64}(::Float64) at ./sysimg.jl:66
...

julia> Point{Float64}(1.0,2.0,3.0)
ERROR: MethodError: no method matching Point{Float64}(::Float64, ::Float64, ::Float64)
Closest candidates are:
  Point{Float64}{T}(::Any, ::Any) at none:3
  Point{Float64}{T}(::Any) at sysimg.jl:66
...
```

Only one default constructor is generated for parametric types, since overriding it is not possible. This constructor accepts any arguments and converts them to the field types.

In many cases, it is redundant to provide the type of `Point` object one wants to construct, since the types of arguments to the constructor call already implicitly provide type information. For that reason, you can also apply `Point` itself as a constructor, provided that the implied value of the parameter type `T` is unambiguous:

```
julia> Point(1.0,2.0)
Point{Float64}(1.0,2.0)

julia> typeof(ans)
Point{Float64}

julia> Point(1,2)
Point{Int64}(1,2)

julia> typeof(ans)
Point{Int64}
```

In the case of `Point`, the type of `T` is unambiguously implied if and only if the two arguments to `Point` have the same type. When this isn't the case, the constructor will fail with a [MethodError](#) (page 375):

```
julia> Point(1,2.5)
ERROR: MethodError: no method matching Point{T} (::Int64, ::Float64)
...
```

Constructor methods to appropriately handle such mixed cases can be defined, but that will not be discussed until later on in [Constructors](#) (page 117).

11.8.2 Parametric Abstract Types

Parametric abstract type declarations declare a collection of abstract types, in much the same way:

```
abstract Pointy{T}
```

With this declaration, `Pointy{T}` is a distinct abstract type for each type or integer value of `T`. As with parametric composite types, each such instance is a subtype of `Pointy`:

```
julia> Pointy{Int64} <: Pointy
true

julia> Pointy{1} <: Pointy
true
```

Parametric abstract types are invariant, much as parametric composite types are:

```
julia> Pointy{Float64} <: Pointy{Real}
false

julia> Pointy{Real} <: Pointy{Float64}
false
```

Much as plain old abstract types serve to create a useful hierarchy of types over concrete types, parametric abstract types serve the same purpose with respect to parametric composite types. We could, for example, have declared `Point{T}` to be a subtype of `Pointy{T}` as follows:

```
type Point{T} <: Pointy{T}
    x::T
    y::T
end
```

Given such a declaration, for each choice of `T`, we have `Point{T}` as a subtype of `Pointy{T}`:

```
julia> Point{Float64} <: Pointy{Float64}
true

julia> Point{Real} <: Pointy{Real}
true

julia> Point{AbstractString} <: Pointy{AbstractString}
true
```

This relationship is also invariant:

```
julia> Point{Float64} <: Pointy{Real}
false
```

What purpose do parametric abstract types like `Pointy` serve? Consider if we create a point-like implementation that only requires a single coordinate because the point is on the diagonal line $x = y$:

```
type DiagPoint{T} <: Pointy{T}
    x::T
end
```

Now both `Point{Float64}` and `DiagPoint{Float64}` are implementations of the `Pointy{Float64}` abstraction, and similarly for every other possible choice of type `T`. This allows programming to a common interface shared by all `Pointy` objects, implemented for both `Point` and `DiagPoint`. This cannot be fully demonstrated, however, until we have introduced methods and dispatch in the next section, [Methods](#) (page 109).

There are situations where it may not make sense for type parameters to range freely over all possible types. In such situations, one can constrain the range of `T` like so:

```
abstract Pointy{T<:Real}
```

With such a declaration, it is acceptable to use any type that is a subtype of `Real` in place of `T`, but not types that are not subtypes of `Real`:

```
julia> Pointy{Float64}
Pointy{Float64}

julia> Pointy{Real}
Pointy{Real}

julia> Pointy{AbstractString}
ERROR: TypeError: Pointy: in T, expected T<:Real, got Type{AbstractString}
...

julia> Pointy{1}
ERROR: TypeError: Pointy: in T, expected T<:Real, got Int64
...
```

Type parameters for parametric composite types can be restricted in the same manner:

```
type Point{T<:Real} <: Pointy{T}
    x::T
    y::T
end
```

To give a real-world example of how all this parametric type machinery can be useful, here is the actual definition of Julia's `Rational` immutable type (except that we omit the constructor here for simplicity), representing an exact ratio of integers:

```
immutable Rational{T<:Integer} <: Real
    num::T
    den::T
end
```

It only makes sense to take ratios of integer values, so the parameter type `T` is restricted to being a subtype of `Integer`, and a ratio of integers represents a value on the real number line, so any `Rational` is an instance of the `Real` abstraction.

11.8.3 Tuple Types

Tuples are an abstraction of the arguments of a function — without the function itself. The salient aspects of a function's arguments are their order and their types. Therefore a tuple type is similar to a parameterized immutable type where each parameter is the type of one field. For example, a 2-element tuple type resembles the following immutable type:

```
immutable Tuple2{A,B}
    a::A
    b::B
end
```

However, there are three key differences:

- Tuple types may have any number of parameters.
- Tuple types are *covariant* in their parameters: `Tuple{Int}` is a subtype of `Tuple{Any}`. Therefore `Tuple{Any}` is considered an abstract type, and tuple types are only concrete if their parameters are.
- Tuples do not have field names; fields are only accessed by index.

Tuple values are written with parentheses and commas. When a tuple is constructed, an appropriate tuple type is generated on demand:

```
julia> typeof((1, "foo", 2.5))
Tuple{Int64, String, Float64}
```

Note the implications of covariance:

```
julia> Tuple{Int, AbstractString} <: Tuple{Real, Any}
true

julia> Tuple{Int, AbstractString} <: Tuple{Real, Real}
false

julia> Tuple{Int, AbstractString} <: Tuple{Real, }
false
```

Intuitively, this corresponds to the type of a function's arguments being a subtype of the function's signature (when the signature matches).

11.8.4 Vararg Tuple Types

The last parameter of a tuple type can be the special type `Vararg`, which denotes any number of trailing elements:

```
julia> isa(("1",), Tuple{AbstractString, Vararg{Int}})
true

julia> isa(("1", 1), Tuple{AbstractString, Vararg{Int}})
```

```
true

julia> isa(("1",1,2), Tuple{AbstractString,Vararg{Int}})
true

julia> isa(("1",1,2,3.0), Tuple{AbstractString,Vararg{Int}})
false
```

Notice that `Vararg{T}` corresponds to zero or more elements of type `T`. `Vararg` tuple types are used to represent the arguments accepted by `varargs` methods (see [Varargs Functions](#) (page 59)).

The type `Vararg{T,N}` corresponds to exactly `N` elements of type `T`. `NTuple{N,T}` is a convenient alias for `Tuple{Vararg{T,N}}`, i.e. a tuple type containing exactly `N` elements of type `T`.

Singleton Types

There is a special kind of abstract parametric type that must be mentioned here: singleton types. For each type, `T`, the “singleton type” `Type{T}` is an abstract type whose only instance is the object `T`. Since the definition is a little difficult to parse, let’s look at some examples:

```
julia> isa(Float64, Type{Float64})
true

julia> isa(Real, Type{Float64})
false

julia> isa(Real, Type{Real})
true

julia> isa(Float64, Type{Real})
false
```

In other words, `isa(A, Type{B})` (page 362) is true if and only if `A` and `B` are the same object and that object is a type. Without the parameter, `Type` is simply an abstract type which has all type objects as its instances, including, of course, singleton types:

```
julia> isa(Type{Float64}, Type)
true

julia> isa(Float64, Type)
true

julia> isa(Real, Type)
true
```

Any object that is not a type is not an instance of `Type`:

```
julia> isa(1, Type)
false

julia> isa("foo", Type)
false
```

Until we discuss [Parametric Methods](#) (page 113) and [conversions](#) (page 128), it is difficult to explain the utility of the singleton type construct, but in short, it allows one to specialize function behavior on specific type *values*. This is useful for writing methods (especially parametric ones) whose behavior depends on a type that is given as an explicit argument rather than implied by the type of one of its arguments.

A few popular languages have singleton types, including Haskell, Scala and Ruby. In general usage, the term “singleton type” refers to a type whose only instance is a single value. This meaning applies to Julia’s singleton types, but with that caveat that only type objects have singleton types.

11.8.5 Parametric Bits Types

Bits types can also be declared parametrically. For example, pointers are represented as boxed bits types which would be declared in Julia like this:

```
# 32-bit system:
bitstype 32 Ptr{T}

# 64-bit system:
bitstype 64 Ptr{T}
```

The slightly odd feature of these declarations as compared to typical parametric composite types, is that the type parameter `T` is not used in the definition of the type itself — it is just an abstract tag, essentially defining an entire family of types with identical structure, differentiated only by their type parameter. Thus, `Ptr{Float64}` and `Ptr{Int64}` are distinct types, even though they have identical representations. And of course, all specific pointer types are subtype of the umbrella `Ptr` type:

```
julia> Ptr{Float64} <: Ptr
true

julia> Ptr{Int64} <: Ptr
true
```

11.9 Type Aliases

Sometimes it is convenient to introduce a new name for an already expressible type. For such occasions, Julia provides the `typealias` mechanism. For example, `UInt` is type aliased to either `UInt32` or `UInt64` as is appropriate for the size of pointers on the system:

```
# 32-bit system:
julia> UInt
UInt32

# 64-bit system:
julia> UInt
UInt64
```

This is accomplished via the following code in `base/boot.jl`:

```
if is(Int, Int64)
    typealias UInt UInt64
else
    typealias UInt UInt32
end
```

Of course, this depends on what `Int` is aliased to — but that is predefined to be the correct type — either `Int32` or `Int64`.

For parametric types, `typealias` can be convenient for providing names for cases where some of the parameter choices are fixed. Julia’s arrays have type `Array{T,N}` where `T` is the element type and `N` is the number of array dimensions. For convenience, writing `Array{Float64}` allows one to specify the element type without specifying the dimension:

```
julia> Array{Float64,1} <: Array{Float64} <: Array{Float64}
true
```

However, there is no way to equally simply restrict just the dimension but not the element type. Yet, one often needs to ensure an object is a vector or a matrix (imposing restrictions on the number of dimensions). For that reason, the following type aliases are provided:

```
typealias Vector{T} Array{T,1}
typealias Matrix{T} Array{T,2}
```

Writing `Vector{Float64}` is equivalent to writing `Array{Float64,1}`, and the umbrella type `Vector` has as instances all `Array` objects where the second parameter — the number of array dimensions — is 1, regardless of what the element type is. In languages where parametric types must always be specified in full, this is not especially helpful, but in Julia, this allows one to write just `Matrix` for the abstract type including all two-dimensional dense arrays of any element type.

This declaration of `Vector` creates a subtype relation `Vector{Int} <: Vector`. However, it is not always the case that a parametric `typealias` statement creates such a relation; for example, the statement:

```
typealias AA{T} Array{Array{T,1},1}
```

does not create the relation `AA{Int} <: AA`. The reason is that `Array{Array{T,1},1}` is not an abstract type at all; in fact, it is a concrete type describing a 1-dimensional array in which each entry is an object of type `Array{T,1}` for some value of `T`.

11.10 Operations on Types

Since types in Julia are themselves objects, ordinary functions can operate on them. Some functions that are particularly useful for working with or exploring types have already been introduced, such as the `<:` operator, which indicates whether its left hand operand is a subtype of its right hand operand.

The `isa` function tests if an object is of a given type and returns true or false:

```
julia> isa(1,Int)
true

julia> isa(1,AbstractFloat)
false
```

The `typeof()` (page 363) function, already used throughout the manual in examples, returns the type of its argument. Since, as noted above, types are objects, they also have types, and we can ask what their types are:

```
julia> typeof(Rational)
DataType

julia> typeof(Union{Real,Float64,Rational})
DataType

julia> typeof(Union{Real,String})
Union
```

What if we repeat the process? What is the type of a type of a type? As it happens, types are all composite values and thus all have a type of `DataType`:

```
julia> typeof(DataType)
DataType
```

```
julia> typeof(Union)
DataType
```

DataType is its own type.

Another operation that applies to some types is `supertype()` (page 365), which reveals a type's supertype. Only declared types (DataType) have unambiguous supertypes:

```
julia> supertype(Float64)
AbstractFloat

julia> supertype(Number)
Any

julia> supertype(AbstractString)
Any

julia> supertype(Any)
Any
```

If you apply `supertype()` (page 365) to other type objects (or non-type objects), a `MethodError` (page 375) is raised:

```
julia> supertype(Union{Float64, Int64})
ERROR: `supertype` has no method matching supertype(::Type{Union{Float64, Int64}})
```

11.11 Custom pretty-printing

Often, one wants to customize how instances of a type are displayed. This is accomplished by overloading the `show()` (page 566) function. For example, suppose we define a type to represent complex numbers in polar form:

```
type Polar{T<:Real} <: Number
    r::T
    θ::T
end
Polar(r::Real, θ::Real) = Polar(promote(r, θ)...)

```

Here, we've added a custom constructor function so that it can take arguments of different Real types and promote them to a common type (see *Constructors* (page 117) and *Conversion and Promotion* (page 127)). (Of course, we would have to define lots of other methods, too, to make it act like a Number, e.g. +, *, one, zero, promotion rules and so on.) By default, instances of this type display rather simply, with information about the type name and the field values, as e.g. `Polar{Float64}(3.0, 4.0)`.

If we want it to display instead as `3.0 * exp(4.0im)`, we would define the following method to print the object to a given output object `io` (representing a file, terminal, buffer, etcetera; see *Networking and Streams* (page 193)):

```
Base.show(io::IO, z::Polar) = print(io, z.r, " * exp(", z.θ, "im")
```

More fine-grained control over display of Polar objects is possible. In particular, sometimes one wants both a verbose multi-line printing format, used for displaying a single object in the REPL and other interactive environments, and also a more compact single-line format used for `print()` (page 566) or for displaying the object as part of another object (e.g. in an array). Although by default the `show(io, z)` function is called in both cases, you can define a *different* multi-line format for displaying an object by overloading a three-argument form of `show` that takes the `text/plain` MIME type as its second argument (see *Multimedia I/O* (page 569)), for example:

```
Base.show{T}(io::IO, ::MIME"text/plain", z::Polar{T}) =
    print(io, "Polar{$T} complex number:\n    ", z)
```

(Note that `print(..., z)` here will call the 2-argument `show(io, z)` method.) This results in:

```
julia> Polar(3, 4.0)
Polar{Float64} complex number:
  3.0 * exp(4.0im)

julia> [Polar(3, 4.0), Polar(4.0, 5.3)]
2-element Array{Polar{Float64},1}:
 3.0 * exp(4.0im)
 4.0 * exp(5.3im)
```

where the single-line `show(io, z)` form is still used for an array of `Polar` values. Technically, the REPL calls `display(z)` to display the result of executing a line, which defaults to `show(STDOUT, MIME("text/plain"), z)`, which in turn defaults to `show(STDOUT, z)`, but you should *not* define new `display()` (page 569) methods unless you are defining a new multimedia display handler (see [Multimedia I/O](#) (page 569)).

Moreover, you can also define `show` methods for other MIME types in order to enable richer display (HTML, images, etcetera) of objects in environments that support this (e.g. JJulia). For example, we can define formatted HTML display of `Polar` objects, with superscripts and italics, via:

```
Base.show{T}(io::IO, ::MIME"text/html", z::Polar{T}) =
    println(io, "<code>Polar{$T}</code> complex number: ",
            z.r, " <i>e</i><sup>4.0</sup>", z.θ, " <i>i</i><sup></sup>")
```

A `Polar` object will then display automatically using HTML in an environment that supports HTML display, but you can call `show` manually to get HTML output if you want:

```
julia> show(STDOUT, "text/html", Polar(3.0, 4.0))
<code>Polar{Float64}</code> complex number: 3.0 <i>e</i><sup>4.0</sup> <i>i</i><sup></sup>
```

11.12 “Value types”

In Julia, you can’t dispatch on a *value* such as `true` or `false`. However, you can dispatch on parametric types, and Julia allows you to include “plain bits” values (Types, Symbols, Integers, floating-point numbers, tuples, etc.) as type parameters. A common example is the dimensionality parameter in `Array{T, N}`, where `T` is a type (e.g., `Float64`) but `N` is just an `Int`.

You can create your own custom types that take values as parameters, and use them to control dispatch of custom types. By way of illustration of this idea, let’s introduce a parametric type, `Val{T}`, which serves as a customary way to exploit this technique for cases where you don’t need a more elaborate hierarchy.

`Val` is defined as:

```
immutable Val{T}
end
```

There is no more to the implementation of `Val` than this. Some functions in Julia’s standard library accept `Val` types as arguments, and you can also use it to write your own functions. For example:

```
firstlast(::Type{Val{true}}) = "First"
firstlast(::Type{Val{false}}) = "Last"

julia> firstlast(Val{true})
```



```
"First"

julia> firstlast(Val{false})

"Last"
```

For consistency across Julia, the call site should always pass a `Val` type rather than creating an *instance*, i.e., use `foo(Val{:bar})` rather than `foo(Val{:bar}())`.

It's worth noting that it's extremely easy to mis-use parametric “value” types, including `Val`; in unfavorable cases, you can easily end up making the performance of your code much *worse*. In particular, you would never want to write actual code as illustrated above. For more information about the proper (and improper) uses of `Val`, please read the more extensive discussion in [the performance tips](#) (page 314).

11.13 Nullable Types: Representing Missing Values

In many settings, you need to interact with a value of type `T` that may or may not exist. To handle these settings, Julia provides a parametric type called `Nullable{T}`, which can be thought of as a specialized container type that can contain either zero or one values. `Nullable{T}` provides a minimal interface designed to ensure that interactions with missing values are safe. At present, the interface consists of four possible interactions:

- Construct a [Nullable](#) (page 369) object.
- Check if a [Nullable](#) (page 369) object has a missing value.
- Access the value of a [Nullable](#) (page 369) object with a guarantee that a [NullException](#) (page 375) will be thrown if the object's value is missing.
- Access the value of a [Nullable](#) (page 369) object with a guarantee that a default value of type `T` will be returned if the object's value is missing.

11.13.1 Constructing Nullable objects

To construct an object representing a missing value of type `T`, use the `Nullable{T}()` function:

```
julia> x1 = Nullable{Int64}()
Nullable{Int64}()

julia> x2 = Nullable{Float64}()
Nullable{Float64}()

julia> x3 = Nullable{Vector{Int64}}()
Nullable{Array{Int64,1}}()
```

To construct an object representing a non-missing value of type `T`, use the `Nullable(x::T)` function:

```
julia> x1 = Nullable(1)
Nullable{Int64}(1)

julia> x2 = Nullable(1.0)
Nullable{Float64}(1.0)

julia> x3 = Nullable([1, 2, 3])
Nullable{Array{Int64,1}}([1, 2, 3])
```

Note the core distinction between these two ways of constructing a [Nullable](#) (page 369) object: in one style, you provide a type, `T`, as a function parameter; in the other style, you provide a single value of type `T` as an argument.

11.13.2 Checking if a `Nullable` object has a value

You can check if a `Nullable` (page 369) object has any value using `isnull()` (page 369):

```
julia> isnull(Nullable{Float64}())
true

julia> isnull(Nullable(0.0))
false
```

11.13.3 Safely accessing the value of a `Nullable` object

You can safely access the value of a `Nullable` (page 369) object using `get()` (page 369):

```
julia> get(Nullable{Float64}())
ERROR:NullException()
  in get(::Nullable{Float64}) at ./nullable.jl:91
...

julia> get(Nullable(1.0))
1.0
```

If the value is not present, as it would be for `Nullable{Float64}`, a `NullException` (page 375) error will be thrown. The error-throwing nature of the `get()` (page 369) function ensures that any attempt to access a missing value immediately fails.

In cases for which a reasonable default value exists that could be used when a `Nullable` (page 369) object's value turns out to be missing, you can provide this default value as a second argument to `get()` (page 369):

```
julia> get(Nullable{Float64}(), 0.0)
0.0

julia> get(Nullable(1.0), 0.0)
1.0
```

Tip: Make sure the type of the default value passed to `get()` (page 369) and that of the `Nullable` (page 369) object match to avoid type instability, which could hurt performance. Use `convert()` (page 364) manually if needed.

Methods

Recall from *Functions* (page 55) that a function is an object that maps a tuple of arguments to a return value, or throws an exception if no appropriate value can be returned. It is common for the same conceptual function or operation to be implemented quite differently for different types of arguments: adding two integers is very different from adding two floating-point numbers, both of which are distinct from adding an integer to a floating-point number. Despite their implementation differences, these operations all fall under the general concept of “addition”. Accordingly, in Julia, these behaviors all belong to a single object: the `+` function.

To facilitate using many different implementations of the same concept smoothly, functions need not be defined all at once, but can rather be defined piecewise by providing specific behaviors for certain combinations of argument types and counts. A definition of one possible behavior for a function is called a *method*. Thus far, we have presented only examples of functions defined with a single method, applicable to all types of arguments. However, the signatures of method definitions can be annotated to indicate the types of arguments in addition to their number, and more than a single method definition may be provided. When a function is applied to a particular tuple of arguments, the most specific method applicable to those arguments is applied. Thus, the overall behavior of a function is a patchwork of the behaviors of its various method definitions. If the patchwork is well designed, even though the implementations of the methods may be quite different, the outward behavior of the function will appear seamless and consistent.

The choice of which method to execute when a function is applied is called *dispatch*. Julia allows the dispatch process to choose which of a function’s methods to call based on the number of arguments given, and on the types of all of the function’s arguments. This is different than traditional object-oriented languages, where dispatch occurs based only on the first argument, which often has a special argument syntax, and is sometimes implied rather than explicitly written as an argument.¹ Using all of a function’s arguments to choose which method should be invoked, rather than just the first, is known as *multiple dispatch*. Multiple dispatch is particularly useful for mathematical code, where it makes little sense to artificially deem the operations to “belong” to one argument more than any of the others: does the addition operation in $x + y$ belong to x any more than it does to y ? The implementation of a mathematical operator generally depends on the types of all of its arguments. Even beyond mathematical operations, however, multiple dispatch ends up being a powerful and convenient paradigm for structuring and organizing programs.

12.1 Defining Methods

Until now, we have, in our examples, defined only functions with a single method having unconstrained argument types. Such functions behave just like they would in traditional dynamically typed languages. Nevertheless, we have used multiple dispatch and methods almost continually without being aware of it: all of Julia’s standard functions and operators, like the aforementioned `+` function, have many methods defining their behavior over various possible combinations of argument type and count.

¹ In C++ or Java, for example, in a method call like `obj.meth(arg1, arg2)`, the object `obj` “receives” the method call and is implicitly passed to the method via the `this` keyword, rather than as an explicit method argument. When the current `this` object is the receiver of a method call, it can be omitted altogether, writing just `meth(arg1, arg2)`, with `this` implied as the receiving object.

When defining a function, one can optionally constrain the types of parameters it is applicable to, using the `::` type-assertion operator, introduced in the section on [Composite Types](#) (page 93):

```
julia> f(x::Float64, y::Float64) = 2x + y;
```

This function definition applies only to calls where `x` and `y` are both values of type `Float64`:

```
julia> f(2.0, 3.0)
7.0
```

Applying it to any other types of arguments will result in a `MethodError`:

```
julia> f(2.0, 3)
ERROR: MethodError: no method matching f(::Float64, ::Int64)
Closest candidates are:
  f(::Float64, !Matched::Float64) at none:1
...

julia> f(Float32(2.0), 3.0)
ERROR: MethodError: no method matching f(::Float32, ::Float64)
Closest candidates are:
  f(!Matched::Float64, ::Float64) at none:1
...

julia> f(2.0, "3.0")
ERROR: MethodError: no method matching f(::Float64, ::String)
Closest candidates are:
  f(::Float64, !Matched::Float64) at none:1
...

julia> f("2.0", "3.0")
ERROR: MethodError: no method matching f(::String, ::String)
...
```

As you can see, the arguments must be precisely of type `Float64`. Other numeric types, such as integers or 32-bit floating-point values, are not automatically converted to 64-bit floating-point, nor are strings parsed as numbers. Because `Float64` is a concrete type and concrete types cannot be subclassed in Julia, such a definition can only be applied to arguments that are exactly of type `Float64`. It may often be useful, however, to write more general methods where the declared parameter types are abstract:

```
julia> f(x::Number, y::Number) = 2x - y;

julia> f(2.0, 3)
1.0
```

This method definition applies to any pair of arguments that are instances of `Number`. They need not be of the same type, so long as they are each numeric values. The problem of handling disparate numeric types is delegated to the arithmetic operations in the expression `2x - y`.

To define a function with multiple methods, one simply defines the function multiple times, with different numbers and types of arguments. The first method definition for a function creates the function object, and subsequent method definitions add new methods to the existing function object. The most specific method definition matching the number and types of the arguments will be executed when the function is applied. Thus, the two method definitions above, taken together, define the behavior for `f` over all pairs of instances of the abstract type `Number` — but with a different behavior specific to pairs of `Float64` values. If one of the arguments is a 64-bit float but the other one is not, then the `f(Float64, Float64)` method cannot be called and the more general `f(Number, Number)` method must be used:

```
julia> f(2.0, 3.0)
7.0

julia> f(2, 3.0)
1.0

julia> f(2.0, 3)
1.0

julia> f(2, 3)
1
```

The $2x + y$ definition is only used in the first case, while the $2x - y$ definition is used in the others. No automatic casting or conversion of function arguments is ever performed: all conversion in Julia is non-magical and completely explicit. [Conversion and Promotion](#) (page 127), however, shows how clever application of sufficiently advanced technology can be indistinguishable from magic. [\[Clarke61\]](#) (page 621)

For non-numeric values, and for fewer or more than two arguments, the function `f` remains undefined, and applying it will still result in a `MethodError`:

```
julia> f("foo", 3)
ERROR: MethodError: no method matching f(::String, ::Int64)
Closest candidates are:
  f(!Matched::Number, ::Number) at none:1
...

julia> f()
ERROR: MethodError: no method matching f()
Closest candidates are:
  f(!Matched::Float64, !Matched::Float64) at none:1
  f(!Matched::Number, !Matched::Number) at none:1
...
```

You can easily see which methods exist for a function by entering the function object itself in an interactive session:

```
julia> f
f (generic function with 2 methods)
```

This output tells us that `f` is a function object with two methods. To find out what the signatures of those methods are, use the `methods()` function:

```
julia> methods(f)
# 2 methods for generic function "f":
f(x::Float64, y::Float64) in Main at none:1
f(x::Number, y::Number) in Main at none:1
```

which shows that `f` has two methods, one taking two `Float64` arguments and one taking arguments of type `Number`. It also indicates the file and line number where the methods were defined: because these methods were defined at the REPL, we get the apparent line number `none:1`.

In the absence of a type declaration with `::`, the type of a method parameter is `Any` by default, meaning that it is unconstrained since all values in Julia are instances of the abstract type `Any`. Thus, we can define a catch-all method for `f` like so:

```
julia> f(x,y) = println("Whoa there, Nelly.");

julia> f("foo", 1)
Whoa there, Nelly.
```

This catch-all is less specific than any other possible method definition for a pair of parameter values, so it is only be called on pairs of arguments to which no other method definition applies.

Although it seems a simple concept, multiple dispatch on the types of values is perhaps the single most powerful and central feature of the Julia language. Core operations typically have dozens of methods:

```
julia> methods(+)
# 166 methods for generic function "+":
+(a::Float16, b::Float16) at float16.jl:136
+(x::Float32, y::Float32) at float.jl:206
+(x::Float64, y::Float64) at float.jl:207
+(x::Bool, z::Complex{Bool}) at complex.jl:126
+(x::Bool, y::Bool) at bool.jl:48
+(x::Bool) at bool.jl:45
+{T<:AbstractFloat}(x::Bool, y::T) at bool.jl:55
+(x::Bool, z::Complex) at complex.jl:133
+(x::Bool, A::AbstractArray{Bool,N<:Any}) at arraymath.jl:105
+(x::Char, y::Integer) at char.jl:40
+{T<:Union{Int128,Int16,Int32,Int64,Int8,UInt128,UInt16,UInt32,UInt64,UInt8}}(x::T, y::T) at int.jl:115
+(z::Complex, w::Complex) at complex.jl:115
+(z::Complex, x::Bool) at complex.jl:134
+(x::Real, z::Complex{Bool}) at complex.jl:140
+(x::Real, z::Complex) at complex.jl:152
+(z::Complex, x::Real) at complex.jl:153
+(x::Rational, y::Rational) at rational.jl:179
...
+(a, b, c, xs...) at operators.jl:119
```

Multiple dispatch together with the flexible parametric type system give Julia its ability to abstractly express high-level algorithms decoupled from implementation details, yet generate efficient, specialized code to handle each case at run time.

12.2 Method Ambiguities

It is possible to define a set of function methods such that there is no unique most specific method applicable to some combinations of arguments:

```
julia> g(x::Float64, y) = 2x + y;

julia> g(x, y::Float64) = x + 2y;

julia> g(2.0, 3)
7.0

julia> g(2, 3.0)
8.0

julia> g(2.0, 3.0)
ERROR: MethodError: g(::Float64, ::Float64) is ambiguous. Candidates:
  g(x, y::Float64) in Main at none:1
  g(x::Float64, y) in Main at none:1
...
```

Here the call `g(2.0, 3.0)` could be handled by either the `g(Float64, Any)` or the `g(Any, Float64)` method, and neither is more specific than the other. In such cases, Julia raises a `MethodError` rather than arbitrarily picking a method. You can avoid method ambiguities by specifying an appropriate method for the intersection case:

```
julia> g(x::Float64, y::Float64) = 2x + 2y;

julia> g(x::Float64, y) = 2x + y;

julia> g(x, y::Float64) = x + 2y;

julia> g(2.0, 3)
7.0

julia> g(2, 3.0)
8.0

julia> g(2.0, 3.0)
10.0
```

It is recommended that the disambiguating method be defined first, since otherwise the ambiguity exists, if transiently, until the more specific method is defined.

12.3 Parametric Methods

Method definitions can optionally have type parameters immediately after the method name and before the parameter tuple:

```
julia> same_type{T}(x::T, y::T) = true;

julia> same_type(x,y) = false;
```

The first method applies whenever both arguments are of the same concrete type, regardless of what type that is, while the second method acts as a catch-all, covering all other cases. Thus, overall, this defines a boolean function that checks whether its two arguments are of the same type:

```
julia> same_type(1, 2)
true

julia> same_type(1, 2.0)
false

julia> same_type(1.0, 2.0)
true

julia> same_type("foo", 2.0)
false

julia> same_type("foo", "bar")
true

julia> same_type{Int32}(1), Int64(2))
false
```

This kind of definition of function behavior by dispatch is quite common — idiomatic, even — in Julia. Method type parameters are not restricted to being used as the types of parameters: they can be used anywhere a value would be in the signature of the function or body of the function. Here's an example where the method type parameter `T` is used as the type parameter to the parametric type `Vector{T}` in the method signature:

```
julia> myappend{T}(v::Vector{T}, x::T) = [v..., x]
myappend (generic function with 1 method)
```

```

julia> myappend([1,2,3],4)
4-element Array{Int64,1}:
 1
 2
 3
 4

julia> myappend([1,2,3],2.5)
ERROR: MethodError: no method matching myappend(::Array{Int64,1}, ::Float64)
Closest candidates are:
  myappend{T} (::Array{T,1}, !Matched::T) at none:1
...

julia> myappend([1.0,2.0,3.0],4.0)
4-element Array{Float64,1}:
 1.0
 2.0
 3.0
 4.0

julia> myappend([1.0,2.0,3.0],4)
ERROR: MethodError: no method matching myappend(::Array{Float64,1}, ::Int64)
Closest candidates are:
  myappend{T} (::Array{T,1}, !Matched::T) at none:1
...

```

As you can see, the type of the appended element must match the element type of the vector it is appended to, or else a `MethodError` is raised. In the following example, the method type parameter `T` is used as the return value:

```

julia> mytypeof{T}(x::T) = T
mytypeof (generic function with 1 method)

julia> mytypeof(1)
Int64

julia> mytypeof(1.0)
Float64

```

Just as you can put subtype constraints on type parameters in type declarations (see [Parametric Types](#) (page 96)), you can also constrain type parameters of methods:

```

same_type_numeric{T<:Number}(x::T, y::T) = true
same_type_numeric(x::Number, y::Number) = false

julia> same_type_numeric(1, 2)
true

julia> same_type_numeric(1, 2.0)
false

julia> same_type_numeric(1.0, 2.0)
true

julia> same_type_numeric("foo", 2.0)
no method same_type_numeric(String,Float64)

julia> same_type_numeric("foo", "bar")
no method same_type_numeric(String,String)

```



```
julia> same_type_numeric(Int32(1), Int64(2))
false
```

The `same_type_numeric` function behaves much like the `same_type` function defined above, but is only defined for pairs of numbers.

12.4 Parametrically-constrained Varargs methods

Function parameters can also be used to constrain the number of arguments that may be supplied to a “varargs” function (*Varargs Functions* (page 59)). The notation `Vararg{T,N}` is used to indicate such a constraint. For example:

```
julia> bar(a,b,x::Vararg{Any,2}) = (a,b,x);

julia> bar(1,2,3)
ERROR: MethodError: no method matching bar(::Int64, ::Int64, ::Int64)
...

julia> bar(1,2,3,4)
(1,2,(3,4))

julia> bar(1,2,3,4,5)
ERROR: MethodError: no method matching bar(::Int64, ::Int64, ::Int64, ::Int64, ::Int64)
...
```

More usefully, it is possible to constrain varargs methods by a parameter. For example:

```
function getindex{T,N}(A::AbstractArray{T,N}, indexes::Vararg{Number,N})
```

would be called only when the number of `indexes` matches the dimensionality of the array.

12.5 Note on Optional and keyword Arguments

As mentioned briefly in *Functions* (page 55), optional arguments are implemented as syntax for multiple method definitions. For example, this definition:

```
f(a=1,b=2) = a+2b
```

translates to the following three methods:

```
f(a,b) = a+2b
f(a) = f(a,2)
f() = f(1,2)
```

This means that calling `f()` is equivalent to calling `f(1,2)`. In this case the result is 5, because `f(1,2)` invokes the first method of `f` above. However, this need not always be the case. If you define a fourth method that is more specialized for integers:

```
f(a::Int,b::Int) = a-2b
```

then the result of both `f()` and `f(1,2)` is `-3`. In other words, optional arguments are tied to a function, not to any specific method of that function. It depends on the types of the optional arguments which method is invoked. When optional arguments are defined in terms of a global variable, the type of the optional argument may even change at run-time.

Keyword arguments behave quite differently from ordinary positional arguments. In particular, they do not participate in method dispatch. Methods are dispatched based only on positional arguments, with keyword arguments processed after the matching method is identified.

12.6 Function-like objects

Methods are associated with types, so it is possible to make any arbitrary Julia object “callable” by adding methods to its type. (Such “callable” objects are sometimes called “functors.”)

For example, you can define a type that stores the coefficients of a polynomial, but behaves like a function evaluating the polynomial:

```
immutable Polynomial{R}
    coeffs::Vector{R}
end

function (p::Polynomial)(x)
    v = p.coeffs[end]
    for i = (length(p.coeffs)-1):-1:1
        v = v*x + p.coeffs[i]
    end
    return v
end
```

Notice that the function is specified by type instead of by name. In the function body, `p` will refer to the object that was called. A `Polynomial` can be used as follows:

```
julia> p = Polynomial([1,10,100])
Polynomial{Int64}([1,10,100])

julia> p(3)
931
```

This mechanism is also the key to how type constructors and closures (inner functions that refer to their surrounding environment) work in Julia, discussed *later in the manual* (page 124).

12.7 Empty generic functions

Occasionally it is useful to introduce a generic function without yet adding methods. This can be used to separate interface definitions from implementations. It might also be done for the purpose of documentation or code readability. The syntax for this is an empty `function` block without a tuple of arguments:

```
function emptyfunc
end
```

Constructors

Constructors ¹ are functions that create new objects — specifically, instances of *Composite Types* (page 93). In Julia, type objects also serve as constructor functions: they create new instances of themselves when applied to an argument tuple as a function. This much was already mentioned briefly when composite types were introduced. For example:

```
type Foo
    bar
    baz
end

julia> foo = Foo(1,2)
Foo(1,2)

julia> foo.bar
1

julia> foo.baz
2
```

For many types, forming new objects by binding their field values together is all that is ever needed to create instances. There are, however, cases where more functionality is required when creating composite objects. Sometimes invariants must be enforced, either by checking arguments or by transforming them. *Recursive data structures*, especially those that may be self-referential, often cannot be constructed cleanly without first being created in an incomplete state and then altered programmatically to be made whole, as a separate step from object creation. Sometimes, it's just convenient to be able to construct objects with fewer or different types of parameters than they have fields. Julia's system for object construction addresses all of these cases and more.

13.1 Outer Constructor Methods

A constructor is just like any other function in Julia in that its overall behavior is defined by the combined behavior of its methods. Accordingly, you can add functionality to a constructor by simply defining new methods. For example, let's say you want to add a constructor method for `Foo` objects that takes only one argument and uses the given value for both the `bar` and `baz` fields. This is simple:

```
Foo(x) = Foo(x, x)
```

¹ Nomenclature: while the term “constructor” generally refers to the entire function which constructs objects of a type, it is common to abuse terminology slightly and refer to specific constructor methods as “constructors”. In such situations, it is generally clear from context that the term is used to mean “constructor method” rather than “constructor function”, especially as it is often used in the sense of singling out a particular method of the constructor from all of the others.

```
julia> Foo(1)
Foo(1,1)
```

You could also add a zero-argument `Foo` constructor method that supplies default values for both of the `bar` and `baz` fields:

```
Foo() = Foo(0)

julia> Foo()
Foo(0,0)
```

Here the zero-argument constructor method calls the single-argument constructor method, which in turn calls the automatically provided two-argument constructor method. For reasons that will become clear very shortly, additional constructor methods declared as normal methods like this are called *outer* constructor methods. Outer constructor methods can only ever create a new instance by calling another constructor method, such as the automatically provided default ones.

13.2 Inner Constructor Methods

While outer constructor methods succeed in addressing the problem of providing additional convenience methods for constructing objects, they fail to address the other two use cases mentioned in the introduction of this chapter: enforcing invariants, and allowing construction of self-referential objects. For these problems, one needs *inner* constructor methods. An inner constructor method is much like an outer constructor method, with two differences:

1. It is declared inside the block of a type declaration, rather than outside of it like normal methods.
2. It has access to a special locally existent function called `new` that creates objects of the block's type.

For example, suppose one wants to declare a type that holds a pair of real numbers, subject to the constraint that the first number is not greater than the second one. One could declare it like this:

```
type OrderedPair
    x::Real
    y::Real

    OrderedPair(x,y) = x > y ? error("out of order") : new(x,y)
end
```

Now `OrderedPair` objects can only be constructed such that `x <= y`:

```
julia> OrderedPair(1,2)
OrderedPair(1,2)

julia> OrderedPair(2,1)
ERROR: out of order
in OrderedPair(::Int64, ::Int64) at ./none:5
...
```

You can still reach in and directly change the field values to violate this invariant, but messing around with an object's internals uninvited is considered poor form. You (or someone else) can also provide additional outer constructor methods at any later point, but once a type is declared, there is no way to add more inner constructor methods. Since outer constructor methods can only create objects by calling other constructor methods, ultimately, some inner constructor must be called to create an object. This guarantees that all objects of the declared type must come into existence by a call to one of the inner constructor methods provided with the type, thereby giving some degree of enforcement of a type's invariants.

Of course, if the type is declared as `immutable`, then its constructor-provided invariants are fully enforced. This is an important consideration when deciding whether a type should be immutable.

If any inner constructor method is defined, no default constructor method is provided: it is presumed that you have supplied yourself with all the inner constructors you need. The default constructor is equivalent to writing your own inner constructor method that takes all of the object's fields as parameters (constrained to be of the correct type, if the corresponding field has a type), and passes them to `new`, returning the resulting object:

```
type Foo
    bar
    baz

    Foo(bar,baz) = new(bar,baz)
end
```

This declaration has the same effect as the earlier definition of the `Foo` type without an explicit inner constructor method. The following two types are equivalent — one with a default constructor, the other with an explicit constructor:

```
type T1
    x::Int64
end

type T2
    x::Int64
    T2(x) = new(x)
end

julia> T1(1)
T1(1)

julia> T2(1)
T2(1)

julia> T1(1.0)
T1(1)

julia> T2(1.0)
T2(1)
```

It is considered good form to provide as few inner constructor methods as possible: only those taking all arguments explicitly and enforcing essential error checking and transformation. Additional convenience constructor methods, supplying default values or auxiliary transformations, should be provided as outer constructors that call the inner constructors to do the heavy lifting. This separation is typically quite natural.

13.3 Incomplete Initialization

The final problem which has still not been addressed is construction of self-referential objects, or more generally, recursive data structures. Since the fundamental difficulty may not be immediately obvious, let us briefly explain it. Consider the following recursive type declaration:

```
type SelfReferential
    obj::SelfReferential
end
```

This type may appear innocuous enough, until one considers how to construct an instance of it. If `a` is an instance of `SelfReferential`, then a second instance can be created by the call:

```
b = SelfReferential(a)
```

But how does one construct the first instance when no instance exists to provide as a valid value for its `obj` field? The only solution is to allow creating an incompletely initialized instance of `SelfReferential` with an unassigned `obj` field, and using that incomplete instance as a valid value for the `obj` field of another instance, such as, for example, itself.

To allow for the creation of incompletely initialized objects, Julia allows the `new` function to be called with fewer than the number of fields that the type has, returning an object with the unspecified fields uninitialized. The inner constructor method can then use the incomplete object, finishing its initialization before returning it. Here, for example, we take another crack at defining the `SelfReferential` type, with a zero-argument inner constructor returning instances having `obj` fields pointing to themselves:

```
type SelfReferential
    obj::SelfReferential

    SelfReferential() = (x = new(); x.obj = x)
end
```

We can verify that this constructor works and constructs objects that are, in fact, self-referential:

```
julia> x = SelfReferential();

julia> x === x
true

julia> x === x.obj
true

julia> x === x.obj.obj
true
```

Although it is generally a good idea to return a fully initialized object from an inner constructor, incompletely initialized objects can be returned:

```
julia> type Incomplete
    xx
    Incomplete() = new()
end

julia> z = Incomplete();
```

While you are allowed to create objects with uninitialized fields, any access to an uninitialized reference is an immediate error:

```
julia> z.xx
ERROR: UndefRefError: access to undefined reference
...
```

This avoids the need to continually check for `null` values. However, not all object fields are references. Julia considers some types to be “plain data”, meaning all of their data is self-contained and does not reference other objects. The plain data types consist of bits types (e.g. `Int`) and immutable structs of other plain data types. The initial contents of a plain data type is undefined:

```
julia> type HasPlain
    n::Int
    HasPlain() = new()
end

julia> HasPlain()
HasPlain(438103441441)
```

Arrays of plain data types exhibit the same behavior.

You can pass incomplete objects to other functions from inner constructors to delegate their completion:

```
type Lazy
    xx

    Lazy(v) = complete_me(new(), v)
end
```

As with incomplete objects returned from constructors, if `complete_me` or any of its callees try to access the `xx` field of the `Lazy` object before it has been initialized, an error will be thrown immediately.

13.4 Parametric Constructors

Parametric types add a few wrinkles to the constructor story. Recall from *Parametric Types* (page 96) that, by default, instances of parametric composite types can be constructed either with explicitly given type parameters or with type parameters implied by the types of the arguments given to the constructor. Here are some examples:

```
julia> type Point{T<:Real}
    x::T
    y::T
end

## implicit T ##

julia> Point(1,2)
Point{Int64}(1,2)

julia> Point(1.0,2.5)
Point{Float64}(1.0,2.5)

julia> Point(1,2.5)
ERROR: MethodError: no method matching Point{T<:Real}(::Int64, ::Float64)
Closest candidates are:
  Point{T<:Real}{T<:Real}(::T<:Real, !Matched::T<:Real) at none:2
  Point{T<:Real}{T}(::Any) at sysimg.jl:66
...

## explicit T ##

julia> Point{Int64}(1,2)
Point{Int64}(1,2)

julia> Point{Int64}(1.0,2.5)
ERROR: InexactError()
in convert(::Type{Int64}, ::Float64) at ./float.jl:656
in Point{Int64}(::Float64, ::Float64) at ./none:2
...

julia> Point{Float64}(1.0,2.5)
Point{Float64}(1.0,2.5)

julia> Point{Float64}(1,2)
Point{Float64}(1.0,2.0)
```

As you can see, for constructor calls with explicit type parameters, the arguments are converted to the implied field types: `Point{Int64}(1,2)` works, but `Point{Int64}(1.0,2.5)` raises an `InexactError` when convert-

ing 2.5 to `Int64`. When the type is implied by the arguments to the constructor call, as in `Point(1,2)`, then the types of the arguments must agree — otherwise the `T` cannot be determined — but any pair of real arguments with matching type may be given to the generic `Point` constructor.

What’s really going on here is that `Point`, `Point{Float64}` and `Point{Int64}` are all different constructor functions. In fact, `Point{T}` is a distinct constructor function for each type `T`. Without any explicitly provided inner constructors, the declaration of the composite type `Point{T<:Real}` automatically provides an inner constructor, `Point{T}`, for each possible type `T<:Real`, that behaves just like non-parametric default inner constructors do. It also provides a single general outer `Point` constructor that takes pairs of real arguments, which must be of the same type. This automatic provision of constructors is equivalent to the following explicit declaration:

```
type Point{T<:Real}
    x::T
    y::T

    Point(x,y) = new(x,y)
end

Point{T<:Real}(x::T, y::T) = Point{T}(x,y)
```

Some features of parametric constructor definitions at work here deserve comment. First, inner constructor declarations always define methods of `Point{T}` rather than methods of the general `Point` constructor function. Since `Point` is not a concrete type, it makes no sense for it to even have inner constructor methods at all. Thus, the inner method declaration `Point(x,y) = new(x,y)` provides an inner constructor method for each value of `T`. It is this method declaration that defines the behavior of constructor calls with explicit type parameters like `Point{Int64}(1,2)` and `Point{Float64}(1.0,2.0)`. The outer constructor declaration, on the other hand, defines a method for the general `Point` constructor which only applies to pairs of values of the same real type. This declaration makes constructor calls without explicit type parameters, like `Point(1,2)` and `Point(1.0,2.5)`, work. Since the method declaration restricts the arguments to being of the same type, calls like `Point(1,2.5)`, with arguments of different types, result in “no method” errors.

Suppose we wanted to make the constructor call `Point(1,2.5)` work by “promoting” the integer value 1 to the floating-point value 1.0. The simplest way to achieve this is to define the following additional outer constructor method:

```
julia> Point(x::Int64, y::Float64) = Point(convert(Float64,x),y);
```

This method uses the `convert()` (page 364) function to explicitly convert `x` to `Float64` (page 437) and then delegates construction to the general constructor for the case where both arguments are `Float64` (page 437). With this method definition what was previously a `MethodError` (page 375) now successfully creates a point of type `Point{Float64}`:

```
julia> Point(1,2.5)
Point{Float64}(1.0,2.5)

julia> typeof(ans)
Point{Float64}
```

However, other similar calls still don’t work:

```
julia> Point(1.5,2)
ERROR: MethodError: no method matching Point{T<:Real}(::Float64, ::Int64)
Closest candidates are:
  Point{T<:Real}{T<:Real}(::T<:Real, !Matched::T<:Real) at none:2
  Point{T<:Real}{T}(::Any) at sysimg.jl:66
...
```

For a much more general way of making all such calls work sensibly, see [Conversion and Promotion](#) (page 127). At the risk of spoiling the suspense, we can reveal here that all it takes is the following outer method definition to make

all calls to the general `Point` constructor work as one would expect:

```
julia> Point(x::Real, y::Real) = Point(promote(x,y)...);
```

The `promote` function converts all its arguments to a common type — in this case `Float64` (page 437). With this method definition, the `Point` constructor promotes its arguments the same way that numeric operators like `+` (page 405) do, and works for all kinds of real numbers:

```
julia> Point(1.5,2)
Point{Float64}(1.5,2.0)

julia> Point(1,1//2)
Point{Rational{Int64}}(1//1,1//2)

julia> Point(1.0,1//2)
Point{Float64}(1.0,0.5)
```

Thus, while the implicit type parameter constructors provided by default in Julia are fairly strict, it is possible to make them behave in a more relaxed but sensible manner quite easily. Moreover, since constructors can leverage all of the power of the type system, methods, and multiple dispatch, defining sophisticated behavior is typically quite simple.

13.5 Case Study: Rational

Perhaps the best way to tie all these pieces together is to present a real world example of a parametric composite type and its constructor methods. To that end, here is beginning of `rational.jl`, which implements Julia's *Rational Numbers* (page 37):

```
immutable Rational{T<:Integer} <: Real
    num::T
    den::T

    function Rational(num::T, den::T)
        if num == 0 && den == 0
            error("invalid rational: 0//0")
        end
        g = gcd(den, num)
        num = div(num, g)
        den = div(den, g)
        new(num, den)
    end
end

Rational{T<:Integer}(n::T, d::T) = Rational{T}(n,d)
Rational(n::Integer, d::Integer) = Rational(promote(n,d)...)
Rational(n::Integer) = Rational(n,one(n))

//(n::Integer, d::Integer) = Rational(n,d)
//(x::Rational, y::Integer) = x.num // (x.den*y)
//(x::Integer, y::Rational) = (x*y.den) // y.num
//(x::Complex, y::Real) = complex(real(x)//y, imag(x)//y)
//(x::Real, y::Complex) = x*y'//real(y*y')

function //(x::Complex, y::Complex)
    xy = x*y'
    yy = real(y*y')
    complex(real(xy)//yy, imag(xy)//yy)
end
```

The first line — `immutable Rational{T<:Integer} <: Real` — declares that `Rational` takes one type parameter of an integer type, and is itself a real type. The field declarations `num::T` and `den::T` indicate that the data held in a `Rational{T}` object are a pair of integers of type `T`, one representing the rational value’s numerator and the other representing its denominator.

Now things get interesting. `Rational` has a single inner constructor method which checks that both of `num` and `den` aren’t zero and ensures that every rational is constructed in “lowest terms” with a non-negative denominator. This is accomplished by dividing the given numerator and denominator values by their greatest common divisor, computed using the `gcd` function. Since `gcd` returns the greatest common divisor of its arguments with sign matching the first argument (`den` here), after this division the new value of `den` is guaranteed to be non-negative. Because this is the only inner constructor for `Rational`, we can be certain that `Rational` objects are always constructed in this normalized form.

`Rational` also provides several outer constructor methods for convenience. The first is the “standard” general constructor that infers the type parameter `T` from the type of the numerator and denominator when they have the same type. The second applies when the given numerator and denominator values have different types: it promotes them to a common type and then delegates construction to the outer constructor for arguments of matching type. The third outer constructor turns integer values into rationals by supplying a value of 1 as the denominator.

Following the outer constructor definitions, we have a number of methods for the `//` (page 408) operator, which provides a syntax for writing rationals. Before these definitions, `//` (page 408) is a completely undefined operator with only syntax and no meaning. Afterwards, it behaves just as described in [Rational Numbers](#) (page 37) — its entire behavior is defined in these few lines. The first and most basic definition just makes `a//b` construct a `Rational` by applying the `Rational` constructor to `a` and `b` when they are integers. When one of the operands of `//` (page 408) is already a rational number, we construct a new rational for the resulting ratio slightly differently; this behavior is actually identical to division of a rational with an integer. Finally, applying `//` (page 408) to complex integral values creates an instance of `Complex{Rational}` — a complex number whose real and imaginary parts are rationals:

```
julia> (1 + 2im)//(1 - 2im)
-3//5 + 4//5*im

julia> typeof(ans)
Complex{Rational{Int64}}

julia> ans <: Complex{Rational}
false
```

Thus, although the `//` (page 408) operator usually returns an instance of `Rational`, if either of its arguments are complex integers, it will return an instance of `Complex{Rational}` instead. The interested reader should consider perusing the rest of `rational.jl`: it is short, self-contained, and implements an entire basic Julia type.

13.6 Constructors and Conversion

Constructors `T(args...)` in Julia are implemented like other callable objects: methods are added to their types. The type of a type is `Type`, so all constructor methods are stored in the method table for the `Type` type. This means that you can declare more flexible constructors, e.g. constructors for abstract types, by explicitly defining methods for the appropriate types.

However, in some cases you could consider adding methods to `Base.convert` *instead* of defining a constructor, because Julia falls back to calling `convert()` (page 364) if no matching constructor is found. For example, if no constructor `T(args...) = ...` exists `Base.convert(::Type{T}, args...) = ...` is called.

`convert` is used extensively throughout Julia whenever one type needs to be converted to another (e.g. in assignment, `ccall`, etcetera), and should generally only be defined (or successful) if the conversion is lossless. For example, `convert{Int, 3.0}` produces 3, but `convert{Int, 3.2}` throws an `InexactError`. If you want to define

a constructor for a lossless conversion from one type to another, you should probably define a `convert` method instead.

On the other hand, if your constructor does not represent a lossless conversion, or doesn't represent "conversion" at all, it is better to leave it as a constructor rather than a `convert` method. For example, the `Array{Int}()` constructor creates a zero-dimensional `Array` of the type `Int`, but is not really a "conversion" from `Int` to an `Array`.

13.7 Outer-only constructors

As we have seen, a typical parametric type has inner constructors that are called when type parameters are known; e.g. they apply to `Point{Int}` but not to `Point`. Optionally, outer constructors that determine type parameters automatically can be added, for example constructing a `Point{Int}` from the call `Point(1, 2)`. Outer constructors call inner constructors to do the core work of making an instance. However, in some cases one would rather not provide inner constructors, so that specific type parameters cannot be requested manually.

For example, say we define a type that stores a vector along with an accurate representation of its sum:

```
type SummedArray{T<:Number, S<:Number}
    data::Vector{T}
    sum::S
end
```

The problem is that we want `S` to be a larger type than `T`, so that we can sum many elements with less information loss. For example, when `T` is `Int32`, we would like `S` to be `Int64`. Therefore we want to avoid an interface that allows the user to construct instances of the type `SummedArray{Int32, Int32}`. One way to do this is to provide only an outer constructor for `SummedArray`. This can be done using method definition by type:

```
type SummedArray{T<:Number, S<:Number}
    data::Vector{T}
    sum::S

    function (::Type{SummedArray}){T}(a::Vector{T})
        S = widen{T}(T)
        new{T, S}(a, sum(S, a))
    end
end
```

This constructor will be invoked by the syntax `SummedArray(a)`. The syntax `new{T, S}` allows specifying parameters for the type to be constructed, i.e. this call will return a `SummedArray{T, S}`.

Conversion and Promotion

Julia has a system for promoting arguments of mathematical operators to a common type, which has been mentioned in various other sections, including *Integers and Floating-Point Numbers* (page 13), *Mathematical Operations and Elementary Functions* (page 25), *Types* (page 89), and *Methods* (page 109). In this section, we explain how this promotion system works, as well as how to extend it to new types and apply it to functions besides built-in mathematical operators. Traditionally, programming languages fall into two camps with respect to promotion of arithmetic arguments:

- **Automatic promotion for built-in arithmetic types and operators.** In most languages, built-in numeric types, when used as operands to arithmetic operators with infix syntax, such as `+`, `-`, `*`, and `/`, are automatically promoted to a common type to produce the expected results. C, Java, Perl, and Python, to name a few, all correctly compute the sum `1 + 1.5` as the floating-point value `2.5`, even though one of the operands to `+` is an integer. These systems are convenient and designed carefully enough that they are generally all-but-invisible to the programmer: hardly anyone consciously thinks of this promotion taking place when writing such an expression, but compilers and interpreters must perform conversion before addition since integers and floating-point values cannot be added as-is. Complex rules for such automatic conversions are thus inevitably part of specifications and implementations for such languages.
- **No automatic promotion.** This camp includes Ada and ML — very “strict” statically typed languages. In these languages, every conversion must be explicitly specified by the programmer. Thus, the example expression `1 + 1.5` would be a compilation error in both Ada and ML. Instead one must write `real(1) + 1.5`, explicitly converting the integer `1` to a floating-point value before performing addition. Explicit conversion everywhere is so inconvenient, however, that even Ada has some degree of automatic conversion: integer literals are promoted to the expected integer type automatically, and floating-point literals are similarly promoted to appropriate floating-point types.

In a sense, Julia falls into the “no automatic promotion” category: mathematical operators are just functions with special syntax, and the arguments of functions are never automatically converted. However, one may observe that applying mathematical operations to a wide variety of mixed argument types is just an extreme case of polymorphic multiple dispatch — something which Julia’s dispatch and type systems are particularly well-suited to handle. “Automatic” promotion of mathematical operands simply emerges as a special application: Julia comes with pre-defined catch-all dispatch rules for mathematical operators, invoked when no specific implementation exists for some combination of operand types. These catch-all rules first promote all operands to a common type using user-definable promotion rules, and then invoke a specialized implementation of the operator in question for the resulting values, now of the same type. User-defined types can easily participate in this promotion system by defining methods for conversion to and from other types, and providing a handful of promotion rules defining what types they should promote to when mixed with other types.

14.1 Conversion

Conversion of values to various types is performed by the `convert` function. The `convert` function generally takes two arguments: the first is a type object while the second is a value to convert to that type; the returned value is the value converted to an instance of given type. The simplest way to understand this function is to see it in action:

```
julia> x = 12
12

julia> typeof(x)
Int64

julia> convert(UInt8, x)
0x0c

julia> typeof(ans)
UInt8

julia> convert(AbstractFloat, x)
12.0

julia> typeof(ans)
Float64

julia> a = Any[1 2 3; 4 5 6]
2x3 Array{Any,2}:
 1  2  3
 4  5  6

julia> convert(Array{Float64}, a)
2x3 Array{Float64,2}:
 1.0  2.0  3.0
 4.0  5.0  6.0
```

Conversion isn't always possible, in which case a `MethodError` is thrown indicating that `convert` doesn't know how to perform the requested conversion:

```
julia> convert(AbstractFloat, "foo")
ERROR: MethodError: Cannot `convert` an object of type String to an object of type AbstractFloat
This may have arisen from a call to the constructor AbstractFloat(...),
since type constructors fall back to convert methods.
...
```

Some languages consider parsing strings as numbers or formatting numbers as strings to be conversions (many dynamic languages will even perform conversion for you automatically), however Julia does not: even though some strings can be parsed as numbers, most strings are not valid representations of numbers, and only a very limited subset of them are. Therefore in Julia the dedicated `parse()` function must be used to perform this operation, making it more explicit.

14.1.1 Defining New Conversions

To define a new conversion, simply provide a new method for `convert()`. That's really all there is to it. For example, the method to convert a real number to a boolean is this:

```
convert{::Type{Bool}, ::Real}() = x==0 ? false : x==1 ? true : throw(InexactError())
```

The type of the first argument of this method is a *singleton type* (page 102), `Type{Bool}`, the only instance of which is `Bool`. Thus, this method is only invoked when the first argument is the type value `Bool`. Notice the syntax used for the first argument: the argument name is omitted prior to the `::` symbol, and only the type is given. This is the syntax in Julia for a function argument whose type is specified but whose value is never used in the function body. In this example, since the type is a singleton, there would never be any reason to use its value within the body. When invoked, the method determines whether a numeric value is true or false as a boolean, by comparing it to one and zero:

```
julia> convert(Bool, 1)
true

julia> convert(Bool, 0)
false

julia> convert(Bool, 1im)
ERROR: InexactError()
  in convert(::Type{Bool}, ::Complex{Int64}) at ./complex.jl:23
...

julia> convert(Bool, 0im)
false
```

The method signatures for conversion methods are often quite a bit more involved than this example, especially for parametric types. The example above is meant to be pedagogical, and is not the actual Julia behaviour. This is the actual implementation in Julia:

```
convert{T<:Real}(::Type{T}, z::Complex) = (imag(z)==0 ? convert(T, real(z)) :
                                           throw(InexactError()))

julia> convert(Bool, 1im)
ERROR: InexactError()
  in convert(::Type{Bool}, ::Complex{Int64}) at ./complex.jl:18
...
```

14.1.2 Case Study: Rational Conversions

To continue our case study of Julia's `Rational` type, here are the conversions declared in `rational.jl`, right after the declaration of the type and its constructors:

```
convert{T<:Integer}(::Type{Rational{T}}, x::Rational) = Rational(convert(T, x.num), convert(T, x.den))
convert{T<:Integer}(::Type{Rational{T}}, x::Integer) = Rational(convert(T, x), convert(T, 1))

function convert{T<:Integer}(::Type{Rational{T}}, x::AbstractFloat, tol::Real)
    if isnan(x); return zero(T)//zero(T); end
    if isinf(x); return sign(x)//zero(T); end
    y = x
    a = d = one(T)
    b = c = zero(T)
    while true
        f = convert(T, round(y)); y -= f
        a, b, c, d = f*a+c, f*b+d, a, b
        if y == 0 || abs(a/b-x) <= tol
            return a//b
        end
        y = 1/y
    end
end
convert{T<:Integer}(rt::Type{Rational{T}}, x::AbstractFloat) = convert(rt, x, eps(x))
```

```
convert{T<:AbstractFloat}(::Type{T}, x::Rational) = convert(T,x.num)/convert(T,x.den)
convert{T<:Integer}(::Type{T}, x::Rational) = div(convert(T,x.num),convert(T,x.den))
```

The initial four `convert` methods provide conversions to rational types. The first method converts one type of rational to another type of rational by converting the numerator and denominator to the appropriate integer type. The second method does the same conversion for integers by taking the denominator to be 1. The third method implements a standard algorithm for approximating a floating-point number by a ratio of integers to within a given tolerance, and the fourth method applies it, using machine epsilon at the given value as the threshold. In general, one should have `a/b == convert(Rational{Int64}, a/b)`.

The last two `convert` methods provide conversions from rational types to floating-point and integer types. To convert to floating point, one simply converts both numerator and denominator to that floating point type and then divides. To convert to integer, one can use the `div` operator for truncated integer division (rounded towards zero).

14.2 Promotion

Promotion refers to converting values of mixed types to a single common type. Although it is not strictly necessary, it is generally implied that the common type to which the values are converted can faithfully represent all of the original values. In this sense, the term “promotion” is appropriate since the values are converted to a “greater” type — i.e. one which can represent all of the input values in a single common type. It is important, however, not to confuse this with object-oriented (structural) super-typing, or Julia’s notion of abstract super-types: promotion has nothing to do with the type hierarchy, and everything to do with converting between alternate representations. For instance, although every `Int32` value can also be represented as a `Float64` value, `Int32` is not a subtype of `Float64`.

Promotion to a common “greater” type is performed in Julia by the `promote` function, which takes any number of arguments, and returns a tuple of the same number of values, converted to a common type, or throws an exception if promotion is not possible. The most common use case for promotion is to convert numeric arguments to a common type:

```
julia> promote(1, 2.5)
(1.0, 2.5)

julia> promote(1, 2.5, 3)
(1.0, 2.5, 3.0)

julia> promote(2, 3//4)
(2//1, 3//4)

julia> promote(1, 2.5, 3, 3//4)
(1.0, 2.5, 3.0, 0.75)

julia> promote(1.5, im)
(1.5 + 0.0im, 0.0 + 1.0im)

julia> promote(1 + 2im, 3//4)
(1//1 + 2//1*im, 3//4 + 0//1*im)
```

Floating-point values are promoted to the largest of the floating-point argument types. Integer values are promoted to the larger of either the native machine word size or the largest integer argument type. Mixtures of integers and floating-point values are promoted to a floating-point type big enough to hold all the values. Integers mixed with rationals are promoted to rationals. Rationals mixed with floats are promoted to floats. Complex values mixed with real values are promoted to the appropriate kind of complex value.

That is really all there is to using promotions. The rest is just a matter of clever application, the most typical “clever” application being the definition of catch-all methods for numeric operations like the arithmetic operators `+`, `-`, `*` and

/. Here are some of the catch-all method definitions given in `promotion.jl`:

```
+ (x::Number, y::Number) = + (promote(x,y)...)
- (x::Number, y::Number) = - (promote(x,y)...)
* (x::Number, y::Number) = * (promote(x,y)...)
/ (x::Number, y::Number) = / (promote(x,y)...)

```

These method definitions say that in the absence of more specific rules for adding, subtracting, multiplying and dividing pairs of numeric values, promote the values to a common type and then try again. That's all there is to it: nowhere else does one ever need to worry about promotion to a common numeric type for arithmetic operations — it just happens automatically. There are definitions of catch-all promotion methods for a number of other arithmetic and mathematical functions in `promotion.jl`, but beyond that, there are hardly any calls to `promote` required in the Julia standard library. The most common usages of `promote` occur in outer constructors methods, provided for convenience, to allow constructor calls with mixed types to delegate to an inner type with fields promoted to an appropriate common type. For example, recall that `rational.jl` provides the following outer constructor method:

```
Rational(n::Integer, d::Integer) = Rational(promote(n,d)...)

```

This allows calls like the following to work:

```
julia> Rational{Int8}(15), Int32(-5)
-3//1

julia> typeof(ans)
Rational{Int32}

```

For most user-defined types, it is better practice to require programmers to supply the expected types to constructor functions explicitly, but sometimes, especially for numeric problems, it can be convenient to do promotion automatically.

14.2.1 Defining Promotion Rules

Although one could, in principle, define methods for the `promote` function directly, this would require many redundant definitions for all possible permutations of argument types. Instead, the behavior of `promote` is defined in terms of an auxiliary function called `promote_rule`, which one can provide methods for. The `promote_rule` function takes a pair of type objects and returns another type object, such that instances of the argument types will be promoted to the returned type. Thus, by defining the rule:

```
promote_rule{::Type{Float64}, ::Type{Float32}}() = Float64

```

one declares that when 64-bit and 32-bit floating-point values are promoted together, they should be promoted to 64-bit floating-point. The promotion type does not need to be one of the argument types, however; the following promotion rules both occur in Julia's standard library:

```
promote_rule{::Type{UInt8}, ::Type{Int8}}() = Int
promote_rule{::Type{BigInt}, ::Type{Int8}}() = BigInt

```

In the latter case, the result type is `BigInt` since `BigInt` is the only type large enough to hold integers for arbitrary-precision integer arithmetic. Also note that one does not need to define both `promote_rule{::Type{A}, ::Type{B}}` and `promote_rule{::Type{B}, ::Type{A}}` — the symmetry is implied by the way `promote_rule` is used in the promotion process.

The `promote_rule` function is used as a building block to define a second function called `promote_type`, which, given any number of type objects, returns the common type to which those values, as arguments to `promote` should be promoted. Thus, if one wants to know, in absence of actual values, what type a collection of values of certain types would promote to, one can use `promote_type`:

```
julia> promote_type{Int8, UInt16}
Int64
```

Internally, `promote_type` is used inside of `promote` to determine what type argument values should be converted to for promotion. It can, however, be useful in its own right. The curious reader can read the code in [promotion.jl](#), which defines the complete promotion mechanism in about 35 lines.

14.2.2 Case Study: Rational Promotions

Finally, we finish off our ongoing case study of Julia’s rational number type, which makes relatively sophisticated use of the promotion mechanism with the following promotion rules:

```
promote_rule{T<:Integer,S<:Integer} (::Type{Rational{T}}, ::Type{S}) = Rational{promote_type{T,S}}
promote_rule{T<:Integer,S<:Integer} (::Type{Rational{T}}, ::Type{Rational{S}}) = Rational{promote_type{T,S}}
promote_rule{T<:Integer,S<:AbstractFloat} (::Type{Rational{T}}, ::Type{S}) = promote_type{T,S}
```

The first rule says that promoting a rational number with any other integer type promotes to a rational type whose numerator/denominator type is the result of promotion of its numerator/denominator type with the other integer type. The second rule applies the same logic to two different types of rational numbers, resulting in a rational of the promotion of their respective numerator/denominator types. The third and final rule dictates that promoting a rational with a float results in the same type as promoting the numerator/denominator type with the float.

This small handful of promotion rules, together with the *conversion methods discussed above* (page 129), are sufficient to make rational numbers interoperate completely naturally with all of Julia’s other numeric types — integers, floating-point numbers, and complex numbers. By providing appropriate conversion methods and promotion rules in the same manner, any user-defined numeric type can interoperate just as naturally with Julia’s predefined numerics.

Interfaces

A lot of the power and extensibility in Julia comes from a collection of informal interfaces. By extending a few specific methods to work for a custom type, objects of that type not only receive those functionalities, but they are also able to be used in other methods that are written to generically build upon those behaviors.

15.1 Iteration

Required methods		Brief description
<code>start(iter)</code>		Returns the initial iteration state
<code>next(iter, state)</code>		Returns the current item and the next state
<code>done(iter, state)</code>		Tests if there are any items remaining
Important optional methods	Default definition	Brief description
<code>itertoolsize(IterType)</code>	<code>HasLength()</code>	One of <i>HasLength()</i> , <i>HasShape()</i> , <i>IsInfinite()</i> , or <i>SizeUnknown()</i> as appropriate
<code>iteratoreltype(IterType)</code>	<code>HasEltype()</code>	Either <i>EltypeUnknown()</i> or <i>HasEltype()</i> as appropriate
<code>eltype(IterType)</code>	<code>Any</code>	The type the items returned by <code>next()</code>
<code>length(iter)</code>	<i>(undefined)</i>	The number of items, if known
<code>size(iter, [dim...])</code>	<i>(undefined)</i>	The number of items in each dimension, if known

Value returned by <code>itertoolsize(IterType)</code>	Required Methods
<i>HasLength()</i>	<code>length(iter)</code>
<i>HasShape()</i>	<code>length(iter)</code> and <code>size(iter, [dim...])</code>
<i>IsInfinite()</i>	<i>(none)</i>
<i>SizeUnknown()</i>	<i>(none)</i>

Value returned by <code>iteratoreltype(IterType)</code>	Required Methods
<i>HasEltype()</i>	<code>eltype(IterType)</code>
<i>EltypeUnknown()</i>	<i>(none)</i>

Sequential iteration is implemented by the methods `start()`, `done()`, and `next()`. Instead of mutating objects as they are iterated over, Julia provides these three methods to keep track of the iteration state externally from the object. The `start(iter)` method returns the initial state for the iterable object `iter`. That state gets passed along to `done(iter, state)`, which tests if there are any elements remaining, and `next(iter, state)`, which returns a tuple containing the current element and an updated state. The state object can be anything, and is generally considered to be an implementation detail private to the iterable object.

Any object defines these three methods is iterable and can be used in the *many functions that rely upon iteration* (page 381). It can also be used directly in a `for` loop since the syntax:

```
for i in iter    # or "for i = iter"
    # body
end
```

is translated into:

```
state = start(iter)
while !done(iter, state)
    (i, state) = next(iter, state)
    # body
end
```

A simple example is an iterable sequence of square numbers with a defined length:

```
julia> immutable Squares
        count::Int
    end
Base.start(::Squares) = 1
Base.next(S::Squares, state) = (state*state, state+1)
Base.done(S::Squares, state) = state > S.count;
Base.eltypes(::Type{Squares}) = Int # Note that this is defined for the type
Base.length(S::Squares) = S.count;
```

With only `start`, `next`, and `done` definitions, the `Squares` type is already pretty powerful. We can iterate over all the elements:

```
julia> for i in Squares(7)
        println(i)
    end
1
4
9
16
25
36
49
```

We can use many of the builtin methods that work with iterables, like `in()`, `mean()` and `std()`:

```
julia> 25 in Squares(10)
true

julia> mean(Squares(100)), std(Squares(100))
(3383.5, 3024.355854282583)
```

There are a few more methods we can extend to give Julia more information about this iterable collection. We know that the elements in a `Squares` sequence will always be `Int`. By extending the `eltype()` method, we can give that information to Julia and help it make more specialized code in the more complicated methods. We also know the number of elements in our sequence, so we can extend `length()`, too.

Now, when we ask Julia to `collect()` all the elements into an array it can preallocate a `Vector{Int}` of the right size instead of blindly pushing each element into a `Vector{Any}`:

```
julia> collect(Squares(100))' # transposed to save space
1×100 Array{Int64,2}:
 1  4  9 16 25 36 49 64 81 100 ... 9025 9216 9409 9604 9801 10000
```

While we can rely upon generic implementations, we can also extend specific methods where we know there is a simpler algorithm. For example, there's a formula to compute the sum of squares, so we can override the generic iterative version with a more performant solution:

```
julia> Base.sum(S::Squares) = (n = S.count; return n*(n+1)*(2n+1)÷6)
      sum(Squares(1803))
1955361914
```

This is a very common pattern throughout the Julia standard library: a small set of required methods define an informal interface that enable many fancier behaviors. In some cases, types will want to additionally specialize those extra behaviors when they know a more efficient algorithm can be used in their specific case.

15.2 Indexing

Methods to implement	Brief description
<code>getindex(X, i)</code>	<code>X[i]</code> , indexed element access
<code>setindex!(X, v, i)</code>	<code>X[i] = v</code> , indexed assignment
<code>endof(X)</code>	The last index, used in <code>X[end]</code>

For the `Squares` iterable above, we can easily compute the i th element of the sequence by squaring it. We can expose this as an indexing expression `S[i]`. To opt into this behavior, `Squares` simply needs to define `getindex()`:

```
julia> function Base.getindex(S::Squares, i::Int)
    1 <= i <= S.count || throw(BoundsError(S, i))
    return i*i
end
Squares(100)[23]
529
```

Additionally, to support the syntax `S[end]`, we must define `endof()` to specify the last valid index:

```
julia> Base.endof(S::Squares) = length(S)
Squares(23)[end]
529
```

Note, though, that the above *only* defines `getindex()` with one integer index. Indexing with anything other than an `Int` will throw a `MethodError` saying that there was no matching method. In order to support indexing with ranges or vectors of `Ints`, separate methods must be written:

```
julia> Base.getindex(S::Squares, i::Number) = S[convert{Int, i}]
      Base.getindex(S::Squares, I) = [S[i] for i in I]
Squares(10)[[3,4,5]]
3-element Array{Int64,1}:
 9
16
25
```

While this is starting to support more of the *indexing operations supported by some of the builtin types* (page 180), there's still quite a number of behaviors missing. This `Squares` sequence is starting to look more and more like a vector as we've added behaviors to it. Instead of defining all these behaviors ourselves, we can officially define it as a subtype of an `AbstractArray`.

15.3 Abstract Arrays

Methods to implement		Brief description
<code>size(A)</code>		Returns a tuple containing the dimensions of <code>A</code>
<code>getindex(A, i::Int)</code>		(if <code>LinearFast</code>) Linear scalar indexing
<code>getindex(A, I::Vararg{Int, N})</code>		(if <code>LinearSlow</code> , where <code>N = ndims(A)</code>) N-dimensional scalar indexing
<code>setindex!(A, v, i::Int)</code>		(if <code>LinearFast</code>) Scalar indexed assignment
<code>setindex!(A, v, I::Vararg{Int, N})</code>		(if <code>LinearSlow</code> , where <code>N = ndims(A)</code>) N-dimensional scalar indexed assignment
Optional methods	Default definition	Brief description
<code>Base.linearindexing(::Type{A})</code>	<code>Base.LinearSlow()</code>	Returns either <code>Base.LinearFast()</code> or <code>Base.LinearSlow()</code> . See the description below.
<code>getindex(A, I...)</code>	defined in terms of scalar <code>getindex()</code>	<i>Multidimensional and nonscalar indexing</i> (page 180)
<code>setindex!(A, I...)</code>	defined in terms of scalar <code>setindex!()</code>	<i>Multidimensional and nonscalar indexed assignment</i> (page 180)
<code>start()/next()/done()</code>	defined in terms of scalar <code>getindex()</code>	Iteration
<code>length(A)</code>	<code>prod(size(A))</code>	Number of elements
<code>similar(A)</code>	<code>similar(A, eltype(A), size(A))</code>	Return a mutable array with the same shape and element type
<code>similar(A, ::Type{S})</code>	<code>similar(A, S, size(A))</code>	Return a mutable array with the same shape and the specified element type
<code>similar(A, dims::NTuple{Int})</code>	<code>similar(A, eltype(A), dims)</code>	Return a mutable array with the same element type and size <i>dims</i>
<code>similar(A, ::Type{S}, dims::NTuple{Int})</code>	<code>Array{S}(dims)</code>	Return a mutable array with the specified element type and size
Non-traditional indices	Default definition	Brief description
<code>indices(A)</code>	<code>map(OneTo, size(A))</code>	Return the <code>AbstractUnitRange</code> of valid indices
<code>Base.similar(A, ::Type{S}, inds::NTuple{Ind})</code>	<code>similar(A, S, Base.to_shape(inds))</code>	Return a mutable array with the specified indices <i>inds</i> (see below)
<code>Base.similar(T::Union{Type{A}, Base.RefValue{A}}, inds)</code>	<code>Base.to_shape(inds)</code>	Return an array similar to <code>T</code> with the specified indices <i>inds</i> (see below)

If a type is defined as a subtype of `AbstractArray`, it inherits a very large set of rich behaviors including iteration and multidimensional indexing built on top of single-element access. See the [arrays manual page](#) (page 177) and [standard library section](#) (page 451) for more supported methods.

A key part in defining an `AbstractArray` subtype is `Base.linearindexing()`. Since indexing is such an important part of an array and often occurs in hot loops, it's important to make both indexing and indexed assignment as efficient as possible. Array data structures are typically defined in one of two ways: either it most efficiently accesses its elements using just one index (linear indexing) or it intrinsically accesses the elements with indices specified for every dimension. These two modalities are identified by Julia as `Base.LinearFast()` and `Base.LinearSlow()`. Converting a linear index to multiple indexing subscripts is typically very expensive, so this provides a traits-based mechanism to enable efficient generic code for all array types.

This distinction determines which scalar indexing methods the type must define. `LinearFast()` arrays are simple: just define `getindex(A::ArrayType, i::Int)`. When the array is subsequently indexed with a mul-

tidimensional set of indices, the fallback `getindex(A::AbstractArray, I...)` efficiently converts the indices into one linear index and then calls the above method. `LinearSlow()` arrays, on the other hand, require methods to be defined for each supported dimensionality with `ndims(A)` `Int` indices. For example, the builtin `SparseMatrixCSC` type only supports two dimensions, so it just defines `getindex(A::SparseMatrixCSC, i::Int, j::Int)`. The same holds for `setindex!`.

Returning to the sequence of squares from above, we could instead define it as a subtype of an `AbstractArray{Int, 1}`:

```
julia> immutable SquaresVector <: AbstractArray{Int, 1}
    count::Int
end
Base.size(S::SquaresVector) = (S.count,)
Base.linearindexing{T<:SquaresVector}{::Type{T}} = Base.LinearFast()
Base.getindex(S::SquaresVector, i::Int) = i*i;
```

Note that it's very important to specify the two parameters of the `AbstractArray`; the first defines the `eltype()`, and the second defines the `ndims()`. That supertype and those three methods are all it takes for `SquaresVector` to be an iterable, indexable, and completely functional array:

```
julia> s = SquaresVector(7)
7-element SquaresVector:
 1
 4
 9
16
25
36
49

julia> s[s .> 20]
3-element Array{Int64,1}:
25
36
49

julia> s \ rand(7,2)
1×2 Array{Float64,2}:
0.0151876 0.0179393
```

As a more complicated example, let's define our own toy N-dimensional sparse-like array type built on top of `Dict`:

```
julia> immutable SparseArray{T,N} <: AbstractArray{T,N}
    data::Dict{NTuple{N,Int}, T}
    dims::NTuple{N,Int}
end
SparseArray{T}{::Type{T}, dims::Int...} = SparseArray{T, dims}
SparseArray{T,N}{::Type{T}, dims::NTuple{N,Int}} = SparseArray{T,N}(Dict{NTuple{N,Int}, T}(),
SparseArray{T,N})

julia> Base.size(A::SparseArray) = A.dims
Base.similar{T}(A::SparseArray, ::Type{T}, dims::Dims) = SparseArray{T, dims}
# Define scalar indexing and indexed assignment
Base.getindex{T,N}(A::SparseArray{T,N}, I::Vararg{Int,N}) = get(A.data, I, zero(T))
Base.setindex!{T,N}(A::SparseArray{T,N}, v, I::Vararg{Int,N}) = (A.data[I] = v)
```

Notice that this is a `LinearSlow` array, so we must manually define `getindex()` and `setindex!()` at the dimensionality of the array. Unlike the `SquaresVector`, we are able to define `setindex!()`, and so we can mutate the array:

```
julia> A = SparseArray{Float64,3,3}
3×3 SparseArray{Float64,2}:
 0.0  0.0  0.0
 0.0  0.0  0.0
 0.0  0.0  0.0

julia> rand!(A)
3×3 SparseArray{Float64,2}:
 0.28119  0.0203749  0.0769509
 0.209472  0.287702  0.640396
 0.251379  0.859512  0.873544

julia> A[:] = 1:length(A); A
3×3 SparseArray{Float64,2}:
 1.0  4.0  7.0
 2.0  5.0  8.0
 3.0  6.0  9.0
```

The result of indexing an `AbstractArray` can itself be an array (for instance when indexing by a `Range`). The `AbstractArray` fallback methods use `similar()` to allocate an `Array` of the appropriate size and element type, which is filled in using the basic indexing method described above. However, when implementing an array wrapper you often want the result to be wrapped as well:

```
julia> A[1:2,:]
2×3 SparseArray{Float64,2}:
 1.0  4.0  7.0
 2.0  5.0  8.0
```

In this example it is accomplished by defining `Base.similar{T}(A::SparseArray, ::Type{T}, dims::Dims)` to create the appropriate wrapped array. (Note that while `similar` supports 1- and 2-argument forms, in most case you only need to specialize the 3-argument form.) For this to work it's important that `SparseArray` is mutable (supports `setindex!`). `similar()` is also used to allocate result arrays for arithmetic on `AbstractArrays`, for instance:

```
julia> A + 4
3×3 SparseArray{Float64,2}:
 5.0  8.0  11.0
 6.0  9.0  12.0
 7.0  10.0  13.0
```

In addition to all the iterable and indexable methods from above, these types can also interact with each other and use all of the methods defined in the standard library for `AbstractArrays`:

```
julia> A[SquaresVector{3}]
3-element SparseArray{Float64,1}:
 1.0
 4.0
 9.0

julia> dot(A[:,1],A[:,2])
32.0
```

If you are defining an array type that allows non-traditional indexing (indices that start at something other than 1), you should specialize `indices`. You should also specialize `similar` so that the `dims` argument (ordinarily a `Dims` size-tuple) can accept `AbstractUnitRange` objects, perhaps range-types `Ind` of your own design. For more information, see `devdocs-offsetarrays`.

Modules

Modules in Julia are separate variable workspaces, i.e. they introduce a new global scope. They are delimited syntactically, inside `module Name ... end`. Modules allow you to create top-level definitions (aka global variables) without worrying about name conflicts when your code is used together with somebody else's. Within a module, you can control which names from other modules are visible (via importing), and specify which of your names are intended to be public (via exporting).

The following example demonstrates the major features of modules. It is not meant to be run, but is shown for illustrative purposes:

```
module MyModule
using Lib

using BigLib: thing1, thing2

import Base.show

importall OtherLib

export MyType, foo

type MyType
    x
end

bar(x) = 2x
foo(a::MyType) = bar(a.x) + 1

show(io::IO, a::MyType) = print(io, "MyType $(a.x)")
end
```

Note that the style is not to indent the body of the module, since that would typically lead to whole files being indented.

This module defines a type `MyType`, and two functions. Function `foo` and type `MyType` are exported, and so will be available for importing into other modules. Function `bar` is private to `MyModule`.

The statement `using Lib` means that a module called `Lib` will be available for resolving names as needed. When a global variable is encountered that has no definition in the current module, the system will search for it among variables exported by `Lib` and import it if it is found there. This means that all uses of that global within the current module will resolve to the definition of that variable in `Lib`.

The statement `using BigLib: thing1, thing2` is a syntactic shortcut for `using BigLib.thing1, BigLib.thing2`.

The `import` keyword supports all the same syntax as `using`, but only operates on a single name at a time. It does

not add modules to be searched the way `using` does. `import` also differs from `using` in that functions must be imported using `import` to be extended with new methods.

In `MyModule` above we wanted to add a method to the standard `show` function, so we had to write `import Base.show`. Functions whose names are only visible via `using` cannot be extended.

The keyword `importall` explicitly imports all names exported by the specified module, as if `import` were individually used on all of them.

Once a variable is made visible via `using` or `import`, a module may not create its own variable with the same name. Imported variables are read-only; assigning to a global variable always affects a variable owned by the current module, or else raises an error.

16.1 Summary of module usage

To load a module, two main keywords can be used: `using` and `import`. To understand their differences, consider the following example:

```
module MyModule

export x, y

x() = "x"
y() = "y"
p() = "p"

end
```

In this module we export the `x` and `y` functions (with the keyword `export`), and also have the non-exported function `p`. There are several different ways to load the Module and its inner functions into the current workspace:

Import Command	What is brought into scope	Available for method extension
<code>using MyModule</code>	All exported names (<code>x</code> and <code>y</code>), <code>MyModule.x</code> , <code>MyModule.y</code> and <code>MyModule.p</code>	<code>MyModule.x</code> , <code>MyModule.y</code> and <code>MyModule.p</code>
<code>using MyModule.x, MyModule.p</code>	<code>x</code> and <code>p</code>	
<code>using MyModule: x, p</code>	<code>x</code> and <code>p</code>	
<code>import MyModule</code>	<code>MyModule.x</code> , <code>MyModule.y</code> and <code>MyModule.p</code>	<code>MyModule.x</code> , <code>MyModule.y</code> and <code>MyModule.p</code>
<code>import MyModule.x, MyModule.p</code>	<code>x</code> and <code>p</code>	<code>x</code> and <code>p</code>
<code>import MyModule: x, p</code>	<code>x</code> and <code>p</code>	<code>x</code> and <code>p</code>
<code>importall MyModule</code>	All exported names (<code>x</code> and <code>y</code>)	<code>x</code> and <code>y</code>

16.1.1 Modules and files

Files and file names are mostly unrelated to modules; modules are associated only with module expressions. One can have multiple files per module, and multiple modules per file:

```

module Foo

include("file1.jl")
include("file2.jl")

end

```

Including the same code in different modules provides mixin-like behavior. One could use this to run the same code with different base definitions, for example testing code by running it with “safe” versions of some operators:

```

module Normal
include("mycode.jl")
end

module Testing
include("safe_operators.jl")
include("mycode.jl")
end

```

16.1.2 Standard modules

There are three important standard modules: Main, Core, and Base.

Main is the top-level module, and Julia starts with Main set as the current module. Variables defined at the prompt go in Main, and `whos()` lists variables in Main.

Core contains all identifiers considered “built in” to the language, i.e. part of the core language and not libraries. Every module implicitly specifies `using Core`, since you can’t do anything without those definitions.

Base is the standard library (the contents of `base/`). All modules implicitly contain `using Base`, since this is needed in the vast majority of cases.

16.1.3 Default top-level definitions and bare modules

In addition to `using Base`, modules also automatically contain a definition of the `eval` function, which evaluates expressions within the context of that module.

If these default definitions are not wanted, modules can be defined using the keyword `baremodule` instead (note: Core is still imported, as per above). In terms of `baremodule`, a standard module looks like this:

```

baremodule Mod

using Base

eval(x) = Core.eval(Mod, x)
eval(m, x) = Core.eval(m, x)

...

end

```

16.1.4 Relative and absolute module paths

Given the statement `using Foo`, the system looks for `Foo` within `Main`. If the module does not exist, the system attempts to `require("Foo")`, which typically results in loading code from an installed package.

However, some modules contain submodules, which means you sometimes need to access a module that is not directly available in `Main`. There are two ways to do this. The first is to use an absolute path, for example using `Base.Sort`. The second is to use a relative path, which makes it easier to import submodules of the current module or any of its enclosing modules:

```
module Parent

module Utils
...
end

using .Utils

...
end
```

Here module `Parent` contains a submodule `Utils`, and code in `Parent` wants the contents of `Utils` to be visible. This is done by starting the `using` path with a period. Adding more leading periods moves up additional levels in the module hierarchy. For example using `..Utils` would look for `Utils` in `Parent`'s enclosing module rather than in `Parent` itself.

Note that relative-import qualifiers are only valid in `using` and `import` statements.

16.1.5 Module file paths

The global variable `LOAD_PATH` contains the directories Julia searches for modules when calling `require`. It can be extended using `push!`:

```
push!(LOAD_PATH, "/Path/To/My/Module/")
```

Putting this statement in the file `~/.juliarc.jl` will extend `LOAD_PATH` on every Julia startup. Alternatively, the module load path can be extended by defining the environment variable `JULIA_LOAD_PATH`.

16.1.6 Namespace miscellanea

If a name is qualified (e.g. `Base.sin`), then it can be accessed even if it is not exported. This is often useful when debugging. It can also have methods added to it by using the qualified name as the function name. However, due to syntactic ambiguities that arise, if you wish to add methods to a function in a different module whose name contains only symbols, such as an operator, `Base. +` for example, you must use `Base. : +` to refer to it. If the operator is more than one character in length you must surround it in brackets, such as: `Base. : (==)`.

Macro names are written with `@` in import and export statements, e.g. `import Mod.@mac`. Macros in other modules can be invoked as `Mod.@mac` or `@Mod.mac`.

The syntax `M.x = y` does not work to assign a global in another module; global assignment is always module-local.

A variable can be “reserved” for the current module without assigning to it by declaring it as `global x` at the top level. This can be used to prevent name conflicts for globals initialized after load time.

16.1.7 Module initialization and precompilation

Large modules can take several seconds to load because executing all of the statements in a module often involves compiling a large amount of code. Julia provides the ability to create precompiled versions of modules to reduce this time.

To create an incremental precompiled module file, add `__precompile__()` at the top of your module file (before the module starts). This will cause it to be automatically compiled the first time it is imported. Alternatively, you can manually call `Base.compilecache(modulename)`. The resulting cache files will be stored in `Base.LOAD_CACHE_PATH[1]`. Subsequently, the module is automatically recompiled upon `import` whenever any of its dependencies change; dependencies are modules it imports, the Julia build, files it includes, or explicit dependencies declared by `include_dependency(path)` in the module file(s).

For file dependencies, a change is determined by examining whether the modification time (mtime) of each file loaded by `include` or added explicitly by `include_dependency` is unchanged, or equal to the modification time truncated to the nearest second (to accommodate systems that can't copy mtime with sub-second accuracy). It also takes into account whether the path to the file chosen by the search logic in `require` matches the path that had created the precompile file.

It also takes into account the set of dependencies already loaded into the current process and won't recompile those modules, even if their files change or disappear, in order to avoid creating incompatibilities between the running system and the precompile cache. If you want to have changes to the source reflected in the running system, you should call `reload("Module")` on the module you changed, and any module that depended on it in which you want to see the change reflected.

Precompiling a module also recursively precompiles any modules that are imported therein. If you know that it is *not* safe to precompile your module (for the reasons described below), you should put `__precompile__(false)` in the module file to cause `Base.compilecache` to throw an error (and thereby prevent the module from being imported by any other precompiled module).

`__precompile__()` should *not* be used in a module unless all of its dependencies are also using `__precompile__()`. Failure to do so can result in a runtime error when loading the module.

In order to make your module work with precompilation, however, you may need to change your module to explicitly separate any initialization steps that must occur at *runtime* from steps that can occur at *compile time*. For this purpose, Julia allows you to define an `__init__()` function in your module that executes any initialization steps that must occur at runtime. This function will not be called during compilation (`--output-*` or `__precompile__()`). You may, of course, call it manually if necessary, but the default is to assume this function deals with computing state for the local machine, which does not need to be – or even should not be – captured in the compiled image. It will be called after the module is loaded into a process, including if it is being loaded into an incremental compile (`--output-incremental=yes`), but not if it is being loaded into a full-compilation process.

In particular, if you define a function `__init__()` in a module, then Julia will call `__init__()` immediately *after* the module is loaded (e.g., by `import`, `using`, or `require`) at runtime for the *first* time (i.e., `__init__` is only called once, and only after all statements in the module have been executed). Because it is called after the module is fully imported, any submodules or other imported modules have their `__init__` functions called *before* the `__init__` of the enclosing module.

Two typical uses of `__init__` are calling runtime initialization functions of external C libraries and initializing global constants that involve pointers returned by external libraries. For example, suppose that we are calling a C library `libfoo` that requires us to call a `foo_init()` initialization function at runtime. Suppose that we also want to define a global constant `foo_data_ptr` that holds the return value of a `void *foo_data()` function defined by `libfoo` — this constant must be initialized at runtime (not at compile time) because the pointer address will change from run to run. You could accomplish this by defining the following `__init__` function in your module:

```
const foo_data_ptr = Ref{Ptr{Void}}{0}
function __init__()
    ccall(:foo_init, :libfoo, Void, ())
    foo_data_ptr[] = ccall(:foo_data, :libfoo, Ptr{Void}, ())
end
```

Notice that it is perfectly possible to define a global inside a function like `__init__`; this is one of the advantages of using a dynamic language. But by making it a constant at global scope, we can ensure that the type is known to the compiler and allow it to generate better optimized code. Obviously, any other globals in your module that depends on

`foo_data_ptr` would also have to be initialized in `__init__`.

Constants involving most Julia objects that are not produced by `ccall` do not need to be placed in `__init__`: their definitions can be precompiled and loaded from the cached module image. This includes complicated heap-allocated objects like arrays. However, any routine that returns a raw pointer value must be called at runtime for precompilation to work (Ptr objects will turn into null pointers unless they are hidden inside an `isbits` object). This includes the return values of the Julia functions `cfunction` and `pointer`.

Dictionary and set types, or in general anything that depends on the output of a `hash(key)` method, are a trickier case. In the common case where the keys are numbers, strings, symbols, ranges, `Expr`, or compositions of these types (via arrays, tuples, sets, pairs, etc.) they are safe to precompile. However, for a few other key types, such as `Function` or `DataType` and generic user-defined types where you haven't defined a `hash` method, the fallback `hash` method depends on the memory address of the object (via its `object_id`) and hence may change from run to run. If you have one of these key types, or if you aren't sure, to be safe you can initialize this dictionary from within your `__init__` function. Alternatively, you can use the `ObjectIdDict` dictionary type, which is specially handled by precompilation so that it is safe to initialize at compile-time.

When using precompilation, it is important to keep a clear sense of the distinction between the compilation phase and the execution phase. In this mode, it will often be much more clearly apparent that Julia is a compiler which allows execution of arbitrary Julia code, not a standalone interpreter that also generates compiled code.

Other known potential failure scenarios include:

1. Global counters (for example, for attempting to uniquely identify objects) Consider the following code snippet:

```
type UniquedById
    myid::Int
    let counter = 0
        UniquedById() = new(counter += 1)
    end
end
```

while the intent of this code was to give every instance a unique id, the counter value is recorded at the end of compilation. All subsequent usages of this incrementally compiled module will start from that same counter value.

Note that `object_id` (which works by hashing the memory pointer) has similar issues (see notes on `Dict` usage below).

One alternative is to store both `current_module()` and the current counter value, however, it may be better to redesign the code to not depend on this global state.

2. Associative collections (such as `Dict` and `Set`) need to be re-hashed in `__init__`. (In the future, a mechanism may be provided to register an initializer function.)
3. Depending on compile-time side-effects persisting through load-time. Example include: modifying arrays or other variables in other Julia modules; maintaining handles to open files or devices; storing pointers to other system resources (including memory);
4. Creating accidental “copies” of global state from another module, by referencing it directly instead of via its lookup path. For example, (in global scope):

```
#mystdout = Base.STDOUT #= will not work correctly, since this will copy Base.STDOUT into this m
# instead use accessor functions:
getstdout() = Base.STDOUT #= best option =#
# or move the assignment into the runtime:
__init__() = global mystdout = Base.STDOUT #= also works =#
```

Several additional restrictions are placed on the operations that can be done while precompiling code to help the user avoid other wrong-behavior situations:

1. Calling `eval` to cause a side-effect in another module. This will also cause a warning to be emitted when the incremental precompile flag is set.
2. `global const` statements from local scope after `__init__()` has been started (see issue #12010 for plans to add an error for this)
3. Replacing a module (or calling `workspace()`) is a runtime error while doing an incremental precompile.

A few other points to be aware of:

1. No code reload / cache invalidation is performed after changes are made to the source files themselves, (including by `Pkg.update`), and no cleanup is done after `Pkg.rm`
2. The memory sharing behavior of a reshaped array is disregarded by precompilation (each view gets its own copy)
3. Expecting the filesystem to be unchanged between compile-time and runtime e.g. `@__FILE__/source_path()` to find resources at runtime, or the `BinDeps @checked_lib` macro. Sometimes this is unavoidable. However, when possible, it can be good practice to copy resources into the module at compile-time so they won't need to be found at runtime.
4. `WeakRef` objects and finalizers are not currently handled properly by the serializer (this will be fixed in an upcoming release).
5. It is usually best to avoid capturing references to instances of internal metadata objects such as `Method`, `MethodInstance`, `MethodTable`, `TypeMapLevel`, `TypeMapEntry` and fields of those objects, as this can confuse the serializer and may not lead to the outcome you desire. It is not necessarily an error to do this, but you simply need to be prepared that the system will try to copy some of these and to create a single unique instance of others.

It is sometimes helpful during module development to turn off incremental precompilation. The command line flag `--compilecache={yes|no}` enables you to toggle module precompilation on and off. When Julia is started with `--compilecache=no` the serialized modules in the compile cache are ignored when loading modules and module dependencies. `Base.compilecache()` can still be called manually and it will respect `__precompile__()` directives for the module. The state of this command line flag is passed to `Pkg.build()` to disable automatic precompilation triggering when installing, updating, and explicitly building packages.

Documentation

Julia enables package developers and users to document functions, types and other objects easily via a built-in documentation system since Julia 0.4.

Tip: This documentation system can also be used in Julia 0.3 via the [Docile.jl](#) package; see the documentation for that package for more details.

The basic syntax is very simple: any string appearing at the top-level right before an object (function, macro, type or instance) will be interpreted as documenting it (these are called *docstrings*). Here is a very simple example:

```
"Tell whether there are too foo items in the array."
foo(xs::Array) = ...
```

Documentation is interpreted as [Markdown](#), so you can use indentation and code fences to delimit code examples from text. Technically, any object can be associated with any other as metadata; Markdown happens to be the default, but one can construct other string macros and pass them to the `@doc` macro just as well.

Here is a more complex example, still using Markdown:

```
"""
    bar(x[, y])

Compute the Bar index between `x` and `y`. If `y` is missing, compute
the Bar index between all pairs of columns of `x`.

# Examples
```julia
julia> bar([1, 2], [1, 2])
1
```
"""
function bar(x, y) ...
```

As in the example above, we recommend following some simple conventions when writing documentation:

1. Always show the signature of a function at the top of the documentation, with a four-space indent so that it is printed as Julia code.

This can be identical to the signature present in the Julia code (like `mean(x::AbstractArray)`), or a simplified form. Optional arguments should be represented with their default values (i.e. `f(x, y=1)`) when possible, following the actual Julia syntax. Optional arguments which do not have a default value should be put in brackets (i.e. `f(x[, y])` and `f(x[, y[, z]])`). An alternative solution is to use several lines: one without optional arguments, the other(s) with them. This solution can also be used to document several related

methods of a given function. When a function accepts many keyword arguments, only include a `<keyword arguments>` placeholder in the signature (i.e. `f(x; <keyword arguments>)`), and give the complete list under an `# Arguments` section (see point 4 below).

2. Include a single one-line sentence describing what the function does or what the object represents after the simplified signature block. If needed, provide more details in a second paragraph, after a blank line.

The one-line sentence should use the imperative form (“Do this”, “Return that”) instead of the third person (do not write “Returns the length...”) when documenting functions. It should end with a period. If the meaning of a function cannot be summarized easily, splitting it into separate composable parts could be beneficial (this should not be taken as an absolute requirement for every single case though).

3. Do not repeat yourself.

Since the function name is given by the signature, there is no need to start the documentation with “The function `bar...`”: go straight to the point. Similarly, if the signature specifies the types of the arguments, mentioning them in the description is redundant.

4. Only provide an argument list when really necessary.

For simple functions, it is often clearer to mention the role of the arguments directly in the description of the function’s purpose. An argument list would only repeat information already provided elsewhere. However, providing an argument list can be a good idea for complex functions with many arguments (in particular keyword arguments). In that case, insert it after the general description of the function, under an `# Arguments` header, with one `*` bullet for each argument. The list should mention the types and default values (if any) of the arguments:

```
"""
...
# Arguments
* `n::Integer`: the number of elements to compute.
* `dim::Integer=1`: the dimensions along which to perform the computation.
...
"""
```

5. Include any code examples in an `# Examples` section.

Examples should, whenever possible, be written as *doctests*. A *doctest* is a fenced code block (see [Code blocks](#) (page 157)) starting with ```jldoctest` and contains any number of `julia>` prompts together with inputs and expected outputs that mimic the Julia REPL.

For example in the following docstring a variable `a` is defined and the expected result, as printed in a Julia REPL, appears afterwards:

```
"""
Some nice documentation here.

# Examples

``jldoctest
julia> a = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4
...
"""
```

Warning: Calling `rand` and other RNG-related functions should be avoided in doctests since they will not produce consistent outputs during different Julia sessions.

Operating system word size (`Int32` or `Int64`) as well as path separator differences (`/` or `\`) will also effect the reproducibility of some doctests.

Note that whitespace in your doctest is significant! The doctest will fail if you misalign the output of pretty-printing an array, for example.

You can then run `make -C doc doctest` to run all the doctests in the Julia Manual, which will ensure that your example works.

Examples that are untestable should be written within fenced code blocks starting with ````julia` so that they are highlighted correctly in the generated documentation.

Tip: Wherever possible examples should be **self-contained** and **runnable** so that readers are able to try them out without having to include any dependencies.

6. Use backticks to identify code and equations.

Julia identifiers and code excerpts should always appear between backticks ``` to enable highlighting. Equations in the LaTeX syntax can be inserted between double backticks ````. Use Unicode characters rather than their LaTeX escape sequence, i.e. ```α = 1``` rather than ```\alpha = 1```.

7. Place the starting and ending `"""` characters on lines by themselves.

That is, write:

```
"""
...

...
"""
f(x, y) = ...
```

rather than:

```
"""...

..."""
f(x, y) = ...
```

This makes it more clear where docstrings start and end.

8. Respect the line length limit used in the surrounding code.

Docstrings are edited using the same tools as code. Therefore, the same conventions should apply. It is advised to add line breaks after 92 characters.

17.1 Accessing Documentation

Documentation can be accessed at the REPL or in IJulia by typing `?` followed by the name of a function or macro, and pressing `Enter`. For example,

```
?fft
?@time
?r"""
```

will bring up docs for the relevant function, macro or string macro respectively. In **Juno** using `Ctrl-J`, `Ctrl-D` will bring up documentation for the object under the cursor.

17.2 Functions & Methods

Functions in Julia may have multiple implementations, known as methods. While it's good practice for generic functions to have a single purpose, Julia allows methods to be documented individually if necessary. In general, only the most generic method should be documented, or even the function itself (i.e. the object created without any methods by `function` `bar` `end`). Specific methods should only be documented if their behaviour differs from the more generic ones. In any case, they should not repeat the information provided elsewhere. For example:

```
"""
Multiplication operator. `x*y*z*...` calls this function with multiple
arguments, i.e. `*(x,y,z...)` .
"""
function *(x, y)
    # ... [implementation sold separately] ...
end

"When applied to strings, concatenates them."
function *(x::AbstractString, y::AbstractString)
    # ... [insert secret sauce here] ...
end

help?>*
Multiplication operator. `x*y*z*...` calls this function with multiple
arguments, i.e. `*(x,y,z...)` .

When applied to strings, concatenates them.
```

When retrieving documentation for a generic function, the metadata for each method is concatenated with the `catdoc` function, which can of course be overridden for custom types.

17.3 Advanced Usage

The `@doc` macro associates its first argument with its second in a per-module dictionary called `META`. By default, documentation is expected to be written in Markdown, and the `doc""` string macro simply creates an object representing the Markdown content. In the future it is likely to do more advanced things such as allowing for relative image or link paths.

When used for retrieving documentation, the `@doc` macro (or equally, the `doc` function) will search all `META` dictionaries for metadata relevant to the given object and return it. The returned object (some Markdown content, for example) will by default display itself intelligently. This design also makes it easy to use the doc system in a programmatic way; for example, to re-use documentation between different versions of a function:

```
@doc "... " foo!
@doc (@doc foo!) foo
```

Or for use with Julia's metaprogramming functionality:

```
for (f, op) in ((:add, :+), (:subtract, :-), (:multiply, :*), (:divide, :/))
    @eval begin
        $f(a,b) = $op(a,b)
    end
end
```

```
end
@doc "`add(a,b)` adds `a` and `b` together" add
@doc "`subtract(a,b)` subtracts `b` from `a`" subtract
```

Documentation written in non-toplevel blocks, such as `if`, `for`, and `let`, are not automatically added to the documentation system. `@doc` must be used in these cases. For example:

```
if VERSION > v"0.4"
    "... "
    f(x) = x
end
```

will not add any documentation to `f` even when the condition is `true` and must instead be written as:

```
if VERSION > v"0.4"
    @doc "... " ->
    f(x) = x
end
```

17.4 Syntax Guide

A comprehensive overview of all documentable Julia syntax.

In the following examples `"... "` is used to illustrate an arbitrary docstring which may be one of the follow four variants and contain arbitrary text:

```
"... "
doc"... "

"""
...
"""

doc"""
...
"""
```

`@doc_str` should only be used when the docstring contains `$` or `\` characters that should not be parsed by Julia such as LaTeX syntax or Julia source code examples containing interpolation.

17.4.1 Functions and Methods

```
"... "
function f end

"... "
f
```

Adds docstring `"... "` to Function `f`. The first version is the preferred syntax, however both are equivalent.

```
"... "
f(x) = x

"... "
function f(x)
```

```
    x
end

"... "
f(x)
```

Adds docstring "... " to Method `f (::Any)`.

```
"..."
f(x, y = 1) = x + y
```

Adds docstring "... " to two Methods, namely `f (::Any)` and `f (::Any, ::Any)`.

17.4.2 Macros

```
"..."
macro m(x) end
```

Adds docstring "... " to the `@m (::Any)` macro definition.

```
"..."
: (@m)
```

Adds docstring "... " to the macro named `@m`.

17.4.3 Types

```
"..."
abstract T1

"... "
type T2
    ...
end

"... "
immutable T3
    ...
end
```

Adds the docstring "... " to types `T1`, `T2`, and `T3`.

```
"..."
type T
    "x"
    x
    "y"
    y
end
```

Adds docstring "... " to type `T`, "`x`" to field `T.x` and "`y`" to field `T.y`. Also applicable to `immutable` types.

```
"..."
typealias A T
```

Adds docstring "... " to the Binding `A`.

Bindings are used to store a reference to a particular `Symbol` in a `Module` without storing the referenced value itself.

17.4.4 Modules

```
"..."
module M end

module M

"..."
M

end
```

Adds docstring "..." to the Module M. Adding the docstring above the Module is the preferred syntax, however both are equivalent.

```
"..."
baremodule M
# ...
end

baremodule M

import Base: @doc

"..."
f(x) = x

end
```

Documenting a baremodule by placing a docstring above the expression automatically imports @doc into the module. These imports must be done manually when the module expression is not documented. Empty baremodules cannot be documented.

17.4.5 Global Variables

```
"..."
const a = 1

"..."
b = 2

"..."
global c = 3
```

Adds docstring "..." to the Bindings a, b, and c.

Note: When a `const` definition is only used to define an alias of another definition, such as is the case with the function `div` and its alias `÷` in `Base`, do not document the alias and instead document the actual function.

If the alias is documented and not the real definition then the docsystem (`? mode`) will not return the docstring attached to the alias when the real definition is searched for.

For example you should write

```
"..."
f(x) = x + 1
const alias = f
```

rather than

```
f(x) = x + 1
"..."
const alias = f
```

```
"..."
sym
```

Adds docstring "..." to the value associated with `sym`. Users should prefer documenting `sym` at its definition.

17.4.6 Multiple Objects

```
"..."
a, b
```

Adds docstring "..." to `a` and `b` each of which should be a documentable expression. This syntax is equivalent to

```
"..."
a

"..."
b
```

Any number of expressions may be documented together in this way. This syntax can be useful when two functions are related, such as non-mutating and mutating versions `f` and `f!`.

17.4.7 Macro-generated code

```
"..."
@m expression
```

Adds docstring "..." to expression generated by expanding `@m expression`. This allows for expressions decorated with `@inline`, `@noinline`, `@generated`, or any other macro to be documented in the same way as undecorated expressions.

Macro authors should take note that only macros that generate a single expression will automatically support docstrings. If a macro returns a block containing multiple subexpressions then the subexpression that should be documented must be marked using the `@__doc__()` (page 154) macro.

The `@enum` macro makes use of `@__doc__` to allow for documenting Enums. Examining its definition should serve as an example of how to use `@__doc__` correctly.

@__doc__(ex)

Low-level macro used to mark expressions returned by a macro that should be documented. If more than one expression is marked then the same docstring is applied to each expression.

```
macro example(f)
    quote
        $(f)() = 0
        @__doc__ $(f)(x) = 1
        $(f)(x, y) = 2
    end
end
```

`@__doc__` has no effect when a macro that uses it is not documented.

17.5 Markdown syntax

The following markdown syntax is supported in Julia.

17.5.1 Inline elements

Here “inline” refers to elements that can be found within blocks of text, i.e. paragraphs. These include the following elements.

Bold

Surround words with two asterisks, `**`, to display the enclosed text in boldface.

```
A paragraph containing a bold word.
```

Italics

Surround words with one asterisk, `*`, to display the enclosed text in italics.

```
A paragraph containing an emphasised word.
```

Literals

Surround text that should be displayed exactly as written with single backticks, ```.

```
A paragraph containing a literal word.
```

Literals should be used when writing text that refers to names of variables, functions, or other parts of a Julia program.

Tip: To include a backtick character within literal text use three backticks rather than one to enclose the text.

```
A paragraph containing a ``backtick` character ``.
```

By extension any odd number of backticks may be used to enclose a lesser number of backticks.

L^AT_EX

Surround text that should be displayed as mathematics using L^AT_EX syntax with double backticks, ````.

```
A paragraph containing some ``\LaTeX`` markup.
```

Tip: As with literals in the previous section, if literal backticks need to be written within double backticks use an even number greater than two. Note that if a single literal backtick needs to be included within L^AT_EX markup then two enclosing backticks is sufficient.

Links

Links to either external or internal addresses can be written using the following syntax, where the text enclosed in square brackets, [], is the name of the link and the text enclosed in parentheses, (), is the URL.

```
A paragraph containing a link to [Julia] (http://www.julialang.org).
```

Footnote references

Named and numbered footnote references can be written using the following syntax. A footnote name must be a single alphanumeric word containing no punctuation.

```
A paragraph containing a numbered footnote [^1] and a named one [^named].
```

Note: The text associated with a footnote can be written anywhere within the same page as the footnote reference. The syntax used to define the footnote text is discussed in the *Footnotes* (page 159) section below.

17.5.2 Toplevel elements

The following elements can be written either at the “toplevel” of a document or within another “toplevel” element.

Paragraphs

A paragraph is a block of plain text, possibly containing any number of inline elements defined in the *Inline elements* (page 155) section above, with one or more blank lines above and below it.

```
This is a paragraph.

And this is *another* one containing some emphasised text.
A new line, but still part of the same paragraph.
```

Headers

A document can be split up into different sections using headers. Headers use the following syntax:

```
# Level One
## Level Two
### Level Three
#### Level Four
##### Level Five
##### Level Six
```

A header line can contain any inline syntax in the same way as a paragraph can.

Tip: Try to avoid using too many levels of header within a single document. A heavily nested document may be indicative of a need to restructure it or split it into several pages covering separate topics.

Code blocks

Source code can be displayed as a literal block using an indent of four spaces as shown in the following example.

```
This is a paragraph.

    function func(x)
        # ...
    end

Another paragraph.
```

Additionally, code blocks can be enclosed using triple backticks with an optional “language” to specify how a block of code should be highlighted.

```
A code block without a "language":

...
function func(x)
    # ...
end
...

and another one with the "language" specified as `julia`:

```julia
function func(x)
 # ...
end
```
```

Note: “Fenced” code blocks, as shown in the last example, should be preferred over indented code blocks since there is no way to specify what language an indented code block is written in.

Block quotes

Text from external sources, such as quotations from books or websites, can be quoted using > characters prepended to each line of the quote as follows.

```
Here's a quote:

> Julia is a high-level, high-performance dynamic programming language for
> technical computing, with syntax that is familiar to users of other
> technical computing environments.
```

Note that a single space must appear after the > character on each line. Quoted blocks may themselves contain other toplevel or inline elements.

Images

The syntax for images is similar to the link syntax mentioned above. Prepending a ! character to a link will display an image from the specified URL rather than a link to it.

```
![alternative text](link/to/image.png)
```

Lists

Unordered lists can be written by prepending each item in a list with either `*`, `+`, or `-`.

```
A list of items:
```

```
* item one
* item two
* item three
```

Note the two spaces before each `*` and the single space after each one.

Lists can contain other nested toplevel elements such as lists, code blocks, or quoteblocks. A blank line should be left between each list item when including any toplevel elements within a list.

```
Another list:
```

```
* item one

* item two
  ``
  f(x) = x
  ``

* And a sublist:
    + sub-item one
    + sub-item two
```

Note: The contents of each item in the list must line up with the first line of the item. In the above example the fenced code block must be indented by four spaces to align with the `i` in `item two`.

Ordered lists are written by replacing the “bullet” character, either `*`, `+`, or `-`, with a positive integer followed by either `.` or `)`.

```
Two ordered lists:
```

```
1. item one
2. item two
3. item three

5) item five
6) item six
7) item seven
```

An ordered list may start from a number other than one, as in the second list of the above example, where it is numbered from five. As with unordered lists, ordered lists can contain nested toplevel elements.

Display equations

Large \LaTeX equations that do not fit inline within a paragraph may be written as display equations using a fenced code block with the “`language`” `math` as in the example below.

```
```math
f(a) = \frac{1}{2\pi} \int_0^{2\pi} (\alpha + R \cos(\theta)) d\theta
```
```

Footnotes

This syntax is paired with the inline syntax for *Footnote references* (page 156). Make sure to read that section as well.

Footnote text is defined using the following syntax, which is similar to footnote reference syntax, aside from the `:` character that is appended to the footnote label.

```
[^1]: Numbered footnote text.

[^note]:

    Named footnote text containing several toplevel elements.

    * item one
    * item two
    * item three

```julia
function func(x)
 # ...
end
```
```

Note: No checks are done during parsing to make sure that all footnote references have matching footnotes.

Horizontal rules

The equivalent of an `<hr>` HTML tag can be written using the following syntax:

```
Text above the line.

---

And text below the line.
```

Tables

Basic tables can be written using the syntax described below. Note that markdown tables have limited features and cannot contain nested toplevel elements unlike other elements discussed above – only inline elements are allowed. Tables must always contain a header row with column names. Cells cannot span multiple rows or columns of the table.

| | | | |
|------------|------------|--------------|--|
| Column One | Column Two | Column Three | |
| :----- | ----- | :----- | |
| Row `1` | Column `2` | | |
| *Row* 2 | **Row** 2 | Column ``3`` | |

Note: As illustrated in the above example each column of `|` characters must be aligned vertically.

A : character on either end of a column's header separator (the row containing – characters) specifies whether the row is left-aligned, right-aligned, or (when : appears on both ends) center-aligned. Providing no : characters will default to right-aligning the column.

Admonitions

Specially formatted blocks with titles such as “Notes”, “Warning”, or “Tips” are known as admonitions and are used when some part of a document needs special attention. They can be defined using the following !!! syntax:

```
!!! note

    This is the content of the note.

!!! warning "Beware!"

    And this is another one.

    This warning admonition has a custom title: `"Beware!"`.
```

Admonitions, like most other toplevel elements, can contain other toplevel elements. When no title text, specified after the admonition type in double quotes, is included then the title used will be the type of the block, i.e. "Note" in the case of the `note` admonition.

17.6 Markdown Syntax Extensions

Julia's markdown supports interpolation in a very similar way to basic string literals, with the difference that it will store the object itself in the Markdown tree (as opposed to converting it to a string). When the Markdown content is rendered the usual `show` methods will be called, and these can be overridden as usual. This design allows the Markdown to be extended with arbitrarily complex features (such as references) without cluttering the basic syntax.

In principle, the Markdown parser itself can also be arbitrarily extended by packages, or an entirely custom flavour of Markdown can be used, but this should generally be unnecessary.

Metaprogramming

The strongest legacy of Lisp in the Julia language is its metaprogramming support. Like Lisp, Julia represents its own code as a data structure of the language itself. Since code is represented by objects that can be created and manipulated from within the language, it is possible for a program to transform and generate its own code. This allows sophisticated code generation without extra build steps, and also allows true Lisp-style macros operating at the level of [abstract syntax trees](#). In contrast, preprocessor “macro” systems, like that of C and C++, perform textual manipulation and substitution before any actual parsing or interpretation occurs. Because all data types and code in Julia are represented by Julia data structures, powerful [reflection](#) capabilities are available to explore the internals of a program and its types just like any other data.

18.1 Program representation

Every Julia program starts life as a string:

```
julia> prog = "1 + 1"
"1 + 1"
```

What happens next?

The next step is to [parse](#) each string into an object called an expression, represented by the Julia type `Expr`:

```
julia> ex1 = parse(prog)
:(1 + 1)

julia> typeof(ex1)
Expr
```

`Expr` objects contain three parts:

- a `Symbol` identifying the kind of expression. A symbol is an [interned string](#) identifier (more discussion below).

```
julia> ex1.head
:call
```

- the expression arguments, which may be symbols, other expressions, or literal values:

```
julia> ex1.args
3-element Array{Any,1}:
 :+
 1
 1
```

- finally, the expression result type, which may be annotated by the user or inferred by the compiler (and may be ignored completely for the purposes of this chapter):

```
julia> ex1.typ
Any
```

Expressions may also be constructed directly in [prefix notation](#):

```
julia> ex2 = Expr(:call, :+, 1, 1)
:(1 + 1)
```

The two expressions constructed above – by parsing and by direct construction – are equivalent:

```
julia> ex1 == ex2
true
```

The key point here is that Julia code is internally represented as a data structure that is accessible from the language itself.

The [dump\(\)](#) (page 567) function provides indented and annotated display of Expr objects:

```
julia> dump(ex2)
Expr
  head: Symbol call
  args: Array{Any}((3,))
    1: Symbol +
    2: Int64 1
    3: Int64 1
  typ: Any
```

Expr objects may also be nested:

```
julia> ex3 = parse("(4 + 4) / 2")
:((4 + 4) / 2)
```

Another way to view expressions is with `Meta.show_sexpr`, which displays the [S-expression](#) form of a given Expr, which may look very familiar to users of Lisp. Here's an example illustrating the display on a nested Expr:

```
julia> Meta.show_sexpr(ex3)
(:call, :/, (:call, :+, 4, 4), 2)
```

18.1.1 Symbols

The `:` character has two syntactic purposes in Julia. The first form creates a [Symbol](#) (page 450), an [interned string](#) used as one building-block of expressions:

```
julia> :foo
:foo

julia> typeof(ans)
Symbol
```

The [Symbol](#) (page 450) constructor takes any number of arguments and creates a new symbol by concatenating their string representations together:

```
julia> :foo == Symbol("foo")
true

julia> Symbol("func", 10)
```



```
:func10

julia> Symbol(:var, '_', "sym")
:var_sym
```

In the context of an expression, symbols are used to indicate access to variables; when an expression is evaluated, a symbol is replaced with the value bound to that symbol in the appropriate *scope* (page 81).

Sometimes extra parentheses around the argument to `:` are needed to avoid ambiguity in parsing.:

```
julia> :(:)
:(:)

julia> :(::)
:(::)
```

18.2 Expressions and evaluation

18.2.1 Quoting

The second syntactic purpose of the `:` character is to create expression objects without using the explicit `Expr` constructor. This is referred to as *quoting*. The `:` character, followed by paired parentheses around a single statement of Julia code, produces an `Expr` object based on the enclosed code. Here is example of the short form used to quote an arithmetic expression:

```
julia> ex = :(a+b*c+1)
:(a + b * c + 1)

julia> typeof(ex)
Expr
```

(to view the structure of this expression, try `ex.head` and `ex.args`, or use `dump()` (page 567) as above)

Note that equivalent expressions may be constructed using `parse()` (page 369) or the direct `Expr` form:

```
julia>      :(a + b*c + 1) ==
      parse("a + b*c + 1") ==
      Expr(:call, :+, :a, Expr(:call, :*, :b, :c), 1)
true
```

Expressions provided by the parser generally only have symbols, other expressions, and literal values as their args, whereas expressions constructed by Julia code can have arbitrary run-time values without literal forms as args. In this specific example, `+` and `a` are symbols, `*(b, c)` is a subexpression, and `1` is a literal 64-bit signed integer.

There is a second syntactic form of quoting for multiple expressions: blocks of code enclosed in `quote ... end`. Note that this form introduces `QuoteNode` elements to the expression tree, which must be considered when directly manipulating an expression tree generated from `quote` blocks. For other purposes, `:(...)` and `quote ... end` blocks are treated identically.

```
julia> ex = quote
      x = 1
      y = 2
      x + y
    end
quote # none, line 2:
  x = 1 # none, line 3:
  y = 2 # none, line 4:
```

```
x + y
end

julia> typeof(ex)
Expr
```

18.2.2 Interpolation

Direct construction of `Expr` objects with value arguments is powerful, but `Expr` constructors can be tedious compared to “normal” Julia syntax. As an alternative, Julia allows “splicing” or interpolation of literals or expressions into quoted expressions. Interpolation is indicated by the `$` prefix.

In this example, the literal value of `a` is interpolated:

```
julia> a = 1;

julia> ex = :($a + b)
:(1 + b)
```

Interpolating into an unquoted expression is not supported and will cause a compile-time error:

```
julia> $a + b
ERROR: unsupported or misplaced expression $
...
```

In this example, the tuple `(1, 2, 3)` is interpolated as an expression into a conditional test:

```
julia> ex = :(a in $((1,2,3)) )
:(a in (1,2,3))
```

Interpolating symbols into a nested expression requires enclosing each symbol in an enclosing quote block:

```
julia> :( :a in $( :(:a + :b) ) )
          ^^^^^^^^^^
          quoted inner expression
```

The use of `$` for expression interpolation is intentionally reminiscent of [string interpolation](#) (page 46) and [command interpolation](#) (page 232). Expression interpolation allows convenient, readable programmatic construction of complex Julia expressions.

18.2.3 `eval()` and effects

Given an expression object, one can cause Julia to evaluate (execute) it at global scope using `eval()` (page 368):

```
julia> :(1 + 2)
:(1 + 2)

julia> eval(ans)
3

julia> ex = :(a + b)
:(a + b)

julia> eval(ex)
ERROR: UndefVarError: b not defined
...
```

```
julia> a = 1; b = 2;

julia> eval(ex)
3
```

Every *module* (page 139) has its own `eval()` (page 368) function that evaluates expressions in its global scope. Expressions passed to `eval()` (page 368) are not limited to returning values — they can also have side-effects that alter the state of the enclosing module’s environment:

```
julia> ex = :(x = 1)
:(x = 1)

julia> x
ERROR: UndefVarError: x not defined
...

julia> eval(ex)
1

julia> x
1
```

Here, the evaluation of an expression object causes a value to be assigned to the global variable `x`.

Since expressions are just `Expr` objects which can be constructed programmatically and then evaluated, it is possible to dynamically generate arbitrary code which can then be run using `eval()` (page 368). Here is a simple example:

```
julia> a = 1;

julia> ex = Expr(:call, :+, a, :b)
:(1 + b)

julia> a = 0; b = 2;

julia> eval(ex)
3
```

The value of `a` is used to construct the expression `ex` which applies the `+` function to the value `1` and the variable `b`. Note the important distinction between the way `a` and `b` are used:

- The value of the *variable* `a` at expression construction time is used as an immediate value in the expression. Thus, the value of `a` when the expression is evaluated no longer matters: the value in the expression is already `1`, independent of whatever the value of `a` might be.
- On the other hand, the *symbol* `:b` is used in the expression construction, so the value of the variable `b` at that time is irrelevant — `:b` is just a symbol and the variable `b` need not even be defined. At expression evaluation time, however, the value of the symbol `:b` is resolved by looking up the value of the variable `b`.

18.2.4 Functions on Expressions

As hinted above, one extremely useful feature of Julia is the capability to generate and manipulate Julia code within Julia itself. We have already seen one example of a function returning `Expr` objects: the `parse()` (page 369) function, which takes a string of Julia code and returns the corresponding `Expr`. A function can also take one or more `Expr` objects as arguments, and return another `Expr`. Here is a simple, motivating example:

```
julia> function math_expr(op, op1, op2)
    expr = Expr(:call, op, op1, op2)
    return expr
end
```

```

end

julia> ex = math_expr(:+, 1, Expr(:call, :*, 4, 5))
:(1 + 4*5)

julia> eval(ex)
21

```

As another example, here is a function that doubles any numeric argument, but leaves expressions alone:

```

julia> function make_expr2(op, opr1, opr2)
    opr1f, opr2f = map(x -> isa(x, Number) ? 2*x : x, (opr1, opr2))
    retexpr = Expr(:call, op, opr1f, opr2f)

    return retexpr
end
make_expr2 (generic function with 1 method)

julia> make_expr2(:+, 1, 2)
:(2 + 4)

julia> ex = make_expr2(:+, 1, Expr(:call, :*, 5, 8))
:(2 + 5 * 8)

julia> eval(ex)
42

```

18.3 Macros

Macros provide a method to include generated code in the final body of a program. A macro maps a tuple of arguments to a returned *expression*, and the resulting expression is compiled directly rather than requiring a runtime `eval()` (page 368) call. Macro arguments may include expressions, literal values, and symbols.

18.3.1 Basics

Here is an extraordinarily simple macro:

```

macro sayhello()
    return :( println("Hello, world!") )
end

```

Macros have a dedicated character in Julia's syntax: the @ (at-sign), followed by the unique name declared in a macro NAME ... end block. In this example, the compiler will replace all instances of @sayhello with:

```

:( println("Hello, world!") )

```

When @sayhello is given at the REPL, the expression executes immediately, thus we only see the evaluation result:

```

julia> @sayhello()
"Hello, world!"

```

Now, consider a slightly more complex macro:

```

julia> macro sayhello(name)
    return :( println("Hello, ", $name) )
end

```

This macro takes one argument: `name`. When `@sayhello` is encountered, the quoted expression is *expanded* to interpolate the value of the argument into the final expression:

```
julia> @sayhello("human")
Hello, human
```

We can view the quoted return expression using the function `macroexpand()` (page 378) (**important note:** this is an extremely useful tool for debugging macros):

```
julia> ex = macroexpand( :(@sayhello("human")) )
:(println("Hello, ", "human"))
      ^^^^^^
      interpolated: now a literal string

julia> typeof(ex)
Expr
```

18.3.2 Hold up: why macros?

We have already seen a function `f (::Expr...) -> Expr` in a previous section. In fact, `macroexpand()` (page 378) is also such a function. So, why do macros exist?

Macros are necessary because they execute when code is parsed, therefore, macros allow the programmer to generate and include fragments of customized code *before* the full program is run. To illustrate the difference, consider the following example:

```
julia> macro twostep(arg)
    println("I execute at parse time. The argument is: ", arg)

    return :(println("I execute at runtime. The argument is: ", $arg))
end

julia> ex = macroexpand( :(@twostep (1, 2, 3)) );
I execute at parse time. The argument is: ((1,2,3))
```

The first call to `println()` (page 566) is executed when `macroexpand()` (page 378) is called. The resulting expression contains *only* the second `println`:

```
julia> typeof(ex)
Expr

julia> ex
:(println("I execute at runtime. The argument is: ", $(Expr(:copyast, :((1,2,3)))))

julia> eval(ex)
I execute at runtime. The argument is: (1,2,3)
```

18.3.3 Macro invocation

Macros are invoked with the following general syntax:

```
@name expr1 expr2 ...
@name(expr1, expr2, ...)
```

Note the distinguishing `@` before the macro name and the lack of commas between the argument expressions in the first form, and the lack of whitespace after `@name` in the second form. The two styles should not be mixed. For

example, the following syntax is different from the examples above; it passes the tuple `(expr1, expr2, ...)` as one argument to the macro:

```
@name (expr1, expr2, ...)
```

It is important to emphasize that macros receive their arguments as expressions, literals, or symbols. One way to explore macro arguments is to call the `show()` (page 566) function within the macro body:

```
julia> macro showarg(x)
    show(x)
    # ... remainder of macro, returning an expression
end

julia> @showarg(a)
(:a,)

julia> @showarg(1+1)
:(1 + 1)

julia> @showarg(println("Yo!"))
:(println("Yo!"))
```

18.3.4 Building an advanced macro

Here is a simplified definition of Julia's `@assert` macro:

```
macro assert(ex)
    return :( $ex ? nothing : throw(AssertionError($(string(ex)))) )
end
```

This macro can be used like this:

```
julia> @assert 1==1.0

julia> @assert 1==0
ERROR: AssertionError: 1 == 0
...
```

In place of the written syntax, the macro call is expanded at parse time to its returned result. This is equivalent to writing:

```
1==1.0 ? nothing : throw(AssertionError("1==1.0"))
1==0 ? nothing : throw(AssertionError("1==0"))
```

That is, in the first call, the expression `:(1==1.0)` is spliced into the test condition slot, while the value of `string(:(1==1.0))` is spliced into the assertion message slot. The entire expression, thus constructed, is placed into the syntax tree where the `@assert` macro call occurs. Then at execution time, if the test expression evaluates to true, then `nothing` is returned, whereas if the test is false, an error is raised indicating the asserted expression that was false. Notice that it would not be possible to write this as a function, since only the *value* of the condition is available and it would be impossible to display the expression that computed it in the error message.

The actual definition of `@assert` in the standard library is more complicated. It allows the user to optionally specify their own error message, instead of just printing the failed expression. Just like in functions with a variable number of arguments, this is specified with an ellipsis following the last argument:

```
macro assert(ex, msgs...)
    msg_body = isempty(msgs) ? ex : msgs[1]
end
```

```

msg = string(msg_body)
return :($ex ? nothing : throw(AssertionError($msg)))
end

```

Now `@assert` has two modes of operation, depending upon the number of arguments it receives! If there's only one argument, the tuple of expressions captured by `msgs` will be empty and it will behave the same as the simpler definition above. But now if the user specifies a second argument, it is printed in the message body instead of the failing expression. You can inspect the result of a macro expansion with the aptly named `macroexpand()` (page 378) function:

```

julia> macroexpand(:(@assert a==b))
:(if a == b
    nothing
else
    (Base.throw) (Base.Main.Base.AssertionError("a == b"))
end)

julia> macroexpand(:(@assert a==b "a should equal b!"))
:(if a == b
    nothing
else
    (Base.throw) (Base.Main.Base.AssertionError("a should equal b!"))
end)

```

There is yet another case that the actual `@assert` macro handles: what if, in addition to printing “a should equal b,” we wanted to print their values? One might naively try to use string interpolation in the custom message, e.g., `@assert a==b "a ($a) should equal b ($b)!"`, but this won't work as expected with the above macro. Can you see why? Recall from *string interpolation* (page 46) that an interpolated string is rewritten to a call to `string()` (page 443). Compare:

```

julia> typeof(:("a should equal b"))
String

julia> typeof(:("a ($a) should equal b ($b)!"))
Expr

julia> dump(:("a ($a) should equal b ($b)!"))
Expr
 head: Symbol string
 args: Array{Any}((5,))
  1: String "a ("
  2: Symbol a
  3: String ") should equal b ("
  4: Symbol b
  5: String ")!"
 typ: Any

```

So now instead of getting a plain string in `msg_body`, the macro is receiving a full expression that will need to be evaluated in order to display as expected. This can be spliced directly into the returned expression as an argument to the `string()` (page 443) call; see `error.jl` for the complete implementation.

The `@assert` macro makes great use of splicing into quoted expressions to simplify the manipulation of expressions inside the macro body.

18.3.5 Hygiene

An issue that arises in more complex macros is that of *hygiene*. In short, macros must ensure that the variables they introduce in their returned expressions do not accidentally clash with existing variables in the surrounding code they expand into. Conversely, the expressions that are passed into a macro as arguments are often *expected* to evaluate in the context of the surrounding code, interacting with and modifying the existing variables. Another concern arises from the fact that a macro may be called in a different module from where it was defined. In this case we need to ensure that all global variables are resolved to the correct module. Julia already has a major advantage over languages with textual macro expansion (like C) in that it only needs to consider the returned expression. All the other variables (such as `msg` in `@assert` above) follow the *normal scoping block behavior* (page 81).

To demonstrate these issues, let us consider writing a `@time` macro that takes an expression as its argument, records the time, evaluates the expression, records the time again, prints the difference between the before and after times, and then has the value of the expression as its final value. The macro might look like this:

```
macro time(ex)
    return quote
        local t0 = time()
        local val = $ex
        local t1 = time()
        println("elapsed time: ", t1-t0, " seconds")
        val
    end
end
```

Here, we want `t0`, `t1`, and `val` to be private temporary variables, and we want `time` to refer to the `time()` (page 372) function in the standard library, not to any `time` variable the user might have (the same applies to `println`). Imagine the problems that could occur if the user expression `ex` also contained assignments to a variable called `t0`, or defined its own `time` variable. We might get errors, or mysteriously incorrect behavior.

Julia's macro expander solves these problems in the following way. First, variables within a macro result are classified as either local or global. A variable is considered local if it is assigned to (and not declared global), declared local, or used as a function argument name. Otherwise, it is considered global. Local variables are then renamed to be unique (using the `gensym()` (page 368) function, which generates new symbols), and global variables are resolved within the macro definition environment. Therefore both of the above concerns are handled; the macro's locals will not conflict with any user variables, and `time` and `println` will refer to the standard library definitions.

One problem remains however. Consider the following use of this macro:

```
module MyModule
import Base.@time

time() = ... # compute something

@time time()
end
```

Here the user expression `ex` is a call to `time`, but not the same `time` function that the macro uses. It clearly refers to `MyModule.time`. Therefore we must arrange for the code in `ex` to be resolved in the macro call environment. This is done by “escaping” the expression with `esc()` (page 368):

```
macro time(ex)
    ...
    local val = $(esc(ex))
    ...
end
```

An expression wrapped in this manner is left alone by the macro expander and simply pasted into the output verbatim. Therefore it will be resolved in the macro call environment.

This escaping mechanism can be used to “violate” hygiene when necessary, in order to introduce or manipulate user variables. For example, the following macro sets `x` to zero in the call environment:

```
macro zerox()
    return esc(:(x = 0))
end

function foo()
    x = 1
    @zerox
    x # is zero
end
```

This kind of manipulation of variables should be used judiciously, but is occasionally quite handy.

18.4 Code Generation

When a significant amount of repetitive boilerplate code is required, it is common to generate it programmatically to avoid redundancy. In most languages, this requires an extra build step, and a separate program to generate the repetitive code. In Julia, expression interpolation and `eval()` (page 368) allow such code generation to take place in the normal course of program execution. For example, the following code defines a series of operators on three arguments in terms of their 2-argument forms:

```
for op = (:+, :*, :&, :|, :$)
    eval(quote
        ($op) (a,b,c) = ($op) (($op) (a,b), c)
    end)
end
```

In this manner, Julia acts as its own **preprocessor**, and allows code generation from inside the language. The above code could be written slightly more tersely using the `:` prefix quoting form:

```
for op = (:+, :*, :&, :|, :$)
    eval(: (($op) (a,b,c) = ($op) (($op) (a,b), c)))
end
```

This sort of in-language code generation, however, using the `eval(quote(...))` pattern, is common enough that Julia comes with a macro to abbreviate this pattern:

```
for op = (:+, :*, :&, :|, :$)
    @eval ($op) (a,b,c) = ($op) (($op) (a,b), c)
end
```

The `@eval` (page 368) macro rewrites this call to be precisely equivalent to the above longer versions. For longer blocks of generated code, the expression argument given to `@eval` (page 368) can be a block:

```
@eval begin
    # multiple lines
end
```

18.5 Non-Standard String Literals

Recall from *Strings* (page 49) that string literals prefixed by an identifier are called non-standard string literals, and can have different semantics than un-prefixed string literals. For example:

- `r"^\s*(?:#|\$)"` produces a regular expression object rather than a string
- `b"DATA\xff\u2200"` is a byte array literal for `[68, 65, 84, 65, 255, 226, 136, 128]`.

Perhaps surprisingly, these behaviors are not hard-coded into the Julia parser or compiler. Instead, they are custom behaviors provided by a general mechanism that anyone can use: prefixed string literals are parsed as calls to specially-named macros. For example, the regular expression macro is just the following:

```
macro r_str(p)
    Regex(p)
end
```

That's all. This macro says that the literal contents of the string literal `r"^\s*(?:#|\$)"` should be passed to the `@r_str` macro and the result of that expansion should be placed in the syntax tree where the string literal occurs. In other words, the expression `r"^\s*(?:#|\$)"` is equivalent to placing the following object directly into the syntax tree:

```
Regex("^\s*(?:#|\\$)")
```

Not only is the string literal form shorter and far more convenient, but it is also more efficient: since the regular expression is compiled and the `Regex` object is actually created *when the code is compiled*, the compilation occurs only once, rather than every time the code is executed. Consider if the regular expression occurs in a loop:

```
for line = lines
    m = match(r"^\s*(?:#|\$)", line)
    if m == nothing
        # non-comment
    else
        # comment
    end
end
```

Since the regular expression `r"^\s*(?:#|\$)"` is compiled and inserted into the syntax tree when this code is parsed, the expression is only compiled once instead of each time the loop is executed. In order to accomplish this without macros, one would have to write this loop like this:

```
re = Regex("^\s*(?:#|\\$)")
for line = lines
    m = match(re, line)
    if m == nothing
        # non-comment
    else
        # comment
    end
end
```

Moreover, if the compiler could not determine that the `regex` object was constant over all loops, certain optimizations might not be possible, making this version still less efficient than the more convenient literal form above. Of course, there are still situations where the non-literal form is more convenient: if one needs to interpolate a variable into the regular expression, one must take this more verbose approach; in cases where the regular expression pattern itself is dynamic, potentially changing upon each loop iteration, a new regular expression object must be constructed on each iteration. In the vast majority of use cases, however, regular expressions are not constructed based on run-time data. In this majority of cases, the ability to write regular expressions as compile-time values is invaluable.

Like non-standard string literals, non-standard command literals exist using a prefixed variant of the command literal syntax. The command literal `custom`literal`` is parsed as `@custom_cmd "literal"`. Julia itself does not contain any non-standard command literals, but packages can make use of this syntax. Aside from the different syntax and the `_cmd` suffix instead of the `_str` suffix, non-standard command literals behave exactly like non-standard string literals.

In the event that two modules provide non-standard string or command literals with the same name, it is possible to qualify the string or command literal with a module name. For instance, if both `Foo` and `Bar` provide non-standard string literal `@x_str`, then one can write `Foo.x"literal"` or `Bar.x"literal"` to disambiguate between the two.

The mechanism for user-defined string literals is deeply, profoundly powerful. Not only are Julia's non-standard literals implemented using it, but also the command literal syntax (``echo "Hello, $person" ``) is implemented with the following innocuous-looking macro:

```
macro cmd(str)
    : (cmd_gen($(shell_parse(str)[1])))
end
```

Of course, a large amount of complexity is hidden in the functions used in this macro definition, but they are just functions, written entirely in Julia. You can read their source and see precisely what they do — and all they do is construct expression objects to be inserted into your program's syntax tree.

18.6 Generated functions

A very special macro is `@generated`, which allows you to define so-called *generated functions*. These have the capability to generate specialized code depending on the types of their arguments with more flexibility and/or less code than what can be achieved with multiple dispatch. While macros work with expressions at parsing-time and cannot access the types of their inputs, a generated function gets expanded at a time when the types of the arguments are known, but the function is not yet compiled.

Instead of performing some calculation or action, a generated function declaration returns a quoted expression which then forms the body for the method corresponding to the types of the arguments. When called, the body expression is first evaluated and compiled, then the returned expression is compiled and run. To make this efficient, the result is often cached. And to make this inferable, only a limited subset of the language is usable. Thus, generated functions provide a flexible framework to move work from run-time to compile-time, at the expense of greater restrictions on the allowable constructs.

When defining generated functions, there are four main differences to ordinary functions:

1. You annotate the function declaration with the `@generated` macro. This adds some information to the AST that lets the compiler know that this is a generated function.
2. In the body of the generated function you only have access to the *types* of the arguments – not their values – and any function that was defined *before* the definition of the generated function.
3. Instead of calculating something or performing some action, you return a *quoted expression* which, when evaluated, does what you want.
4. Generated functions must not have any side-effects or examine any non-constant global state (including, for example, IO, locks, or non-local dictionaries). In other words, they must be completely pure. Due to an implementation limitation, this also means that they currently cannot define a closure or untyped generator.

It's easiest to illustrate this with an example. We can declare a generated function `foo` as

```
julia> @generated function foo(x)
    Core.println(x)
    return : (x * x)
end
foo (generic function with 1 method)
```

Note that the body returns a quoted expression, namely `: (x * x)`, rather than just the value of `x * x`.

From the caller's perspective, they are very similar to regular functions; in fact, you don't have to know if you're calling a regular or generated function - the syntax and result of the call is just the same. Let's see how `foo` behaves:

```
# note: output is from println() statement in the body
julia> x = foo(2);
Int64

julia> x           # now we print x
4

julia> y = foo("bar");
String

julia> y
"barbar"
```

So, we see that in the body of the generated function, `x` is the *type* of the passed argument, and the value returned by the generated function, is the result of evaluating the quoted expression we returned from the definition, now with the *value* of `x`.

What happens if we evaluate `foo` again with a type that we have already used?

```
julia> foo(4)
16
```

Note that there is no printout of `Int64`. We can see that the body of the generated function was only executed once here, for the specific set of argument types, and the result was cached. After that, for this example, the expression returned from the generated function on the first invocation was re-used as the method body. However, the actual caching behavior is an implementation-defined performance optimization, so it is invalid to depend too closely on this behavior.

The number of times a generated function is generated *might* be only once, but it *might* also be more often, or appear to not happen at all. As a consequence, you should *never* write a generated function with side effects - when, and how often, the side effects occur is undefined. (This is true for macros too - and just like for macros, the use of `eval()` (page 368) in a generated function is a sign that you're doing something the wrong way.) However, unlike macros, the runtime system cannot correctly handle a call to `eval()` (page 368), so it is disallowed.

The example generated function `foo` above did not do anything a normal function `foo(x) = x * x` could not do (except printing the type on the first invocation, and incurring higher overhead). However, the power of a generated function lies in its ability to compute different quoted expressions depending on the types passed to it:

```
julia> @generated function bar(x)
    if x <: Integer
        return : (x ^ 2)
    else
        return : (x)
    end
end
bar (generic function with 1 method)

julia> bar(4)
16

julia> bar("baz")
"baz"
```

(although of course this contrived example would be more easily implemented using multiple dispatch...)

Abusing this will corrupt the runtime system and cause undefined behavior:

```
julia> @generated function baz(x)
    if rand() < .9
        return : (x^2)
    end
end
```

```

        else
            return :("boo!")
        end
    end
end
baz (generic function with 1 method)

```

Since the body of the generated function is non-deterministic, its behavior, *and the behavior of all subsequent code* is undefined.

Don't copy these examples!

These examples are hopefully helpful to illustrate how generated functions work, both in the definition end and at the call site; however, *don't copy them*, for the following reasons:

- the `foo` function has side-effects (the call to `Core.println`), and it is undefined exactly when, how often or how many times these side-effects will occur
- the `bar` function solves a problem that is better solved with multiple dispatch - defining `bar(x) = x` and `bar(x::Integer) = x ^ 2` will do the same thing, but it is both simpler and faster.
- the `baz` function is pathologically insane

Note that the set of operations that should not be attempted in a generated function is unbounded, and the runtime system can currently only detect a subset of the invalid operations. There are many other operations that will simply corrupt the runtime system without notification, usually in subtle ways not obviously connected to the bad definition. Because the function generator is run during inference, it must respect all of the limitations of that code.

Some operations that should not be attempted include:

1. Caching of native pointers.
2. Interacting with the contents or methods of `Core.Inference` in any way.
3. Observing any mutable state.
 - Inference on the generated function may be run at *any* time, including while your code is attempting to observe or mutate this state.
4. Taking any locks: C code you call out to may use locks internally, (for example, it is not problematic to call `malloc`, even though most implementations require locks internally) but don't attempt to hold or acquire any while executing Julia code.
5. Calling any function that is defined after the body of the generated function. This condition is relaxed for incrementally-loaded precompiled modules to allow calling any function in the module.

Alright, now that we have a better understanding of how generated functions work, let's use them to build some more advanced (and valid) functionality...

18.6.1 An advanced example

Julia's base library has a `sub2ind()` (page 454) function to calculate a linear index into an n-dimensional array, based on a set of n multilinear indices - in other words, to calculate the index `i` that can be used to index into an array `A` using `A[i]`, instead of `A[x, y, z, ...]`. One possible implementation is the following:

```

function sub2ind_loop{N}(dims::NTuple{N}, I::Integer...)
    ind = I[N] - 1
    for i = N-1:-1:1
        ind = I[i]-1 + dims[i]*ind
    end
    return ind + 1
end

```

The same thing can be done using recursion:

```
sub2ind_rec(dims::Tuple{}) = 1
sub2ind_rec(dims::Tuple{}, i1::Integer, I::Integer...) =
    i1 == 1 ? sub2ind_rec(dims, I...) : throw(BoundsError())
sub2ind_rec(dims::Tuple{Integer, Vararg{Integer}}, i1::Integer) = i1
sub2ind_rec(dims::Tuple{Integer, Vararg{Integer}}, i1::Integer, I::Integer...) =
    i1 + dims[1] * (sub2ind_rec(tail(dims), I...) - 1)
```

Both these implementations, although different, do essentially the same thing: a runtime loop over the dimensions of the array, collecting the offset in each dimension into the final index.

However, all the information we need for the loop is embedded in the type information of the arguments. Thus, we can utilize generated functions to move the iteration to compile-time; in compiler parlance, we use generated functions to manually unroll the loop. The body becomes almost identical, but instead of calculating the linear index, we build up an *expression* that calculates the index:

```
@generated function sub2ind_gen{N}(dims::NTuple{N}, I::Integer...)
    ex = :(I[$N] - 1)
    for i = (N - 1):-1:1
        ex = :(I[$i] - 1 + dims[$i] * $ex)
    end
    return :($ex + 1)
end
```

What code will this generate?

An easy way to find out, is to extract the body into another (regular) function:

```
julia> @generated function sub2ind_gen{N}(dims::NTuple{N}, I::Integer...)
    return sub2ind_gen_impl(dims, I...)
end
sub2ind_gen (generic function with 1 method)

julia> function sub2ind_gen_impl{N}(dims::Type{NTuple{N}}, I...)
    length(I) == N || return :(error("partial indexing is unsupported"))
    ex = :(I[$N] - 1)
    for i = (N - 1):-1:1
        ex = :(I[$i] - 1 + dims[$i] * $ex)
    end
    return :($ex + 1)
end
sub2ind_gen_impl (generic function with 1 method)
```

We can now execute `sub2ind_gen_impl` and examine the expression it returns:

```
julia> sub2ind_gen_impl(Tuple{Int,Int}, Int, Int)
:(((I[1] - 1) + dims[1] * (I[2] - 1)) + 1)
```

So, the method body that will be used here doesn't include a loop at all - just indexing into the two tuples, multiplication and addition/subtraction. All the looping is performed compile-time, and we avoid looping during execution entirely. Thus, we only loop *once per type*, in this case once per `N` (except in edge cases where the function is generated more than once - see disclaimer above).

Multi-dimensional Arrays

Julia, like most technical computing languages, provides a first-class array implementation. Most technical computing languages pay a lot of attention to their array implementation at the expense of other containers. Julia does not treat arrays in any special way. The array library is implemented almost completely in Julia itself, and derives its performance from the compiler, just like any other code written in Julia. As such, it's also possible to define custom array types by inheriting from `AbstractArray`. See the [manual section on the `AbstractArray` interface](#) (page 136) for more details on implementing a custom array type.

An array is a collection of objects stored in a multi-dimensional grid. In the most general case, an array may contain objects of type `Any`. For most computational purposes, arrays should contain objects of a more specific type, such as `Float64` or `Int32`.

In general, unlike many other technical computing languages, Julia does not expect programs to be written in a vectorized style for performance. Julia's compiler uses type inference and generates optimized code for scalar array indexing, allowing programs to be written in a style that is convenient and readable, without sacrificing performance, and using less memory at times.

In Julia, all arguments to functions are passed by reference. Some technical computing languages pass arrays by value, and this is convenient in many cases. In Julia, modifications made to input arrays within a function will be visible in the parent function. The entire Julia array library ensures that inputs are not modified by library functions. User code, if it needs to exhibit similar behavior, should take care to create a copy of inputs that it may modify.

19.1 Arrays

19.1.1 Basic Functions

| Function | Description |
|---|---|
| <code>eltype(A)</code> (page 384) | the type of the elements contained in <code>A</code> |
| <code>length(A)</code> (page 443) | the number of elements in <code>A</code> |
| <code>ndims(A)</code> (page 451) | the number of dimensions of <code>A</code> |
| <code>size(A)</code> (page 451) | a tuple containing the dimensions of <code>A</code> |
| <code>size(A, n)</code> (page 451) | the size of <code>A</code> along a particular dimension |
| <code>indices(A)</code> (page 451) | a tuple containing the valid indices of <code>A</code> |
| <code>indices(A, n)</code> (page 451) | a range expressing the valid indices along dimension <code>n</code> |
| <code>eachindex(A)</code> (page 452) | an efficient iterator for visiting each position in <code>A</code> |
| <code>stride(A, k)</code> (page 453) | the stride (linear index distance between adjacent elements) along dimension <code>k</code> |
| <code>strides(A)</code> (page 453) | a tuple of the strides in each dimension |

19.1.2 Construction and Initialization

Many functions for constructing and initializing arrays are provided. In the following list of such functions, calls with a `dims...` argument can either take a single tuple of dimension sizes or a series of dimension sizes passed as a variable number of arguments.

| Function | Description |
|---|---|
| <code>Array{type}(dims...)</code> (page 454) | an uninitialized dense array |
| <code>zeros(type, dims...)</code> (page 454) | an array of all zeros of specified type, defaults to <code>Float64</code> if <code>type</code> not specified |
| <code>zeros(A)</code> (page 454) | an array of all zeros of same element type and shape of <code>A</code> |
| <code>ones(type, dims...)</code> (page 455) | an array of all ones of specified type, defaults to <code>Float64</code> if <code>type</code> not specified |
| <code>ones(A)</code> (page 455) | an array of all ones of same element type and shape of <code>A</code> |
| <code>trues(dims...)</code> (page 455) | a <code>Bool</code> array with all values <code>true</code> |
| <code>trues(A)</code> (page 455) | a <code>Bool</code> array with all values <code>true</code> and the shape of <code>A</code> |
| <code>false(dims...)</code> (page 456) | a <code>Bool</code> array with all values <code>false</code> |
| <code>false(A)</code> (page 456) | a <code>Bool</code> array with all values <code>false</code> and the shape of <code>A</code> |
| <code>reshape(A, dims...)</code> (page 456) | an array with the same data as the given array, but with different dimensions. |
| <code>copy(A)</code> (page 363) | copy <code>A</code> |
| <code>deepcopy(A)</code> (page 363) | copy <code>A</code> , recursively copying its elements |
| <code>similar(A, element_type, dims...)</code> (page 457) | an uninitialized array of the same type as the given array (dense, sparse, etc.), but with the specified element type and dimensions. The second and third arguments are both optional, defaulting to the element type and dimensions of <code>A</code> if omitted. |
| <code>reinterpret(type, A)</code> (page 458) | an array with the same binary data as the given array, but with the specified element type |
| <code>rand(dims)</code> (page 440) | <code>Array</code> (page 454) of <code>Float64</code> s with random, iid ¹ and uniformly distributed values in the half-open interval $[0, 1)$ |
| <code>randn(dims)</code> (page 441) | <code>Array</code> (page 454) of <code>Float64</code> s with random, iid and standard normally distributed random values |
| <code>eye(n)</code> (page 458) | <code>n</code> -by- <code>n</code> identity matrix |
| <code>eye(m, n)</code> (page 458) | <code>m</code> -by- <code>n</code> identity matrix |
| <code>linspace(start, stop, n)</code> (page 458) | range of <code>n</code> linearly spaced elements from <code>start</code> to <code>stop</code> |
| <code>fill!(A, x)</code> (page 456) | fill the array <code>A</code> with the value <code>x</code> |
| <code>fill(x, dims)</code> (page 456) | create an array filled with the value <code>x</code> |

The syntax `[A, B, C, ...]` constructs a 1-d array (vector) of its arguments.

19.1.3 Concatenation

Arrays can be constructed and also concatenated using the following functions:

| Function | Description |
|--------------------------------------|---|
| <code>cat(k, A...)</code> (page 462) | concatenate input <code>n</code> -d arrays along the dimension <code>k</code> |
| <code>vcat(A...)</code> (page 462) | shorthand for <code>cat(1, A...)</code> |
| <code>hcat(A...)</code> (page 462) | shorthand for <code>cat(2, A...)</code> |

¹ *iid*, independently and identically distributed.

Scalar values passed to these functions are treated as 1-element arrays.

The concatenation functions are used so often that they have special syntax:

| Expression | Calls |
|------------------------------|---------------------------------|
| <code>[A; B; C; ...]</code> | <code>vcat()</code> (page 462) |
| <code>[A B C ...]</code> | <code>hcat()</code> (page 462) |
| <code>[A B; C D; ...]</code> | <code>hvcat()</code> (page 463) |

`hvcat()` (page 463) concatenates in both dimension 1 (with semicolons) and dimension 2 (with spaces).

19.1.4 Typed array initializers

An array with a specific element type can be constructed using the syntax `T[A, B, C, ...]`. This will construct a 1-d array with element type `T`, initialized to contain elements `A`, `B`, `C`, etc. For example `Any[x, y, z]` constructs a heterogeneous array that can contain any values.

Concatenation syntax can similarly be prefixed with a type to specify the element type of the result.

```
julia> [[1 2] [3 4]]
1×4 Array{Int64,2}:
 1  2  3  4

julia> Int8[[1 2] [3 4]]
1×4 Array{Int8,2}:
 1  2  3  4
```

19.1.5 Comprehensions

Comprehensions provide a general and powerful way to construct arrays. Comprehension syntax is similar to set construction notation in mathematics:

```
A = [ F(x,y,...) for x=rx, y=ry, ... ]
```

The meaning of this form is that `F(x,y,...)` is evaluated with the variables `x`, `y`, etc. taking on each value in their given list of values. Values can be specified as any iterable object, but will commonly be ranges like `1:n` or `2:(n-1)`, or explicit arrays of values like `[1.2, 3.4, 5.7]`. The result is an N-d dense array with dimensions that are the concatenation of the dimensions of the variable ranges `rx`, `ry`, etc. and each `F(x,y,...)` evaluation returns a scalar.

The following example computes a weighted average of the current element and its left and right neighbor along a 1-d grid. :

```
julia> x = rand(8)
8-element Array{Float64,1}:
 0.843025
 0.869052
 0.365105
 0.699456
 0.977653
 0.994953
 0.41084
 0.809411

julia> [ 0.25*x[i-1] + 0.5*x[i] + 0.25*x[i+1] for i=2:length(x)-1 ]
6-element Array{Float64,1}:
 0.736559
```

```
0.57468
0.685417
0.912429
0.8446
0.656511
```

The resulting array type depends on the types of the computed elements. In order to control the type explicitly, a type can be prepended to the comprehension. For example, we could have requested the result in single precision by writing:

```
Float32[ 0.25*x[i-1] + 0.5*x[i] + 0.25*x[i+1] for i=2:length(x)-1 ]
```

19.1.6 Generator Expressions

Comprehensions can also be written without the enclosing square brackets, producing an object known as a generator. This object can be iterated to produce values on demand, instead of allocating an array and storing them in advance (see [Iteration](#) (page 133)). For example, the following expression sums a series without allocating memory:

```
julia> sum(1/n^2 for n=1:1000)
1.6439345666815615
```

When writing a generator expression with multiple dimensions inside an argument list, parentheses are needed to separate the generator from subsequent arguments:

```
julia> map(tuple, 1/(i+j) for i=1:2, j=1:2, [1:4;])
ERROR: syntax: invalid iteration specification
```

All comma-separated expressions after `for` are interpreted as ranges. Adding parentheses lets us add a third argument to `map`:

```
julia> map(tuple, (1/(i+j) for i=1:2, j=1:2), [1 3; 2 4])
2×2 Array{Tuple{Float64,Int64},2}:
 (0.5,1)      (0.333333,3)
 (0.333333,2) (0.25,4)
```

Ranges in generators and comprehensions can depend on previous ranges by writing multiple `for` keywords:

```
julia> [(i,j) for i=1:3 for j=1:i]
6-element Array{Tuple{Int64,Int64},1}:
 (1,1)
 (2,1)
 (2,2)
 (3,1)
 (3,2)
 (3,3)
```

In such cases, the result is always 1-d.

Generated values can be filtered using the `if` keyword:

```
julia> [(i,j) for i=1:3 for j=1:i if i+j == 4]
2-element Array{Tuple{Int64,Int64},1}:
 (2,2)
 (3,1)
```

19.1.7 Indexing

The general syntax for indexing into an n -dimensional array `A` is:

```
X = A[I_1, I_2, ..., I_n]
```

where each I_k may be:

1. A scalar integer
2. A Range of the form `a:b`, or `a:b:c`
3. A `:` or `Colon()` to select entire dimensions
4. An arbitrary integer array, including the empty array `[]`
5. A boolean array to select a vector of elements at its `true` indices

If all the indices are scalars, then the result X is a single element from the array A . Otherwise, X is an array with the same number of dimensions as the sum of the dimensionalities of all the indices.

If all indices are vectors, for example, then the shape of X would be `(length(I_1), length(I_2), ..., length(I_n))`, with location `(i_1, i_2, ..., i_n)` of X containing the value `A[I_1[i_1], I_2[i_2], ..., I_n[i_n]]`. If I_1 is changed to a two-dimensional matrix, then X becomes an $n+1$ -dimensional array of shape `(size(I_1, 1), size(I_1, 2), length(I_2), ..., length(I_n))`. The matrix adds a dimension. The location `(i_1, i_2, i_3, ..., i_{n+1})` contains the value at `A[I_1[i_1, i_2], I_2[i_3], ..., I_n[i_{n+1}]]`. All dimensions indexed with scalars are dropped. For example, the result of `A[2, I, 3]` is an array with size `size(I)`. Its i th element is populated by `A[2, I[i], 3]`.

Indexing by a boolean array B is effectively the same as indexing by the vector that is returned by `find(B)` (page 465). Often referred to as logical indexing, this selects elements at the indices where the values are `true`, akin to a mask. A logical index must be a vector of the same length as the dimension it indexes into, or it must be the only index provided and match the size and dimensionality of the array it indexes into. It is generally more efficient to use boolean arrays as indices directly instead of first calling `find()` (page 465).

Additionally, single elements of a multidimensional array can be indexed as `x = A[I]`, where I is a `CartesianIndex`. It effectively behaves like an n -tuple of integers spanning multiple dimensions of A . See *Iteration* (page 183) below.

As a special part of this syntax, the `end` keyword may be used to represent the last index of each dimension within the indexing brackets, as determined by the size of the innermost array being indexed. Indexing syntax without the `end` keyword is equivalent to a call to `getindex`:

```
X = getindex(A, I_1, I_2, ..., I_n)
```

Example:

```
julia> x = reshape(1:16, 4, 4)
4×4 Base.ReshapedArray{Int64,2,UnitRange{Int64},Tuple{}}:
 1  5   9  13
 2  6  10  14
 3  7  11  15
 4  8  12  16

julia> x[2:3, 2:end-1]
2×2 Array{Int64,2}:
 6  10
 7  11

julia> x[map(ispow2, x)]
5-element Array{Int64,1}:
 1
 2
 4
```

```
8
16

julia> x[1, [2 3; 4 1]]
2×2 Array{Int64,2}:
 5  9
13  1
```

Empty ranges of the form `n:n-1` are sometimes used to indicate the inter-index location between `n-1` and `n`. For example, the `searchsorted()` (page 579) function uses this convention to indicate the insertion point of a value not found in a sorted array:

```
julia> a = [1,2,5,6,7];

julia> searchsorted(a, 3)
3:2
```

19.1.8 Assignment

The general syntax for assigning values in an `n`-dimensional array `A` is:

```
A[I_1, I_2, ..., I_n] = X
```

where each `I_k` may be:

1. A scalar integer
2. A Range of the form `a:b`, or `a:b:c`
3. A `:` or `Colon()` to select entire dimensions
4. An arbitrary integer array, including the empty array `[]`
5. A boolean array to select elements at its `true` indices

If `X` is an array, it must have the same number of elements as the product of the lengths of the indices: `prod(length(I_1), length(I_2), ..., length(I_n))`. The value in location `I_1[i_1], I_2[i_2], ..., I_n[i_n]` of `A` is overwritten with the value `X[i_1, i_2, ..., i_n]`. If `X` is not an array, its value is written to all referenced locations of `A`.

A boolean array used as an index behaves as in `getindex()` (page 454), behaving as though it is first transformed with `find()` (page 465).

Index assignment syntax is equivalent to a call to `setindex!()` (page 461):

```
setindex!(A, X, I_1, I_2, ..., I_n)
```

Example:

```
julia> x = collect(reshape(1:9, 3, 3))
3×3 Array{Int64,2}:
 1  4  7
 2  5  8
 3  6  9

julia> x[1:2, 2:3] = -1
-1

julia> x
3×3 Array{Int64,2}:
```

```
1  -1  -1
2  -1  -1
3   6   9
```

19.1.9 Iteration

The recommended ways to iterate over a whole array are

```
for a in A
    # Do something with the element a
end

for i in eachindex(A)
    # Do something with i and/or A[i]
end
```

The first construct is used when you need the value, but not index, of each element. In the second construct, `i` will be an `Int` if `A` is an array type with fast linear indexing; otherwise, it will be a `CartesianIndex`:

```
A = rand(4,3)
B = view(A, 1:3, 2:3)
julia> for i in eachindex(B)
           @show i
           end
           i = Base.IteratorsMD.CartesianIndex_2{1,1}
           i = Base.IteratorsMD.CartesianIndex_2{2,1}
           i = Base.IteratorsMD.CartesianIndex_2{3,1}
           i = Base.IteratorsMD.CartesianIndex_2{1,2}
           i = Base.IteratorsMD.CartesianIndex_2{2,2}
           i = Base.IteratorsMD.CartesianIndex_2{3,2}
```

In contrast with `for i = 1:length(A)`, iterating with `eachindex` provides an efficient way to iterate over any array type.

19.1.10 Array traits

If you write a custom `AbstractArray` type, you can specify that it has fast linear indexing using

```
Base.linearindexing{T<:MyArray}{::Type{T}} = LinearFast()
```

This setting will cause `eachindex` iteration over a `MyArray` to use integers. If you don't specify this trait, the default value `LinearSlow()` is used.

19.1.11 Vectorized Operators and Functions

The following operators are supported for arrays. The dot version of a binary operator should be used for elementwise operations.

1. Unary arithmetic — `-`, `+`, `!`
2. Binary arithmetic — `+`, `-`, `*`, `.*`, `/`, `./`, `\`, `.\`, `^`, `.^`, `div`, `mod`
3. Comparison — `==`, `!=`, `<`, `<=`, `>`, `>=`
4. Unary Boolean or bitwise — `~`
5. Binary Boolean or bitwise — `&`, `|`, `$`

Some operators without dots operate elementwise anyway when one argument is a scalar. These operators are `*`, `+`, `-`, and the bitwise operators. The operators `/` and `\` operate elementwise when the denominator is a scalar.

Note that comparisons such as `==` operate on whole arrays, giving a single boolean answer. Use dot operators for elementwise comparisons.

To enable convenient vectorization of mathematical and other operations, Julia provides the compact syntax `f.(args...)`, e.g. `sin.(x)` or `min.(x,y)`, for elementwise operations over arrays or mixtures of arrays and scalars (a *broadcast* ([page 459](#)) operation). See *Dot Syntax for Vectorizing Functions* ([page 63](#)).

Note that there is a difference between `max.(a,b)`, which broadcasts *max* ([page 418](#)) elementwise over `a` and `b`, and `maximum(a)`, which finds the largest value within `a`. The same statements hold for `min.(a,b)` and `minimum(a)`.

19.1.12 Broadcasting

It is sometimes useful to perform element-by-element binary operations on arrays of different sizes, such as adding a vector to each column of a matrix. An inefficient way to do this would be to replicate the vector to the size of the matrix:

```
julia> a = rand(2,1); A = rand(2,3);

julia> repmat(a,1,3)+A
2×3 Array{Float64,2}:
 1.20813  1.82068  1.25387
 1.56851  1.86401  1.67846
```

This is wasteful when dimensions get large, so Julia offers *broadcast* ([page 459](#)), which expands singleton dimensions in array arguments to match the corresponding dimension in the other array without using extra memory, and applies the given function elementwise:

```
julia> broadcast(+, a, A)
2×3 Array{Float64,2}:
 1.20813  1.82068  1.25387
 1.56851  1.86401  1.67846

julia> b = rand(1,2)
1×2 Array{Float64,2}:
 0.867535  0.00457906

julia> broadcast(+, a, b)
2×2 Array{Float64,2}:
 1.71056  0.847604
 1.73659  0.873631
```

Elementwise operators such as `.+` and `.*` perform broadcasting if necessary. There is also a *broadcast!* ([page 460](#)) function to specify an explicit destination, and *broadcast_getindex* ([page 461](#)) and *broadcast_setindex!* ([page 462](#)) that broadcast the indices before indexing. Moreover, `f.(args...)` is equivalent to `broadcast(f, args...)`, providing a convenient syntax to broadcast any function (*Dot Syntax for Vectorizing Functions* ([page 63](#))).

Additionally, *broadcast* ([page 459](#)) is not limited to arrays (see the function documentation), it also handles tuples and treats any argument that is not an array or a tuple as a “scalar”.

```
julia> convert.(Float32, [1, 2])
2-element Array{Float32,1}:
 1.0
 2.0
```

```
julia> ceil.((UInt8,), [1.2 3.4; 5.6 6.7])
2×2 Array{UInt8,2}:
 0x02  0x04
 0x06  0x07

julia> string.(1:3, ". ", ["First", "Second", "Third"])
3-element Array{String,1}:
"1. First"
"2. Second"
"3. Third"
```

19.1.13 Implementation

The base array type in Julia is the abstract type `AbstractArray{T,N}`. It is parametrized by the number of dimensions `N` and the element type `T`. `AbstractVector` and `AbstractMatrix` are aliases for the 1-d and 2-d cases. Operations on `AbstractArray` objects are defined using higher level operators and functions, in a way that is independent of the underlying storage. These operations generally work correctly as a fallback for any specific array implementation.

The `AbstractArray` type includes anything vaguely array-like, and implementations of it might be quite different from conventional arrays. For example, elements might be computed on request rather than stored. However, any concrete `AbstractArray{T,N}` type should generally implement at least `size(A)` (page 451) (returning an `Int` tuple), `getindex(A,i)` (page 454) and `getindex(A,i1,...,iN)` (page 454); mutable arrays should also implement `setindex!()` (page 461). It is recommended that these operations have nearly constant time complexity, or technically $\tilde{O}(1)$ complexity, as otherwise some array functions may be unexpectedly slow. Concrete types should also typically provide a `similar(A, T=eltype(A), dims=size(A))` (page 457) method, which is used to allocate a similar array for `copy()` (page 363) and other out-of-place operations. No matter how an `AbstractArray{T,N}` is represented internally, `T` is the type of object returned by `integer` indexing (`A[1, ..., 1]`, when `A` is not empty) and `N` should be the length of the tuple returned by `size()` (page 451).

`DenseArray` is an abstract subtype of `AbstractArray` intended to include all arrays that are laid out at regular offsets in memory, and which can therefore be passed to external C and Fortran functions expecting this memory layout. Subtypes should provide a method `stride(A,k)` (page 453) that returns the “stride” of dimension `k`: increasing the index of dimension `k` by 1 should increase the index `i` of `getindex(A,i)` (page 454) by `stride(A,k)` (page 453). If a pointer conversion method `Base.unsafe_convert{Ptr{T}, A}` (page 607) is provided, the memory layout should correspond in the same way to these strides.

The `Array` (page 454) type is a specific instance of `DenseArray` where elements are stored in column-major order (see additional notes in *Performance Tips* (page 305)). `Vector` and `Matrix` are aliases for the 1-d and 2-d cases. Specific operations such as scalar indexing, assignment, and a few other basic storage-specific operations are all that have to be implemented for `Array` (page 454), so that the rest of the array library can be implemented in a generic manner.

`SubArray` is a specialization of `AbstractArray` that performs indexing by reference rather than by copying. A `SubArray` is created with the `view()` (page 461) function, which is called the same way as `getindex()` (page 454) (with an array and a series of index arguments). The result of `view()` (page 461) looks the same as the result of `getindex()` (page 454), except the data is left in place. `view()` (page 461) stores the input index vectors in a `SubArray` object, which can later be used to index the original array indirectly.

`StridedVector` and `StridedMatrix` are convenient aliases defined to make it possible for Julia to call a wider range of BLAS and LAPACK functions by passing them either `Array` (page 454) or `SubArray` objects, and thus saving inefficiencies from memory allocation and copying.

The following example computes the QR decomposition of a small section of a larger array, without creating any temporaries, and by calling the appropriate LAPACK function with the right leading dimension size and stride parameters.

```

julia> a = rand(10,10)
10×10 Array{Float64,2}:
 0.561255  0.226678  0.203391  0.308912  ...  0.750307  0.235023  0.217964
 0.718915  0.537192  0.556946  0.996234  ...  0.666232  0.509423  0.660788
 0.493501  0.0565622 0.118392  0.493498  ...  0.262048  0.940693  0.252965
 0.0470779 0.736979  0.264822  0.228787  ...  0.161441  0.897023  0.567641
 0.343935  0.32327  0.795673  0.452242  ...  0.468819  0.628507  0.511528
 0.935597  0.991511 0.571297  0.74485  ...  0.84589  0.178834  0.284413
 0.160706  0.672252 0.133158  0.65554  ...  0.371826  0.770628  0.0531208
 0.306617  0.836126 0.301198  0.0224702 ...  0.39344  0.0370205 0.536062
 0.890947  0.168877 0.32002  0.486136  ...  0.096078  0.172048  0.77672
 0.507762  0.573567 0.220124  0.165816  ...  0.211049  0.433277  0.539476

julia> b = view(a, 2:2:8,2:2:4)
4×2 SubArray{Float64,2,Array{Float64,2},Tuple{StepRange{Int64,Int64},StepRange{Int64,Int64}},false}:
 0.537192  0.996234
 0.736979  0.228787
 0.991511  0.74485
 0.836126  0.0224702

julia> (q,r) = qr(b);

julia> q
4×2 Array{Float64,2}:
-0.338809  0.78934
-0.464815 -0.230274
-0.625349  0.194538
-0.527347 -0.534856

julia> r
2×2 Array{Float64,2}:
-1.58553 -0.921517
 0.0      0.866567

```

19.2 Sparse Matrices

Sparse matrices are matrices that contain enough zeros that storing them in a special data structure leads to savings in space and execution time. Sparse matrices may be used when operations on the sparse representation of a matrix lead to considerable gains in either time or space when compared to performing the same operations on a dense matrix.

19.2.1 Compressed Sparse Column (CSC) Storage

In Julia, sparse matrices are stored in the [Compressed Sparse Column \(CSC\)](#) format. Julia sparse matrices have the type `SparseMatrixCSC{Tv,Ti}`, where `Tv` is the type of the nonzero values, and `Ti` is the integer type for storing column pointers and row indices.:

```

type SparseMatrixCSC{Tv,Ti<:Integer} <: AbstractSparseMatrix{Tv,Ti}
    m::Int          # Number of rows
    n::Int          # Number of columns
    colptr::Vector{Ti} # Column i is in colptr[i]:(colptr[i+1]-1)
    rowval::Vector{Ti} # Row values of nonzeros
    nzval::Vector{Tv}  # Nonzero values
end

```


The compressed sparse column storage makes it easy and quick to access the elements in the column of a sparse matrix, whereas accessing the sparse matrix by rows is considerably slower. Operations such as insertion of nonzero values one at a time in the CSC structure tend to be slow. This is because all elements of the sparse matrix that are beyond the point of insertion have to be moved one place over.

All operations on sparse matrices are carefully implemented to exploit the CSC data structure for performance, and to avoid expensive operations.

If you have data in CSC format from a different application or library, and wish to import it in Julia, make sure that you use 1-based indexing. The row indices in every column need to be sorted. If your `SparseMatrixCSC` object contains unsorted row indices, one quick way to sort them is by doing a double transpose.

In some applications, it is convenient to store explicit zero values in a `SparseMatrixCSC`. These *are* accepted by functions in `Base` (but there is no guarantee that they will be preserved in mutating operations). Such explicitly stored zeros are treated as structural nonzeros by many routines. The `nnz()` (page 480) function returns the number of elements explicitly stored in the sparse data structure, including structural nonzeros. In order to count the exact number of actual values that are nonzero, use `countnz()` (page 453), which inspects every stored element of a sparse matrix.

19.2.2 Sparse matrix constructors

The simplest way to create sparse matrices is to use functions equivalent to the `zeros()` (page 454) and `eye()` (page 458) functions that Julia provides for working with dense matrices. To produce sparse matrices instead, you can use the same names with an `sp` prefix:

```
julia> spzeros(3,5)
3×5 sparse matrix with 0 Float64 nonzero entries

julia> speye(3,5)
3×5 sparse matrix with 3 Float64 nonzero entries:
 [1, 1] = 1.0
 [2, 2] = 1.0
 [3, 3] = 1.0
```

The `sparse()` (page 479) function is often a handy way to construct sparse matrices. It takes as its input a vector `I` of row indices, a vector `J` of column indices, and a vector `V` of nonzero values. `sparse(I, J, V)` constructs a sparse matrix such that $S[I[k], J[k]] = V[k]$.

```
julia> I = [1, 4, 3, 5]; J = [4, 7, 18, 9]; V = [1, 2, -5, 3];

julia> S = sparse(I,J,V)
5×18 sparse matrix with 4 Int64 nonzero entries:
 [1, 4] = 1
 [4, 7] = 2
 [5, 9] = 3
 [3, 18] = -5
```

The inverse of the `sparse()` (page 479) function is `findn()` (page 465), which retrieves the inputs used to create the sparse matrix.

```
julia> findn(S)
([1,4,5,3],[4,7,9,18])

julia> findnz(S)
([1,4,5,3],[4,7,9,18],[1,2,3,-5])
```

Another way to create sparse matrices is to convert a dense matrix into a sparse matrix using the `sparse()` (page 479) function:

```
julia> sparse(eye(5))
5×5 sparse matrix with 5 Float64 nonzero entries:
 [1, 1] = 1.0
 [2, 2] = 1.0
 [3, 3] = 1.0
 [4, 4] = 1.0
 [5, 5] = 1.0
```

You can go in the other direction using the `full()` (page 503) function. The `issparse()` (page 480) function can be used to query if a matrix is sparse.

```
julia> issparse(speye(5))
true
```

19.2.3 Sparse matrix operations

Arithmetic operations on sparse matrices also work as they do on dense matrices. Indexing of, assignment into, and concatenation of sparse matrices work in the same way as dense matrices. Indexing operations, especially assignment, are expensive, when carried out one element at a time. In many cases it may be better to convert the sparse matrix into (I, J, V) format using `findnz()` (page 466), manipulate the non-zeroes or the structure in the dense vectors (I, J, V) , and then reconstruct the sparse matrix.

19.2.4 Correspondence of dense and sparse methods

The following table gives a correspondence between built-in methods on sparse matrices and their corresponding methods on dense matrix types. In general, methods that generate sparse matrices differ from their dense counterparts in that the resulting matrix follows the same sparsity pattern as a given sparse matrix S , or that the resulting sparse matrix has density d , i.e. each matrix element has a probability d of being non-zero.

Details can be found in the *Sparse Vectors and Matrices* (page 479) section of the standard library reference.

| Sparse | Dense | Description |
|---|--|---|
| <code>spzeros(m, n)</code> (page 480) | <code>zeros(m, n)</code> (page 454) | Creates a m -by- n matrix of zeros. (<code>spzeros(m, n)</code> (page 480) is empty.) |
| <code>spones(S)</code> (page 480) | <code>ones(m, n)</code> (page 455) | Creates a matrix filled with ones. Unlike the dense version, <code>spones()</code> (page 480) has the same sparsity pattern as S . |
| <code>speye(n)</code> (page 481) | <code>eye(n)</code> (page 458) | Creates a n -by- n identity matrix. |
| <code>full(S)</code> (page 503) | <code>sparse(A)</code> (page 479) | Interconverts between dense and sparse formats. |
| <code>sprand(m, n, d)</code> (page 481) | <code>rand(m, n)</code> (page 440) | Creates a m -by- n random matrix (of density d) with iid non-zero elements distributed uniformly on the half-open interval $[0, 1)$. |
| <code>sprandn(m, n, d)</code> (page 482) | <code>randn(m, n)</code> (page 441) | Creates a m -by- n random matrix (of density d) with iid non-zero elements distributed according to the standard normal (Gaussian) distribution. |
| <code>sprandn(m, n, d, X)</code> (page 482) | <code>randn(m, n, X)</code> (page 441) | Creates a m -by- n random matrix (of density d) with iid non-zero elements distributed according to the X distribution. (Requires the <code>Distributions</code> package.) |

Linear algebra

20.1 Matrix factorizations

Matrix factorizations (a.k.a. matrix decompositions) compute the factorization of a matrix into a product of matrices, and are one of the central concepts in linear algebra.

The following table summarizes the types of matrix factorizations that have been implemented in Julia. Details of their associated methods can be found in the *Linear Algebra* (page 501) section of the standard library documentation.

| | |
|-----------------|--|
| Cholesky | Cholesky factorization |
| CholeskyPivoted | Pivoted Cholesky factorization |
| LU | LU factorization |
| LUTridiagonal | LU factorization for Tridiagonal matrices |
| UmfpackLU | LU factorization for sparse matrices (computed by UMFPack) |
| QR | QR factorization |
| QRCompactWY | Compact WY form of the QR factorization |
| QRPivoted | Pivoted QR factorization |
| Hessenberg | Hessenberg decomposition |
| Eigen | Spectral decomposition |
| SVD | Singular value decomposition |
| GeneralizedSVD | Generalized SVD |

20.2 Special matrices

Matrices with special symmetries and structures arise often in linear algebra and are frequently associated with various matrix factorizations. Julia features a rich collection of special matrix types, which allow for fast computation with specialized routines that are specially developed for particular matrix types.

The following tables summarize the types of special matrices that have been implemented in Julia, as well as whether hooks to various optimized methods for them in LAPACK are available.

| | |
|----------------------------------|-------------------------------|
| <i>Hermitian</i> (page 507) | Hermitian matrix |
| UpperTriangular | Upper triangular matrix |
| LowerTriangular | Lower triangular matrix |
| <i>Tridiagonal</i> (page 505) | Tridiagonal matrix |
| <i>SymTridiagonal</i> (page 505) | Symmetric tridiagonal matrix |
| <i>Bidiagonal</i> (page 503) | Upper/lower bidiagonal matrix |
| <i>Diagonal</i> (page 503) | Diagonal matrix |
| UniformScaling | Uniform scaling operator |

20.2.1 Elementary operations

| Matrix type | + | - | * | \ | Other functions with optimized methods |
|-------------------------------------|---|---|-----|-----|---|
| <i>Hermitian</i>
(page 507) | | | | MV | <i>inv()</i> (page 528), <i>sqrtn()</i> (page 532), <i>expm()</i> (page 531) |
| UpperTriangular | | | MV | MV | <i>inv()</i> (page 528), <i>det()</i> (page 528) |
| LowerTriangular | | | MV | MV | <i>inv()</i> (page 528), <i>det()</i> (page 528) |
| <i>SymTridiagonal</i>
(page 505) | M | M | MS | MV | <i>eigmax()</i> (page 515), <i>eigmin()</i> (page 515) |
| <i>Tridiagonal</i>
(page 505) | M | M | MS | MV | |
| <i>Bidiagonal</i>
(page 503) | M | M | MS | MV | |
| <i>Diagonal</i> (page 503) | M | M | MV | MV | <i>inv()</i> (page 528), <i>det()</i> (page 528), <i>logdet()</i> (page 528), <i>/()</i> (page 405) |
| UniformScaling | M | M | MVS | MVS | <i>/()</i> (page 405) |

Legend:

| | |
|------------|---|
| M (matrix) | An optimized method for matrix-matrix operations is available |
| V (vector) | An optimized method for matrix-vector operations is available |
| S (scalar) | An optimized method for matrix-scalar operations is available |

20.2.2 Matrix factorizations

| Matrix type | LA-PACK | <i>eig()</i>
(page 514) | <i>eigvals()</i>
(page 515) | <i>eigvecs()</i>
(page 516) | <i>svd()</i>
(page 520) | <i>svdvals()</i>
(page 520) |
|-------------------------------------|---------|----------------------------|--------------------------------|--------------------------------|----------------------------|--------------------------------|
| <i>Hermitian</i>
(page 507) | HE | | ARI | | | |
| UpperTriangular | TR | A | A | A | | |
| LowerTriangular | TR | A | A | A | | |
| <i>SymTridiagonal</i>
(page 505) | ST | A | ARI | AV | | |
| <i>Tridiagonal</i>
(page 505) | GT | | | | | |
| <i>Bidiagonal</i>
(page 503) | BD | | | | A | A |
| <i>Diagonal</i>
(page 503) | DI | | A | | | |

Legend:

| | | |
|--------------|--|------------------------------------|
| A (all) | An optimized method to find all the characteristic values and/or vectors is available | e.g.
<i>eigvals</i> (M) |
| R (range) | An optimized method to find the i_1^{th} through the i_h^{th} characteristic values are available | <i>eigvals</i> (M, i_1 , i_h) |
| I (interval) | An optimized method to find the characteristic values in the interval $[v_l, v_h]$ is available | <i>eigvals</i> (M, v_l , v_h) |
| V (vectors) | An optimized method to find the characteristic vectors corresponding to the characteristic values $x=[x_1, x_2, \dots]$ is available | <i>eigvecs</i> (M, x) |

20.2.3 The uniform scaling operator

A `UniformScaling` operator represents a scalar times the identity operator, $\lambda \mathbf{I}$. The identity operator \mathbf{I} is defined as a constant and is an instance of `UniformScaling`. The size of these operators are generic and match the other matrix in the binary operations `+` (page 405), `-` (page 405), `*` (page 443) and `\` (page 501). For $\mathbf{A} + \mathbf{I}$ and $\mathbf{A} - \mathbf{I}$ this means that \mathbf{A} must be square. Multiplication with the identity operator \mathbf{I} is a noop (except for checking that the scaling factor is one) and therefore almost without overhead.

Networking and Streams

Julia provides a rich interface to deal with streaming I/O objects such as terminals, pipes and TCP sockets. This interface, though asynchronous at the system level, is presented in a synchronous manner to the programmer and it is usually unnecessary to think about the underlying asynchronous operation. This is achieved by making heavy use of Julia cooperative threading (*coroutine* (page 78)) functionality.

21.1 Basic Stream I/O

All Julia streams expose at least a `read()` (page 562) and a `write()` (page 562) method, taking the stream as their first argument, e.g.:

```
julia> write(STDOUT, "Hello World"); # suppress return value 11 with ;
Hello World
julia> read(STDIN, Char)

'\n'
```

Note that `write()` (page 562) returns 11, the number of bytes (in "Hello World") written to `STDOUT` (page 561), but this return value is suppressed with the `;`.

Here Enter was pressed again so that Julia would read the newline. Now, as you can see from this example, `write()` (page 562) takes the data to write as its second argument, while `read()` (page 562) takes the type of the data to be read as the second argument.

For example, to read a simple byte array, we could do:

```
julia> x = zeros{UInt8, 4}
4-element Array{UInt8,1}:
 0x00
 0x00
 0x00
 0x00

julia> read!(STDIN, x)
abcd
4-element Array{UInt8,1}:
 0x61
 0x62
 0x63
 0x64
```

However, since this is slightly cumbersome, there are several convenience methods provided. For example, we could have written the above as:

```
julia> read(STDIN,4)
abcd
4-element Array{UInt8,1}:
 0x61
 0x62
 0x63
 0x64
```

or if we had wanted to read the entire line instead:

```
julia> readline(STDIN)
abcd
"abcd\n"
```

Note that depending on your terminal settings, your TTY may be line buffered and might thus require an additional enter before the data is sent to Julia.

To read every line from `STDIN` (page 561) you can use `eachline()` (page 568):

```
for line in eachline(STDIN)
    print("Found $line")
end
```

or `read()` (page 562) if you wanted to read by character instead:

```
while !eof(STDIN)
    x = read(STDIN, Char)
    println("Found: $x")
end
```

21.2 Text I/O

Note that the `write()` (page 562) method mentioned above operates on binary streams. In particular, values do not get converted to any canonical text representation but are written out as is:

```
julia> write(STDOUT,0x61); # suppress return value 1 with ;
a
```

Note that `a` is written to `STDOUT` (page 561) by the `write()` (page 562) function and that the returned value is `1` (since `0x61` is one byte).

For text I/O, use the `print()` (page 566) or `show()` (page 566) methods, depending on your needs (see the standard library reference for a detailed discussion of the difference between the two):

```
julia> print(STDOUT,0x61)
97
```

21.3 IO Output Contextual Properties

Sometimes IO output can benefit from the ability to pass contextual information into show methods. The `IOContext` object provides this framework for associating arbitrary metadata with an IO object. For example, `showcompact` adds a hinting parameter to the IO object that the invoked show method should print a shorter output (if applicable).

21.4 Working with Files

Like many other environments, Julia has an `open()` (page 561) function, which takes a filename and returns an `IOStream` object that you can use to read and write things from the file. For example if we have a file, `hello.txt`, whose contents are `Hello, World!`:

```
julia> f = open("hello.txt")
IOStream(<file hello.txt>)

julia> readlines(f)
1-element Array{String,1}:
 "Hello, World!\n"
```

If you want to write to a file, you can open it with the write ("`w`") flag:

```
julia> f = open("hello.txt", "w")
IOStream(<file hello.txt>)

julia> write(f, "Hello again.")
12
```

If you examine the contents of `hello.txt` at this point, you will notice that it is empty; nothing has actually been written to disk yet. This is because the `IOStream` must be closed before the write is actually flushed to disk:

```
julia> close(f)
```

Examining `hello.txt` again will show its contents have been changed.

Opening a file, doing something to its contents, and closing it again is a very common pattern. To make this easier, there exists another invocation of `open()` (page 561) which takes a function as its first argument and filename as its second, opens the file, calls the function with the file as an argument, and then closes it again. For example, given a function:

```
function read_and_capitalize(f::IOStream)
    return uppercase(readstring(f))
end
```

You can call:

```
julia> open(read_and_capitalize, "hello.txt")
"HELLO AGAIN."
```

to open `hello.txt`, call `read_and_capitalize` on it, close `hello.txt` and return the capitalized contents.

To avoid even having to define a named function, you can use the `do` syntax, which creates an anonymous function on the fly:

```
julia> open("hello.txt") do f
    uppercase(readstring(f))
end
"HELLO AGAIN."
```

21.5 A simple TCP example

Let's jump right in with a simple example involving TCP sockets. Let's first create a simple server:

```
julia> @async begin
    server = listen(2000)
    while true
        sock = accept(server)
        println("Hello World\n")
    end
end
Task (runnable) @0x00007fd31dc11ae0

julia>
```

To those familiar with the Unix socket API, the method names will feel familiar, though their usage is somewhat simpler than the raw Unix socket API. The first call to `listen()` (page 573) will create a server waiting for incoming connections on the specified port (2000) in this case. The same function may also be used to create various other kinds of servers:

```
julia> listen(2000) # Listens on localhost:2000 (IPv4)
TCPServer(active)

julia> listen(ip"127.0.0.1",2000) # Equivalent to the first
TCPServer(active)

julia> listen(ip "::1",2000) # Listens on localhost:2000 (IPv6)
TCPServer(active)

julia> listen(IPv4(0),2001) # Listens on port 2001 on all IPv4 interfaces
TCPServer(active)

julia> listen(IPv6(0),2001) # Listens on port 2001 on all IPv6 interfaces
TCPServer(active)

julia> listen("testsocket") # Listens on a UNIX domain socket/named pipe
PipeServer(active)
```

Note that the return type of the last invocation is different. This is because this server does not listen on TCP, but rather on a named pipe (Windows) or UNIX domain socket. The difference is subtle and has to do with the `accept()` (page 573) and `connect()` (page 499) methods. The `accept()` (page 573) method retrieves a connection to the client that is connecting on the server we just created, while the `connect()` (page 499) function connects to a server using the specified method. The `connect()` (page 499) function takes the same arguments as `listen()` (page 573), so, assuming the environment (i.e. host, cwd, etc.) is the same you should be able to pass the same arguments to `connect()` (page 499) as you did to listen to establish the connection. So let's try that out (after having created the server above):

```
julia> connect(2000)
TCPSocket(open, 0 bytes waiting)

julia> Hello World
```

As expected we saw “Hello World” printed. So, let's actually analyze what happened behind the scenes. When we called `connect()` (page 499), we connect to the server we had just created. Meanwhile, the `accept` function returns a server-side connection to the newly created socket and prints “Hello World” to indicate that the connection was successful.

A great strength of Julia is that since the API is exposed synchronously even though the I/O is actually happening asynchronously, we didn't have to worry callbacks or even making sure that the server gets to run. When we called `connect()` (page 499) the current task waited for the connection to be established and only continued executing after that was done. In this pause, the server task resumed execution (because a connection request was now available), accepted the connection, printed the message and waited for the next client. Reading and writing works in the same

way. To see this, consider the following simple echo server:

```
julia> @async begin
    server = listen(2001)
    while true
        sock = accept(server)
        @async while isopen(sock)
            write(sock, readline(sock))
        end
    end
end
Task (runnable) @0x00007fd31dc12e60

julia> clientside=connect(2001)
TCPSocket(open, 0 bytes waiting)

julia> @async while true
    write(STDOUT, readline(clientside))
end
Task (runnable) @0x00007fd31dc11870

julia> println(clientside, "Hello World from the Echo Server")
Hello World from the Echo Server
```

As with other streams, use `close()` (page 487) to disconnect the socket:

```
julia> close(clientside)
```

21.6 Resolving IP Addresses

One of the `connect()` (page 499) methods that does not follow the `listen()` (page 573) methods is `connect(host::String, port)`, which will attempt to connect to the host given by the `host` parameter on the port given by the `port` parameter. It allows you to do things like:

```
julia> connect("google.com", 80)
TCPSocket(open, 0 bytes waiting)
```

At the base of this functionality is `getaddrinfo()` (page 573), which will do the appropriate address resolution:

```
julia> getaddrinfo("google.com")
ip"74.125.226.225"
```

Parallel Computing

Most modern computers possess more than one CPU, and several computers can be combined together in a cluster. Harnessing the power of these multiple CPUs allows many computations to be completed more quickly. There are two major factors that influence performance: the speed of the CPUs themselves, and the speed of their access to memory. In a cluster, it's fairly obvious that a given CPU will have fastest access to the RAM within the same computer (node). Perhaps more surprisingly, similar issues are relevant on a typical multicore laptop, due to differences in the speed of main memory and the [cache](#). Consequently, a good multiprocessing environment should allow control over the “ownership” of a chunk of memory by a particular CPU. Julia provides a multiprocessing environment based on message passing to allow programs to run on multiple processes in separate memory domains at once.

Julia's implementation of message passing is different from other environments such as MPI¹. Communication in Julia is generally “one-sided”, meaning that the programmer needs to explicitly manage only one process in a two-process operation. Furthermore, these operations typically do not look like “message send” and “message receive” but rather resemble higher-level operations like calls to user functions.

Parallel programming in Julia is built on two primitives: *remote references* and *remote calls*. A remote reference is an object that can be used from any process to refer to an object stored on a particular process. A remote call is a request by one process to call a certain function on certain arguments on another (possibly the same) process.

Remote references come in two flavors: [Future](#) (page 490) and [RemoteChannel](#) (page 490).

A remote call returns a [Future](#) (page 490) to its result. Remote calls return immediately; the process that made the call proceeds to its next operation while the remote call happens somewhere else. You can wait for a remote call to finish by calling [wait\(\)](#) (page 490) on the returned [Future](#) (page 490), and you can obtain the full value of the result using [fetch\(\)](#) (page 487).

On the other hand, [RemoteChannel](#) (page 490) s are rewritable. For example, multiple processes can co-ordinate their processing by referencing the same remote Channel.

Each process has an associated identifier. The process providing the interactive Julia prompt always has an `id` equal to 1. The processes used by default for parallel operations are referred to as “workers”. When there is only one process, process 1 is considered a worker. Otherwise, workers are considered to be all processes other than process 1.

Let's try this out. Starting with `julia -p n` provides `n` worker processes on the local machine. Generally it makes sense for `n` to equal the number of CPU cores on the machine.

```
$ ./julia -p 2

julia> r = remotecall(rand, 2, 2, 2)
Future{2,1,3,Nullable{Any}}()
```

¹ In this context, MPI refers to the MPI-1 standard. Beginning with MPI-2, the MPI standards committee introduced a new set of communication mechanisms, collectively referred to as Remote Memory Access (RMA). The motivation for adding RMA to the MPI standard was to facilitate one-sided communication patterns. For additional information on the latest MPI standard, see <http://mpi-forum.org/docs>.

```
julia> s = @spawnat 2 1 .+ fetch(r)
Future{2,1,6,Nullable{Any}}()

julia> fetch(s)
2×2 Array{Float64,2}:
 1.60401  1.50111
 1.17457  1.15741
```

The first argument to `remotecall()` (page 491) is the function to call. Most parallel programming in Julia does not reference specific processes or the number of processes available, but `remotecall()` (page 491) is considered a low-level interface providing finer control. The second argument to `remotecall()` (page 491) is the `id` of the process that will do the work, and the remaining arguments will be passed to the function being called.

As you can see, in the first line we asked process 2 to construct a 2-by-2 random matrix, and in the second line we asked it to add 1 to it. The result of both calculations is available in the two futures, `r` and `s`. The `@spawnat` (page 493) macro evaluates the expression in the second argument on the process specified by the first argument.

Occasionally you might want a remotely-computed value immediately. This typically happens when you read from a remote object to obtain data needed by the next local operation. The function `remotecall_fetch()` (page 491) exists for this purpose. It is equivalent to `fetch(remotecall(...))` but is more efficient.

```
julia> remotecall_fetch(getindex, 2, r, 1, 1)
0.10824216411304866
```

Remember that `getindex(r, 1, 1)` (page 454) is *equivalent* (page 180) to `r[1, 1]`, so this call fetches the first element of the future `r`.

The syntax of `remotecall()` (page 491) is not especially convenient. The macro `@spawn` (page 493) makes things easier. It operates on an expression rather than a function, and picks where to do the operation for you:

```
julia> r = @spawn rand(2,2)
Future{2,1,4,Nullable{Any}}()

julia> s = @spawn 1 .+ fetch(r)
Future{3,1,5,Nullable{Any}}()

julia> fetch(s)
1.10824216411304866 1.13798233877923116
1.12376292706355074 1.18750497916607167
```

Note that we used `1 .+ fetch(r)` instead of `1 .+ r`. This is because we do not know where the code will run, so in general a `fetch()` (page 487) might be required to move `r` to the process doing the addition. In this case, `@spawn` (page 493) is smart enough to perform the computation on the process that owns `r`, so the `fetch()` (page 487) will be a no-op (no work is done).

(It is worth noting that `@spawn` (page 493) is not built-in but defined in Julia as a *macro* (page 166). It is possible to define your own such constructs.)

An important thing to remember is that, once fetched, a `Future` (page 490) will cache its value locally. Further `fetch()` (page 487) calls do not entail a network hop. Once all referencing `Future` (page 490)s have fetched, the remote stored value is deleted.

22.1 Code Availability and Loading Packages

Your code must be available on any process that runs it. For example, type the following into the Julia prompt:

```
julia> function rand2(dims...)
    return 2*rand(dims...)
end

julia> rand2(2,2)
2×2 Array{Float64,2}:
 0.153756  0.368514
 1.15119   0.918912

julia> fetch(@spawn rand2(2,2))
ERROR: On worker 2:
function rand2 not defined on process 2
```

Process 1 knew about the function `rand2`, but process 2 did not.

Most commonly you'll be loading code from files or packages, and you have a considerable amount of flexibility in controlling which processes load code. Consider a file, `"DummyModule.jl"`, containing the following code:

```
module DummyModule

export MyType, f

type MyType
    a::Int
end

f(x) = x^2+1

println("loaded")

end
```

Starting Julia with `julia -p 2`, you can use this to verify the following:

- `include("DummyModule.jl")` (page 361) loads the file on just a single process (whichever one executes the statement).
- `using DummyModule` causes the module to be loaded on all processes; however, the module is brought into scope only on the one executing the statement.
- As long as `DummyModule` is loaded on process 2, commands like

```
rr = RemoteChannel{2}
put!(rr, MyType(7))
```

allow you to store an object of type `MyType` on process 2 even if `DummyModule` is not in scope on process 2.

You can force a command to run on all processes using the `@everywhere` macro. For example, `@everywhere` can also be used to directly define a function on all processes:

```
julia> @everywhere id = myid()

julia> remotecall_fetch(()->id, 2)
2
```

A file can also be preloaded on multiple processes at startup, and a driver script can be used to drive the computation:

```
julia -p <n> -L file1.jl -L file2.jl driver.jl
```

The Julia process running the driver script in the example above has an `id` equal to 1, just like a process providing an interactive prompt.

The base Julia installation has in-built support for two types of clusters:

- A local cluster specified with the `-p` option as shown above.
- A cluster spanning machines using the `--machinefile` option. This uses a passwordless `ssh` login to start Julia worker processes (from the same path as the current host) on the specified machines.

Functions `addprocs()` (page 488), `rmprocs()` (page 489), `workers()` (page 489), and others are available as a programmatic means of adding, removing and querying the processes in a cluster.

Note that workers do not run a `.juliarc.jl` startup script, nor do they synchronize their global state (such as global variables, new method definitions, and loaded modules) with any of the other running processes.

Other types of clusters can be supported by writing your own custom `ClusterManager`, as described below in the *ClusterManagers* (page 213) section.

22.2 Data Movement

Sending messages and moving data constitute most of the overhead in a parallel program. Reducing the number of messages and the amount of data sent is critical to achieving performance and scalability. To this end, it is important to understand the data movement performed by Julia's various parallel programming constructs.

`fetch()` (page 487) can be considered an explicit data movement operation, since it directly asks that an object be moved to the local machine. `@spawn` (page 493) (and a few related constructs) also moves data, but this is not as obvious, hence it can be called an implicit data movement operation. Consider these two approaches to constructing and squaring a random matrix:

```
# method 1
A = rand(1000,1000)
Bref = @spawn A^2
...
fetch(Bref)

# method 2
Bref = @spawn rand(1000,1000)^2
...
fetch(Bref)
```

The difference seems trivial, but in fact is quite significant due to the behavior of `@spawn` (page 493). In the first method, a random matrix is constructed locally, then sent to another process where it is squared. In the second method, a random matrix is both constructed and squared on another process. Therefore the second method sends much less data than the first.

In this toy example, the two methods are easy to distinguish and choose from. However, in a real program designing data movement might require more thought and likely some measurement. For example, if the first process needs matrix `A` then the first method might be better. Or, if computing `A` is expensive and only the current process has it, then moving it to another process might be unavoidable. Or, if the current process has very little to do between the `@spawn` (page 493) and `fetch(Bref)`, it might be better to eliminate the parallelism altogether. Or imagine `rand(1000,1000)` is replaced with a more expensive operation. Then it might make sense to add another `@spawn` (page 493) statement just for this step.

22.3 Parallel Map and Loops

Fortunately, many useful parallel computations do not require data movement. A common example is a Monte Carlo simulation, where multiple processes can handle independent simulation trials simultaneously. We can use `@spawn` (page 493) to flip coins on two processes. First, write the following function in `count_heads.jl`:


```
function count_heads(n)
    c::Int = 0
    for i=1:n
        c += rand(Bool)
    end
    c
end
```

The function `count_heads` simply adds together `n` random bits. Here is how we can perform some trials on two machines, and add together the results:

```
@everywhere include("count_heads.jl")

a = @spawn count_heads(100000000)
b = @spawn count_heads(100000000)
fetch(a)+fetch(b)
```

This example demonstrates a powerful and often-used parallel programming pattern. Many iterations run independently over several processes, and then their results are combined using some function. The combination process is called a *reduction*, since it is generally tensor-rank-reducing: a vector of numbers is reduced to a single number, or a matrix is reduced to a single row or column, etc. In code, this typically looks like the pattern `x = f(x, v[i])`, where `x` is the accumulator, `f` is the reduction function, and the `v[i]` are the elements being reduced. It is desirable for `f` to be associative, so that it does not matter what order the operations are performed in.

Notice that our use of this pattern with `count_heads` can be generalized. We used two explicit `@spawn` (page 493) statements, which limits the parallelism to two processes. To run on any number of processes, we can use a *parallel for loop*, which can be written in Julia using `@parallel` (page 493) like this:

```
nheads = @parallel (+) for i=1:200000000
    Int(rand(Bool))
end
```

This construct implements the pattern of assigning iterations to multiple processes, and combining them with a specified reduction (in this case `(+)`). The result of each iteration is taken as the value of the last expression inside the loop. The whole parallel loop expression itself evaluates to the final answer.

Note that although parallel for loops look like serial for loops, their behavior is dramatically different. In particular, the iterations do not happen in a specified order, and writes to variables or arrays will not be globally visible since iterations run on different processes. Any variables used inside the parallel loop will be copied and broadcast to each process.

For example, the following code will not work as intended:

```
a = zeros(100000)
@parallel for i=1:100000
    a[i] = i
end
```

This code will not initialize all of `a`, since each process will have a separate copy of it. Parallel for loops like these must be avoided. Fortunately, *Shared Arrays* (page ??) can be used to get around this limitation:

```
a = SharedArray{Float64,10}
@parallel for i=1:10
    a[i] = i
end
```

Using “outside” variables in parallel loops is perfectly reasonable if the variables are read-only:

```
a = randn(1000)
@parallel (+) for i=1:100000
```

```
f(a[rand(1:end)])
end
```

Here each iteration applies `f` to a randomly-chosen sample from a vector `a` shared by all processes.

As you could see, the reduction operator can be omitted if it is not needed. In that case, the loop executes asynchronously, i.e. it spawns independent tasks on all available workers and returns an array of *Future* (page 490) immediately without waiting for completion. The caller can wait for the *Future* (page 490) completions at a later point by calling `fetch()` (page 487) on them, or wait for completion at the end of the loop by prefixing it with `@sync` (page 493), like `@sync @parallel for`.

In some cases no reduction operator is needed, and we merely wish to apply a function to all integers in some range (or, more generally, to all elements in some collection). This is another useful operation called *parallel map*, implemented in Julia as the `pmap()` (page 489) function. For example, we could compute the singular values of several large random matrices in parallel as follows:

```
M = Matrix{Float64}[rand(1000,1000) for i=1:10]
pmap(svd, M)
```

Julia’s `pmap()` (page 489) is designed for the case where each function call does a large amount of work. In contrast, `@parallel for` can handle situations where each iteration is tiny, perhaps merely summing two numbers. Only worker processes are used by both `pmap()` (page 489) and `@parallel for` for the parallel computation. In case of `@parallel for`, the final reduction is done on the calling process.

22.4 Synchronization With Remote References

22.5 Scheduling

Julia’s parallel programming platform uses *Tasks (aka Coroutines)* (page 78) to switch among multiple computations. Whenever code performs a communication operation like `fetch()` (page 487) or `wait()` (page 490), the current task is suspended and a scheduler picks another task to run. A task is restarted when the event it is waiting for completes.

For many problems, it is not necessary to think about tasks directly. However, they can be used to wait for multiple events at the same time, which provides for *dynamic scheduling*. In dynamic scheduling, a program decides what to compute or where to compute it based on when other jobs finish. This is needed for unpredictable or unbalanced workloads, where we want to assign more work to processes only when they finish their current tasks.

As an example, consider computing the singular values of matrices of different sizes:

```
M = Matrix{Float64}[rand(800,800), rand(600,600), rand(800,800), rand(600,600)]
pmap(svd, M)
```

If one process handles both 800×800 matrices and another handles both 600×600 matrices, we will not get as much scalability as we could. The solution is to make a local task to “feed” work to each process when it completes its current task. For example, consider a simple `pmap()` (page 489) implementation:

```
function pmap(f, lst)
    np = nprocs() # determine the number of processes available
    n = length(lst)
    results = Vector{Any}(n)
    i = 1
    # function to produce the next work item from the queue.
    # in this case it's just an index.
    nextidx() = (idx=i; i+=1; idx)
    @sync begin
```

```

    for p=1:np
        if p != myid() || np == 1
            @async begin
                while true
                    idx = nextidx()
                    if idx > n
                        break
                    end
                    results[idx] = remotecall_fetch(f, p, lst[idx])
                end
            end
        end
    end
end
results
end

```

`@async` (page 493) is similar to `@spawn` (page 493), but only runs tasks on the local process. We use it to create a “feeder” task for each process. Each task picks the next index that needs to be computed, then waits for its process to finish, then repeats until we run out of indexes. Note that the feeder tasks do not begin to execute until the main task reaches the end of the `@sync` (page 493) block, at which point it surrenders control and waits for all the local tasks to complete before returning from the function. The feeder tasks are able to share state via `nextidx()` because they all run on the same process. No locking is required, since the threads are scheduled cooperatively and not preemptively. This means context switches only occur at well-defined points: in this case, when `remotecall_fetch()` (page 491) is called.

22.6 Channels

The section on Tasks in *Control Flow* (page 65) discussed the execution of multiple functions in a co-operative manner. Channels can be quite useful to pass data between running tasks, particularly those involving I/O operations.

Examples of operations involving I/O include reading/writing to files, accessing web services, executing external programs, etc. In all these cases, overall execution time can be improved if other tasks can be run while a file is being read, or while waiting for an external service/program to complete.

A channel can be visualized as a pipe, i.e., it has a write end and read end.

- Multiple writers in different tasks can write to the same channel concurrently via `put!()` (page 487) calls.
- Multiple readers in different tasks can read data concurrently via `take!()` (page 487) calls.
- As an example:

```

# Given Channels c1 and c2,
c1 = Channel{32}
c2 = Channel{32}

# and a function `foo()` which reads items from c1, processes the item read
# and writes a result to c2,
function foo()
    while true
        data = take!(c1)
        ..... # process data
        put!(c2, result) # write out result
    end
end
end

```

```
# we can schedule `n` instances of `foo()` to be active concurrently.
for _ in 1:n
    @schedule foo()
end
```

- Channels are created via the `Channel{T}(sz)` constructor. The channel will only hold objects of type `T`. If the type is not specified, the channel can hold objects of any type. `sz` refers to the maximum number of elements that can be held in the channel at any time. For example, `Channel{32}` creates a channel that can hold a maximum of 32 objects of any type. A `Channel{MyType}(64)` can hold up to 64 objects of `MyType` at any time.
- If a Channel is empty, readers (on a `take!()` (page 487) call) will block until data is available.
- If a Channel is full, writers (on a `put!()` (page 487) call) will block until space becomes available.
- `isready()` (page 487) tests for the presence of any object in the channel, while `wait()` (page 490) waits for an object to become available.
- A Channel is in an open state initially. This means that it can be read from and written to freely via `take!()` (page 487) and `put!()` (page 487) calls. `close()` (page 487) closes a Channel. On a closed Channel, `put!()` (page 487) will fail. For example:

```
julia> c=Channel{2};

julia> put!(c,1)    # `put!` on an open channel succeeds
1

julia> close(c);

julia> put!(c,2)    # `put!` on a closed channel throws an exception.
ERROR: InvalidStateException("Channel is closed.":closed)
...
```

- `take!()` (page 487) and `fetch()` (page 487) (which retrieves but does not remove the value) on a closed channel successfully return any existing values until it is emptied. Continuing the above example:

```
julia> fetch(c)     # Any number of `fetch` calls succeed.
1

julia> fetch(c)
1

julia> take!(c)     # The first `take!` removes the value.
1

julia> take!(c)     # No more data available on a closed channel.
ERROR: InvalidStateException("Channel is closed.":closed)
...
```

A Channel can be used as an iterable object in a `for` loop, in which case the loop runs as long as the Channel has data or is open. The loop variable takes on all values added to the Channel. The `for` loop is terminated once the Channel is closed and emptied.

For example, the following would cause the `for` loop to wait for more data:

```
c=Channel{Int}(10)
foreach(i->put!(c, i), 1:3)    # add a few entries
data = [i for i in c]
```

while this will return after reading all data:

```
julia> c=Channel{Int}(10);

julia> foreach(i->put!(c, i), 1:3);    # add a few entries

julia> close(c);                      # `for` loops can exit

julia> data = [i for i in c]
3-element Array{Int64,1}:
 1
 2
 3
```

Consider a simple example using channels for inter-task communication. We start 4 tasks to process data from a single jobs channel. Jobs, identified by an id (`job_id`), are written to the channel. Each task in this simulation reads a `job_id`, waits for a random amount of time and writes back a tuple of `job_id` and the simulated time to the results channel. Finally all the results are printed out.

```
const jobs = Channel{Int}(32)
const results = Channel{Tuple}(32)

function do_work()
    for job_id in jobs
        exec_time = rand()
        sleep(exec_time)                # simulates elapsed time doing actual work
                                       # typically performed externally.
        put!(results, (job_id, exec_time))
    end
end

function make_jobs(n)
    for i in 1:n
        put!(jobs, i)
    end
end

# feed the jobs channel with "n" jobs
n = 12
@schedule make_jobs(n)

# start 4 tasks to process requests in parallel
for i in 1:4
    @schedule do_work()
end

# print out results
@elapsed while n > 0
    job_id, exec_time = take!(results)
    println("$job_id finished in $(round(exec_time,2)) seconds")
    n = n - 1
end
```

The current version of Julia multiplexes all tasks onto a single OS thread. Thus, while tasks involving I/O operations benefit from parallel execution, compute bound tasks are effectively executed sequentially on a single OS thread. Future versions of Julia may support scheduling of tasks on multiple threads, in which case compute bound tasks will see benefits of parallel execution too.

22.7 Remote References and AbstractChannels

Remote references always refer to an implementation of an `AbstractChannel`.

A concrete implementation of an `AbstractChannel` (like `Channel`), is required to implement `put!()` (page 487), `take!()` (page 487), `fetch()` (page 487), `isready()` (page 487) and `wait()` (page 490). The remote object referred to by a `Future` (page 490) is stored in a `Channel{Any}(1)`, i.e., a `Channel` of size 1 capable of holding objects of `Any` type.

`RemoteChannel` (page 490), which is rewritable, can point to any type and size of channels, or any other implementation of an `AbstractChannel`.

The constructor `RemoteChannel(f::Function, pid)()` allows us to construct references to channels holding more than one value of a specific type. `f()` is a function executed on `pid` and it must return an `AbstractChannel`.

For example, `RemoteChannel(()->Channel{Int}(10), pid)`, will return a reference to a channel of type `Int` and size 10. The channel exists on worker `pid`.

Methods `put!()` (page 487), `take!()` (page 487), `fetch()` (page 487), `isready()` (page 487) and `wait()` (page 490) on a `RemoteChannel` (page 490) are proxied onto the backing store on the remote process.

`RemoteChannel` (page 490) can thus be used to refer to user implemented `AbstractChannel` objects. A simple example of this is provided in `examples/dictchannel.jl` which uses a dictionary as its remote store.

22.8 Channels and RemoteChannels

- A `Channel` is local to a process. Worker 2 cannot directly refer to a `Channel` on worker 3 and vice-versa. A `RemoteChannel` (page 490), however, can put and take values across workers.
- A `RemoteChannel` (page 490) can be thought of as a *handle* to a `Channel`.
- The process id, `pid`, associated with a `RemoteChannel` (page 490) identifies the process where the backing store, i.e., the backing `Channel` exists.
- Any process with a reference to a `RemoteChannel` (page 490) can put and take items from the channel. Data is automatically sent to (or retrieved from) the process a `RemoteChannel` (page 490) is associated with.
- Serializing a `Channel` also serializes any data present in the channel. Deserializing it therefore effectively makes a copy of the original object.
- On the other hand, serializing a `RemoteChannel` (page 490) only involves the serialization of an identifier that identifies the location and instance of `Channel` referred to by the handle. A deserialized `RemoteChannel` (page 490) object (on any worker), therefore also points to the same backing store as the original.

The channels example from above *Channels Example* (page 207) can be modified for interprocess communication, as shown below.

We start 4 workers to process a single `jobs` remote channel. Jobs, identified by an id (`job_id`), are written to the channel. Each remotely executing task in this simulation reads a `job_id`, waits for a random amount of time and writes back a tuple of `job_id`, time taken and its own `pid` to the results channel. Finally all the results are printed out on the master process.

```
addprocs(4)           # add worker processes

const jobs = RemoteChannel(()->Channel{Int}(32))
const results = RemoteChannel(()->Channel{Tuple}(32))

# define work function everywhere
```

```

@everywhere function do_work(jobs, results)
    while true
        job_id = take!(jobs)
        exec_time = rand()
        sleep(exec_time)           # simulates elapsed time doing actual work
        put!(results, (job_id, exec_time, myid()))
    end
end

function make_jobs(n)
    for i in 1:n
        put!(jobs, i)
    end
end

# feed the jobs channel with "n" jobs
n = 12
@schedule make_jobs(n)

# start tasks on the workers to process requests in parallel
for p in workers()
    @async remote_do(do_work, p, jobs, results)
end

# print out results
@elapsed while n > 0
    job_id, exec_time, where = take!(results)
    println("$job_id finished in $(round(exec_time,2)) seconds on worker $where")
    n = n - 1
end

```

22.9 Remote References and Distributed Garbage Collection

Objects referred to by remote references can be freed only when *all* held references in the cluster are deleted.

The node where the value is stored keeps track of which of the workers have a reference to it. Every time a *RemoteChannel* (page 490) or a (unfetched) *Future* (page 490) is serialized to a worker, the node pointed to by the reference is notified. And every time a *RemoteChannel* (page 490) or a (unfetched) *Future* (page 490) is garbage collected locally, the node owning the value is again notified.

The notifications are done via sending of “tracking” messages—an “add reference” message when a reference is serialized to a different process and a “delete reference” message when a reference is locally garbage collected.

Since *Future* (page 490)s are write-once and cached locally, the act of *fetch()* (page 487)ing a *Future* (page 490) also updates reference tracking information on the node owning the value.

The node which owns the value frees it once all references to it are cleared.

With *Future* (page 490)s, serializing an already fetched *Future* (page 490) to a different node also sends the value since the original remote store may have collected the value by this time.

It is important to note that *when* an object is locally garbage collected depends on the size of the object and the current memory pressure in the system.

In case of remote references, the size of the local reference object is quite small, while the value stored on the remote node may be quite large. Since the local object may not be collected immediately, it is a good practice to explicitly call *finalize()* (page 363) on local instances of a *RemoteChannel* (page 490), or on unfetched *Future* (page 490)s. Since calling *fetch()* (page 487) on a *Future* (page 490) also removes its reference from the remote

store, this is not required on fetched *Future* (page 490)s. Explicitly calling *finalize()* (page 363) results in an immediate message sent to the remote node to go ahead and remove its reference to the value.

Once finalized, a reference becomes invalid and cannot be used in any further calls.

22.10 Shared Arrays

Shared Arrays use system shared memory to map the same array across many processes. While there are some similarities to a *DArray*, the behavior of a *SharedArray* (page 494) is quite different. In a *DArray*, each process has local access to just a chunk of the data, and no two processes share the same chunk; in contrast, in a *SharedArray* (page 494) each “participating” process has access to the entire array. A *SharedArray* (page 494) is a good choice when you want to have a large amount of data jointly accessible to two or more processes on the same machine.

SharedArray (page 494) indexing (assignment and accessing values) works just as with regular arrays, and is efficient because the underlying memory is available to the local process. Therefore, most algorithms work naturally on *SharedArray* (page 494)s, albeit in single-process mode. In cases where an algorithm insists on an *Array* (page 454) input, the underlying array can be retrieved from a *SharedArray* (page 494) by calling *sdata()* (page 495). For other *AbstractArray* types, *sdata()* (page 495) just returns the object itself, so it’s safe to use *sdata()* (page 495) on any *Array*-type object.

The constructor for a shared array is of the form:

```
SharedArray{T::Type, dims::NTuple; init=false, pids=Int[]}
```

which creates a shared array of a bits type *T* and size *dims* across the processes specified by *pids*. Unlike distributed arrays, a shared array is accessible only from those participating workers specified by the *pids* named argument (and the creating process too, if it is on the same host).

If an *init* function, of signature *initfn(S::SharedArray)*, is specified, it is called on all the participating workers. You can specify that each worker runs the *init* function on a distinct portion of the array, thereby parallelizing initialization.

Here’s a brief example:

```
julia> addprocs(3)
3-element Array{Int64,1}:
 2
 3
 4

julia> S = SharedArray{Int, (3,4), init = S -> S[Base.localindexes(S)] = myid()}
3×4 SharedArray{Int64,2}:
 2  2  3  4
 2  3  3  4
 2  3  4  4

julia> S[3,2] = 7
7

julia> S
3×4 SharedArray{Int64,2}:
 2  2  3  4
 2  3  3  4
 2  7  4  4
```

Base.localindexes() (page 495) provides disjoint one-dimensional ranges of indexes, and is sometimes convenient for splitting up tasks among processes. You can, of course, divide the work any way you wish:


```
julia> S = SharedArray{Int, (3,4), init = S -> S[indexpids(S):length(procs(S)):length(S)] = myid()}
3×4 SharedArray{Int64,2}:
 2  2  2  2
 3  3  3  3
 4  4  4  4
```

Since all processes have access to the underlying data, you do have to be careful not to set up conflicts. For example:

```
@sync begin
    for p in procs(S)
        @async begin
            remotecall_wait(fill!, p, S, p)
        end
    end
end
```

would result in undefined behavior. Because each process fills the *entire* array with its own `pid`, whichever process is the last to execute (for any particular element of `S`) will have its `pid` retained.

As a more extended and complex example, consider running the following “kernel” in parallel:

```
q[i,j,t+1] = q[i,j,t] + u[i,j,t]
```

In this case, if we try to split up the work using a one-dimensional index, we are likely to run into trouble: if `q[i, j, t]` is near the end of the block assigned to one worker and `q[i, j, t+1]` is near the beginning of the block assigned to another, it’s very likely that `q[i, j, t]` will not be ready at the time it’s needed for computing `q[i, j, t+1]`. In such cases, one is better off chunking the array manually. Let’s split along the second dimension:

```
# This function returns the (irange,jrange) indexes assigned to this worker
@everywhere function myrange(q::SharedArray)
    idx = indexpids(q)
    if idx == 0
        # This worker is not assigned a piece
        return 1:0, 1:0
    end
    nchunks = length(procs(q))
    splits = [round{Int, s} for s in linspace(0, size(q,2), nchunks+1)]
    1:size(q,1), splits[idx]+1:splits[idx+1]
end

# Here's the kernel
@everywhere function advection_chunk!(q, u, irange, jrange, trange)
    @show (irange, jrange, trange) # display so we can see what's happening
    for t in trange, j in jrange, i in irange
        q[i,j,t+1] = q[i,j,t] + u[i,j,t]
    end
    q
end

# Here's a convenience wrapper for a SharedArray implementation
@everywhere advection_shared_chunk!(q, u) = advection_chunk!(q, u, myrange(q)..., 1:size(q,3)-1)
```

Now let’s compare three different versions, one that runs in a single process:

```
advection_serial!(q, u) = advection_chunk!(q, u, 1:size(q,1), 1:size(q,2), 1:size(q,3)-1)
```

one that uses `@parallel` (page 493):

```
function advection_parallel!(q, u)
    for t = 1:size(q,3)-1
```

```
@sync @parallel for j = 1:size(q,2)
    for i = 1:size(q,1)
        q[i,j,t+1] = q[i,j,t] + u[i,j,t]
    end
end
end
q
end
```

and one that delegates in chunks:

```
function advection_shared!(q, u)
    @sync begin
        for p in procs(q)
            @async remotecall_wait(advection_shared_chunk!, p, q, u)
        end
    end
    q
end
```

If we create SharedArrays and time these functions, we get the following results (with `julia -p 4`):

```
q = SharedArray{Float64, (500,500,500)}
u = SharedArray{Float64, (500,500,500)}

# Run once to JIT-compile
advection_serial!(q, u)
advection_parallel!(q, u)
advection_shared!(q, u)

# Now the real results:
julia> @time advection_serial!(q, u);
(irange,jrange,trange) = (1:500,1:500,1:499)
830.220 milliseconds (216 allocations: 13820 bytes)

julia> @time advection_parallel!(q, u);
2.495 seconds (3999 k allocations: 289 MB, 2.09% gc time)

julia> @time advection_shared!(q, u);
From worker 2: (irange,jrange,trange) = (1:500,1:125,1:499)
From worker 4: (irange,jrange,trange) = (1:500,251:375,1:499)
From worker 3: (irange,jrange,trange) = (1:500,126:250,1:499)
From worker 5: (irange,jrange,trange) = (1:500,376:500,1:499)
238.119 milliseconds (2264 allocations: 169 KB)
```

The biggest advantage of `advection_shared!` is that it minimizes traffic among the workers, allowing each to compute for an extended time on the assigned piece.

22.11 Shared Arrays and Distributed Garbage Collection

Like remote references, shared arrays are also dependent on garbage collection on the creating node to release references from all participating workers. Code which creates many short lived shared array objects would benefit from explicitly finalizing these objects as soon as possible. This results in both memory and file handles mapping the shared segment being released sooner.

22.12 ClusterManagers

The launching, management and networking of Julia processes into a logical cluster is done via cluster managers. A `ClusterManager` is responsible for

- launching worker processes in a cluster environment
- managing events during the lifetime of each worker
- optionally, providing data transport

A Julia cluster has the following characteristics:

- The initial Julia process, also called the `master`, is special and has an `id` of 1.
- Only the `master` process can add or remove worker processes.
- All processes can directly communicate with each other.

Connections between workers (using the in-built TCP/IP transport) is established in the following manner:

- `addprocs()` (page 488) is called on the master process with a `ClusterManager` object.
- `addprocs()` (page 488) calls the appropriate `launch()` (page 498) method which spawns required number of worker processes on appropriate machines.
- Each worker starts listening on a free port and writes out its host and port information to `STDOUT` (page 561).
- The cluster manager captures the `STDOUT` (page 561) of each worker and makes it available to the master process.
- The master process parses this information and sets up TCP/IP connections to each worker.
- Every worker is also notified of other workers in the cluster.
- Each worker connects to all workers whose `id` is less than the worker's own `id`.
- In this way a mesh network is established, wherein every worker is directly connected with every other worker.

While the default transport layer uses plain `TCPSocket`, it is possible for a Julia cluster to provide its own transport.

Julia provides two in-built cluster managers:

- `LocalManager`, used when `addprocs()` (page 488) or `addprocs(np::Integer)` (page 488) are called
- `SSHManager`, used when `addprocs(hostnames::Array)` (page 488) is called with a list of hostnames

`LocalManager` is used to launch additional workers on the same host, thereby leveraging multi-core and multi-processor hardware.

Thus, a minimal cluster manager would need to:

- be a subtype of the abstract `ClusterManager`
- implement `launch()` (page 498), a method responsible for launching new workers
- implement `manage()` (page 499), which is called at various events during a worker's lifetime (for example, sending an interrupt signal)

`addprocs(manager::FooManager)` (page 488) requires `FooManager` to implement:

```
function launch(manager::FooManager, params::Dict, launched::Array, c::Condition)
    ...
end

function manage(manager::FooManager, id::Integer, config::WorkerConfig, op::Symbol)
```

```
...  
end
```

As an example let us see how the `LocalManager`, the manager responsible for starting workers on the same host, is implemented:

```
immutable LocalManager <: ClusterManager  
    np::Integer  
end  
  
function launch(manager::LocalManager, params::Dict, launched::Array, c::Condition)  
    ...  
end  
  
function manage(manager::LocalManager, id::Integer, config::WorkerConfig, op::Symbol)  
    ...  
end
```

The `launch()` (page 498) method takes the following arguments:

- `manager::ClusterManager`: the cluster manager that `addprocs()` (page 488) is called with
- `params::Dict`: all the keyword arguments passed to `addprocs()` (page 488)
- `launched::Array`: the array to append one or more `WorkerConfig` objects to
- `c::Condition`: the condition variable to be notified as and when workers are launched

The `launch()` (page 498) method is called asynchronously in a separate task. The termination of this task signals that all requested workers have been launched. Hence the `launch()` (page 498) function MUST exit as soon as all the requested workers have been launched.

Newly launched workers are connected to each other, and the master process, in an all-to-all manner. Specifying the command argument `--worker <cookie>` results in the launched processes initializing themselves as workers and connections being set up via TCP/IP sockets. Optionally, `--bind-to bind_addr[:port]` may also be specified to enable other workers to connect to it at the specified `bind_addr` and `port`. This is useful for multi-homed hosts.

As an example of a non-TCP/IP transport, an implementation may choose to use MPI, in which case `--worker` must NOT be specified. Instead, newly launched workers should call `init_worker(cookie)` before using any of the parallel constructs.

For every worker launched, the `launch()` (page 498) method must add a `WorkerConfig` object (with appropriate fields initialized) to `launched`

```
type WorkerConfig  
    # Common fields relevant to all cluster managers  
    io::Nullable{IO}  
    host::Nullable{AbstractString}  
    port::Nullable{Integer}  
  
    # Used when launching additional workers at a host  
    count::Nullable{Union{Int, Symbol}}  
    exename::Nullable{AbstractString}  
    exeflags::Nullable{Cmd}  
  
    # External cluster managers can use this to store information at a per-worker level  
    # Can be a dict if multiple fields need to be stored.  
    userdata::Nullable{Any}  
  
    # SSHManager / SSH tunnel connections to workers  
    tunnel::Nullable{Bool}
```

```

bind_addr::Nullable{AbstractString}
sshflags::Nullable{Cmd}
max_parallel::Nullable{Integer}

connect_at::Nullable{Any}

.....
end

```

Most of the fields in `WorkerConfig` are used by the inbuilt managers. Custom cluster managers would typically specify only `io` or `host / port`:

- If `io` is specified, it is used to read host/port information. A Julia worker prints out its bind address and port at startup. This allows Julia workers to listen on any free port available instead of requiring worker ports to be configured manually.
- If `io` is not specified, `host` and `port` are used to connect.
- `count`, `exename` and `exeflags` are relevant for launching additional workers from a worker. For example, a cluster manager may launch a single worker per node, and use that to launch additional workers.
 - `count` with an integer value `n` will launch a total of `n` workers.
 - `count` with a value of `:auto` will launch as many workers as the number of cores on that machine.
 - `exename` is the name of the `julia` executable including the full path.
 - `exeflags` should be set to the required command line arguments for new workers.
- `tunnel`, `bind_addr`, `sshflags` and `max_parallel` are used when a ssh tunnel is required to connect to the workers from the master process.
- `userdata` is provided for custom cluster managers to store their own worker-specific information.

`manage(manager::FooManager, id::Integer, config::WorkerConfig, op::Symbol)` is called at different times during the worker's lifetime with appropriate `op` values:

- with `:register/:deregister` when a worker is added / removed from the Julia worker pool.
- with `:interrupt` when `interrupt(workers)` is called. The `ClusterManager` should signal the appropriate worker with an interrupt signal.
- with `:finalize` for cleanup purposes.

22.13 Cluster Managers with Custom Transports

Replacing the default TCP/IP all-to-all socket connections with a custom transport layer is a little more involved. Each Julia process has as many communication tasks as the workers it is connected to. For example, consider a Julia cluster of 32 processes in an all-to-all mesh network:

- Each Julia process thus has 31 communication tasks.
- Each task handles all incoming messages from a single remote worker in a message-processing loop.
- The message-processing loop waits on an `IO` object (for example, a `TCP Socket` in the default implementation), reads an entire message, processes it and waits for the next one.
- Sending messages to a process is done directly from any Julia task—not just communication tasks—again, via the appropriate `IO` object.

Replacing the default transport requires the new implementation to set up connections to remote workers and to provide appropriate IO objects that the message-processing loops can wait on. The manager-specific callbacks to be implemented are:

```
connect(manager::FooManager, pid::Integer, config::WorkerConfig)
kill(manager::FooManager, pid::Int, config::WorkerConfig)
```

The default implementation (which uses TCP/IP sockets) is implemented as `connect(manager::ClusterManager, pid::Integer, config::WorkerConfig)`.

`connect` should return a pair of IO objects, one for reading data sent from worker `pid`, and the other to write data that needs to be sent to worker `pid`. Custom cluster managers can use an in-memory `BufferStream` as the plumbing to proxy data between the custom, possibly non-IO transport and Julia's in-built parallel infrastructure.

A `BufferStream` is an in-memory `IOBuffer` which behaves like an IO—it is a stream which can be handled asynchronously.

Folder `examples/clustermanager/0mq` contains an example of using ZeroMQ to connect Julia workers in a star topology with a 0MQ broker in the middle. Note: The Julia processes are still all *logically* connected to each other—any worker can message any other worker directly without any awareness of 0MQ being used as the transport layer.

When using custom transports:

- Julia workers must NOT be started with `--worker`. Starting with `--worker` will result in the newly launched workers defaulting to the TCP/IP socket transport implementation.
- For every incoming logical connection with a worker, `Base.process_messages(rd::IO, wr::IO)()` must be called. This launches a new task that handles reading and writing of messages from/to the worker represented by the IO objects.
- `init_worker(cookie, manager::FooManager)` MUST be called as part of worker process initialization.
- Field `connect_at::Any` in `WorkerConfig` can be set by the cluster manager when `launch()` (page 498) is called. The value of this field is passed in in all `connect()` (page 499) callbacks. Typically, it carries information on *how to connect* to a worker. For example, the TCP/IP socket transport uses this field to specify the `(host, port)` tuple at which to connect to a worker.

`kill(manager, pid, config)` is called to remove a worker from the cluster. On the master process, the corresponding IO objects must be closed by the implementation to ensure proper cleanup. The default implementation simply executes an `exit()` call on the specified remote worker.

`examples/clustermanager/simple` is an example that shows a simple implementation using UNIX domain sockets for cluster setup.

22.14 Network Requirements for LocalManager and SSHManager

Julia clusters are designed to be executed on already secured environments on infrastructure such as local laptops, departmental clusters, or even the cloud. This section covers network security requirements for the inbuilt `LocalManager` and `SSHManager`:

- The master process does not listen on any port. It only connects out to the workers.
- Each worker binds to only one of the local interfaces and listens on the first free port starting from 9009.
- `LocalManager`, used by `addprocs(N)`, by default binds only to the loopback interface. This means that workers started later on remote hosts (or by anyone with malicious intentions) are unable to connect to the cluster. An `addprocs(4)` followed by an `addprocs(["remote_host"])` will fail. Some users may

need to create a cluster comprising their local system and a few remote systems. This can be done by explicitly requesting `LocalManager` to bind to an external network interface via the `restrict` keyword argument: `addprocs(4; restrict=false)`.

- `SSHManager`, used by `addprocs(list_of_remote_hosts)`, launches workers on remote hosts via SSH. By default SSH is only used to launch Julia workers. Subsequent master-worker and worker-worker connections use plain, unencrypted TCP/IP sockets. The remote hosts must have passwordless login enabled. Additional SSH flags or credentials may be specified via keyword argument `sshflags`.
- `addprocs(list_of_remote_hosts; tunnel=true, sshflags=<ssh keys and other flags>)` is useful when we wish to use SSH connections for master-worker too. A typical scenario for this is a local laptop running the Julia REPL (i.e., the master) with the rest of the cluster on the cloud, say on Amazon EC2. In this case only port 22 needs to be opened at the remote cluster coupled with SSH client authenticated via public key infrastructure (PKI). Authentication credentials can be supplied via `sshflags`, for example `sshflags='-e <keyfile>'`.

Note that worker-worker connections are still plain TCP and the local security policy on the remote cluster must allow for free connections between worker nodes, at least for ports 9009 and above.

Securing and encrypting all worker-worker traffic (via SSH) or encrypting individual messages can be done via a custom `ClusterManager`.

22.15 Cluster Cookie

All processes in a cluster share the same cookie which, by default, is a randomly generated string on the master process:

- `Base.cluster_cookie()` (page 494) returns the cookie, while `Base.cluster_cookie(cookie)` () sets it and returns the new cookie.
- All connections are authenticated on both sides to ensure that only workers started by the master are allowed to connect to each other.
- The cookie must be passed to the workers at startup via argument `--worker <cookie>`. Custom `ClusterManagers` can retrieve the cookie on the master by calling `Base.cluster_cookie()` (page 494). Cluster managers not using the default TCP/IP transport (and hence not specifying `--worker`) must call `init_worker(cookie, manager)` with the same cookie as on the master.

Note that environments requiring higher levels of security can implement this via a custom `ClusterManager`. For example, cookies can be pre-shared and hence not specified as a startup argument.

22.16 Specifying Network Topology (Experimental)

The keyword argument `topology` passed to `addprocs` is used to specify how the workers must be connected to each other:

- `:all_to_all`, the default: all workers are connected to each other.
- `:master_slave`: only the driver process, i.e. `pid 1`, has connections to the workers.
- `:custom`: the `launch` method of the cluster manager specifies the connection topology via the fields `ident` and `connect_ids` in `WorkerConfig`. A worker with a cluster-manager-provided identity `ident` will connect to all workers specified in `connect_ids`.

Currently, sending a message between unconnected workers results in an error. This behaviour, as with the functionality and interface, should be considered experimental in nature and may change in future releases.

22.17 Multi-Threading (Experimental)

In addition to tasks, remote calls, and remote references, Julia from v0.5 forwards will natively support multi-threading. Note that this section is experimental and the interfaces may change in the future.

22.17.1 Setup

By default, Julia starts up with a single thread of execution. This can be verified by using the command `Threads.nthreads()` (page 495):

```
julia> Threads.nthreads()
1
```

The number of threads Julia starts up with is controlled by an environment variable called `JULIA_NUM_THREADS`. Now, let's start up Julia with 4 threads:

```
export JULIA_NUM_THREADS=4
```

(The above command works on bourne shells on Linux and OSX. Note that if you're using a C shell on these platforms, you should use the keyword `set` instead of `export`. If you're on Windows, start up the command line in the location of `julia.exe` and use `set` instead of `export`.)

Let's verify there are 4 threads at our disposal.

```
julia> Threads.nthreads()
4
```

But we are currently on the master thread. To check, we use the command `Threads.threadid()` (page 495)

```
julia> Threads.threadid()
1
```

22.17.2 The `@threads` Macro

Let's work a simple example using our native threads. Let us create an array of zeros:

```
julia> a = zeros{Float64,1}(10)
10-element Array{Float64,1}:
 0.0
 0.0
 0.0
 0.0
 0.0
 0.0
 0.0
 0.0
 0.0
 0.0
```

Let us operate on this array simultaneously using 4 threads. We'll have each thread write its thread ID into each location.

Julia supports parallel loops using the `Threads.@threads` (page 495) macro. This macro is affixed in front of a `for` loop to indicate to Julia that the loop is a multi-threaded region:


```
Threads.@threads for i = 1:10
    a[i] = Threads.threadid()
end
```

The iteration space is split amongst the threads, after which each thread writes its thread ID to its assigned locations:

```
julia> a
10-element Array{Float64,1}:
 1.0
 1.0
 1.0
 2.0
 2.0
 2.0
 3.0
 3.0
 4.0
 4.0
```

Note that `Threads.@threads` (page 495) does not have an optional reduction parameter like `@parallel` (page 493).

22.18 @threadcall (Experimental)

All I/O tasks, timers, REPL commands, etc are multiplexed onto a single OS thread via an event loop. A patched version of libuv (<http://docs.libuv.org/en/v1.x/>) provides this functionality. Yield points provide for co-operatively scheduling multiple tasks onto the same OS thread. I/O tasks and timers yield implicitly while waiting for the event to occur. Calling `yield()` explicitly allows for other tasks to be scheduled.

Thus, a task executing a `ccall` effectively prevents the Julia scheduler from executing any other tasks till the call returns. This is true for all calls into external libraries. Exceptions are calls into custom C code that call back into Julia (which may then yield) or C code that calls `jl_yield()` (C equivalent of `yield()`).

Note that while Julia code runs on a single thread (by default), libraries used by Julia may launch their own internal threads. For example, the BLAS library may start as many threads as there are cores on a machine.

The `@threadcall` macro addresses scenarios where we do not want a `ccall` to block the main Julia event loop. It schedules a C function for execution in a separate thread. A threadpool with a default size of 4 is used for this. The size of the threadpool is controlled via environment variable `UV_THREADPOOL_SIZE`. While waiting for a free thread, and during function execution once a thread is available, the requesting task (on the main Julia event loop) yields to other tasks. Note that `@threadcall` does not return till the execution is complete. From a user point of view, it is therefore a blocking call like other Julia APIs.

It is very important that the called function does not call back into Julia.

`@threadcall` may be removed/changed in future versions of Julia.

Date and DateTime

The *Dates* (page 585) module provides two types for working with dates: *Date* (page 587) and *DateTime* (page 587), representing day and millisecond precision, respectively; both are subtypes of the abstract *TimeType* (page 587). The motivation for distinct types is simple: some operations are much simpler, both in terms of code and mental reasoning, when the complexities of greater precision don't have to be dealt with. For example, since the *Date* (page 587) type only resolves to the precision of a single date (i.e. no hours, minutes, or seconds), normal considerations for time zones, daylight savings/summer time, and leap seconds are unnecessary and avoided.

Both *Date* (page 587) and *DateTime* (page 587) are basically immutable `Int64` wrappers. The single instant field of either type is actually a `UTInstant{P}` type, which represents a continuously increasing machine timeline based on the UT second¹. The *DateTime* (page 587) type is *timezone-unaware* (in Python parlance) or is analogous to a *LocalDateTime* in Java 8. Additional time zone functionality can be added through the *TimeZones.jl* package, which compiles the *IANA time zone database*. Both *Date* (page 587) and *DateTime* (page 587) are based on the ISO 8601 standard, which follows the proleptic Gregorian calendar. One note is that the ISO 8601 standard is particular about BC/BCE dates. In general, the last day of the BC/BCE era, 1-12-31 BC/BCE, was followed by 1-1-1 AD/CE, thus no year zero exists. The ISO standard, however, states that 1 BC/BCE is year zero, so 0000-12-31 is the day before 0001-01-01, and year -0001 (yes, negative one for the year) is 2 BC/BCE, year -0002 is 3 BC/BCE, etc.

23.1 Constructors

Date (page 587) and *DateTime* (page 587) types can be constructed by integer or *Period* (page 587) types, by parsing, or through adjusters (more on those later):

```
julia> DateTime(2013)
2013-01-01T00:00:00

julia> DateTime(2013, 7)
2013-07-01T00:00:00

julia> DateTime(2013, 7, 1)
2013-07-01T00:00:00

julia> DateTime(2013, 7, 1, 12)
2013-07-01T12:00:00
```

¹ The notion of the UT second is actually quite fundamental. There are basically two different notions of time generally accepted, one based on the physical rotation of the earth (one full rotation = 1 day), the other based on the SI second (a fixed, constant value). These are radically different! Think about it, a “UT second”, as defined relative to the rotation of the earth, may have a different absolute length depending on the day! Anyway, the fact that *Date* (page 587) and *DateTime* (page 587) are based on UT seconds is a simplifying, yet honest assumption so that things like leap seconds and all their complexity can be avoided. This basis of time is formally called *UT* or *UT1*. Basing types on the UT second basically means that every minute has 60 seconds and every day has 24 hours and leads to more natural calculations when working with calendar dates.

```
julia> DateTime(2013, 7, 1, 12, 30)
2013-07-01T12:30:00

julia> DateTime(2013, 7, 1, 12, 30, 59)
2013-07-01T12:30:59

julia> DateTime(2013, 7, 1, 12, 30, 59, 1)
2013-07-01T12:30:59.001

julia> Date(2013)
2013-01-01

julia> Date(2013, 7)
2013-07-01

julia> Date(2013, 7, 1)
2013-07-01

julia> Date(Dates.Year(2013), Dates.Month(7), Dates.Day(1))
2013-07-01

julia> Date(Dates.Month(7), Dates.Year(2013))
2013-07-01
```

Date (page 587) or *DateTime* (page 587) parsing is accomplished by the use of format strings. Format strings work by the notion of defining *delimited* or *fixed-width* “slots” that contain a period to parse and passing the text to parse and format string to a *Date* (page 587) or *DateTime* (page 587) constructor, of the form `Date("2015-01-01", "y-m-d")` or `DateTime("20150101", "yyyymmdd")`.

Delimited slots are marked by specifying the delimiter the parser should expect between two subsequent periods; so "y-m-d" lets the parser know that between the first and second slots in a date string like "2014-07-16", it should find the - character. The y, m, and d characters let the parser know which periods to parse in each slot.

Fixed-width slots are specified by repeating the period character the number of times corresponding to the width with no delimiter between characters. So "yyyymmdd" would correspond to a date string like "20140716". The parser distinguishes a fixed-width slot by the absence of a delimiter, noting the transition "yyyymm" from one period character to the next.

Support for text-form month parsing is also supported through the u and U characters, for abbreviated and full-length month names, respectively. By default, only English month names are supported, so u corresponds to “Jan”, “Feb”, “Mar”, etc. And U corresponds to “January”, “February”, “March”, etc. Similar to other name=>value mapping functions *dayname()* (page 591) and *monthname()* (page 591), custom locales can be loaded by passing in the `locale=>Dict{String, Int}` mapping to the `MONTHTOVALUEABBR` and `MONTHTOVALUE` dicts for abbreviated and full-name month names, respectively.

One note on parsing performance: using the `Date(date_string, format_string)` function is fine if only called a few times. If there are many similarly formatted date strings to parse however, it is much more efficient to first create a *Dates.DateFormat* (page 589), and pass it instead of a raw format string.

```
julia> df = Dates.DateFormat("y-m-d");

julia> dt = Date("2015-01-01", df)
2015-01-01

julia> dt2 = Date("2015-01-02", df)
2015-01-02
```

A full suite of parsing and formatting tests and examples is available in [tests/dates/io.jl](https://github.com/JuliaLang/julia/blob/master/test/dates/io.jl).

23.2 Durations/Comparisons

Finding the length of time between two *Date* (page 587) or *DateTime* (page 587) is straightforward given their underlying representation as `UTInstant{Day}` and `UTInstant{Millisecond}`, respectively. The difference between *Date* (page 587) is returned in the number of *Day* (page 587), and *DateTime* (page 587) in the number of *Millisecond* (page 587). Similarly, comparing *TimeType* (page 587) is a simple matter of comparing the underlying machine instants (which in turn compares the internal `Int64` values).

```
julia> dt = Date(2012,2,29)
2012-02-29

julia> dt2 = Date(2000,2,1)
2000-02-01

julia> dump(dt)
Date
  instant: UTInstant{Day}
    periods: Day
      value: Int64 734562

julia> dump(dt2)
Date
  instant: UTInstant{Day}
    periods: Day
      value: Int64 730151

julia> dt > dt2
true

julia> dt != dt2
true

julia> dt + dt2
Operation not defined for TimeTypes

julia> dt * dt2
Operation not defined for TimeTypes

julia> dt / dt2
Operation not defined for TimeTypes

julia> dt - dt2
4411 days

julia> dt2 - dt
-4411 days

julia> dt = DateTime(2012,2,29)
2012-02-29T00:00:00

julia> dt2 = DateTime(2000,2,1)
2000-02-01T00:00:00

julia> dt - dt2
381110402000 milliseconds
```

23.3 Accessor Functions

Because the *Date* (page 587) and *DateTime* (page 587) types are stored as single `Int64` values, date parts or fields can be retrieved through accessor functions. The lowercase accessors return the field as an integer:

```
julia> t = Date(2014,1,31)
2014-01-31

julia> Dates.year(t)
2014

julia> Dates.month(t)
1

julia> Dates.week(t)
5

julia> Dates.day(t)
31
```

While propercase return the same value in the corresponding *Period* (page 587) type:

```
julia> Dates.Year(t)
2014 years

julia> Dates.Day(t)
31 days
```

Compound methods are provided, as they provide a measure of efficiency if multiple fields are needed at the same time:

```
julia> Dates.yearmonth(t)
(2014,1)

julia> Dates.monthday(t)
(1,31)

julia> Dates.yearmonthday(t)
(2014,1,31)
```

One may also access the underlying `UTInstant` or integer value:

```
julia> dump(t)
Date
instant: UTInstant{Day}
  periods: Day
    value: Int64 735264

julia> t.instant
UTInstant{Day} (735264 days)

julia> Dates.value(t)
735264
```

23.4 Query Functions

Query functions provide calendrical information about a *TimeType* (page 587). They include information about the day of the week:

```
julia> t = Date(2014,1,31)
2014-01-31

julia> Dates.dayofweek(t)
5

julia> Dates.dayname(t)
"Friday"

julia> Dates.dayofweekofmonth(t)
5 # 5th Friday of January
```

Month of the year:

```
julia> Dates.monthname(t)
"January"

julia> Dates.daysinmonth(t)
31
```

As well as information about the *TimeType* (page 587)'s year and quarter:

```
julia> Dates.isleapyear(t)
false

julia> Dates.dayofyear(t)
31

julia> Dates.quarterofyear(t)
1

julia> Dates.dayofquarter(t)
31
```

The *dayname()* (page 591) and *monthname()* (page 591) methods can also take an optional *locale* keyword that can be used to return the name of the day or month of the year for other languages/locales:

```
julia> const french_daysofweek = Dict{Int,String}(1=>"Lundi", 2=>"Mardi", 3=>"Mercredi", 4=>"Jeudi", 5=>"Vendredi", 6=>"Samedi", 7=>"Dimanche")

# Load the mapping into the Dates module under locale name "french"
julia> Dates.VALUETODAYOFWEEK["french"] = french_daysofweek;

julia> Dates.dayname(t; locale="french")
"Vendredi"
```

Similarly for the *monthname()* (page 591) function, a mapping of *locale*=>*Dict{Int,String}* should be loaded in *VALUETOMONTH*.

23.5 TimeType-Period Arithmetic

It's good practice when using any language/date framework to be familiar with how date-period arithmetic is handled as there are some *tricky issues* to deal with (though much less so for day-precision types).

The *Dates* (page 585) module approach tries to follow the simple principle of trying to change as little as possible when doing *Period* (page 587) arithmetic. This approach is also often known as *calendrical* arithmetic or what you would probably guess if someone were to ask you the same calculation in a conversation. Why all the fuss about this? Let's take a classic example: add 1 month to January 31st, 2014. What's the answer? Javascript will say *March 3* (assumes 31 days). PHP says *March 2* (assumes 30 days). The fact is, there is no right answer. In the *Dates* (page 585) module, it gives the result of February 28th. How does it figure that out? I like to think of the classic 7-7-7 gambling game in casinos.

Now just imagine that instead of 7-7-7, the slots are Year-Month-Day, or in our example, 2014-01-31. When you ask to add 1 month to this date, the month slot is incremented, so now we have 2014-02-31. Then the day number is checked if it is greater than the last valid day of the new month; if it is (as in the case above), the day number is adjusted down to the last valid day (28). What are the ramifications with this approach? Go ahead and add another month to our date, 2014-02-28 + Month(1) == 2014-03-28. What? Were you expecting the last day of March? Nope, sorry, remember the 7-7-7 slots. As few slots as possible are going to change, so we first increment the month slot by 1, 2014-03-28, and boom, we're done because that's a valid date. On the other hand, if we were to add 2 months to our original date, 2014-01-31, then we end up with 2014-03-31, as expected. The other ramification of this approach is a loss in associativity when a specific ordering is forced (i.e. adding things in different orders results in different outcomes). For example:

```
julia> (Date(2014,1,29)+Dates.Day(1)) + Dates.Month(1)
2014-02-28

julia> (Date(2014,1,29)+Dates.Month(1)) + Dates.Day(1)
2014-03-01
```

What's going on there? In the first line, we're adding 1 day to January 29th, which results in 2014-01-30; then we add 1 month, so we get 2014-02-30, which then adjusts down to 2014-02-28. In the second example, we add 1 month *first*, where we get 2014-02-29, which adjusts down to 2014-02-28, and *then* add 1 day, which results in 2014-03-01. One design principle that helps in this case is that, in the presence of multiple Periods, the operations will be ordered by the Periods' *types*, not their value or positional order; this means Year will always be added first, then Month, then Week, etc. Hence the following *does* result in associativity and Just Works:

```
julia> Date(2014,1,29) + Dates.Day(1) + Dates.Month(1)
2014-03-01

julia> Date(2014,1,29) + Dates.Month(1) + Dates.Day(1)
2014-03-01
```

Tricky? Perhaps. What is an innocent *Dates* (page 585) user to do? The bottom line is to be aware that explicitly forcing a certain associativity, when dealing with months, may lead to some unexpected results, but otherwise, everything should work as expected. Thankfully, that's pretty much the extent of the odd cases in date-period arithmetic when dealing with time in UT (avoiding the "joys" of dealing with daylight savings, leap seconds, etc.).

As a bonus, all period arithmetic objects work directly with ranges:

```
julia> dr = Date(2014,1,29):Date(2014,2,3)
2014-01-29:1 day:2014-02-03

julia> collect(dr)
6-element Array{Date,1}:
 2014-01-29
 2014-01-30
 2014-01-31
 2014-02-01
 2014-02-02
 2014-02-03

julia> dr = Date(2014,1,29):Dates.Month(1):Date(2014,07,29)
```



```
2014-01-29:1 month:2014-07-29
```

```
julia> collect(dr)
7-element Array{Date,1}:
 2014-01-29
 2014-02-28
 2014-03-29
 2014-04-29
 2014-05-29
 2014-06-29
 2014-07-29
```

23.6 Adjuster Functions

As convenient as date-period arithmetics are, often the kinds of calculations needed on dates take on a *calendrical* or *temporal* nature rather than a fixed number of periods. Holidays are a perfect example; most follow rules such as “Memorial Day = Last Monday of May”, or “Thanksgiving = 4th Thursday of November”. These kinds of temporal expressions deal with rules relative to the calendar, like first or last of the month, next Tuesday, or the first and third Wednesdays, etc.

The *Dates* (page 585) module provides the *adjuster* API through several convenient methods that aid in simply and succinctly expressing temporal rules. The first group of adjuster methods deal with the first and last of weeks, months, quarters, and years. They each take a single *TimeType* (page 587) as input and return or *adjust to* the first or last of the desired period relative to the input.

```
# Adjusts the input to the Monday of the input's week
julia> Dates.firstdayofweek(Date(2014,7,16))
2014-07-14

# Adjusts to the last day of the input's month
julia> Dates.lastdayofmonth(Date(2014,7,16))
2014-07-31

# Adjusts to the last day of the input's quarter
julia> Dates.lastdayofquarter(Date(2014,7,16))
2014-09-30
```

The next two higher-order methods, *tonext()* (page 592), and *toprev()* (page 592), generalize working with temporal expressions by taking a *DateFunction* as first argument, along with a starting *TimeType* (page 587). A *DateFunction* is just a function, usually anonymous, that takes a single *TimeType* (page 587) as input and returns a *Bool*, true indicating a satisfied adjustment criterion. For example:

```
julia> istuesday = x->Dates.dayofweek(x) == Dates.Tuesday # Returns true if the day of the week of x is Tuesday
(anonymous function)

julia> Dates.tonext(istuesday, Date(2014,7,13)) # 2014-07-13 is a Sunday
2014-07-15

# Convenience method provided for day of the week adjustments
julia> Dates.tonext(Date(2014,7,13), Dates.Tuesday)
2014-07-15
```

This is useful with the do-block syntax for more complex temporal expressions:

```
julia> Dates.tonext(Date(2014,7,13)) do x
    # Return true on the 4th Thursday of November (Thanksgiving)
```

```
Dates.dayofweek(x) == Dates.Thursday &&
Dates.dayofweekofmonth(x) == 4 &&
Dates.month(x) == Dates.November
end
2014-11-27
```

The `Base.filter()` (page 392) method can be used to obtain all valid dates/moments in a specified range:

```
# Pittsburgh street cleaning; Every 2nd Tuesday from April to November
# Date range from January 1st, 2014 to January 1st, 2015
julia> dr = Dates.Date(2014):Dates.Date(2015);
julia> filter(dr) do x
    Dates.dayofweek(x) == Dates.Tue &&
    Dates.April <= Dates.month(x) <= Dates.Nov &&
    Dates.dayofweekofmonth(x) == 2
end
8-element Array{Date,1}:
2014-04-08
2014-05-13
2014-06-10
2014-07-08
2014-08-12
2014-09-09
2014-10-14
2014-11-11
```

Additional examples and tests are available in `test/dates/adjusters.jl`.

23.7 Period Types

Periods are a human view of discrete, sometimes irregular durations of time. Consider 1 month; it could represent, in days, a value of 28, 29, 30, or 31 depending on the year and month context. Or a year could represent 365 or 366 days in the case of a leap year. *Period* (page 587) types are simple `Int64` wrappers and are constructed by wrapping any `Int64` convertible type, i.e. `Year(1)` or `Month(3.0)`. Arithmetic between *Period* (page 587) of the same type behave like integers, and limited `Period`-`Real` arithmetic is available.

```
julia> y1 = Dates.Year(1)
1 year

julia> y2 = Dates.Year(2)
2 years

julia> y3 = Dates.Year(10)
10 years

julia> y1 + y2
3 years

julia> div(y3,y2)
5 years

julia> y3 - y2
8 years

julia> y3 * y2
20 years
```

```
julia> y3 % y2
0 years

julia> y1 + 20
21 years

julia> div(y3,3) # mirrors integer division
3 years
```

23.8 Rounding

Date (page 587) and *DateTime* (page 587) values can be rounded to a specified resolution (e.g., 1 month or 15 minutes) with *floor()* (page 593), *ceil()* (page 593), or *round()* (page 594):

```
julia> floor(Date(1985, 8, 16), Dates.Month)
1985-08-01

julia> ceil(DateTime(2013, 2, 13, 0, 31, 20), Dates.Minute(15))
2013-02-13T00:45:00

julia> round(DateTime(2016, 8, 6, 20, 15), Dates.Day)
2016-08-07T00:00:00
```

Unlike the numeric *round()* (page 594) method, which breaks ties toward the even number by default, the *TimeType* (page 587) *round()* (page 594) method uses the `RoundNearestTiesUp` rounding mode. (It’s difficult to guess what breaking ties to nearest “even” *TimeType* (page 587) would entail.) Further details on the available `RoundingMode`s can be found in the [API reference](#).

Rounding should generally behave as expected, but there are a few cases in which the expected behaviour is not obvious.

23.8.1 Rounding Epoch

In many cases, the resolution specified for rounding (e.g., `Dates.Second(30)`) divides evenly into the next largest period (in this case, `Dates.Minute(1)`). But rounding behaviour in cases in which this is not true may lead to confusion. What is the expected result of rounding a *DateTime* (page 587) to the nearest 10 hours?

```
julia> round(DateTime(2016, 7, 17, 11, 55), Dates.Hour(10))
2016-07-17T12:00:00
```

That may seem confusing, given that the hour (12) is not divisible by 10. The reason that `2016-07-17T12:00:00` was chosen is that it is 17,676,660 hours after `0000-01-01T00:00:00`, and 17,676,660 is divisible by 10.

As Julia *Date* (page 587) and *DateTime* (page 587) values are represented according to the ISO 8601 standard, `0000-01-01T00:00:00` was chosen as base (or “rounding epoch”) from which to begin the count of days (and milliseconds) used in rounding calculations. (Note that this differs slightly from Julia’s internal representation of *Date* (page 587)s using Rata Die notation; but since the ISO 8601 standard is most visible to the end user, `0000-01-01T00:00:00` was chosen as the rounding epoch instead of the `0000-12-31T00:00:00` used internally to minimize confusion.)

The only exception to the use of `0000-01-01T00:00:00` as the rounding epoch is when rounding to weeks. Rounding to the nearest week will always return a Monday (the first day of the week as specified by ISO 8601). For this reason, we use `0000-01-03T00:00:00` (the first day of the first week of year 0000, as defined by ISO 8601) as the base when rounding to a number of weeks.

Here is a related case in which the expected behaviour is not necessarily obvious: What happens when we round to the nearest $P(2)$, where P is a *Period* (page 587) type? In some cases (specifically, when $P <: \text{Dates.TimePeriod}$) the answer is clear:

```
julia> round(DateTime(2016, 7, 17, 8, 55, 30), Dates.Hour(2))
2016-07-17T08:00:00

julia> round(DateTime(2016, 7, 17, 8, 55, 30), Dates.Minute(2))
2016-07-17T08:56:00
```

This seems obvious, because two of each of these periods still divides evenly into the next larger order period. But in the case of two months (which still divides evenly into one year), the answer may be surprising:

```
julia> round(DateTime(2016, 7, 17, 8, 55, 30), Dates.Month(2))
2016-07-01T00:00:00
```

Why round to the first day in July, even though it is month 7 (an odd number)? The key is that months are 1-indexed (the first month is assigned 1), unlike hours, minutes, seconds, and milliseconds (the first of which are assigned 0).

This means that rounding a *DateTime* (page 587) to an even multiple of seconds, minutes, hours, or years (because the ISO 8601 specification includes a year zero) will result in a *DateTime* (page 587) with an even value in that field, while rounding a *DateTime* (page 587) to an even multiple of months will result in the months field having an odd value. Because both months and years may contain an irregular number of days, whether rounding to an even number of days will result in an even value in the days field is uncertain.

See the [API reference](#) for additional information on methods exported from the *Dates* (page 585) module.

Running External Programs

Julia borrows backtick notation for commands from the shell, Perl, and Ruby. However, in Julia, writing

```
julia> `echo hello`
`echo hello`
```

differs in several aspects from the behavior in various shells, Perl, or Ruby:

- Instead of immediately running the command, backticks create a `Cmd` object to represent the command. You can use this object to connect the command to others via pipes, run it, and read or write to it.
- When the command is run, Julia does not capture its output unless you specifically arrange for it to. Instead, the output of the command by default goes to `STDOUT` (page 561) as it would using `libc`'s `system` call.
- The command is never run with a shell. Instead, Julia parses the command syntax directly, appropriately interpolating variables and splitting on words as the shell would, respecting shell quoting syntax. The command is run as `julia`'s immediate child process, using `fork` and `exec` calls.

Here's a simple example of running an external program:

```
julia> run(`echo hello`)
hello
```

The `hello` is the output of the `echo` command, sent to `STDOUT` (page 561). The `run` method itself returns `nothing`, and throws an `ErrorException` (page 375) if the external command fails to run successfully.

If you want to read the output of the external command, `readstring()` (page 567) can be used instead:

```
julia> a=readstring(`echo hello`)
"hello\n"

julia> (chomp(a)) == "hello"
true
```

More generally, you can use `open()` (page 561) to read from or write to an external command.

```
julia> open(`less`, "w", STDOUT) do io
    for i = 1:3
        println(io, i)
    end
end

1
2
3
```

24.1 Interpolation

Suppose you want to do something a bit more complicated and use the name of a file in the variable `file` as an argument to a command. You can use `$` for interpolation much as you would in a string literal (see [Strings](#) (page 41)):

```
julia> file = "/etc/passwd"
"/etc/passwd"

julia> `sort $file`
`sort /etc/passwd`
```

A common pitfall when running external programs via a shell is that if a file name contains characters that are special to the shell, they may cause undesirable behavior. Suppose, for example, rather than `/etc/passwd`, we wanted to sort the contents of the file `/Volumes/External HD/data.csv`. Let's try it:

```
julia> file = "/Volumes/External HD/data.csv"
"/Volumes/External HD/data.csv"

julia> `sort $file`
`sort '/Volumes/External HD/data.csv`
```

How did the file name get quoted? Julia knows that `file` is meant to be interpolated as a single argument, so it quotes the word for you. Actually, that is not quite accurate: the value of `file` is never interpreted by a shell, so there's no need for actual quoting; the quotes are inserted only for presentation to the user. This will even work if you interpolate a value as part of a shell word:

```
julia> path = "/Volumes/External HD"
"/Volumes/External HD"

julia> name = "data"
"data"

julia> ext = "csv"
"csv"

julia> `sort $path/$name.$ext`
`sort '/Volumes/External HD/data.csv`
```

As you can see, the space in the `path` variable is appropriately escaped. But what if you *want* to interpolate multiple words? In that case, just use an array (or any other iterable container):

```
julia> files = ["/etc/passwd", "/Volumes/External HD/data.csv"]
2-element Array{String,1}:
"/etc/passwd"
"/Volumes/External HD/data.csv"

julia> `grep foo $files`
`grep foo /etc/passwd '/Volumes/External HD/data.csv`
```

If you interpolate an array as part of a shell word, Julia emulates the shell's `{a,b,c}` argument generation:

```
julia> names = ["foo", "bar", "baz"]
3-element Array{String,1}:
"foo"
"bar"
"baz"
```

```
julia> `grep xylophone $names.txt`
`grep xylophone foo.txt bar.txt baz.txt`
```

Moreover, if you interpolate multiple arrays into the same word, the shell's Cartesian product generation behavior is emulated:

```
julia> names = ["foo", "bar", "baz"]
3-element Array{String,1}:
"foo"
"bar"
"baz"

julia> exts = ["aux", "log"]
2-element Array{String,1}:
"aux"
"log"

julia> `rm -f $names.$exts`
`rm -f foo.aux foo.log bar.aux bar.log baz.aux baz.log`
```

Since you can interpolate literal arrays, you can use this generative functionality without needing to create temporary array objects first:

```
julia> `rm -rf $["foo", "bar", "baz", "qux"].$["aux", "log", "pdf"]`
`rm -rf foo.aux foo.log foo.pdf bar.aux bar.log bar.pdf baz.aux baz.log baz.pdf qux.aux qux.log qux.pdf`
```

24.2 Quoting

Inevitably, one wants to write commands that aren't quite so simple, and it becomes necessary to use quotes. Here's a simple example of a Perl one-liner at a shell prompt:

```
sh$ perl -le '$|=1; for (0..3) { print }'
0
1
2
3
```

The Perl expression needs to be in single quotes for two reasons: so that spaces don't break the expression into multiple shell words, and so that uses of Perl variables like `$|` (yes, that's the name of a variable in Perl), don't cause interpolation. In other instances, you may want to use double quotes so that interpolation *does* occur:

```
sh$ first="A"
sh$ second="B"
sh$ perl -le '$|=1; print for @ARGV' "1: $first" "2: $second"
1: A
2: B
```

In general, the Julia backtick syntax is carefully designed so that you can just cut-and-paste shell commands as is into backticks and they will work: the escaping, quoting, and interpolation behaviors are the same as the shell's. The only difference is that the interpolation is integrated and aware of Julia's notion of what is a single string value, and what is a container for multiple values. Let's try the above two examples in Julia:

```
julia> `perl -le '$|=1; for (0..3) { print }'`
`perl -le '$|=1; for (0..3) { print }'`

julia> run(ans)
```

```
0
1
2
3

julia> first = "A"; second = "B";

julia> `perl -le 'print for @ARGV' "1: $first" "2: $second"`
`perl -le 'print for @ARGV' '1: A' '2: B'`

julia> run(ans)
1: A
2: B
```

The results are identical, and Julia’s interpolation behavior mimics the shell’s with some improvements due to the fact that Julia supports first-class iterable objects while most shells use strings split on spaces for this, which introduces ambiguities. When trying to port shell commands to Julia, try cut and pasting first. Since Julia shows commands to you before running them, you can easily and safely just examine its interpretation without doing any damage.

24.3 Pipelines

Shell metacharacters, such as `|`, `&`, and `>`, are not special inside of Julia’s backticks: unlike in the shell, inside of Julia’s backticks, a pipe is always just a pipe:

```
julia> run(`echo hello | sort`)
hello | sort
```

This expression invokes the `echo` command with three words as arguments: “hello”, “|”, and “sort”. The result is that a single line is printed: “hello | sort”. Inside of backticks, a “|” is just a literal pipe character. How, then, does one construct a pipeline? Instead of using “|” inside of backticks, one uses `pipeline()` (page 372):

```
julia> run(pipeline(`echo hello`, `sort`))
hello
```

This pipes the output of the `echo` command to the `sort` command. Of course, this isn’t terribly interesting since there’s only one line to sort, but we can certainly do much more interesting things:

```
julia> run(pipeline(`cut -d: -f3 /etc/passwd`, `sort -n`, `tail -n5`))
210
211
212
213
214
```

This prints the highest five user IDs on a UNIX system. The `cut`, `sort` and `tail` commands are all spawned as immediate children of the current `julia` process, with no intervening shell process. Julia itself does the work to setup pipes and connect file descriptors that is normally done by the shell. Since Julia does this itself, it retains better control and can do some things that shells cannot.

Julia can run multiple commands in parallel:

```
julia> run(`echo hello` & `echo world`)
world
hello
```

The order of the output here is non-deterministic because the two `echo` processes are started nearly simultaneously, and race to make the first write to the `STDOUT` (page 561) descriptor they share with each other and the `julia` parent

process. Julia lets you pipe the output from both of these processes to another program:

```
julia> run(pipeline(`echo world` & `echo hello`, `sort`))
hello
world
```

In terms of UNIX plumbing, what’s happening here is that a single UNIX pipe object is created and written to by both `echo` processes, and the other end of the pipe is read from by the `sort` command.

IO redirection can be accomplished by passing keyword arguments `stdin`, `stdout`, and `stderr` to the `pipeline` function:

```
pipeline(`do_work`, stdout=pipeline(`sort`, "out.txt"), stderr="errs.txt")
```

24.3.1 Avoiding Deadlock in Pipelines

When reading and writing to both ends of a pipeline from a single process, it is important to avoid forcing the kernel to buffer all of the data.

For example, when reading all of the output from a command, call `readstring(out)`, not `wait(process)`, since the former will actively consume all of the data written by the process, whereas the latter will attempt to store the data in the kernel’s buffers while waiting for a reader to be connected.

Another common solution is to separate the reader and writer of the pipeline into separate Tasks:

```
writer = @async writeall(process, "data")
reader = @async do_compute(readstring(process))
wait(process)
fetch(reader)
```

24.3.2 Complex Example

The combination of a high-level programming language, a first-class command abstraction, and automatic setup of pipes between processes is a powerful one. To give some sense of the complex pipelines that can be created easily, here are some more sophisticated examples, with apologies for the excessive use of Perl one-liners:

```
julia> prefixer(prefix, sleep) = `perl -nle '$|=1; print "'$prefix' ", $_; sleep '$sleep';`;

julia> run(pipeline(`perl -le '$|=1; for(0..9){ print; sleep 1 }'`, prefixer("A",2) & prefixer("B",2)))
A 0
B 1
A 2
B 3
A 4
B 5
A 6
B 7
A 8
B 9
```

This is a classic example of a single producer feeding two concurrent consumers: one `perl` process generates lines with the numbers 0 through 9 on them, while two parallel processes consume that output, one prefixing lines with the letter “A”, the other with the letter “B”. Which consumer gets the first line is non-deterministic, but once that race has been won, the lines are consumed alternately by one process and then the other. (Setting `$|=1` in Perl causes each `print` statement to flush the `STDOUT` (page 561) handle, which is necessary for this example to work. Otherwise all the output is buffered and printed to the pipe at once, to be read by just one consumer process.)

Here is an even more complex multi-stage producer-consumer example:

```
julia> run(pipeline(`perl -le '$|=1; for(0..9){ print; sleep 1 }'`,
    prefixer("X",3) & prefixer("Y",3) & prefixer("Z",3),
    prefixer("A",2) & prefixer("B",2)))
A X 0
B Y 1
A Z 2
B X 3
A Y 4
B Z 5
A X 6
B Y 7
A Z 8
B X 9
```

This example is similar to the previous one, except there are two stages of consumers, and the stages have different latency so they use a different number of parallel workers, to maintain saturated throughput.

We strongly encourage you to try all these examples to see how they work.

Calling C and Fortran Code

Though most code can be written in Julia, there are many high-quality, mature libraries for numerical computing already written in C and Fortran. To allow easy use of this existing code, Julia makes it simple and efficient to call C and Fortran functions. Julia has a “no boilerplate” philosophy: functions can be called directly from Julia without any “glue” code, code generation, or compilation — even from the interactive prompt. This is accomplished just by making an appropriate call with `ccall()` (page 607) syntax, which looks like an ordinary function call.

The code to be called must be available as a shared library. Most C and Fortran libraries ship compiled as shared libraries already, but if you are compiling the code yourself using GCC (or Clang), you will need to use the `-shared` and `-fPIC` options. The machine instructions generated by Julia’s JIT are the same as a native C call would be, so the resulting overhead is the same as calling a library function from C code. (Non-library function calls in both C and Julia can be inlined and thus may have even less overhead than calls to shared library functions. When both libraries and executables are generated by LLVM, it is possible to perform whole-program optimizations that can even optimize across this boundary, but Julia does not yet support that. In the future, however, it may do so, yielding even greater performance gains.)

Shared libraries and functions are referenced by a tuple of the form `(:function, "library")` or `("function", "library")` where `function` is the C-exported function name. `library` refers to the shared library name: shared libraries available in the (platform-specific) load path will be resolved by name, and if necessary a direct path may be specified.

A function name may be used alone in place of the tuple (just `:function` or `"function"`). In this case the name is resolved within the current process. This form can be used to call C library functions, functions in the Julia runtime, or functions in an application linked to Julia.

By default, Fortran compilers [generate mangled names](#) (for example, converting function names to lowercase or uppercase, often appending an underscore), and so to call a Fortran function via `ccall()` (page 607) you must pass the mangled identifier corresponding to the rule followed by your Fortran compiler. Also, when calling a Fortran function, all inputs must be passed by reference.

Finally, you can use `ccall()` (page 607) to actually generate a call to the library function. Arguments to `ccall()` (page 607) are as follows:

1. `(:function, "library")` pair (must be a constant, but see below).
2. Return type (see below for mapping the declared C type to Julia)
 - This argument will be evaluated at compile-time.
3. A tuple of input types. The input types must be written as a literal tuple, not a tuple-valued variable or expression.
 - This argument will be evaluated at compile-time.
4. The following arguments, if any, are the actual argument values passed to the function.

As a complete but simple example, the following calls the `clock` function from the standard C library:

```
julia> t = ccall( (:clock, "libc"), Int32, ())
2292761

julia> t
2292761

julia> typeof(ans)
Int32
```

`clock` takes no arguments and returns an `Int32`. One common gotcha is that a 1-tuple must be written with a trailing comma. For example, to call the `getenv` function to get a pointer to the value of an environment variable, one makes a call like this:

```
julia> path = ccall( (:getenv, "libc"), Cstring, (Cstring,), "SHELL")
Cstring(@0x00007fff5fbffc45)

julia> unsafe_string(path)
"/bin/bash"
```

Note that the argument type tuple must be written as `(Cstring,)`, rather than `(Cstring)`. This is because `(Cstring)` is just the expression `Cstring` surrounded by parentheses, rather than a 1-tuple containing `Cstring`:

```
julia> (Cstring)
Cstring

julia> (Cstring,)
(Cstring,)
```

In practice, especially when providing reusable functionality, one generally wraps `ccall()` (page 607) uses in Julia functions that set up arguments and then check for errors in whatever manner the C or Fortran function indicates them, propagating to the Julia caller as exceptions. This is especially important since C and Fortran APIs are notoriously inconsistent about how they indicate error conditions. For example, the `getenv` C library function is wrapped in the following Julia function, which is a simplified version of the actual definition from `env.jl`:

```
function getenv(var::AbstractString)
    val = ccall( (:getenv, "libc"),
                  Cstring, (Cstring,), var)
    if val == C_NULL
        error("getenv: undefined variable: ", var)
    end
    unsafe_string(val)
end
```

The C `getenv` function indicates an error by returning `NULL`, but other standard C functions indicate errors in various different ways, including by returning `-1`, `0`, `1` and other special values. This wrapper throws an exception clearly indicating the problem if the caller tries to get a non-existent environment variable:

```
julia> getenv("SHELL")
"/bin/bash"

julia> getenv("FOOBAR")
getenv: undefined variable: FOOBAR
```

Here is a slightly more complex example that discovers the local machine's hostname:

```
function gethostname()
    hostname = Array{UInt8}(128)
    ccall( (:gethostname, "libc"), Int32,
           (Ptr{UInt8}, Csize_t),
```

```

    hostname, sizeof(hostname))
hostname[end] = 0; # ensure null-termination
return unsafe_string(pointer(hostname))
end

```

This example first allocates an array of bytes, then calls the C library function `gethostname` to fill the array in with the hostname, takes a pointer to the hostname buffer, and converts the pointer to a Julia string, assuming that it is a NUL-terminated C string. It is common for C libraries to use this pattern of requiring the caller to allocate memory to be passed to the callee and filled in. Allocation of memory from Julia like this is generally accomplished by creating an uninitialized array and passing a pointer to its data to the C function. This is why we don't use the `Cstring` type here: as the array is uninitialized, it could contain NUL bytes. Converting to a `Cstring` as part of the `ccall()` (page 607) checks for contained NUL bytes and could therefore throw a conversion error.

25.1 Creating C-Compatible Julia Function Pointers

It is possible to pass Julia functions to native C functions that accept function pointer arguments. For example, to match C prototypes of the form:

```
typedef returntype (*functiontype)(argumenttype,...)
```

The function `cfunction()` (page 607) generates the C-compatible function pointer for a call to a Julia library function. Arguments to `cfunction()` (page 607) are as follows:

1. A Julia Function
2. Return type
3. A tuple of input types

A classic example is the standard C library `qsort` function, declared as:

```
void qsort(void *base, size_t nmemb, size_t size,
          int(*compare)(const void *a, const void *b));
```

The `base` argument is a pointer to an array of length `nmemb`, with elements of `size` bytes each. `compare` is a callback function which takes pointers to two elements `a` and `b` and returns an integer less/greater than zero if `a` should appear before/after `b` (or zero if any order is permitted). Now, suppose that we have a 1d array `A` of values in Julia that we want to sort using the `qsort` function (rather than Julia's built-in `sort` function). Before we worry about calling `qsort` and passing arguments, we need to write a comparison function that works for some arbitrary type `T`:

```

function mycompare{T}(a::T, b::T)
    return convert{Cint, a < b ? -1 : a > b ? +1 : 0}::Cint
end

```

Notice that we have to be careful about the return type: `qsort` expects a function returning a C `int`, so we must be sure to return `Cint` via a call to `convert` and a `typeassert`.

In order to pass this function to C, we obtain its address using the function `cfunction`:

```
const mycompare_c = cfunction(mycompare, Cint, (Ref{Cdouble}, Ref{Cdouble}))
```

`cfunction()` (page 607) accepts three arguments: the Julia function (`mycompare`), the return type (`Cint`), and a tuple of the argument types, in this case to sort an array of `Cdouble` (`Float64`) elements.

The final call to `qsort` looks like this:

```
A = [1.3, -2.7, 4.4, 3.1]
ccall(:qsort, Void, (Ptr{Cdouble}, Csize_t, Csize_t, Ptr{Void}),
      A, length(A), sizeof(elttype(A)), mycompare_c)
```

After this executes, `A` is changed to the sorted array `[-2.7, 1.3, 3.1, 4.4]`. Note that Julia knows how to convert an array into a `Ptr{Cdouble}`, how to compute the size of a type in bytes (identical to C's `sizeof` operator), and so on. For fun, try inserting a `println("mycompare($a, $b)")` line into `mycompare`, which will allow you to see the comparisons that `qsort` is performing (and to verify that it is really calling the Julia function that you passed to it).

25.2 Mapping C Types to Julia

It is critical to exactly match the declared C type with its declaration in Julia. Inconsistencies can cause code that works correctly on one system to fail or produce indeterminate results on a different system.

Note that no C header files are used anywhere in the process of calling C functions: you are responsible for making sure that your Julia types and call signatures accurately reflect those in the C header file. (The [Clang package](#) can be used to auto-generate Julia code from a C header file.)

25.2.1 Auto-conversion:

Julia automatically inserts calls to the `cconvert()` (page 607) function to convert each argument to the specified type. For example, the following call:

```
ccall((:foo, "libfoo"), Void, (Int32, Float64), x, y)
```

will behave as if the following were written:

```
ccall((:foo, "libfoo"), Void, (Int32, Float64),
      Base.unsafe_convert{Int32, Base.cconvert{Int32, x}},
      Base.unsafe_convert{Float64, Base.cconvert{Float64, y}})
```

`cconvert()` (page 607) normally just calls `convert()` (page 364), but can be defined to return an arbitrary new object more appropriate for passing to C. For example, this is used to convert an `Array` of objects (e.g. strings) to an array of pointers.

`unsafe_convert()` (page 607) handles conversion to `Ptr` types. It is considered unsafe because converting an object to a native pointer can hide the object from the garbage collector, causing it to be freed prematurely.

25.2.2 Type Correspondences:

First, a review of some relevant Julia type terminology:

| Syntax / Keyword | Example | Description |
|------------------------------------|---|--|
| <code>type</code> | <code>String</code> | “Leaf Type” :: A group of related data that includes a type-tag, is managed by the Julia GC, and is defined by object-identity. The type parameters of a leaf type must be fully defined (no <code>TypeVars</code> are allowed) in order for the instance to be constructed. |
| <code>abstract</code> | <code>Any, AbstractArray{T,N}, Complex{T}</code> | “Super Type” :: A super-type (not a leaf-type) that cannot be instantiated, but can be used to describe a group of types. |
| <code>{T}</code> | <code>Vector{Int}</code> | “Type Parameter” :: A specialization of a type (typically used for dispatch or storage optimization).
“TypeVar” :: The <code>T</code> in the type parameter declaration is referred to as a <code>TypeVar</code> (short for type variable). |
| <code>bitstype</code> | <code>Int, Float64</code> | “Bits Type” :: A type with no fields, but a size. It is stored and defined by-value. |
| <code>immutable</code> | <code>Pair{Int,Int}</code>
<code>Complex128(isbits)</code> | “Immutable” :: A type with all fields defined to be constant. It is defined by-value. And may be stored with a type-tag.
“Is-Bits” :: A <code>bitstype</code> , or an <code>immutable</code> type where all fields are other <code>isbits</code> types. It is defined by-value, and is stored without a type-tag. |
| <code>type ...; end</code> | <code>nothing</code> | “Singleton” :: a Leaf Type or Immutable with no fields. |
| <code>(...) or tuple(...) `</code> | <code>(1,2,3)</code> | “Tuple” :: an immutable data-structure similar to an anonymous immutable type, or a constant array. Represented as either an array or a struct. |
| <code>typealias</code> | Not applicable here | Type aliases, and other similar mechanisms of doing type indirection, are resolved to their base type (this includes assigning a type to another name, or getting the type out of a function call). |

25.2.3 Bits Types:

There are several special types to be aware of, as no other type can be defined to behave the same:

Float32 Exactly corresponds to the `float` type in C (or `REAL*4` in Fortran).

Float64 Exactly corresponds to the `double` type in C (or `REAL*8` in Fortran).

Complex64 Exactly corresponds to the `complex float` type in C (or `COMPLEX*8` in Fortran).

Complex128 Exactly corresponds to the `complex double` type in C (or `COMPLEX*16` in Fortran).

Signed Exactly corresponds to the `signed` type annotation in C (or any `INTEGER` type in Fortran). Any Julia type that is not a subtype of `Signed` is assumed to be unsigned.

Ref{T} Behaves like a `Ptr{T}` that owns its memory.

Array{T,N} When an array is passed to C as a `Ptr{T}` argument, it is not reinterpret-cast: Julia requires that the element type of the array matches `T`, and the address of the first element is passed.

Therefore, if an `Array` contains data in the wrong format, it will have to be explicitly converted using a call such as `trunc{Int32,a}`.

To pass an array `A` as a pointer of a different type *without* converting the data beforehand (for example, to pass a `Float64` array to a function that operates on uninterpreted bytes), you can declare the argument as `Ptr{Void}`.

If an array of eltype `Ptr{T}` is passed as a `Ptr{Ptr{T}}` argument, `Base.cconvert()` (page 607) will attempt to first make a null-terminated copy of the array with each element replaced by its `cconvert()` (page 607) version. This allows, for example, passing an `argv` pointer array of type `Vector{String}` to an argument of type `Ptr{Ptr{Cchar}}`.

On all systems we currently support, basic C/C++ value types may be translated to Julia types as follows. Every C type also has a corresponding Julia type with the same name, prefixed by `C`. This can help for writing portable code (and remembering that an `int` in C is not the same as an `Int` in Julia).

System Independent:

| C name | Fortran name | Standard Julia Alias | Julia Base Type |
|--|------------------------|----------------------|--|
| unsigned char
bool (C++) | CHARACTER | Cuchar | UInt8 |
| short | INTEGER*2
LOGICAL*2 | Cshort | Int16 |
| unsigned short | | Cushort | UInt16 |
| int
BOOL (C, typical) | INTEGER*4
LOGICAL*4 | Cint | Int32 |
| unsigned int | | Cuint | UInt32 |
| long long | INTEGER*8
LOGICAL*8 | Clonglong | Int64 |
| unsigned long long | | Culonglong | UInt64 |
| intmax_t | | Cintmax_t | Int64 |
| uintmax_t | | Cuintmax_t | UInt64 |
| float | REAL*4i | Cfloat | Float32 |
| double | REAL*8 | Cdouble | Float64 |
| complex float | COMPLEX*8 | Complex64 | Complex{Float32} |
| complex double | COMPLEX*16 | Complex128 | Complex{Float64} |
| ptrdiff_t | | Cptrdiff_t | Int |
| ssize_t | | Cssize_t | Int |
| size_t | | Csize_t | UInt |
| void | | | Void |
| void and
[[noreturn]] or
_Noreturn | | | Union{} |
| void* | | | Ptr{Void} |
| T* (where T represents an
appropriately defined type) | | | Ref{T} |
| char* (or char[], e.g. a
string) | CHARACTER*N | | Cstring if NUL-
terminated, or
Ptr{UInt8} if not |
| char** (or *char[]) | | | Ptr{Ptr{UInt8}} |
| j1_value_t* (any Julia
Type) | | | Any |
| j1_value_t** (a refer-
ence to a Julia Type) | | | Ref{Any} |
| va_arg | | | Not supported |
| ... (variadic function
specification) | | | T... (where T is one
of the above types, vari-
adic functions of different
argument types are not sup-
ported) |

The `Cstring` type is essentially a synonym for `Ptr{UInt8}`, except the conversion to `Cstring` throws an error if the Julia string contains any embedded NUL characters (which would cause the string to be silently truncated if the C routine treats NUL as the terminator). If you are passing a `char*` to a C routine that does not assume NUL termination (e.g. because you pass an explicit string length), or if you know for certain that your Julia string does not contain NUL and want to skip the check, you can use `Ptr{UInt8}` as the argument type. `Cstring` can also be used as the `ccall()` (page 607) return type, but in that case it obviously does not introduce any extra checks and is only meant to improve readability of the call.

System-dependent:

| C name | Standard Julia Alias | Julia Base Type |
|---------------|----------------------|--|
| char | Cchar | Int8 (x86, x86_64)
UInt8 (powerpc, arm) |
| long | Clong | Int (UNIX)
Int32 (Windows) |
| unsigned long | Culong | UInt (UNIX)
UInt32 (Windows) |
| wchar_t | Cwchar_t | Int32 (UNIX)
UInt16 (Windows) |

Note: When calling a Fortran function, all inputs must be passed by reference, so all type correspondences above should contain an additional `Ptr{...}` or `Ref{...}` wrapper around their type specification.

Warning: For string arguments (`char*`) the Julia type should be `Cstring` (if NUL-terminated data is expected) or either `Ptr{Cchar}` or `Ptr{UInt8}` otherwise (these two pointer types have the same effect), as described above, not `String`. Similarly, for array arguments (`T[]` or `T*`), the Julia type should again be `Ptr{T}`, not `Vector{T}`.

Warning: Julia's `Char` type is 32 bits, which is not the same as the wide character type (`wchar_t` or `wint_t`) on all platforms.

Warning: A return type of `Union{}` means the function will not return i.e. C++11 `[[noreturn]]` or C11 `_Noreturn` (e.g. `jl_throw` or `longjmp`). Do not use this for functions that return no value (`void`) but do return.

Note: For `wchar_t*` arguments, the Julia type should be `Cwstring` (if the C routine expects a NUL-terminated string) or `Ptr{Cwchar_t}` otherwise. Note also that UTF-8 string data in Julia is internally NUL-terminated, so it can be passed to C functions expecting NUL-terminated data without making a copy (but using the `Cwstring` type will cause an error to be thrown if the string itself contains NUL characters).

Note: C functions that take an argument of the type `char**` can be called by using a `Ptr{Ptr{UInt8}}` type within Julia. For example, C functions of the form:

```
int main(int argc, char **argv);
```

can be called via the following Julia code:

```
argv = [ "a.out", "arg1", "arg2" ]
ccall(:main, Int32, (Int32, Ptr{Ptr{UInt8}}), length(argv), argv)
```

Note: A C function declared to return `Void` will return the value `nothing` in Julia.

25.2.4 Struct Type correspondences

Composite types, aka `struct` in C or `TYPE` in Fortran90 (or `STRUCTURE` / `RECORD` in some variants of F77), can be mirrored in Julia by creating a `type` or `immutable` definition with the same field layout.

When used recursively, `isbits` types are stored inline. All other types are stored as a pointer to the data. When mirroring a struct used by-value inside another struct in C, it is imperative that you do not attempt to manually copy the fields over, as this will not preserve the correct field alignment. Instead, declare an immutable `isbits` type and use that instead. Unnamed structs are not possible in the translation to Julia.

Packed structs and union declarations are not supported by Julia.

You can get a near approximation of a `union` if you know, a priori, the field that will have the greatest size (potentially including padding). When translating your fields to Julia, declare the Julia field to be only of that type.

Arrays of parameters must be expanded manually, currently (either inline, or in an immutable helper type). For example:

```
in C:
struct B {
    int A[3];
};
b_a_2 = B.A[2];

in Julia:
immutable B_A
    A_1::Cint
    A_2::Cint
    A_3::Cint
end
type B
    A::B_A
end
b_a_2 = B.A.(2)
```

Arrays of unknown size are not supported.

In the future, some of these restrictions may be reduced or eliminated.

25.2.5 SIMD Values

Note: This feature is currently implemented on 64-bit x86 and AArch64 platforms only.

If a C/C++ routine has an argument or return value that is a native SIMD type, the corresponding Julia type is a homogeneous tuple of `VecElement` that naturally maps to the SIMD type. Specifically:

- The tuple must be the same size as the SIMD type. For example, a tuple representing an `__m128` on x86 must have a size of 16 bytes.
- The element type of the tuple must be an instance of `VecElement{T}` where `T` is a bitstype that is 1, 2, 4 or 8 bytes.

For instance, consider this C routine that uses AVX intrinsics:

```
#include <immintrin.h>

__m256 dist( __m256 a, __m256 b ) {
    return _mm256_sqrt_ps(_mm256_add_ps(_mm256_mul_ps(a, a),
                                         _mm256_mul_ps(b, b)));
}
```

The following Julia code calls `dist` using `ccall`:

```
 typealias m256 NTuple{8,VecElement{Float32}}
```

```
a = m256(ntuple(i->VecElement(sin(Float32(i))),8))
b = m256(ntuple(i->VecElement(cos(Float32(i))),8))

function call_dist(a::m256, b::m256)
    ccall(:dist, "libdist", m256, (m256, m256), a, b)
end

println(call_dist(a,b))
```

The host machine must have the requisite SIMD registers. For example, the code above will not work on hosts without AVX support.

25.2.6 Memory Ownership

malloc/free

Memory allocation and deallocation of such objects must be handled by calls to the appropriate cleanup routines in the libraries being used, just like in any C program. Do not try to free an object received from a C library with `Libc.free` in Julia, as this may result in the `free` function being called via the wrong `libc` library and cause Julia to crash. The reverse (passing an object allocated in Julia to be freed by an external library) is equally invalid.

25.2.7 When to use `T`, `Ptr{T}` and `Ref{T}`

In Julia code wrapping calls to external C routines, ordinary (non-pointer) data should be declared to be of type `T` inside the `ccall()` (page 607), as they are passed by value. For C code accepting pointers, `Ref{T}` should generally be used for the types of input arguments, allowing the use of pointers to memory managed by either Julia or C through the implicit call to `cconvert()` (page 607). In contrast, pointers returned by the C function called should be declared to be of output type `Ptr{T}`, reflecting that the memory pointed to is managed by C only. Pointers contained in C structs should be represented as fields of type `Ptr{T}` within the corresponding Julia immutable types designed to mimic the internal structure of corresponding C structs.

In Julia code wrapping calls to external Fortran routines, all input arguments should be declared as of type `Ref{T}`, as Fortran passes all variables by reference. The return type should either be `Void` for Fortran subroutines, or a `T` for Fortran functions returning the type `T`.

25.3 Mapping C Functions to Julia

25.3.1 `ccall/cfunction` argument translation guide

For translating a C argument list to Julia:

- `T`, where `T` is one of the primitive types: `char`, `int`, `long`, `short`, `float`, `double`, `complex`, `enum` or any of their `typedef` equivalents
 - `T`, where `T` is an equivalent Julia Bits Type (per the table above)
 - if `T` is an `enum`, the argument type should be equivalent to `Cint` or `Cuint`
 - argument value will be copied (passed by value)
- `struct T` (including `typedef` to a struct)
 - `T`, where `T` is a Julia leaf type
 - argument value will be copied (passed by value)

- `void*`
 - depends on how this parameter is used, first translate this to the intended pointer type, then determine the Julia equivalent using the remaining rules in this list
 - this argument may be declared as `Ptr{Void}`, if it really is just an unknown pointer
- `jl_value_t*`
 - `Any`
 - argument value must be a valid Julia object
 - currently unsupported by `cfunction()` (page 607)
- `jl_value_t**`
 - `Ref{Any}`
 - argument value must be a valid Julia object (or `C_NULL`)
 - currently unsupported by `cfunction()` (page 607)
- `T*`
 - `Ref{T}`, where `T` is the Julia type corresponding to `T`
 - argument value will be copied if it is an `isbits` type otherwise, the value must be a valid Julia object
- `(T*) (...)` (e.g. a pointer to a function)
 - `Ptr{Void}` (you may need to use `cfunction()` (page 607) explicitly to create this pointer)
- `...` (e.g. a vararg)
 - `T...`, where `T` is the Julia type
- `va_arg`
 - not supported

25.3.2 `ccall/cfunction` return type translation guide

For translating a C return type to Julia:

- `void`
 - `Void` (this will return the singleton instance `nothing::Void`)
- `T`, where `T` is one of the primitive types: `char`, `int`, `long`, `short`, `float`, `double`, `complex`, `enum` or any of their `typedef` equivalents
 - `T`, where `T` is an equivalent Julia Bits Type (per the table above)
 - if `T` is an `enum`, the argument type should be equivalent to `Cint` or `Cuint`
 - argument value will be copied (returned by-value)
- `struct T` (including `typedef` to a struct)
 - `T`, where `T` is a Julia Leaf Type
 - argument value will be copied (returned by-value)
- `void*`
 - depends on how this parameter is used, first translate this to the intended pointer type, then determine the Julia equivalent using the remaining rules in this list

- this argument may be declared as `Ptr{Void}`, if it really is just an unknown pointer
- `jl_value_t*`
 - `Any`
 - argument value must be a valid Julia object
- `jl_value_t**`
 - `Ref{Any}`
 - argument value must be a valid Julia object (or `C_NULL`)
- `T*`
 - If the memory is already owned by Julia, or is an `isbits` type, and is known to be non-null:
 - * `Ref{T}`, where `T` is the Julia type corresponding to `T`
 - * a return type of `Ref{Any}` is invalid, it should either be `Any` (corresponding to `jl_value_t*`) or `Ptr{Any}` (corresponding to `Ptr{Any}`)
 - * **C MUST NOT** modify the memory returned via `Ref{T}` if `T` is an `isbits` type
 - If the memory is owned by C:
 - * `Ptr{T}`, where `T` is the Julia type corresponding to `T`
- `(T*) (...)` (e.g. a pointer to a function)
 - `Ptr{Void}` (you may need to use `cfunction()` (page 607) explicitly to create this pointer)

25.3.3 Passing Pointers for Modifying Inputs

Because C doesn't support multiple return values, often C functions will take pointers to data that the function will modify. To accomplish this within a `ccall()` (page 607), you need to first encapsulate the value inside an `Ref{T}` of the appropriate type. When you pass this `Ref` object as an argument, Julia will automatically pass a C pointer to the encapsulated data:

```
width = Ref{Cint}(0)
range = Ref{Cfloat}(0)
ccall(:foo, Void, (Ref{Cint}, Ref{Cfloat}), width, range)
```

Upon return, the contents of `width` and `range` can be retrieved (if they were changed by `foo`) by `width[]` and `range[]`; that is, they act like zero-dimensional arrays.

25.3.4 Special Reference Syntax for `ccall` (deprecated):

The `&` syntax is deprecated, use the `Ref{T}` argument type instead.

A prefix `&` is used on an argument to `ccall()` (page 607) to indicate that a pointer to a scalar argument should be passed instead of the scalar value itself (required for all Fortran function arguments, as noted above). The following example computes a dot product using a BLAS function.

```
function compute_dot(DX::Vector{Float64}, DY::Vector{Float64})
    assert(length(DX) == length(DY))
    n = length(DX)
    incx = incy = 1
    product = ccall((:ddot_, "libLAPACK"),
                    Float64,
                    (Ptr{Int32}, Ptr{Float64}, Ptr{Int32}, Ptr{Float64}, Ptr{Int32}),
                    (Ptr{Int32}, Ptr{Float64}, Ptr{Int32}, Ptr{Float64}, Ptr{Int32}),
```

```

        &n, DX, &incx, DY, &incy)
    return product
end

```

The meaning of prefix `&` is not quite the same as in C. In particular, any changes to the referenced variables will not be visible in Julia unless the type is mutable (declared via `type`). However, even for immutable types it will not cause any harm for called functions to attempt such modifications (that is, writing through the passed pointers). Moreover, `&` may be used with any expression, such as `&0` or `&f(x)`.

When a scalar value is passed with `&` as an argument of type `Ptr{T}`, the value will first be converted to type `T`.

25.4 Some Examples of C Wrappers

Here is a simple example of a C wrapper that returns a `Ptr` type:

```

type gsl_permutation
end

# The corresponding C signature is
#   gsl_permutation * gsl_permutation_alloc (size_t n);
function permutation_alloc(n::Integer)
    output_ptr = ccall(
        (:gsl_permutation_alloc, :libgsl), #name of C function and library
        Ptr{gsl_permutation},             #output type
        (Csize_t,),                       #tuple of input types
        n                                 #name of Julia variable to pass in
    )
    if output_ptr==C_NULL # Could not allocate memory
        throw(OutOfMemoryError())
    end
    return output_ptr
end

```

The [GNU Scientific Library](#) (here assumed to be accessible through `:libgsl`) defines an opaque pointer, `gsl_permutation *`, as the return type of the C function `gsl_permutation_alloc()`. As user code never has to look inside the `gsl_permutation` struct, the corresponding Julia wrapper simply needs a new type declaration, `gsl_permutation`, that has no internal fields and whose sole purpose is to be placed in the type parameter of a `Ptr` type. The return type of the `ccall()` (page 607) is declared as `Ptr{gsl_permutation}`, since the memory allocated and pointed to by `output_ptr` is controlled by C (and not Julia).

The input `n` is passed by value, and so the function's input signature is simply declared as `(Csize_t,)` without any `Ref` or `Ptr` necessary. (If the wrapper was calling a Fortran function instead, the corresponding function input signature should instead be `(Ref{Csize_t},)`, since Fortran variables are passed by reference.) Furthermore, `n` can be any type that is convertible to a `Csize_t` integer; the `ccall()` (page 607) implicitly calls `Base.cconvert(Csize_t, n)` (page 607).

Here is a second example wrapping the corresponding destructor:

```

# The corresponding C signature is
#   void gsl_permutation_free (gsl_permutation * p);
function permutation_free(p::Ref{gsl_permutation})
    ccall(
        (:gsl_permutation_free, :libgsl), #name of C function and library
        Void,                             #output type
        (Ref{gsl_permutation},),          #tuple of input types
        p                                 #name of Julia variable to pass in
    )
end

```

```
)
end
```

Here, the input `p` is declared to be of type `Ref{gsl_permutation}`, meaning that the memory that `p` points to may be managed by Julia or by C. A pointer to memory allocated by C should be of type `Ptr{gsl_permutation}`, but it is convertible using `cconvert()` (page 607) and therefore can be used in the same (covariant) context of the input argument to a `ccall()` (page 607). A pointer to memory allocated by Julia must be of type `Ref{gsl_permutation}`, to ensure that the memory address pointed to is valid and that Julia's garbage collector manages the chunk of memory pointed to correctly. Therefore, the `Ref{gsl_permutation}` declaration allows pointers managed by C or Julia to be used.

If the C wrapper never expects the user to pass pointers to memory managed by Julia, then using `p::Ptr{gsl_permutation}` for the method signature of the wrapper and similarly in the `ccall()` (page 607) is also acceptable.

Here is a third example passing Julia arrays:

```
# The corresponding C signature is
#   int gsl_sf_bessel_Jn_array (int nmin, int nmax, double x,
#                               double result_array[])
#
function sf_bessel_Jn_array(nmin::Integer, nmax::Integer, x::Real)
    if nmax < nmin throw(DomainError()) end
    result_array = Array{Cdouble}(nmax-nmin+1)
    errorcode = ccall(
        (:gsl_sf_bessel_Jn_array, :libgsl), #name of C function and library
        Cint,                               #output type
        (Cint, Cint, Cdouble, Ref{Cdouble}), #tuple of input types
        nmin, nmax, x, result_array         #names of Julia variables to pass in
    )
    if errorcode != 0 error("GSL error code $errorcode") end
    return result_array
end
```

The C function wrapped returns an integer error code; the results of the actual evaluation of the Bessel J function populate the Julia array `result_array`. This variable can only be used with corresponding input type declaration `Ref{Cdouble}`, since its memory is allocated and managed by Julia, not C. The implicit call to `Base.cconvert(Ref{Cdouble}, result_array)` (page 607) unpacks the Julia pointer to a Julia array data structure into a form understandable by C.

Note that for this code to work correctly, `result_array` must be declared to be of type `Ref{Cdouble}` and not `Ptr{Cdouble}`. The memory is managed by Julia and the `Ref` signature alerts Julia's garbage collector to keep managing the memory for `result_array` while the `ccall()` (page 607) executes. If `Ptr{Cdouble}` were used instead, the `ccall()` (page 607) may still work, but Julia's garbage collector would not be aware that the memory declared for `result_array` is being used by the external C function. As a result, the code may produce a memory leak if `result_array` never gets freed by the garbage collector, or if the garbage collector prematurely frees `result_array`, the C function may end up throwing an invalid memory access exception.

25.5 Garbage Collection Safety

When passing data to a `ccall()` (page 607), it is best to avoid using the `pointer()` (page 608) function. Instead define a convert method and pass the variables directly to the `ccall()` (page 607). `ccall()` (page 607) automatically arranges that all of its arguments will be preserved from garbage collection until the call returns. If a C API will store a reference to memory allocated by Julia, after the `ccall()` (page 607) returns, you must arrange that the object remains visible to the garbage collector. The suggested way to handle this is to make a global variable of type `Array{Ref, 1}` to hold these values, until the C library notifies you that it is finished with them.

Whenever you have created a pointer to Julia data, you must ensure the original data exists until you are done with using the pointer. Many methods in Julia such as `unsafe_load()` (page 608) and `String()` (page 443) make copies of data instead of taking ownership of the buffer, so that it is safe to free (or alter) the original data without affecting Julia. A notable exception is `unsafe_wrap()` (page 444) which, for performance reasons, shares (or can be told to take ownership of) the underlying buffer.

The garbage collector does not guarantee any order of finalization. That is, if `a` contained a reference to `b` and both `a` and `b` are due for garbage collection, there is no guarantee that `b` would be finalized after `a`. If proper finalization of `a` depends on `b` being valid, it must be handled in other ways.

25.6 Non-constant Function Specifications

A `(name, library)` function specification must be a constant expression. However, it is possible to use computed values as function names by staging through `eval` as follows:

```
@eval ccall(($ (string("a", "b")), "lib"), ...
```

This expression constructs a name using `string`, then substitutes this name into a new `ccall()` (page 607) expression, which is then evaluated. Keep in mind that `eval` only operates at the top level, so within this expression local variables will not be available (unless their values are substituted with `$`). For this reason, `eval` is typically only used to form top-level definitions, for example when wrapping libraries that contain many similar functions.

If your usage is more dynamic, use indirect calls as described in the next section.

25.7 Indirect Calls

The first argument to `ccall()` (page 607) can also be an expression evaluated at run time. In this case, the expression must evaluate to a `Ptr`, which will be used as the address of the native function to call. This behavior occurs when the first `ccall()` (page 607) argument contains references to non-constants, such as local variables, function arguments, or non-constant globals.

For example, you might lookup the function via `dlsym`, then cache it in a global variable for that session. For example:

```
macro dlsym(func, lib)
    z, zlocal = gensym(string(func)), gensym()
    eval(current_module(), :(global $z = C_NULL))
    z = esc(z)
    quote
        let $zlocal::Ptr{Void} = $z::Ptr{Void}
            if $zlocal == C_NULL
                $zlocal = dlsym($(esc(lib))::Ptr{Void}, $(esc(func)))
                global $z = $zlocal
            end
            $zlocal
        end
    end
end

mylibvar = dlopen("mylib")
ccall(@dlsym("myfunc", mylibvar), Void, ())
```

25.8 Calling Convention

The second argument to `ccall()` (page 607) can optionally be a calling convention specifier (immediately preceding return type). Without any specifier, the platform-default C calling convention is used. Other supported conventions are: `stdcall`, `cdecl`, `fastcall`, and `thiscall`. For example (from `base/libc.jl`) we see the same `gethostname` `ccall()` (page 607) as above, but with the correct signature for Windows:

```
hn = Array{UInt8}(256)
err = ccall(:gethostname, stdcall, Int32, (Ptr{UInt8}, UInt32), hn, length(hn))
```

For more information, please see the [LLVM Language Reference](#).

25.9 Accessing Global Variables

Global variables exported by native libraries can be accessed by name using the `cglobal()` (page 607) function. The arguments to `cglobal()` (page 607) are a symbol specification identical to that used by `ccall()` (page 607), and a type describing the value stored in the variable:

```
julia> cglobal((:errno,:libc), Int32)
Ptr{Int32} @0x00007f418d0816b8
```

The result is a pointer giving the address of the value. The value can be manipulated through this pointer using `unsafe_load()` (page 608) and `unsafe_store!()` (page 608).

25.10 Accessing Data through a Pointer

The following methods are described as “unsafe” because a bad pointer or type declaration can cause Julia to terminate abruptly.

Given a `Ptr{T}`, the contents of type `T` can generally be copied from the referenced memory into a Julia object using `unsafe_load(ptr, [index])`. The index argument is optional (default is 1), and follows the Julia-convention of 1-based indexing. This function is intentionally similar to the behavior of `getindex()` (page 454) and `setindex!()` (page 461) (e.g. `[]` access syntax).

The return value will be a new object initialized to contain a copy of the contents of the referenced memory. The referenced memory can safely be freed or released.

If `T` is `Any`, then the memory is assumed to contain a reference to a Julia object (a `jl_value_t*`), the result will be a reference to this object, and the object will not be copied. You must be careful in this case to ensure that the object was always visible to the garbage collector (pointers do not count, but the new reference does) to ensure the memory is not prematurely freed. Note that if the object was not originally allocated by Julia, the new object will never be finalized by Julia’s garbage collector. If the `Ptr` itself is actually a `jl_value_t*`, it can be converted back to a Julia object reference by `unsafe_pointer_to_objref(ptr)` (page 608). (Julia values `v` can be converted to `jl_value_t*` pointers, as `Ptr{Void}`, by calling `pointer_from_objref(v)` (page 608).)

The reverse operation (writing data to a `Ptr{T}`), can be performed using `unsafe_store!(ptr, value, [index])` (page 608). Currently, this is only supported for bittypes or other pointer-free (`isbits`) immutable types.

Any operation that throws an error is probably currently unimplemented and should be posted as a bug so that it can be resolved.

If the pointer of interest is a plain-data array (bitstype or immutable), the function `unsafe_wrap(Array, ptr, dims, [own])` (page 444) may be more useful. The final parameter should be true if Julia should “take

ownership” of the underlying buffer and call `free(ptr)` when the returned `Array` object is finalized. If the `own` parameter is omitted or false, the caller must ensure the buffer remains in existence until all access is complete.

Arithmetic on the `Ptr` type in Julia (e.g. using `+`) does not behave the same as C’s pointer arithmetic. Adding an integer to a `Ptr` in Julia always moves the pointer by some number of *bytes*, not elements. This way, the address values obtained from pointer arithmetic do not depend on the element types of pointers.

25.11 Thread-safety

Some C libraries execute their callbacks from a different thread, and since Julia isn’t thread-safe you’ll need to take some extra precautions. In particular, you’ll need to set up a two-layered system: the C callback should only *schedule* (via Julia’s event loop) the execution of your “real” callback. To do this, create a `AsyncCondition` object and wait on it:

```
cond = Base.AsyncCondition()
wait(cond)
```

The callback you pass to C should only execute a `ccall()` (page 607) to `:uv_async_send`, passing `cb.handle` as the argument, taking care to avoid any allocations or other interactions with the Julia runtime.

Note that events may be coalesced, so multiple calls to `uv_async_send` may result in a single wakeup notification to the condition.

25.12 More About Callbacks

For more details on how to pass callbacks to C libraries, see this [blog post](#).

25.13 C++

Limited support for C++ is provided by the `Cxx`, `Clang`, and `Cxx` packages.

Handling Operating System Variation

When dealing with platform libraries, it is often necessary to provide special cases for various platforms. The variable `Sys.KERNEL` can be used to write these special cases. There are several functions intended to make this easier: `is_unix`, `is_linux`, `is_apple`, `is_bsd`, and `is_windows`. These may be used as follows:

```
if is_windows()
    some_complicated_thing(a)
end
```

Note that `is_linux` and `is_apple` are mutually exclusive subsets of `is_unix`. Additionally, there is a macro `@static` which makes it possible to use these functions to conditionally hide invalid code, as demonstrated in the following examples.

Simple blocks:

```
ccall( (@static is_windows() ? :_fopen : :fopen), ...)
```

Complex blocks:

```
@static if is_linux()
    some_complicated_thing(a)
else
    some_different_thing(a)
end
```

When chaining conditionals (including `if/elseif/end`), the `@static` must be repeated for each level (parentheses optional, but recommended for readability):

```
@static is_windows() ? :a : (@static is_apple() ? :b : :c)
```

Interacting With Julia

Julia comes with a full-featured interactive command-line REPL (read-eval-print loop) built into the `julia` executable. In addition to allowing quick and easy evaluation of Julia statements, it has a searchable history, tab-completion, many helpful keybindings, and dedicated help and shell modes. The REPL can be started by simply calling `julia` with no arguments or double-clicking on the executable:

```
$ julia

      _       _ _(_)_      | A fresh approach to technical computing
  (_)_      | (_)_(_)      | Documentation: http://docs.julialang.org
    _ _    _|_|_ _ _ _    | Type "help()" to list help topics
  | | | | | | | / _` | |   |
  | | | | | | | (_| | |   | Version 0.3.0-prerelease+2834 (2014-04-30 03:13 UTC)
 _/ | \_ ' _|_|_| \_ ' _| | Commit 64f437b (0 days old master)
|___/                          | x86_64-apple-darwin13.1.0

julia>
```

To exit the interactive session, type `^D` — the control key together with the `d` key on a blank line — or type `quit()` followed by the return or enter key. The REPL greets you with a banner and a `julia>` prompt.

27.1 The different prompt modes

27.1.1 The Julian mode

The REPL has four main modes of operation. The first and most common is the Julian prompt. It is the default mode of operation; each new line initially starts with `julia>`. It is here that you can enter Julia expressions. Hitting return or enter after a complete expression has been entered will evaluate the entry and show the result of the last expression.

```
julia> string(1 + 2)
"3"
```

There are a number useful features unique to interactive work. In addition to showing the result, the REPL also binds the result to the variable `ans`. A trailing semicolon on the line can be used as a flag to suppress showing the result.

```
julia> string(3 * 4);

julia> ans
"12"
```

In Julia mode, the REPL supports something called *prompt pasting*. This activates when pasting text that starts with `julia>` into the REPL. In that case, only expressions starting with `julia>` are parsed, others are removed.

This makes it possible to paste a chunk of code that has been copied from a REPL session without having to scrub away prompts and outputs. This feature is enabled by default but can be disabled or enabled at will with `Base.REPL.enable_promptpaste (::Bool)`. If it is enabled, you can try it out by pasting the code block above this paragraph straight into the REPL. This feature does not work on the standard Windows command prompt due to its limitation at detecting when a paste occurs.

27.1.2 Help mode

When the cursor is at the beginning of the line, the prompt can be changed to a help mode by typing `?`. Julia will attempt to print help or documentation for anything entered in help mode:

```
julia> ? # upon typing ?, the prompt changes (in place) to: help>

help> string
Base.string(xs...)

    Create a string from any values using the "print" function.
```

In addition to function names, complete function calls may be entered to see which method is called for the given argument(s). Macros, types and variables can also be queried:

```
help> string(1)
string(x::Union{Int16, Int128, Int8, Int32, Int64}) at string.jl:1553

help> @printf
Base.@printf([io::IOStream], "%Fmt", args...)

    Print arg(s) using C "printf()" style format specification
    string. Optionally, an IOStream may be passed as the first argument
    to redirect output.

help> AbstractString
DataType      : AbstractString
supertype: Any
subtypes  : Any[DirectIndexString, RepString, RevString{T<:AbstractString}, SubString{T<:AbstractString}, ...]
```

Help mode can be exited by pressing backspace at the beginning of the line.

27.1.3 Shell mode

Just as help mode is useful for quick access to documentation, another common task is to use the system shell to execute system commands. Just as `?` entered help mode when at the beginning of the line, a semicolon (`;`) will enter the shell mode. And it can be exited by pressing backspace at the beginning of the line.

```
julia> ; # upon typing ;, the prompt changes (in place) to: shell>

shell> echo hello
hello
```

27.1.4 Search modes

In all of the above modes, the executed lines get saved to a history file, which can be searched. To initiate an incremental search through the previous history, type `^R` — the control key together with the `r` key. The prompt will change to (reverse-i-search) `''`, and as you type the search query will appear in the quotes. The most recent result

that matches the query will dynamically update to the right of the colon as more is typed. To find an older result using the same query, simply type `^R` again.

Just as `^R` is a reverse search, `^S` is a forward search, with the prompt (i-search) `''`:. The two may be used in conjunction with each other to move through the previous or next matching results, respectively.

27.2 Key bindings

The Julia REPL makes great use of key bindings. Several control-key bindings were already introduced above (`^D` to exit, `^R` and `^S` for searching), but there are many more. In addition to the control-key, there are also meta-key bindings. These vary more by platform, but most terminals default to using alt- or option- held down with a key to send the meta-key (or can be configured to do so).

| Program control | |
|-----------------------------------|--|
| <code>^D</code> | Exit (when buffer is empty) |
| <code>^C</code> | Interrupt or cancel |
| <code>^L</code> | Clear console screen |
| Return/Enter, <code>^J</code> | New line, executing if it is complete |
| meta-Return/Enter | Insert new line without executing it |
| <code>?</code> or <code>;</code> | Enter help or shell mode (when at start of a line) |
| <code>^R</code> , <code>^S</code> | Incremental history search, described above |
| Cursor movement | |
| Right arrow, <code>^F</code> | Move right one character |
| Left arrow, <code>^B</code> | Move left one character |
| Home, <code>^A</code> | Move to beginning of line |
| End, <code>^E</code> | Move to end of line |
| <code>^P</code> | Change to the previous or next history entry |
| <code>^N</code> | Change to the next history entry |
| Up arrow | Move up one line (or to the previous history entry) |
| Down arrow | Move down one line (or to the next history entry) |
| Page-up | Change to the previous history entry that matches the text before the cursor |
| Page-down | Change to the next history entry that matches the text before the cursor |
| meta-F | Move right one word |
| meta-B | Move left one word |
| Editing | |
| Backspace, <code>^H</code> | Delete the previous character |
| Delete, <code>^D</code> | Forward delete one character (when buffer has text) |
| meta-Backspace | Delete the previous word |
| meta-D | Forward delete the next word |
| <code>^W</code> | Delete previous text up to the nearest whitespace |
| <code>^K</code> | “Kill” to end of line, placing the text in a buffer |
| <code>^Y</code> | “Yank” insert the text from the kill buffer |
| <code>^T</code> | Transpose the characters about the cursor |

27.2.1 Customizing keybindings

Julia’s REPL keybindings may be fully customized to a user’s preferences by passing a dictionary to `REPL.setup_interface()`. The keys of this dictionary may be characters or strings. The key `'*'` refers to the default action. Control plus character `x` bindings are indicated with `"^x"`. Meta plus `x` can be written `"\Mx"`. The values of the custom keymap must be `nothing` (indicating that the input should be ignored) or functions that accept the signature `(PromptState, AbstractREPL, Char)`. The `REPL.setup_interface()` function

must be called before the REPL is initialized, by registering the operation with `atreplinit()`. For example, to bind the up and down arrow keys to move through history without prefix search, one could put the following code in `.juliarc.jl`:

```
import Base: LineEdit, REPL

const mykeys = Dict{Any, Any}()
# Up Arrow
"\e[A" => (s,o...) -> (LineEdit.edit_move_up(s) || LineEdit.history_prev(s, LineEdit.mode(s).hist)),
# Down Arrow
"\e[B" => (s,o...) -> (LineEdit.edit_move_up(s) || LineEdit.history_next(s, LineEdit.mode(s).hist))
)

function customize_keys(repl)
    repl.interface = REPL.setup_interface(repl; extra_repl_keymap = mykeys)
end

atreplinit(customize_keys)
```

Users should refer to `base/LineEdit.jl` to discover the available actions on key input.

27.3 Tab completion

In both the Julian and help modes of the REPL, one can enter the first few characters of a function or type and then press the tab key to get a list all matches:

```
julia> stri
stride      strides      string      stringmime  strip

julia> Stri
StridedArray  StridedVecOrMat  AbstractString
StridedMatrix  StridedVector
```

The tab key can also be used to substitute LaTeX math symbols with their Unicode equivalents, and get a list of LaTeX matches as well:

```
julia> \pi[TAB]
julia> π
π = 3.1415926535897...

julia> e\_1[TAB] = [1,0]
julia> e1 = [1,0]
2-element Array{Int64,1}:
 1
 0

julia> e^1[TAB] = [1 0]
julia> e¹ = [1 0]
1×2 Array{Int64,2}:
 1  0

julia> \sqrt[TAB]2      # √ is equivalent to the sqrt() function
julia> √2
1.4142135623730951

julia> \hbar[TAB](h) = h / 2\pi[TAB]
julia> ħ(h) = h / 2π
```

```

h (generic function with 1 method)

julia> \h[TAB]
\hat                \heartsuit          \hksearrow          \hookleftarrow      \hslash
\hbar               \hermitconjmatrix    \hkswarrow          \hookrightarrow     \hspace

julia> α="\alpha[TAB]"    # LaTeX completion also works in strings
julia> α="α"

```

A full list of tab-completions can be found in the *Unicode Input* (page 355) section of the manual.

Completion of paths works for strings and julia's shell mode:

```

julia> path="/[TAB]"
.dockerenv  .juliabox/  boot/      etc/      lib/      media/    opt/      root/
.dockerinit bin/      dev/      home/     lib64/    mnt/      proc/     run/
shell> /[TAB]
.dockerenv  .juliabox/  boot/      etc/      lib/      media/    opt/      root/
.dockerinit bin/      dev/      home/     lib64/    mnt/      proc/     run/

```

Tab completion can help with investigation of the available methods matching the input arguments:

```

julia> max([TAB] # All methods are displayed, not shown here due to size of the list

julia> max([1,2],[TAB] # All methods where `Vector{Int}` matches as first argument
max{T1<:Real,T2<:Real}(x::AbstractArray{T1,N<:Any}, y::T2) at operators.jl:544
max{Tx<:Real,Ty<:Real}(x::Union{Base.ReshapedArray{Tx,1,A<:DenseArray,MI<:Tuple{Vararg{Base.Multipli
max{T1<:Real,T2<:Real}(x::AbstractArray{T1,N<:Any}, y::AbstractArray{T2,N<:Any}) at operators.jl:548
max(x, y) at operators.jl:78
max(a, b, c, xs...) at operators.jl:119

julia> max([1,2], max(1,2),[TAB] # All methods matching the arguments.
max{T1<:Real,T2<:Real}(x::AbstractArray{T1,N<:Any}, y::T2) at operators.jl:544
max(x, y) at operators.jl:78
max(a, b, c, xs...) at operators.jl:119

julia> split("1 1 1", # Keywords are also displayed in the suggested methods, see second line after
split(str::AbstractString) at strings/util.jl:151
split{T<:AbstractString}(str::T, splitter; limit, keep) at strings/util.jl:127

```

The completion of the methods uses type inference and can therefore see if the arguments match even if the arguments are output from functions. The function needs to be type stable for the completion to be able to remove non-matching methods.

Tab completion can also help completing fields:

```

julia> Pkg.a
add          available

```

Fields for output from functions can also be completed:

```

julia> split("", "")[1].[TAB]
endof  offset  string

```

The completion of fields for output from functions uses type inference, and it can only suggest fields if the function is type stable.

27.4 Customizing Colors

The colors used by Julia and the REPL can be customized, as well. To change the color of the Julia prompt you can add something like the following to your `juliarc.jl` file:

```
function customize_colors(repl)
    repl.prompt_color = Base.text_colors[:cyan]
end

atreplinit(customize_colors)
```

The available color keys in `Base.text_colors` are `:black`, `:red`, `:green`, `:yellow`, `:blue`, `:magenta`, `:cyan`, `:white`, `:normal`, and `:bold` as well as the integers 0 to 255 for terminals with 256 color support. Similarly, you can change the colors for the help and shell prompts and input and answer text by setting the appropriate field of `repl` in the `customize_colors` function above (respectively, `help_color`, `shell_color`, `input_color`, and `answer_color`). For the latter two, be sure that the `envcolors` field is also set to `false`.

You can also customize the color used to render warning and informational messages by setting the appropriate environment variable. For instance, to render warning messages in yellow and informational messages in cyan you can add the following to your `juliarc.jl` file:

```
ENV["JULIA_WARN_COLOR"] = :yellow
ENV["JULIA_INFO_COLOR"] = :cyan
```

Embedding Julia

As we have seen in *Calling C and Fortran Code* (page 237), Julia has a simple and efficient way to call functions written in C. But there are situations where the opposite is needed: calling Julia function from C code. This can be used to integrate Julia code into a larger C/C++ project, without the need to rewrite everything in C/C++. Julia has a C API to make this possible. As almost all programming languages have some way to call C functions, the Julia C API can also be used to build further language bridges (e.g. calling Julia from Python or C#).

28.1 High-Level Embedding

We start with a simple C program that initializes Julia and calls some Julia code:

```
#include <julia.h>

int main(int argc, char *argv[])
{
    /* required: setup the Julia context */
    jl_init(NULL);

    /* run Julia commands */
    jl_eval_string("print(sqrt(2.0))");

    /* strongly recommended: notify Julia that the
       program is about to terminate. this allows
       Julia time to cleanup pending write requests
       and run all finalizers */
    /*
    jl_atexit_hook(0);
    return 0;
    */
}
```

In order to build this program you have to put the path to the Julia header into the include path and link against `libjulia`. For instance, when Julia is installed to `$JULIA_DIR`, one can compile the above test program `test.c` with `gcc` using:

```
gcc -o test -fPIC -I$JULIA_DIR/include/julia -L$JULIA_DIR/lib test.c -ljulia $JULIA_DIR/lib/julia/lib
```

Then if the environment variable `JULIA_HOME` is set to `$JULIA_DIR/bin`, the output test program can be executed.

Alternatively, look at the `embedding.c` program in the Julia source tree in the `examples/` folder. The file `ui/repl.c` program is another simple example of how to set `jl_options` options while linking against `libjulia`.

The first thing that has to be done before calling any other Julia C function is to initialize Julia. This is done by calling `jl_init`, which takes as argument a C string (`const char*`) to the location where Julia is installed. When the argument is `NULL`, Julia tries to determine the install location automatically.

The second statement in the test program evaluates a Julia statement using a call to `jl_eval_string`.

Before the program terminates, it is strongly recommended to call `jl_atexit_hook`. The above example program calls this before returning from `main`.

Note: Currently, dynamically linking with the `libjulia` shared library requires passing the `RTLD_GLOBAL` option. In Python, this looks like:

```
>>> julia=CDLL('./libjulia.dylib',RTLD_GLOBAL)
>>> julia.jl_init.argtypes = [c_char_p]
>>> julia.jl_init('.')
250593296
```

Note: If the `julia` program needs to access symbols from the main executable, it may be necessary to add `-Wl,--export-dynamic` linker flag at compile time on Linux in addition to the ones generated by `julia-config.jl` described below. This is not necessary when compiling a shared library.

28.1.1 Using `julia-config` to automatically determine build parameters

The script `julia-config.jl` was created to aid in determining what build parameters are required by a program that uses embedded Julia. This script uses the build parameters and system configuration of the particular Julia distribution it is invoked by to export the necessary compiler flags for an embedding program to interact with that distribution. This script is located in the Julia shared data directory.

Example

Below is essentially the same as above with one small change; the argument to `jl_init` is now `JULIA_INIT_DIR` which is defined by `julia-config.jl`:

```
#include <julia.h>

int main(int argc, char *argv[])
{
    jl_init(JULIA_INIT_DIR);
    (void) jl_eval_string("println(sqrt(2.0))");
    jl_atexit_hook(0);
    return 0;
}
```

On the command line

A simple use of this script is from the command line. Assuming that `julia-config.jl` is located in `/usr/local/julia/share/julia`, it can be invoked on the command line directly and takes any combination of 3 flags:

```
/usr/local/julia/share/julia/julia-config.jl
Usage: julia-config [--cflags|--ldflags|--ldlibs]
```

If the above example source is saved in the file *embed_example.c*, then the following command will compile it into a running program on Linux and Windows (MSYS2 environment), or if on OS/X, then substitute `clang` for `gcc`:

```
/usr/local/julia/share/julia/julia-config.jl --cflags --ldflags --ldlibs | xargs gcc embed_example.c
```

Use in Makefiles

But in general, embedding projects will be more complicated than the above, and so the following allows general makefile support as well – assuming GNU make because of the use of the **shell** macro expansions. Additionally, though many times *julia-config.jl* may be found in the directory */usr/local*, this is not necessarily the case, but Julia can be used to locate *julia-config.jl* too, and the makefile can be used to take advantage of that. The above example is extended to use a Makefile:

```
JL_SHARE = $(shell julia -e 'print(joinpath(JULIA_HOME,Base.DATAROOTDIR,"julia"))')
CFLAGS   += $(shell $(JL_SHARE)/julia-config.jl --cflags)
CXXFLAGS += $(shell $(JL_SHARE)/julia-config.jl --cflags)
LDFLAGS  += $(shell $(JL_SHARE)/julia-config.jl --ldflags)
LDLIBS   += $(shell $(JL_SHARE)/julia-config.jl --ldlibs)

all: embed_example
```

Now the build command is simply **make**.

28.2 Converting Types

Real applications will not just need to execute expressions, but also return their values to the host program. `jl_eval_string` returns a `jl_value_t*`, which is a pointer to a heap-allocated Julia object. Storing simple data types like `Float64` in this way is called **boxing**, and extracting the stored primitive data is called **unboxing**. Our improved sample program that calculates the square root of 2 in Julia and reads back the result in C looks as follows:

```
jl_value_t *ret = jl_eval_string("sqrt(2.0)");

if (jl_is_float64(ret)) {
    double ret_unboxed = jl_unbox_float64(ret);
    printf("sqrt(2.0) in C: %e \n", ret_unboxed);
}
```

In order to check whether `ret` is of a specific Julia type, we can use the `jl_is_...` functions. By typing `typeof(sqrt(2.0))` into the Julia shell we can see that the return type is `Float64` (double in C). To convert the boxed Julia value into a C double the `jl_unbox_float64` function is used in the above code snippet.

Corresponding `jl_box_...` functions are used to convert the other way:

```
jl_value_t *a = jl_box_float64(3.0);
jl_value_t *b = jl_box_float32(3.0f);
jl_value_t *c = jl_box_int32(3);
```

As we will see next, boxing is required to call Julia functions with specific arguments.

28.3 Calling Julia Functions

While `jl_eval_string` allows C to obtain the result of a Julia expression, it does not allow passing arguments computed in C to Julia. For this you will need to invoke Julia functions directly, using `jl_call`:

```
jl_function_t *func = jl_get_function(jl_base_module, "sqrt");
jl_value_t *argument = jl_box_float64(2.0);
jl_value_t *ret = jl_call1(func, argument);
```

In the first step, a handle to the Julia function `sqrt` is retrieved by calling `jl_get_function`. The first argument passed to `jl_get_function` is a pointer to the Base module in which `sqrt` is defined. Then, the double value is boxed using `jl_box_float64`. Finally, in the last step, the function is called using `jl_call1`. `jl_call0`, `jl_call2`, and `jl_call3` functions also exist, to conveniently handle different numbers of arguments. To pass more arguments, use `jl_call`:

```
jl_value_t *jl_call(jl_function_t *f, jl_value_t **args, int32_t nargs)
```

Its second argument `args` is an array of `jl_value_t*` arguments and `nargs` is the number of arguments.

28.4 Memory Management

As we have seen, Julia objects are represented in C as pointers. This raises the question of who is responsible for freeing these objects.

Typically, Julia objects are freed by a garbage collector (GC), but the GC does not automatically know that we are holding a reference to a Julia value from C. This means the GC can free objects out from under you, rendering pointers invalid.

The GC can only run when Julia objects are allocated. Calls like `jl_box_float64` perform allocation, and allocation might also happen at any point in running Julia code. However, it is generally safe to use pointers in between `jl_...` calls. But in order to make sure that values can survive `jl_...` calls, we have to tell Julia that we hold a reference to a Julia value. This can be done using the `JL_GC_PUSH` macros:

```
jl_value_t *ret = jl_eval_string("sqrt(2.0)");
JL_GC_PUSH1(&ret);
// Do something with ret
JL_GC_POP();
```

The `JL_GC_POP` call releases the references established by the previous `JL_GC_PUSH`. Note that `JL_GC_PUSH` is working on the stack, so it must be exactly paired with a `JL_GC_POP` before the stack frame is destroyed.

Several Julia values can be pushed at once using the `JL_GC_PUSH2`, `JL_GC_PUSH3`, and `JL_GC_PUSH4` macros. To push an array of Julia values one can use the `JL_GC_PUSHARGS` macro, which can be used as follows:

```
jl_value_t **args;
JL_GC_PUSHARGS(args, 2); // args can now hold 2 `jl_value_t*` objects
args[0] = some_value;
args[1] = some_other_value;
// Do something with args (e.g. call jl_... functions)
JL_GC_POP();
```

The garbage collector also operates under the assumption that it is aware of every old-generation object pointing to a young-generation one. Any time a pointer is updated breaking that assumption, it must be signaled to the collector with the `jl_gc_wb` (write barrier) function like so:

```
jl_value_t *parent = some_old_value, *child = some_young_value;
((some_specific_type*)parent)->field = child;
jl_gc_wb(parent, child);
```

It is in general impossible to predict which values will be old at runtime, so the write barrier must be inserted after all explicit stores. One notable exception is if the `parent` object was just allocated and garbage collection was not run since then. Remember that most `jl_...` functions can sometimes invoke garbage collection.

The write barrier is also necessary for arrays of pointers when updating their data directly. For example:

```
jl_array_t *some_array = ...; // e.g. a Vector{Any}
void **data = (void**) jl_array_data(some_array);
jl_value_t *some_value = ...;
data[0] = some_value;
jl_gc_wb(some_array, some_value);
```

28.4.1 Manipulating the Garbage Collector

There are some functions to control the GC. In normal use cases, these should not be necessary.

| | |
|---------------------------------|--|
| <code>jl_gc_collect()</code> | Force a GC run |
| <code>jl_gc_enable(0)</code> | Disable the GC, return previous state as int |
| <code>jl_gc_enable(1)</code> | Enable the GC, return previous state as int |
| <code>jl_gc_is_enabled()</code> | Return current state as int |

28.5 Working with Arrays

Julia and C can share array data without copying. The next example will show how this works.

Julia arrays are represented in C by the datatype `jl_array_t*`. Basically, `jl_array_t` is a struct that contains:

- Information about the datatype
- A pointer to the data block
- Information about the sizes of the array

To keep things simple, we start with a 1D array. Creating an array containing Float64 elements of length 10 is done by:

```
jl_value_t* array_type = jl_apply_array_type(jl_float64_type, 1);
jl_array_t* x = jl_alloc_array_1d(array_type, 10);
```

Alternatively, if you have already allocated the array you can generate a thin wrapper around its data:

```
double *existingArray = (double*) malloc(sizeof(double) * 10);
jl_array_t *x = jl_ptr_to_array_1d(array_type, existingArray, 10, 0);
```

The last argument is a boolean indicating whether Julia should take ownership of the data. If this argument is non-zero, the GC will call `free` on the data pointer when the array is no longer referenced.

In order to access the data of `x`, we can use `jl_array_data`:

```
double *xData = (double*) jl_array_data(x);
```

Now we can fill the array:

```
for(size_t i=0; i<jl_array_len(x); i++)
    xData[i] = i;
```

Now let us call a Julia function that performs an in-place operation on `x`:

```
jl_function_t *func = jl_get_function(jl_base_module, "reverse!");
jl_call1(func, (jl_value_t*)x);
```

By printing the array, one can verify that the elements of `x` are now reversed.

28.5.1 Accessing Returned Arrays

If a Julia function returns an array, the return value of `jl_eval_string` and `jl_call` can be cast to a `jl_array_t`:

```
jl_function_t *func = jl_get_function(jl_base_module, "reverse");
jl_array_t *y = (jl_array_t*)jl_call1(func, (jl_value_t*)x);
```

Now the content of `y` can be accessed as before using `jl_array_data`. As always, be sure to keep a reference to the array while it is in use.

28.5.2 Multidimensional Arrays

Julia's multidimensional arrays are stored in memory in column-major order. Here is some code that creates a 2D array and accesses its properties:

```
// Create 2D array of float64 type
jl_value_t *array_type = jl_apply_array_type(jl_float64_type, 2);
jl_array_t *x = jl_alloc_array_2d(array_type, 10, 5);

// Get array pointer
double *p = (double*)jl_array_data(x);
// Get number of dimensions
int ndims = jl_array_ndims(x);
// Get the size of the i-th dim
size_t size0 = jl_array_dim(x, 0);
size_t size1 = jl_array_dim(x, 1);

// Fill array with data
for(size_t i=0; i<size1; i++)
    for(size_t j=0; j<size0; j++)
        p[j + size0*i] = i + j;
```

Notice that while Julia arrays use 1-based indexing, the C API uses 0-based indexing (for example in calling `jl_array_dim`) in order to read as idiomatic C code.

28.6 Exceptions

Julia code can throw exceptions. For example, consider:

```
jl_eval_string("this_function_does_not_exist()");
```

This call will appear to do nothing. However, it is possible to check whether an exception was thrown:

```
if (jl_exception_occurred())
    printf("%s\n", jl_typeof_str(jl_exception_occurred()));
```

If you are using the Julia C API from a language that supports exceptions (e.g. Python, C#, C++), it makes sense to wrap each call into `libjulia` with a function that checks whether an exception was thrown, and then rethrows the exception in the host language.

28.6.1 Throwing Julia Exceptions

When writing Julia callable functions, it might be necessary to validate arguments and throw exceptions to indicate errors. A typical type check looks like:

```
if (!jl_is_float64(val)) {  
    jl_type_error(function_name, (jl_value_t*)jl_float64_type, val);  
}
```

General exceptions can be raised using the functions:

```
void jl_error(const char *str);  
void jl_errorf(const char *fmt, ...);
```

`jl_error` takes a C string, and `jl_errorf` is called like `printf`:

```
jl_errorf("argument x = %d is too large", x);
```

where in this example `x` is assumed to be an integer.

Packages

Julia has a built-in package manager for installing add-on functionality written in Julia. It can also install external libraries using your operating system's standard system for doing so, or by compiling from source. The list of registered Julia packages can be found at <http://pkg.julialang.org>. All package manager commands are found in the *Pkg* (page 581) module, included in Julia's Base install.

First we'll go over the mechanics of the *Pkg* family of commands and then we'll provide some guidance on how to get your package registered. Be sure to read the section below on package naming conventions, tagging versions and the importance of a `REQUIRE` file for when you're ready to add your code to the curated METADATA repository.

29.1 Package Status

The *Pkg.status()* (page 584) function prints out a summary of the state of packages you have installed. Initially, you'll have no packages installed:

```
julia> Pkg.status()
INFO: Initializing package repository /Users/stefan/.julia/v0.6
INFO: Cloning METADATA from git://github.com/JuliaLang/METADATA.jl
No packages installed.
```

Your package directory is automatically initialized the first time you run a *Pkg* (page 581) command that expects it to exist – which includes *Pkg.status()* (page 584). Here's an example non-trivial set of required and additional packages:

```
julia> Pkg.status()
Required packages:
 - Distributions          0.2.8
 - UTF16                 0.2.0
Additional packages:
 - NumericExtensions     0.2.17
 - Stats                 0.2.6
```

These packages are all on registered versions, managed by *Pkg* (page 581). Packages can be in more complicated states, indicated by annotations to the right of the installed package version; we will explain these states and annotations as we encounter them. For programmatic usage, *Pkg.installed()* (page 584) returns a dictionary, mapping installed package names to the version of that package which is installed:

```
julia> Pkg.installed()
Dict{String,VersionNumber} with 4 entries:
"Distributions" => v"0.2.8"
"Stats"        => v"0.2.6"
```

```
"UTF16"          => v"0.2.0"
"NumericExtensions" => v"0.2.17"
```

29.2 Adding and Removing Packages

Julia’s package manager is a little unusual in that it is declarative rather than imperative. This means that you tell it what you want and it figures out what versions to install (or remove) to satisfy those requirements optimally – and minimally. So rather than installing a package, you just add it to the list of requirements and then “resolve” what needs to be installed. In particular, this means that if some package had been installed because it was needed by a previous version of something you wanted, and a newer version doesn’t have that requirement anymore, updating will actually remove that package.

Your package requirements are in the file `~/.julia/v0.6/REQUIRE`. You can edit this file by hand and then call `Pkg.resolve()` (page 583) to install, upgrade or remove packages to optimally satisfy the requirements, or you can do `Pkg.edit()` (page 583), which will open `REQUIRE` in your editor (configured via the `EDITOR` or `VISUAL` environment variables), and then automatically call `Pkg.resolve()` (page 583) afterwards if necessary. If you only want to add or remove the requirement for a single package, you can also use the non-interactive `Pkg.add()` (page 583) and `Pkg.rm()` (page 583) commands, which add or remove a single requirement to `REQUIRE` and then call `Pkg.resolve()` (page 583).

You can add a package to the list of requirements with the `Pkg.add()` (page 583) function, and the package and all the packages that it depends on will be installed:

```
julia> Pkg.status()
No packages installed.

julia> Pkg.add("Distributions")
INFO: Cloning cache of Distributions from git://github.com/JuliaStats/Distributions.jl.git
INFO: Cloning cache of NumericExtensions from git://github.com/lindahua/NumericExtensions.jl.git
INFO: Cloning cache of Stats from git://github.com/JuliaStats/Stats.jl.git
INFO: Installing Distributions v0.2.7
INFO: Installing NumericExtensions v0.2.17
INFO: Installing Stats v0.2.6
INFO: REQUIRE updated.

julia> Pkg.status()
Required packages:
- Distributions          0.2.7
Additional packages:
- NumericExtensions     0.2.17
- Stats                 0.2.6
```

What this is doing is first adding `Distributions` to your `~/.julia/v0.6/REQUIRE` file:

```
$ cat ~/.julia/v0.6/REQUIRE
Distributions
```

It then runs `Pkg.resolve()` (page 583) using these new requirements, which leads to the conclusion that the `Distributions` package should be installed since it is required but not installed. As stated before, you can accomplish the same thing by editing your `~/.julia/v0.6/REQUIRE` file by hand and then running `Pkg.resolve()` (page 583) yourself:

```
$ echo UTF16 >> ~/.julia/v0.6/REQUIRE

julia> Pkg.resolve()
INFO: Cloning cache of UTF16 from git://github.com/nolta/UTF16.jl.git
```

```
INFO: Installing UTF16 v0.2.0
```

```
julia> Pkg.status()
Required packages:
- Distributions          0.2.7
- UTF16                  0.2.0
Additional packages:
- NumericExtensions     0.2.17
- Stats                  0.2.6
```

This is functionally equivalent to calling `Pkg.add("UTF16")` (page 583), except that `Pkg.add()` (page 583) doesn't change `REQUIRE` until *after* installation has completed, so if there are problems, `REQUIRE` will be left as it was before calling `Pkg.add()` (page 583). The format of the `REQUIRE` file is described in [Requirements Specification](#) (page 288); it allows, among other things, requiring specific ranges of versions of packages.

When you decide that you don't want to have a package around any more, you can use `Pkg.rm()` (page 583) to remove the requirement for it from the `REQUIRE` file:

```
julia> Pkg.rm("Distributions")
INFO: Removing Distributions v0.2.7
INFO: Removing Stats v0.2.6
INFO: Removing NumericExtensions v0.2.17
INFO: REQUIRE updated.

julia> Pkg.status()
Required packages:
- UTF16                  0.2.0

julia> Pkg.rm("UTF16")
INFO: Removing UTF16 v0.2.0
INFO: REQUIRE updated.

julia> Pkg.status()
No packages installed.
```

Once again, this is equivalent to editing the `REQUIRE` file to remove the line with each package name on it then running `Pkg.resolve()` (page 583) to update the set of installed packages to match. While `Pkg.add()` (page 583) and `Pkg.rm()` (page 583) are convenient for adding and removing requirements for a single package, when you want to add or remove multiple packages, you can call `Pkg.edit()` (page 583) to manually change the contents of `REQUIRE` and then update your packages accordingly. `Pkg.edit()` (page 583) does not roll back the contents of `REQUIRE` if `Pkg.resolve()` (page 583) fails – rather, you have to run `Pkg.edit()` (page 583) again to fix the files contents yourself.

Because the package manager uses libgit2 internally to manage the package git repositories, users may run into protocol issues (if behind a firewall, for example), when running `Pkg.add()` (page 583). By default, all GitHub-hosted packages will be accessed via 'https'; this default can be modified by calling `Pkg.setprotocol!()` (page 584). The following command can be run from the command line in order to tell git to use 'https' instead of the 'git' protocol when cloning all repositories, wherever they are hosted:

```
git config --global url."https://".insteadOf git://
```

However, this change will be system-wide and thus the use of `Pkg.setprotocol!()` (page 584) is preferable.

29.3 Offline Installation of Packages

For machines with no Internet connection, packages may be installed by copying the package root directory (given by `Pkg.dir()` (page 583)) from a machine with the same operating system and environment.

`Pkg.add()` (page 583) does the following within the package root directory:

1. Adds the name of the package to `REQUIRE`.
2. Downloads the package to `.cache`, then copies the package to the package root directory.
3. Recursively performs step 2 against all the packages listed in the package's `REQUIRE` file.
4. Runs `Pkg.build()` (page 584)

Warning: Copying installed packages from a different machine is brittle for packages requiring binary external dependencies. Such packages may break due to differences in operating system versions, build environments, and/or absolute path dependencies.

29.4 Installing Unregistered Packages

Julia packages are simply git repositories, clonable via any of the [protocols](#) that git supports, and containing Julia code that follows certain layout conventions. Official Julia packages are registered in the `METADATA.jl` repository, available at a well-known location ¹. The `Pkg.add()` (page 583) and `Pkg.rm()` (page 583) commands in the previous section interact with registered packages, but the package manager can install and work with unregistered packages too. To install an unregistered package, use `Pkg.clone(url)` (page 583), where `url` is a git URL from which the package can be cloned:

```
julia> Pkg.clone("git://example.com/path/to/Package.jl.git")
INFO: Cloning Package from git://example.com/path/to/Package.jl.git
Cloning into 'Package'...
remote: Counting objects: 22, done.
remote: Compressing objects: 100% (10/10), done.
remote: Total 22 (delta 8), reused 22 (delta 8)
Receiving objects: 100% (22/22), 2.64 KiB, done.
Resolving deltas: 100% (8/8), done.
```

By convention, Julia repository names end with `.jl` (the additional `.git` indicates a “bare” git repository), which keeps them from colliding with repositories for other languages, and also makes Julia packages easy to find in search engines. When packages are installed in your `.julia/v0.6` directory, however, the extension is redundant so we leave it off.

If unregistered packages contain a `REQUIRE` file at the top of their source tree, that file will be used to determine which registered packages the unregistered package depends on, and they will automatically be installed. Unregistered packages participate in the same version resolution logic as registered packages, so installed package versions will be adjusted as necessary to satisfy the requirements of both registered and unregistered packages.

29.5 Updating Packages

When package developers publish new registered versions of packages that you’re using, you will, of course, want the new shiny versions. To get the latest and greatest versions of all your packages, just do `Pkg.update()` (page 584):

¹ The official set of packages is at <https://github.com/JuliaLang/METADATA.jl>, but individuals and organizations can easily use a different metadata repository. This allows control which packages are available for automatic installation. One can allow only audited and approved package versions, and make private packages or forks available. See [Custom METADATA](#) (page 277) for details.


```
julia> Pkg.update()
INFO: Updating METADATA...
INFO: Computing changes...
INFO: Upgrading Distributions: v0.2.8 => v0.2.10
INFO: Upgrading Stats: v0.2.7 => v0.2.8
```

The first step of updating packages is to pull new changes to `~/.julia/v0.6/METADATA` and see if any new registered package versions have been published. After this, `Pkg.update()` (page 584) attempts to update packages that are checked out on a branch and not dirty (i.e. no changes have been made to files tracked by git) by pulling changes from the package’s upstream repository. Upstream changes will only be applied if no merging or rebasing is necessary – i.e. if the branch can be “fast-forwarded”. If the branch cannot be fast-forwarded, it is assumed that you’re working on it and will update the repository yourself.

Finally, the update process recomputes an optimal set of package versions to have installed to satisfy your top-level requirements and the requirements of “fixed” packages. A package is considered fixed if it is one of the following:

1. **Unregistered:** the package is not in METADATA – you installed it with `Pkg.clone()` (page 583).
2. **Checked out:** the package repo is on a development branch.
3. **Dirty:** changes have been made to files in the repo.

If any of these are the case, the package manager cannot freely change the installed version of the package, so its requirements must be satisfied by whatever other package versions it picks. The combination of top-level requirements in `~/.julia/v0.6/REQUIRE` and the requirement of fixed packages are used to determine what should be installed.

You can also update only a subset of the installed packages, by providing arguments to the `Pkg.update` function. In that case, only the packages provided as arguments and their dependencies will be updated:

```
julia> Pkg.update("Example")
INFO: Updating METADATA...
INFO: Computing changes...
INFO: Upgrading Example: v0.4.0 => 0.4.1
```

This partial update process still computes the new set of package versions according to top-level requirements and “fixed” packages, but it additionally considers all other packages except those explicitly provided, and their dependencies, as fixed.

29.6 Checkout, Pin and Free

You may want to use the `master` version of a package rather than one of its registered versions. There might be fixes or functionality on `master` that you need that aren’t yet published in any registered versions, or you may be a developer of the package and need to make changes on `master` or some other development branch. In such cases, you can do `Pkg.checkout(pkg)` (page 584) to checkout the `master` branch of `pkg` or `Pkg.checkout(pkg, branch)` (page 584) to checkout some other branch:

```
julia> Pkg.add("Distributions")
INFO: Installing Distributions v0.2.9
INFO: Installing NumericExtensions v0.2.17
INFO: Installing Stats v0.2.7
INFO: REQUIRE updated.

julia> Pkg.status()
Required packages:
- Distributions          0.2.9
Additional packages:
```

```
- NumericExtensions      0.2.17
- Stats                  0.2.7

julia> Pkg.checkout("Distributions")
INFO: Checking out Distributions master...
INFO: No packages to install, update or remove.

julia> Pkg.status()
Required packages:
- Distributions          0.2.9+          master
Additional packages:
- NumericExtensions      0.2.17
- Stats                  0.2.7
```

Immediately after installing `Distributions` with `Pkg.add()` (page 583) it is on the current most recent registered version – 0.2.9 at the time of writing this. Then after running `Pkg.checkout("Distributions")` (page 584), you can see from the output of `Pkg.status()` (page 584) that `Distributions` is on an unregistered version greater than 0.2.9, indicated by the “pseudo-version” number 0.2.9+.

When you checkout an unregistered version of a package, the copy of the `REQUIRE` file in the package repo takes precedence over any requirements registered in `METADATA`, so it is important that developers keep this file accurate and up-to-date, reflecting the actual requirements of the current version of the package. If the `REQUIRE` file in the package repo is incorrect or missing, dependencies may be removed when the package is checked out. This file is also used to populate newly published versions of the package if you use the API that `Pkg` (page 581) provides for this (described below).

When you decide that you no longer want to have a package checked out on a branch, you can “free” it back to the control of the package manager with `Pkg.free(pkg)` (page 584):

```
julia> Pkg.free("Distributions")
INFO: Freeing Distributions...
INFO: No packages to install, update or remove.

julia> Pkg.status()
Required packages:
- Distributions          0.2.9
Additional packages:
- NumericExtensions      0.2.17
- Stats                  0.2.7
```

After this, since the package is on a registered version and not on a branch, its version will be updated as new registered versions of the package are published.

If you want to pin a package at a specific version so that calling `Pkg.update()` (page 584) won’t change the version the package is on, you can use the `Pkg.pin()` (page 584) function:

```
julia> Pkg.pin("Stats")
INFO: Creating Stats branch pinned.47c198b1.tmp

julia> Pkg.status()
Required packages:
- Distributions          0.2.9
Additional packages:
- NumericExtensions      0.2.17
- Stats                  0.2.7          pinned.47c198b1.tmp
```

After this, the `Stats` package will remain pinned at version 0.2.7 – or more specifically, at commit 47c198b1, but since versions are permanently associated a given git hash, this is the same thing. `Pkg.pin()` (page 584) works by creating a throw-away branch for the commit you want to pin the package at and then checking that branch out. By

default, it pins a package at the current commit, but you can choose a different version by passing a second argument:

```
julia> Pkg.pin("Stats",v"0.2.5")
INFO: Creating Stats branch pinned.1fd0983b.tmp
INFO: No packages to install, update or remove.

julia> Pkg.status()
Required packages:
- Distributions                0.2.9
Additional packages:
- NumericExtensions           0.2.17
- Stats                        0.2.5                pinned.1fd0983b.tmp
```

Now the `Stats` package is pinned at commit `1fd0983b`, which corresponds to version `0.2.5`. When you decide to “unpin” a package and let the package manager update it again, you can use `Pkg.free()` (page 584) like you would to move off of any branch:

```
julia> Pkg.free("Stats")
INFO: Freeing Stats...
INFO: No packages to install, update or remove.

julia> Pkg.status()
Required packages:
- Distributions                0.2.9
Additional packages:
- NumericExtensions           0.2.17
- Stats                        0.2.7
```

After this, the `Stats` package is managed by the package manager again, and future calls to `Pkg.update()` (page 584) will upgrade it to newer versions when they are published. The throw-away `pinned.1fd0983b.tmp` branch remains in your local `Stats` repo, but since git branches are extremely lightweight, this doesn’t really matter; if you feel like cleaning them up, you can go into the repo and delete those branches ².

29.7 Custom METADATA Repository

By default, Julia assumes you will be using the [official METADATA.jl](#) repository for downloading and installing packages. You can also provide a different metadata repository location. A common approach is to keep your `metadata-v2` branch up to date with the Julia official branch and add another branch with your custom packages. You can initialize your local metadata repository using that custom location and branch and then periodically rebase your custom branch with the official `metadata-v2` branch. In order to use a custom repository and branch, issue the following command:

```
julia> Pkg.init("https://me.example.com/METADATA.jl.git", "branch")
```

The branch argument is optional and defaults to `metadata-v2`. Once initialized, a file named `META_BRANCH` in your `~/julia/vX.Y/` path will track the branch that your METADATA repository was initialized with. If you want to change branches, you will need to either modify the `META_BRANCH` file directly (be careful!) or remove the `vX.Y` directory and re-initialize your METADATA repository using the `Pkg.init` command.

² Packages that aren’t on branches will also be marked as dirty if you make changes in the repo, but that’s a less common thing to do.

Package Development

Julia's package manager is designed so that when you have a package installed, you are already in a position to look at its source code and full development history. You are also able to make changes to packages, commit them using git, and easily contribute fixes and enhancements upstream. Similarly, the system is designed so that if you want to create a new package, the simplest way to do so is within the infrastructure provided by the package manager.

30.1 Initial Setup

Since packages are git repositories, before doing any package development you should setup the following standard global git configuration settings:

```
$ git config --global user.name "FULL NAME"
$ git config --global user.email "EMAIL"
```

where `FULL NAME` is your actual full name (spaces are allowed between the double quotes) and `EMAIL` is your actual email address. Although it isn't necessary to use [GitHub](#) to create or publish Julia packages, most Julia packages as of writing this are hosted on GitHub and the package manager knows how to format origin URLs correctly and otherwise work with the service smoothly. We recommend that you create a [free account](#) on GitHub and then do:

```
$ git config --global github.user "USERNAME"
```

where `USERNAME` is your actual GitHub user name. Once you do this, the package manager knows your GitHub user name and can configure things accordingly. You should also [upload](#) your public SSH key to GitHub and set up an [SSH agent](#) on your development machine so that you can push changes with minimal hassle. In the future, we will make this system extensible and support other common git hosting options like [BitBucket](#) and allow developers to choose their favorite. Since the package development functions has been moved to the [PkgDev](#) package, you need to run `Pkg.add("PkgDev"); import PkgDev` to access the functions starting with `PkgDev.` in the document below.

30.2 Making changes to an existing package

30.2.1 Documentation changes

If you want to improve the online documentation of a package, the easiest approach (at least for small changes) is to use GitHub's online editing functionality. First, navigate to the repository's GitHub "home page," find the file (e.g., `README.md`) within the repository's folder structure, and click on it. You'll see the contents displayed, along with a small "pencil" icon in the upper right hand corner. Clicking that icon opens the file in edit mode. Make your changes,

write a brief summary describing the changes you want to make (this is your *commit message*), and then hit “Propose file change.” Your changes will be submitted for consideration by the package owner(s) and collaborators.

For larger documentation changes—and especially ones that you expect to have to update in response to feedback—you might find it easier to use the procedure for code changes described below.

30.2.2 Code changes

Executive summary

Here we assume you’ve already set up git on your local machine and have a GitHub account (see above). Let’s imagine you’re fixing a bug in the Images package:

```
Pkg.checkout("Images")           # check out the master branch
<here, make sure your bug is still a bug and hasn't been fixed already>
cd(Pkg.dir("Images"))
;git checkout -b myfixes         # create a branch for your changes
<edit code>                     # be sure to add a test for your bug
Pkg.test("Images")              # make sure everything works now
;git commit -a -m "Fix foo by calling bar" # write a descriptive message
using PkgDev
PkgDev.submit("Images")
```

The last line will present you with a link to submit a pull request to incorporate your changes.

Detailed description

If you want to fix a bug or add new functionality, you want to be able to test your changes before you submit them for consideration. You also need to have an easy way to update your proposal in response to the package owner’s feedback. Consequently, in this case the strategy is to work locally on your own machine; once you are satisfied with your changes, you submit them for consideration. This process is called a *pull request* because you are asking to “pull” your changes into the project’s main repository. Because the online repository can’t see the code on your private machine, you first *push* your changes to a publicly-visible location, your own online *fork* of the package (hosted on your own personal GitHub account).

Let’s assume you already have the `Foo` package installed. In the description below, anything starting with `Pkg.` or `PkgDev.` is meant to be typed at the Julia prompt; anything starting with `git` is meant to be typed in *Julia’s shell mode* (page 258) (or using the shell that comes with your operating system). Within Julia, you can combine these two modes:

```
julia> cd(Pkg.dir("Foo"))           # go to Foo's folder
shell> git command arguments...    # command will apply to Foo
```

Now suppose you’re ready to make some changes to `Foo`. While there are several possible approaches, here is one that is widely used:

- From the Julia prompt, type `Pkg.checkout("Foo")` (page 584). This ensures you’re running the latest code (the *master* branch), rather than just whatever “official release” version you have installed. (If you’re planning to fix a bug, at this point it’s a good idea to check again whether the bug has already been fixed by someone else. If it has, you can request that a new official release be tagged so that the fix gets distributed to the rest of the community.) If you receive an error `Foo is dirty`, bailing, see *Dirty packages* (page 281) below.
- Create a branch for your changes: navigate to the package folder (the one that Julia reports from `Pkg.dir("Foo")` (page 583)) and (in shell mode) create a new branch using `git checkout -b`

`<newbranch>`, where `<newbranch>` might be some descriptive name (e.g., `fixbar`). By creating a branch, you ensure that you can easily go back and forth between your new work and the current `master` branch (see <https://git-scm.com/book/en/v2/Git-Branching-Branches-in-a-Nutshell>).

If you forget to do this step until after you’ve already made some changes, don’t worry: see [more detail about branching](#) (page 282) below.

- Make your changes. Whether it’s fixing a bug or adding new functionality, in most cases your change should include updates to both the `src/` and `test/` folders. If you’re fixing a bug, add your minimal example demonstrating the bug (on the current code) to the test suite; by contributing a test for the bug, you ensure that the bug won’t accidentally reappear at some later time due to other changes. If you’re adding new functionality, creating tests demonstrates to the package owner that you’ve made sure your code works as intended.
- Run the package’s tests and make sure they pass. There are several ways to run the tests:
 - From Julia, run `Pkg.test("Foo")` (page 584): this will run your tests in a separate (new) julia process.
 - From Julia, `include("runtests.jl")` from the package’s `test/` folder (it’s possible the file has a different name, look for one that runs all the tests): this allows you to run the tests repeatedly in the same session without reloading all the package code; for packages that take a while to load, this can be much faster. With this approach, you do have to do some extra work to make [changes in the package code](#) (page 325).
 - From the shell, run `julia ../test/runtests.jl` from within the package’s `src/` folder.
- Commit your changes: see <https://git-scm.com/book/en/v2/Git-Basics-Recording-Changes-to-the-Repository>.
- Submit your changes: From the Julia prompt, type `PkgDev.submit("Foo")`. This will push your changes to your GitHub fork, creating it if it doesn’t already exist. (If you encounter an error, [make sure you’ve set up your SSH keys](#) (page 279).) Julia will then give you a hyperlink; open that link, edit the message, and then click “submit.” At that point, the package owner will be notified of your changes and may initiate discussion. (If you are comfortable with git, you can also do these steps manually from the shell.)
- The package owner may suggest additional improvements. To respond to those suggestions, you can easily update the pull request (this only works for changes that have not already been merged; for merged pull requests, make new changes by starting a new branch):
 - If you’ve changed branches in the meantime, make sure you go back to the same branch with `git checkout fixbar` (from shell mode) or `Pkg.checkout("Foo", "fixbar")` (page 584) (from the Julia prompt).
 - As above, make your changes, run the tests, and commit your changes.
 - From the shell, type `git push`. This will add your new commit(s) to the same pull request; you should see them appear automatically on the page holding the discussion of your pull request.

One potential type of change the owner may request is that you squash your commits. See [Squashing](#) (page 282) below.

30.2.3 Dirty packages

If you can’t change branches because the package manager complains that your package is dirty, it means you have some changes that have not been committed. From the shell, use `git diff` to see what these changes are; you can either discard them (`git checkout changedfile.jl`) or commit them before switching branches. If you can’t easily resolve the problems manually, as a last resort you can delete the entire `"Foo"` folder and reinstall a fresh copy with `Pkg.add("Foo")` (page 583). Naturally, this deletes any changes you’ve made.

30.2.4 Making a branch *post hoc*

Especially for newcomers to git, one often forgets to create a new branch until after some changes have already been made. If you haven't yet staged or committed your changes, you can create a new branch with `git checkout -b <newbranch>` just as usual—git will kindly show you that some files have been modified and create the new branch for you. *Your changes have not yet been committed to this new branch*, so the normal work rules still apply.

However, if you've already made a commit to `master` but wish to go back to the official `master` (called `origin/master`), use the following procedure:

- Create a new branch. This branch will hold your changes.
- Make sure everything is committed to this branch.
- `git checkout master`. If this fails, *do not* proceed further until you have resolved the problems, or you may lose your changes.
- *Reset* `master` (your current branch) back to an earlier state with `git reset --hard origin/master` (see <https://git-scm.com/blog/2011/07/11/reset.html>).

This requires a bit more familiarity with git, so it's much better to get in the habit of creating a branch at the outset.

30.2.5 Squashing and rebasing

Depending on the tastes of the package owner (s)he may ask you to “squash” your commits. This is especially likely if your change is quite simple but your commit history looks like this:

```
WIP: add new 1-line whizbang function (currently breaks package)
Finish whizbang function
Fix typo in variable name
Oops, don't forget to supply default argument
Split into two 1-line functions
Rats, forgot to export the second function
...
```

This gets into the territory of more advanced git usage, and you're encouraged to do some reading (<https://git-scm.com/book/en/v2/Git-Branching-Rebasing>). However, a brief summary of the procedure is as follows:

- To protect yourself from error, start from your `fixbar` branch and create a new branch with `git checkout -b fixbar_backup`. Since you started from `fixbar`, this will be a copy. Now go back to the one you intend to modify with `git checkout fixbar`.
- From the shell, type `git rebase -i origin/master`.
- To combine commits, change `pick` to `squash` (for additional options, consult other sources). Save the file and close the editor window.
- Edit the combined commit message.

If the rebase goes badly, you can go back to the beginning to try again like this:

```
git checkout fixbar
git reset --hard fixbar_backup
```

Now let's assume you've rebased successfully. Since your `fixbar` repository has now diverged from the one in your GitHub fork, you're going to have to do a *force push*:

- To make it easy to refer to your GitHub fork, create a “handle” for it with `git remote add myfork https://github.com/myaccount/Foo.jl.git`, where the URL comes from the “clone URL” on your GitHub fork's page.

- Force-push to your fork with `git push myfork +fixbar`. The `+` indicates that this should replace the `fixbar` branch found at `myfork`.

30.3 Creating a new Package

30.3.1 REQUIRE speaks for itself

You should have a `REQUIRE` file in your package repository, with a bare minimum directive of what Julia version you expect your users to be running for the package to work. Putting a floor on what Julia version your package supports is done by simply adding `julia 0.x` in this file. While this line is partly informational, it also has the consequence of whether `Pkg.update()` will update code found in `.julia` version directories. It will not update code found in version directories beneath the floor of what's specified in your `REQUIRE`.

As the development version `0.y` matures, you may find yourself using it more frequently, and wanting your package to support it. Be warned, the development branch of Julia is the land of breakage, and you can expect things to break. When you go about fixing whatever broke your package in the development `0.y` branch, you will likely find that you just broke your package on the stable version.

There is a mechanism found in the `Compat` package that will enable you to support both the stable version and breaking changes found in the development version. Should you decide to use this solution, you will need to add `Compat` to your `REQUIRE` file. In this case, you will still have `julia 0.x` in your `REQUIRE`. The `x` is the floor version of what your package supports.

You might also have no interest in supporting the development version of Julia. Just as you can add a floor to the version you expect your users to be on, you can set an upper bound. In this case, you would put `julia 0.x 0.y-` in your `REQUIRE` file. The `-` at the end of the version number means pre-release versions of that specific version from the very first commit. By setting it as the ceiling, you mean the code supports everything up to but not including the ceiling version.

Another scenario is that you are writing the bulk of the code for your package with Julia `0.y` and do not want to support the current stable version of Julia. If you choose to do this, simply add `julia 0.y-` to your `REQUIRE`. Just remember to change the `julia 0.y-` to `julia 0.y` in your `REQUIRE` file once `0.y` is officially released. If you don't edit the dash cruft you are suggesting that you support both the development and stable versions of the same version number! That would be madness. See the *Requirements Specification* (page 288) for the full format of `REQUIRE`.

30.3.2 Guidelines for naming a package

Package names should be sensible to most Julia users, *even to those who are not domain experts*. When you submit your package to METADATA, you can expect a little back and forth about the package name with collaborators, especially if it's ambiguous or can be confused with something other than what it is. During this bike-shedding, it's not uncommon to get a range of *different* name suggestions. These are only suggestions though, with the intent being to keep a tidy namespace in the curated METADATA repository. Since this repository belongs to the entire community, there will likely be a few collaborators who care about your package name. Here are some guidelines to follow in naming your package:

1. Avoid jargon. In particular, avoid acronyms unless there is minimal possibility of confusion.
 - It's ok to say `USA` if you're talking about the USA.
 - It's not ok to say `PMA`, even if you're talking about positive mental attitude.
2. Avoid using `Julia` in your package name.
 - It is usually clear from context and to your users that the package is a Julia package.

- Having Julia in the name can imply that the package is connected to, or endorsed by, contributors to the Julia language itself.
3. Packages that provide most of their functionality in association with a new type should have pluralized names.
 - `DataFrames` provides the `DataFrame` type.
 - `BloomFilters` provides the `BloomFilter` type.
 - In contrast, `JuliaParser` provides no new type, but instead new functionality in the `JuliaParser.parse()` function.
 4. Err on the side of clarity, even if clarity seems long-winded to you.
 - `RandomMatrices` is a less ambiguous name than `RndMat` or `RMT`, even though the latter are shorter.
 5. A less systematic name may suit a package that implements one of several possible approaches to its domain.
 - Julia does not have a single comprehensive plotting package. Instead, `Gadfly`, `PyPlot`, `Winston` and other packages each implement a unique approach based on a particular design philosophy.
 - In contrast, `SortingAlgorithms` provides a consistent interface to use many well-established sorting algorithms.
 6. Packages that wrap external libraries or programs should be named after those libraries or programs.
 - `CPLEX.jl` wraps the `CPLEX` library, which can be identified easily in a web search.
 - `MATLAB.jl` provides an interface to call the `MATLAB` engine from within Julia.

30.3.3 Generating the package

Suppose you want to create a new Julia package called `FooBar`. To get started, do `PkgDev.generate(pkg, license)` where `pkg` is the new package name and `license` is the name of a license that the package generator knows about:

```
julia> PkgDev.generate("FooBar", "MIT")
INFO: Initializing FooBar repo: /Users/stefan/.julia/v0.6/FooBar
INFO: Origin: git://github.com/StefanKarpinski/FooBar.jl.git
INFO: Generating LICENSE.md
INFO: Generating README.md
INFO: Generating src/FooBar.jl
INFO: Generating test/runtests.jl
INFO: Generating REQUIRE
INFO: Generating .travis.yml
INFO: Generating appveyor.yml
INFO: Generating .gitignore
INFO: Committing FooBar generated files
```

This creates the directory `~/.julia/v0.6/FooBar`, initializes it as a git repository, generates a bunch of files that all packages should have, and commits them to the repository:

```
$ cd ~/.julia/v0.6/FooBar && git show --stat

commit 84b8e266dae6de30ab9703150b3bf771ec7b6285
Author: Stefan Karpinski <stefan@karpinski.org>
Date:   Wed Oct 16 17:57:58 2013 -0400

    FooBar.jl generated files.

    license: MIT
```

```

authors: Stefan Karpinski
years:   2013
user:    StefanKarpinski

Julia Version 0.3.0-prerelease+3217 [5fcfb13*]

.gitignore      |  2 ++
.travis.yml     | 13 ++++++
LICENSE.md      | 22 ++++++
README.md       |  3 +++
REQUIRED       |  1 +
appveyor.yml    | 34 ++++++
src/FooBar.jl   |  5 +++++
test/runtests.jl |  5 +++++
8 files changed, 85 insertions(+)

```

At the moment, the package manager knows about the MIT “Expat” License, indicated by "MIT", the Simplified BSD License, indicated by "BSD", and version 2.0 of the Apache Software License, indicated by "ASL". If you want to use a different license, you can ask us to add it to the package generator, or just pick one of these three and then modify the `~/.julia/v0.6/PACKAGE/LICENSE.md` file after it has been generated.

If you created a GitHub account and configured git to know about it, `PkgDev.generate()` will set an appropriate origin URL for you. It will also automatically generate a `.travis.yml` file for using the [Travis](#) automated testing service, and an `appveyor.yml` file for using [AppVeyor](#). You will have to enable testing on the Travis and AppVeyor websites for your package repository, but once you’ve done that, it will already have working tests. Of course, all the default testing does is verify that using `FooBar` in Julia works.

30.3.4 Loading Static Non-Julia Files

If your package code needs to load static files which are not Julia code, e.g. an external library or data files, and are located within the package directory, use the `@__DIR__` macro to determine the directory of the current source file. For example if `FooBar/src/FooBar.jl` needs to load `FooBar/data/foo.csv`, use the following code:

```

datapath = joinpath(@__DIR__, "..", "data")
foo = readcsv(joinpath(datapath, "foo.csv"))

```

30.3.5 Making Your Package Available

Once you’ve made some commits and you’re happy with how `FooBar` is working, you may want to get some other people to try it out. First you’ll need to create the remote repository and push your code to it; we don’t yet automatically do this for you, but we will in the future and it’s not too hard to figure out³. Once you’ve done this, letting people try out your code is as simple as sending them the URL of the published repo – in this case:

```
git://github.com/StefanKarpinski/FooBar.jl.git
```

For your package, it will be your GitHub user name and the name of your package, but you get the idea. People you send this URL to can use `Pkg.clone()` (page 583) to install the package and try it out:

```

julia> Pkg.clone("git://github.com/StefanKarpinski/FooBar.jl.git")
INFO: Cloning FooBar from git@github.com:StefanKarpinski/FooBar.jl.git

```

³ Installing and using GitHub’s “hub” tool is highly recommended. It allows you to do things like run `hub create` in the package repo and have it automatically created via GitHub’s API.

30.3.6 Tagging and Publishing Your Package

Once you’ve decided that `FooBar` is ready to be registered as an official package, you can add it to your local copy of METADATA using `PkgDev.register()`:

```
julia> PkgDev.register("FooBar")
INFO: Registering FooBar at git://github.com/StefanKarpinski/FooBar.jl.git
INFO: Committing METADATA for FooBar
```

This creates a commit in the `~/.julia/v0.6/METADATA` repo:

```
$ cd ~/.julia/v0.6/METADATA && git show

commit 9f71f4becb05cadacb983c54a72eed744e5c019d
Author: Stefan Karpinski <stefan@karpinski.org>
Date:   Wed Oct 16 18:46:02 2013 -0400

    Register FooBar

diff --git a/FooBar/url b/FooBar/url
new file mode 100644
index 0000000..30e525e
--- /dev/null
+++ b/FooBar/url
@@ -0,0 +1 @@
+git://github.com/StefanKarpinski/FooBar.jl.git
```

This commit is only locally visible, however. To make it visible to the Julia community, you need to merge your local METADATA upstream into the official repo. The `PkgDev.publish()` command will fork the METADATA repository on GitHub, push your changes to your fork, and open a pull request:

```
julia> PkgDev.publish()
INFO: Validating METADATA
INFO: No new package versions to publish
INFO: Submitting METADATA changes
INFO: Forking JuliaLang/METADATA.jl to StefanKarpinski
INFO: Pushing changes as branch pull-request/ef45f54b
INFO: To create a pull-request open:

https://github.com/StefanKarpinski/METADATA.jl/compare/pull-request/ef45f54b
```

Tip: If `PkgDev.publish()` fails with error:

```
ERROR: key not found: "token"
```

then you may have encountered an issue from using the GitHub API on multiple systems. The solution is to delete the “Julia Package Manager” personal access token [from your Github account](#) and try again.

Other failures may require you to circumvent `PkgDev.publish()` by [creating a pull request on GitHub](#). See: [Publishing METADATA manually](#) (page 287) below.

Once the package URL for `FooBar` is registered in the official METADATA repo, people know where to clone the package from, but there still aren’t any registered versions available. You can tag and register it with the `PkgDev.tag()` command:

```
julia> PkgDev.tag("FooBar")
INFO: Tagging FooBar v0.0.1
INFO: Committing METADATA for FooBar
```

This tags `v0.0.1` in the `FooBar` repo:

```
$ cd ~/.julia/v0.6/FooBar && git tag
v0.0.1
```

It also creates a new version entry in your local METADATA repo for `FooBar`:

```
$ cd ~/.julia/v0.6/FooBar && git show
commit de77ee4dc0689b12c5e8b574aef7f70e8b311b0e
Author: Stefan Karpinski <stefan@karpinski.org>
Date:   Wed Oct 16 23:06:18 2013 -0400

    Tag FooBar v0.0.1

diff --git a/FooBar/versions/0.0.1/sha1 b/FooBar/versions/0.0.1/sha1
new file mode 100644
index 0000000..c1cb1c1
--- /dev/null
+++ b/FooBar/versions/0.0.1/sha1
@@ -0,0 +1 @@
+84b8e266dae6de30ab9703150b3bf771ec7b6285
```

The `PkgDev.tag()` command takes an optional second argument that is either an explicit version number object like `v"0.0.1"` or one of the symbols `:patch`, `:minor` or `:major`. These increment the patch, minor or major version number of your package intelligently.

Adding a tagged version of your package will expedite the official registration into `METADATA.jl` by collaborators. It is strongly recommended that you complete this process, regardless if your package is completely ready for an official release.

As a general rule, packages should be tagged `0.0.1` first. Since Julia itself hasn't achieved `1.0` status, it's best to be conservative in your package's tagged versions.

As with `PkgDev.register()`, these changes to `METADATA` aren't available to anyone else until they've been included upstream. Again, use the `PkgDev.publish()` command, which first makes sure that individual package repos have been tagged, pushes them if they haven't already been, and then opens a pull request to `METADATA`:

```
julia> PkgDev.publish()
INFO: Validating METADATA
INFO: Pushing FooBar permanent tags: v0.0.1
INFO: Submitting METADATA changes
INFO: Forking JuliaLang/METADATA.jl to StefanKarpinski
INFO: Pushing changes as branch pull-request/3ef4f5c4
INFO: To create a pull-request open:

https://github.com/StefanKarpinski/METADATA.jl/compare/pull-request/3ef4f5c4
```

Publishing METADATA manually

If `PkgDev.publish()` fails you can follow these instructions to manually publish your package.

By “forking” the main `METADATA` repository, you can create a personal copy (of `METADATA.jl`) under your GitHub account. Once that copy exists, you can push your local changes to your copy (just like any other GitHub project).

1. go to https://github.com/login?return_to=https%3A%2F%2Fgithub.com%2FJuliaLang%2FMETADATA.jl%2Ffork and create your own fork.
2. add your fork as a remote repository for the `METADATA` repository on your local computer (in the terminal where `USERNAME` is your github username):

```
cd ~/.julia/v0.6/METADATA
git remote add USERNAME https://github.com/USERNAME/METADATA.jl.git
```

3. push your changes to your fork:

```
git push USERNAME metadata-v2
```

4. If all of that works, then go back to the GitHub page for your fork, and click the “pull request” link.

30.4 Fixing Package Requirements

If you need to fix the registered requirements of an already-published package version, you can do so just by editing the metadata for that version, which will still have the same commit hash – the hash associated with a version is permanent:

```
$ cd ~/.julia/v0.6/METADATA/FooBar/versions/0.0.1 && cat requires
julia 0.3-
$ vi requires
```

Since the commit hash stays the same, the contents of the `REQUIRE` file that will be checked out in the repo will **not** match the requirements in `METADATA` after such a change; this is unavoidable. When you fix the requirements in `METADATA` for a previous version of a package, however, you should also fix the `REQUIRE` file in the current version of the package.

30.5 Requirements Specification

The `~/.julia/v0.6/REQUIRE` file, the `REQUIRE` file inside packages, and the `METADATA` package `requires` files use a simple line-based format to express the ranges of package versions which need to be installed. Package `REQUIRE` and `METADATA requires` files should also include the range of versions of `julia` the package is expected to work with.

Here’s how these files are parsed and interpreted.

- Everything after a `#` mark is stripped from each line as a comment.
- If nothing but whitespace is left, the line is ignored.
- If there are non-whitespace characters remaining, the line is a requirement and the is split on whitespace into words.

The simplest possible requirement is just the name of a package name on a line by itself:

```
Distributions
```

This requirement is satisfied by any version of the `Distributions` package. The package name can be followed by zero or more version numbers in ascending order, indicating acceptable intervals of versions of that package. One version opens an interval, while the next closes it, and the next opens a new interval, and so on; if an odd number of version numbers are given, then arbitrarily large versions will satisfy; if an even number of version numbers are given, the last one is an upper limit on acceptable version numbers. For example, the line:

```
Distributions 0.1
```

is satisfied by any version of `Distributions` greater than or equal to `0.1.0`. Suffixing a version with `-` allows any pre-release versions as well. For example:

```
Distributions 0.1-
```

is satisfied by pre-release versions such as 0.1-dev or 0.1-rc1, or by any version greater than or equal to 0.1.0.

This requirement entry:

```
Distributions 0.1 0.2.5
```

is satisfied by versions from 0.1.0 up to, but not including 0.2.5. If you want to indicate that any 0.1.x version will do, you will want to write:

```
Distributions 0.1 0.2-
```

If you want to start accepting versions after 0.2.7, you can write:

```
Distributions 0.1 0.2- 0.2.7
```

If a requirement line has leading words that begin with @, it is a system-dependent requirement. If your system matches these system conditionals, the requirement is included, if not, the requirement is ignored. For example:

```
@osx Homebrew
```

will require the Homebrew package only on systems where the operating system is OS X. The system conditions that are currently supported are (hierarchically):

- @unix
 - @linux
 - @bsd
 - * @osx
- @windows

The @unix condition is satisfied on all UNIX systems, including Linux and BSD. Negated system conditionals are also supported by adding a ! after the leading @. Examples:

```
@!windows
@unix @!osx
```

The first condition applies to any system but Windows and the second condition applies to any UNIX system besides OS X.

Runtime checks for the current version of Julia can be made using the built-in VERSION variable, which is of type VersionNumber. Such code is occasionally necessary to keep track of new or deprecated functionality between various releases of Julia. Examples of runtime checks:

```
VERSION < v"0.3-" #exclude all pre-release versions of 0.3

v"0.2-" <= VERSION < v"0.3-" #get all 0.2 versions, including pre-releases, up to the above

v"0.2" <= VERSION < v"0.3-" #To get only stable 0.2 versions (Note v"0.2" == v"0.2.0")

VERSION >= v"0.2.1" #get at least version 0.2.1
```

See the section on *version number literals* (page 53) for a more complete description.

Profiling

The *Profile* (page 616) module provides tools to help developers improve the performance of their code. When used, it takes measurements on running code, and produces output that helps you understand how much time is spent on individual line(s). The most common usage is to identify “bottlenecks” as targets for optimization.

Profile (page 616) implements what is known as a “sampling” or *statistical profiler*. It works by periodically taking a backtrace during the execution of any task. Each backtrace captures the currently-running function and line number, plus the complete chain of function calls that led to this line, and hence is a “snapshot” of the current state of execution.

If much of your run time is spent executing a particular line of code, this line will show up frequently in the set of all backtraces. In other words, the “cost” of a given line—or really, the cost of the sequence of function calls up to and including this line—is proportional to how often it appears in the set of all backtraces.

A sampling profiler does not provide complete line-by-line coverage, because the backtraces occur at intervals (by default, 1 ms on Unix systems and 10 ms on Windows, although the actual scheduling is subject to operating system load). Moreover, as discussed further below, because samples are collected at a sparse subset of all execution points, the data collected by a sampling profiler is subject to statistical noise.

Despite these limitations, sampling profilers have substantial strengths:

- You do not have to make any modifications to your code to take timing measurements (in contrast to the alternative *instrumenting profiler*).
- It can profile into Julia’s core code and even (optionally) into C and Fortran libraries.
- By running “infrequently” there is very little performance overhead; while profiling, your code can run at nearly native speed.

For these reasons, it’s recommended that you try using the built-in sampling profiler before considering any alternatives.

31.1 Basic usage

Let’s work with a simple test case:

```
function myfunc()  
    A = rand(100, 100, 200)  
    maximum(A)  
end
```

It’s a good idea to first run the code you intend to profile at least once (unless you want to profile Julia’s JIT-compiler):

```
julia> myfunc() # run once to force compilation
```

Now we’re ready to profile this function:

```
julia> @profile myfunc()
```

To see the profiling results, there is a [graphical browser](#) available, but here we’ll use the text-based display that comes with the standard library:

```
julia> Profile.print()
 23 client.jl; _start; line: 373
 23 client.jl; run_repl; line: 166
 23 client.jl; eval_user_input; line: 91
 23 profile.jl; anonymous; line: 14
   8 none; myfunc; line: 2
   8 dsfmt.jl; dsfmt_gv_fill_array_close_open!; line: 128
 15 none; myfunc; line: 3
   2 reduce.jl; max; line: 35
   2 reduce.jl; max; line: 36
 11 reduce.jl; max; line: 37
```

Each line of this display represents a particular spot (line number) in the code. Indentation is used to indicate the nested sequence of function calls, with more-indented lines being deeper in the sequence of calls. In each line, the first “field” indicates the number of backtraces (samples) taken *at this line or in any functions executed by this line*. The second field is the file name, followed by a semicolon; the third is the function name followed by a semicolon, and the fourth is the line number. Note that the specific line numbers may change as Julia’s code changes; if you want to follow along, it’s best to run this example yourself.

In this example, we can see that the top level is `client.jl`’s `_start` function. This is the first Julia function that gets called when you launch Julia. If you examine line 373 of `client.jl`, you’ll see that (at the time of this writing) it calls `run_repl()`, mentioned on the second line. This in turn calls `eval_user_input()`. These are the functions in `client.jl` that interpret what you type at the REPL, and since we’re working interactively these functions were invoked when we entered `@profile myfunc()`. The next line reflects actions taken in the `@profile` (page 617) macro.

The first line shows that 23 backtraces were taken at line 373 of `client.jl`, but it’s not that this line was “expensive” on its own: the second line reveals that all 23 of these backtraces were actually triggered inside its call to `run_repl`, and so on. To find out which operations are actually taking the time, we need to look deeper in the call chain.

The first “important” line in this output is this one:

```
8 none; myfunc; line: 2
```

`none` refers to the fact that we defined `myfunc` in the REPL, rather than putting it in a file; if we had used a file, this would show the file name. Line 2 of `myfunc()` contains the call to `rand`, and there were 8 (out of 23) backtraces that occurred at this line. Below that, you can see a call to `dsfmt_gv_fill_array_close_open!()` inside `dsfmt.jl`. You might be surprised not to see the `rand` function listed explicitly: that’s because `rand` is *inlined*, and hence doesn’t appear in the backtraces.

A little further down, you see:

```
15 none; myfunc; line: 3
```

Line 3 of `myfunc` contains the call to `max`, and there were 15 (out of 23) backtraces taken here. Below that, you can see the specific places in `base/reduce.jl` that carry out the time-consuming operations in the `max` function for this type of input data.

Overall, we can tentatively conclude that finding the maximum element is approximately twice as expensive as generating the random numbers. We could increase our confidence in this result by collecting more samples:

```
julia> @profile (for i = 1:100; myfunc(); end)
```

```
julia> Profile.print()
3121 client.jl; _start; line: 373
3121 client.jl; run_repl; line: 166
3121 client.jl; eval_user_input; line: 91
3121 profile.jl; anonymous; line: 1
848 none; myfunc; line: 2
842 dsfmt.jl; dsfmt_gv_fill_array_close_open!; line: 128
1510 none; myfunc; line: 3
74 reduce.jl; max; line: 35
122 reduce.jl; max; line: 36
1314 reduce.jl; max; line: 37
```

In general, if you have N samples collected at a line, you can expect an uncertainty on the order of \sqrt{N} (barring other sources of noise, like how busy the computer is with other tasks). The major exception to this rule is garbage collection, which runs infrequently but tends to be quite expensive. (Since Julia’s garbage collector is written in C, such events can be detected using the `C=true` output mode described below, or by using [ProfileView.jl](#).)

This illustrates the default “tree” dump; an alternative is the “flat” dump, which accumulates counts independent of their nesting:

```
julia> Profile.print(format=:flat)
Count File           Function           Line
3121 client.jl        _start            373
3121 client.jl        eval_user_input   91
3121 client.jl        run_repl          166
842  dsfmt.jl         dsfmt_gv_fill_array_close_open! 128
848  none             myfunc            2
1510 none             myfunc            3
3121 profile.jl       anonymous          1
74   reduce.jl        max               35
122  reduce.jl        max               36
1314 reduce.jl        max               37
```

If your code has recursion, one potentially-confusing point is that a line in a “child” function can accumulate more counts than there are total backtraces. Consider the following function definitions:

```
dumbsum(n::Integer) = n == 1 ? 1 : 1 + dumbsum(n-1)
dumbsum3() = dumbsum(3)
```

If you were to profile `dumbsum3`, and a backtrace was taken while it was executing `dumbsum(1)`, the backtrace would look like this:

```
dumbsum3
  dumbsum(3)
    dumbsum(2)
      dumbsum(1)
```

Consequently, this child function gets 3 counts, even though the parent only gets one. The “tree” representation makes this much clearer, and for this reason (among others) is probably the most useful way to view the results.

31.2 Accumulation and clearing

Results from `@profile` (page 617) accumulate in a buffer; if you run multiple pieces of code under `@profile` (page 617), then `Profile.print()` (page 617) will show you the combined results. This can be very useful, but sometimes you want to start fresh; you can do so with `Profile.clear()` (page 617).

31.3 Options for controlling the display of profile results

`Profile.print()` (page 617) has more options than we’ve described so far. Let’s see the full declaration:

```
function print(io::IO = STDOUT, data = fetch(); format = :tree, C = false, combine = true, cols = tt
```

Let’s discuss these arguments in order:

- The first argument allows you to save the results to a file, but the default is to print to STDOUT (the console).
- The second argument contains the data you want to analyze; by default that is obtained from `Profile.fetch()` (page 617), which pulls out the backtraces from a pre-allocated buffer. For example, if you want to profile the profiler, you could say:

```
data = copy(Profile.fetch())
Profile.clear()
@profile Profile.print(STDOUT, data) # Prints the previous results
Profile.print()                    # Prints results from Profile.print()
```

- The first keyword argument, `format`, was introduced above. The possible choices are `:tree` and `:flat`.
- `C`, if set to `true`, allows you to see even the calls to C code. Try running the introductory example with `Profile.print(C = true)`. This can be extremely helpful in deciding whether it’s Julia code or C code that is causing a bottleneck; setting `C=true` also improves the interpretability of the nesting, at the cost of longer profile dumps.
- Some lines of code contain multiple operations; for example, `s += A[i]` contains both an array reference (`A[i]`) and a sum operation. These correspond to different lines in the generated machine code, and hence there may be two or more different addresses captured during backtraces on this line. `combine=true` lumps them together, and is probably what you typically want, but you can generate an output separately for each unique instruction pointer with `combine=false`.
- `cols` allows you to control the number of columns that you are willing to use for display. When the text would be wider than the display, you might see output like this:

```
33 inference.jl; abstract_call; line: 645
 33 inference.jl; abstract_call; line: 645
   33 ..ence.jl; abstract_call_gf; line: 567
    33 ..nce.jl; typeinf; line: 1201
     +1 5 ..nce.jl; ..t_interpret; line: 900
     +3 5 ..ence.jl; abstract_eval; line: 758
     +4 5 ..ence.jl; ..ct_eval_call; line: 733
     +6 5 ..ence.jl; abstract_call; line: 645
```

File/function names are sometimes truncated (with `...`), and indentation is truncated with a `+n` at the beginning, where `n` is the number of extra spaces that would have been inserted, had there been room. If you want a complete profile of deeply-nested code, often a good idea is to save to a file and use a very wide `cols` setting:

```
s = open("/tmp/prof.txt", "w")
Profile.print(s, cols = 500)
close(s)
```

- `maxdepth` can be used to limit the size of the output in `:tree` format (it nests only up to level `maxdepth`)
- `sortedby = :count` sorts the `:flat` format in order of increasing counts

31.4 Configuration

`@profile` (page 617) just accumulates backtraces, and the analysis happens when you call `Profile.print()` (page 617). For a long-running computation, it's entirely possible that the pre-allocated buffer for storing backtraces will be filled. If that happens, the backtraces stop but your computation continues. As a consequence, you may miss some important profiling data (you will get a warning when that happens).

You can obtain and configure the relevant parameters this way:

```
Profile.init()           # returns the current settings
Profile.init(n, delay)
Profile.init(delay = 0.01)
```

`n` is the total number of instruction pointers you can store, with a default value of 10^6 . If your typical backtrace is 20 instruction pointers, then you can collect 50000 backtraces, which suggests a statistical uncertainty of less than 1%. This may be good enough for most applications.

Consequently, you are more likely to need to modify `delay`, expressed in seconds, which sets the amount of time that Julia gets between snapshots to perform the requested computations. A very long-running job might not need frequent backtraces. The default setting is `delay = 0.001`. Of course, you can decrease the delay as well as increase it; however, the overhead of profiling grows once the delay becomes similar to the amount of time needed to take a backtrace (~30 microseconds on the author's laptop).

Memory allocation analysis

One of the most common techniques to improve performance is to reduce memory allocation. The total amount of allocation can be measured with `@time` (page 373) and `@allocated` (page 373), and specific lines triggering allocation can often be inferred from profiling via the cost of garbage collection that these lines incur. However, sometimes it is more efficient to directly measure the amount of memory allocated by each line of code.

To measure allocation line-by-line, start Julia with the `--track-allocation=<setting>` command-line option, for which you can choose `none` (the default, do not measure allocation), `user` (measure memory allocation everywhere except Julia's core code), or `all` (measure memory allocation at each line of Julia code). Allocation gets measured for each line of compiled code. When you quit Julia, the cumulative results are written to text files with `.mem` appended after the file name, residing in the same directory as the source file. Each line lists the total number of bytes allocated. The `Coverage` package contains some elementary analysis tools, for example to sort the lines in order of number of bytes allocated.

In interpreting the results, there are a few important details. Under the `user` setting, the first line of any function directly called from the REPL will exhibit allocation due to events that happen in the REPL code itself. More significantly, JIT-compilation also adds to allocation counts, because much of Julia's compiler is written in Julia (and compilation usually requires memory allocation). The recommended procedure is to force compilation by executing all the commands you want to analyze, then call `Profile.clear_malloc_data()` (page 618) to reset all allocation counters. Finally, execute the desired commands and quit Julia to trigger the generation of the `.mem` files.

Stack Traces

The `StackTraces` module provides simple stack traces that are both human readable and easy to use programmatically.

33.1 Viewing a stack trace

The primary function used to obtain a stack trace is `stacktrace()`:

```
julia> stacktrace()
3-element Array{StackFrame,1}:
 eval at boot.jl:265
 [inlined code from REPL.jl:3] eval_user_input at REPL.jl:62
 [inlined code from REPL.jl:92] anonymous at task.jl:63
```

Calling `stacktrace()` returns a vector of `StackFrame`s. For ease of use, the alias `StackTrace` can be used in place of `Vector{StackFrame}`. (Examples with `...` indicate that output may vary depending on how the code is run.)

```
julia> example() = stacktrace()
example (generic function with 1 method)

julia> example()
11-element Array{StackFrame,1}:
  in example() at none:1
  in eval(::Module, ::Any) at boot.jl:236
  ...

julia> @noinline child() = stacktrace()
child (generic function with 1 method)

julia> @noinline parent() = child()
parent (generic function with 1 method)

julia> grandparent() = parent()
grandparent (generic function with 1 method)

julia> grandparent()
13-element Array{StackFrame,1}:
  in child() at none:1
  in parent() at none:1
  in grandparent() at none:1
  ...
```

Note that when calling `stacktrace()` you'll typically see a frame with `eval(...)` at `boot.jl`. When calling `stacktrace()` from the REPL you'll also have a few extra frames in the stack from `REPL.jl`, usually looking something like this:

```
julia> example() = stacktrace()
example (generic function with 1 method)

julia> example()
5-element Array{StackFrame,1}:
  in example() at REPL[1]:1
  in eval(::Module, ::Any) at boot.jl:234
  in eval_user_input(::Any, ::Base.REPL.REPLBackend) at REPL.jl:62
  in macro expansion at REPL.jl:92 [inlined]
  in (::Base.REPL.##1#2{Base.REPL.REPLBackend}) () at event.jl:46
```

33.2 Extracting useful information

Each `StackFrame` contains the function name, file name, line number, lambda info, a flag indicating whether the frame has been inlined, a flag indicating whether it is a C function (by default C functions do not appear in the stack trace), and an integer representation of the pointer returned by `backtrace()` (page 374):

```
julia> top_frame = stacktrace()[1]
  in eval(::Module, ::Any) at boot.jl:236

julia> top_frame.func
:eval

julia> top_frame.file
Symbol("./boot.jl")

julia> top_frame.line
236

julia> top_frame.lininfo
Nullable{Core.MethodInstance}(MethodInstance for eval(::Module, ::Any))

julia> top_frame.inlined
false

julia> top_frame.from_c
false
```

```
julia> top_frame.pointer
13203085684
```

This makes stack trace information available programmatically for logging, error handling, and more.

33.3 Error handling

While having easy access to information about the current state of the callstack can be helpful in many places, the most obvious application is in error handling and debugging.

```
julia> @noinline bad_function() = undeclared_variable
bad_function (generic function with 1 method)
```

```
julia> @noinline example() = try
    bad_function()
    catch
        stacktrace()
    end
example (generic function with 1 method)

julia> example()
11-element Array{StackFrame,1}:
 in example() at none:4
 in eval(::Module, ::Any) at boot.jl:236
 ...
```

You may notice that in the example above the first stack frame points points at line 4, where `stacktrace()` is called, rather than line 2, where `bad_function` is called, and `bad_function`'s frame is missing entirely. This is understandable, given that `stacktrace()` is called from the context of the `catch`. While in this example it's fairly easy to find the actual source of the error, in complex cases tracking down the source of the error becomes nontrivial.

This can be remedied by calling `catch_stacktrace()` instead of `stacktrace()`. Instead of returning callstack information for the current context, `catch_stacktrace()` returns stack information for the context of the most recent exception:

```
julia> @noinline bad_function() = undeclared_variable
bad_function (generic function with 1 method)

julia> @noinline example() = try
    bad_function()
    catch
        catch_stacktrace()
    end
example (generic function with 1 method)

julia> example()
12-element Array{StackFrame,1}:
 in bad_function() at none:1
 in example() at none:2
 ...
```

Notice that the stack trace now indicates the appropriate line number and the missing frame.

```
julia> @noinline child() = error("Whoops!")
child (generic function with 1 method)

julia> @noinline parent() = child()
parent (generic function with 1 method)

julia> @noinline function grandparent()
    try
        parent()
    catch err
        println("ERROR: ", err.msg)
        catch_stacktrace()
    end
end
grandparent (generic function with 1 method)

julia> grandparent()
ERROR: Whoops!
```

```
13-element Array{StackFrame,1}:
  in child() at none:1
  in parent() at none:1
  in grandparent() at none:3
  ...
```

33.4 Comparison with `backtrace()`

A call to `backtrace()` (page 374) returns a vector of `Ptr{Void}`, which may then be passed into `stacktrace()` for translation:

```
julia> trace = backtrace()
20-element Array{Ptr{Void},1}:
 Ptr{Void} @0x0000000100a26fc2
 Ptr{Void} @0x00000001029435df
 Ptr{Void} @0x0000000102943635
 Ptr{Void} @0x00000001009e9620
 Ptr{Void} @0x00000001009fe1e8
 Ptr{Void} @0x00000001009fc7b6
 Ptr{Void} @0x00000001009fdae3
 Ptr{Void} @0x00000001009fe0d2
 Ptr{Void} @0x0000000100a1321b
 Ptr{Void} @0x00000001009f64e7
 Ptr{Void} @0x000000010265ac5d
 Ptr{Void} @0x000000010265acc1
 Ptr{Void} @0x00000001009e9620
 Ptr{Void} @0x000000031007744b
 Ptr{Void} @0x0000000310077537
 Ptr{Void} @0x00000001009e9620
 Ptr{Void} @0x000000031006feec
 Ptr{Void} @0x00000003100701b0
 Ptr{Void} @0x00000001009e9635
 Ptr{Void} @0x0000000100a06418

julia> stacktrace(trace)
5-element Array{StackFrame,1}:
  in backtrace() at error.jl:26
  in eval(::Module, ::Any) at boot.jl:231
  in eval_user_input(::Any, ::Base.REPL.REPLBackend) at REPL.jl:62
  in macro expansion at REPL.jl:92 [inlined]
  in (::Base.REPL.##1#2{Base.REPL.REPLBackend}) () at event.jl:46
```

Notice that the vector returned by `backtrace()` (page 374) had 15 pointers, while the vector returned by `stacktrace()` only has 4. This is because, by default, `stacktrace()` removes any lower-level C functions from the stack. If you want to include stack frames from C calls, you can do it like this:

```
julia> stacktrace(trace, true)
26-element Array{StackFrame,1}:
  in jl_backtrace_from_here at stackwalk.c:104
  in backtrace() at error.jl:26
  in ip:0x102943635
  in jl_call_method_internal at julia_internal.h:86 [inlined]
  in jl_apply_generic at gf.c:1805
  in do_call at interpreter.c:65
  in eval at interpreter.c:188
  in eval_body at interpreter.c:469
```

```

in jl_interpret_call at interpreter.c:573
in jl_toplevel_eval_flex at toplevel.c:543
in jl_toplevel_eval_in_warn at builtins.c:571
in eval(::Module, ::Any) at boot.jl:231
in ip:0x10265acc1
in jl_call_method_internal at julia_internal.h:86 [inlined]
in jl_apply_generic at gf.c:1805
in eval_user_input(::Any, ::Base.REPL.REPLBackend) at REPL.jl:62
in ip:0x310077537
in jl_call_method_internal at julia_internal.h:86 [inlined]
in jl_apply_generic at gf.c:1805
in macro expansion at REPL.jl:92 [inlined]
in (::Base.REPL.##1#2{Base.REPL.REPLBackend})() at event.jl:46
in ip:0x3100701b0
in jl_call_method_internal at julia_internal.h:86 [inlined]
in jl_apply_generic at gf.c:1795
in jl_apply at julia.h:1388 [inlined]
in start_task at task.c:247

```

Individual pointers returned by `backtrace()` (page 374) can be translated into `StackFrame`s by passing them into `StackTraces.lookup()`:

```

julia> pointer = backtrace()[1];

julia> frame = StackTraces.lookup(pointer)
1-element Array{StackFrame,1}:
  in jl_backtrace_from_here at stackwalk.c:105

julia> println("The top frame is from $(frame[1].func)!")
The top frame is from jl_backtrace_from_here!

```

Performance Tips

In the following sections, we briefly go through a few techniques that can help make your Julia code run as fast as possible.

34.1 Avoid global variables

A global variable might have its value, and therefore its type, change at any point. This makes it difficult for the compiler to optimize code using global variables. Variables should be local, or passed as arguments to functions, whenever possible.

Any code that is performance critical or being benchmarked should be inside a function.

We find that global names are frequently constants, and declaring them as such greatly improves performance:

```
const DEFAULT_VAL = 0
```

Uses of non-constant globals can be optimized by annotating their types at the point of use:

```
global x  
y = f(x::Int + 1)
```

Writing functions is better style. It leads to more reusable code and clarifies what steps are being done, and what their inputs and outputs are.

NOTE: All code in the REPL is evaluated in global scope, so a variable defined and assigned at toplevel will be a **global** variable.

In the following REPL session:

```
julia> x = 1.0
```

is equivalent to:

```
julia> global x = 1.0
```

so all the performance issues discussed previously apply.

34.2 Measure performance with `@time` and pay attention to memory allocation

The most useful tool for measuring performance is the `@time` (page 373) macro. The following example illustrates good working style:

```
julia> function f(n)
    s = 0
    for i = 1:n
        s += i/2
    end
    s
end
f (generic function with 1 method)

julia> @time f(1)
elapsed time: 0.004710563 seconds (93504 bytes allocated)
0.5

julia> @time f(10^6)
elapsed time: 0.04123202 seconds (32002136 bytes allocated)
2.5000025e11
```

On the first call (`@time f(1)`), `f` gets compiled. (If you’ve not yet used `@time` (page 373) in this session, it will also compile functions needed for timing.) You should not take the results of this run seriously. For the second run, note that in addition to reporting the time, it also indicated that a large amount of memory was allocated. This is the single biggest advantage of `@time` (page 373) vs. functions like `tict()` (page 372) and `toc()` (page 373), which only report time.

Unexpected memory allocation is almost always a sign of some problem with your code, usually a problem with type-stability. Consequently, in addition to the allocation itself, it’s very likely that the code generated for your function is far from optimal. Take such indications seriously and follow the advice below.

As a teaser, note that an improved version of this function allocates no memory (except to pass back the result back to the REPL) and has an order of magnitude faster execution after the first call:

```
julia> @time f_improved(1)    # first call
elapsed time: 0.003702172 seconds (78944 bytes allocated)
0.5

julia> @time f_improved(10^6)
elapsed time: 0.004313644 seconds (112 bytes allocated)
2.5000025e11
```

Below you’ll learn how to spot the problem with `f` and how to fix it.

In some situations, your function may need to allocate memory as part of its operation, and this can complicate the simple picture above. In such cases, consider using one of the [tools](#) (page 306) below to diagnose problems, or write a version of your function that separates allocation from its algorithmic aspects (see [Pre-allocating outputs](#) (page 317)).

34.3 Tools

Julia and its package ecosystem includes tools that may help you diagnose problems and improve the performance of your code:

- *Profiling* (page 617) allows you to measure the performance of your running code and identify lines that serve as bottlenecks. For complex projects, the *ProfileView* package can help you visualize your profiling results.
- Unexpectedly-large memory allocations—as reported by *@time* (page 373), *@allocated* (page 373), or the profiler (through calls to the garbage-collection routines)—hint that there might be issues with your code. If you don't see another reason for the allocations, suspect a type problem. You can also start Julia with the `--track-allocation=user` option and examine the resulting `*.mem` files to see information about where those allocations occur. See *Memory allocation analysis* (page 297).
- *@code_warntype* generates a representation of your code that can be helpful in finding expressions that result in type uncertainty. See *@code_warntype* (page 323) below.
- The *Lint* and *TypeCheck* packages can also warn you of certain types of programming errors.

34.4 Avoid containers with abstract type parameters

When working with parameterized types, including arrays, it is best to avoid parameterizing with abstract types where possible.

Consider the following:

```
a = Real[]      # typeof(a) = Array{Real,1}
if (f = rand()) < .8
    push!(a, f)
end
```

Because `a` is an array of abstract type `Real`, it must be able to hold any `Real` value. Since `Real` objects can be of arbitrary size and structure, `a` must be represented as an array of pointers to individually allocated `Real` objects. Because `f` will always be a `Float64` (page 437), we should instead, use:

```
a = Float64[] # typeof(a) = Array{Float64,1}
```

which will create a contiguous block of 64-bit floating-point values that can be manipulated efficiently.

See also the discussion under *Parametric Types* (page 96).

34.5 Type declarations

In many languages with optional type declarations, adding declarations is the principal way to make code run faster. This is *not* the case in Julia. In Julia, the compiler generally knows the types of all function arguments, local variables, and expressions. However, there are a few specific instances where declarations are helpful.

34.5.1 Avoid fields with abstract type

Types can be declared without specifying the types of their fields:

```
julia> type MyAmbiguousType
    a
end
```

This allows `a` to be of any type. This can often be useful, but it does have a downside: for objects of type `MyAmbiguousType`, the compiler will not be able to generate high-performance code. The reason is that the compiler uses the types of objects, not their values, to determine how to build code. Unfortunately, very little can be inferred about an object of type `MyAmbiguousType`:

```
julia> b = MyAmbiguousType("Hello")
MyAmbiguousType("Hello")

julia> c = MyAmbiguousType(17)
MyAmbiguousType(17)

julia> typeof(b)
MyAmbiguousType

julia> typeof(c)
MyAmbiguousType
```

`b` and `c` have the same type, yet their underlying representation of data in memory is very different. Even if you stored just numeric values in field `a`, the fact that the memory representation of a `UInt8` differs from a `Float64` also means that the CPU needs to handle them using two different kinds of instructions. Since the required information is not available in the type, such decisions have to be made at run-time. This slows performance.

You can do better by declaring the type of `a`. Here, we are focused on the case where `a` might be any one of several types, in which case the natural solution is to use parameters. For example:

```
julia> type MyType{T<:AbstractFloat}
    a::T
end
```

This is a better choice than

```
julia> type MyStillAmbiguousType
    a::AbstractFloat
end
```

because the first version specifies the type of `a` from the type of the wrapper object. For example:

```
julia> m = MyType(3.2)
MyType{Float64}(3.2)

julia> t = MyStillAmbiguousType(3.2)
MyStillAmbiguousType(3.2)

julia> typeof(m)
MyType{Float64}

julia> typeof(t)
MyStillAmbiguousType
```

The type of field `a` can be readily determined from the type of `m`, but not from the type of `t`. Indeed, in `t` it's possible to change the type of field `a`:

```
julia> typeof(t.a)
Float64

julia> t.a = 4.5f0
4.5f0

julia> typeof(t.a)
Float32
```

In contrast, once `m` is constructed, the type of `m.a` cannot change:

```
julia> m.a = 4.5f0
4.5f0
```

```
julia> typeof(m.a)
Float64
```

The fact that the type of `m.a` is known from `m`'s type—coupled with the fact that its type cannot change mid-function—allows the compiler to generate highly-optimized code for objects like `m` but not for objects like `t`.

Of course, all of this is true only if we construct `m` with a concrete type. We can break this by explicitly constructing it with an abstract type:

```
julia> m = MyType{AbstractFloat}(3.2)
MyType{AbstractFloat}(3.2)

julia> typeof(m.a)
Float64

julia> m.a = 4.5f0
4.5f0

julia> typeof(m.a)
Float32
```

For all practical purposes, such objects behave identically to those of `MyStillAmbiguousType`.

It's quite instructive to compare the sheer amount code generated for a simple function

```
func(m::MyType) = m.a+1
```

using

```
code_llvm(func, (MyType{Float64},))
code_llvm(func, (MyType{AbstractFloat},))
code_llvm(func, (MyType,))
```

For reasons of length the results are not shown here, but you may wish to try this yourself. Because the type is fully-specified in the first case, the compiler doesn't need to generate any code to resolve the type at run-time. This results in shorter and faster code.

34.5.2 Avoid fields with abstract containers

The same best practices also work for container types:

```
julia> type MySimpleContainer{A<:AbstractVector}
    a::A
end

julia> type MyAmbiguousContainer{T}
    a::AbstractVector{T}
end
```

For example:

```
julia> c = MySimpleContainer(1:3);

julia> typeof(c)
MySimpleContainer{UnitRange{Int64}}

julia> c = MySimpleContainer([1:3;]);

julia> typeof(c)
```

```
MySimpleContainer{Array{Int64,1}}

julia> b = MyAmbiguousContainer(1:3);

julia> typeof(b)
MyAmbiguousContainer{Int64}

julia> b = MyAmbiguousContainer([1:3;]);

julia> typeof(b)
MyAmbiguousContainer{Int64}
```

For `MySimpleContainer`, the object is fully-specified by its type and parameters, so the compiler can generate optimized functions. In most instances, this will probably suffice.

While the compiler can now do its job perfectly well, there are cases where *you* might wish that your code could do different things depending on the *element type* of `a`. Usually the best way to achieve this is to wrap your specific operation (here, `foo`) in a separate function:

```
function sumfoo(c::MySimpleContainer)
    s = 0
    for x in c.a
        s += foo(x)
    end
    s
end

foo(x::Integer) = x
foo(x::AbstractFloat) = round(x)
```

This keeps things simple, while allowing the compiler to generate optimized code in all cases.

However, there are cases where you may need to declare different versions of the outer function for different element types of `a`. You could do it like this:

```
function myfun{T<:AbstractFloat}(c::MySimpleContainer{Vector{T}})
    ...
end
function myfun{T<:Integer}(c::MySimpleContainer{Vector{T}})
    ...
end
```

This works fine for `Vector{T}`, but we'd also have to write explicit versions for `UnitRange{T}` or other abstract types. To prevent such tedium, you can use two parameters in the declaration of `MyContainer`:

```
type MyContainer{T, A<:AbstractVector}
    a::A
end
MyContainer(v::AbstractVector) = MyContainer{eltype(v), typeof(v)}(v)

julia> b = MyContainer(1.3:5);

julia> typeof(b)
MyContainer{Float64, UnitRange{Float64}}
```

Note the somewhat surprising fact that `T` doesn't appear in the declaration of field `a`, a point that we'll return to in a moment. With this approach, one can write functions such as:

```

function myfunc{T<:Integer, A<:AbstractArray}(c::MyContainer{T,A})
    return c.a[1]+1
end
# Note: because we can only define MyContainer for
# A<:AbstractArray, and any unspecified parameters are arbitrary,
# the previous could have been written more succinctly as
#     function myfunc{T<:Integer}(c::MyContainer{T})

function myfunc{T<:AbstractFloat}(c::MyContainer{T})
    return c.a[1]+2
end

function myfunc{T<:Integer}(c::MyContainer{T,Vector{T}})
    return c.a[1]+3
end

julia> myfunc(MyContainer{1:3})
2

julia> myfunc(MyContainer{1.0:3})
3.0

julia> myfunc(MyContainer{[1:3]})
4

```

As you can see, with this approach it's possible to specialize on both the element type `T` and the array type `A`.

However, there's one remaining hole: we haven't enforced that `A` has element type `T`, so it's perfectly possible to construct an object like this:

```

julia> b = MyContainer{Int64, UnitRange{Float64}}(1.3:5);

julia> typeof(b)
MyContainer{Int64,UnitRange{Float64}}

```

To prevent this, we can add an inner constructor:

```

type MyBetterContainer{T<:Real, A<:AbstractVector}
    a::A

    MyBetterContainer(v::AbstractVector{T}) = new(v)
end
MyBetterContainer(v::AbstractVector) = MyBetterContainer{eltype(v),typeof(v)}(v)

julia> b = MyBetterContainer(1.3:5);

julia> typeof(b)
MyBetterContainer{Float64,UnitRange{Float64}}

julia> b = MyBetterContainer{Int64, UnitRange{Float64}}(1.3:5);
ERROR: no method MyBetterContainer{UnitRange{Float64},}

```

The inner constructor requires that the element type of `A` be `T`.

34.5.3 Annotate values taken from untyped locations

It is often convenient to work with data structures that may contain values of any type (arrays of type `Array{Any}`). But, if you're using one of these structures and happen to know the type of an element, it helps to share this knowledge with the compiler:

```
function foo(a::Array{Any,1})
    x = a[1]::Int32
    b = x+1
    ...
end
```

Here, we happened to know that the first element of `a` would be an `Int32`. Making an annotation like this has the added benefit that it will raise a run-time error if the value is not of the expected type, potentially catching certain bugs earlier.

34.5.4 Declare types of keyword arguments

Keyword arguments can have declared types:

```
function with_keyword(x; name::Int = 1)
    ...
end
```

Functions are specialized on the types of keyword arguments, so these declarations will not affect performance of code inside the function. However, they will reduce the overhead of calls to the function that include keyword arguments.

Functions with keyword arguments have near-zero overhead for call sites that pass only positional arguments.

Passing dynamic lists of keyword arguments, as in `f(x; keywords...)`, can be slow and should be avoided in performance-sensitive code.

34.6 Break functions into multiple definitions

Writing a function as many small definitions allows the compiler to directly call the most applicable code, or even inline it.

Here is an example of a “compound function” that should really be written as multiple definitions:

```
function norm(A)
    if isa(A, Vector)
        return sqrt(real(dot(A,A)))
    elseif isa(A, Matrix)
        return max(svd(A)[2])
    else
        error("norm: invalid argument")
    end
end
```

This can be written more concisely and efficiently as:

```
norm(x::Vector) = sqrt(real(dot(x,x)))
norm(A::Matrix) = max(svd(A)[2])
```

34.7 Write “type-stable” functions

When possible, it helps to ensure that a function always returns a value of the same type. Consider the following definition:

```
pos(x) = x < 0 ? 0 : x
```

Although this seems innocent enough, the problem is that 0 is an integer (of type `Int`) and `x` might be of any type. Thus, depending on the value of `x`, this function might return a value of either of two types. This behavior is allowed, and may be desirable in some cases. But it can easily be fixed as follows:

```
pos(x) = x < 0 ? zero(x) : x
```

There is also a `one()` (page 435) function, and a more general `oftype(x, y)` (page 365) function, which returns `y` converted to the type of `x`.

34.8 Avoid changing the type of a variable

An analogous “type-stability” problem exists for variables used repeatedly within a function:

```
function foo()
    x = 1
    for i = 1:10
        x = x/bar()
    end
    return x
end
```

Local variable `x` starts as an integer, and after one loop iteration becomes a floating-point number (the result of `/` (page 405) operator). This makes it more difficult for the compiler to optimize the body of the loop. There are several possible fixes:

- Initialize `x` with `x = 1.0`
- Declare the type of `x`: `x::Float64 = 1`
- Use an explicit conversion: `x = one{T}`

34.9 Separate kernel functions (aka, function barriers)

Many functions follow a pattern of performing some set-up work, and then running many iterations to perform a core computation. Where possible, it is a good idea to put these core computations in separate functions. For example, the following contrived function returns an array of a randomly-chosen type:

```
function strange_twos(n)
    a = Array{rand{Bool} ? Int64 : Float64, n}
    for i = 1:n
        a[i] = 2
    end
    return a
end
```

This should be written as:

```
function fill_twos!(a)
    for i=1:length(a)
        a[i] = 2
    end
end

function strange_twos(n)
    a = Array{rand{Bool} ? Int64 : Float64, n}
    fill_twos!(a)
    return a
end
```

Julia’s compiler specializes code for argument types at function boundaries, so in the original implementation it does not know the type of `a` during the loop (since it is chosen randomly). Therefore the second version is generally faster since the inner loop can be recompiled as part of `fill_twos!` for different types of `a`.

The second form is also often better style and can lead to more code reuse.

This pattern is used in several places in the standard library. For example, see `hvcats_fill` in `abstractarray.jl`, or the `fill!` function, which we could have used instead of writing our own `fill_twos!`.

Functions like `strange_twos` occur when dealing with data of uncertain type, for example data loaded from an input file that might contain either integers, floats, strings, or something else.

34.10 Types with values-as-parameters

Let’s say you want to create an N -dimensional array that has size 3 along each axis. Such arrays can be created like this:

```
A = fill{5.0, (3, 3)}
```

This approach works very well: the compiler can figure out that `A` is an `Array{Float64, 2}` because it knows the type of the fill value (`5.0::Float64`) and the dimensionality (`(3, 3)::NTuple{2, Int}`). This implies that the compiler can generate very efficient code for any future usage of `A` in the same function.

But now let’s say you want to write a function that creates a $3 \times 3 \times \dots$ array in arbitrary dimensions; you might be tempted to write a function

```
function array3(fillval, N)
    fill(fillval, ntuple(d->3, N))
end
```

This works, but (as you can verify for yourself using `@code_warntype array3(5.0, 2)`) the problem is that the output type cannot be inferred: the argument `N` is a *value* of type `Int`, and type-inference does not (and cannot) predict its value in advance. This means that code using the output of this function has to be conservative, checking the type on each access of `A`; such code will be very slow.

Now, one very good way to solve such problems is by using the *function-barrier technique* (page 313). However, in some cases you might want to eliminate the type-instability altogether. In such cases, one approach is to pass the dimensionality as a parameter, for example through `Val{T}` (see “*Value types*” (page 106)):

```
function array3{N}(fillval, ::Type{Val{N}})
    fill(fillval, ntuple(d->3, Val{N}))
end
```

Julia has a specialized version of `ntuple` that accepts a `Val{::Int}` as the second parameter; by passing `N` as a type-parameter, you make its “value” known to the compiler. Consequently, this version of `array3` allows the compiler to predict the return type.

However, making use of such techniques can be surprisingly subtle. For example, it would be of no help if you called `array3` from a function like this:

```
function call_array3(fillval, n)
    A = array3(fillval, Val{n})
end
```

Here, you’ve created the same problem all over again: the compiler can’t guess the type of `n`, so it doesn’t know the type of `Val{n}`. Attempting to use `Val`, but doing so incorrectly, can easily make performance *worse* in many situations. (Only in situations where you’re effectively combining `Val` with the function-barrier trick, to make the kernel function more efficient, should code like the above be used.)

An example of correct usage of `Val` would be:

```
function filter3{T,N}(A::AbstractArray{T,N})
    kernel = array3(1, Val{N})
    filter(A, kernel)
end
```

In this example, `N` is passed as a parameter, so its “value” is known to the compiler. Essentially, `Val{T}` works only when `T` is either hard-coded (`Val{3}`) or already specified in the type-domain.

34.11 The dangers of abusing multiple dispatch (aka, more on types with values-as-parameters)

Once one learns to appreciate multiple dispatch, there’s an understandable tendency to go crazy and try to use it for everything. For example, you might imagine using it to store information, e.g.

```
immutable Car{Make,Model}
    year::Int
    ...more fields...
end
```

and then dispatch on objects like `Car{:Honda,:Accord}(year, args...)`.

This might be worthwhile when the following are true:

- You require CPU-intensive processing on each `Car`, and it becomes vastly more efficient if you know the `Make` and `Model` at compile time.
- You have homogenous lists of the same type of `Car` to process, so that you can store them all in an `Array{Car{:Honda,:Accord},N}`.

When the latter holds, a function processing such a homogenous array can be productively specialized: Julia knows the type of each element in advance (all objects in the container have the same concrete type), so Julia can “look up” the correct method calls when the function is being compiled (obviating the need to check at run-time) and thereby emit efficient code for processing the whole list.

When these do not hold, then it’s likely that you’ll get no benefit; worse, the resulting “combinatorial explosion of types” will be counterproductive. If `items[i+1]` has a different type than `item[i]`, Julia has to look up the type at run-time, search for the appropriate method in method tables, decide (via type intersection) which one matches, determine whether it has been JIT-compiled yet (and do so if not), and then make the call. In essence, you’re asking the full type- system and JIT-compilation machinery to basically execute the equivalent of a switch statement or dictionary lookup in your own code.

Some run-time benchmarks comparing (1) type dispatch, (2) dictionary lookup, and (3) a “switch” statement can be found [on the mailing list](#).

Perhaps even worse than the run-time impact is the compile-time impact: Julia will compile specialized functions for each different `Car{Make, Model}`; if you have hundreds or thousands of such types, then every function that accepts such an object as a parameter (from a custom `get_year` function you might write yourself, to the generic `push!` function in the standard library) will have hundreds or thousands of variants compiled for it. Each of these increases the size of the cache of compiled code, the length of internal lists of methods, etc. Excess enthusiasm for values-as-parameters can easily waste enormous resources.

34.12 Access arrays in memory order, along columns

Multidimensional arrays in Julia are stored in column-major order. This means that arrays are stacked one column at a time. This can be verified using the `vec` function or the syntax `[:]` as shown below (notice that the array is ordered `[1 3 2 4]`, not `[1 2 3 4]`):

```
julia> x = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> x[:]
4-element Array{Int64,1}:
 1
 3
 2
 4
```

This convention for ordering arrays is common in many languages like Fortran, Matlab, and R (to name a few). The alternative to column-major ordering is row-major ordering, which is the convention adopted by C and Python (numpy) among other languages. Remembering the ordering of arrays can have significant performance effects when looping over arrays. A rule of thumb to keep in mind is that with column-major arrays, the first index changes most rapidly. Essentially this means that looping will be faster if the inner-most loop index is the first to appear in a slice expression.

Consider the following contrived example. Imagine we wanted to write a function that accepts a `Vector` and returns a square `Matrix` with either the rows or the columns filled with copies of the input vector. Assume that it is not important whether rows or columns are filled with these copies (perhaps the rest of the code can be easily adapted accordingly). We could conceivably do this in at least four ways (in addition to the recommended call to the built-in `repmat()` (page 529)):

```
function copy_cols{T}(x::Vector{T})
    n = size(x, 1)
    out = Array{T}(n, n)
    for i=1:n
        out[:, i] = x
    end
    out
end

function copy_rows{T}(x::Vector{T})
    n = size(x, 1)
    out = Array{T}(n, n)
    for i=1:n
        out[i, :] = x
    end
    out
end
```

```

function copy_col_row{T}(x::Vector{T})
    n = size(x, 1)
    out = Array{T}(n, n)
    for col=1:n, row=1:n
        out[row, col] = x[row]
    end
    out
end

function copy_row_col{T}(x::Vector{T})
    n = size(x, 1)
    out = Array{T}(n, n)
    for row=1:n, col=1:n
        out[row, col] = x[col]
    end
    out
end

```

Now we will time each of these functions using the same random 10000 by 1 input vector:

```

julia> x = randn(10000);

julia> fmt(f) = println(rpad(string(f)*": ", 14, ' '), @elapsed f(x))

julia> map(fmt, Any[copy_cols, copy_rows, copy_col_row, copy_row_col]);
copy_cols:      0.331706323
copy_rows:      1.799009911
copy_col_row:   0.415630047
copy_row_col:   1.721531501

```

Notice that `copy_cols` is much faster than `copy_rows`. This is expected because `copy_cols` respects the column-based memory layout of the `Matrix` and fills it one column at a time. Additionally, `copy_col_row` is much faster than `copy_row_col` because it follows our rule of thumb that the first element to appear in a slice expression should be coupled with the inner-most loop.

34.13 Pre-allocating outputs

If your function returns an `Array` or some other complex type, it may have to allocate memory. Unfortunately, often-times allocation and its converse, garbage collection, are substantial bottlenecks.

Sometimes you can circumvent the need to allocate memory on each function call by preallocating the output. As a trivial example, compare

```

function xinc(x)
    return [x, x+1, x+2]
end

function loopinc()
    y = 0
    for i = 1:10^7
        ret = xinc(i)
        y += ret[2]
    end
    y
end

```

with

```
function xinc!{T}(ret::AbstractVector{T}, x::T)
    ret[1] = x
    ret[2] = x+1
    ret[3] = x+2
    nothing
end

function loopinc_prealloc()
    ret = Array{Int}(3)
    y = 0
    for i = 1:10^7
        xinc!(ret, i)
        y += ret[2]
    end
    y
end
```

Timing results:

```
julia> @time loopinc()
elapsed time: 1.955026528 seconds (1279975584 bytes allocated)
50000015000000

julia> @time loopinc_prealloc()
elapsed time: 0.078639163 seconds (144 bytes allocated)
50000015000000
```

Preallocation has other advantages, for example by allowing the caller to control the “output” type from an algorithm. In the example above, we could have passed a `SubArray` rather than an `Array` (page 454), had we so desired.

Taken to its extreme, pre-allocation can make your code uglier, so performance measurements and some judgment may be required. However, for “vectorized” (element-wise) functions, the convenient syntax `x .= f.(y)` can be used for in-place operations with fused loops and no temporary arrays (*Dot Syntax for Vectorizing Functions* (page 63)).

34.14 Avoid string interpolation for I/O

When writing data to a file (or other I/O device), forming extra intermediate strings is a source of overhead. Instead of:

```
println(file, "$a $b")
```

use:

```
println(file, a, " ", b)
```

The first version of the code forms a string, then writes it to the file, while the second version writes values directly to the file. Also notice that in some cases string interpolation can be harder to read. Consider:

```
println(file, "$(f(a))$(f(b))")
```

versus:

```
println(file, f(a), f(b))
```

34.15 Optimize network I/O during parallel execution

When executing a remote function in parallel:

```
responses = Vector{Any}(nworkers())
@sync begin
    for (idx, pid) in enumerate(workers())
        @async responses[idx] = remotecall_fetch(pid, foo, args...)
    end
end
```

is faster than:

```
refs = Vector{Any}(nworkers())
for (idx, pid) in enumerate(workers())
    refs[idx] = @spawnat pid foo(args...)
end
responses = [fetch(r) for r in refs]
```

The former results in a single network round-trip to every worker, while the latter results in two network calls - first by the `@spawnat` and the second due to the `fetch` (or even a `wait`). The `fetch/wait` is also being executed serially resulting in an overall poorer performance.

34.16 Fix deprecation warnings

A deprecated function internally performs a lookup in order to print a relevant warning only once. This extra lookup can cause a significant slowdown, so all uses of deprecated functions should be modified as suggested by the warnings.

34.17 Tweaks

These are some minor points that might help in tight inner loops.

- Avoid unnecessary arrays. For example, instead of `sum([x, y, z])` (page 388) use `x+y+z`.
- Use `abs2(z)` (page 419) instead of `abs(z)^2` (page 418) for complex `z`. In general, try to rewrite code to use `abs2()` (page 419) instead of `abs()` (page 418) for complex arguments.
- Use `div(x, y)` (page 407) for truncating division of integers instead of `trunc(x/y)` (page 417), `fld(x, y)` (page 407) instead of `floor(x/y)` (page 417), and `cld(x, y)` (page 407) instead of `ceil(x/y)` (page 417).

34.18 Performance Annotations

Sometimes you can enable better optimization by promising certain program properties.

- Use `@inbounds` to eliminate array bounds checking within expressions. Be certain before doing this. If the subscripts are ever out of bounds, you may suffer crashes or silent corruption.
- Use `@fastmath` to allow floating point optimizations that are correct for real numbers, but lead to differences for IEEE numbers. Be careful when doing this, as this may change numerical results. This corresponds to the `-ffast-math` option of clang.

- Write `@simd` in front of `for` loops that are amenable to vectorization. **This feature is experimental** and could change or disappear in future versions of Julia.

Note: While `@simd` needs to be placed directly in front of a loop, both `@inbounds` and `@fastmath` can be applied to several statements at once, e.g. using `begin ... end`, or even to a whole function.

Here is an example with both `@inbounds` and `@simd` markup:

```
function inner( x, y )
    s = zero(eltype(x))
    for i=1:length(x)
        @inbounds s += x[i]*y[i]
    end
    s
end

function innersimd( x, y )
    s = zero(eltype(x))
    @simd for i=1:length(x)
        @inbounds s += x[i]*y[i]
    end
    s
end

function timeit( n, reps )
    x = rand(Float32,n)
    y = rand(Float32,n)
    s = zero(Float64)
    time = @elapsed for j in 1:reps
        s+=inner(x,y)
    end
    println("GFlop/sec          = ",2.0*n*reps/time*1E-9)
    time = @elapsed for j in 1:reps
        s+=innersimd(x,y)
    end
    println("GFlop/sec (SIMD) = ",2.0*n*reps/time*1E-9)
end

timeit(1000,1000)
```

On a computer with a 2.4GHz Intel Core i5 processor, this produces:

```
GFlop/sec          = 1.9467069505224963
GFlop/sec (SIMD) = 17.578554163920018
```

(GFlop/sec measures the performance, and larger numbers are better.) The range for a `@simd` `for` loop should be a one-dimensional range. A variable used for accumulating, such as `s` in the example, is called a *reduction variable*. By using `@simd`, you are asserting several properties of the loop:

- It is safe to execute iterations in arbitrary or overlapping order, with special consideration for reduction variables.
- Floating-point operations on reduction variables can be reordered, possibly causing different results than without `@simd`.
- No iteration ever waits on another iteration to make forward progress.

A loop containing `break`, `continue`, or `@goto` will cause a compile-time error.

Using `@simd` merely gives the compiler license to vectorize. Whether it actually does so depends on the compiler. To actually benefit from the current implementation, your loop should have the following additional properties:

- The loop must be an innermost loop.

- The loop body must be straight-line code. This is why `@inbounds` is currently needed for all array accesses. The compiler can sometimes turn short `&&`, `||`, and `?:` expressions into straight-line code, if it is safe to evaluate all operands unconditionally. Consider using `ifelse()` (page 362) instead of `?:` in the loop if it is safe to do so.
- Accesses must have a stride pattern and cannot be “gathers” (random-index reads) or “scatters” (random-index writes).
- The stride should be unit stride.
- In some simple cases, for example with 2-3 arrays accessed in a loop, the LLVM auto-vectorization may kick in automatically, leading to no further speedup with `@simd`.

Here is an example with all three kinds of markup. This program first calculates the finite difference of a one-dimensional array, and then evaluates the L2-norm of the result:

```
function init!(u)
    n = length(u)
    dx = 1.0 / (n-1)
    @fastmath @inbounds @simd for i in 1:n
        u[i] = sin(2pi*dx*i)
    end
end

function deriv!(u, du)
    n = length(u)
    dx = 1.0 / (n-1)
    @fastmath @inbounds du[1] = (u[2] - u[1]) / dx
    @fastmath @inbounds @simd for i in 2:n-1
        du[i] = (u[i+1] - u[i-1]) / (2*dx)
    end
    @fastmath @inbounds du[n] = (u[n] - u[n-1]) / dx
end

function norm(u)
    n = length(u)
    T = eltype(u)
    s = zero(T)
    @fastmath @inbounds @simd for i in 1:n
        s += u[i]^2
    end
    @fastmath @inbounds return sqrt(s/n)
end

function main()
    n = 2000
    u = Array{Float64}(n)
    init!(u)
    du = similar(u)

    deriv!(u, du)
    nu = norm(du)

    @time for i in 1:10^6
        deriv!(u, du)
        nu = norm(du)
    end

    println(nu)
end
```

```
main()
```

On a computer with a 2.7 GHz Intel Core i7 processor, this produces:

```
$ julia wave.jl;
elapsed time: 1.207814709 seconds (0 bytes allocated)

$ julia --math-mode=ieee wave.jl;
elapsed time: 4.487083643 seconds (0 bytes allocated)
```

Here, the option `--math-mode=ieee` disables the `@fastmath` macro, so that we can compare results.

In this case, the speedup due to `@fastmath` is a factor of about 3.7. This is unusually large – in general, the speedup will be smaller. (In this particular example, the working set of the benchmark is small enough to fit into the L1 cache of the processor, so that memory access latency does not play a role, and computing time is dominated by CPU usage. In many real world programs this is not the case.) Also, in this case this optimization does not change the result – in general, the result will be slightly different. In some cases, especially for numerically unstable algorithms, the result can be very different.

The annotation `@fastmath` re-arranges floating point expressions, e.g. changing the order of evaluation, or assuming that certain special cases (inf, nan) cannot occur. In this case (and on this particular computer), the main difference is that the expression `1 / (2*dx)` in the function `deriv` is hoisted out of the loop (i.e. calculated outside the loop), as if one had written `idx = 1 / (2*dx)`. In the loop, the expression `... / (2*dx)` then becomes `... * idx`, which is much faster to evaluate. Of course, both the actual optimization that is applied by the compiler as well as the resulting speedup depend very much on the hardware. You can examine the change in generated code by using Julia's `code_native()` (page 379) function.

34.19 Treat Subnormal Numbers as Zeros

Subnormal numbers, formerly called *denormal numbers*, are useful in many contexts, but incur a performance penalty on some hardware. A call `set_zero_subnormals(true)` (page 438) grants permission for floating-point operations to treat subnormal inputs or outputs as zeros, which may improve performance on some hardware. A call `set_zero_subnormals(false)` (page 438) enforces strict IEEE behavior for subnormal numbers.

Below is an example where subnormals noticeably impact performance on some hardware:

```
function timestep{T}( b::Vector{T}, a::Vector{T}, Δt::T )
    @assert length(a)==length(b)
    n = length(b)
    b[1] = 1                                # Boundary condition
    for i=2:n-1
        b[i] = a[i] + (a[i-1] - T(2)*a[i] + a[i+1]) * Δt
    end
    b[n] = 0                                # Boundary condition
end

function heatflow{T}( a::Vector{T}, nstep::Integer )
    b = similar(a)
    for t=1:div(nstep,2)                    # Assume nstep is even
        timestep(b,a,T(0.1))
        timestep(a,b,T(0.1))
    end
end

heatflow(zeros{Float32,10},2)              # Force compilation
for trial=1:6
```



```

a = zeros(Float32,1000)
set_zero_subnormals(iseven(trial))  # Odd trials use strict IEEE arithmetic
@time heatflow(a,1000)
end

```

This example generates many subnormal numbers because the values in `a` become an exponentially decreasing curve, which slowly flattens out over time.

Treating subnormals as zeros should be used with caution, because doing so breaks some identities, such as `x-y==0` implies `x==y`:

```

julia> x=3f-38; y=2f-38;

julia> set_zero_subnormals(false); (x-y,x==y)
(1.0000001f-38,false)

julia> set_zero_subnormals(true); (x-y,x==y)
(0.0f0,false)

```

In some applications, an alternative to zeroing subnormal numbers is to inject a tiny bit of noise. For example, instead of initializing `a` with zeros, initialize it with:

```

a = rand(Float32,1000) * 1.f-9

```

34.20 @code_warntype

The macro `@code_warntype` (page 379) (or its function variant `code_warntype()` (page 379)) can sometimes be helpful in diagnosing type-related problems. Here's an example:

```

pos(x) = x < 0 ? 0 : x

function f(x)
    y = pos(x)
    sin(y*x+1)
end

julia> @code_warntype f(3.2)
Variables:
  x::Float64
  y::UNION{INT64,FLOAT64}
  _var0::Float64
  _var3::Tuple{Int64}
  _var4::UNION{INT64,FLOAT64}
  _var1::Float64
  _var2::Float64

Body:
begin # none, line 2:
    _var0 = (top(box)) (Float64, (top(sitofp)) (Float64,0))
    unless (top(box)) (Bool, (top(or_int)) ((top(lt_float)) (x::Float64,_var0::Float64))::Bool, (top(box)
    _var4 = 0
    goto 2
  1:
    _var4 = x::Float64
  2:
    y = _var4::UNION{INT64,FLOAT64} # line 3:
    _var1 = y::UNION{INT64,FLOAT64} * x::Float64::Float64

```

```
_var2 = (top(box)) (Float64, (top(add_float)) (_var1::Float64, (top(box)) (Float64, (top(sitofp)) (Fl  
return (GlobalRef(Base.Math,:nan_dom_err)) ((top(ccall)) ($(Expr(:call1, : (top(tuple)), "sin", G  
end::Float64
```

Interpreting the output of `@code_warntype` (page 379), like that of its cousins `@code_lowered` (page 378), `@code_typed` (page 379), `@code_llvm` (page 379), and `@code_native` (page 379), takes a little practice. Your code is being presented in form that has been partially digested on its way to generating compiled machine code. Most of the expressions are annotated by a type, indicated by the `::T` (where `T` might be `Float64` (page 437), for example). The most important characteristic of `@code_warntype` (page 379) is that non-concrete types are displayed in red; in the above example, such output is shown in all-caps.

The top part of the output summarizes the type information for the different variables internal to the function. You can see that `y`, one of the variables you created, is a `Union{Int64,Float64}`, due to the type-instability of `pos`. There is another variable, `_var4`, which you can see also has the same type.

The next lines represent the body of `f`. The lines starting with a number followed by a colon (`1 :`, `2 :`) are labels, and represent targets for jumps (via `goto`) in your code. Looking at the body, you can see that `pos` has been *inlined* into `f`—everything before `2 :` comes from code defined in `pos`.

Starting at 2:, the variable `y` is defined, and again annotated as a Union type. Next, we see that the compiler created the temporary variable `_var1` to hold the result of `y*x`. Because a `Float64` (page 437) times *either* an `Int64` or `Float64` (page 437) yields a `Float64` (page 437), all type-instability ends here. The net result is that `f(x::Float64)` will not be type-unstable in its output, even if some of the intermediate computations are type-unstable.

How you use this information is up to you. Obviously, it would be far and away best to fix `pos` to be type-stable: if you did so, all of the variables in `f` would be concrete, and its performance would be optimal. However, there are circumstances where this kind of *ephemeral* type instability might not matter too much: for example, if `pos` is never used in isolation, the fact that `f`'s output is type-stable (for [Float64](#) (page 437) inputs) will shield later code from the propagating effects of type instability. This is particularly relevant in cases where fixing the type instability is difficult or impossible: for example, currently it's not possible to infer the return type of an anonymous function. In such cases, the tips above (e.g., adding type annotations and/or breaking up functions) are your best tools to contain the “damage” from type instability.

The following examples may help you interpret expressions marked as containing non-leaf types:

- Function body ending in `end :: Union{T1, T2}`
 - Interpretation: function with unstable return type
 - Suggestion: make the return value type-stable, even if you have to annotate it
- `f(x :: T) :: Union{T1, T2}`
 - Interpretation: call to a type-unstable function
 - Suggestion: fix the function, or if necessary annotate the return value
- `(top(arrayref))(A :: Array{Any, 1}, 1) :: Any`
 - Interpretation: accessing elements of poorly-typed arrays
 - Suggestion: use arrays with better-defined types, or if necessary annotate the type of individual element accesses
- `(top(getfield))(A :: ArrayContainer{Float64}, :data) :: Array{Float64, N}`
 - Interpretation: getting a field that is of non-leaf type. In this case, `ArrayContainer` had a field `data :: Array{T}`. But `Array` needs the dimension `N`, too, to be a concrete type.
 - Suggestion: use concrete types like `Array{T, 3}` or `Array{T, N}`, where `N` is now a parameter of `ArrayContainer`

Workflow Tips

Here are some tips for working with Julia efficiently.

35.1 REPL-based workflow

As already elaborated in *Interacting With Julia* (page 257), Julia’s REPL provides rich functionality that facilitates an efficient interactive workflow. Here are some tips that might further enhance your experience at the command line.

35.1.1 A basic editor/REPL workflow

The most basic Julia workflows involve using a text editor in conjunction with the `julia` command line. A common pattern includes the following elements:

- **Put code under development in a temporary module.** Create a file, say `Tmp.jl`, and include within it

```
module Tmp

    <your definitions here>

end
```

- **Put your test code in another file.** Create another file, say `tst.jl`, which begins with

```
import Tmp
```

and includes tests for the contents of `Tmp`. The value of using `import` versus `using` is that you can call `reload("Tmp")` instead of having to restart the REPL when your definitions change. Of course, the cost is the need to prepend `Tmp.` to uses of names defined in your module. (You can lower that cost by keeping your module name short.)

Alternatively, you can wrap the contents of your test file in a module, as

```
module Tst
    using Tmp

    <scratch work>

end
```

The advantage is that you can now do `using Tmp` in your test code and can therefore avoid prepending `Tmp.` everywhere. The disadvantage is that code can no longer be selectively copied to the REPL without some tweaking.

- **Lather. Rinse. Repeat.** Explore ideas at the `julia` command prompt. Save good ideas in `tst.jl`. Occasionally restart the REPL, issuing

```
reload("Tmp")
include("tst.jl")
```

35.1.2 Simplify initialization

To simplify restarting the REPL, put project-specific initialization code in a file, say `_init.jl`, which you can run on startup by issuing the command:

```
julia -L _init.jl
```

If you further add the following to your `.juliarc.jl` file

```
isfile("_init.jl") && include(joinpath(pwd(), "_init.jl"))
```

then calling `julia` from that directory will run the initialization code without the additional command line argument.

35.2 Browser-based workflow

It is also possible to interact with a Julia REPL in the browser via [IJulia](#). See the package home for details.

Style Guide

The following sections explain a few aspects of idiomatic Julia coding style. None of these rules are absolute; they are only suggestions to help familiarize you with the language and to help you choose among alternative designs.

36.1 Write functions, not just scripts

Writing code as a series of steps at the top level is a quick way to get started solving a problem, but you should try to divide a program into functions as soon as possible. Functions are more reusable and testable, and clarify what steps are being done and what their inputs and outputs are. Furthermore, code inside functions tends to run much faster than top level code, due to how Julia's compiler works.

It is also worth emphasizing that functions should take arguments, instead of operating directly on global variables (aside from constants like *pi* (page 435)).

36.2 Avoid writing overly-specific types

Code should be as generic as possible. Instead of writing:

```
convert(Complex{Float64}, x)
```

it's better to use available generic functions:

```
complex(float(x))
```

The second version will convert *x* to an appropriate type, instead of always the same type.

This style point is especially relevant to function arguments. For example, don't declare an argument to be of type *Int* or *Int32* if it really could be any integer, expressed with the abstract type *Integer*. In fact, in many cases you can omit the argument type altogether, unless it is needed to disambiguate from other method definitions, since a *MethodError* (page 375) will be thrown anyway if a type is passed that does not support any of the requisite operations. (This is known as *duck typing*.)

For example, consider the following definitions of a function *addone* that returns one plus its argument:

```
addone(x::Int) = x + 1           # works only for Int
addone(x::Integer) = x + one(x) # any integer type
addone(x::Number) = x + one(x)  # any numeric type
addone(x) = x + one(x)          # any type supporting + and one
```

The last definition of `addone` handles any type supporting `one()` (page 435) (which returns 1 in the same type as `x`, which avoids unwanted type promotion) and the `+` (page 405) function with those arguments. The key thing to realize is that there is *no performance penalty* to defining *only* the general `addone(x) = x + one(x)`, because Julia will automatically compile specialized versions as needed. For example, the first time you call `addone(12)`, Julia will automatically compile a specialized `addone` function for `x::Int` arguments, with the call to `one()` (page 435) replaced by its inlined value 1. Therefore, the first three definitions of `addone` above are completely redundant.

36.3 Handle excess argument diversity in the caller

Instead of:

```
function foo(x, y)
    x = Int(x); y = Int(y)
    ...
end
foo(x, y)
```

use:

```
function foo(x::Int, y::Int)
    ...
end
foo(Int(x), Int(y))
```

This is better style because `foo` does not really accept numbers of all types; it really needs `Int`s.

One issue here is that if a function inherently requires integers, it might be better to force the caller to decide how non-integers should be converted (e.g. `floor` or `ceiling`). Another issue is that declaring more specific types leaves more “space” for future method definitions.

36.4 Append ! to names of functions that modify their arguments

Instead of:

```
function double{T<:Number}(a::AbstractArray{T})
    for i = 1:endof(a); a[i] *= 2; end
    a
end
```

use:

```
function double!{T<:Number}(a::AbstractArray{T})
    for i = 1:endof(a); a[i] *= 2; end
    a
end
```

The Julia standard library uses this convention throughout and contains examples of functions with both copying and modifying forms (e.g., `sort()` (page 578) and `sort!()` (page 578)), and others which are just modifying (e.g., `push!()` (page 398), `pop!()` (page 395), `splice!()` (page 399)). It is typical for such functions to also return the modified array for convenience.

36.5 Avoid strange type Unions

Types such as `Union{Function, AbstractString}` are often a sign that some design could be cleaner.

36.6 Avoid type Unions in fields

When creating a type such as:

```
type MyType
    ...
    x::Union{Void, T}
end
```

ask whether the option for `x` to be `nothing` (of type `Void`) is really necessary. Here are some alternatives to consider:

- Find a safe default value to initialize `x` with
- Introduce another type that lacks `x`
- If there are many fields like `x`, store them in a dictionary
- Determine whether there is a simple rule for when `x` is `nothing`. For example, often the field will start as `nothing` but get initialized at some well-defined point. In that case, consider leaving it undefined at first.
- If `x` really needs to hold no value at some times, define it as `::Nullable{T}` instead, as this guarantees type-stability in the code accessing this field (see [Nullable types](#) (page 107))

36.7 Avoid elaborate container types

It is usually not much help to construct arrays like the following:

```
a = Array{Union{Int, AbstractString, Tuple, Array}}(n)
```

In this case `Array{Any}(n)` is better. It is also more helpful to the compiler to annotate specific uses (e.g. `a[i]::Int`) than to try to pack many alternatives into one type.

36.8 Use naming conventions consistent with Julia's base/

- modules and type names use capitalization and camel case: module `SparseArrays`, immutable `UnitRange`.
- functions are lowercase (`maximum()` (page 385), `convert()` (page 364)) and, when readable, with multiple words squashed together (`isequal()` (page 362), `haskey()` (page 394)). When necessary, use underscores as word separators. Underscores are also used to indicate a combination of concepts (`remotecall_fetch()` (page 491) as a more efficient implementation of `fetch(remotecall(...))`) or as modifiers (`sum_kbn()` (page 476)).
- conciseness is valued, but avoid abbreviation (`indexin()` (page 384) rather than `indxin()`) as it becomes difficult to remember whether and how particular words are abbreviated.

If a function name requires multiple words, consider whether it might represent more than one concept and might be better split into pieces.

36.9 Don't overuse try-catch

It is better to avoid errors than to rely on catching them.

36.10 Don't parenthesize conditions

Julia doesn't require parens around conditions in `if` and `while`. Write:

```
if a == b
```

instead of:

```
if (a == b)
```

36.11 Don't overuse ...

Splicing function arguments can be addictive. Instead of `[a..., b...]`, use simply `[a; b]`, which already concatenates arrays. `collect(a)` (page 392) is better than `[a...]`, but since `a` is already iterable it is often even better to leave it alone, and not convert it to an array.

36.12 Don't use unnecessary static parameters

A function signature:

```
foo{T<:Real}(x::T) = ...
```

should be written as:

```
foo(x::Real) = ...
```

instead, especially if `T` is not used in the function body. Even if `T` is used, it can be replaced with `typeof(x)` (page 363) if convenient. There is no performance difference. Note that this is not a general caution against static parameters, just against uses where they are not needed.

Note also that container types, specifically may need type parameters in function calls. See the FAQ [Avoid fields with abstract containers](#) (page 309) for more information.

36.13 Avoid confusion about whether something is an instance or a type

Sets of definitions like the following are confusing:

```
foo(::Type{MyType}) = ...  
foo(::MyType) = foo(MyType)
```

Decide whether the concept in question will be written as `MyType` or `MyType()`, and stick to it.

The preferred style is to use instances by default, and only add methods involving `Type{MyType}` later if they become necessary to solve some problem.

If a type is effectively an enumeration, it should be defined as a single (ideally `immutable`) type, with the enumeration values being instances of it. Constructors and conversions can check whether values are valid. This design is preferred over making the enumeration an abstract type, with the “values” as subtypes.

36.14 Don’t overuse macros

Be aware of when a macro could really be a function instead.

Calling `eval()` (page 368) inside a macro is a particularly dangerous warning sign; it means the macro will only work when called at the top level. If such a macro is written as a function instead, it will naturally have access to the run-time values it needs.

36.15 Don’t expose unsafe operations at the interface level

If you have a type that uses a native pointer:

```
type NativeType
    p::Ptr{UInt8}
    ...
end
```

don’t write definitions like the following:

```
getindex(x::NativeType, i) = unsafe_load(x.p, i)
```

The problem is that users of this type can write `x[i]` without realizing that the operation is unsafe, and then be susceptible to memory bugs.

Such a function should either check the operation to ensure it is safe, or have `unsafe` somewhere in its name to alert callers.

36.16 Don’t overload methods of base container types

It is possible to write definitions like the following:

```
show(io::IO, v::Vector{MyType}) = ...
```

This would provide custom showing of vectors with a specific new element type. While tempting, this should be avoided. The trouble is that users will expect a well-known type like `Vector()` to behave in a certain way, and overly customizing its behavior can make it harder to work with.

36.17 Be careful with type equality

You generally want to use `isa()` (page 362) and `<: (issubtype())` (page 365)) for testing types, not `==`. Checking types for exact equality typically only makes sense when comparing to a known concrete type (e.g. `T == Float64`), or if you *really, really* know what you’re doing.

36.18 Do not write `x->f(x)`

Since higher-order functions are often called with anonymous functions, it is easy to conclude that this is desirable or even necessary. But any function can be passed directly, without being “wrapped” in an anonymous function. Instead of writing `map(x->f(x), a)`, write `map(f, a)` (page 390).

36.19 Avoid using floats for numeric literals in generic code when possible

If you write generic code which handles numbers, and which can be expected to run with many different numeric type arguments, try using literals of a numeric type that will affect the arguments as little as possible through promotion.

For example,

```
julia> f(x) = 2.0 * x
f (generic function with 1 method)

julia> f(1//2)
1.0

julia> f(1/2)
1.0

julia> f(1)
2.0
```

while

```
julia> g(x) = 2 * x
g (generic function with 1 method)

julia> g(1//2)
1//1

julia> g(1/2)
1.0

julia> g(2)
4
```

As you can see, the second version, where we used an `Int` literal, preserved the type of the input argument, while the first didn't. This is because e.g. `promote_type{Int, Float64} == Float64`, and promotion happens with the multiplication. Similarly, `Rational` literals are less type disruptive than `Float64` (page 437) literals, but more disruptive than `Ints`:

```
julia> h(x) = 2//1 * x
h (generic function with 1 method)

julia> h(1//2)
1//1

julia> h(1/2)
1.0

julia> h(1)
2//1
```

Thus, use `Int` literals when possible, with `Rational{Int}` for literal non-integer numbers, in order to make it easier to use your code.

Frequently Asked Questions

37.1 Sessions and the REPL

37.1.1 How do I delete an object in memory?

Julia does not have an analog of MATLAB's `clear` function; once a name is defined in a Julia session (technically, in module `Main`), it is always present.

If memory usage is your concern, you can always replace objects with ones that consume less memory. For example, if `A` is a gigabyte-sized array that you no longer need, you can free the memory with `A = 0`. The memory will be released the next time the garbage collector runs; you can force this to happen with `gc()` (page 378).

37.1.2 How can I modify the declaration of a type/immutable in my session?

Perhaps you've defined a type and then realize you need to add a new field. If you try this at the REPL, you get the error:

```
ERROR: invalid redefinition of constant MyType
```

Types in module `Main` cannot be redefined.

While this can be inconvenient when you are developing new code, there's an excellent workaround. Modules can be replaced by redefining them, and so if you wrap all your new code inside a module you can redefine types and constants. You can't import the type names into `Main` and then expect to be able to redefine them there, but you can use the module name to resolve the scope. In other words, while developing you might use a workflow something like this:

```
include("mynewcode.jl")           # this defines a module MyModule
obj1 = MyModule.ObjConstructor(a, b)
obj2 = MyModule.somefunction(obj1)
# Got an error. Change something in "mynewcode.jl"
include("mynewcode.jl")           # reload the module
obj1 = MyModule.ObjConstructor(a, b) # old objects are no longer valid, must reconstruct
obj2 = MyModule.somefunction(obj1)  # this time it worked!
obj3 = MyModule.someotherfunction(obj2, c)
...
```

37.2 Functions

37.2.1 I passed an argument `x` to a function, modified it inside that function, but on the outside, the variable `x` is still unchanged. Why?

Suppose you call a function like this:

```
julia> x = 10
10

julia> function change_value!(y)
           y = 17
       end
change_value! (generic function with 1 method)

julia> change_value!(x)
17

julia> x # x is unchanged!
10
```

In Julia, the binding of a variable `x` cannot be changed by passing `x` as an argument to a function. When calling `change_value!(x)` in the above example, `y` is a newly created variable, bound initially to the value of `x`, i.e. 10; then `y` is rebound to the constant 17, while the variable `x` of the outer scope is left untouched.

But here is a thing you should pay attention to: suppose `x` is bound to an object of type `Array` (or any other *mutable* type). From within the function, you cannot “unbind” `x` from this `Array`, but you can change its content. For example:

```
julia> x = [1,2,3]
3-element Array{Int64,1}:
 1
 2
 3

julia> function change_array!(A)
           A[1] = 5
       end
change_array! (generic function with 1 method)

julia> change_array!(x)
5

julia> x
3-element Array{Int64,1}:
 5
 2
 3
```

Here we created a function `change_array!()`, that assigns 5 to the first element of the passed array (bound to `x` at the call site, and bound to `A` within the function). Notice that, after the function call, `x` is still bound to the same array, but the content of that array changed: the variables `A` and `x` were distinct bindings referring to the same mutable `Array` object.

37.2.2 Can I use `using` or `import` inside a function?

No, you are not allowed to have a `using` or `import` statement inside a function. If you want to import a module but only use its symbols inside a specific function or set of functions, you have two options:

1. Use `import`:

```
import Foo
function bar(...)
    ... refer to Foo symbols via Foo.baz ...
end
```

This loads the module `Foo` and defines a variable `Foo` that refers to the module, but does not import any of the other symbols from the module into the current namespace. You refer to the `Foo` symbols by their qualified names `Foo.bar` etc.

2. Wrap your function in a module:

```
module Bar
export bar
using Foo
function bar(...)
    ... refer to Foo.baz as simply baz ....
end
end
using Bar
```

This imports all the symbols from `Foo`, but only inside the module `Bar`.

37.2.3 What does the `...` operator do?

37.2.4 The two uses of the `...` operator: slurping and splatting

Many newcomers to Julia find the use of `...` operator confusing. Part of what makes the `...` operator confusing is that it means two different things depending on context.

37.2.5 `...` combines many arguments into one argument in function definitions

In the context of function definitions, the `...` operator is used to combine many different arguments into a single argument. This use of `...` for combining many different arguments into a single argument is called slurping:

```
julia> function printargs(args...)
    @printf("%s\n", typeof(args))
    for (i, arg) in enumerate(args)
        @printf("Arg %d = %s\n", i, arg)
    end
end
printargs (generic function with 1 method)

julia> printargs(1, 2, 3)
(Int64,Int64,Int64)
Arg 1 = 1
Arg 2 = 2
Arg 3 = 3
```

If Julia were a language that made more liberal use of ASCII characters, the slurping operator might have been written as `<-...` instead of `...`.

37.2.6 ... splits one argument into many different arguments in function calls

In contrast to the use of the ... operator to denote slurping many different arguments into one argument when defining a function, the ... operator is also used to cause a single function argument to be split apart into many different arguments when used in the context of a function call. This use of ... is called *splatting*:

```
julia> function threeargs(a, b, c)
    @printf("a = %s::%s\n", a, typeof(a))
    @printf("b = %s::%s\n", b, typeof(b))
    @printf("c = %s::%s\n", c, typeof(c))
end
threeargs (generic function with 1 method)

julia> vec = [1, 2, 3]
3-element Array{Int64,1}:
 1
 2
 3

julia> threeargs(vec...)
a = 1::Int64
b = 2::Int64
c = 3::Int64
```

If Julia were a language that made more liberal use of ASCII characters, the splatting operator might have been written as ...-> instead of

37.3 Types, type declarations, and constructors

37.3.1 What does “type-stable” mean?

It means that the type of the output is predictable from the types of the inputs. In particular, it means that the type of the output cannot vary depending on the *values* of the inputs. The following code is *not* type-stable:

```
function unstable(flag::Bool)
    if flag
        return 1
    else
        return 1.0
    end
end
```

It returns either an `Int` or a `Float64` depending on the value of its argument. Since Julia can’t predict the return type of this function at compile-time, any computation that uses it will have to guard against both types possibly occurring, making generation of fast machine code difficult.

37.3.2 Why does Julia give a `DomainError` for certain seemingly-sensible operations?

Certain operations make mathematical sense but result in errors:

```
julia> sqrt(-2.0)
ERROR: DomainError
 in sqrt at math.jl:128
```



```
julia> 2^-5
ERROR: DomainError
  in power_by_squaring at intfuncs.jl:70
  in ^ at intfuncs.jl:84
```

This behavior is an inconvenient consequence of the requirement for type-stability. In the case of `sqrt()` (page 420), most users want `sqrt(2.0)` to give a real number, and would be unhappy if it produced the complex number `1.4142135623730951 + 0.0im`. One could write the `sqrt()` (page 420) function to switch to a complex-valued output only when passed a negative number (which is what `sqrt()` (page 420) does in some other languages), but then the result would not be *type-stable* (page 338) and the `sqrt()` (page 420) function would have poor performance.

In these and other cases, you can get the result you want by choosing an *input type* that conveys your willingness to accept an *output type* in which the result can be represented:

```
julia> sqrt(-2.0+0im)
0.0 + 1.4142135623730951im

julia> 2.0^-5
0.03125
```

37.3.3 Why does Julia use native machine integer arithmetic?

Julia uses machine arithmetic for integer computations. This means that the range of `Int` values is bounded and wraps around at either end so that adding, subtracting and multiplying integers can overflow or underflow, leading to some results that can be unsettling at first:

```
julia> typemax{Int}
9223372036854775807

julia> ans+1
-9223372036854775808

julia> -ans
-9223372036854775808

julia> 2*ans
0
```

Clearly, this is far from the way mathematical integers behave, and you might think it less than ideal for a high-level programming language to expose this to the user. For numerical work where efficiency and transparency are at a premium, however, the alternatives are worse.

One alternative to consider would be to check each integer operation for overflow and promote results to bigger integer types such as `Int128` or `BigInt` (page 437) in the case of overflow. Unfortunately, this introduces major overhead on every integer operation (think incrementing a loop counter) – it requires emitting code to perform run-time overflow checks after arithmetic instructions and branches to handle potential overflows. Worse still, this would cause every computation involving integers to be type-unstable. As we mentioned above, *type-stability is crucial* (page 338) for effective generation of efficient code. If you can't count on the results of integer operations being integers, it's impossible to generate fast, simple code the way C and Fortran compilers do.

A variation on this approach, which avoids the appearance of type instability is to merge the `Int` and `BigInt` (page 437) types into a single hybrid integer type, that internally changes representation when a result no longer fits into the size of a machine integer. While this superficially avoids type-instability at the level of Julia code, it just sweeps the problem under the rug by foisting all of the same difficulties onto the C code implementing this hybrid integer type. This approach *can* be made to work and can even be made quite fast in many cases, but has several drawbacks. One problem is that the in-memory representation of integers and arrays of integers no longer match the

natural representation used by C, Fortran and other languages with native machine integers. Thus, to interoperate with those languages, we would ultimately need to introduce native integer types anyway. Any unbounded representation of integers cannot have a fixed number of bits, and thus cannot be stored inline in an array with fixed-size slots – large integer values will always require separate heap-allocated storage. And of course, no matter how clever a hybrid integer implementation one uses, there are always performance traps – situations where performance degrades unexpectedly. Complex representation, lack of interoperability with C and Fortran, the inability to represent integer arrays without additional heap storage, and unpredictable performance characteristics make even the cleverest hybrid integer implementations a poor choice for high-performance numerical work.

An alternative to using hybrid integers or promoting to BigInts is to use saturating integer arithmetic, where adding to the largest integer value leaves it unchanged and likewise for subtracting from the smallest integer value. This is precisely what Matlab™ does:

```
>> int64(9223372036854775807)
ans =
    9223372036854775807
>> int64(9223372036854775807) + 1
ans =
    9223372036854775807
>> int64(-9223372036854775808)
ans =
   -9223372036854775808
>> int64(-9223372036854775808) - 1
ans =
   -9223372036854775808
```

At first blush, this seems reasonable enough since 9223372036854775807 is much closer to 9223372036854775808 than -9223372036854775808 is and integers are still represented with a fixed size in a natural way that is compatible with C and Fortran. Saturated integer arithmetic, however, is deeply problematic. The first and most obvious issue is that this is not the way machine integer arithmetic works, so implementing saturated operations requires emitting instructions after each machine integer operation to check for underflow or overflow and replace the result with `typemin(Int)` (page 365) or `typemax(Int)` (page 365) as appropriate. This alone expands each integer operation from a single, fast instruction into half a dozen instructions, probably including branches. Ouch. But it gets worse – saturating integer arithmetic isn’t associative. Consider this Matlab computation:

```
>> n = int64(2)^62
4611686018427387904
>> n + (n - 1)
9223372036854775807
>> (n + n) - 1
9223372036854775806
```

This makes it hard to write many basic integer algorithms since a lot of common techniques depend on the fact that machine addition with overflow *is* associative. Consider finding the midpoint between integer values `lo` and `hi` in Julia using the expression `(lo + hi) >>> 1`:

```
julia> n = 2^62
4611686018427387904

julia> (n + 2n) >>> 1
6917529027641081856
```

See? No problem. That's the correct midpoint between 2^{62} and 2^{63} , despite the fact that $n + 2n$ is - 4611686018427387904. Now try it in Matlab:

```
>> (n + 2*n)/2

ans =

    4611686018427387904
```

Oops. Adding a `>>>` operator to Matlab wouldn't help, because saturation that occurs when adding n and $2n$ has already destroyed the information necessary to compute the correct midpoint.

Not only is lack of associativity unfortunate for programmers who cannot rely it for techniques like this, but it also defeats almost anything compilers might want to do to optimize integer arithmetic. For example, since Julia integers use normal machine integer arithmetic, LLVM is free to aggressively optimize simple little functions like $f(k) = 5k-1$. The machine code for this function is just this:

```
julia> code_native(f, (Int,))
.section    __TEXT,__text,regular,pure_instructions
Filename: none
Source line: 1
    push    RBP
    mov RBP, RSP
Source line: 1
    lea RAX, QWORD PTR [RDI + 4*RDI - 1]
    pop RBP
    ret
```

The actual body of the function is a single `lea` instruction, which computes the integer multiply and add at once. This is even more beneficial when `f` gets inlined into another function:

```
julia> function g(k,n)
    for i = 1:n
        k = f(k)
    end
    return k
end
g (generic function with 2 methods)

julia> code_native(g, (Int, Int))
.section    __TEXT,__text,regular,pure_instructions
Filename: none
Source line: 3
    push    RBP
    mov RBP, RSP
    test    RSI, RSI
    jle 22
    mov EAX, 1
Source line: 3
    lea RDI, QWORD PTR [RDI + 4*RDI - 1]
    inc RAX
    cmp RAX, RSI
Source line: 2
```

```
jle -17
Source line: 5
    mov RAX, RDI
    pop RBP
    ret
```

Since the call to `f` gets inlined, the loop body ends up being just a single `lea` instruction. Next, consider what happens if we make the number of loop iterations fixed:

```
julia> function g(k)
    for i = 1:10
        k = f(k)
    end
    return k
end
g (generic function with 2 methods)

julia> code_native(g, (Int,))
.section    __TEXT,__text,regular,pure_instructions
Filename: none
Source line: 3
    push    RBP
    mov RBP, RSP
Source line: 3
    imul    RAX, RDI, 9765625
    add RAX, -2441406
Source line: 5
    pop RBP
    ret
```

Because the compiler knows that integer addition and multiplication are associative and that multiplication distributes over addition – neither of which is true of saturating arithmetic – it can optimize the entire loop down to just a multiply and an add. Saturated arithmetic completely defeats this kind of optimization since associativity and distributivity can fail at each loop iteration, causing different outcomes depending on which iteration the failure occurs in. The compiler can unroll the loop, but it cannot algebraically reduce multiple operations into fewer equivalent operations.

The most reasonable alternative to having integer arithmetic silently overflow is to do checked arithmetic everywhere, raising errors when adds, subtracts, and multiplies overflow, producing values that are not value-correct. In this [blog post](#), Dan Luu analyzes this and finds that rather than the trivial cost that this approach should in theory have, it ends up having a substantial cost due to compilers (LLVM and GCC) not gracefully optimizing around the added overflow checks. If this improves in the future, we could consider defaulting to checked integer arithmetic in Julia, but for now, we have to live with the possibility of overflow.

37.4 Packages and Modules

37.4.1 What is the difference between “using” and “importall”?

There is only one difference, and on the surface (syntax-wise) it may seem very minor. The difference between `using` and `importall` is that with `using` you need to say `function Foo.bar(...` to extend module `Foo`’s function `bar` with a new method, but with `importall` or `import Foo.bar`, you only need to say `function bar(...` and it automatically extends module `Foo`’s function `bar`.

If you use `importall`, then `function Foo.bar(...` and `function bar(...` become equivalent. If you use `using`, then they are different.

The reason this is important enough to have been given separate syntax is that you don't want to accidentally extend a function that you didn't know existed, because that could easily cause a bug. This is most likely to happen with a method that takes a common type like a string or integer, because both you and the other module could define a method to handle such a common type. If you use `importall`, then you'll replace the other module's implementation of `bar(s::AbstractString)` with your new implementation, which could easily do something completely different (and break all/many future usages of the other functions in module `Foo` that depend on calling `bar`).

37.5 Nothingness and missing values

37.5.1 How does “null” or “nothingness” work in Julia?

Unlike many languages (for example, C and Java), Julia does not have a “null” value. When a reference (variable, object field, or array element) is uninitialized, accessing it will immediately throw an error. This situation can be detected using the `isdefined` function.

Some functions are used only for their side effects, and do not need to return a value. In these cases, the convention is to return the value `nothing`, which is just a singleton object of type `Void`. This is an ordinary type with no fields; there is nothing special about it except for this convention, and that the REPL does not print anything for it. Some language constructs that would not otherwise have a value also yield `nothing`, for example `if false; end`.

For situations where a value exists only sometimes (for example, missing statistical data), it is best to use the `Nullable{T}` type, which allows specifying the type of a missing value.

The empty tuple `()` is another form of nothingness. But, it should not really be thought of as nothing but rather a tuple of zero values.

In code written for Julia prior to version 0.4 you may occasionally see `None`, which is quite different. It is the empty (or “bottom”) type, a type with no values and no subtypes (except itself). This is now written as `Union{}` (an empty union type). You will generally not need to use this type.

37.6 Memory

37.6.1 Why does `x += y` allocate memory when `x` and `y` are arrays?

In Julia, `x += y` gets replaced during parsing by `x = x + y`. For arrays, this has the consequence that, rather than storing the result in the same location in memory as `x`, it allocates a new array to store the result.

While this behavior might surprise some, the choice is deliberate. The main reason is the presence of `immutable` objects within Julia, which cannot change their value once created. Indeed, a number is an immutable object; the statements `x = 5; x += 1` do not modify the meaning of 5, they modify the value bound to `x`. For an immutable, the only way to change the value is to reassign it.

To amplify a bit further, consider the following function:

```
function power_by_squaring(x, n::Int)
    ispow2(n) || error("This implementation only works for powers of 2")
    while n >= 2
        x *= x
        n >>= 1
    end
    x
end
```

After a call like `x = 5; y = power_by_squaring(x, 4)`, you would get the expected result: `x == 5` && `y == 625`. However, now suppose that `*`, when used with matrices, instead mutated the left hand side. There would be two problems:

- For general square matrices, `A = A*B` cannot be implemented without temporary storage: `A[1,1]` gets computed and stored on the left hand side before you're done using it on the right hand side.
- Suppose you were willing to allocate a temporary for the computation (which would eliminate most of the point of making `*` work in-place); if you took advantage of the mutability of `x`, then this function would behave differently for mutable vs. immutable inputs. In particular, for immutable `x`, after the call you'd have (in general) `y != x`, but for mutable `x` you'd have `y == x`.

Because supporting generic programming is deemed more important than potential performance optimizations that can be achieved by other means (e.g., using explicit loops), operators like `+=` and `*` work by rebinding new values.

37.7 Asynchronous IO and concurrent synchronous writes

37.7.1 Why do concurrent writes to the same stream result in inter-mixed output?

While the streaming I/O API is synchronous, the underlying implementation is fully asynchronous.

The following:

```
@sync for i in 1:3
    @async write(STDOUT, string(i), " Foo ", " Bar ")
end
```

results in:

```
123 Foo Foo Foo Bar Bar Bar
```

This is happening because, while the `write` call is synchronous, the writing of each argument yields to other tasks while waiting for that part of the I/O to complete.

`print` and `println` “lock” the stream during a call. Consequently changing `write` to `println` in the above example results in:

```
1 Foo Bar
2 Foo Bar
3 Foo Bar
```

You can lock your writes with a `ReentrantLock` like this:

```
l = ReentrantLock()
@sync for i in 1:3
    @async begin
        lock(l)
        try
            write(STDOUT, string(i), " Foo ", " Bar ")
        finally
            unlock(l)
        end
    end
end
```

37.8 Julia Releases

37.8.1 Do I want to use a release, beta, or nightly version of Julia?

You may prefer the release version of Julia if you are looking for a stable code base. Releases generally occur every 6 months, giving you a stable platform for writing code.

You may prefer the beta version of Julia if you don't mind being slightly behind the latest bugfixes and changes, but find the slightly faster rate of changes more appealing. Additionally, these binaries are tested before they are published to ensure they are fully functional.

You may prefer the nightly version of Julia if you want to take advantage of the latest updates to the language, and don't mind if the version available today occasionally doesn't actually work.

Finally, you may also consider building Julia from source for yourself. This option is mainly for those individuals who are comfortable at the command line, or interested in learning. If this describes you, you may also be interested in reading our [guidelines for contributing](#).

Links to each of these download types can be found on the download page at <http://julialang.org/downloads/>. Note that not all versions of Julia are available for all platforms.

37.8.2 When are deprecated functions removed?

Deprecated functions are removed after the subsequent release. For example, functions marked as deprecated in the 0.1 release will not be available starting with the 0.2 release.

Noteworthy Differences from other Languages

38.1 Noteworthy differences from MATLAB

Although MATLAB users may find Julia’s syntax familiar, Julia is not a MATLAB clone. There are major syntactic and functional differences. The following are some noteworthy differences that may trip up Julia users accustomed to MATLAB:

- Julia arrays are indexed with square brackets, `A[i, j]`.
- Julia arrays are assigned by reference. After `A=B`, changing elements of `B` will modify `A` as well.
- Julia values are passed and assigned by reference. If a function modifies an array, the changes will be visible in the caller.
- Julia does not automatically grow arrays in an assignment statement. Whereas in MATLAB `a(4) = 3.2` can create the array `a = [0 0 0 3.2]` and `a(5) = 7` can grow it into `a = [0 0 0 3.2 7]`, the corresponding Julia statement `a[5] = 7` throws an error if the length of `a` is less than 5 or if this statement is the first use of the identifier `a`. Julia has `push!()` (page 398) and `append!()` (page 401), which grow `Vectors` much more efficiently than MATLAB’s `a(end+1) = val`.
- The imaginary unit `sqrt(-1)` is represented in Julia as `im` (page 435), not `i` or `j` as in MATLAB.
- In Julia, literal numbers without a decimal point (such as `42`) create integers instead of floating point numbers. Arbitrarily large integer literals are supported. As a result, some operations such as `2^-1` will throw a domain error as the result is not an integer (see *the FAQ entry on domain errors* (page 338) for details).
- In Julia, multiple values are returned and assigned as tuples, e.g. `(a, b) = (1, 2)` or `a, b = 1, 2`. MATLAB’s `nargout`, which is often used in MATLAB to do optional work based on the number of returned values, does not exist in Julia. Instead, users can use optional and keyword arguments to achieve similar capabilities.
- Julia has true one-dimensional arrays. Column vectors are of size `N`, not `N×1`. For example, `rand(N)` (page 440) makes a 1-dimensional array.
- In Julia v0.3, concatenating scalars and arrays with the syntax `[x, y, z]` concatenates in the first dimension (“vertically”). For concatenation in the second dimension (“horizontally”), use spaces as in `[x y z]`. To construct block matrices (concatenating in the first two dimensions), the syntax `[a b; c d]` is used to avoid confusion. In Julia v0.4, the concatenation syntax `[x, [y, z]]` is deprecated in favor of `[x; [y, z]]`.
- In Julia, `a:b` and `a:b:c` construct `Range` objects. To construct a full vector like in MATLAB, use `collect(a:b)` (page 392). Generally, there is no need to call `collect` though. `Range` will act like a normal array in most cases but is more efficient because it lazily computes its values. This pattern of creating specialized objects instead of full arrays is used frequently, and is also seen in functions such as `linspace`

(page 458), or with iterators such as `enumerate`, and `zip`. The special objects can mostly be used as if they were normal arrays.

- Functions in Julia return values from their last expression or the `return` keyword instead of listing the names of variables to return in the function definition (see [The `return` Keyword](#) (page 56) for details).
- A Julia script may contain any number of functions, and all definitions will be externally visible when the file is loaded. Function definitions can be loaded from files outside the current working directory.
- In Julia, reductions such as `sum()` (page 388), `prod()` (page 388), and `max()` (page 418) are performed over every element of an array when called with a single argument, as in `sum(A)`, even if `A` has more than one dimension.
- In Julia, functions such as `sort()` (page 578) that operate column-wise by default (`sort(A)` is equivalent to `sort(A, 1)`) do not have special behavior for $1 \times N$ arrays; the argument is returned unmodified since it still performs `sort(A, 1)`. To sort a $1 \times N$ matrix like a vector, use `sort(A, 2)`.
- In Julia, if `A` is a 2-dimensional array, `fft(A)` computes a 2D FFT. In particular, it is not equivalent to `fft(A, 1)`, which computes a 1D FFT acting column-wise.
- In Julia, parentheses must be used to call a function with zero arguments, like in `tic()` (page 372) and `toc()` (page 373).
- Julia discourages the use of semicolons to end statements. The results of statements are not automatically printed (except at the interactive prompt), and lines of code do not need to end with semicolons. `println()` (page 566) or `@printf()` (page 567) can be used to print specific output.
- In Julia, if `A` and `B` are arrays, logical comparison operations like `A == B` do not return an array of booleans. Instead, use `A .== B`, and similarly for the other boolean operators like `<` (page 411), `>` (page 411) and `=`.
- In Julia, the operators `&` (page 411), `|` (page 411), and `^` (page 411) (`xor` (page 411)) perform the bitwise operations equivalent to `and`, `or`, and `xor` respectively in MATLAB, and have precedence similar to Python's bitwise operators (unlike C). They can operate on scalars or element-wise across arrays and can be used to combine logical arrays, but note the difference in order of operations: parentheses may be required (e.g., to select elements of `A` equal to 1 or 2 use `(A .== 1) | (A .== 2)`).
- In Julia, the elements of a collection can be passed as arguments to a function using the splat operator `...`, as in `xs=[1,2]; f(xs...)`.
- Julia's `svd()` (page 520) returns singular values as a vector instead of as a dense diagonal matrix.
- In Julia, `...` is not used to continue lines of code. Instead, incomplete expressions automatically continue onto the next line.
- In both Julia and MATLAB, the variable `ans` is set to the value of the last expression issued in an interactive session. In Julia, unlike MATLAB, `ans` is not set when Julia code is run in non-interactive mode.
- Julia's types do not support dynamically adding fields at runtime, unlike MATLAB's classes. Instead, use a `Dict` (page 394).
- In Julia each module has its own global scope/namespace, whereas in MATLAB there is just one global scope.
- In MATLAB, an idiomatic way to remove unwanted values is to use logical indexing, like in the expression `x(x>3)` or in the statement `x(x>3) = []` to modify `x` in-place. In contrast, Julia provides the higher order functions `filter()` (page 392) and `filter!()` (page 392), allowing users to write `filter(z->z>3, x)` and `filter!(z->z>3, x)` as alternatives to the corresponding transliterations `x[x.>3]` and `x = x[x.>3]`. Using `filter!()` (page 392) reduces the use of temporary arrays.
- The analogue of extracting (or “dereferencing”) all elements of a cell array, e.g. in `vertcat(A{:})` in MATLAB, is written using the splat operator in Julia, e.g. as `vcate(A...)`.

38.2 Noteworthy differences from R

One of Julia's goals is to provide an effective language for data analysis and statistical programming. For users coming to Julia from R, these are some noteworthy differences:

- Julia's single quotes enclose characters, not strings.
- Julia can create substrings by indexing into strings. In R, strings must be converted into character vectors before creating substrings.
- In Julia, like Python but unlike R, strings can be created with triple quotes `""" ... """`. This syntax is convenient for constructing strings that contain line breaks.
- In Julia, varargs are specified using the splat operator `...`, which always follows the name of a specific variable, unlike R, for which `...` can occur in isolation.
- In Julia, modulus is `mod(a, b)`, not `a %% b`. `%` in Julia is the remainder operator.
- In Julia, not all data structures support logical indexing. Furthermore, logical indexing in Julia is supported only with vectors of length equal to the object being indexed. For example: - In R, `c(1, 2, 3, 4)[c(TRUE, FALSE)]` is equivalent to `c(1, 3)`. - In R, `c(1, 2, 3, 4)[c(TRUE, FALSE, TRUE, FALSE)]` is equivalent to `c(1, 3)`. - In Julia, `[1, 2, 3, 4][true, false]` throws a *BoundsError* (page 374). - In Julia, `[1, 2, 3, 4][[true, false, true, false]]` produces `[1, 3]`.
- Like many languages, Julia does not always allow operations on vectors of different lengths, unlike R where the vectors only need to share a common index range. For example, `c(1, 2, 3, 4) + c(1, 2)` is valid R but the equivalent `[1, 2, 3, 4] + [1, 2]` will throw an error in Julia.
- Julia's `map()` (page 390) takes the function first, then its arguments, unlike `lapply(<structure>, function, ...)` in R. Similarly Julia's equivalent of `apply(X, MARGIN, FUN, ...)` in R is `maplices()` (page 475) where the function is the first argument.
- Multivariate `apply` in R, e.g. `mapply(choose, 11:13, 1:3)`, can be written as `broadcast(binomial, 11:13, 1:3)` in Julia. Equivalently Julia offers a shorter dot syntax for vectorizing functions `binomial.(11:13, 1:3)`.
- Julia uses `end` to denote the end of conditional blocks, like `if`, loop blocks, like `while/for`, and functions. In lieu of the one-line `if (cond) statement`, Julia allows statements of the form `if cond; statement; end`, `cond && statement` and `!cond || statement`. Assignment statements in the latter two syntaxes must be explicitly wrapped in parentheses, e.g. `cond && (x = value)`.
- In Julia, `<-`, `<<-` and `->` are not assignment operators.
- Julia's `->` creates an anonymous function, like Python.
- Julia constructs vectors using brackets. Julia's `[1, 2, 3]` is the equivalent of R's `c(1, 2, 3)`.
- Julia's `*` (page 443) operator can perform matrix multiplication, unlike in R. If `A` and `B` are matrices, then `A * B` denotes a matrix multiplication in Julia, equivalent to R's `A %*% B`. In R, this same notation would perform an element-wise (Hadamard) product. To get the element-wise multiplication operation, you need to write `A .* B` in Julia.
- Julia performs matrix transposition using the `.'` operator and conjugated transposition using the `'` operator. Julia's `A.'` is therefore equivalent to R's `t(A)`.
- Julia does not require parentheses when writing `if` statements or `for/while` loops: use `for i in [1, 2, 3]` instead of `for (i in c(1, 2, 3))` and `if i == 1` instead of `if (i == 1)`.
- Julia does not treat the numbers 0 and 1 as Booleans. You cannot write `if (1)` in Julia, because `if` statements accept only booleans. Instead, you can write `if true`, `if Bool(1)`, or `if 1==1`.

- Julia does not provide `nrow` and `ncol`. Instead, use `size(M, 1)` for `nrow(M)` and `size(M, 2)` for `ncol(M)`.
- Julia is careful to distinguish scalars, vectors and matrices. In R, `1` and `c(1)` are the same. In Julia, they can not be used interchangeably. One potentially confusing result of this is that `x' * y` for vectors `x` and `y` is a 1-element vector, not a scalar. To get a scalar, use `dot(x, y)` (page 501).
- Julia's `diag()` (page 524) and `diagm()` (page 525) are not like R's.
- Julia cannot assign to the results of function calls on the left hand side of an assignment operation: you cannot write `diag(M) = ones(n)`.
- Julia discourages populating the main namespace with functions. Most statistical functionality for Julia is found in [packages](#) under the [JuliaStats organization](#). For example:
 - Functions pertaining to probability distributions are provided by the [Distributions package](#).
 - The [DataFrames package](#) provides data frames.
 - Generalized linear models are provided by the [GLM package](#).
- Julia provides tuples and real hash tables, but not R-style lists. When returning multiple items, you should typically use a tuple: instead of `list(a = 1, b = 2)`, use `(1, 2)`.
- Julia encourages users to write their own types, which are easier to use than S3 or S4 objects in R. Julia's multiple dispatch system means that `table(x::TypeA)` and `table(x::TypeB)` act like R's `table.TypeA(x)` and `table.TypeB(x)`.
- In Julia, values are passed and assigned by reference. If a function modifies an array, the changes will be visible in the caller. This is very different from R and allows new functions to operate on large data structures much more efficiently.
- In Julia, vectors and matrices are concatenated using `hcat()` (page 462), `vcat()` (page 462) and `hvcate()` (page 463), not `c`, `rbind` and `cbind` like in R.
- In Julia, a range like `a:b` is not shorthand for a vector like in R, but is a specialized `Range` that is used for iteration without high memory overhead. To convert a range into a vector, use `collect(a:b)` (page 392).
- Julia's `max()` (page 418) and `min()` (page 417) are the equivalent of `pmax` and `pmin` respectively in R, but both arguments need to have the same dimensions. While `maximum()` (page 385) and `minimum()` (page 385) replace `max` and `min` in R, there are important differences.
- Julia's `sum()` (page 388), `prod()` (page 388), `maximum()` (page 385), and `minimum()` (page 385) are different from their counterparts in R. They all accept one or two arguments. The first argument is an iterable collection such as an array. If there is a second argument, then this argument indicates the dimensions, over which the operation is carried out. For instance, let `A=[[1 2], [3 4]]` in Julia and `B=rbind(c(1,2), c(3,4))` be the same matrix in R. Then `sum(A)` gives the same result as `sum(B)`, but `sum(A, 1)` is a row vector containing the sum over each column and `sum(A, 2)` is a column vector containing the sum over each row. This contrasts to the behavior of R, where `sum(B, 1)=11` and `sum(B, 2)=12`. If the second argument is a vector, then it specifies all the dimensions over which the sum is performed, e.g., `sum(A, [1,2])=10`. It should be noted that there is no error checking regarding the second argument.
- Julia has several functions that can mutate their arguments. For example, it has both `sort()` (page 578) and `sort!()` (page 578).
- In R, performance requires vectorization. In Julia, almost the opposite is true: the best performing code is often achieved by using devectorized loops.
- Julia is eagerly evaluated and does not support R-style lazy evaluation. For most users, this means that there are very few unquoted expressions or column names.
- Julia does not support the `NULL` type.

- Julia lacks the equivalent of R's `assign` or `get`.
- In Julia, `return` does not require parentheses.
- In R, an idiomatic way to remove unwanted values is to use logical indexing, like in the expression `x[x>3]` or in the statement `x = x[x>3]` to modify `x` in-place. In contrast, Julia provides the higher order functions `filter()` (page 392) and `filter!()` (page 392), allowing users to write `filter(z->z>3, x)` and `filter!(z->z>3, x)` as alternatives to the corresponding transliterations `x[x.>3]` and `x = x[x.>3]`. Using `filter!()` (page 392) reduces the use of temporary arrays.

38.3 Noteworthy differences from Python

- Julia requires `end` to end a block. Unlike Python, Julia has no `pass` keyword.
- In Julia, indexing of arrays, strings, etc. is 1-based not 0-based.
- Julia's slice indexing includes the last element, unlike in Python. `a[2:3]` in Julia is `a[1:3]` in Python.
- Julia does not support negative indexes. In particular, the last element of a list or array is indexed with `end` in Julia, not `-1` as in Python.
- Julia's `for`, `if`, `while`, etc. blocks are terminated by the `end` keyword. Indentation level is not significant as it is in Python.
- Julia has no line continuation syntax: if, at the end of a line, the input so far is a complete expression, it is considered done; otherwise the input continues. One way to force an expression to continue is to wrap it in parentheses.
- Julia arrays are column major (Fortran ordered) whereas NumPy arrays are row major (C-ordered) by default. To get optimal performance when looping over arrays, the order of the loops should be reversed in Julia relative to NumPy (see relevant section of *Performance Tips* (page 305)).
- Julia's updating operators (e.g. `+=`, `-=`, ...) are *not in-place* whereas NumPy's are. This means `A = ones(4); B = A; B += 3` doesn't change values in `A`, it rather rebinds the name `B` to the result of the right-hand side `B = B + 3`, which is a new array. Use `B[:] += 3`, explicit loops, or `InplaceOps.jl`.
- Julia evaluates default values of function arguments every time the method is invoked, unlike in Python where the default values are evaluated only once when the function is defined. For example, the function `f(x=rand()) = x` returns a new random number every time it is invoked without argument. On the other hand, the function `g(x=[1,2]) = push!(x,3)` returns `[1,2,3]` every time it is called as `g()`.
- In Julia `%` is the remainder operator, whereas in Python it is the modulus.

38.4 Noteworthy differences from C/C++

- Julia arrays are indexed with square brackets, and can have more than one dimension `A[i,j]`. This syntax is not just syntactic sugar for a reference to a pointer or address as in C/C++. See the Julia documentation for the syntax for array construction (it has changed between versions).
- In Julia, indexing of arrays, strings, etc. is 1-based not 0-based.
- Julia arrays are assigned by reference. After `A=B`, changing elements of `B` will modify `A` as well. Updating operators like `+=` do not operate in-place, they are equivalent to `A = A + B` which rebinds the left-hand side to the result of the right-hand side expression.
- Julia arrays are column major (Fortran ordered) whereas C/C++ arrays are row major ordered by default. To get optimal performance when looping over arrays, the order of the loops should be reversed in Julia relative to C/C++ (see relevant section of *Performance Tips* (page 305)).

- Julia values are passed and assigned by reference. If a function modifies an array, the changes will be visible in the caller.
- In Julia, whitespace is significant, unlike C/C++, so care must be taken when adding/removing whitespace from a Julia program.
- In Julia, literal numbers without a decimal point (such as 42) create signed integers, of type `Int`, but literals too large to fit in the machine word size will automatically be promoted to a larger size type, such as `Int64` (if `Int` is `Int32`), `Int128`, or the arbitrarily large `BigInt` type. There are no numeric literal suffixes, such as `L`, `LL`, `U`, `UL`, `ULL` to indicate unsigned and/or signed vs. unsigned. Decimal literals are always signed, and hexadecimal literals (which start with `0x` like C/C++), are unsigned. Hexadecimal literals also, unlike C/C++/Java and unlike decimal literals in Julia, have a type based on the *length* of the literal, including leading 0s. For example, `0x0` and `0x00` have type `UInt8`, `0x000` and `0x0000` have type `UInt16`, then literals with 5 to 8 hex digits have type `UInt32`, 9 to 16 hex digits type `UInt64` and 17 to 32 hex digits type `UInt128`. This needs to be taken into account when defining hexadecimal masks, for example `~0xf == 0xf0` is very different from `~0x000f == 0xffff0`. 64 bit `Float64` and 32 bit `Float32` bit literals are expressed as `1.0` and `1.0f0` respectively. Floating point literals are rounded (and not promoted to the `BigFloat` type) if they can not be exactly represented. Floating point literals are closer in behavior to C/C++. Octal (prefixed with `0o`) and binary (prefixed with `0b`) literals are also treated as unsigned.
- String literals can be delimited with either `"` or `" "`, `" "` delimited literals can contain `"` characters without quoting it like `"\""` String literals can have values of other variables or expressions interpolated into them, indicated by `$variablename` or `$(expression)`, which evaluates the variable name or the expression in the context of the function.
- `//` indicates a `Rational` number, and not a single-line comment (which is `#` in Julia)
- `#=` indicates the start of a multiline comment, and `#=` ends it.
- Functions in Julia return values from their last expression(s) or the `return` keyword. Multiple values can be returned from functions and assigned as tuples, e.g. `(a, b) = myfunction()` or `a, b = myfunction()`, instead of having to pass pointers to values as one would have to do in C/C++ (i.e. `a = myfunction(&b)`).
- Julia does not require the use of semicolons to end statements. The results of expressions are not automatically printed (except at the interactive prompt, i.e. the REPL), and lines of code do not need to end with semicolons. `println()` (page 566) or `@printf()` (page 567) can be used to print specific output. In the REPL, `;` can be used to suppress output. `;` also has a different meaning within `[]`, something to watch out for. `;` can be used to separate expressions on a single line, but are not strictly necessary in many cases, and are more an aid to readability.
- In Julia, the operator (page 411) (`xor` (page 411)) performs the bitwise XOR operation, i.e. `^` (page 443) in C/C++. Also, the bitwise operators do not have the same precedence as C/C++, so parenthesis may be required.
- Julia's `^` (page 443) is exponentiation (pow), not bitwise XOR as in C/C++ (use (page 411), or `xor` (page 411), in Julia)
- Julia has two right-shift operators, `>>` and `>>>`. `>>>` performs an arithmetic shift, `>>` always performs a logical shift, unlike C/C++, where the meaning of `>>` depends on the type of the value being shifted.
- Julia's `->` creates an anonymous function, it does not access a member via a pointer.
- Julia does not require parentheses when writing `if` statements or `for/while` loops: use `for i in [1, 2, 3]` instead of `for (int i=1; i <= 3; i++)` and `if i == 1` instead of `if (i == 1)`.
- Julia does not treat the numbers 0 and 1 as Booleans. You cannot write `if (1)` in Julia, because `if` statements accept only booleans. Instead, you can write `if true`, `if Bool(1)`, or `if 1==1`.
- Julia uses `end` to denote the end of conditional blocks, like `if`, loop blocks, like `while/for`, and functions. In lieu of the one-line `if (cond) statement`, Julia allows statements of the form `if cond; statement; end`, `cond && statement` and `!cond || statement`. Assignment statements in the

latter two syntaxes must be explicitly wrapped in parentheses, e.g. `cond && (x = value)`, because of the operator precedence.

- Julia has no line continuation syntax: if, at the end of a line, the input so far is a complete expression, it is considered done; otherwise the input continues. One way to force an expression to continue is to wrap it in parentheses.
- Julia macros operate on parsed expressions, rather than the text of the program, which allows them to perform sophisticated transformations of Julia code. Macro names start with the `@` character, and have both a function-like syntax, `@mymacro(arg1, arg2, arg3)`, and a statement-like syntax, `@mymacro arg1 arg2 arg3`. The forms are interchangeable; the function-like form is particularly useful if the macro appears within another expression, and is often clearest. The statement-like form is often used to annotate blocks, as in the `parallel for` construct: `@parallel for i in 1:n; body =#; end`. Where the end of the macro construct may be unclear, use the function-like form.
- Julia now has an enumeration type, expressed using the macro `@enum(name, value1, value2, ...)`. For example: `@enum(Fruit, Banana=1, Apple, Pear)`
- By convention, functions that modify their arguments have a `!` at the end of the name, for example `push!`.
- In C++, by default, you have static dispatch, i.e. you need to annotate a function as `virtual`, in order to have dynamic dispatch. On the other hand, in Julia every method is “virtual” (although it’s more general than that since methods are dispatched on every argument type, not only `this`, using the most-specific-declaration rule).

Unicode Input

Please see the online documentation.

Part II

The Julia Standard Library

40.1 Introduction

The Julia standard library contains a range of functions and macros appropriate for performing scientific and numerical computing, but is also as broad as those of many general purpose programming languages. Additional functionality is available from a growing collection of available packages. Functions are grouped by topic below.

Some general notes:

- Except for functions in built-in modules (*Pkg* (page 581), *Collections* (page 401), *Test* (page 601) and *Profile*), all functions documented here are directly available for use in programs.
- To use module functions, use `import Module` to import the module, and `Module.fn(x)` to use the functions.
- Alternatively, `using Module` will import all exported `Module` functions into the current namespace.
- By convention, function names ending with an exclamation point (!) modify their arguments. Some functions have both modifying (e.g., `sort!`) and non-modifying (`sort`) versions.

40.2 Getting Around

exit ([*code*])

Quit (or control-D at the prompt). The default exit code is zero, indicating that the processes completed successfully.

quit ()

Quit the program indicating that the processes completed successfully. This function calls `exit(0)` (see [exit\(\)](#) (page 359)).

atexit (*f*)

Register a zero-argument function `f()` to be called at process exit. `atexit()` hooks are called in last in first out (LIFO) order and run before object finalizers.

atreplinit (*f*)

Register a one-argument function to be called before the REPL interface is initialized in interactive sessions; this is useful to customize the interface. The argument of `f` is the REPL object. This function should be called from within the `.juliarc.jl` initialization file.

isinteractive () → Bool

Determine whether Julia is running an interactive session.

whos (*io::IO=STDOUT, m::Module=current_module(), pattern::Regex=r""*)

Print information about exported global variables in a module, optionally restricted to those matching *pattern*.

The memory consumption estimate is an approximate lower bound on the size of the internal structure of the object.

Base **.summarysize** (*obj; exclude=Union{Module, DataType, TypeName}*) → Int

Compute the amount of memory used by all unique objects reachable from the argument. Keyword argument *exclude* specifies a type of objects to exclude from the traversal.

edit (*path::AbstractString, line::Integer=0*)

Edit a file or directory optionally providing a line number to edit the file at. Returns to the `julia` prompt when you quit the editor. The editor can be changed by setting `JULIA_EDITOR`, `VISUAL` or `EDITOR` as an environmental variable.

edit (*function[, types]*)

Edit the definition of a function, optionally specifying a tuple of types to indicate which method to edit. The editor can be changed by setting `JULIA_EDITOR`, `VISUAL` or `EDITOR` as an environmental variable.

@edit ()

Evaluates the arguments to the function or macro call, determines their types, and calls the `edit` function on the resulting expression.

less (*file::AbstractString[, line::Integer]*)

Show a file using the default pager, optionally providing a starting line number. Returns to the `julia` prompt when you quit the pager.

less (*function[, types]*)

Show the definition of a function using the default pager, optionally specifying a tuple of types to indicate which method to see.

@less ()

Evaluates the arguments to the function or macro call, determines their types, and calls the `less` function on the resulting expression.

clipboard (*x*)

Send a printed form of *x* to the operating system clipboard (“copy”).

clipboard () → AbstractString

Return a string with the contents of the operating system clipboard (“paste”).

reload (*name::AbstractString*)

Force reloading of a package, even if it has been loaded before. This is intended for use during package development as code is modified.

require (*module::Symbol*)

This function is part of the implementation of `using / import`, if a module is not already defined in `Main`. It can also be called directly to force reloading a module, regardless of whether it has been loaded before (for example, when interactively developing libraries).

Loads a source file, in the context of the `Main` module, on every active node, searching standard locations for files. `require` is considered a top-level operation, so it sets the current `include` path but does not use it to search for files (see help for `include`). This function is typically used to load library code, and is implicitly called by `using` to load packages.

When searching for files, `require` first looks for package code under `Pkg.dir()`, then tries paths in the global array `LOAD_PATH`. `require` is case-sensitive on all platforms, including those with case-insensitive filesystems like macOS and Windows.

Base **.compilecache** (*module::String*)

Creates a *Module initialization and precompilation* (page 142) for a module and all of its dependencies. This

can be used to reduce package load times. Cache files are stored in `LOAD_CACHE_PATH[1]`, which defaults to `~/.julia/lib/VERSION`. See [Module initialization and precompilation](#) (page 142) for important notes.

__precompile__ (*isprecompilable::Bool=true*)

Specify whether the file calling this function is precompilable. If `isprecompilable` is `true`, then `__precompile__` throws an exception when the file is loaded by using `using/import/require` *unless* the file is being precompiled, and in a module file it causes the module to be automatically precompiled when it is imported. Typically, `__precompile__()` should occur before the module declaration in the file, or better yet `VERSION >= v"0.4" && __precompile__()` in order to be backward-compatible with Julia 0.3.

If a module or file is *not* safely precompilable, it should call `__precompile__(false)` in order to throw an error if Julia attempts to precompile it.

`__precompile__()` should *not* be used in a module unless all of its dependencies are also using `__precompile__()`. Failure to do so can result in a runtime error when loading the module.

include (*path::AbstractString...*)

Evaluate the contents of the input source file(s) in the current context. Returns the result of the last evaluated argument (of the last input file). During including, a task-local include path is set to the directory containing the file. Nested calls to `include` will search relative to that path. All paths refer to files on node 1 when running in parallel, and files will be fetched from node 1. This function is typically used to load source interactively, or to combine files in packages that are broken into multiple source files.

include_string (*code::AbstractString, filename::AbstractString="string"*)

Like `include`, except reads code from the given string rather than from a file. Since there is no file path involved, no path processing or fetching from node 1 is done.

include_dependency (*path::AbstractString*)

In a module, declare that the file specified by `path` (relative or absolute) is a dependency for precompilation; that is, the module will need to be recompiled if this file changes.

This is only needed if your module depends on a file that is not used via `include`. It has no effect outside of compilation.

apropos (*string*)

Search through all documentation for a string, ignoring case.

which (*f, types*)

Returns the method of `f` (a `Method` object) that would be called for arguments of the given `types`.

If `types` is an abstract type, then the method that would be called by `invoke` is returned.

which (*symbol*)

Return the module in which the binding for the variable referenced by `symbol` was created.

@which ()

Applied to a function or macro call, it evaluates the arguments to the specified call, and returns the `Method` object for the method that would be called for those arguments. Applied to a variable, it returns the module in which the variable was bound. It calls out to the `which` function.

methods (*f[, types]*)

Returns the method table for `f`.

If `types` is specified, returns an array of methods whose types match.

methodswith (*typ[, module or function][, showparents::Bool=false]*)

Return an array of methods with an argument of type `typ`.

The optional second argument restricts the search to a particular module or function (the default is all modules, starting from `Main`).

If optional `showparents` is `true`, also return arguments with a parent type of `typ`, excluding type `Any`.

@show ()

Show an expression and result, returning the result.

versioninfo (*io::IO=STDOUT, verbose::Bool=false*)

Print information about the version of Julia in use. If the `verbose` argument is `true`, detailed system information is shown as well.

workspace ()

Replace the top-level module (`Main`) with a new one, providing a clean workspace. The previous `Main` module is made available as `LastMain`. A previously-loaded package can be accessed using a statement such as `using LastMain.Package`.

This function should only be used interactively.

ans

A variable referring to the last computed value, automatically set at the interactive prompt.

40.3 All Objects

=== (*x, y*) → *Bool*

\equiv (*x, y*) → *Bool*

Determine whether *x* and *y* are identical, in the sense that no program could distinguish them. Compares mutable objects by address in memory, and compares immutable objects (such as numbers) by contents at the bit level. This function is sometimes called `egal`.

isa (*x, type*) → *Bool*

Determine whether *x* is of the given *type*.

isequal (*x, y*)

Similar to `==`, except treats all floating-point NaN values as equal to each other, and treats `-0.0` as unequal to `0.0`. The default implementation of `isequal` calls `==`, so if you have a type that doesn't have these floating-point subtleties then you probably only need to define `==`.

`isequal` is the comparison function used by hash tables (`Dict`). `isequal(x, y)` must imply that `hash(x) == hash(y)`.

This typically means that if you define your own `==` function then you must define a corresponding `hash` (and vice versa). Collections typically implement `isequal` by calling `isequal` recursively on all contents.

Scalar types generally do not need to implement `isequal` separate from `==`, unless they represent floating-point numbers amenable to a more efficient implementation than that provided as a generic fallback (based on `isnan`, `signbit`, and `==`).

isequal (*x::Nullable, y::Nullable*)

If neither *x* nor *y* is null, compare them according to their values (i.e. `isequal(get(x), get(y))`). Else, return `true` if both arguments are null, and `false` if one is null but not the other: nulls are considered equal.

isless (*x, y*)

Test whether *x* is less than *y*, according to a canonical total order. Values that are normally unordered, such as NaN, are ordered in an arbitrary but consistent fashion. This is the default comparison used by `sort`. Non-numeric types with a canonical total order should implement this function. Numeric types only need to implement it if they have special values such as NaN.

isless (*x::Nullable, y::Nullable*)

If neither *x* nor *y* is null, compare them according to their values (i.e. `isless(get(x), get(y))`). Else, return `true` if only *y* is null, and `false` otherwise: nulls are always considered greater than non-nulls, but not greater than another null.

ifelse (*condition::Bool*, *x*, *y*)

Return *x* if *condition* is `true`, otherwise return *y*. This differs from `? or if` in that it is an ordinary function, so all the arguments are evaluated first. In some cases, using `ifelse` instead of an `if` statement can eliminate the branch in generated code and provide higher performance in tight loops.

lexcmp (*x*, *y*)

Compare *x* and *y* lexicographically and return -1, 0, or 1 depending on whether *x* is less than, equal to, or greater than *y*, respectively. This function should be defined for lexicographically comparable types, and `lexless` will call `lexcmp` by default.

lexless (*x*, *y*)

Determine whether *x* is lexicographically less than *y*.

typeof (*x*)

Get the concrete type of *x*.

tuple (*xs...*)

Construct a tuple of the given objects.

ntuple (*f::Function*, *n::Integer*)

Create a tuple of length *n*, computing each element as `f(i)`, where *i* is the index of the element.

object_id (*x*)

Get a hash value for *x* based on object identity. `object_id(x) == object_id(y)` if `x === y`.

hash (*x* [, *h::UInt*])

Compute an integer hash code such that `isequal(x, y)` implies `hash(x) == hash(y)`. The optional second argument *h* is a hash code to be mixed with the result.

New types should implement the 2-argument form, typically by calling the 2-argument `hash` method recursively in order to mix hashes of the contents with each other (and with *h*). Typically, any type that implements `hash` should also implement its own `==` (hence `isequal`) to guarantee the property mentioned above.

finalizer (*x*, *function*)

Register a function `f(x)` to be called when there are no program-accessible references to *x*. The behavior of this function is unpredictable if *x* is of a bits type.

finalize (*x*)

Immediately run finalizers registered for object *x*.

copy (*x*)

Create a shallow copy of *x*: the outer structure is copied, but not all internal values. For example, copying an array produces a new array with identically-same elements as the original.

deepcopy (*x*)

Create a deep copy of *x*: everything is copied recursively, resulting in a fully independent object. For example, deep-copying an array produces a new array whose elements are deep copies of the original elements. Calling `deepcopy` on an object should generally have the same effect as serializing and then deserializing it.

As a special case, functions can only be actually deep-copied if they are anonymous, otherwise they are just copied. The difference is only relevant in the case of closures, i.e. functions which may contain hidden internal references.

While it isn't normally necessary, user-defined types can override the default `deepcopy` behavior by defining a specialized version of the function `deepcopy_internal(x::T, dict::ObjectIdDict)` (which shouldn't otherwise be used), where *T* is the type to be specialized for, and `dict` keeps track of objects copied so far within the recursion. Within the definition, `deepcopy_internal` should be used in place of `deepcopy`, and the `dict` variable should be updated as appropriate before returning.

isdefined (*[m::Module]*, *s::Symbol*)

isdefined (*object*, *s::Symbol*)

isdefined (*object*, *index::Int*)

Tests whether an assignable location is defined. The arguments can be a module and a symbol or a composite object and field name (as a symbol) or index. With a single symbol argument, tests whether a global variable with that name is defined in `current_module()`.

convert (*T*, *x*)

Convert *x* to a value of type *T*.

If *T* is an Integer type, an *InexactError* (page 375) will be raised if *x* is not representable by *T*, for example if *x* is not integer-valued, or is outside the range supported by *T*.

```
julia> convert{Int, 3.0}
3

julia> convert{Int, 3.5}
ERROR: InexactError{Int64} at ./float.jl:656
...

```

If *T* is a *AbstractFloat* or *Rational* type, then it will return the closest value to *x* representable by *T*.

```
julia> x = 1/3
0.3333333333333333

julia> convert{Float32, x}
0.33333334f0

julia> convert{Rational{Int32}, x}
1//3

julia> convert{Rational{Int64}, x}
6004799503160661//18014398509481984

```

If *T* is a collection type and *x* a collection, the result of `convert(T, x)` may alias *x*.

```
julia> x = Int[1,2,3];

julia> y = convert{Vector{Int}, x};

julia> y === x
true

```

Similarly, if *T* is a composite type and *x* a related instance, the result of `convert(T, x)` may alias part or all of *x*.

```
julia> x = speye{5};

julia> typeof(x)
SparseMatrixCSC{Float64,Int64}

julia> y = convert{SparseMatrixCSC{Float64,Int64}, x};

julia> z = convert{SparseMatrixCSC{Float32,Int64}, y};

julia> y === x
true

julia> z === x
false

```

```
julia> z.colptr === x.colptr
true
```

promote (*xs...*)

Convert all arguments to their common promotion type (if any), and return them all (as a tuple).

oftype (*x, y*)

Convert *y* to the type of *x* (`convert (typeof (x) , y)`).

widen (*x*)

If *x* is a type, return a “larger” type (for numeric types, this will be a type with at least as much range and precision as the argument, and usually more). Otherwise *x* is converted to `widen (typeof (x))`.

```
julia> widen(Int32)
Int64

julia> widen(1.5f0)
1.5
```

identity (*x*)

The identity function. Returns its argument.

40.4 Types

supertype (*T::DataType*)

Return the supertype of *DataType* *T*.

```
julia> supertype(Int32)
Signed
```

issubtype (*type1, type2*)

Return `true` if and only if all values of *type1* are also of *type2*. Can also be written using the `<:` infix operator as *type1* `<:` *type2*.

```
julia> issubtype(Int8, Int32)
false

julia> Int8 <: Integer
true
```

`<:` (*T1, T2*)

Subtype operator, equivalent to `issubtype (T1, T2)`.

subtypes (*T::DataType*)

Return a list of immediate subtypes of *DataType* *T*. Note that all currently loaded subtypes are included, including those not visible in the current module.

```
julia> subtypes(Integer)
4-element Array{DataType,1}:
  BigInt
  Bool
  Signed
  Unsigned
```

typemin (*T*)

The lowest value representable by the given (real) numeric *DataType* *T*.

typemax(*T*)

The highest value representable by the given (real) numeric `DataType`.

realmin(*T*)

The smallest in absolute value non-subnormal value representable by the given floating-point `DataType` *T*.

realmax(*T*)

The highest finite value representable by the given floating-point `DataType` *T*.

maxintfloat(*T*)

The largest integer losslessly representable by the given floating-point `DataType` *T*.

sizeof(*T*)

Size, in bytes, of the canonical binary representation of the given `DataType` *T*, if any.

```
julia> sizeof(Float32)
4

julia> sizeof(Complex128)
16
```

If *T* is not a bitstype, an error is thrown.

```
julia> sizeof(Base.LinAlg.LU)
ERROR: argument is an abstract type; size is indeterminate
in sizeof(::Type{T}) at ./essentials.jl:99
...
```

eps(*T*)

The distance between 1.0 and the next larger representable floating-point value of `DataType` *T*. Only floating-point types are sensible arguments.

eps()

The distance between 1.0 and the next larger representable floating-point value of `Float64`.

eps(*x*)

The distance between *x* and the next larger representable floating-point value of the same `DataType` as *x*.

promote_type(*type1*, *type2*)

Determine a type big enough to hold values of each argument type without loss, whenever possible. In some cases, where no type exists to which both types can be promoted losslessly, some loss is tolerated; for example, `promote_type{Int64,Float64}` returns `Float64` even though strictly, not all `Int64` values can be represented exactly as `Float64` values.

promote_rule(*type1*, *type2*)

Specifies what type should be used by `promote()` (page 365) when given values of types *type1* and *type2*. This function should not be called directly, but should have definitions added to it for new types as appropriate.

getfield(*value*, *name::Symbol*)

Extract a named field from a value of composite type. The syntax `a.b` calls `getfield(a, :b)`.

setfield!(*value*, *name::Symbol*, *x*)

Assign *x* to a named field in *value* of composite type. The syntax `a.b = c` calls `setfield!(a, :b, c)`.

fieldoffset(*type*, *i*)

The byte offset of field *i* of a type relative to the data start. For example, we could use it in the following manner to summarize information about a struct type:

```
julia> structinfo(T) = [(fieldoffset(T,i), fieldname(T,i), fieldtype(T,i)) for i = 1:nfields(T)]

julia> structinfo(Base.Filesystem.StatStruct)
```

```
12-element Array{Tuple{UInt64,Symbol,DataType},1}:
 (0x0000000000000000, :device, UInt64)
 (0x0000000000000008, :inode, UInt64)
 (0x0000000000000010, :mode, UInt64)
 (0x0000000000000018, :nlink, Int64)
 (0x0000000000000020, :uid, UInt64)
 (0x0000000000000028, :gid, UInt64)
 (0x0000000000000030, :rdev, UInt64)
 (0x0000000000000038, :size, Int64)
 (0x0000000000000040, :blksize, Int64)
 (0x0000000000000048, :blocks, Int64)
 (0x0000000000000050, :mtime, Float64)
 (0x0000000000000058, :ctime, Float64)
```

fieldtype (*T*, *name::Symbol* | *index::Int*)

Determine the declared type of a field (specified by name or index) in a composite `DataType` *T*.

isimmutable (*v*)

Return `true` iff value *v* is immutable. See *Immutable Composite Types* (page 95) for a discussion of immutability. Note that this function works on values, so if you give it a type, it will tell you that a value of `DataType` is mutable.

isbits (*T*)

Return `true` if *T* is a “plain data” type, meaning it is immutable and contains no references to other values. Typical examples are numeric types such as `UInt8`, `Float64`, and `Complex{Float64}`.

```
julia> isbits(Complex{Float64})
true

julia> isbits(Complex)
false
```

isleafytype (*T*)

Determine whether *T* is a concrete type that can have instances, meaning its only subtypes are itself and `Union{}` (but *T* itself is not `Union{}`).

typejoin (*T*, *S*)

Compute a type that contains both *T* and *S*.

typeintersect (*T*, *S*)

Compute a type that contains the intersection of *T* and *S*. Usually this will be the smallest such type or one close to it.

Val{c}()

Create a “value type” out of *c*, which must be an `isbits` value. The intent of this construct is to be able to dispatch on constants, e.g., `f(Val{false})` allows you to dispatch directly (at compile-time) to an implementation `f(::Type{Val{false}})`, without having to test the boolean value at runtime.

@enum EnumName EnumValue1[=x] EnumValue2[=y]

Create an `Enum` type with name `EnumName` and enum member values of `EnumValue1` and `EnumValue2` with optional assigned values of *x* and *y*, respectively. `EnumName` can be used just like other types and enum member values as regular values, such as

```
julia> @enum FRUIT apple=1 orange=2 kiwi=3

julia> f(x::FRUIT) = "I'm a FRUIT with value: $(Int(x))"
f (generic function with 1 method)

julia> f(apple)
"I'm a FRUIT with value: 1"
```

instances ($T::Type$)

Return a collection of all instances of the given type, if applicable. Mostly used for enumerated types (see `@enum`).

40.5 Generic Functions

method_exists ($f, Tuple\ type$) $\rightarrow Bool$

Determine whether the given generic function has a method matching the given `Tuple` of argument types.

```
julia> method_exists(length, Tuple{Array})
true
```

applicable ($f, args\dots$) $\rightarrow Bool$

Determine whether the given generic function has a method applicable to the given arguments.

```
julia> function f(x, y)
           x + y
       end;

julia> applicable(f, 1)
false

julia> applicable(f, 1, 2)
true
```

invoke ($f, (types\dots), args\dots$)

Invoke a method for the given generic function matching the specified types (as a tuple), on the specified arguments. The arguments must be compatible with the specified types. This allows invoking a method other than the most specific matching method, which is useful when the behavior of a more general definition is explicitly needed (often as part of the implementation of a more specific method of the same function).

|> (x, f)

Applies a function to the preceding argument. This allows for easy function chaining.

```
julia> [1:5;] |> x->x.^2 |> sum |> inv
0.01818181818181818
```

40.6 Syntax

eval ($[m::Module], expr::Expr$)

Evaluate an expression in the given module and return the result. Every `Module` (except those defined with `baremodule`) has its own 1-argument definition of `eval`, which evaluates expressions in that module.

@eval ()

Evaluate an expression and return the value.

evalfile ($path::AbstractString, args::Vector{String}=String[]$)

Load the file using `include()` (page 361), evaluate all expressions, and return the value of the last one.

esc ($e::ANY$)

Only valid in the context of an `Expr` returned from a macro. Prevents the macro hygiene pass from turning embedded variables into gensym variables. See the *Macros* (page 166) section of the Metaprogramming chapter of the manual for more details and examples.

gensym(*[tag]*)

Generates a symbol which will not conflict with other variable names.

@gensym()

Generates a gensym symbol for a variable. For example, `@gensym x y` is transformed into `x = gensym("x"); y = gensym("y")`.

@polly()

Tells the compiler to apply the polyhedral optimizer Polly to a function.

parse(*str, start; greedy=true, raise=true*)

Parse the expression string and return an expression (which could later be passed to `eval` for execution). *start* is the index of the first character to start parsing. If *greedy* is `true` (default), `parse` will try to consume as much input as it can; otherwise, it will stop as soon as it has parsed a valid expression. Incomplete but otherwise syntactically valid expressions will return `Expr(:incomplete, "(error message)")`. If *raise* is `true` (default), syntax errors other than incomplete expressions will raise an error. If *raise* is `false`, `parse` will return an expression that will raise an error upon evaluation.

parse(*str; raise=true*)

Parse the expression string greedily, returning a single expression. An error is thrown if there are additional characters after the first expression. If *raise* is `true` (default), syntax errors will raise an error; otherwise, `parse` will return an expression that will raise an error upon evaluation.

40.7 Nullables

Nullable(*x, hasvalue::Bool=true*)

Wrap value *x* in an object of type `Nullable`, which indicates whether a value is present. `Nullable(x)` yields a non-empty wrapper and `Nullable{T}()` yields an empty instance of a wrapper that might contain a value of type *T*.

`Nullable(x, false)` yields `Nullable{typeof(x)}()` with *x* stored in the result's value field.

Examples

```
julia> Nullable(1)
Nullable{Int64}(1)

julia> Nullable{Int64}()
Nullable{Int64}()

julia> Nullable(1, false)
Nullable{Int64}()

julia> dump(Nullable(1, false))
Nullable{Int64}
  hasvalue: Bool false
  value: Int64 1
```

get(*x::Nullable{y}*)

Attempt to access the value of *x*. Returns the value if it is present; otherwise, returns *y* if provided, or throws a `NullException` if not.

isnull(*x*)

Return whether or not *x* is null for `Nullable` (page 369) *x*; return `false` for all other *x*.

Examples

```
julia> x = Nullable{1, false}
Nullable{Int64}()

julia> isnull(x)
true

julia> x = Nullable{1, true}
Nullable{Int64}(1)

julia> isnull(x)
false

julia> x = 1
1

julia> isnull(x)
false
```

unsafe_get(*x*)

Return the value of *x* for *Nullable* (page 369) *x*; return *x* for all other *x*.

This method does not check whether or not *x* is null before attempting to access the value of *x* for *x::Nullable* (hence “unsafe”).

```
julia> x = Nullable{1}
Nullable{Int64}(1)

julia> unsafe_get(x)
1

julia> x = Nullable{String}()
Nullable{String}()

julia> unsafe_get(x)
ERROR: UndefRefError: access to undefined reference
  in unsafe_get(::Nullable{String}) at ./nullable.jl:124
...

julia> x = 1
1

julia> unsafe_get(x)
1
```

40.8 System

run(*command*, *args...*)

Run a command object, constructed with backticks. Throws an error if anything goes wrong, including the process exiting with a non-zero status.

spawn(*command*)

Run a command object asynchronously, returning the resulting `Process` object.

DevNull

Used in a stream redirect to discard all data written to it. Essentially equivalent to `/dev/null` on Unix or `NUL` on Windows. Usage:


```
run(pipeline(`cat test.txt`, DevNull))
```

success (*command*)

Run a command object, constructed with backticks, and tell whether it was successful (exited with a code of 0). An exception is raised if the process cannot be started.

process_running (*p::Process*)

Determine whether a process is currently running.

process_exited (*p::Process*)

Determine whether a process has exited.

kill (*p::Process*, *signum=SIGTERM*)

Send a signal to a process. The default is to terminate the process.

Sys.set_process_title (*title::AbstractString*)

Set the process title. No-op on some operating systems. (not exported)

Sys.get_process_title ()

Get the process title. On some systems, will always return empty string. (not exported)

readandwrite (*command*)

Starts running a command asynchronously, and returns a tuple (stdout,stdin,process) of the output stream and input stream of the process, and the process object itself.

ignorestatus (*command*)

Mark a command object so that running it will not throw an error if the result code is non-zero.

detach (*command*)

Mark a command object so that it will be run in a new process group, allowing it to outlive the julia process, and not have Ctrl-C interrupts passed to it.

Cmd(*cmd::Cmd*; *ignorestatus*, *detach*, *windows_verbatim*, *windows_hide*, *env*, *dir*)

Construct a new Cmd object, representing an external program and arguments, from *cmd*, while changing the settings of the optional keyword arguments:

- *ignorestatus::Bool*: If *true* (defaults to *false*), then the Cmd will not throw an error if the return code is nonzero.
- *detach::Bool*: If *true* (defaults to *false*), then the Cmd will be run in a new process group, allowing it to outlive the julia process and not have Ctrl-C passed to it.
- *windows_verbatim::Bool*: If *true* (defaults to *false*), then on Windows the Cmd will send a command-line string to the process with no quoting or escaping of arguments, even arguments containing spaces. (On Windows, arguments are sent to a program as a single “command-line” string, and programs are responsible for parsing it into arguments. By default, empty arguments and arguments with spaces or tabs are quoted with double quotes " in the command line, and \ or " are preceded by backslashes. *windows_verbatim=true* is useful for launching programs that parse their command line in non-standard ways.) Has no effect on non-Windows systems.
- *windows_hide::Bool*: If *true* (defaults to *false*), then on Windows no new console window is displayed when the Cmd is executed. This has no effect if a console is already open or on non-Windows systems.
- *env*: Set environment variables to use when running the Cmd. *env* is either a dictionary mapping strings to strings, an array of strings of the form "var=val", an array or tuple of "var"=>val pairs, or nothing. In order to modify (rather than replace) the existing environment, create *env* by copy (ENV) and then set *env["var"]*=val as desired.

- `dir::AbstractString`: Specify a working directory for the command (instead of the current directory).

For any keywords that are not specified, the current settings from `cmd` are used. Normally, to create a `Cmd` object in the first place, one uses backticks, e.g.

```
Cmd(`echo "Hello world"`, ignorestatus=true, detach=false)
```

setenv (*command::Cmd, env; dir=""*)

Set environment variables to use when running the given `command`. `env` is either a dictionary mapping strings to strings, an array of strings of the form `"var=val"`, or zero or more `"var"=>val` pair arguments. In order to modify (rather than replace) the existing environment, create `env` by `copy(ENV)` and then setting `env["var"]=val` as desired, or use `withenv`.

The `dir` keyword argument can be used to specify a working directory for the command.

withenv (*f::Function, kv::Pair...*)

Execute `f()` in an environment that is temporarily modified (not replaced as in `setenv`) by zero or more `"var"=>val` arguments `kv`. `withenv` is generally used via the `withenv(kv...) do ... end` syntax. A value of `nothing` can be used to temporarily unset an environment variable (if it is set). When `withenv` returns, the original environment has been restored.

pipeline (*from, to, ...*)

Create a pipeline from a data source to a destination. The source and destination can be commands, I/O streams, strings, or results of other pipeline calls. At least one argument must be a command. Strings refer to filenames. When called with more than two arguments, they are chained together from left to right. For example `pipeline(a,b,c)` is equivalent to `pipeline(pipeline(a,b),c)`. This provides a more concise way to specify multi-stage pipelines.

Examples:

```
run(pipeline(`ls`, `grep xyz`))
run(pipeline(`ls`, "out.txt"))
run(pipeline("out.txt", `grep xyz`))
```

pipeline (*command; stdin, stdout, stderr, append=false*)

Redirect I/O to or from the given `command`. Keyword arguments specify which of the command's streams should be redirected. `append` controls whether file output appends to the file. This is a more general version of the 2-argument pipeline function. `pipeline(from, to)` is equivalent to `pipeline(from, stdout=to)` when `from` is a command, and to `pipeline(to, stdin=from)` when `from` is another kind of data source.

Examples:

```
run(pipeline(`dothings`, stdout="out.txt", stderr="errs.txt"))
run(pipeline(`update`, stdout="log.txt", append=true))
```

gethostname () → `AbstractString`

Get the local machine's host name.

getipaddr () → `IPAddr`

Get the IP address of the local machine.

getpid () → `Int32`

Get Julia's process ID.

time ()

Get the system time in seconds since the epoch, with fairly high (typically, microsecond) resolution.

time_ns ()

Get the time in nanoseconds. The time corresponding to 0 is undefined, and wraps every 5.8 years.

tic()

Set a timer to be read by the next call to `toc()` (page 373) or `toq()` (page 373). The macro call `@time expr` can also be used to time evaluation.

toc()

Print and return the time elapsed since the last `tic()` (page 372). The macro call `@time expr` can also be used to time evaluation.

toq()

Return, but do not print, the time elapsed since the last `tic()` (page 372). The macro calls `@timed expr` and `@elapsed expr` also return evaluation time.

@time()

A macro to execute an expression, printing the time it took to execute, the number of allocations, and the total number of bytes its execution caused to be allocated, before returning the value of the expression.

See also `@timev()` (page 373), `@timed()` (page 373), `@elapsed()` (page 373), and `@allocated()` (page 373).

@timev()

This is a verbose version of the `@time` macro. It first prints the same information as `@time`, then any non-zero memory allocation counters, and then returns the value of the expression.

See also `@time()` (page 373), `@timed()` (page 373), `@elapsed()` (page 373), and `@allocated()` (page 373).

@timed()

A macro to execute an expression, and return the value of the expression, elapsed time, total bytes allocated, garbage collection time, and an object with various memory allocation counters.

See also `@time()` (page 373), `@timev()` (page 373), `@elapsed()` (page 373), and `@allocated()` (page 373).

@elapsed()

A macro to evaluate an expression, discarding the resulting value, instead returning the number of seconds it took to execute as a floating-point number.

See also `@time()` (page 373), `@timev()` (page 373), `@timed()` (page 373), and `@allocated()` (page 373).

@allocated()

A macro to evaluate an expression, discarding the resulting value, instead returning the total number of bytes allocated during evaluation of the expression. Note: the expression is evaluated inside a local function, instead of the current context, in order to eliminate the effects of compilation, however, there still may be some allocations due to JIT compilation. This also makes the results inconsistent with the `@time` macros, which do not try to adjust for the effects of compilation.

See also `@time()` (page 373), `@timev()` (page 373), `@timed()` (page 373), and `@elapsed()` (page 373).

EnvHash() → EnvHash

A singleton of this type provides a hash table interface to environment variables.

ENV

Reference to the singleton EnvHash, providing a dictionary interface to system environment variables.

is_unix([os])

Predicate for testing if the OS provides a Unix-like interface. See documentation in *Handling Operating System Variation* (page 255).

is_apple([os])

Predicate for testing if the OS is a derivative of Apple Macintosh OS X or Darwin. See documentation in *Handling Operating System Variation* (page 255).

is_linux(*[os]*)

Predicate for testing if the OS is a derivative of Linux. See documentation in *Handling Operating System Variation* (page 255).

is_bsd(*[os]*)

Predicate for testing if the OS is a derivative of BSD. See documentation in *Handling Operating System Variation* (page 255).

is_windows(*[os]*)

Predicate for testing if the OS is a derivative of Microsoft Windows NT. See documentation in *Handling Operating System Variation* (page 255).

windows_version()

Returns the version number for the Windows NT Kernel as a (major, minor) pair, or (0, 0) if this is not running on Windows.

@static()

Partially evaluates an expression at parse time.

For example, `@static is_windows() ? foo : bar` will evaluate `is_windows()` and insert either `foo` or `bar` into the expression. This is useful in cases where a construct would be invalid on other platforms, such as a `ccall` to a non-existent function.

40.9 Errors

error(*message::AbstractString*)

Raise an `ErrorException` with the given message.

throw(*e*)

Throw an object as an exception.

rethrow(*[e]*)

Throw an object without changing the current exception backtrace. The default argument is the current exception (if called within a `catch` block).

backtrace()

Get a backtrace object for the current program point.

catch_backtrace()

Get the backtrace of the current exception, for use within `catch` blocks.

assert(*cond*)

Throw an `AssertionError` (page 374) if *cond* is false. Also available as the macro `@assert expr`.

@assert cond [text]

Throw an `AssertionError` if *cond* is false. Preferred syntax for writing assertions. Message *text* is optionally displayed upon assertion failure.

ArgumentError(*msg*)

The parameters to a function call do not match a valid signature. Argument *msg* is a descriptive error string.

AssertionError(*[msg]*)

The asserted condition did not evaluate to `true`. Optional argument *msg* is a descriptive error string.

BoundsError(*[a][, i]*)

An indexing operation into an array, *a*, tried to access an out-of-bounds element, *i*.

DimensionMismatch(*[msg]*)

The objects called do not have matching dimensionality. Optional argument *msg* is a descriptive error string.

DivideError ()

Integer division was attempted with a denominator value of 0.

DomainError ()

The arguments to a function or constructor are outside the valid domain.

EOFError ()

No more data was available to read from a file or stream.

ErrorException (*msg*)

Generic error type. The error message, in the `.msg` field, may provide more specific details.

InexactError ()

Type conversion cannot be done exactly.

InterruptException ()

The process was stopped by a terminal interrupt (CTRL+C).

KeyError (*key*)

An indexing operation into an Associative (Dict) or Set like object tried to access or delete a non-existent element.

LoadError (*file::AbstractString, line::Int, error*)

An error occurred while including, requiring, or using a file. The error specifics should be available in the `.error` field.

MethodError (*f, args*)

A method with the required type signature does not exist in the given generic function. Alternatively, there is no unique most-specific method.

NullException ()

An attempted access to a [Nullable](#) (page 369) with no defined value.

OutOfMemoryError ()

An operation allocated too much memory for either the system or the garbage collector to handle properly.

ReadOnlyMemoryError ()

An operation tried to write to memory that is read-only.

OverflowError ()

The result of an expression is too large for the specified type and will cause a wraparound.

ParseError (*msg*)

The expression passed to the `parse` function could not be interpreted as a valid Julia expression.

ProcessExitedException ()

After a client Julia process has exited, further attempts to reference the dead child will throw this exception.

StackOverflowError ()

The function call grew beyond the size of the call stack. This usually happens when a call recurses infinitely.

SystemError (*prefix::AbstractString* [, *errno::Int32*])

A system call failed with an error code (in the `errno` global variable).

TypeError (*func::Symbol, context::AbstractString, expected::Type, got*)

A type assertion failure, or calling an intrinsic function with an incorrect argument type.

UndefRefError ()

The item or field is not defined for the given object.

UndefVarError (*var::Symbol*)

A symbol in the current scope is not defined.

InitError (*mod::Symbol, error*)

An error occurred when running a module's `__init__` function. The actual error thrown is available in the `.error` field.

retry (*f, [retry_on]; n=1, max_delay=10.0*) → Function

Returns a lambda that retries function `f` up to `n` times in the event of an exception. If `retry_on` is a `Type` then retry only for exceptions of that type. If `retry_on` is a function `test_error(::Exception) -> Bool` then retry only if it is true.

The first retry happens after a gap of 50 milliseconds or `max_delay`, whichever is lower. Subsequently, the delays between retries are exponentially increased with a random factor up to `max_delay`.

Examples

```
retry(http_get, e -> e.status == "503")(url)
retry(read, UVEError)(io)
```

40.10 Events

Timer (*callback::Function, delay, repeat=0*)

Create a timer to call the given `callback` function. The `callback` is passed one argument, the timer object itself. The callback will be invoked after the specified initial `delay`, and then repeating with the given `repeat` interval. If `repeat` is 0, the timer is only triggered once. Times are in seconds. A timer is stopped and has its resources freed by calling `close()` (page 487) on it.

Timer (*delay, repeat=0*)

Create a timer that wakes up tasks waiting for it (by calling `wait()` (page 490) on the timer object) at a specified interval. Times are in seconds. Waiting tasks are woken with an error when the timer is closed (by `close()` (page 487)). Use `isopen()` (page 564) to check whether a timer is still active.

AsyncCondition ()

Create an async condition that wakes up tasks waiting for it (by calling `wait()` (page 490) on the object) when notified from C by a call to `uv_async_send`. Waiting tasks are woken with an error when the object is closed (by `close()` (page 487)). Use `isopen()` (page 564) to check whether it is still active.

AsyncCondition (*callback::Function*)

Create an async condition that calls the given `callback` function. The `callback` is passed one argument, the async condition object itself.

40.11 Reflection

module_name (*m::Module*) → Symbol

Get the name of a `Module` as a `Symbol`.

```
julia> module_name(Base.LinAlg)
:LinAlg
```

module_parent (*m::Module*) → Module

Get a module's enclosing `Module`. `Main` is its own parent, as is `LastMain` after `workspace()`.

```
julia> module_parent(Main)
Main

julia> module_parent(Base.LinAlg.BLAS)
Base.LinAlg
```

current_module() → Module

Get the *dynamically* current Module, which is the Module code is currently being read from. In general, this is not the same as the module containing the call to this function.

fullname(*m::Module*)

Get the fully-qualified name of a module as a tuple of symbols. For example,

```
julia> fullname(Base.Pkg)
(:Base, :Pkg)

julia> fullname(Main)
()
```

names(*x::Module*, *all::Bool=false*, *imported::Bool=false*)

Get an array of the names exported by a Module, excluding deprecated names. If *all* is true, then the list also includes non-exported names defined in the module, deprecated names, and compiler-generated names. If *imported* is true, then names explicitly imported from other modules are also included.

As a special case, all names defined in `Main` are considered “exported”, since it is not idiomatic to explicitly export names from `Main`.

nfields(*x::DataType*) → Int

Get the number of fields of a DataType.

fieldnames(*x::DataType*)

Get an array of the fields of a DataType.

```
julia> fieldnames(Hermitian)
2-element Array{Symbol,1}:
 :data
 :uplo
```

fieldname(*x::DataType*, *i::Integer*)

Get the name of field *i* of a DataType.

```
julia> fieldname(SparseMatrixCSC, 1)
:m

julia> fieldname(SparseMatrixCSC, 5)
:nzval
```

Base.datatype_module(*t::DataType*) → Module

Determine the module containing the definition of a DataType.

Base.datatype_name(*t::DataType*) → Symbol

Get the name of a DataType (without its parent module) as a symbol.

isconst(*[m::Module]*, *s::Symbol*) → Bool

Determine whether a global is declared `const` in a given Module. The default Module argument is `current_module()` (page 376).

Base.function_name(*f::Function*) → Symbol

Get the name of a generic Function as a symbol, or `:anonymous`.

Base.function_module(*f::Function*) → Module

Determine the module containing the (first) definition of a generic function.

Base.function_module(*f::Function*, *types*) → Module

Determine the module containing a given definition of a generic function.

functionloc(*f::Function*, *types*)

Returns a tuple (*filename*, *line*) giving the location of a generic Function definition.

functionloc (*m::Method*)

Returns a tuple (*filename*, *line*) giving the location of a Method definition.

@functionloc ()

Applied to a function or macro call, it evaluates the arguments to the specified call, and returns a tuple (*filename*, *line*) giving the location for the method that would be called for those arguments. It calls out to the `functionloc` function.

40.12 Internals

gc ()

Perform garbage collection. This should not generally be used.

gc_enable (*on::Bool*)

Control whether garbage collection is enabled using a boolean argument (`true` for enabled, `false` for disabled). Returns previous GC state. Disabling garbage collection should be used only with extreme caution, as it can cause memory use to grow without bound.

macroexpand (*x*)

Takes the expression *x* and returns an equivalent expression with all macros removed (expanded).

@macroexpand ()

Return equivalent expression with all macros removed (expanded).

There is a subtle difference between `@macroexpand` and `macroexpand` in that expansion takes place in different contexts. This is best seen in the following example:

```
julia> module M
           macro m()
               1
           end
           function f()
               (@macroexpand(@m), macroexpand(:(@m)))
           end
       end
M

julia> macro m()
           2
       end
@m (macro with 1 method)

julia> M.f()
(1,2)
```

With `@macroexpand` the expression expands where `@macroexpand` appears in the code (module `M` in the example). With `macroexpand` the expression expands in the current module where the code was finally called (REPL in the example). Note that when calling `macroexpand` or `@macroexpand` directly from the REPL, both of these contexts coincide, hence there is no difference.

expand (*x*)

Takes the expression *x* and returns an equivalent expression in lowered form. See also `code_lowered()` (page 378).

code_lowered (*f*, *types*)

Returns an array of lowered ASTs for the methods matching the given generic function and type signature.

@code_lowered()

Evaluates the arguments to the function or macro call, determines their types, and calls `code_lowered()` (page 378) on the resulting expression.

code_typed(*f*, *types*; *optimize=true*)

Returns an array of lowered and type-inferred ASTs for the methods matching the given generic function and type signature. The keyword argument `optimize` controls whether additional optimizations, such as inlining, are also applied.

@code_typed()

Evaluates the arguments to the function or macro call, determines their types, and calls `code_typed()` (page 379) on the resulting expression.

code_warntype(*[io::IO]*, *f*, *types*)

Prints lowered and type-inferred ASTs for the methods matching the given generic function and type signature to `io` which defaults to `STDOUT`. The ASTs are annotated in such a way as to cause “non-leaf” types to be emphasized (if color is available, displayed in red). This serves as a warning of potential type instability. Not all non-leaf types are particularly problematic for performance, so the results need to be used judiciously. See [@code_warntype](#) (page 323) for more information.

@code_warntype()

Evaluates the arguments to the function or macro call, determines their types, and calls `code_warntype()` (page 379) on the resulting expression.

code_llvm(*[io]*, *f*, *types*)

Prints the LLVM bitcodes generated for running the method matching the given generic function and type signature to `io` which defaults to `STDOUT`.

All metadata and `dbg.*` calls are removed from the printed bitcode. Use `code_llvm_raw` for the full IR.

@code_llvm()

Evaluates the arguments to the function or macro call, determines their types, and calls `code_llvm()` (page 379) on the resulting expression.

code_native(*[io]*, *f*, *types*[, *syntax*])

Prints the native assembly instructions generated for running the method matching the given generic function and type signature to `io` which defaults to `STDOUT`. Switch assembly syntax using `syntax` symbol parameter set to `:att` for AT&T syntax or `:intel` for Intel syntax. Output is AT&T syntax by default.

@code_native()

Evaluates the arguments to the function or macro call, determines their types, and calls `code_native()` (page 379) on the resulting expression.

precompile(*f*, *args::Tuple{Vararg{Any}}*)

Compile the given function `f` for the argument tuple (of types) `args`, but do not execute it.

Collections and Data Structures

41.1 Iteration

Sequential iteration is implemented by the methods `start()` (page 381), `done()` (page 381), and `next()` (page 381). The general `for` loop:

```
for i = I    # or "for i in I"
  # body
end
```

is translated into:

```
state = start(I)
while !done(I, state)
  (i, state) = next(I, state)
  # body
end
```

The `state` object may be anything, and should be chosen appropriately for each iterable type. See the [manual section on the iteration interface](#) (page 133) for more details about defining a custom iterable type.

start (*iter*) → *state*

Get initial iteration state for an iterable object.

done (*iter, state*) → *Bool*

Test whether we are done iterating.

next (*iter, state*) → *item, state*

For a given iterable object and iteration state, return the current item and the next iteration state.

iterator_size (*iter_type::Type*) → *IteratorSize*

Given the type of an iterator, returns one of the following values:

- `SizeUnknown()` if the length (number of elements) cannot be determined in advance.
- `HasLength()` if there is a fixed, finite length.
- `HasShape()` if there is a known length plus a notion of multidimensional shape (as for an array). In this case the `size` function is valid for the iterator.
- `IsInfinite()` if the iterator yields values forever.

The default value (for iterators that do not define this function) is `HasLength()`. This means that most iterators are assumed to implement `length`.

This trait is generally used to select between algorithms that pre-allocate space for their result, and algorithms that resize their result incrementally.

iteratoreltype (*itertype::Type*) → *IteratorEltype*

Given the type of an iterator, returns one of the following values:

- `EltypeUnknown()` if the type of elements yielded by the iterator is not known in advance.
- `HasEltype()` if the element type is known, and `eltype` would return a meaningful value.

`HasEltype()` is the default, since iterators are assumed to implement `eltype`.

This trait is generally used to select between algorithms that pre-allocate a specific type of result, and algorithms that pick a result type based on the types of yielded values.

Fully implemented by:

- `Range`
- `UnitRange`
- `Tuple`
- `Number`
- `AbstractArray`
- `IntSet` (page 396)
- `ObjectIdDict`
- `Dict` (page 394)
- `WeakKeyDict`
- `EachLine`
- `AbstractString`
- `Set` (page 396)
- `Task` (page 485)

41.2 General Collections

isempty (*collection*) → `Bool`

Determine whether a collection is empty (has no elements).

```
julia> isempty([])
true

julia> isempty([1 2 3])
false
```

empty! (*collection*) → *collection*

Remove all elements from a collection.

```
julia> A = Dict{"a" => 1, "b" => 2}
Dict{String, Int64} with 2 entries:
  "b" => 2
  "a" => 1

julia> empty!(A);
```

```
julia> A
Dict{String,Int64} with 0 entries
```

length (*collection*) → Integer

For ordered, indexable collections, returns the maximum index *i* for which `getindex(collection, i)` is valid. For unordered collections, returns the number of elements.

```
julia> length(1:5)
5

julia> length([1; 2; 3; 4])
4
```

endof (*collection*) → Integer

Returns the last index of the collection.

```
julia> endof([1,2,4])
3
```

Fully implemented by:

- Range
- UnitRange
- Tuple
- Number
- AbstractArray
- *IntSet* (page 396)
- *Dict* (page 394)
- WeakKeyDict
- AbstractString
- *Set* (page 396)

41.3 Iterable Collections

in (*item*, *collection*) → Bool

\in (*item*, *collection*) → Bool

\ni (*collection*, *item*) → Bool

\notin (*item*, *collection*) → Bool

\nexists (*collection*, *item*) → Bool

Determine whether an item is in the given collection, in the sense that it is `==` to one of the values generated by iterating over the collection. Some collections need a slightly different definition; for example *Set* (page 396)s check whether the item *isequal*() (page 362) to one of the elements. *Dict* (page 394)s look for (key, value) pairs, and the key is compared using *isequal*() (page 362). To test for the presence of a key in a dictionary, use *haskey*() (page 394) or `k in keys(dict)`.

```
julia> a = 1:3:20
1:3:19

julia> 4 in a
true
```

```
julia> 5 in a
false
```

eltype(*type*)

Determine the type of the elements generated by iterating a collection of the given *type*. For associative collection types, this will be a `Pair{KeyType,ValType}`. The definition `eltype(x) = eltype(typeof(x))` is provided for convenience so that instances can be passed instead of types. However the form that accepts a type argument should be defined for new types.

indexin(*a, b*)

Returns a vector containing the highest index in *b* for each value in *a* that is a member of *b*. The output vector contains 0 wherever *a* is not a member of *b*.

```
julia> a = ['a', 'b', 'c', 'b', 'd', 'a'];

julia> b = ['a','b','c'];

julia> indexin(a,b)
6-element Array{Int64,1}:
 1
 2
 3
 2
 0
 1

julia> indexin(b,a)
3-element Array{Int64,1}:
 6
 4
 3
```

findin(*a, b*)

Returns the indices of elements in collection *a* that appear in collection *b*.

```
julia> a = collect(1:3:15)
5-element Array{Int64,1}:
 1
 4
 7
10
13

julia> b = collect(2:4:10)
3-element Array{Int64,1}:
 2
 6
10

julia> findin(a,b) # 10 is the only common element
1-element Array{Int64,1}:
 4
```

unique(*itr*[, *dim*])

Returns an array containing only the unique elements of the iterable *itr*, in the order that the first of each set of equivalent elements originally appears. If *dim* is specified, returns unique regions of the array *itr* along *dim*.

unique(*itr*)

Returns an array containing one value from *itr* for each unique value, as determined by `isequal`.

unique (*f*, *itr*)

Returns an array containing one value from *itr* for each unique value produced by *f* applied to elements of *itr*.

allunique (*itr*)

Return `true` if all values from *itr* are distinct when compared with `isequal`.

reduce (*op*, *v0*, *itr*)

Reduce the given collection *itr* with the given binary operator *op*. *v0* must be a neutral element for *op* that will be returned for empty collections. It is unspecified whether *v0* is used for non-empty collections.

Reductions for certain commonly-used operators have special implementations which should be used instead: `maximum(itr)`, `minimum(itr)`, `sum(itr)`, `prod(itr)`, `any(itr)`, `all(itr)`.

The associativity of the reduction is implementation dependent. This means that you can't use non-associative operations like `-` because it is undefined whether `reduce(-, [1, 2, 3])` should be evaluated as $(1-2)-3$ or $1-(2-3)$. Use `foldl()` (page 385) or `foldr()` (page 385) instead for guaranteed left or right associativity.

Some operations accumulate error, and parallelism will also be easier if the reduction can be executed in groups. Future versions of Julia might change the algorithm. Note that the elements are not reordered if you use an ordered collection.

reduce (*op*, *itr*)

Like `reduce(op, v0, itr)`. This cannot be used with empty collections, except for some special cases (e.g. when *op* is one of `+`, `*`, `max`, `min`, `&`, `|`) when Julia can determine the neutral element of *op*.

foldl (*op*, *v0*, *itr*)

Like `reduce()` (page 385), but with guaranteed left associativity. *v0* will be used exactly once.

foldl (*op*, *itr*)

Like `foldl(op, v0, itr)`, but using the first element of *itr* as *v0*. In general, this cannot be used with empty collections (see `reduce(op, itr)`).

foldr (*op*, *v0*, *itr*)

Like `reduce()` (page 385), but with guaranteed right associativity. *v0* will be used exactly once.

foldr (*op*, *itr*)

Like `foldr(op, v0, itr)`, but using the last element of *itr* as *v0*. In general, this cannot be used with empty collections (see `reduce(op, itr)`).

maximum (*itr*)

Returns the largest element in a collection.

```
julia> maximum(-20.5:10)
9.5

julia> maximum([1, 2, 3])
3
```

maximum (*A*, *dims*)

Compute the maximum value of an array over the given dimensions. See also the `max()` (page 418) function to take the maximum of two or more arguments, which can be applied elementwise to arrays via `max.(a, b)`.

maximum! (*r*, *A*)

Compute the maximum value of *A* over the singleton dimensions of *r*, and write results to *r*.

minimum (*itr*)

Returns the smallest element in a collection.

```
julia> minimum(-20.5:10)
-20.5
```

```
julia> minimum([1,2,3])
1
```

minimum(*A*, *dims*)

Compute the minimum value of an array over the given dimensions. See also the [min\(\)](#) (page 417) function to take the minimum of two or more arguments, which can be applied elementwise to arrays via `min.(a,b)`.

```
julia> A = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> minimum(A, 1)
1×2 Array{Int64,2}:
 1  2

julia> minimum(A, 2)
2×1 Array{Int64,2}:
 1
 3
```

minimum!(*r*, *A*)

Compute the minimum value of *A* over the singleton dimensions of *r*, and write results to *r*.

extrema(*itr*) → Tuple

Compute both the minimum and maximum element in a single pass, and return them as a 2-tuple.

```
julia> extrema(2:10)
(2,10)

julia> extrema([9,pi,4.5])
(3.141592653589793,9.0)
```

extrema(*A*, *dims*) → Array{Tuple}

Compute the minimum and maximum elements of an array over the given dimensions.

indmax(*itr*) → Integer

Returns the index of the maximum element in a collection. If there are multiple maximal elements, then the first one will be returned. NaN values are ignored, unless all elements are NaN.

The collection must not be empty.

```
julia> indmax([8,0.1,-9,pi])
1

julia> indmax([1,7,7,6])
2

julia> indmax([1,7,7,NaN])
2
```

indmin(*itr*) → Integer

Returns the index of the minimum element in a collection. If there are multiple minimal elements, then the first one will be returned. NaN values are ignored, unless all elements are NaN.

The collection must not be empty.

```
julia> indmin([8,0.1,-9,pi])
3
```



```
julia> indmin([7,1,1,6])
2

julia> indmin([7,1,1,NaN])
2
```

findmax (*itr*) → (x, index)

Returns the maximum element of the collection *itr* and its index. If there are multiple maximal elements, then the first one will be returned. NaN values are ignored, unless all elements are NaN.

The collection must not be empty.

```
julia> findmax([8,0.1,-9,pi])
(8.0,1)

julia> findmax([1,7,7,6])
(7,2)

julia> findmax([1,7,7,NaN])
(7.0,2)
```

findmax (*A*, *region*) → (maxval, index)

For an array input, returns the value and index of the maximum over the given region.

findmin (*itr*) → (x, index)

Returns the minimum element of the collection *itr* and its index. If there are multiple minimal elements, then the first one will be returned. NaN values are ignored, unless all elements are NaN.

The collection must not be empty.

```
julia> findmin([8,0.1,-9,pi])
(-9.0,3)

julia> findmin([7,1,1,6])
(1,2)

julia> findmin([7,1,1,NaN])
(1.0,2)
```

findmin (*A*, *region*) → (minval, index)

For an array input, returns the value and index of the minimum over the given region.

findmax! (*rval*, *rind*, *A* [, *init=true*]) → (maxval, index)

Find the maximum of *A* and the corresponding linear index along singleton dimensions of *rval* and *rind*, and store the results in *rval* and *rind*.

findmin! (*rval*, *rind*, *A* [, *init=true*]) → (minval, index)

Find the minimum of *A* and the corresponding linear index along singleton dimensions of *rval* and *rind*, and store the results in *rval* and *rind*.

maxabs (*itr*)

Compute the maximum absolute value of a collection of values.

```
julia> maxabs([-1, 3, 4*im])
4.0
```

maxabs (*A*, *dims*)

Compute the maximum absolute values over given dimensions.

maxabs! (*r*, *A*)

Compute the maximum absolute values over the singleton dimensions of *r*, and write values to *r*.

minabs (*itr*)

Compute the minimum absolute value of a collection of values.

```
julia> minabs([-1, 3, 4*im])
1.0
```

minabs (*A*, *dims*)

Compute the minimum absolute values over given dimensions.

minabs! (*r*, *A*)

Compute the minimum absolute values over the singleton dimensions of *r*, and write values to *r*.

sum (*itr*)

Returns the sum of all elements in a collection.

sum (*A*, *dims*)

Sum elements of an array over the given dimensions.

```
julia> A = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> sum(A, 1)
1×2 Array{Int64,2}:
 4  6

julia> sum(A, 2)
2×1 Array{Int64,2}:
 3
 7
```

sum! (*r*, *A*)

Sum elements of *A* over the singleton dimensions of *r*, and write results to *r*.

sum (*f*, *itr*)

Sum the results of calling function *f* on each element of *itr*.

sumabs (*itr*)

Sum absolute values of all elements in a collection. This is equivalent to `sum(abs(itr))` but faster.

sumabs (*A*, *dims*)

Sum absolute values of elements of an array over the given dimensions.

sumabs! (*r*, *A*)

Sum absolute values of elements of *A* over the singleton dimensions of *r*, and write results to *r*.

sumabs2 (*itr*)

Sum squared absolute values of all elements in a collection. This is equivalent to `sum(abs2(itr))` but faster.

sumabs2 (*A*, *dims*)

Sum squared absolute values of elements of an array over the given dimensions.

sumabs2! (*r*, *A*)

Sum squared absolute values of elements of *A* over the singleton dimensions of *r*, and write results to *r*.

prod (*itr*)

Returns the product of all elements of a collection.

prod (*A*, *dims*)

Multiply elements of an array over the given dimensions.

```
julia> A = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> prod(A, 1)
1×2 Array{Int64,2}:
 3  8

julia> prod(A, 2)
2×1 Array{Int64,2}:
 2
12
```

prod! (*r*, *A*)

Multiply elements of *A* over the singleton dimensions of *r*, and write results to *r*.

any (*itr*) → Bool

Test whether any elements of a boolean collection are `true`.

```
julia> a = [true, false, false, true]
4-element Array{Bool,1}:
 true
 false
 false
 true

julia> any(a)
true
```

any (*A*, *dims*)

Test whether any values along the given dimensions of an array are `true`.

any! (*r*, *A*)

Test whether any values in *A* along the singleton dimensions of *r* are `true`, and write results to *r*.

all (*itr*) → Bool

Test whether all elements of a boolean collection are `true`.

```
julia> a = [true, false, false, true]
4-element Array{Bool,1}:
 true
 false
 false
 true

julia> all(a)
false
```

all (*A*, *dims*)

Test whether all values along the given dimensions of an array are `true`.

```
julia> A = [true false; true true]
2×2 Array{Bool,2}:
 true  false
 true   true

julia> all(A, 1)
1×2 Array{Bool,2}:
 true  false
```

```
julia> all(A, 2)
2×1 Array{Bool,2}:
 false
  true
```

all! (*r*, *A*)

Test whether all values in *A* along the singleton dimensions of *r* are `true`, and write results to *r*.

count (*p*, *itr*) → Integer

Count the number of elements in *itr* for which predicate *p* returns `true`.

```
julia> count(i->(4<=i<=6), [2,3,4,5,6])
3
```

any (*p*, *itr*) → Bool

Determine whether predicate *p* returns `true` for any elements of *itr*.

```
julia> any(i->(4<=i<=6), [3,5,7])
true
```

all (*p*, *itr*) → Bool

Determine whether predicate *p* returns `true` for all elements of *itr*.

```
julia> all(i->(4<=i<=6), [4,5,6])
true
```

foreach (*f*, *c...*) → Void

Call function *f* on each element of iterable *c*. For multiple iterable arguments, *f* is called element-wise. `foreach` should be used instead of `map` when the results of *f* are not needed, for example in `foreach println, array`.

```
julia> a = 1:3:7;

julia> foreach(x->println(x^2), a)
1
16
49
```

map (*f*, *c...*) → collection

Transform collection *c* by applying *f* to each element. For multiple collection arguments, apply *f* elementwise.

```
julia> map((x) -> x * 2, [1, 2, 3])
3-element Array{Int64,1}:
 2
 4
 6

julia> map(+, [1, 2, 3], [10, 20, 30])
3-element Array{Int64,1}:
11
22
33
```

map! (*function*, *collection*)

In-place version of `map()` (page 390).

map! (*function*, *destination*, *collection...*)

Like `map()` (page 390), but stores the result in *destination* rather than a new collection. *destination* must be at least as large as the first collection.

mapreduce (*f*, *op*, *v0*, *itr*)

Apply function *f* to each element in *itr*, and then reduce the result using the binary function *op*. *v0* must be a neutral element for *op* that will be returned for empty collections. It is unspecified whether *v0* is used for non-empty collections.

`mapreduce()` (page 390) is functionally equivalent to calling `reduce(op, v0, map(f, itr))`, but will in general execute faster since no intermediate collection needs to be created. See documentation for `reduce()` (page 385) and `map()` (page 390).

```
julia> mapreduce(x->x^2, +, [1:3;]) # == 1 + 4 + 9
14
```

The associativity of the reduction is implementation-dependent. Additionally, some implementations may reuse the return value of *f* for elements that appear multiple times in *itr*. Use `mapfoldl()` (page 391) or `mapfoldr()` (page 391) instead for guaranteed left or right associativity and invocation of *f* for every value.

mapreduce (*f*, *op*, *itr*)

Like `mapreduce(f, op, v0, itr)`. In general, this cannot be used with empty collections (see `reduce(op, itr)`).

mapfoldl (*f*, *op*, *v0*, *itr*)

Like `mapreduce()` (page 390), but with guaranteed left associativity. *v0* will be used exactly once.

mapfoldl (*f*, *op*, *itr*)

Like `mapfoldl(f, op, v0, itr)`, but using the first element of *itr* as *v0*. In general, this cannot be used with empty collections (see `reduce(op, itr)`).

mapfoldr (*f*, *op*, *v0*, *itr*)

Like `mapreduce()` (page 390), but with guaranteed right associativity. *v0* will be used exactly once.

mapfoldr (*f*, *op*, *itr*)

Like `mapfoldr(f, op, v0, itr)`, but using the first element of *itr* as *v0*. In general, this cannot be used with empty collections (see `reduce(op, itr)`).

first (*coll*)

Get the first element of an iterable collection. Returns the start point of a `Range` even if it is empty.

```
julia> first(2:2:10)
2

julia> first([1; 2; 3; 4])
1
```

last (*coll*)

Get the last element of an ordered collection, if it can be computed in $O(1)$ time. This is accomplished by calling `endof()` (page 383) to get the last index. Returns the end point of a `Range` even if it is empty.

```
julia> last(1:2:10)
9

julia> last([1; 2; 3; 4])
4
```

step (*r*)

Get the step size of a `Range` object.

```
julia> step(1:10)
1

julia> step(1:2:10)
```

```
2

julia> step(2.5:0.3:10.9)
0.3

julia> step(linspace(2.5,10.9,85))
0.1
```

collect (*collection*)

Return an `Array` of all items in a collection or iterator. For associative collections, returns `Pair{KeyType, ValType}`. If the argument is array-like or is an iterator with the `HasShape()` trait, the result will have the same shape and number of dimensions as the argument.

```
julia> collect(1:2:13)
7-element Array{Int64,1}:
 1
 3
 5
 7
 9
11
13
```

collect (*element_type, collection*)

Return an `Array` with the given element type of all items in a collection or iterable. The result has the same shape and number of dimensions as *collection*.

issubset (*a, b*)

$\subseteq(a, b) \rightarrow \text{Bool}$

$\not\subseteq(a, b) \rightarrow \text{Bool}$

$\subsetneq(a, b) \rightarrow \text{Bool}$

Determine whether every element of *a* is also in *b*, using `in()` (page 383).

filter (*function, collection*)

Return a copy of *collection*, removing elements for which *function* is false. For associative collections, the function is passed two arguments (key and value).

```
julia> a = 1:10
1:10

julia> filter(isodd, a)
5-element Array{Int64,1}:
 1
 3
 5
 7
 9
```

filter! (*function, collection*)

Update *collection*, removing elements for which *function* is false. For associative collections, the function is passed two arguments (key and value).

```
julia> filter!(isodd, collect(1:10))
5-element Array{Int64,1}:
 1
 3
 5
 7
 9
```

41.4 Indexable Collections

getindex (*collection, key...*)

Retrieve the value(s) stored at the given key or index within a collection. The syntax `a[i, j, ...]` is converted by the compiler to `getindex(a, i, j, ...)`.

```
julia> A = Dict{"a" => 1, "b" => 2}
Dict{String, Int64} with 2 entries:
  "b" => 2
  "a" => 1

julia> getindex(A, "a")
1
```

setindex! (*collection, value, key...*)

Store the given value at the given key or index within a collection. The syntax `a[i, j, ...] = x` is converted by the compiler to `(setindex!(a, x, i, j, ...); x)`.

Fully implemented by:

- [Array](#) (page 454)
- [BitArray](#) (page 455)
- [AbstractArray](#)
- [SubArray](#)
- [ObjectIdDict](#)
- [Dict](#) (page 394)
- [WeakKeyDict](#)
- [AbstractString](#)

Partially implemented by:

- [Range](#)
- [UnitRange](#)
- [Tuple](#)

41.5 Associative Collections

[Dict](#) (page 394) is the standard associative collection. Its implementation uses [hash\(\)](#) (page 363) as the hashing function for the key, and [isequal\(\)](#) (page 362) to determine equality. Define these two functions for custom types to override how they are stored in a hash table.

[ObjectIdDict](#) is a special hash table where the keys are always object identities.

[WeakKeyDict](#) is a hash table implementation where the keys are weak references to objects, and thus may be garbage collected even when referenced in a hash table.

[Dict](#) (page 394)s can be created by passing pair objects constructed with `=>()` to a [Dict](#) (page 394) constructor: `Dict("A"=>1, "B"=>2)`. This call will attempt to infer type information from the keys and values (i.e. this example creates a `Dict{String, Int64}`). To explicitly specify types use the syntax `Dict{KeyType, ValueType}(...)`. For example, `Dict{String, Int32}("A"=>1, "B"=>2)`.

[Dict](#) (page 394)s may also be created with generators. For example, `Dict{i => f(i) for i = 1:10}`.

Given a dictionary `D`, the syntax `D[x]` returns the value of key `x` (if it exists) or throws an error, and `D[x] = y` stores the key-value pair `x => y` in `D` (replacing any existing value for the key `x`). Multiple arguments to `D[...]` are converted to tuples; for example, the syntax `D[x, y]` is equivalent to `D[(x, y)]`, i.e. it refers to the value keyed by the tuple `(x, y)`.

Dict (*itr*)

`Dict{K,V}()` constructs a hash table with keys of type `K` and values of type `V`.

Given a single iterable argument, constructs a *Dict* (page 394) whose key-value pairs are taken from 2-tuples `(key, value)` generated by the argument.

```
julia> Dict{String, Int64}([("A", 1), ("B", 2)])
Dict{String, Int64} with 2 entries:
  "B" => 2
  "A" => 1
```

Alternatively, a sequence of pair arguments may be passed.

```
julia> Dict{String, Int64}("A"=>1, "B"=>2)
Dict{String, Int64} with 2 entries:
  "B" => 2
  "A" => 1
```

haskey (*collection, key*) → Bool

Determine whether a collection has a mapping for a given key.

```
julia> a = Dict{Char, Int64}('a'=>2, 'b'=>3)
Dict{Char, Int64} with 2 entries:
  'b' => 3
  'a' => 2

julia> haskey(a, 'a')
true

julia> haskey(a, 'c')
false
```

get (*collection, key, default*)

Return the value stored for the given key, or the given default value if no mapping for the key is present.

get (*f::Function, collection, key*)

Return the value stored for the given key, or if no mapping for the key is present, return `f()`. Use *get!()* (page 394) to also store the default value in the dictionary.

This is intended to be called using `do` block syntax

```
get(dict, key) do
    # default value calculated here
    time()
end
```

get! (*collection, key, default*)

Return the value stored for the given key, or if no mapping for the key is present, store `key => default`, and return `default`.

get! (*f::Function, collection, key*)

Return the value stored for the given key, or if no mapping for the key is present, store `key => f()`, and return `f()`.

This is intended to be called using `do` block syntax:


```

get!(dict, key) do
    # default value calculated here
    time()
end

```

getkey (*collection, key, default*)

Return the key matching argument *key* if one exists in *collection*, otherwise return *default*.

```

julia> a = Dict{'a'=>2, 'b'=>3}
Dict{Char,Int64} with 2 entries:
  'b' => 3
  'a' => 2

julia> getkey(a, 'a', 1)
'a'

julia> getkey(a, 'd', 'a')
'a'

```

delete! (*collection, key*)

Delete the mapping for the given *key* in a collection, and return the collection.

pop! (*collection, key* [, *default*])

Delete and return the mapping for *key* if it exists in *collection*, otherwise return *default*, or throw an error if *default* is not specified.

keys (*a::Associative*)

Return an iterator over all keys in a collection. `collect(keys(d))` returns an array of keys. Since the keys are stored internally in a hash table, the order in which they are returned may vary.

```

julia> a = Dict{'a'=>2, 'b'=>3}
Dict{Char,Int64} with 2 entries:
  'b' => 3
  'a' => 2

julia> collect(keys(a))
2-element Array{Char,1}:
 'b'
 'a'

```

values (*a::Associative*)

Return an iterator over all values in a collection. `collect(values(d))` returns an array of values.

```

julia> a = Dict{'a'=>2, 'b'=>3}
Dict{Char,Int64} with 2 entries:
  'b' => 3
  'a' => 2

julia> collect(values(a))
2-element Array{Int64,1}:
 3
 2

```

merge (*d::Associative, others::Associative...*)

Construct a merged collection from the given collections. If necessary, the types of the resulting collection will be promoted to accommodate the types of the merged collections. If the same key is present in another collection, the value for that key will be the value it has in the last collection listed.

```
julia> a = Dict{"foo" => 0.0, "bar" => 42.0}
Dict{String,Float64} with 2 entries:
  "bar" => 42.0
  "foo" => 0.0

julia> b = Dict{"baz" => 17, "bar" => 4711}
Dict{String,Int64} with 2 entries:
  "bar" => 4711
  "baz" => 17

julia> merge(a, b)
Dict{String,Float64} with 3 entries:
  "bar" => 4711.0
  "baz" => 17.0
  "foo" => 0.0

julia> merge(b, a)
Dict{String,Float64} with 3 entries:
  "bar" => 42.0
  "baz" => 17.0
  "foo" => 0.0
```

merge! (*d::Associative, others::Associative...*)

Update collection with pairs from the other collections. See also [merge\(\)](#) (page 395).

sizehint! (*s, n*)

Suggest that collection *s* reserve capacity for at least *n* elements. This can improve performance.

keytype (*type*)

Get the key type of an associative collection type. Behaves similarly to [eltype\(\)](#) (page 384).

valtype (*type*)

Get the value type of an associative collection type. Behaves similarly to [eltype\(\)](#) (page 384).

Fully implemented by:

- `ObjectIdDict`
- [Dict](#) (page 394)
- `WeakKeyDict`

Partially implemented by:

- [IntSet](#) (page 396)
- [Set](#) (page 396)
- [EnvHash](#) (page 373)
- [Array](#) (page 454)
- [BitArray](#) (page 455)

41.6 Set-Like Collections

Set (*[itr]*)

Construct a [Set](#) (page 396) of the values generated by the given iterable object, or an empty set. Should be used instead of [IntSet](#) (page 396) for sparse integer sets, or for sets of arbitrary objects.

IntSet (*itr*)

Construct a sorted set of positive `Ints` generated by the given iterable object, or an empty set. Implemented as a bit string, and therefore designed for dense integer sets. Only `Ints` greater than 0 can be stored. If the set will be sparse (for example holding a few very large integers), use `Set` (page 396) instead.

union (*s1, s2...*)
 $\cup (s1, s2...)$

Construct the union of two or more sets. Maintains order with arrays.

union! (*s, iterable*)

Union each element of *iterable* into set *s* in-place.

intersect (*s1, s2...*)
 $\cap (s1, s2)$

Construct the intersection of two or more sets. Maintains order and multiplicity of the first argument for arrays and ranges.

setdiff (*a, b*)

Construct the set of elements in *a* but not *b*. Maintains order with arrays. Note that both arguments must be collections, and both will be iterated over. In particular, `setdiff(set, element)` where *element* is a potential member of *set*, will not work in general.

```
julia> setdiff([1,2,3],[3,4,5])
2-element Array{Int64,1}:
 1
 2
```

setdiff! (*s, iterable*)

Remove each element of *iterable* from set *s* in-place.

symdiff (*a, b, rest...*)

Construct the symmetric difference of elements in the passed in sets or arrays. Maintains order with arrays.

```
julia> symdiff([1,2,3],[3,4,5],[4,5,6])
3-element Array{Int64,1}:
 1
 2
 6
```

symdiff! (*s, n*)

The set *s* is destructively modified to toggle the inclusion of integer *n*.

symdiff! (*s, itr*)

For each element in *itr*, destructively toggle its inclusion in set *s*.

symdiff! (*s1, s2*)

Construct the symmetric difference of sets *s1* and *s2*, storing the result in *s1*.

intersect! (*s1::IntSet, s2::IntSet*)

Intersects sets *s1* and *s2* and overwrites the set *s1* with the result. If needed, *s1* will be expanded to the size of *s2*.

issubset (*A, S*) \rightarrow Bool
 $\subseteq (A, S) \rightarrow \text{Bool}$

Return `true` if *A* is a subset of or equal to *S*.

Fully implemented by:

- `IntSet` (page 396)
- `Set` (page 396)

Partially implemented by:

- [Array](#) (page 454)

41.7 Dequeues

push! (*collection*, *items...*) → *collection*

Insert one or more *items* at the end of *collection*.

```
julia> push!([1, 2, 3], 4, 5, 6)
6-element Array{Int64,1}:
 1
 2
 3
 4
 5
 6
```

Use [append!\(\)](#) (page 401) to add all the elements of another collection to *collection*. The result of the preceding example is equivalent to `append!([1, 2, 3], [4, 5, 6])`.

pop! (*collection*) → *item*

Remove the last item in *collection* and return it.

```
julia> A=[1, 2, 3, 4, 5, 6]
6-element Array{Int64,1}:
 1
 2
 3
 4
 5
 6

julia> pop!(A)
6

julia> A
5-element Array{Int64,1}:
 1
 2
 3
 4
 5
```

unshift! (*collection*, *items...*) → *collection*

Insert one or more *items* at the beginning of *collection*.

```
julia> unshift!([1, 2, 3, 4], 5, 6)
6-element Array{Int64,1}:
 5
 6
 1
 2
 3
 4
```

shift! (*collection*) → *item*

Remove the first item from *collection*.

```
julia> A = [1, 2, 3, 4, 5, 6]
6-element Array{Int64,1}:
 1
 2
 3
 4
 5
 6

julia> shift!(A)
1

julia> A
5-element Array{Int64,1}:
 2
 3
 4
 5
 6
```

insert! (*a::Vector*, *index::Integer*, *item*)

Insert an item into *a* at the given *index*. *index* is the index of *item* in the resulting *a*.

```
julia> insert!([6, 5, 4, 2, 1], 4, 3)
6-element Array{Int64,1}:
 6
 5
 4
 3
 2
 1
```

deleteat! (*a::Vector*, *i::Integer*)

Remove the item at the given *i* and return the modified *a*. Subsequent items are shifted to fill the resulting gap.

```
julia> deleteat!([6, 5, 4, 3, 2, 1], 2)
5-element Array{Int64,1}:
 6
 4
 3
 2
 1
```

deleteat! (*a::Vector*, *inds*)

Remove the items at the indices given by *inds*, and return the modified *a*. Subsequent items are shifted to fill the resulting gap. *inds* must be sorted and unique.

```
julia> deleteat!([6, 5, 4, 3, 2, 1], 1:2:5)
3-element Array{Int64,1}:
 5
 3
 1

julia> deleteat!([6, 5, 4, 3, 2, 1], (2, 2))
ERROR: ArgumentError: indices must be unique and sorted
in deleteat! (::Array{Int64,1}, ::Tuple{Int64,Int64}) at ./array.jl:748
...
```

splice! (*a::Vector*, *index::Integer*[, *replacement*]) → *item*

Remove the item at the given index, and return the removed item. Subsequent items are shifted left to fill the resulting gap. If specified, replacement values from an ordered collection will be spliced in place of the removed item.

```
julia> A = [6, 5, 4, 3, 2, 1]; splice!(A, 5)
2
```

```
julia> A
5-element Array{Int64,1}:
 6
 5
 4
 3
 1
```

```
julia> splice!(A, 5, -1)
1
```

```
julia> A
5-element Array{Int64,1}:
 6
 5
 4
 3
-1
```

```
julia> splice!(A, 1, [-1, -2, -3])
6
```

```
julia> A
7-element Array{Int64,1}:
-1
-2
-3
 5
 4
 3
-1
```

To insert replacement before an index n without removing any items, use `splice!(collection, n:n-1, replacement)`.

splice! (*a::Vector*, *range*[, *replacement*]) → *items*

Remove items in the specified index range, and return a collection containing the removed items. Subsequent items are shifted left to fill the resulting gap. If specified, replacement values from an ordered collection will be spliced in place of the removed items.

To insert replacement before an index n without removing any items, use `splice!(collection, n:n-1, replacement)`.

```
julia> splice!(A, 4:3, 2)
0-element Array{Int64,1}
```

```
julia> A
8-element Array{Int64,1}:
-1
-2
-3
 2
 5
 4
 3
-1
```

```

5
4
3
-1

```

resize! (*a::Vector*, *n::Integer*) → *Vector*

Resize *a* to contain *n* elements. If *n* is smaller than the current collection length, the first *n* elements will be retained. If *n* is larger, the new elements are not guaranteed to be initialized.

```

julia> resize!([6, 5, 4, 3, 2, 1], 3)
3-element Array{Int64,1}:
 6
 5
 4

```

```

julia> resize!([6, 5, 4, 3, 2, 1], 8)
8-element Array{Int64,1}:
 6
 5
 4
 3
 2
 1
 0
 0

```

append! (*collection*, *collection2*) → *collection*.

Add the elements of *collection2* to the end of *collection*.

```

julia> append!([1], [2, 3])
3-element Array{Int64,1}:
 1
 2
 3

```

```

julia> append!([1, 2, 3], [4, 5, 6])
6-element Array{Int64,1}:
 1
 2
 3
 4
 5
 6

```

Use [push!\(\)](#) (page 398) to add individual items to *collection* which are not already themselves in another collection. The result of the preceding example is equivalent to `push!([1, 2, 3], 4, 5, 6)`.

prepend! (*a::Vector*, *items*) → *collection*

Insert the elements of *items* to the beginning of *a*.

```

julia> prepend!([3], [1, 2])
3-element Array{Int64,1}:
 1
 2
 3

```

Fully implemented by:

- `Vector` (a.k.a. 1-dimensional [Array](#) (page 454))
- `BitVector` (a.k.a. 1-dimensional [BitArray](#) (page 455))

41.8 PriorityQueue

The *PriorityQueue* (page 402) type is available from the `Collections` module. It provides a basic priority queue implementation allowing for arbitrary key and priority types. Multiple identical keys are not permitted, but the priority of existing keys can be changed efficiently.

PriorityQueue (*K*, *V*[, *ord*])

Construct a new *PriorityQueue* (page 402), with keys of type *K* and values/priorities of type *V*. If an order is not given, the priority queue is min-ordered using the default comparison for *V*.

A *PriorityQueue* acts like a `Dict`, mapping values to their priorities, with the addition of a `dequeue!` function to remove the lowest priority element.

```
julia> a = Base.Collections.PriorityQueue(["a", "b", "c"], [2, 3, 1], Base.Order.Forward)
Base.Collections.PriorityQueue{String, Int64, Base.Order.ForwardOrdering} with 3 entries:
  "c" => 1
  "b" => 3
  "a" => 2
```

enqueue! (*pq*, *k*, *v*)

Insert the a key *k* into a priority queue *pq* with priority *v*.

```
julia> a = Base.Collections.PriorityQueue(["a", "b", "c"], [2, 3, 1], Base.Order.Forward)
Base.Collections.PriorityQueue{String, Int64, Base.Order.ForwardOrdering} with 3 entries:
  "c" => 1
  "b" => 3
  "a" => 2

julia> Base.Collections.enqueue!(a, "d", 4)
Base.Collections.PriorityQueue{String, Int64, Base.Order.ForwardOrdering} with 4 entries:
  "c" => 1
  "b" => 3
  "a" => 2
  "d" => 4
```

dequeue! (*pq*)

Remove and return the lowest priority key from a priority queue.

```
julia> a = Base.Collections.PriorityQueue(["a", "b", "c"], [2, 3, 1], Base.Order.Forward)
Base.Collections.PriorityQueue{String, Int64, Base.Order.ForwardOrdering} with 3 entries:
  "c" => 1
  "b" => 3
  "a" => 2

julia> Base.Collections.dequeue!(a)
"c"

julia> a
Base.Collections.PriorityQueue{String, Int64, Base.Order.ForwardOrdering} with 2 entries:
  "b" => 3
  "a" => 2
```

peek (*pq*)

Return the lowest priority key from a priority queue without removing that key from the queue.

PriorityQueue (page 402) also behaves similarly to a `Dict` in that keys can be inserted and priorities accessed or changed using indexing notation.


```
julia> # Julia code
      pq = Collections.PriorityQueue();

julia> # Insert keys with associated priorities
      pq["a"] = 10; pq["b"] = 5; pq["c"] = 15; pq
Base.Collections.PriorityQueue{Any, Any, Base.Order.ForwardOrdering} with 3 entries:
  "c" => 15
  "b" => 5
  "a" => 10

julia> # Change the priority of an existing key
      pq["a"] = 0; pq
Base.Collections.PriorityQueue{Any, Any, Base.Order.ForwardOrdering} with 3 entries:
  "c" => 15
  "b" => 5
  "a" => 0
```

41.9 Heap Functions

Along with the *PriorityQueue* (page 402) type, the `Collections` module provides lower level functions for performing binary heap operations on arrays. Each function takes an optional ordering argument. If not given, default ordering is used, so that elements popped from the heap are given in ascending order.

heapify (*v*, *ord::Ordering=Forward*)

Returns a new vector in binary heap order, optionally using the given ordering.

```
julia> a = [1,3,4,5,2];

julia> Base.Collections.heapify(a)
5-element Array{Int64,1}:
 1
 2
 4
 5
 3

julia> Base.Collections.heapify(a, Base.Order.Reverse)
5-element Array{Int64,1}:
 5
 3
 4
 1
 2
```

heapify! (*v*, *ord::Ordering=Forward*)

In-place *heapify()* (page 403).

isheap (*v*, *ord::Ordering=Forward*)

Return `true` if an array is heap-ordered according to the given order.

```
julia> a = [1,2,3]
3-element Array{Int64,1}:
 1
 2
 3

julia> Base.Collections.isheap(a, Base.Order.Forward)
```

```
true

julia> Base.Collections.isheap(a,Base.Order.Reverse)
false
```

heappush! (*v*, *x*[, *ord*])

Given a binary heap-ordered array, push a new element *x*, preserving the heap property. For efficiency, this function does not check that the array is indeed heap-ordered.

heappop! (*v*[, *ord*])

Given a binary heap-ordered array, remove and return the lowest ordered element. For efficiency, this function does not check that the array is indeed heap-ordered.

Mathematics

42.1 Mathematical Operators

$- (x)$

Unary minus operator.

$+ (x, y\dots)$

Addition operator. $x+y+z+\dots$ calls this function with all arguments, i.e. $+(x, y, z, \dots)$.

$- (x, y)$

Subtraction operator.

$\star (x, y\dots)$

Multiplication operator. $x\star y\star z\star\dots$ calls this function with all arguments, i.e. $\star(x, y, z, \dots)$.

$/ (x, y)$

Right division operator: multiplication of x by the inverse of y on the right. Gives floating-point results for integer arguments.

$\backslash (x, y)$

Left division operator: multiplication of y by the inverse of x on the left. Gives floating-point results for integer arguments.

$^ (x, y)$

Exponentiation operator.

$.\star (x, y)$

Element-wise addition operator.

```
julia> A = [1 2; 3 4];

julia> B = [5 6; 7 8];

julia> C = [A, B]
2-element Array{Array{Int64,2},1}:
 [1 2; 3 4]
 [5 6; 7 8]

julia> C .+ [[1; 2] [3; 4]]
2×2 Array{Array{Int64,2},2}:
 [2 3; 4 5]  [4 5; 6 7]
 [7 8; 9 10] [9 10; 11 12]
```

See also `broadcast()` (page 459).

`.- (x, y)`

Element-wise subtraction operator.

```
julia> [4; 5; 6] .- [1; 2; 4]
3-element Array{Int64,1}:
 3
 3
 2
```

`.* (x, y)`

Element-wise multiplication operator.

```
julia> [1 2 3] .* [1 2 3]
1×3 Array{Int64,2}:
 1  4  9
```

`./ (x, y)`

Element-wise right division operator.

```
julia> [1 2 3] ./ [1 2 3]
1×3 Array{Float64,2}:
 1.0  1.0  1.0
```

`.\ (x, y)`

Element-wise left division operator.

```
julia> A = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4
```

```
julia> A .\ [1 2]
2×2 Array{Float64,2}:
 1.0  1.0
 0.333333  0.5
```

```
julia> A = [1 0; 0 -1];

julia> B = [0 1; 1 0];

julia> C = [A, B]
2-element Array{Array{Int64,2},1}:
 [1 0; 0 -1]
 [0 1; 1 0]

julia> x = [1; 0];

julia> y = [0; 1];

julia> D = [x, y]
2-element Array{Array{Int64,1},1}:
 [1,0]
 [0,1]

julia> C .\ D
2-element Array{Array{Float64,1},1}:
 [1.0,-0.0]
 [1.0,0.0]
```

See also [broadcast \(\)](#) (page 459).

.[^](*x*, *y*)

Element-wise exponentiation operator.

```
julia> [1 2 3] .^ [1 2 3]
1×3 Array{Int64,2}:
 1  4 27
```

fma(*x*, *y*, *z*)

Computes $x*y+z$ without rounding the intermediate result $x*y$. On some systems this is significantly more expensive than $x*y+z$. `fma` is used to improve accuracy in certain algorithms. See [muladd\(\)](#) (page 407).

muladd(*x*, *y*, *z*)

Combined multiply-add, computes $x*y+z$ in an efficient manner. This may on some systems be equivalent to $x*y+z$, or to `fma(x, y, z)`. `muladd` is used to improve performance. See [fma\(\)](#) (page 407).

div(*x*, *y*)**÷**(*x*, *y*)

The quotient from Euclidean division. Computes x/y , truncated to an integer.

fld(*x*, *y*)

Largest integer less than or equal to x/y .

```
julia> fld(7.3, 5.5)
1.0
```

cld(*x*, *y*)

Smallest integer larger than or equal to x/y .

```
julia> cld(5.5, 2.2)
3.0
```

mod(*x*, *y*)

Modulus after flooring division, returning in the range $[0, y)$, if y is positive, or $(y, 0]$ if y is negative.

```
x == fld(x, y) * y + mod(x, y)
```

mod2pi(*x*)

Modulus after division by 2π , returning in the range $[0, 2\pi)$.

This function computes a floating point representation of the modulus after division by numerically exact 2π , and is therefore not exactly the same as `mod(x, 2π)`, which would compute the modulus of x relative to division by the floating-point number 2π .

```
julia> mod2pi(9*pi/4)
0.7853981633974481
```

rem(*x*, *y*)**%**(*x*, *y*)

Remainder from Euclidean division, returning a value of the same sign as x , and smaller in magnitude than y . This value is always exact.

```
x == div(x, y) * y + rem(x, y)
```

divrem(*x*, *y*)

The quotient and remainder from Euclidean division. Equivalent to `(div(x, y), rem(x, y))` or `(x÷y, x%y)`.

```
julia> divrem(3, 7)
(0, 3)
```

```
julia> divrem(7,3)
(2,1)
```

fldmod(*x*, *y*)

The floored quotient and modulus after division. Equivalent to `(fld(x, y), mod(x, y))`.

fld1(*x*, *y*)

Flooring division, returning a value consistent with `mod1(x, y)`

```
x == fld(x,y)*y + mod(x,y)
x == (fld1(x,y)-1)*y + mod1(x,y)
```

mod1(*x*, *y*)

Modulus after flooring division, returning a value *r* such that `mod(r, y) == mod(x, y)` in the range `(0, y]` for positive *y* and in the range `[y, 0)` for negative *y*.

fldmod1(*x*, *y*)

Return `(fld1(x, y), mod1(x, y))`.

//(*num*, *den*)

Divide two integers or rational numbers, giving a `Rational` result.

rationalize(`[T::Integer=Int]`, *x*; `tol::Real=eps(x)`)

Approximate floating point number *x* as a `Rational` number with components of the given integer type. The result will differ from *x* by no more than *tol*. If *T* is not provided, it defaults to `Int`.

```
julia> rationalize(5.6)
28//5

julia> a = rationalize{BigInt}(10.3)
103//10

julia> typeof(numerator(a))
BigInt
```

numerator(*x*)

Numerator of the rational representation of *x*.

denominator(*x*)

Denominator of the rational representation of *x*.

<<(*x*, *n*)

Left bit shift operator, `x << n`. For `n ≥ 0`, the result is *x* shifted left by *n* bits, filling with 0s. This is equivalent to `x * 2n`. For `n < 0`, this is equivalent to `x >> -n`.

```
julia> Int8(3) << 2
12

julia> bits{Int8}(3)
"00000011"

julia> bits{Int8}(12)
"00001100"
```

See also `>>()` (page 408), `>>>()` (page 409).

>>(*x*, *n*)

Right bit shift operator, `x >> n`. For `n ≥ 0`, the result is *x* shifted right by *n* bits, where `n ≥ 0`, filling with 0s if `x ≥ 0`, 1s if `x < 0`, preserving the sign of *x*. This is equivalent to `fld(x, 2n)`. For `n < 0`, this is equivalent to `x << -n`.

```
julia> Int8(13) >> 2
3

julia> bits(Int8(13))
"00001101"

julia> bits(Int8(3))
"00000011"

julia> Int8(-14) >> 2
-4

julia> bits(Int8(-14))
"11110010"

julia> bits(Int8(-4))
"11111100"
```

See also `>>>()` (page 409), `<<()` (page 408).

>>>(*x*, *n*)

Unsigned right bit shift operator, `x >>> n`. For `n >= 0`, the result is `x` shifted right by `n` bits, where `n >= 0`, filling with 0s. For `n < 0`, this is equivalent to `x << -n`.

For Unsigned integer types, this is equivalent to `>>()` (page 408). For Signed integer types, this is equivalent to `signed(unsigned(x) >> n)`.

```
julia> Int8(-14) >>> 2
60

julia> bits(Int8(-14))
"11110010"

julia> bits(Int8(60))
"00111100"
```

BigInts are treated as if having infinite size, so no filling is required and this is equivalent to `>>()` (page 408).

See also `>>()` (page 408), `<<()` (page 408).

:(*start*[, *step*], *stop*)

Range operator. `a:b` constructs a range from `a` to `b` with a step size of 1, and `a:s:b` is similar but uses a step size of `s`. These syntaxes call the function `colon`. The colon is also used in indexing to select whole dimensions.

colon(*start*[, *step*], *stop*)

Called by `:` syntax for constructing ranges.

range(*start*[, *step*], *length*)

Construct a range by length, given a starting value and optional step (defaults to 1).

Base.**OneTo**(*n*)

Define an `AbstractUnitRange` that behaves like `1:n`, with the added distinction that the lower limit is guaranteed (by the type system) to be 1.

==(*x*, *y*)

Generic equality operator, giving a single `Bool` result. Falls back to `===`. Should be implemented for all types with a notion of equality, based on the abstract value that an instance represents. For example, all numeric types are compared by numeric value, ignoring type. Strings are compared as sequences of characters, ignoring encoding.

Follows IEEE semantics for floating-point numbers.

Collections should generally implement `==` by calling `==` recursively on all contents.

New numeric types should implement this function for two arguments of the new type, and handle comparison to other types via promotion rules where possible.

`!= (x, y)`

$\neq (x, y)$

Not-equals comparison operator. Always gives the opposite answer as `==`. New types should generally not implement this, and rely on the fallback definition `!= (x, y) = ! (x==y)` instead.

`!== (x, y)`

$\neq (x, y)$

Equivalent to `! (x === y)`.

`< (x, y)`

Less-than comparison operator. New numeric types should implement this function for two arguments of the new type. Because of the behavior of floating-point NaN values, `<` implements a partial order. Types with a canonical partial order should implement `<`, and types with a canonical total order should implement `isless`.

`<= (x, y)`

$\leq (x, y)$

Less-than-or-equals comparison operator.

`> (x, y)`

Greater-than comparison operator. Generally, new types should implement `<` instead of this function, and rely on the fallback definition `> (x, y) = y < x`.

`>= (x, y)`

$\geq (x, y)$

Greater-than-or-equals comparison operator.

`.== (x, y)`

Element-wise equality comparison operator.

```
julia> [1 2 3] .== [1 2 4]
1×3 BitArray{2}:
 true  true  false
```

`.!= (x, y)`

$\neq (x, y)$

Element-wise not-equals comparison operator.

`.< (x, y)`

Element-wise less-than comparison operator.

```
julia> [1; 2; 3] .< [2; 1; 4]
3-element BitArray{1}:
 true
 false
 true
```

`.<= (x, y)`

$\leq (x, y)$

Element-wise less-than-or-equals comparison operator.

`.> (x, y)`

Element-wise greater-than comparison operator.

`.>= (x, y)`

. \geq (*x*, *y*)

Element-wise greater-than-or-equals comparison operator.

cmp (*x*, *y*)Return -1, 0, or 1 depending on whether *x* is less than, equal to, or greater than *y*, respectively. Uses the total order implemented by `isless`. For floating-point numbers, uses `<` but throws an error for unordered arguments.**~** (*x*)

Bitwise not.

```
julia> ~4
-5

julia> ~10
-11

julia> ~true
false
```

& (*x*, *y*)

Bitwise and.

```
julia> 4 & 10
0

julia> 4 & 12
4
```

| (*x*, *y*)

Bitwise or.

```
julia> 4 | 10
14

julia> 4 | 1
5
```

xor (*x*, *y*)(*x*, *y*)

Bitwise exclusive or.

```
julia> [true; true; false] [true; false; false]
3-element Array{Bool,1}:
 false
  true
  false
```

! (*x*)

Boolean not.

```
julia> !true
false

julia> !false
true

julia> ![true false true]
1×3 Array{Bool,2}:
 false true false
```

x && y

Short-circuiting boolean AND.

x || y

Short-circuiting boolean OR.

42.2 Mathematical Functions

isapprox(*x*, *y*; *rtol*::*Real*=*sqrt(eps)*, *atol*::*Real*=0)

Inexact equality comparison: `true` if `norm(x-y) <= atol + rtol*max(norm(x), norm(y))`. The default `atol` is zero and the default `rtol` depends on the types of `x` and `y`.

For real or complex floating-point values, `rtol` defaults to `sqrt(eps(typeof(real(x-y))))`. This corresponds to requiring equality of about half of the significand digits. For other types, `rtol` defaults to zero.

`x` and `y` may also be arrays of numbers, in which case `norm` defaults to `vecnorm` but may be changed by passing a `norm::Function` keyword argument. (For numbers, `norm` is the same thing as `abs`.) When `x` and `y` are arrays, if `norm(x-y)` is not finite (i.e. `±Inf` or `NaN`), the comparison falls back to checking whether all elements of `x` and `y` are approximately equal component-wise.

The binary operator `≈` is equivalent to `isapprox` with the default arguments, and `x ≈ y` is equivalent to `!isapprox(x, y)`.

sin(*x*)

Compute sine of `x`, where `x` is in radians.

cos(*x*)

Compute cosine of `x`, where `x` is in radians.

tan(*x*)

Compute tangent of `x`, where `x` is in radians.

sind(*x*)

Compute sine of `x`, where `x` is in degrees.

cosd(*x*)

Compute cosine of `x`, where `x` is in degrees.

tand(*x*)

Compute tangent of `x`, where `x` is in degrees.

sinpi(*x*)

Compute `sin(πx)` more accurately than `sin(pi*x)`, especially for large `x`.

cospi(*x*)

Compute `cos(πx)` more accurately than `cos(pi*x)`, especially for large `x`.

sinh(*x*)

Compute hyperbolic sine of `x`.

cosh(*x*)

Compute hyperbolic cosine of `x`.

tanh(*x*)

Compute hyperbolic tangent of `x`.

asin(*x*)

Compute the inverse sine of `x`, where the output is in radians.

acos(*x*)

Compute the inverse cosine of `x`, where the output is in radians

atan(x)Compute the inverse tangent of x , where the output is in radians.**atan2**(y, x)Compute the inverse tangent of y/x , using the signs of both x and y to determine the quadrant of the return value.**asind**(x)Compute the inverse sine of x , where the output is in degrees.**acosd**(x)Compute the inverse cosine of x , where the output is in degrees.**atand**(x)Compute the inverse tangent of x , where the output is in degrees.**sec**(x)Compute the secant of x , where x is in radians.**csc**(x)Compute the cosecant of x , where x is in radians.**cot**(x)Compute the cotangent of x , where x is in radians.**secd**(x)Compute the secant of x , where x is in degrees.**cscd**(x)Compute the cosecant of x , where x is in degrees.**cotd**(x)Compute the cotangent of x , where x is in degrees.**asec**(x)Compute the inverse secant of x , where the output is in radians.**acsc**(x)Compute the inverse cosecant of x , where the output is in radians.**acot**(x)Compute the inverse cotangent of x , where the output is in radians.**asecd**(x)Compute the inverse secant of x , where the output is in degrees.**acscd**(x)Compute the inverse cosecant of x , where the output is in degrees.**acotd**(x)Compute the inverse cotangent of x , where the output is in degrees.**sech**(x)Compute the hyperbolic secant of x **csch**(x)Compute the hyperbolic cosecant of x .**coth**(x)Compute the hyperbolic cotangent of x .**asinh**(x)Compute the inverse hyperbolic sine of x .

acosh (*x*)

Compute the inverse hyperbolic cosine of *x*.

atanh (*x*)

Compute the inverse hyperbolic tangent of *x*.

asech (*x*)

Compute the inverse hyperbolic secant of *x*.

acsch (*x*)

Compute the inverse hyperbolic cosecant of *x*.

acoth (*x*)

Compute the inverse hyperbolic cotangent of *x*.

sinc (*x*)

Compute $\sin(\pi x)/(\pi x)$ if $x \neq 0$, and 1 if $x = 0$.

cosc (*x*)

Compute $\cos(\pi x)/x - \sin(\pi x)/(\pi x^2)$ if $x \neq 0$, and 0 if $x = 0$. This is the derivative of `sinc(x)`.

deg2rad (*x*)

Convert *x* from degrees to radians.

```
julia> deg2rad(90)
1.5707963267948966
```

rad2deg (*x*)

Convert *x* from radians to degrees.

```
julia> rad2deg(pi)
180.0
```

hypot (*x*, *y*)

Compute the hypotenuse $\sqrt{x^2 + y^2}$ avoiding overflow and underflow.

hypot (*x...*)

Compute the hypotenuse $\sqrt{\sum x_i^2}$ avoiding overflow and underflow.

log (*x*)

Compute the natural logarithm of *x*. Throws `DomainError` (page 375) for negative Real arguments. Use complex negative arguments to obtain complex results.

There is an experimental variant in the `Base.Math.JuliaLibm` module, which is typically faster and more accurate.

log (*b*, *x*)

Compute the base *b* logarithm of *x*. Throws `DomainError` (page 375) for negative Real arguments.

```
julia> log(4, 8)
1.5

julia> log(4, 2)
0.5
```

Note: If *b* is a power of 2 or 10, `log2()` (page 415) or `log10()` (page 415) should be used, as these will typically be faster and more accurate. For example,

```
julia> log(100, 1000000)
2.9999999999999996
```

```
julia> log10(1000000)/2
3.0
```

log2(x)

Compute the logarithm of x to base 2. Throws *DomainError* (page 375) for negative *Real* arguments.

```
julia> log2(4)
2.0

julia> log2(10)
3.321928094887362
```

log10(x)

Compute the logarithm of x to base 10. Throws *DomainError* (page 375) for negative *Real* arguments.

```
julia> log10(100)
2.0

julia> log10(2)
0.3010299956639812
```

log1p(x)

Accurate natural logarithm of $1+x$. Throws *DomainError* (page 375) for *Real* arguments less than -1.

There is an experimental variant in the `Base.Math.Libm` module, which is typically faster and more accurate.

frexp(val)

Return (x, exp) such that x has a magnitude in the interval $[1/2, 1)$ or 0, and val is equal to $x \times 2^{\text{exp}}$.

exp(x)

Compute e^x .

exp2(x)

Compute 2^x .

```
julia> exp2(5)
32.0
```

exp10(x)

Compute 10^x .

```
julia> exp10(2)
100.0

julia> exp10(0.2)
1.5848931924611136
```

ldexp(x, n)

Compute $x \times 2^n$.

```
julia> ldexp(5., 2)
20.0
```

modf(x)

Return a tuple (fpart, ipart) of the fractional and integral parts of a number. Both parts have the same sign as the argument.

```
julia> modf(3.5)
(0.5, 3.0)
```

expm1(*x*)

Accurately compute $e^x - 1$.

round(*T*, *x*, [*digits*, *base*], [*r::RoundingMode*])

Rounds *x* to an integer value according to the provided [RoundingMode](#) (page 416), returning a value of the same type as *x*. When not specifying a rounding mode the global mode will be used (see [rounding\(\)](#) (page 437)), which by default is round to the nearest integer ([RoundNearest](#) (page 417) mode), with ties (fractional values of 0.5) being rounded to the nearest even integer.

```
julia> round(1.7)
2.0

julia> round(1.5)
2.0

julia> round(2.5)
2.0
```

The optional [RoundingMode](#) (page 416) argument will change how the number gets rounded.

`round(T, x, [r::RoundingMode])` converts the result to type *T*, throwing an [InexactError](#) (page 375) if the value is not representable.

`round(x, digits)` rounds to the specified number of digits after the decimal place (or before if negative).

`round(x, digits, base)` rounds using a base other than 10.

```
julia> round(pi, 2)
3.14

julia> round(pi, 3, 2)
3.125
```

Note: Rounding to specified digits in bases other than 2 can be inexact when operating on binary floating point numbers. For example, the `Float64` value represented by `1.15` is actually *less* than 1.15, yet will be rounded to 1.2.

```
julia> x = 1.15
1.15

julia> @sprintf "%.20f" x
"1.14999999999999991118"

julia> x < 115//100
true

julia> round(x, 1)
1.2
```

RoundingMode

A type used for controlling the rounding mode of floating point operations (via [rounding\(\)](#) (page 437)/[setrounding\(\)](#) (page 437) functions), or as optional arguments for rounding to the nearest integer (via the [round\(\)](#) (page 416) function).

Currently supported rounding modes are:

- [RoundNearest](#) (page 417) (default)
- [RoundNearestTiesAway](#) (page 417)

- [RoundNearestTiesUp](#) (page 417)
- [RoundToZero](#) (page 417)
- [RoundFromZero](#) (BigFloat only)
- [RoundUp](#) (page 417)
- [RoundDown](#) (page 417)

RoundNearest

The default rounding mode. Rounds to the nearest integer, with ties (fractional values of 0.5) being rounded to the nearest even integer.

RoundNearestTiesAway

Rounds to nearest integer, with ties rounded away from zero (C/C++ [round\(\)](#) (page 416) behaviour).

RoundNearestTiesUp

Rounds to nearest integer, with ties rounded toward positive infinity (Java/JavaScript [round\(\)](#) (page 416) behaviour).

RoundToZero

[round\(\)](#) (page 416) using this rounding mode is an alias for [trunc\(\)](#) (page 417).

RoundUp

[round\(\)](#) (page 416) using this rounding mode is an alias for [ceil\(\)](#) (page 417).

RoundDown

[round\(\)](#) (page 416) using this rounding mode is an alias for [floor\(\)](#) (page 417).

round(*z*, *RoundingModeReal*, *RoundingModeImaginary*)

Returns the nearest integral value of the same type as the complex-valued *z* to *z*, breaking ties using the specified *RoundingMode* (page 416)s. The first *RoundingMode* (page 416) is used for rounding the real components while the second is used for rounding the imaginary components.

ceil(*[T]*, *x*, *digits*, *base*)

`ceil(x)` returns the nearest integral value of the same type as *x* that is greater than or equal to *x*.

`ceil(T, x)` converts the result to type *T*, throwing an `InexactError` if the value is not representable.

digits and *base* work as for [round\(\)](#) (page 416).

floor(*[T]*, *x*, *digits*, *base*)

`floor(x)` returns the nearest integral value of the same type as *x* that is less than or equal to *x*.

`floor(T, x)` converts the result to type *T*, throwing an `InexactError` if the value is not representable.

digits and *base* work as for [round\(\)](#) (page 416).

trunc(*[T]*, *x*, *digits*, *base*)

`trunc(x)` returns the nearest integral value of the same type as *x* whose absolute value is less than or equal to *x*.

`trunc(T, x)` converts the result to type *T*, throwing an `InexactError` if the value is not representable.

digits and *base* work as for [round\(\)](#) (page 416).

unsafe_trunc(*T*, *x*)

`unsafe_trunc(T, x)` returns the nearest integral value of type *T* whose absolute value is less than or equal to *x*. If the value is not representable by *T*, an arbitrary value will be returned.

signif(*x*, *digits*, *base*)

Rounds (in the sense of [round\(\)](#) (page 416)) *x* so that there are *digits* significant digits, under a *base* representation, default 10. E.g., `signif(123.456, 2)` is 120.0, and `signif(357.913, 4, 2)` is 352.0.

min(*x*, *y*, ...)

Return the minimum of the arguments. See also the [minimum\(\)](#) (page 385) function to take the minimum element from a collection.

max(*x*, *y*, ...)

Return the maximum of the arguments. See also the [maximum\(\)](#) (page 385) function to take the maximum element from a collection.

minmax(*x*, *y*)

Return `(min(x, y), max(x, y))`. See also: [extrema\(\)](#) (page 386) that returns `(minimum(x), maximum(x))`.

```
julia> minmax('c', 'b')
('b', 'c')
```

clamp(*x*, *lo*, *hi*)

Return *x* if `lo <= x <= hi`. If `x < lo`, return *lo*. If `x > hi`, return *hi*. Arguments are promoted to a common type. Operates elementwise over *x* if *x* is an array.

```
julia> clamp([pi, 1.0, big(10.)], 2., 9.)
3-element Array{BigFloat,1}:
 3.141592653589793238462643383279502884197169399375105820974944592307816406286198
 2.0000000000000000000000000000000000000000000000000000000000000000000000000000000000
 9.0000000000000000000000000000000000000000000000000000000000000000000000000000000000
```

clamp!(*array::AbstractArray*, *lo*, *hi*)

Restrict values in *array* to the specified range, in-place. See also [clamp\(\)](#) (page 418).

abs(*x*)

The absolute value of *x*.

When `abs` is applied to signed integers, overflow may occur, resulting in the return of a negative value. This overflow occurs only when `abs` is applied to the minimum representable value of a signed integer. That is, when `x == typemin(typeof(x))`, `abs(x) == x < 0`, not `-x` as might be expected.

Base. **checked_abs**(*x*)

Calculates `abs(x)`, checking for overflow errors where applicable. For example, standard two's complement signed integers (e.g. `Int`) cannot represent `abs(typemin(Int))`, thus leading to an overflow.

The overflow protection may impose a perceptible performance penalty.

Base. **checked_neg**(*x*)

Calculates `-x`, checking for overflow errors where applicable. For example, standard two's complement signed integers (e.g. `Int`) cannot represent `-typemin(Int)`, thus leading to an overflow.

The overflow protection may impose a perceptible performance penalty.

Base. **checked_add**(*x*, *y*)

Calculates `x+y`, checking for overflow errors where applicable.

The overflow protection may impose a perceptible performance penalty.

Base. **checked_sub**(*x*, *y*)

Calculates `x-y`, checking for overflow errors where applicable.

The overflow protection may impose a perceptible performance penalty.

Base. **checked_mul**(*x*, *y*)

Calculates `x*y`, checking for overflow errors where applicable.

The overflow protection may impose a perceptible performance penalty.

`Base.checked_div(x, y)`

Calculates `div(x, y)`, checking for overflow errors where applicable.

The overflow protection may impose a perceptible performance penalty.

`Base.checked_rem(x, y)`

Calculates `x%y`, checking for overflow errors where applicable.

The overflow protection may impose a perceptible performance penalty.

`Base.checked_fld(x, y)`

Calculates `fld(x, y)`, checking for overflow errors where applicable.

The overflow protection may impose a perceptible performance penalty.

`Base.checked_mod(x, y)`

Calculates `mod(x, y)`, checking for overflow errors where applicable.

The overflow protection may impose a perceptible performance penalty.

`Base.checked_cld(x, y)`

Calculates `cld(x, y)`, checking for overflow errors where applicable.

The overflow protection may impose a perceptible performance penalty.

`Base.add_with_overflow(x, y) → (r, f)`

Calculates `r = x+y`, with the flag `f` indicating whether overflow has occurred.

`Base.sub_with_overflow(x, y) → (r, f)`

Calculates `r = x-y`, with the flag `f` indicating whether overflow has occurred.

`Base.mul_with_overflow(x, y) → (r, f)`

Calculates `r = x*y`, with the flag `f` indicating whether overflow has occurred.

`abs2(x)`

Squared absolute value of `x`.

`copysign(x, y) → z`

Return `z` which has the magnitude of `x` and the same sign as `y`.

```
julia> copysign(1, -2)
-1

julia> copysign(-1, 2)
1
```

`sign(x)`

Return zero if `x==0` and `x/|x|` otherwise (i.e., ± 1 for real `x`).

`signbit(x)`

Returns `true` if the value of the sign of `x` is negative, otherwise `false`.

```
julia> signbit(-4)
true

julia> signbit(5)
false

julia> signbit(5.5)
false

julia> signbit(-4.1)
true
```

$\text{sqrt}(x)$ **isqrt** (*n::Integer*)

```
julia> isqrt(5)
2
```

Return the cube root of `x`, i.e. $x^{1/3}$. Negative values are accepted (returning the negative real root when $x < 0$).

[illegible]

Compute the error function of x , defined by $\frac{2}{\sqrt{\pi}} \int_0^x e^{-t^2} dt$ for arbitrary complex x .

Compute the complementary error function of `x`, defined by $1 - \text{erf}(x)$.

Compute the scaled complementary error function of `x`, defined by $e^{x^2} \operatorname{erfc}(x)$. Note also that `erfcx(-ix)` computes the Faddeeva function $w(x)$.

Compute the imaginary error function of `x`, defined by $-i \operatorname{erf}(ix)$.

Compute the Dawson function (scaled imaginary error function) of x , defined by $\frac{\sqrt{\pi}}{2} e^{-x^2} \operatorname{erfi}(x)$.

Compute the inverse error function of a real x , defined by $\text{erf}(\text{erfinv}(x)) = x$.

Compute the inverse error complementary function of a real x , defined by $\operatorname{erfc}(\operatorname{erfcinv}(x)) = x$.

Return the real part of the complex number z .

Return the imaginary part of the complex number z .

Return both the real and imaginary parts of the complex number z .

Compute the complex conjugate of a complex number z .

Compute the phase angle in radians of a complex number z .

Return $\exp(iz)$.

Number of ways to choose k out of n items.

factorial (*n*)

Factorial of *n*. If *n* is an `Integer`, the factorial is computed as an integer (promoted to at least 64 bits). Note that this may overflow if *n* is not small, but you can use `factorial(big(n))` to compute the result exactly in arbitrary precision. If *n* is not an `Integer`, `factorial(n)` is equivalent to `gamma(n+1)` (page 422).

gcd (*x*, *y*)

Greatest common (positive) divisor (or zero if *x* and *y* are both zero).

```
julia> gcd(6, 9)
3

julia> gcd(6, -9)
3
```

lcm (*x*, *y*)

Least common (non-negative) multiple.

```
julia> lcm(2, 3)
6

julia> lcm(-2, 3)
6
```

gcdx (*x*, *y*)

Computes the greatest common (positive) divisor of *x* and *y* and their Bézout coefficients, i.e. the integer coefficients *u* and *v* that satisfy $ux + vy = d = \text{gcd}(x, y)$. `gcdx(x, y)` returns (*d*, *u*, *v*).

```
julia> gcdx(12, 42)
(6, -3, 1)
```

```
julia> gcdx(240, 46)
(2, -9, 47)
```

Note: Bézout coefficients are *not* uniquely defined. `gcdx` returns the minimal Bézout coefficients that are computed by the extended Euclidean algorithm. (Ref: D. Knuth, TAOCP, 2/e, p. 325, Algorithm X.) For signed integers, these coefficients *u* and *v* are minimal in the sense that $|u| < |y/d|$ and $|v| < |x/d|$. Furthermore, the signs of *u* and *v* are chosen so that *d* is positive. For unsigned integers, the coefficients *u* and *v* might be near their `typemax`, and the identity then holds only via the unsigned integers' modulo arithmetic.

ispow2 (*n::Integer*) → Bool

Test whether *n* is a power of two.

```
julia> ispow2(4)
true

julia> ispow2(5)
false
```

nextpow2 (*n::Integer*)

The smallest power of two not less than *n*. Returns 0 for *n*==0, and returns `-nextpow2(-n)` for negative arguments.

```
julia> nextpow2(16)
16

julia> nextpow2(17)
32
```

prevpow2 (*n::Integer*)

The largest power of two not greater than *n*. Returns 0 for *n*==0, and returns `-prevpow2(-n)` for negative arguments.

```
julia> prevpow2(5)
4
```

nextpow (*a*, *x*)

The smallest a^n not less than *x*, where *n* is a non-negative integer. *a* must be greater than 1, and *x* must be greater than 0.

prevpow (*a*, *x*)

The largest a^n not greater than *x*, where *n* is a non-negative integer. *a* must be greater than 1, and *x* must not be less than 1.

nextprod (*[k_1, k_2, ...]*, *n*)

Next integer not less than *n* that can be written as $\prod k_i^{p_i}$ for integers p_1, p_2 , etc.

invmod (*x*, *m*)

Take the inverse of *x* modulo *m*: *y* such that $xy = 1 \pmod{m}$, with $\text{div}(x, y) = 0$. This is undefined for $m = 0$, or if $\text{gcd}(x, m) \neq 1$.

```
julia> invmod(2,5)
3

julia> invmod(2,3)
2

julia> invmod(5,6)
5
```

powermod (*x::Integer*, *p::Integer*, *m*)

Compute $x^p \pmod{m}$.

gamma (*x*)

Compute the gamma function of *x*.

lgamma (*x*)

Compute the logarithm of the absolute value of [gamma\(\)](#) (page 422) for Real *x*, while for Complex *x* it computes the principal branch cut of the logarithm of `gamma(x)` (defined for negative `real(x)` by analytic continuation from positive `real(x)`).

lfact (*x*)

Compute the logarithmic factorial of *x*

digamma (*x*)

Compute the digamma function of *x* (the logarithmic derivative of [gamma\(\)](#) (page 422)).

invdigamma (*x*)

Compute the inverse [digamma\(\)](#) (page 422) function of *x*.

trigamma (*x*)

Compute the trigamma function of *x* (the logarithmic second derivative of [gamma\(\)](#) (page 422)).

polygamma (*m*, *x*)

Compute the polygamma function of order *m* of argument *x* (the $(m+1)$ th derivative of the logarithm of [gamma\(\)](#) (page 422)).

airy (*k*, *x*)

The *k*th derivative of the Airy function $\text{Ai}(x)$.

airyai(x)
Airy function $\text{Ai}(x)$.

airyprime(x)
Airy function derivative $\text{Ai}'(x)$.

airyaiprime(x)
Airy function derivative $\text{Ai}'(x)$.

airybi(x)
Airy function $\text{Bi}(x)$.

airybiprime(x)
Airy function derivative $\text{Bi}'(x)$.

airyx(k, x)
scaled k th derivative of the Airy function, return $\text{Ai}(x)e^{\frac{2}{3}x\sqrt{x}}$ for $k == 0$ || $k == 1$, and $\text{Ai}(x)e^{-|\text{Re}(\frac{2}{3}x\sqrt{x})|}$ for $k == 2$ || $k == 3$.

besselj0(x)
Bessel function of the first kind of order 0, $J_0(x)$.

besselj1(x)
Bessel function of the first kind of order 1, $J_1(x)$.

besselj(ν, x)
Bessel function of the first kind of order ν , $J_\nu(x)$.

besseljx(ν, x)
Scaled Bessel function of the first kind of order ν , $J_\nu(x)e^{-|\text{Im}(x)|}$.

bessely0(x)
Bessel function of the second kind of order 0, $Y_0(x)$.

bessely1(x)
Bessel function of the second kind of order 1, $Y_1(x)$.

bessely(ν, x)
Bessel function of the second kind of order ν , $Y_\nu(x)$.

besselyx(ν, x)
Scaled Bessel function of the second kind of order ν , $Y_\nu(x)e^{-|\text{Im}(x)|}$.

hankelh1(ν, x)
Bessel function of the third kind of order ν , $H_\nu^{(1)}(x)$.

hankelh1x(ν, x)
Scaled Bessel function of the third kind of order ν , $H_\nu^{(1)}(x)e^{-xi}$.

hankelh2(ν, x)
Bessel function of the third kind of order ν , $H_\nu^{(2)}(x)$.

hankelh2x(ν, x)
Scaled Bessel function of the third kind of order ν , $H_\nu^{(2)}(x)e^{xi}$.

besselh($\nu, k=1$), x)
Bessel function of the third kind of order ν (the Hankel function). k is either 1 or 2, selecting [hankelh1\(\)](#) (page 423) or [hankelh2\(\)](#) (page 423), respectively. k defaults to 1 if it is omitted. (See also [besselhx\(\)](#) (page 423) for an exponentially scaled variant.)

besselhx($\nu, k=1$), z)
Compute the scaled Hankel function $\exp(\mp iz)H_\nu^{(k)}(z)$, where k is 1 or 2, $H_\nu^{(k)}(z)$ is [besselh](#)(ν , k , z), and \mp is $-$ for $k = 1$ and $+$ for $k = 2$. k defaults to 1 if it is omitted.

The reason for this function is that $H_\nu^{(k)}(z)$ is asymptotically proportional to $\exp(\mp iz)/\sqrt{z}$ for large $|z|$, and so the `besselh()` (page 423) function is susceptible to overflow or underflow when z has a large imaginary part. The `besselhx` function cancels this exponential factor (analytically), so it avoids these problems.

besseli (nu, x)

Modified Bessel function of the first kind of order ν , $I_\nu(x)$.

besselix (nu, x)Scaled modified Bessel function of the first kind of order nu, $I_\nu(x)e^{-|\operatorname{Re}(x)|}$.**besselk** (nu, x)

Modified Bessel function of the second kind of order nu, $K_\nu(x)$.

besselkx (*nu*, *x*)

Scaled modified Bessel function of the second kind of order ν , $K_\nu(x)e^x$.

beta (x, y)

Euler integral of the first kind $B(x, y) = \Gamma(x)\Gamma(y)/\Gamma(x + y)$.

$$\text{lbeta}(x, y)$$

Natural logarithm of the absolute value of the `beta()` (page 424) function $\log(|B(x, y)|)$.

 $\eta(x)$

Dirichlet eta function $\eta(s) = \sum_{n=1}^{\infty} (-1)^{n-1}/n^s$.

 $\zeta(s)$

Riemann zeta function $\zeta(s)$.

zeta (s, z)

Generalized zeta function $\zeta(s, z)$, defined by the sum $\sum_{k=0}^{\infty} ((k+z)^2)^{-s/2}$, where any term with $k+z=0$ is excluded. For $\Re z > 0$, this definition is equivalent to the Hurwitz zeta function $\sum_{k=0}^{\infty} (k+z)^{-s}$. For $z=1$, it yields the Riemann zeta function $\zeta(s)$.

ndigits ($n::Integer, b::Integer=10$)

Compute the number of digits in integer n written in base b .

widemul (x, y)

Multiply x and y , giving the result as a larger type.

[illegible]

@evalpoly ($z, c...$)

Evaluate the polynomial $\sum_k c[k]z^{k-1}$ for the coefficients $c[1], c[2], \dots$; that is, the coefficients are given in ascending order by power of z . This macro expands to efficient inline code that uses either Horner's method or, for complex z , a more efficient Goertzel-like algorithm.

```
julia> @evalpoly(3, 1, 0, 1)
10

julia> @evalpoly(2, 1, 0, 1)
5

julia> @evalpoly(2, 1, 1, 1)
7
```

42.3 Statistics

mean (v [, *region*])

Compute the mean of whole array v , or optionally along the dimensions in *region*.

Note: Julia does not ignore NaN values in the computation. For applications requiring the handling of missing data, the `DataArrays.jl` package is recommended.

mean ($f::Function, v$)

Apply the function f to each element of v and take the mean.

mean! (r, v)

Compute the mean of v over the singleton dimensions of r , and write results to r .

std (v [, *region*]; *corrected::Bool*=true, *mean=nothing*)

Compute the sample standard deviation of a vector or array v , optionally along dimensions in *region*. The algorithm returns an estimator of the generative distribution's standard deviation under the assumption that each entry of v is an IID drawn from that generative distribution. This computation is equivalent to calculating $\sqrt{\text{sum}((v - \text{mean}(v)).^2) / (\text{length}(v) - 1)}$. A pre-computed mean may be provided. If *corrected* is true, then the sum is scaled with $n-1$, whereas the sum is scaled with n if *corrected* is false where $n = \text{length}(x)$.

Note: Julia does not ignore NaN values in the computation. For applications requiring the handling of missing data, the `DataArrays.jl` package is recommended.

stdm ($v, m::Number$; *corrected::Bool*=true)

Compute the sample standard deviation of a vector v with known mean m . If *corrected* is true, then the sum is scaled with $n-1$, whereas the sum is scaled with n if *corrected* is false where $n = \text{length}(x)$.

Note: Julia does not ignore NaN values in the computation. For applications requiring the handling of missing data, the `DataArrays.jl` package is recommended.

var (v [, *region*]; *corrected::Bool*=true, *mean=nothing*)

Compute the sample variance of a vector or array v , optionally along dimensions in *region*. The algorithm will return an estimator of the generative distribution's variance under the assumption that each entry of v is an IID drawn from that generative distribution. This computation is equivalent to calculating $\text{sumabs2}(v - \text{mean}(v)) / (\text{length}(v) - 1)$. If *corrected* is true, then the sum is scaled with $n-1$, whereas the sum is scaled with n if *corrected* is false where $n = \text{length}(x)$. The mean over the region may be provided.

Note: Julia does not ignore NaN values in the computation. For applications requiring the handling of missing data, the `DataArrays.jl` package is recommended.

varm (v, m [, *region*]; *corrected::Bool*=true)

Compute the sample variance of a collection v with known mean(s) m , optionally over *region*. m may contain means for each dimension of v . If *corrected* is true, then the sum is scaled with $n-1$, whereas the sum is scaled with n if *corrected* is false where $n = \text{length}(x)$.

Note: Julia does not ignore NaN values in the computation. For applications requiring the handling of missing data, the `DataArrays.jl` package is recommended.

data, the `DataArrays.jl` package is recommended.

middle(*x*)

Compute the middle of a scalar value, which is equivalent to *x* itself, but of the type of `middle(x, x)` for consistency.

middle(*x*, *y*)

Compute the middle of two reals *x* and *y*, which is equivalent in both value and type to computing their mean $((x + y) / 2)$.

middle(*range*)

Compute the middle of a range, which consists of computing the mean of its extrema. Since a range is sorted, the mean is performed with the first and last element.

```
julia> middle(1:10)
5.5
```

middle(*a*)

Compute the middle of an array *a*, which consists of finding its extrema and then computing their mean.

```
julia> a = [1,2,3.6,10.9]
4-element Array{Float64,1}:
 1.0
 2.0
 3.6
10.9

julia> middle(a)
5.95
```

median(*v*[, *region*])

Compute the median of an entire array *v*, or, optionally, along the dimensions in *region*. For an even number of elements no exact median element exists, so the result is equivalent to calculating mean of two median elements.

Note: Julia does not ignore NaN values in the computation. For applications requiring the handling of missing data, the `DataArrays.jl` package is recommended.

median!(*v*)

Like `median()` (page 426), but may overwrite the input vector.

midpoints(*e*)

Compute the midpoints of the bins with edges *e*. The result is a vector/range of length `length(e) - 1`. Note: Julia does not ignore NaN values in the computation.

quantile(*v*, *p*; *sorted=false*)

Compute the quantile(s) of a vector *v* at a specified probability or vector *p*. The keyword argument *sorted* indicates whether *v* can be assumed to be sorted.

The *p* should be on the interval [0,1], and *v* should not have any NaN values.

Quantiles are computed via linear interpolation between the points $((k-1)/(n-1), v[k])$, for $k = 1:n$ where $n = \text{length}(v)$. This corresponds to Definition 7 of Hyndman and Fan (1996), and is the same as the R default.

Note: Julia does not ignore NaN values in the computation. For applications requiring the handling of missing data, the `DataArrays.jl` package is recommended. `quantile` will throw an `ArgumentError` in the

presence of NaN values in the data array.

- Hyndman, R.J and Fan, Y. (1996) “Sample Quantiles in Statistical Packages”, *The American Statistician*, Vol. 50, No. 4, pp. 361-365

quantile! ($[q]$, v , p ; $sorted=false$)

Compute the quantile(s) of a vector v at the probabilities p , with optional output into array q (if not provided, a new output array is created). The keyword argument `sorted` indicates whether v can be assumed to be sorted; if `false` (the default), then the elements of v may be partially sorted.

The elements of p should be on the interval $[0,1]$, and v should not have any NaN values.

Quantiles are computed via linear interpolation between the points $((k-1)/(n-1), v[k])$, for $k = 1:n$ where $n = \text{length}(v)$. This corresponds to Definition 7 of Hyndman and Fan (1996), and is the same as the R default.

Note: Julia does not ignore NaN values in the computation. For applications requiring the handling of missing data, the `DataArrays.jl` package is recommended. `quantile!` will throw an `ArgumentError` in the presence of NaN values in the data array.

- Hyndman, R.J and Fan, Y. (1996) “Sample Quantiles in Statistical Packages”, *The American Statistician*, Vol. 50, No. 4, pp. 361-365

cov (x [, $corrected=true$])

Compute the variance of the vector x . If `corrected` is `true` (the default) then the sum is scaled with $n-1$, whereas the sum is scaled with n if `corrected` is `false` where $n = \text{length}(x)$.

cov (X [, $varlim=1$, $corrected=true$])

Compute the covariance matrix of the matrix X along the dimension `varlim`. If `corrected` is `true` (the default) then the sum is scaled with $n-1$, whereas the sum is scaled with n if `corrected` is `false` where $n = \text{size}(X, \text{varlim})$.

cov (x , y [, $corrected=true$])

Compute the covariance between the vectors x and y . If `corrected` is `true` (the default) then the sum is scaled with $n-1$, whereas the sum is scaled with n if `corrected` is `false` where $n = \text{length}(x) = \text{length}(y)$.

cov (X , Y [, $varlim=1$, $corrected=true$])

Compute the covariance between the vectors or matrices X and Y along the dimension `varlim`. If `corrected` is `true` (the default) then the sum is scaled with $n-1$, whereas the sum is scaled with n if `corrected` is `false` where $n = \text{size}(X, \text{varlim}) = \text{size}(Y, \text{varlim})$.

cor (x)

Return the number one.

cor (X [, $varlim=1$])

Compute the Pearson correlation matrix of the matrix X along the dimension `varlim`.

cor (x , y)

Compute the Pearson correlation between the vectors x and y .

cor (X , Y [, $varlim=1$])

Compute the Pearson correlation between the vectors or matrices X and Y along the dimension `varlim`.

42.4 Signal Processing

Fast Fourier transform (FFT) functions in Julia are implemented by calling functions from `FFTW`.

fft ($A[, \text{dims}]$)

Performs a multidimensional FFT of the array A . The optional `dims` argument specifies an iterable subset of dimensions (e.g. an integer, range, tuple, or array) to transform along. Most efficient if the size of A along the transformed dimensions is a product of small primes; see `nextprod()`. See also `plan_fft()` for even greater efficiency.

A one-dimensional FFT computes the one-dimensional discrete Fourier transform (DFT) as defined by

$$\text{DFT}(A)[k] = \sum_{n=1}^{\text{length}(A)} \exp\left(-i \frac{2\pi(n-1)(k-1)}{\text{length}(A)}\right) A[n].$$

A multidimensional FFT simply performs this operation along each transformed dimension of A .

Note:

- Julia starts FFTW up with 1 thread by default. Higher performance is usually possible by increasing number of threads. Use `FFTW.set_num_threads(Sys.CPU_CORES)` to use as many threads as cores on your system.
 - This performs a multidimensional FFT by default. FFT libraries in other languages such as Python and Octave perform a one-dimensional FFT along the first non-singleton dimension of the array. This is worth noting while performing comparisons. For more details, refer to the *Noteworthy Differences from other Languages* (page 347) section of the manual.
-

fft! ($A[, \text{dims}]$)

Same as `fft()` (page 428), but operates in-place on A , which must be an array of complex floating-point numbers.

ifft ($A[, \text{dims}]$)

Multidimensional inverse FFT.

A one-dimensional inverse FFT computes

$$\text{IDFT}(A)[k] = \frac{1}{\text{length}(A)} \sum_{n=1}^{\text{length}(A)} \exp\left(+i \frac{2\pi(n-1)(k-1)}{\text{length}(A)}\right) A[n].$$

A multidimensional inverse FFT simply performs this operation along each transformed dimension of A .

ifft! ($A[, \text{dims}]$)

Same as `ifft()` (page 428), but operates in-place on A .

bfft ($A[, \text{dims}]$)

Similar to `ifft()` (page 428), but computes an unnormalized inverse (backward) transform, which must be divided by the product of the sizes of the transformed dimensions in order to obtain the inverse. (This is slightly more efficient than `ifft()` (page 428) because it omits a scaling step, which in some applications can be combined with other computational steps elsewhere.)

$$\text{BDFT}(A)[k] = \text{length}(A) \text{IDFT}(A)[k]$$

bfft! ($A[, \text{dims}]$)

Same as `bfft()` (page 428), but operates in-place on A .

plan_fft (*A* [, *dims*]; *flags*=FFTW.ESTIMATE; *timelimit*=Inf)

Pre-plan an optimized FFT along given dimensions (*dims*) of arrays matching the shape and type of *A*. (The first two arguments have the same meaning as for `fft()` (page 428).) Returns an object *P* which represents the linear operator computed by the FFT, and which contains all of the information needed to compute `fft(A, dims)` quickly.

To apply *P* to an array *A*, use `P * A`; in general, the syntax for applying plans is much like that of matrices. (A plan can only be applied to arrays of the same size as the *A* for which the plan was created.) You can also apply a plan with a preallocated output array \hat{A} by calling `A_mul_B!(\hat{A} , plan, A)`. (For `A_mul_B!`, however, the input array *A* must be a complex floating-point array like the output \hat{A} .) You can compute the inverse-transform plan by `inv(P)` and apply the inverse plan with `P \ \hat{A}` (the inverse plan is cached and reused for subsequent calls to `inv` or `\`), and apply the inverse plan to a pre-allocated output array *A* with `A_ldiv_B!(A, P, \hat{A})`.

The *flags* argument is a bitwise-or of FFTW planner flags, defaulting to `FFTW.ESTIMATE`. e.g. passing `FFTW.MEASURE` or `FFTW.PATIENT` will instead spend several seconds (or more) benchmarking different possible FFT algorithms and picking the fastest one; see the FFTW manual for more information on planner flags. The optional *timelimit* argument specifies a rough upper bound on the allowed planning time, in seconds. Passing `FFTW.MEASURE` or `FFTW.PATIENT` may cause the input array *A* to be overwritten with zeros during plan creation.

`plan_fft!()` (page 429) is the same as `plan_fft()` (page 428) but creates a plan that operates in-place on its argument (which must be an array of complex floating-point numbers). `plan_iff()` (page 429) and so on are similar but produce plans that perform the equivalent of the inverse transforms `iff()` (page 428) and so on.

plan_iff (*A* [, *dims*]; *flags*=FFTW.ESTIMATE; *timelimit*=Inf)

Same as `plan_fft()` (page 428), but produces a plan that performs inverse transforms `iff()` (page 428).

plan_bfft (*A* [, *dims*]; *flags*=FFTW.ESTIMATE; *timelimit*=Inf)

Same as `plan_fft()` (page 428), but produces a plan that performs an unnormalized backwards transform `bfft()` (page 428).

plan_fft! (*A* [, *dims*]; *flags*=FFTW.ESTIMATE; *timelimit*=Inf)

Same as `plan_fft()` (page 428), but operates in-place on *A*.

plan_iff! (*A* [, *dims*]; *flags*=FFTW.ESTIMATE; *timelimit*=Inf)

Same as `plan_iff()` (page 429), but operates in-place on *A*.

plan_bfft! (*A* [, *dims*]; *flags*=FFTW.ESTIMATE; *timelimit*=Inf)

Same as `plan_bfft()` (page 429), but operates in-place on *A*.

rfft (*A* [, *dims*])

Multidimensional FFT of a real array *A*, exploiting the fact that the transform has conjugate symmetry in order to save roughly half the computational time and storage costs compared with `fft()` (page 428). If *A* has size (*n*₁, ..., *n*_{*d*}), the result has size (*div*(*n*₁, 2)+1, ..., *n*_{*d*}).

The optional *dims* argument specifies an iterable subset of one or more dimensions of *A* to transform, similar to `fft()` (page 428). Instead of (roughly) halving the first dimension of *A* in the result, the *dims*[1] dimension is (roughly) halved in the same way.

irfft (*A*, *d* [, *dims*])

Inverse of `rfft()` (page 429): for a complex array *A*, gives the corresponding real array whose FFT yields *A* in the first half. As for `rfft()` (page 429), *dims* is an optional subset of dimensions to transform, defaulting to `1:ndims(A)`.

d is the length of the transformed real array along the *dims*[1] dimension, which must satisfy `div(d, 2)+1 == size(A, dims[1])`. (This parameter cannot be inferred from `size(A)` since both `2*size(A, dims[1])-2` as well as `2*size(A, dims[1])-1` are valid sizes for the transformed real array.)

brfft (*A*, *d* [, *dims*])

Similar to *irfft* () (page 429) but computes an unnormalized inverse transform (similar to *bfft* () (page 428)), which must be divided by the product of the sizes of the transformed dimensions (of the real output array) in order to obtain the inverse transform.

plan_rfft (*A* [, *dims*]; *flags*=*FFTW.ESTIMATE*; *timelimit*=*Inf*)

Pre-plan an optimized real-input FFT, similar to *plan_fft* () (page 428) except for *rfft* () (page 429) instead of *fft* () (page 428). The first two arguments, and the size of the transformed result, are the same as for *rfft* () (page 429).

plan_brfft (*A*, *d* [, *dims*]; *flags*=*FFTW.ESTIMATE*; *timelimit*=*Inf*)

Pre-plan an optimized real-input unnormalized transform, similar to *plan_rfft* () (page 430) except for *brfft* () (page 429) instead of *rfft* () (page 429). The first two arguments and the size of the transformed result, are the same as for *brfft* () (page 429).

plan_irfft (*A*, *d* [, *dims*]; *flags*=*FFTW.ESTIMATE*; *timelimit*=*Inf*)

Pre-plan an optimized inverse real-input FFT, similar to *plan_rfft* () (page 430) except for *irfft* () (page 429) and *brfft* () (page 429), respectively. The first three arguments have the same meaning as for *irfft* () (page 429).

dct (*A* [, *dims*])

Performs a multidimensional type-II discrete cosine transform (DCT) of the array *A*, using the unitary normalization of the DCT. The optional *dims* argument specifies an iterable subset of dimensions (e.g. an integer, range, tuple, or array) to transform along. Most efficient if the size of *A* along the transformed dimensions is a product of small primes; see *nextprod* () (page 422). See also *plan_dct* () (page 430) for even greater efficiency.

dct! (*A* [, *dims*])

Same as *dct*! () (page 430), except that it operates in-place on *A*, which must be an array of real or complex floating-point values.

idct (*A* [, *dims*])

Computes the multidimensional inverse discrete cosine transform (DCT) of the array *A* (technically, a type-III DCT with the unitary normalization). The optional *dims* argument specifies an iterable subset of dimensions (e.g. an integer, range, tuple, or array) to transform along. Most efficient if the size of *A* along the transformed dimensions is a product of small primes; see *nextprod* () (page 422). See also *plan_idct* () (page 430) for even greater efficiency.

idct! (*A* [, *dims*])

Same as *idct*! () (page 430), but operates in-place on *A*.

plan_dct (*A* [, *dims* [, *flags* [, *timelimit*]]])

Pre-plan an optimized discrete cosine transform (DCT), similar to *plan_fft* () (page 428) except producing a function that computes *dct* () (page 430). The first two arguments have the same meaning as for *dct* () (page 430).

plan_dct! (*A* [, *dims* [, *flags* [, *timelimit*]]])

Same as *plan_dct* () (page 430), but operates in-place on *A*.

plan_idct (*A* [, *dims* [, *flags* [, *timelimit*]]])

Pre-plan an optimized inverse discrete cosine transform (DCT), similar to *plan_fft* () (page 428) except producing a function that computes *idct* () (page 430). The first two arguments have the same meaning as for *idct* () (page 430).

plan_idct! (*A* [, *dims* [, *flags* [, *timelimit*]]])

Same as *plan_idct* () (page 430), but operates in-place on *A*.

fftshift (*x*)

Swap the first and second halves of each dimension of *x*.

fftshift (*x*, *dim*)

Swap the first and second halves of the given dimension of array *x*.

ifftshift (*x*[, *dim*])

Undoes the effect of `fftshift`.

filt (*b*, *a*, *x*[, *si*])

Apply filter described by vectors *a* and *b* to vector *x*, with an optional initial filter state vector *si* (defaults to zeros).

filt! (*out*, *b*, *a*, *x*[, *si*])

Same as `filt()` (page 431) but writes the result into the *out* argument, which may alias the input *x* to modify it in-place.

deconv (*b*, *a*) → *c*

Construct vector *c* such that $b = \text{conv}(a, c) + r$. Equivalent to polynomial division.

conv (*u*, *v*)

Convolution of two vectors. Uses FFT algorithm.

conv2 (*u*, *v*, *A*)

2-D convolution of the matrix *A* with the 2-D separable kernel generated by the vectors *u* and *v*. Uses 2-D FFT algorithm.

conv2 (*B*, *A*)

2-D convolution of the matrix *B* with the matrix *A*. Uses 2-D FFT algorithm.

xcorr (*u*, *v*)

Compute the cross-correlation of two vectors.

The following functions are defined within the `Base.FFTW` module.

r2r (*A*, *kind*[, *dims*])

Performs a multidimensional real-input/real-output (r2r) transform of type *kind* of the array *A*, as defined in the FFTW manual. *kind* specifies either a discrete cosine transform of various types (`FFTW.REDFT00`, `FFTW.REDFT01`, `FFTW.REDFT10`, or `FFTW.REDFT11`), a discrete sine transform of various types (`FFTW.RODFT00`, `FFTW.RODFT01`, `FFTW.RODFT10`, or `FFTW.RODFT11`), a real-input DFT with halfcomplex-format output (`FFTW.R2HC` and its inverse `FFTW.HC2R`), or a discrete Hartley transform (`FFTW.DHT`). The *kind* argument may be an array or tuple in order to specify different transform types along the different dimensions of *A*; *kind*[*end*] is used for any unspecified dimensions. See the FFTW manual for precise definitions of these transform types, at <http://www.fftw.org/doc>.

The optional *dims* argument specifies an iterable subset of dimensions (e.g. an integer, range, tuple, or array) to transform along. *kind*[*i*] is then the transform type for *dims*[*i*], with *kind*[*end*] being used for *i* > `length(kind)`.

See also `plan_r2r()` (page 431) to pre-plan optimized r2r transforms.

r2r! (*A*, *kind*[, *dims*])

Same as `r2r()` (page 431), but operates in-place on *A*, which must be an array of real or complex floating-point numbers.

plan_r2r (*A*, *kind*[, *dims*[, *flags*[, *timelimit*]]])

Pre-plan an optimized r2r transform, similar to `Base.plan_fft()` (page 428) except that the transforms (and the first three arguments) correspond to `r2r()` (page 431) and `r2r!` (page 431), respectively.

plan_r2r! (*A*, *kind*[, *dims*[, *flags*[, *timelimit*]]])

Similar to `Base.plan_fft()` (page 428), but corresponds to `r2r!` (page 431).

42.5 Numerical Integration

Although several external packages are available for numeric integration and solution of ordinary differential equations, we also provide some built-in integration support in Julia.

quadgk (*f*, *a*, *b*, *c*,...; *reltol*=*sqrt(eps)*, *abstol*=0, *maxevals*= 10^7 , *order*=7, *norm*=*vecnorm*)

Numerically integrate the function $f(x)$ from *a* to *b*, and optionally over additional intervals *b* to *c* and so on. Keyword options include a relative error tolerance *reltol* (defaults to *sqrt(eps)* in the precision of the endpoints), an absolute error tolerance *abstol* (defaults to 0), a maximum number of function evaluations *maxevals* (defaults to 10^7), and the *order* of the integration rule (defaults to 7).

Returns a pair (*I*,*E*) of the estimated integral *I* and an estimated upper bound on the absolute error *E*. If *maxevals* is not exceeded then $E \leq \max(\text{abstol}, \text{reltol} \cdot \text{norm}(I))$ will hold. (Note that it is useful to specify a positive *abstol* in cases where *norm(I)* may be zero.)

The endpoints *a* et cetera can also be complex (in which case the integral is performed over straight-line segments in the complex plane). If the endpoints are `BigFloat`, then the integration will be performed in `BigFloat` precision as well.

Note: It is advisable to increase the integration *order* in rough proportion to the precision, for smooth integrands.

More generally, the precision is set by the precision of the integration endpoints (promoted to floating-point types).

The integrand $f(x)$ can return any numeric scalar, vector, or matrix type, or in fact any type supporting $+$, $-$, multiplication by real values, and a *norm* (i.e., any normed vector space). Alternatively, a different norm can be specified by passing a norm-like function as the *norm* keyword argument (which defaults to *vecnorm*).

Note: Only one-dimensional integrals are provided by this function. For multi-dimensional integration (cubature), there are many different algorithms (often much better than simple nested 1d integrals) and the optimal choice tends to be very problem-dependent. See the Julia external-package listing for available algorithms for multidimensional integration or other specialized tasks (such as integrals of highly oscillatory or singular functions).

The algorithm is an adaptive Gauss-Kronrod integration technique: the integral in each interval is estimated using a Kronrod rule ($2 \cdot \text{order} + 1$ points) and the error is estimated using an embedded Gauss rule (*order* points). The interval with the largest error is then subdivided into two intervals and the process is repeated until the desired error tolerance is achieved.

These quadrature rules work best for smooth functions within each interval, so if your function has a known discontinuity or other singularity, it is best to subdivide your interval to put the singularity at an endpoint. For example, if f has a discontinuity at $x=0.7$ and you want to integrate from 0 to 1, you should use `quadgk(f, 0, 0.7, 1)` to subdivide the interval at the point of discontinuity. The integrand is never evaluated exactly at the endpoints of the intervals, so it is possible to integrate functions that diverge at the endpoints as long as the singularity is integrable (for example, a $\log(x)$ or $1/\sqrt{x}$ singularity).

For real-valued endpoints, the starting and/or ending points may be infinite. (A coordinate transformation is performed internally to map the infinite interval to a finite one.)

Numbers

43.1 Standard Numeric Types

Bool Int8 UInt8 Int16 UInt16 Int32 UInt32 Int64 UInt64 Int128 UInt128 Float16 Float32
Float64 Complex64 Complex128

43.2 Data Formats

bin (*n*, *pad*::Int=1)

Convert an integer to a binary string, optionally specifying a number of digits to pad to.

```
julia> bin(10,2)
"1010"

julia> bin(10,8)
"00001010"
```

hex (*n*, *pad*::Int=1)

Convert an integer to a hexadecimal string, optionally specifying a number of digits to pad to.

dec (*n*, *pad*::Int=1)

Convert an integer to a decimal string, optionally specifying a number of digits to pad to.

oct (*n*, *pad*::Int=1)

Convert an integer to an octal string, optionally specifying a number of digits to pad to.

base (*base*::Integer, *n*::Integer, *pad*::Integer=1)

Convert an integer *n* to a string in the given *base*, optionally specifying a number of digits to pad to.

```
julia> base(13,5,4)
"0005"

julia> base(5,13,4)
"0023"
```

digits (*T*<::Integer, *n*::Integer, *base*::T=10, *pad*::Integer=1)

Returns an array with element type *T* (default Int) of the digits of *n* in the given base, optionally padded with zeros to a specified size. More significant digits are at higher indexes, such that $n == \sum([digits[k] * base^{(k-1)} \text{ for } k=1:\text{length}(digits)])$.

digits! (*array*, *n*::Integer, *base*::Integer=10)

Fills an array of the digits of *n* in the given base. More significant digits are at higher indexes. If the array length

is insufficient, the least significant digits are filled up to the array length. If the array length is excessive, the excess portion is filled with zeros.

bits (n)

A string giving the literal bit representation of a number.

[illegible]

parse (*type*, *str*_[, *base*])

Parse a string as a number. If the type is an integer type, then a base can be specified (the default is 10). If the type is a floating point type, the string is parsed as a decimal floating point number. If the string does not contain a valid number, an error is raised.

tryparse (*type*, *str*_[, *base*])

Like `parse()` (page 369), but returns a `Nullable` (page 369) of the requested type. The result will be null if the string does not contain a valid number.

$$\mathbf{big}(x)$$

Convert a number to a maximum precision representation (typically `BigInt` or `BigFloat`). See [BigFloat](#) (page 437) for information about some pitfalls with floating-point numbers.

signed (x)

Convert a number to a signed integer. If the argument is unsigned, it is reinterpreted as signed without checking for overflow.

$$\text{unsigned}(x) \rightarrow \text{Unsigned}$$

Convert a number to an unsigned integer. If the argument is signed, it is reinterpreted as unsigned without checking for negative values.

float (x)

Convert a number or array to a floating point data type. When passed a string, this function is equivalent to `parse(Float64, x)`.

 $\text{significand}(x)$

Extract the significand(s) (a.k.a. mantissa), in binary representation, of a floating-point number. If x is a non-zero finite number, then the result will be a number of the same type on the interval $[1, 2)$. Otherwise x is returned.

```
julia> signifcand(15.2)/15.2
0.125

julia> signifcand(15.2)*8
15.2
```

$$\text{exponent}(x) \rightarrow \text{Int}$$

Get the exponent of a normalized floating-point number.

$$\text{complex}(r[, i])$$

Convert real numbers or arrays to complex. `i` defaults to zero.

bswap (n)

Byte-swap an integer.

 $\text{num2hex}(f)$

Get a hexadecimal string of the binary representation of a floating point number.


```
julia> num2hex(2.2)
"4001999999999999a"
```

hex2num (*str*)

Convert a hexadecimal string to the floating point number it represents.

hex2bytes (*s::AbstractString*)

Convert an arbitrarily long hexadecimal string to its binary representation. Returns an `Array{UInt8,1}`, i.e. an array of bytes.

```
julia> a = hex(12345)
"3039"

julia> hex2bytes(a)
2-element Array{UInt8,1}:
 0x30
 0x39
```

bytes2hex (*bin_arr::Array{UInt8,1}*) → *String*

Convert an array of bytes to its hexadecimal representation. All characters are in lower-case.

```
julia> a = hex(12345)
"3039"

julia> b = hex2bytes(a)
2-element Array{UInt8,1}:
 0x30
 0x39

julia> bytes2hex(b)
"3039"
```

43.3 General Number Functions and Constants

one (*x*)

Get the multiplicative identity element for the type of *x* (*x* can also specify the type itself). For matrices, returns an identity matrix of the appropriate size and type.

zero (*x*)

Get the additive identity element for the type of *x* (*x* can also specify the type itself).

pi
 π

The constant pi.

im

The imaginary unit.

e**eu**

The constant e.

catalan

Catalan's constant.

 γ
eulergamma

Euler's constant.

ϕ **golden**

The golden ratio.

InfPositive infinity of type `Float64`.**Inf32**Positive infinity of type `Float32`.**Inf16**Positive infinity of type `Float16`.**NaN**A not-a-number value of type `Float64`.**NaN32**A not-a-number value of type `Float32`.**NaN16**A not-a-number value of type `Float16`.**issubnormal** (*f*) → `Bool`

Test whether a floating point number is subnormal.

isfinite (*f*) → `Bool`

Test whether a number is finite.

```
julia> isfinite(5)
true

julia> isfinite(NaN32)
false
```

isinf (*f*) → `Bool`

Test whether a number is infinite.

isnan (*f*) → `Bool`

Test whether a floating point number is not a number (NaN).

nextfloat (*x::AbstractFloat*)Returns the smallest floating point number *y* of the same type as *x* such *x* < *y*. If no such *y* exists (e.g. if *x* is `Inf` or `NaN`), then returns *x*.**prevfloat** (*x::AbstractFloat*)Returns the largest floating point number *y* of the same type as *x* such *y* < *x*. If no such *y* exists (e.g. if *x* is `-Inf` or `NaN`), then returns *x*.**nextfloat** (*x::AbstractFloat*, *n::Integer*)The result of *n* iterative applications of `nextfloat` to *x* if *n* ≥ 0, or `-n` applications of `prevfloat` if *n* < 0.**isinteger** (*x*) → `Bool`Test whether *x* or all its elements are numerically equal to some integer.

```
julia> isinteger(4.0)
true

julia> isinteger([1; 2; 5.5])
false
```

isreal (*x*) → `Bool`Test whether *x* or all its elements are numerically equal to some real number.

```
julia> isreal(5.)
true

julia> isreal([4.; complex(0,1)])
false
```

$$\mathbf{isimag}(z) \rightarrow \text{Bool}$$

Test whether z is purely imaginary, i.e. has a real part equal to 0.

Float32 (x [, *mode::RoundingMode*])

Create a Float32 from `x`. If `x` is not exactly representable then `mode` determines how `x` is rounded.

```
julia> Float32(1/3, RoundDown)
0.3333333f0

julia> Float32(1/3, RoundUp)
0.33333334f0
```

See *RoundingMode* (page 416) for available rounding modes.

```
Float64 (x[, mode::RoundingMode])
```

Create a Float64 from `x`. If `x` is not exactly representable then `mode` determines how `x` is rounded.

```
julia> Float64(pi, RoundDown)
3.141592653589793

julia> Float64(pi, RoundUp)
3.1415926535897936
```

See *RoundingMode* (page 416) for available rounding modes.

 $\text{BigInt}(x)$

Create an arbitrary precision integer. `x` may be an `Int` (or anything that can be converted to an `Int`). The usual mathematical operators are defined for this type, and results are promoted to a `BigInt`.

Instances can be constructed from strings via `parse()` (page 369), or using the `big` string literal.

BigFloat (x)

Create an arbitrary precision floating point number. `x` may be an `Integer`, a `Float64` or a `BigInt`. The usual mathematical operators are defined for this type, and results are promoted to a `BigFloat`.

Note that because decimal literals are converted to floating point numbers when parsed, `BigFloat(2.1)` may not yield what you expect. You may instead prefer to initialize constants from strings via `parse()` (page 369), or using the `big` string literal.

[illegible] $\text{rounding}(T)$

Get the current floating point rounding mode for type T, controlling the rounding of basic arithmetic functions (`+`) (page 405), `-` (page 405), `*` (page 443), `/` (page 405) and `sqrt` (page 420)) and type conversion.

See *RoundingMode* (page 416) for available modes.

`setrounding(T, mode)`

Set the rounding mode of floating point type T, controlling the rounding of basic arithmetic functions (`+` (page 405), `-` (page 405), `*` (page 443), `/` (page 405) and `sqrt` (page 420)) and type conversion.

Other numerical functions may give incorrect or invalid values when using rounding modes other than the default `RoundNearest`.

Note that this may affect other types, for instance changing the rounding mode of `Float64` will change the rounding mode of `Float32`. See [RoundingMode](#) (page 416) for available modes.

Warning: This feature is still experimental, and may give unexpected or incorrect values.

setrounding (*f::Function, T, mode*)

Change the rounding mode of floating point type `T` for the duration of `f`. It is logically equivalent to:

```
old = rounding(T)
setrounding(T, mode)
f()
setrounding(T, old)
```

See [RoundingMode](#) (page 416) for available rounding modes.

Warning: This feature is still experimental, and may give unexpected or incorrect values. A known problem is the interaction with compiler optimisations, e.g.

```
julia> setrounding(Float64, RoundDown) do
    1.1 + 0.1
end
1.2000000000000002
```

Here the compiler is *constant folding*, that is evaluating a known constant expression at compile time, however the rounding mode is only changed at runtime, so this is not reflected in the function result. This can be avoided by moving constants outside the expression, e.g.

```
julia> x = 1.1; y = 0.1;

julia> setrounding(Float64, RoundDown) do
    x + y
end
1.2
```

get_zero_subnormals () → Bool

Returns `false` if operations on subnormal floating-point values (“denormals”) obey rules for IEEE arithmetic, and `true` if they might be converted to zeros.

set_zero_subnormals (*yes::Bool*) → Bool

If `yes` is `false`, subsequent floating-point operations follow rules for IEEE arithmetic on subnormal values (“denormals”). Otherwise, floating-point operations are permitted (but not required) to convert subnormal inputs or outputs to zero. Returns `true` unless `yes==true` but the hardware does not support zeroing of subnormal numbers.

`set_zero_subnormals(true)` can speed up some computations on some hardware. However, it can break identities such as `(x-y==0) == (x==y)`.

43.3.1 Integers

count_ones (*x::Integer*) → Integer

Number of ones in the binary representation of `x`.

```
julia> count_ones(7)
3
```

count_zeros ($x::Integer$) \rightarrow Integer
Number of zeros in the binary representation of x .

```
julia> count_zeros(Int32(2 ^ 16 - 1))
16
```

leading_zeros ($x::Integer$) \rightarrow Integer
Number of zeros leading the binary representation of x .

```
julia> leading_zeros(Int32(1))
31
```

leading_ones ($x::Integer$) \rightarrow Integer
Number of ones leading the binary representation of x .

```
julia> leading_ones(UInt32(2 ^ 32 - 2))
31
```

trailing_zeros ($x::Integer$) \rightarrow Integer
Number of zeros trailing the binary representation of x .

```
julia> trailing_zeros(2)
1
```

trailing_ones ($x::Integer$) \rightarrow Integer
Number of ones trailing the binary representation of x .

```
julia> trailing_ones(3)
2
```

isodd ($x::Integer$) \rightarrow Bool
Returns `true` if x is odd (that is, not divisible by 2), and `false` otherwise.

```
julia> isodd(9)
true

julia> isodd(10)
false
```

iseven ($x::Integer$) \rightarrow Bool
Returns `true` if x is even (that is, divisible by 2), and `false` otherwise.

```
julia> iseven(9)
false

julia> iseven(10)
true
```

43.4 BigFloats

The `BigFloat` type implements arbitrary-precision floating-point arithmetic using the [GNU MPFR library](#).

precision ($num::AbstractFloat$)
Get the precision of a floating point number, as defined by the effective number of bits in the mantissa.

precision (*BigFloat*)

Get the precision (in bits) currently used for *BigFloat* arithmetic.

setprecision (*[T=BigFloat]*, *precision::Int*)

Set the precision (in bits) to be used for *T* arithmetic.

setprecision (*f::Function* [*T=BigFloat*], *precision::Integer*)

Change the *T* arithmetic precision (in bits) for the duration of *f*. It is logically equivalent to:

```
old = precision(BigFloat)
setprecision(BigFloat, precision)
f()
setprecision(BigFloat, old)
```

Often used as `setprecision(T, precision) do ... end`

43.5 Random Numbers

Random number generation in Julia uses the [Mersenne Twister library](#) via *MersenneTwister* objects. Julia has a global RNG, which is used by default. Other RNG types can be plugged in by inheriting the *AbstractRNG* type; they can then be used to have multiple streams of random numbers. Besides *MersenneTwister*, Julia also provides the *RandomDevice* RNG type, which is a wrapper over the OS provided entropy.

Most functions related to random generation accept an optional *AbstractRNG* as the first argument, *rng*, which defaults to the global one if not provided. Moreover, some of them accept optionally dimension specifications *dims...* (which can be given as a tuple) to generate arrays of random values.

A *MersenneTwister* or *RandomDevice* RNG can generate random numbers of the following types: *Float16*, *Float32*, *Float64*, *Bool*, *Int8*, *UInt8*, *Int16*, *UInt16*, *Int32*, *UInt32*, *Int64*, *UInt64*, *Int128*, *UInt128*, *BigInt* (or complex numbers of those types). Random floating point numbers are generated uniformly in $[0, 1)$. As *BigInt* represents unbounded integers, the interval must be specified (e.g. `rand(big(1:6))`).

srand (*[rng=GLOBAL_RNG]* [*, seed*])

Reseed the random number generator. If a *seed* is provided, the RNG will give a reproducible sequence of numbers, otherwise Julia will get entropy from the system. For *MersenneTwister*, the *seed* may be a non-negative integer, a vector of *UInt32* integers or a filename, in which case the seed is read from a file. *RandomDevice* does not support seeding.

MersenneTwister (*seed=0*)

Create a *MersenneTwister* RNG object. Different RNG objects can have their own seeds, which may be useful for generating different streams of random numbers.

RandomDevice ()

Create a *RandomDevice* RNG object. Two such objects will always generate different streams of random numbers.

rand (*[rng=GLOBAL_RNG]* [*, S*] [*, dims...*])

Pick a random element or array of random elements from the set of values specified by *S*; *S* can be

- an indexable collection (for example `1:n` or `['x', 'y', 'z']`), or
- a type: the set of values to pick from is then equivalent to `typemin(S) : typemax(S)` for integers (this is not applicable to *BigInt*), and to $[0, 1)$ for floating point numbers;

S defaults to *Float64*.

rand! (*[rng=GLOBAL_RNG]*, *A* [*, coll*])

Populate the array *A* with random values. If the indexable collection *coll* is specified, the values are picked

randomly from `coll`. This is equivalent to `copy!(A, rand(rng, coll, size(A)))` or `copy!(A, rand(rng, eltype(A), size(A)))` but without allocating a new array.

bitrand (`[rng=GLOBAL_RNG][, dims...]`)

Generate a `BitArray` of random boolean values.

randn (`[rng=GLOBAL_RNG][, T=Float64][, dims...]`)

Generate a normally-distributed random number of type `T` with mean 0 and standard deviation 1. Optionally generate an array of normally-distributed random numbers. The `Base` module currently provides an implementation for the types `Float16`, `Float32`, and `Float64` (the default).

randn! (`[rng=GLOBAL_RNG], A::AbstractArray`) \rightarrow `A`

Fill the array `A` with normally-distributed (mean 0, standard deviation 1) random numbers. Also see the `rand()` (page 440) function.

randexp (`[rng=GLOBAL_RNG][, T=Float64][, dims...]`)

Generate a random number of type `T` according to the exponential distribution with scale 1. Optionally generate an array of such random numbers. The `Base` module currently provides an implementation for the types `Float16`, `Float32`, and `Float64` (the default).

randexp! (`[rng=GLOBAL_RNG], A::AbstractArray`) \rightarrow `A`

Fill the array `A` with random numbers following the exponential distribution (with scale 1).

randjump (`(r::MersenneTwister, jumps::Integer[, jumppoly::AbstractString=dSFMT.JPOLY1e21])` \rightarrow `Vector{MersenneTwister}`)

Create an array of the size `jumps` of initialized `MersenneTwister` RNG objects. The first RNG object given as a parameter and following `MersenneTwister` RNGs in the array are initialized such that a state of the RNG object in the array would be moved forward (without generating numbers) from a previous RNG object array element on a particular number of steps encoded by the jump polynomial `jumppoly`.

Default jump polynomial moves forward `MersenneTwister` RNG state by 10^{20} steps.

Strings

length (*s::AbstractString*)

The number of characters in string *s*.

sizeof (*s::AbstractString*)

The number of bytes in string *s*.

```
julia> sizeof("")
3
```

***** (*s::AbstractString, t::AbstractString*)

Concatenate strings. The `*` operator is an alias to this function.

```
julia> "Hello " * "world"
"Hello world"
```

^ (*s::AbstractString, n::Integer*)

Repeat *n* times the string *s*. The `repeat()` (page 530) function is an alias to this operator.

```
julia> "Test " ^ 3
"Test Test Test "
```

string (*xs...*)

Create a string from any values using the `print()` (page 566) function.

repr (*x*)

Create a string from any value using the `showall()` (page 566) function.

String (*s::AbstractString*)

Convert a string to a contiguous byte array representation encoded as UTF-8 bytes. This representation is often appropriate for passing strings to C.

transcode (*T, src*)

Convert string data between Unicode encodings. *src* is either a `String` or a `Vector{UIntXX}` of UTF-XX code units, where XX is 8, 16, or 32. *T* indicates the encoding of the return value: `String` to return a (UTF-8 encoded) `String` or `UIntXX` to return a `Vector{UIntXX}` of UTF-XX data. (The alias `Cwchar_t` can also be used as the integer type, for converting `wchar_t*` strings used by external C libraries.)

The `transcode` function succeeds as long as the input data can be reasonably represented in the target encoding; it always succeeds for conversions between UTF-XX encodings, even for invalid Unicode data.

Only conversion to/from UTF-8 is currently supported.

unsafe_string (*p::Ptr{UInt8}[, length::Integer]*)

Copy a string from the address of a C-style (NUL-terminated) string encoded as UTF-8. (The pointer can be

safely freed afterwards.) If `length` is specified (the length of the data in bytes), the string does not have to be NUL-terminated.

This function is labelled “unsafe” because it will crash if `p` is not a valid memory address to data of the requested length.

See also `unsafe_wrap()` (page 444), which takes a pointer and wraps a string object around it without making a copy.

unsafe_wrap (*String*, *p::Ptr{UInt8}*[, *length*], *own=false*)

Wrap a pointer `p` to an array of bytes in a `String` object, interpreting the bytes as UTF-8 encoded characters *without making a copy*. The optional `length` argument indicates the length in bytes of the pointer’s data; if it is omitted, the data is assumed to be NUL-terminated. The `own` argument optionally specifies whether Julia should take ownership of the memory, calling `free` on the pointer when the array is no longer referenced.

This function is labelled “unsafe” because it will crash if `p` is not a valid memory address to data of the requested length.

See also `unsafe_string()` (page 443), which takes a pointer and makes a copy of the data.

ascii (*s::AbstractString*)

Convert a string to `String` type and check that it contains only ASCII data, otherwise throwing an `ArgumentError` indicating the position of the first non-ASCII byte.

@r_str () → `Regex`

Construct a regex, such as `r"^[a-z]*$"`. The regex also accepts one or more flags, listed after the ending quote, to change its behaviour:

- `i` enables case-insensitive matching
- `m` treats the `^` and `$` tokens as matching the start and end of individual lines, as opposed to the whole string.
- `s` allows the `.` modifier to match newlines.
- `x` enables “comment mode”: whitespace is enabled except when escaped with `\`, and `#` is treated as starting a comment.

For example, this regex has all three flags enabled:

```
julia> match(r"a+.*b+.*?d$"ism, "Goodbye,\nOh, angry,\nBad world\n")
RegexMatch("angry,\nBad world")
```

@html_str () → `Docs.HTML`

Create an `HTML` object from a literal string.

@text_str () → `Docs.Text`

Create a `Text` object from a literal string.

normalize_string (*s::AbstractString*, *normalform::Symbol*)

Normalize the string `s` according to one of the four “normal forms” of the Unicode standard: `normalform` can be `:NFC`, `:NFD`, `:NFKC`, or `:NFKD`. Normal forms C (canonical composition) and D (canonical decomposition) convert different visually identical representations of the same abstract string into a single canonical form, with form C being more compact. Normal forms KC and KD additionally canonicalize “compatibility equivalents”: they convert characters that are abstractly similar but visually distinct into a single canonical choice (e.g. they expand ligatures into the individual characters), with form KC being more compact.

Alternatively, finer control and additional transformations may be obtained by calling `normalize_string(s; keywords...)`, where any number of the following boolean keywords options (which all default to `false` except for `compose`) are specified:

- `compose=false`: do not perform canonical composition

- `decompose=true`: do canonical decomposition instead of canonical composition (`compose=true` is ignored if present)
- `compat=true`: compatibility equivalents are canonicalized
- `casefold=true`: perform Unicode case folding, e.g. for case-insensitive string comparison
- `newline2lf=true`, `newline2ls=true`, or `newline2ps=true`: convert various newline sequences (LF, CRLF, CR, NEL) into a linefeed (LF), line-separation (LS), or paragraph-separation (PS) character, respectively
- `stripmark=true`: strip diacritical marks (e.g. accents)
- `stripignore=true`: strip Unicode’s “default ignorable” characters (e.g. the soft hyphen or the left-to-right marker)
- `stripcc=true`: strip control characters; horizontal tabs and form feeds are converted to spaces; newlines are also converted to spaces unless a newline-conversion flag was specified
- `rejectna=true`: throw an error if unassigned code points are found
- `stable=true`: enforce Unicode Versioning Stability

For example, NFKC corresponds to the options `compose=true`, `compat=true`, `stable=true`.

graphemes (*s::AbstractString*) → GraphemeIterator

Returns an iterator over substrings of *s* that correspond to the extended graphemes in the string, as defined by Unicode UAX #29. (Roughly, these are what users would perceive as single characters, even though they may contain more than one codepoint; for example a letter combined with an accent mark is a single grapheme.)

isvalid (*value*) → Bool

Returns `true` if the given value is valid for its type, which currently can be either `Char` or `String`.

isvalid (*T*, *value*) → Bool

Returns `true` if the given value is valid for that type. Types currently can be either `Char` or `String`. Values for `Char` can be of type `Char` or `UInt32`. Values for `String` can be of that type, or `Vector{UInt8}`.

isvalid (*str::AbstractString*, *i::Integer*)

Tells whether index *i* is valid for the given string.

is_assigned_char (*c*) → Bool

Returns `true` if the given char or integer is an assigned Unicode code point.

ismatch (*r::Regex*, *s::AbstractString*) → Bool

Test whether a string contains a match of the given regular expression.

match (*r::Regex*, *s::AbstractString*[, *idx::Integer*[, *addopts*]])

Search for the first match of the regular expression *r* in *s* and return a `RegexMatch` object containing the match, or nothing if the match failed. The matching substring can be retrieved by accessing `m.match` and the captured sequences can be retrieved by accessing `m.captures`. The optional *idx* argument specifies an index at which to start the search.

eachmatch (*r::Regex*, *s::AbstractString*[, *overlap::Bool=false*])

Search for all matches of a the regular expression *r* in *s* and return an iterator over the matches. If *overlap* is `true`, the matching sequences are allowed to overlap indices in the original string, otherwise they must be from distinct character ranges.

matchall (*r::Regex*, *s::AbstractString*[, *overlap::Bool=false*]) → Vector{AbstractString}

Return a vector of the matching substrings from `eachmatch()` (page 445).

lpad (*s*, *n::Integer*, *p::AbstractString*=" ")

Make a string at least *n* columns wide when printed by padding *s* on the left with copies of *p*.

```
julia> lpad("March",10)
"      March"
```

rpadd (*s*, *n*::Integer, *p*::AbstractString=" ")

Make a string at least *n* columns wide when printed by padding *s* on the right with copies of *p*.

```
julia> rpadd("March",20)
"March"
```

search (*string*::AbstractString, *chars*::Chars[, *start*::Integer])

Search for the first occurrence of the given characters within the given string. The second argument may be a single character, a vector or a set of characters, a string, or a regular expression (though regular expressions are only allowed on contiguous strings, such as ASCII or UTF-8 strings). The third argument optionally specifies a starting index. The return value is a range of indexes where the matching sequence is found, such that `s[search(s,x)] == x`:

`search(string, "substring") = start:end` such that `string[start:end] == "substring"`, or `0:-1` if unmatched.

`search(string, 'c') = index` such that `string[index] == 'c'`, or 0 if unmatched.

```
julia> search("Hello to the world", "z")
0:-1

julia> search("JuliaLang", "Julia")
1:5
```

rsearch (*s*::AbstractString, *chars*::Chars[, *start*::Integer])

Similar to `search()` (page 446), but returning the last occurrence of the given characters within the given string, searching in reverse from *start*.

```
julia> rsearch("aaabbb", "b")
6:6
```

searchindex (*s*::AbstractString, *substring*[, *start*::Integer])

Similar to `search()` (page 446), but return only the start index at which the substring is found, or 0 if it is not.

rsearchindex (*s*::AbstractString, *substring*[, *start*::Integer])

Similar to `rsearch()` (page 446), but return only the start index at which the substring is found, or 0 if it is not.

contains (*haystack*::AbstractString, *needle*::AbstractString)

Determine whether the second argument is a substring of the first.

```
julia> contains("JuliaLang is pretty cool!", "Julia")
true
```

reverse (*s*::AbstractString) → AbstractString

Reverses a string.

```
julia> reverse("JuliaLang")
"gnailuJ"
```

replace (*string*::AbstractString, *pat*, *r*[, *n*::Integer=0])

Search for the given pattern *pat*, and replace each occurrence with *r*. If *n* is provided, replace at most *n* occurrences. As with `search`, the second argument may be a single character, a vector or a set of characters, a string, or a regular expression. If *r* is a function, each occurrence is replaced with `r(s)` where *s* is the matched substring. If *pat* is a regular expression and *r* is a `SubstitutionString`, then capture group references in *r* are replaced with the corresponding matched text.

split (*s::AbstractString*, [*chars*]; *limit::Integer*=0, *keep::Bool*=true)

Return an array of substrings by splitting the given string on occurrences of the given character delimiters, which may be specified in any of the formats allowed by `search`'s second argument (i.e. a single character, collection of characters, string, or regular expression). If *chars* is omitted, it defaults to the set of all space characters, and *keep* is taken to be `false`. The two keyword arguments are optional: they are a maximum size for the result and a flag determining whether empty fields should be kept in the result.

```
julia> a = "Ma.rch"
"Ma.rch"

julia> split(a, ".")
2-element Array{SubString{String},1}:
 "Ma"
 "rch"
```

rsplit (*s::AbstractString*, [*chars*]; *limit::Integer*=0, *keep::Bool*=true)

Similar to `split()` (page 446), but starting from the end of the string.

```
julia> a = "M.a.r.c.h"
"M.a.r.c.h"

julia> rsplit(a, ".")
5-element Array{SubString{String},1}:
 "M"
 "a"
 "r"
 "c"
 "h"

julia> rsplit(a, "."; limit=1)
1-element Array{SubString{String},1}:
 "M.a.r.c.h"

julia> rsplit(a, "."; limit=2)
2-element Array{SubString{String},1}:
 "M.a.r.c"
 "h"
```

strip (*s::AbstractString* [, *chars::Chars*])

Return *s* with any leading and trailing whitespace removed. If *chars* (a character, or vector or set of characters) is provided, instead remove characters contained in it.

lstrip (*s::AbstractString* [, *chars::Chars*])

Return *s* with any leading whitespace and delimiters removed. The default delimiters to remove are ' ', \t, \n, \v, \f, and \r. If *chars* (a character, or vector or set of characters) is provided, instead remove characters contained in it.

rstrip (*s::AbstractString* [, *chars::Chars*])

Return *s* with any trailing whitespace and delimiters removed. The default delimiters to remove are ' ', \t, \n, \v, \f, and \r. If *chars* (a character, or vector or set of characters) is provided, instead remove characters contained in it.

```
julia> a = rpad("March", 20)
"March          "

julia> rstrip(a)
"March"
```

startswith (*s::AbstractString*, *prefix::AbstractString*)

Returns `true` if `s` starts with `prefix`. If `prefix` is a vector or set of characters, tests whether the first character of `s` belongs to that set.

```
julia> startswith("JuliaLang", "Julia")
true
```

endswith (*s::AbstractString*, *suffix::AbstractString*)

Returns `true` if `s` ends with `suffix`. If `suffix` is a vector or set of characters, tests whether the last character of `s` belongs to that set.

```
julia> endswith("Sunday", "day")
true
```

uppercase (*s::AbstractString*)

Returns `s` with all characters converted to uppercase.

```
julia> uppercase("Julia")
"JULIA"
```

lowercase (*s::AbstractString*)

Returns `s` with all characters converted to lowercase.

```
julia> lowercase("STRINGS AND THINGS")
"strings and things"
```

titlecase (*s::AbstractString*)

Capitalizes the first character of each word in `s`.

```
julia> titlecase("the julia programming language")
"The Julia Programming Language"
```

ucfirst (*s::AbstractString*)

Returns string with the first character converted to uppercase.

```
julia> ucfirst("python")
"Python"
```

lcfirst (*s::AbstractString*)

Returns string with the first character converted to lowercase.

```
julia> lcfirst("Julia")
"julia"
```

join (*io::IO*, *strings*, *delim*_[, last])

Join an array of strings into a single string, inserting the given delimiter between adjacent strings. If `last` is given, it will be used instead of `delim` between the last two strings. For example,

```
julia> join(["apples", "bananas", "pineapples"], ", ", " and ")
"apples, bananas and pineapples"
```

strings can be any iterable over elements `x` which are convertible to strings via `print(io::IOBuffer, x)`. strings will be printed to `io`.

chop (*s::AbstractString*)

Remove the last character from `s`.

```
julia> a = "March"
"March"

julia> chop(a)
"Marc"
```

chomp (*s::AbstractString*)

Remove a single trailing newline from a string.

ind2chr (*s::AbstractString, i::Integer*)

Convert a byte index *i* to a character index with respect to string *s*.

chr2ind (*s::AbstractString, i::Integer*)

Convert a character index *i* to a byte index.

nextind (*str::AbstractString, i::Integer*)

Get the next valid string index after *i*. Returns a value greater than `endof(str)` at or after the end of the string.

prevind (*str::AbstractString, i::Integer*)

Get the previous valid string index before *i*. Returns a value less than 1 at the beginning of the string.

randstring (*[rng]*, *len=8*)

Create a random ASCII string of length *len*, consisting of upper- and lower-case letters and the digits 0-9. The optional *rng* argument specifies a random number generator, see [Random Numbers](#) (page 440).

charwidth (*c*)

Gives the number of columns needed to print a character.

strwidth (*s::AbstractString*)

Gives the number of columns needed to print a string.

```
julia> strwidth("March")
5
```

isalnum (*c::Union{Char, AbstractString}*) → Bool

Tests whether a character is alphanumeric, or whether this is true for all elements of a string. A character is classified as alphabetic if it belongs to the Unicode general category Letter or Number, i.e. a character whose category code begins with ‘L’ or ‘N’.

isalpha (*c::Union{Char, AbstractString}*) → Bool

Tests whether a character is alphabetic, or whether this is true for all elements of a string. A character is classified as alphabetic if it belongs to the Unicode general category Letter, i.e. a character whose category code begins with ‘L’.

isascii (*c::Union{Char, AbstractString}*) → Bool

Tests whether a character belongs to the ASCII character set, or whether this is true for all elements of a string.

iscntrl (*c::Union{Char, AbstractString}*) → Bool

Tests whether a character is a control character, or whether this is true for all elements of a string. Control characters are the non-printing characters of the Latin-1 subset of Unicode.

isdigit (*c::Union{Char, AbstractString}*) → Bool

Tests whether a character is a numeric digit (0-9), or whether this is true for all elements of a string.

isgraph (*c::Union{Char, AbstractString}*) → Bool

Tests whether a character is printable, and not a space, or whether this is true for all elements of a string. Any character that would cause a printer to use ink should be classified with `isgraph(c) == true`.

islower (*c::Union{Char, AbstractString}*) → Bool

Tests whether a character is a lowercase letter, or whether this is true for all elements of a string. A character is classified as lowercase if it belongs to Unicode category Ll, Letter: Lowercase.

isnumber (*c::Union{Char, AbstractString}*) → Bool

Tests whether a character is numeric, or whether this is true for all elements of a string. A character is classified as numeric if it belongs to the Unicode general category Number, i.e. a character whose category code begins with ‘N’.

isprint (*c::Union{Char, AbstractString}*) → Bool

Tests whether a character is printable, including spaces, but not a control character. For strings, tests whether this is true for all elements of the string.

ispunct (*c::Union{Char, AbstractString}*) → Bool

Tests whether a character belongs to the Unicode general category Punctuation, i.e. a character whose category code begins with ‘P’. For strings, tests whether this is true for all elements of the string.

isspace (*c::Union{Char, AbstractString}*) → Bool

Tests whether a character is any whitespace character. Includes ASCII characters ‘\t’, ‘\n’, ‘\v’, ‘\f’, ‘\r’, and ‘ ’, Latin-1 character U+0085, and characters in Unicode category Zs. For strings, tests whether this is true for all elements of the string.

isupper (*c::Union{Char, AbstractString}*) → Bool

Tests whether a character is an uppercase letter, or whether this is true for all elements of a string. A character is classified as uppercase if it belongs to Unicode category Lu, Letter: Uppercase, or Lt, Letter: Titlecase.

isxdigit (*c::Union{Char, AbstractString}*) → Bool

Tests whether a character is a valid hexadecimal digit, or whether this is true for all elements of a string.

```
julia> isxdigit("abc")
true

julia> isxdigit("0x9")
false
```

Symbol (*x...*) → Symbol

Create a Symbol by concatenating the string representations of the arguments together.

escape_string (*[io]*, *str::AbstractString*, *esc::AbstractString*) → AbstractString

General escaping of traditional C and Unicode escape sequences. Any characters in *esc* are also escaped (with a backslash). See also *unescape_string()* (page 450).

unescape_string (*[io]*, *s::AbstractString*) → AbstractString

General unescaping of traditional C and Unicode escape sequences. Reverse of *escape_string()* (page 450).

Arrays

45.1 Basic functions

ndims (*A::AbstractArray*) → Integer

Returns the number of dimensions of *A*.

```
julia> A = ones(3,4,5);  
  
julia> ndims(A)  
3
```

size (*A::AbstractArray*[, *dim...*])

Returns a tuple containing the dimensions of *A*. Optionally you can specify the dimension(s) you want the length of, and get the length of that dimension, or a tuple of the lengths of dimensions you asked for.

```
julia> A = ones(2,3,4);  
  
julia> size(A, 2)  
3  
  
julia> size(A,3,2)  
(4,3)
```

indices (*A*)

Returns the tuple of valid indices for array *A*.

```
julia> A = ones(5,6,7);  
  
julia> indices(A)  
(Base.OneTo(5), Base.OneTo(6), Base.OneTo(7))
```

indices (*A*, *d*)

Returns the valid range of indices for array *A* along dimension *d*.

```
julia> A = ones(5,6,7);  
  
julia> indices(A,2)  
Base.OneTo(6)
```

length (*A::AbstractArray*) → Integer

Returns the number of elements in *A*.

```
julia> A = ones(3,4,5);

julia> length(A)
60
```

eachindex (A...)

Creates an iterable object for visiting each index of an AbstractArray A in an efficient manner. For array types that have opted into fast linear indexing (like Array), this is simply the range `1:length(A)`. For other array types, this returns a specialized Cartesian range to efficiently index into the array with indices specified for every dimension. For other iterables, including strings and dictionaries, this returns an iterator object supporting arbitrary index types (e.g. unevenly spaced or non-integer indices).

Example for a sparse 2-d array:

```
julia> A = sparse([1, 1, 2], [1, 3, 1], [1, 2, -5])
2×3 sparse matrix with 3 Int64 nonzero entries:
 [1, 1] = 1
 [2, 1] = -5
 [1, 3] = 2

julia> for iter in eachindex(A)
        @show iter.I[1], iter.I[2]
        @show A[iter]

    end
(iter.I[1],iter.I[2]) = (1,1)
A[iter] = 1
(iter.I[1],iter.I[2]) = (2,1)
A[iter] = -5
(iter.I[1],iter.I[2]) = (1,2)
A[iter] = 0
(iter.I[1],iter.I[2]) = (2,2)
A[iter] = 0
(iter.I[1],iter.I[2]) = (1,3)
A[iter] = 2
(iter.I[1],iter.I[2]) = (2,3)
A[iter] = 0
```

If you supply more than one AbstractArray argument, `eachindex` will create an iterable object that is fast for all arguments (a `UnitRange` if all inputs have fast linear indexing, a `CartesianRange` otherwise). If the arrays have different sizes and/or dimensionalities, `eachindex` returns an iterable that spans the largest range along each dimension.

linearindices (A)

Returns a `UnitRange` specifying the valid range of indices for `A[i]` where `i` is an `Int`. For arrays with conventional indexing (indices start at 1), or any multidimensional array, this is `1:length(A)`; however, for one-dimensional arrays with unconventional indices, this is `indices(A, 1)`.

Calling this function is the “safe” way to write algorithms that exploit linear indexing.

```
julia> A = ones(5,6,7);

julia> b = linearindices(A);

julia> extrema(b)
(1,210)
```

Base.linearindexing (A)

`linearindexing` defines how an AbstractArray most efficiently accesses its elements. If `Base.linearindexing(A)` returns `Base.LinearFast()`, this means that linear indexing with only

one index is an efficient operation. If it instead returns `Base.LinearSlow()` (by default), this means that the array intrinsically accesses its elements with indices specified for every dimension. Since converting a linear index to multiple indexing subscripts is typically very expensive, this provides a traits-based mechanism to enable efficient generic code for all array types.

An abstract array subtype `MyArray` that wishes to opt into fast linear indexing behaviors should define `linearindexing` in the type-domain:

```
Base.linearindexing{T<:MyArray}{::Type{T}} = Base.LinearFast()
```

countnz (A)

Counts the number of nonzero values in array A (dense or sparse). Note that this is not a constant-time operation. For sparse matrices, one should usually use `nnz()` (page 480), which returns the number of stored values.

```
julia> A = [1 2 4; 0 0 1; 1 1 0]
3×3 Array{Int64,2}:
 1  2  4
 0  0  1
 1  1  0

julia> countnz(A)
6
```

conj! (A)

Transform an array to its complex conjugate in-place.

See also `conj()` (page 420).

stride (A, k::Integer)

Returns the distance in memory (in number of elements) between adjacent elements in dimension k.

```
julia> A = ones(3,4,5);

julia> stride(A,2)
3

julia> stride(A,3)
12
```

strides (A)

Returns a tuple of the memory strides in each dimension.

```
julia> A = ones(3,4,5);

julia> strides(A)
(1, 3, 12)
```

ind2sub (dims, index) → subscripts

Returns a tuple of subscripts into an array with dimensions `dims`, corresponding to the linear index `index`.

Example:

```
i, j, ... = ind2sub(size(A), indmax(A))
```

provides the indices of the maximum element.

```
julia> ind2sub((3,4), 2)
(2, 1)

julia> ind2sub((3,4), 3)
(3, 1)
```

```
julia> ind2sub((3,4),4)
(1,2)
```

ind2sub (*a, index*) → subscripts

Returns a tuple of subscripts into array *a* corresponding to the linear index *index*.

```
julia> A = ones(5,6,7);

julia> ind2sub(A,35)
(5,1,2)

julia> ind2sub(A,70)
(5,2,3)
```

sub2ind (*dims, i, j, k...*) → index

The inverse of *ind2sub()* (page 453), returns the linear index corresponding to the provided subscripts.

```
julia> sub2ind((5,6,7),1,2,3)
66

julia> sub2ind((5,6,7),1,6,3)
86
```

LinAlg.checksquare (*A*)

Check that a matrix is square, then return its common dimension. For multiple arguments, return a vector.

Example

```
julia> A = ones(4,4); B = zeros(5,5);

julia> LinAlg.checksquare(A, B)
2-element Array{Int64,1}:
 4
 5
```

45.2 Constructors

Array (*dims*)

`Array{T}(dims)` constructs an uninitialized dense array with element type *T*. *dims* may be a tuple or a series of integer arguments. The syntax `Array{T, dims}` is also available, but deprecated.

getindex (*type*[, *elements...*])

Construct a 1-d array of the specified type. This is usually called with the syntax `Type[]`. Element values can be specified using `Type[a,b,c,...]`.

zeros (*type, dims*)

Create an array of all zeros of specified type. The type defaults to `Float64` if not specified.

```
julia> zeros{Int8, 2, 3}
2×3 Array{Int8,2}:
 0  0  0
 0  0  0
```

zeros (*A*)

Create an array of all zeros with the same element type and shape as *A*.

```
julia> A = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> zeros(A)
2×2 Array{Int64,2}:
 0  0
 0  0
```

ones (*type, dims*)

Create an array of all ones of specified type. The type defaults to `Float64` if not specified.

```
julia> ones(Complex128, 2, 3)
2×3 Array{Complex{Float64},2}:
 1.0+0.0im  1.0+0.0im  1.0+0.0im
 1.0+0.0im  1.0+0.0im  1.0+0.0im
```

ones (*A*)

Create an array of all ones with the same element type and shape as *A*.

```
julia> A = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> ones(A)
2×2 Array{Int64,2}:
 1  1
 1  1
```

BitArray (*dims::Integer...*)**BitArray{N}** (*dims::NTuple{N, Int}*)

Construct an uninitialized `BitArray` with the given dimensions. Behaves identically to the `Array{}` (page 454) constructor.

BitArray (*itr*)

Construct a `BitArray` generated by the given iterable object. The shape is inferred from the `itr` object.

```
julia> BitArray([1 0; 0 1])
2×2 BitArray{2}:
 true  false
 false  true

julia> BitArray(x+y == 3 for x = 1:2, y = 1:3)
2×3 BitArray{2}:
 false  true  false
 true  false  false

julia> BitArray(x+y == 3 for x = 1:2 for y = 1:3)
6-element BitArray{1}:
 false
  true
 false
  true
 false
 false
```

trues (*dims*)

Create a `BitArray` with all values set to `true`.

```
julia> trues(2,3)
2×3 BitArray{2}:
 true  true  true
 true  true  true
```

trues (*A*)

Create a BitArray with all values set to `true` of the same shape as *A*.

```
julia> A = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> trues(A)
2×2 BitArray{2}:
 true  true
 true  true
```

falses (*dims*)

Create a BitArray with all values set to `false`.

```
julia> falses(2,3)
2×3 BitArray{2}:
 false false false
 false false false
```

falses (*A*)

Create a BitArray with all values set to `false` of the same shape as *A*.

```
julia> A = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> falses(A)
2×2 BitArray{2}:
 false false
 false false
```

fill (*x*, *dims*)

Create an array filled with the value *x*. For example, `fill(1.0, (5,5))` returns a 5×5 array of floats, with each element initialized to 1.0.

```
julia> fill(1.0, (5,5))
5×5 Array{Float64,2}:
 1.0  1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0  1.0
```

If *x* is an object reference, all elements will refer to the same object. `fill(Foo(), dims)` will return an array filled with the result of evaluating `Foo()` once.

fill! (*A*, *x*)

Fill array *A* with the value *x*. If *x* is an object reference, all elements will refer to the same object. `fill!(A, Foo())` will return *A* filled with the result of evaluating `Foo()` once.

reshape (*A*, *dims*)

Create an array with the same data as the given array, but with different dimensions.

```
julia> A = collect(1:16)
16-element Array{Int64,1}:
 1
 2
 3
 4
 5
 6
 7
 8
 9
10
11
12
13
14
15
16

julia> reshape(A, (2, 8))
2×8 Array{Int64,2}:
 1  3  5  7  9 11 13 15
 2  4  6  8 10 12 14 16
```

similar(array[, element_type=eltype(array)][, dims=size(array)])

Create an uninitialized mutable array with the given element type and size, based upon the given source array. The second and third arguments are both optional, defaulting to the given array's `eltype` and `size`. The dimensions may be specified either as a single tuple argument or as a series of integer arguments.

Custom `AbstractArray` subtypes may choose which specific array type is best-suited to return for the given element type and dimensionality. If they do not specialize this method, the default is an `Array{element_type}(dims...)`.

For example, `similar(1:10, 1, 4)` returns an uninitialized `Array{Int,2}` since ranges are neither mutable nor support 2 dimensions:

```
julia> similar(1:10, 1, 4)
1×4 Array{Int64,2}:
4419743872 4374413872 4419743888 0
```

Conversely, `similar(trues(10,10), 2)` returns an uninitialized `BitVector` with two elements since `BitArrays` are both mutable and can support 1-dimensional arrays:

```
julia> similar(trues(10,10), 2)
2-element BitArray{1}:
 false
 false
```

Since `BitArrays` can only store elements of type `Bool`, however, if you request a different element type it will create a regular `Array` instead:

```
julia> similar(falses(10), Float64, 2, 4)
2×4 Array{Float64,2}:
 2.18425e-314  2.18425e-314  2.18425e-314  2.18425e-314
 2.18425e-314  2.18425e-314  2.18425e-314  2.18425e-314
```

similar(storage_type, indices)

Create an uninitialized mutable array analogous to that specified by `storage_type`, but with indices specified by the last argument. `storage_type` might be a type or a function.

Examples:

```
similar(Array{Int}, indices(A))
```

creates an array that “acts like” an `Array{Int}` (and might indeed be backed by one), but which is indexed identically to `A`. If `A` has conventional indexing, this will be identical to `Array{Int}(size(A))`, but if `A` has unconventional indexing then the indices of the result will match `A`.

```
similar(BitArray, (indices(A, 2),))
```

would create a 1-dimensional logical array whose indices match those of the columns of `A`.

```
similar(dims->zeros{Int, dims}, indices(A))
```

would create an array of `Int`, initialized to zero, matching the indices of `A`.

reinterpret (*type*, *A*)

Change the type-interpretation of a block of memory. For arrays, this constructs an array with the same binary data as the given array, but with the specified element type. For example, `reinterpret{Float32, UInt32(7)}` interprets the 4 bytes corresponding to `UInt32(7)` as a `Float32`.

```
julia> reinterpret{Float32, UInt32(7)}
1.0f-44

julia> reinterpret{Float32, UInt32[1 2 3 4 5]}
1×5 Array{Float32,2}:
 1.4013f-45  2.8026f-45  4.2039f-45  5.60519f-45  7.00649f-45
```

eye (*T::Type=Float64*, *n::Integer*)

n-by-*n* identity matrix. The default element type is `Float64`.

eye (*T::Type=Float64*, *m::Integer*, *n::Integer*)

m-by-*n* identity matrix. The default element type is `Float64`.

eye (*A*)

Constructs an identity matrix of the same dimensions and type as `A`.

```
julia> A = [1 2 3; 4 5 6; 7 8 9]
3×3 Array{Int64,2}:
 1  2  3
 4  5  6
 7  8  9

julia> eye(A)
3×3 Array{Int64,2}:
 1  0  0
 0  1  0
 0  0  1
```

Note the difference from `ones()` (page 455).

linspace (*start::Real*, *stop::Real*, *n::Real=50*)

Construct a range of *n* linearly spaced elements from *start* to *stop*.

```
julia> linspace(1.3, 2.9, 9)
9-element LinSpace{Float64}:
 1.3, 1.5, 1.7, 1.9, 2.1, 2.3, 2.5, 2.7, 2.9
```

logspace (*start::Real*, *stop::Real*, *n::Integer=50*)

Construct a vector of *n* logarithmically spaced numbers from 10^{start} to 10^{stop} .


```
julia> logspace(1.,10.,5)
5-element Array{Float64,1}:
 10.0
1778.28
 3.16228e5
 5.62341e7
 1.0e10
```

45.3 Mathematical operators and functions

All mathematical operations and functions are supported for arrays

broadcast (*f*, *As...*)

Broadcasts the arrays, tuples and/or scalars *As* to a container of the appropriate type and dimensions. In this context, anything that is not a subtype of `AbstractArray` or `Tuple` is considered a scalar. The resulting container is established by the following rules:

- If all the arguments are scalars, it returns a scalar.
- If the arguments are tuples and zero or more scalars, it returns a tuple.
- If there is at least an array in the arguments, it returns an array (and treats tuples as 1-dimensional arrays) expanding singleton dimensions.

A special syntax exists for broadcasting: `f. (args...)` is equivalent to `broadcast(f, args...)`, and nested `f. (g. (args...))` calls are fused into a single broadcast loop.

```
julia> A = [1, 2, 3, 4, 5]
5-element Array{Int64,1}:
 1
 2
 3
 4
 5

julia> B = [1 2; 3 4; 5 6; 7 8; 9 10]
5×2 Array{Int64,2}:
 1  2
 3  4
 5  6
 7  8
 9 10

julia> broadcast(+, A, B)
5×2 Array{Int64,2}:
 2  3
 5  6
 8  9
11 12
14 15

julia> parse.(Int, ["1", "2"])
2-element Array{Int64,1}:
 1
 2

julia> abs.((1, -2))
```

```
(1,2)

julia> broadcast(+, 1.0, (0, -2.0))
(1.0,-1.0)

julia> broadcast(+, 1.0, (0, -2.0), [1])
2-element Array{Float64,1}:
 2.0
 0.0

julia> string.(("one","two","three","four"), ": ", 1:4)
4-element Array{String,1}:
 "one: 1"
 "two: 2"
 "three: 3"
 "four: 4"
```

broadcast! (*f, dest, As...*)

Like *broadcast()* (page 459), but store the result of *broadcast(f, As...)* in the *dest* array. Note that *dest* is only used to store the result, and does not supply arguments to *f* unless it is also listed in the *As*, as in *broadcast!(f, A, A, B)* to perform *A[:] = broadcast(f, A, B)*.

bitbroadcast (*f, As...*)

Like *broadcast()* (page 459), but allocates a *BitArray* to store the result, rather than an *Array*.

```
julia> bitbroadcast(isodd, [1,2,3,4,5])
5-element BitArray{1}:
 true
 false
 true
 false
 true
```

45.4 Indexing, Assignment, and Concatenation

getindex (*A, inds...*)

Returns a subset of array *A* as specified by *inds*, where each *ind* may be an *Int*, a *Range*, or a *Vector*. See the manual section on *array indexing* (page 180) for details.

```
julia> A = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> getindex(A, 1)
1

julia> getindex(A, [2, 1])
2-element Array{Int64,1}:
 3
 1

julia> getindex(A, 2:4)
3-element Array{Int64,1}:
 3
 2
```

4

view (*A*, *inds...*)

Like `getindex()` (page 454), but returns a view into the parent array *A* with the given indices instead of making a copy. Calling `getindex()` (page 454) or `setindex!()` (page 461) on the returned `SubArray` computes the indices to the parent array on the fly without checking bounds.

@view *A*[*inds...*]

Creates a `SubArray` from an indexing expression. This can only be applied directly to a reference expression (e.g. `@view A[1,2:end]`), and should *not* be used as the target of an assignment (e.g. `@view(A[1,2:end]) = ...`).

parent (*A*)

Returns the “parent array” of an array view type (e.g., `SubArray`), or the array itself if it is not a view.

parentindexes (*A*)

From an array view *A*, returns the corresponding indexes in the parent.

slicedim (*A*, *d::Integer*, *i*)

Return all the data of *A* where the index for dimension *d* equals *i*. Equivalent to `A[:, :, ..., i, :, :, ...]` where *i* is in position *d*.

```
julia> A = [1 2 3 4; 5 6 7 8]
2×4 Array{Int64,2}:
 1  2  3  4
 5  6  7  8

julia> slicedim(A,2,3)
2-element Array{Int64,1}:
 3
 7
```

setindex! (*A*, *X*, *inds...*)

Store values from array *X* within some subset of *A* as specified by *inds*.

broadcast_getindex (*A*, *inds...*)

Broadcasts the *inds* arrays to a common size like `broadcast()` (page 459) and returns an array of the results `A[ks...]`, where *ks* goes over the positions in the broadcast result *A*.

```
julia> A = [1, 2, 3, 4, 5]
5-element Array{Int64,1}:
 1
 2
 3
 4
 5

julia> B = [1 2; 3 4; 5 6; 7 8; 9 10]
5×2 Array{Int64,2}:
 1  2
 3  4
 5  6
 7  8
 9 10

julia> C = broadcast(+,A,B)
5×2 Array{Int64,2}:
 2  3
 5  6
 8  9
```

```

11 12
14 15

julia> broadcast_getindex(C, [1,2,10])
3-element Array{Int64,1}:
 2
 5
15

```

broadcast_setindex! (*A*, *X*, *inds*...)

Broadcasts the *X* and *inds* arrays to a common size and stores the value from each position in *X* at the indices in *A* given by the same positions in *inds*.

isassigned (*array*, *i*) → Bool

Tests whether the given array has a value associated with index *i*. Returns *false* if the index is out of bounds, or has an undefined reference.

cat (*dims*, *A*...)

Concatenate the input arrays along the specified dimensions in the iterable *dims*. For dimensions not in *dims*, all input arrays should have the same size, which will also be the size of the output array along that dimension. For dimensions in *dims*, the size of the output array is the sum of the sizes of the input arrays along that dimension. If *dims* is a single number, the different arrays are tightly stacked along that dimension. If *dims* is an iterable containing several dimensions, this allows one to construct block diagonal matrices and their higher-dimensional analogues by simultaneously increasing several dimensions for every new input array and putting zero blocks elsewhere. For example, `cat([1,2], matrices...)` builds a block diagonal matrix, i.e. a block matrix with `matrices[1]`, `matrices[2]`, ... as diagonal blocks and matching zero blocks away from the diagonal.

vcat (*A*...)

Concatenate along dimension 1.

```

julia> a = [1 2 3 4 5]
1×5 Array{Int64,2}:
 1  2  3  4  5

julia> b = [6 7 8 9 10; 11 12 13 14 15]
2×5 Array{Int64,2}:
 6  7  8  9 10
11 12 13 14 15

julia> vcat(a,b)
3×5 Array{Int64,2}:
 1  2  3  4  5
 6  7  8  9 10
11 12 13 14 15

julia> c = ([1 2 3], [4 5 6])
(
 [1 2 3],

 [4 5 6])

julia> vcat(c...)
2×3 Array{Int64,2}:
 1  2  3
 4  5  6

```

hcat (*A*...)

Concatenate along dimension 2.

```

julia> a = [1; 2; 3; 4; 5]
5-element Array{Int64,1}:
 1
 2
 3
 4
 5

julia> b = [6 7; 8 9; 10 11; 12 13; 14 15]
5×2 Array{Int64,2}:
 6  7
 8  9
10 11
12 13
14 15

julia> hcat(a,b)
5×3 Array{Int64,2}:
 1  6  7
 2  8  9
 3 10 11
 4 12 13
 5 14 15

julia> c = ([1; 2; 3], [4; 5; 6])
([1,2,3],[4,5,6])

julia> hcat(c...)
3×2 Array{Int64,2}:
 1  4
 2  5
 3  6

```

hvcat (*rows::Tuple{Vararg{Int}}*, *values...*)

Horizontal and vertical concatenation in one call. This function is called for block matrix syntax. The first argument specifies the number of arguments to concatenate in each block row.

```

julia> a, b, c, d, e, f = 1, 2, 3, 4, 5, 6
(1,2,3,4,5,6)

julia> [a b c; d e f]
2×3 Array{Int64,2}:
 1  2  3
 4  5  6

julia> hvcat((3,3), a,b,c,d,e,f)
2×3 Array{Int64,2}:
 1  2  3
 4  5  6

julia> [a b;c d; e f]
3×2 Array{Int64,2}:
 1  2
 3  4
 5  6

julia> hvcat((2,2,2), a,b,c,d,e,f)
3×2 Array{Int64,2}:

```

```
1  2
3  4
5  6
```

If the first argument is a single integer `n`, then all block rows are assumed to have `n` block columns.

flipdim(*A*, *d::Integer*)

Reverse *A* in dimension *d*.

```
julia> b = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> flipdim(b,2)
2×2 Array{Int64,2}:
 2  1
 4  3
```

circshift(*A*, *shifts*)

Circularly shift the data in an array. The second argument is a vector giving the amount to shift in each dimension.

```
julia> b = reshape(collect(1:16), (4,4))
4×4 Array{Int64,2}:
 1  5  9 13
 2  6 10 14
 3  7 11 15
 4  8 12 16

julia> circshift(b, (0,2))
4×4 Array{Int64,2}:
 9 13  1  5
10 14  2  6
11 15  3  7
12 16  4  8

julia> circshift(b, (-1,0))
4×4 Array{Int64,2}:
 2  6 10 14
 3  7 11 15
 4  8 12 16
 1  5  9 13
```

See also [`circshift!\(\)`](#) (page 464).

circshift!(*dest*, *src*, *shifts*)

Circularly shift the data in *src*, storing the result in *dest*. *shifts* specifies the amount to shift in each dimension.

The *dest* array must be distinct from the *src* array (they cannot alias each other).

See also `circshift`.

circcopy!(*dest*, *src*)

Copy *src* to *dest*, indexing each dimension modulo its length. *src* and *dest* must have the same size, but can be offset in their indices; any offset results in a (circular) wraparound. If the arrays have overlapping indices, then on the domain of the overlap *dest* agrees with *src*.

```

julia> src = reshape(collect(1:16), (4,4))
4×4 Array{Int64,2}:
 1  5  9 13
 2  6 10 14
 3  7 11 15
 4  8 12 16

julia> dest = OffsetArray{Int}((0:3,2:5))

julia> circcopy!(dest, src)
OffsetArrays.OffsetArray{Int64,2,Array{Int64,2}} with indices 0:3×2:5:
 8 12 16  4
 5  9 13  1
 6 10 14  2
 7 11 15  3

julia> dest[1:3,2:4] == src[1:3,2:4]
true

```

find(A)

Return a vector of the linear indexes of the non-zeros in *A* (determined by *A*[*i*] != 0). A common use of this is to convert a boolean array to an array of indexes of the `true` elements. If there are no non-zero elements of *A*, `find` returns an empty array.

```

julia> A = [true false; false true]
2×2 Array{Bool,2}:
 true  false
 false  true

julia> find(A)
2-element Array{Int64,1}:
 1
 4

```

find(f::Function, A)

Return a vector *I* of the linear indexes of *A* where *f*(*A*[*I*]) returns `true`. If there are no such elements of *A*, `find` returns an empty array.

```

julia> A = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> find(isodd, A)
2-element Array{Int64,1}:
 1
 2

```

findn(A)

Return a vector of indexes for each dimension giving the locations of the non-zeros in *A* (determined by *A*[*i*] != 0). If there are no non-zero elements of *A*, `findn` returns a 2-tuple of empty arrays.

```

julia> A = [1 2 0; 0 0 3; 0 4 0]
3×3 Array{Int64,2}:
 1  2  0
 0  0  3
 0  4  0

julia> findn(A)

```

```
([1, 1, 3, 2], [1, 2, 2, 3])

julia> A = zeros(2,2)
2×2 Array{Float64,2}:
 0.0  0.0
 0.0  0.0

julia> findn(A)
(Int64[], Int64[])
```

findnz(A)

Return a tuple (I, J, V) where I and J are the row and column indexes of the non-zero values in matrix A, and V is a vector of the non-zero values.

```
julia> A = [1 2 0; 0 0 3; 0 4 0]
3×3 Array{Int64,2}:
 1  2  0
 0  0  3
 0  4  0

julia> findnz(A)
([1, 1, 3, 2], [1, 2, 2, 3], [1, 2, 4, 3])
```

findfirst(A)

Return the linear index of the first non-zero value in A (determined by $A[i] \neq 0$). Returns 0 if no such value is found.

```
julia> A = [0 0; 1 0]
2×2 Array{Int64,2}:
 0  0
 1  0

julia> findfirst(A)
2
```

findfirst(A, v)

Return the linear index of the first element equal to v in A. Returns 0 if v is not found.

```
julia> A = [4 6; 2 2]
2×2 Array{Int64,2}:
 4  6
 2  2

julia> findfirst(A, 2)
2

julia> findfirst(A, 3)
0
```

findfirst(predicate::Function, A)

Return the linear index of the first element of A for which predicate returns true. Returns 0 if there is no such element.

```
julia> A = [1 4; 2 2]
2×2 Array{Int64,2}:
 1  4
 2  2

julia> findfirst(iseven, A)
```



```
2

julia> findfirst(x -> x>10, A)
0
```

findlast (*A*)

Return the linear index of the last non-zero value in *A* (determined by *A*[*i*] != 0). Returns 0 if there is no non-zero value in *A*.

```
julia> A = [1 0; 1 0]
2×2 Array{Int64,2}:
 1  0
 1  0

julia> findlast(A)
2

julia> A = zeros(2,2)
2×2 Array{Float64,2}:
 0.0  0.0
 0.0  0.0

julia> findlast(A)
0
```

findlast (*A*, *v*)

Return the linear index of the last element equal to *v* in *A*. Returns 0 if there is no element of *A* equal to *v*.

```
julia> A = [1 2; 2 1]
2×2 Array{Int64,2}:
 1  2
 2  1

julia> findlast(A,1)
4

julia> findlast(A,2)
3

julia> findlast(A,3)
0
```

findlast (*predicate::Function*, *A*)

Return the linear index of the last element of *A* for which *predicate* returns `true`. Returns 0 if there is no such element.

```
julia> A = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> findlast(isodd, A)
2

julia> findlast(x -> x > 5, A)
0
```

findnext (*A*, *i::Integer*)

Find the next linear index $\geq i$ of a non-zero element of *A*, or 0 if not found.

```
julia> A = [0 0; 1 0]
2×2 Array{Int64,2}:
 0  0
 1  0

julia> findnext(A,1)
2

julia> findnext(A,3)
0
```

findnext (*predicate::Function*, *A*, *i::Integer*)

Find the next linear index $\geq i$ of an element of *A* for which *predicate* returns `true`, or 0 if not found.

```
julia> A = [1 4; 2 2]
2×2 Array{Int64,2}:
 1  4
 2  2

julia> findnext(isodd, A, 1)
1

julia> findnext(isodd, A, 2)
0
```

findnext (*A*, *v*, *i::Integer*)

Find the next linear index $\geq i$ of an element of *A* equal to *v* (using `==`), or 0 if not found.

```
julia> A = [1 4; 2 2]
2×2 Array{Int64,2}:
 1  4
 2  2

julia> findnext(A,4,4)
0

julia> findnext(A,4,3)
3
```

findprev (*A*, *i::Integer*)

Find the previous linear index $\leq i$ of a non-zero element of *A*, or 0 if not found.

```
julia> A = [0 0; 1 2]
2×2 Array{Int64,2}:
 0  0
 1  2

julia> findprev(A,2)
2

julia> findprev(A,1)
0
```

findprev (*predicate::Function*, *A*, *i::Integer*)

Find the previous linear index $\leq i$ of an element of *A* for which *predicate* returns `true`, or 0 if not found.

```
julia> A = [4 6; 1 2]
2×2 Array{Int64,2}:
 4  6
 1  2
```

```
julia> findprev(isodd, A, 1)
0

julia> findprev(isodd, A, 3)
2
```

findprev (*A*, *v*, *i::Integer*)

Find the previous linear index $\leq i$ of an element of *A* equal to *v* (using `==`), or 0 if not found.

```
julia> A = [0 0; 1 2]
2×2 Array{Int64,2}:
 0  0
 1  2

julia> findprev(A, 1, 4)
2

julia> findprev(A, 1, 1)
0
```

permutedims (*A*, *perm*)

Permute the dimensions of array *A*. *perm* is a vector specifying a permutation of length `ndims(A)`. This is a generalization of transpose for multi-dimensional arrays. Transpose is equivalent to `permutedims(A, [2,1])`.

```
julia> A = reshape(collect(1:8), (2,2,2))
2×2×2 Array{Int64,3}:
[:, :, 1] =
 1  3
 2  4

[:, :, 2] =
 5  7
 6  8

julia> permutedims(A, [3, 2, 1])
2×2×2 Array{Int64,3}:
[:, :, 1] =
 1  3
 5  7

[:, :, 2] =
 2  4
 6  8
```

permutedims! (*dest*, *src*, *perm*)

Permute the dimensions of array *src* and store the result in the array *dest*. *perm* is a vector specifying a permutation of length `ndims(src)`. The preallocated array *dest* should have `size(dest) == size(src)[perm]` and is completely overwritten. No in-place permutation is supported and unexpected results will happen if *src* and *dest* have overlapping memory regions.

squeeze (*A*, *dims*)

Remove the dimensions specified by *dims* from array *A*. Elements of *dims* must be unique and within the range `1:ndims(A)`. `size(A,i)` must equal 1 for all *i* in *dims*.

```
julia> a = reshape(collect(1:4), (2,2,1,1))
2×2×1×1 Array{Int64,4}:
[:, :, 1, 1] =
```

```
1 3
2 4

julia> squeeze(a,3)
2×2×1 Array{Int64,3}:
[:, :, 1] =
1 3
2 4
```

vec (*a::AbstractArray*) → Vector

Reshape array *a* as a one-dimensional column vector.

```
julia> a = [1 2 3; 4 5 6]
2×3 Array{Int64,2}:
1 2 3
4 5 6

julia> vec(a)
6-element Array{Int64,1}:
1
4
2
5
3
6
```

promote_shape (*s1, s2*)

Check two array shapes for compatibility, allowing trailing singleton dimensions, and return whichever shape has more dimensions.

```
julia> a = ones(3,4,1,1,1);

julia> b = ones(3,4);

julia> promote_shape(a,b)
(Base.OneTo(3),Base.OneTo(4),Base.OneTo(1),Base.OneTo(1),Base.OneTo(1))

julia> promote_shape((2,3,1,4), (2,3,1,4,1))
(2,3,1,4,1)
```

checkbounds (*A, I...*)

Throw an error if the specified indices *I* are not in bounds for the given array *A*.

checkbounds (*Bool, A, I...*)

Return `true` if the specified indices *I* are in bounds for the given array *A*. Subtypes of `AbstractArray` should specialize this method if they need to provide custom bounds checking behaviors; however, in many cases one can rely on *A*'s indices and `checkindex()` (page 470).

See also `checkindex()` (page 470).

checkindex (*Bool, inds::AbstractUnitRange, index*)

Return `true` if the given *index* is within the bounds of *inds*. Custom types that would like to behave as indices for all arrays can extend this method in order to provide a specialized bounds checking implementation.

```
julia> checkindex(Bool, 1:20, 8)
true

julia> checkindex(Bool, 1:20, 21)
false
```

randsubseq (*A*, *p*) → Vector

Return a vector consisting of a random subsequence of the given array *A*, where each element of *A* is included (in order) with independent probability *p*. (Complexity is linear in $p \cdot \text{length}(A)$, so this function is efficient even if *p* is small and *A* is large.) Technically, this process is known as “Bernoulli sampling” of *A*.

randsubseq! (*S*, *A*, *p*)

Like [randsubseq\(\)](#) (page 470), but the results are stored in *S* (which is resized as needed).

45.5 Array functions

accumulate (*op*, *A*, *dim=1*)

Cumulative operation *op* along a dimension *dim* (defaults to 1). See also [accumulate!\(\)](#) (page 471) to use a preallocated output array, both for performance and to control the precision of the output (e.g. to avoid overflow). For common operations there are specialized variants of [accumulate](#), see: [cumsum\(\)](#) (page 471), [cumprod\(\)](#) (page 471)

```
julia> accumulate(+, [1,2,3])
3-element Array{Int64,1}:
 1
 3
 6

julia> accumulate(*, [1,2,3])
3-element Array{Int64,1}:
 1
 2
 6
```

accumulate! (*op*, *B*, *A*, *dim=1*)

Cumulative operation *op* on *A* along a dimension, storing the result in *B*. The dimension defaults to 1. See also [accumulate\(\)](#) (page 471).

cumprod (*A*, *dim=1*)

Cumulative product along a dimension *dim* (defaults to 1). See also [cumprod!\(\)](#) (page 471) to use a preallocated output array, both for performance and to control the precision of the output (e.g. to avoid overflow).

```
julia> a = [1 2 3; 4 5 6]
2×3 Array{Int64,2}:
 1  2  3
 4  5  6

julia> cumprod(a,1)
2×3 Array{Int64,2}:
 1  2  3
 4 10 18

julia> cumprod(a,2)
2×3 Array{Int64,2}:
 1  2  6
 4 20 120
```

cumprod! (*B*, *A*[, *dim*])

Cumulative product of *A* along a dimension, storing the result in *B*. The dimension defaults to 1.

cumsum (*A*, *dim=1*)

Cumulative sum along a dimension *dim* (defaults to 1). See also [cumsum!\(\)](#) (page 472) to use a preallocated output array, both for performance and to control the precision of the output (e.g. to avoid overflow).

```
julia> a = [1 2 3; 4 5 6]
2×3 Array{Int64,2}:
 1  2  3
 4  5  6

julia> cumsum(a,1)
2×3 Array{Int64,2}:
 1  2  3
 5  7  9

julia> cumsum(a,2)
2×3 Array{Int64,2}:
 1  3  6
 4  9 15
```

cumsum! (*B*, *A*[, *dim*])

Cumulative sum of *A* along a dimension, storing the result in *B*. The dimension defaults to 1.

cumsum_kbn (*A*[, *dim::Integer=1*])

Cumulative sum along a dimension, using the Kahan-Babuska-Neumaier compensated summation algorithm for additional accuracy. The dimension defaults to 1.

cummin (*A*[, *dim*])

Cumulative minimum along a dimension. The dimension defaults to 1.

cummax (*A*[, *dim*])

Cumulative maximum along a dimension. The dimension defaults to 1.

diff (*A*[, *dim::Integer=1*])

Finite difference operator of matrix or vector *A*. If *A* is a matrix, compute the finite difference over a dimension *dim* (default 1).

Example

```
julia> a = [2 4; 6 16]
2×2 Array{Int64,2}:
 2  4
 6 16

julia> diff(a,2)
2×1 Array{Int64,2}:
 2
10
```

gradient (*F::AbstractVector*[, *h::Real*])

Compute differences along vector *F*, using *h* as the spacing between points. The default spacing is one.

Example

```
julia> a = [2,4,6,8];

julia> gradient(a)
4-element Array{Float64,1}:
 2.0
 2.0
 2.0
 2.0
```

rot180 (*A*)

Rotate matrix *A* 180 degrees.

```
julia> a = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> rot180(a)
2×2 Array{Int64,2}:
 4  3
 2  1
```

rot180(A, k)

Rotate matrix A 180 degrees an integer k number of times. If k is even, this is equivalent to a copy.

```
julia> a = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> rot180(a,1)
2×2 Array{Int64,2}:
 4  3
 2  1

julia> rot180(a,2)
2×2 Array{Int64,2}:
 1  2
 3  4
```

rot190(A)

Rotate matrix A left 90 degrees.

```
julia> a = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> rot190(a)
2×2 Array{Int64,2}:
 2  4
 1  3
```

rot190(A, k)

Rotate matrix A left 90 degrees an integer k number of times. If k is zero or a multiple of four, this is equivalent to a copy.

```
julia> a = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> rot190(a,1)
2×2 Array{Int64,2}:
 2  4
 1  3

julia> rot190(a,2)
2×2 Array{Int64,2}:
 4  3
 2  1
```

```
julia> rotl90(a,3)
2×2 Array{Int64,2}:
 3  1
 4  2

julia> rotl90(a,4)
2×2 Array{Int64,2}:
 1  2
 3  4
```

rotr90(A)

Rotate matrix A right 90 degrees.

```
julia> a = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> rotr90(a)
2×2 Array{Int64,2}:
 3  1
 4  2
```

rotr90(A, k)

Rotate matrix A right 90 degrees an integer k number of times. If k is zero or a multiple of four, this is equivalent to a copy.

```
julia> a = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> rotr90(a,1)
2×2 Array{Int64,2}:
 3  1
 4  2

julia> rotr90(a,2)
2×2 Array{Int64,2}:
 4  3
 2  1

julia> rotr90(a,3)
2×2 Array{Int64,2}:
 2  4
 1  3

julia> rotr90(a,4)
2×2 Array{Int64,2}:
 1  2
 3  4
```

reducedim(f, A, region[, v0])

Reduce 2-argument function *f* along dimensions of A. *region* is a vector specifying the dimensions to reduce, and *v0* is the initial value to use in the reductions. For +, *, max and min the *v0* argument is optional.

The associativity of the reduction is implementation-dependent; if you need a particular associativity, e.g. left-to-right, you should write your own loop. See documentation for [reduce\(\)](#) (page 385).


```

julia> a = reshape(collect(1:16), (4,4))
4×4 Array{Int64,2}:
 1  5  9 13
 2  6 10 14
 3  7 11 15
 4  8 12 16

julia> reducedim(max, a, 2)
4×1 Array{Int64,2}:
13
14
15
16

julia> reducedim(max, a, 1)
1×4 Array{Int64,2}:
 4  8 12 16

```

mapreducedim(*f*, *op*, *A*, *region*[, *v0*])

Evaluates to the same as `reducedim(op, map(f, A), region, f(v0))`, but is generally faster because the intermediate array is avoided.

```

julia> a = reshape(collect(1:16), (4,4))
4×4 Array{Int64,2}:
 1  5  9 13
 2  6 10 14
 3  7 11 15
 4  8 12 16

julia> mapreducedim(isodd, *, a, 1)
1×4 Array{Bool,2}:
false false false false

julia> mapreducedim(isodd, |, a, 1, true)
1×4 Array{Bool,2}:
true true true true

```

mapslices(*f*, *A*, *dims*)

Transform the given dimensions of array *A* using function *f*. *f* is called on each slice of *A* of the form *A*[..., :, ..., :, ...]. *dims* is an integer vector specifying where the colons go in this expression. The results are concatenated along the remaining dimensions. For example, if *dims* is `[1, 2]` and *A* is 4-dimensional, *f* is called on *A*[:, :, *i*, *j*] for all *i* and *j*.

```

julia> a = reshape(collect(1:16), (2,2,2,2))
2×2×2×2 Array{Int64,4}:
[:, :, 1, 1] =
 1  3
 2  4

[:, :, 2, 1] =
 5  7
 6  8

[:, :, 1, 2] =
 9 11
10 12

[:, :, 2, 2] =
13 15

```

```

14 16

julia> mapslices(sum, a, [1,2])
1×1×2×2 Array{Int64,4}:
[:, :, 1, 1] =
 10

[:, :, 2, 1] =
 26

[:, :, 1, 2] =
 42

[:, :, 2, 2] =
 58

```

sum_kbn(A)

Returns the sum of all array elements, using the Kahan-Babuska-Neumaier compensated summation algorithm for additional accuracy.

45.6 Combinatorics

randperm([rng=GLOBAL_RNG], n::Integer)

Construct a random permutation of length n . The optional `rng` argument specifies a random number generator (see [Random Numbers](#) (page 440)). To randomly permute an arbitrary vector, see [shuffle\(\)](#) (page 477) or [shuffle!\(\)](#) (page 477).

invperm(v)

Return the inverse permutation of v . If $B = A[v]$, then $A == B[\text{invperm}(v)]$.

```

julia> v = [2; 4; 3; 1];

julia> invperm(v)
4-element Array{Int64,1}:
 4
 1
 3
 2

julia> A = ['a', 'b', 'c', 'd'];

julia> B = A[v]
4-element Array{Char,1}:
 'b'
 'd'
 'c'
 'a'

julia> B[invperm(v)]
4-element Array{Char,1}:
 'a'
 'b'
 'c'
 'd'

```

isperm(v) → Bool

Returns `true` if `v` is a valid permutation.

```
julia> isperm([1; 2])
true

julia> isperm([1; 3])
false
```

permute! (`v, p`)

Permute vector `v` in-place, according to permutation `p`. No checking is done to verify that `p` is a permutation.

To return a new permutation, use `v[p]`. Note that this is generally faster than `permute!(v, p)` for large vectors.

ipermute! (`v, p`)

Like `permute!`, but the inverse of the given permutation is applied.

randcycle (`[rng=GLOBAL_RNG]`, `n::Integer`)

Construct a random cyclic permutation of length `n`. The optional `rng` argument specifies a random number generator, see [Random Numbers](#) (page 440).

shuffle (`[rng=GLOBAL_RNG]`, `v`)

Return a randomly permuted copy of `v`. The optional `rng` argument specifies a random number generator (see [Random Numbers](#) (page 440)). To permute `v` in-place, see `shuffle!()` (page 477). To obtain randomly permuted indices, see `randperm()` (page 476).

shuffle! (`[rng=GLOBAL_RNG]`, `v`)

In-place version of `shuffle()` (page 477): randomly permute the array `v` in-place, optionally supplying the random-number generator `rng`.

reverse (`v`, `start=1`, `stop=length(v)`)

Return a copy of `v` reversed from `start` to `stop`.

reverseind (`v`, `i`)

Given an index `i` in `reverse(v)`, return the corresponding index in `v` so that `v[reverseind(v, i)] == reverse(v)[i]`. (This can be nontrivial in the case where `v` is a Unicode string.)

reverse! (`v`, `start=1`, `stop=length(v)`) → `v`

In-place version of `reverse()` (page 446).

45.7 BitArrays

BitArrays are space-efficient “packed” boolean arrays, which store one bit per boolean value. They can be used similarly to `Array{Bool}` arrays (which store one byte per boolean value), and can be converted to/from the latter via `Array{Bool}(bitarray)` and `BitArray(array)`, respectively.

flipbits! (`B::BitArray{N}`) → `BitArray{N}`

Performs a bitwise not operation on `B`. See [~ operator](#) (page 411).

```
julia> A = trues(2,2)
2×2 BitArray{2}:
 true  true
 true  true

julia> flipbits!(A)
2×2 BitArray{2}:
 false false
 false false
```

rol! (*dest::BitVector, src::BitVector, i::Integer*) → *BitVector*

Performs a left rotation operation on *src* and puts the result into *dest*. *i* controls how far to rotate the bits.

rol! (*B::BitVector, i::Integer*) → *BitVector*

Performs a left rotation operation in-place on *B*. *i* controls how far to rotate the bits.

rol (*B::BitVector, i::Integer*) → *BitVector*

Performs a left rotation operation, returning a new *BitVector*. *i* controls how far to rotate the bits. See also [rol!\(\)](#) (page 477).

```
julia> A = BitArray([true, true, false, false, true])
5-element BitArray{1}:
 true
 true
 false
 false
 true

julia> rol(A, 1)
5-element BitArray{1}:
 true
 false
 false
 true
 true

julia> rol(A, 2)
5-element BitArray{1}:
 false
 false
 true
 true
 true

julia> rol(A, 5)
5-element BitArray{1}:
 true
 true
 false
 false
 true
```

ror! (*dest::BitVector, src::BitVector, i::Integer*) → *BitVector*

Performs a right rotation operation on *src* and puts the result into *dest*. *i* controls how far to rotate the bits.

ror! (*B::BitVector, i::Integer*) → *BitVector*

Performs a right rotation operation in-place on *B*. *i* controls how far to rotate the bits.

ror (*B::BitVector, i::Integer*) → *BitVector*

Performs a right rotation operation on *B*, returning a new *BitVector*. *i* controls how far to rotate the bits. See also [ror!\(\)](#) (page 478).

```
julia> A = BitArray([true, true, false, false, true])
5-element BitArray{1}:
 true
 true
 false
 false
 true
```

```

julia> ror(A,1)
5-element BitArray{1}:
 true
 true
 true
 false
 false

julia> ror(A,2)
5-element BitArray{1}:
 false
 true
 true
 true
 false

julia> ror(A,5)
5-element BitArray{1}:
 true
 true
 false
 false
 true

```

45.8 Sparse Vectors and Matrices

Sparse vectors and matrices largely support the same set of operations as their dense counterparts. The following functions are specific to sparse arrays.

sparse (*I*, *J*, *V* [, *m*, *n*, *combine*])

Create a sparse matrix *S* of dimensions *m* × *n* such that *S*[*I*[*k*], *J*[*k*]] = *V*[*k*]. The *combine* function is used to combine duplicates. If *m* and *n* are not specified, they are set to `maximum(I)` and `maximum(J)` respectively. If the *combine* function is not supplied, *combine* defaults to `+` unless the elements of *V* are Booleans in which case *combine* defaults to `|`. All elements of *I* must satisfy `1 <= I[k] <= m`, and all elements of *J* must satisfy `1 <= J[k] <= n`. Numerical zeros in (*I*, *J*, *V*) are retained as structural nonzeros; to drop numerical zeros, use `dropzeros!()` (page 483).

For additional documentation and an expert driver, see `Base.SparseArrays.sparse!`.

```

julia> Is = [1; 2; 3];

julia> Js = [1; 2; 3];

julia> Vs = [1; 2; 3];

julia> sparse(Is, Js, Vs)
3×3 sparse matrix with 3 Int64 nonzero entries:
 [1, 1] = 1
 [2, 2] = 2
 [3, 3] = 3

```

sparsevec (*I*, *V* [, *m*, *combine*])

Create a sparse vector *S* of length *m* such that *S*[*I*[*k*]] = *V*[*k*]. Duplicates are combined using the *combine* function, which defaults to `+` if no *combine* argument is provided, unless the elements of *V* are Booleans in which case *combine* defaults to `|`.

sparsevec (*D::Dict{*m*}*)

Create a sparse vector of length *m* where the nonzero indices are keys from the dictionary, and the nonzero values are the values from the dictionary.

issparse (*S*)

Returns `true` if *S* is sparse, and `false` otherwise.

sparse (*A*)

Convert an AbstractMatrix *A* into a sparse matrix.

```
julia> A = eye(3)
3×3 Array{Float64,2}:
 1.0  0.0  0.0
 0.0  1.0  0.0
 0.0  0.0  1.0

julia> sparse(A)
3×3 sparse matrix with 3 Float64 nonzero entries:
 [1, 1] = 1.0
 [2, 2] = 1.0
 [3, 3] = 1.0
```

sparsevec (*A*)

Convert a vector *A* into a sparse vector of length *m*.

full (*S*)

Convert a sparse matrix or vector *S* into a dense matrix or vector.

```
julia> A = speye(3)
3×3 sparse matrix with 3 Float64 nonzero entries:
 [1, 1] = 1.0
 [2, 2] = 1.0
 [3, 3] = 1.0

julia> full(A)
3×3 Array{Float64,2}:
 1.0  0.0  0.0
 0.0  1.0  0.0
 0.0  0.0  1.0
```

nnz (*A*)

Returns the number of stored (filled) elements in a sparse array.

```
julia> A = speye(3)
3×3 sparse matrix with 3 Float64 nonzero entries:
 [1, 1] = 1.0
 [2, 2] = 1.0
 [3, 3] = 1.0

julia> nnz(A)
3
```

spzeros (*[type]*, *m*, *n*)

Create a sparse vector of length *m* or sparse matrix of size *m* × *n*. This sparse array will not contain any nonzero values. No storage will be allocated for nonzero values during construction. The type defaults to `Float64` if not specified.

spones (*S*)

Create a sparse array with the same structure as that of *S*, but with every nonzero element having the value `1.0`.

```
julia> A = sparse([1,2,3,4],[2,4,3,1],[5.,4.,3.,2.])
4×4 sparse matrix with 4 Float64 nonzero entries:
 [4, 1] = 2.0
 [1, 2] = 5.0
 [3, 3] = 3.0
 [2, 4] = 4.0

julia> spones(A)
4×4 sparse matrix with 4 Float64 nonzero entries:
 [4, 1] = 1.0
 [1, 2] = 1.0
 [3, 3] = 1.0
 [2, 4] = 1.0
```

Note the difference from `speye()` (page 481).

`speye([type], m[, n])`

Create a sparse identity matrix of size $m \times m$. When n is supplied, create a sparse identity matrix of size $m \times n$. The type defaults to `Float64` if not specified.

`sparse(I, m, n)` is equivalent to `speye{Int}(m, n)`, and `sparse(α *I, m, n)` can be used to efficiently create a sparse multiple α of the identity matrix.

`speye(S)`

Create a sparse identity matrix with the same size as S .

```
julia> A = sparse([1,2,3,4],[2,4,3,1],[5.,4.,3.,2.])
4×4 sparse matrix with 4 Float64 nonzero entries:
 [4, 1] = 2.0
 [1, 2] = 5.0
 [3, 3] = 3.0
 [2, 4] = 4.0

julia> speye(A)
4×4 sparse matrix with 4 Float64 nonzero entries:
 [1, 1] = 1.0
 [2, 2] = 1.0
 [3, 3] = 1.0
 [4, 4] = 1.0
```

Note the difference from `spones()` (page 480).

`spdiagm(B, d[, m, n])`

Construct a sparse diagonal matrix. B is a tuple of vectors containing the diagonals and d is a tuple containing the positions of the diagonals. In the case the input contains only one diagonal, B can be a vector (instead of a tuple) and d can be the diagonal position (instead of a tuple), defaulting to 0 (diagonal). Optionally, m and n specify the size of the resulting sparse matrix.

```
julia> spdiagm(([1,2,3,4],[4,3,2,1]),(-1,1))
5×5 sparse matrix with 8 Int64 nonzero entries:
 [2, 1] = 1
 [1, 2] = 4
 [3, 2] = 2
 [2, 3] = 3
 [4, 3] = 3
 [3, 4] = 2
 [5, 4] = 4
 [4, 5] = 1
```

`sprand([rng], [type], m[, n], p::AbstractFloat[, rfn])`

Create a random length m sparse vector or m by n sparse matrix, in which the probability of any element being nonzero is independently given by p (and hence the mean density of nonzeros is also exactly p). Nonzero values are sampled from the distribution specified by `rfn` and have the type `type`. The uniform distribution is used in case `rfn` is not specified. The optional `rng` argument specifies a random number generator, see [Random Numbers](#) (page 440).

sprandn (`[rng]`, `m`, `n`, `p::AbstractFloat`)

Create a random sparse vector of length m or sparse matrix of size m by n with the specified (independent) probability p of any entry being nonzero, where nonzero values are sampled from the normal distribution. The optional `rng` argument specifies a random number generator, see [Random Numbers](#) (page 440).

nonzeros (`A`)

Return a vector of the structural nonzero values in sparse array `A`. This includes zeros that are explicitly stored in the sparse array. The returned vector points directly to the internal nonzero storage of `A`, and any modifications to the returned vector will mutate `A` as well. See [rowvals\(\)](#) (page 482) and [nzrange\(\)](#) (page 482).

```
julia> A = speye(3)
3×3 sparse matrix with 3 Float64 nonzero entries:
 [1, 1] = 1.0
 [2, 2] = 1.0
 [3, 3] = 1.0

julia> nonzeros(A)
3-element Array{Float64,1}:
 1.0
 1.0
 1.0
```

rowvals (`A::SparseMatrixCSC`)

Return a vector of the row indices of `A`. Any modifications to the returned vector will mutate `A` as well. Providing access to how the row indices are stored internally can be useful in conjunction with iterating over structural nonzero values. See also [nonzeros\(\)](#) (page 482) and [nzrange\(\)](#) (page 482).

```
julia> A = speye(3)
3×3 sparse matrix with 3 Float64 nonzero entries:
 [1, 1] = 1.0
 [2, 2] = 1.0
 [3, 3] = 1.0

julia> rowvals(A)
3-element Array{Int64,1}:
 1
 2
 3
```

nzrange (`A::SparseMatrixCSC`, `col::Integer`)

Return the range of indices to the structural nonzero values of a sparse matrix column. In conjunction with [nonzeros\(\)](#) (page 482) and [rowvals\(\)](#) (page 482), this allows for convenient iterating over a sparse matrix:

```
A = sparse(I, J, V)
rows = rowvals(A)
vals = nonzeros(A)
m, n = size(A)
for i = 1:n
    for j in nzrange(A, i)
        row = rows[j]
        val = vals[j]
        # perform sparse wizardry...
```



```
end
end
```

dropzeros! (*A::SparseMatrixCSC*, *trim::Bool* = *true*)

Removes stored numerical zeros from *A*, optionally trimming resulting excess space from *A.rowval* and *A.nzval* when *trim* is *true*.

For an out-of-place version, see [dropzeros\(\)](#) (page 483). For algorithmic information, see `Base.SparseArrays.fkeep!()`.

dropzeros (*A::SparseMatrixCSC*, *trim::Bool* = *true*)

Generates a copy of *A* and removes stored numerical zeros from that copy, optionally trimming excess space from the result's *rowval* and *nzval* arrays when *trim* is *true*.

For an in-place version and algorithmic information, see [dropzeros!\(\)](#) (page 483).

dropzeros! (*x::SparseVector*, *trim::Bool* = *true*)

Removes stored numerical zeros from *x*, optionally trimming resulting excess space from *x.nzind* and *x.nzval* when *trim* is *true*.

For an out-of-place version, see `Base.SparseArrays.dropzeros()`. For algorithmic information, see `Base.SparseArrays.fkeep!()`.

dropzeros (*x::SparseVector*, *trim::Bool* = *true*)

Generates a copy of *x* and removes numerical zeros from that copy, optionally trimming excess space from the result's *nzind* and *nzval* arrays when *trim* is *true*.

For an in-place version and algorithmic information, see `Base.SparseArrays.dropzeros!()`.

permute{*Tv*, *Ti*, *Tp*<:Integer, *Tq*<:Integer} (*A::SparseMatrixCSC*{*Tv*, *Ti*}, *p::AbstractVector*{*Tp*}, *q::AbstractVector*{*Tq*}) ()

Bilaterally permute *A*, returning *PAQ* (*A*[*p*, *q*]). Column-permutation *q*'s length must match *A*'s column count (`length(q) == A.n`). Row-permutation *p*'s length must match *A*'s row count (`length(p) == A.m`).

For expert drivers and additional information, see `Base.SparseArrays.permute!()`.

permute!{*Tv*, *Ti*, *Tp*<:Integer, *Tq*<:Integer} (*X::SparseMatrixCSC*{*Tv*, *Ti*}, *A::SparseMatrixCSC*{*Tv*, *Ti*}, *p::AbstractVector*{*Tp*}, *q::AbstractVector*{*Tq*}[, *C::SparseMatrixCSC*{*Tv*, *Ti*}])

Bilaterally permute *A*, storing result *PAQ* (*A*[*p*, *q*]) in *X*. Stores intermediate result $(AQ)^T$ (`transpose(A[:, q])`) in optional argument *C* if present. Requires that none of *X*, *A*, and, if present, *C* alias each other; to store result *PAQ* back into *A*, use the following method lacking *X*:

```
permute!{Tv,Ti,Tp<:Integer,Tq<:Integer}(A::SparseMatrixCSC{Tv,Ti}, p::AbstractVector{Tp},
    q::AbstractVector{Tq}[, C::SparseMatrixCSC{Tv,Ti}[, workcolptr::Vector{Ti}]])
```

X's dimensions must match those of *A* (`X.m == A.m` and `X.n == A.n`), and *X* must have enough storage to accommodate all allocated entries in *A* (`length(X.rowval) >= nnz(A)` and `length(X.nzval) >= nnz(A)`). Column-permutation *q*'s length must match *A*'s column count (`length(q) == A.n`). Row-permutation *p*'s length must match *A*'s row count (`length(p) == A.m`).

C's dimensions must match those of `transpose(A)` (`C.m == A.n` and `C.n == A.m`), and *C* must have enough storage to accommodate all allocated entries in *A* (`length(C.rowval) >= nnz(A)` and `length(C.nzval) >= nnz(A)`).

For additional (algorithmic) information, and for versions of these methods that forgo argument checking, see (unexported) parent methods `Base.SparseArrays.unchecked_noalias_permute!()` and `Base.SparseArrays.unchecked_aliasing_permute!()`.

See also: `Base.SparseArrays.permute()`

Tasks and Parallel Computing

46.1 Tasks

Task (*func*)

Create a `Task` (i.e. coroutine) to execute the given function (which must be callable with no arguments). The task exits when this function returns.

```
julia> a() = det(rand(1000, 1000));

julia> b = Task(a);
```

In this example, `b` is a runnable `Task` that hasn't started yet.

yieldto (*t::Task, arg = nothing*)

Switch to the given task. The first time a task is switched to, the task's function is called with no arguments. On subsequent switches, `arg` is returned from the task's last call to `yieldto`. This is a low-level call that only switches tasks, not considering states or scheduling in any way. Its use is discouraged.

current_task ()

Get the currently running `Task` (page 485).

istaskdone (*t::Task*) → Bool

Determine whether a task has exited.

```
julia> a2() = det(rand(1000, 1000));

julia> b = Task(a2);

julia> istaskdone(b)
false

julia> schedule(b);

julia> yield();

julia> istaskdone(b)
true
```

istaskstarted (*t::Task*) → Bool

Determine whether a task has started executing.

```
julia> a3() = det(rand(1000, 1000));

julia> b = Task(a3);
```

```
julia> istaskstarted(b)
false
```

consume (*task*, *values...*)

Receive the next value passed to `produce` by the specified task. Additional arguments may be passed, to be returned from the last `produce` call in the producer.

produce (*value*)

Send the given value to the last `consume` call, switching to the consumer task. If the next `consume` call passes any values, they are returned by `produce`.

yield ()

Switch to the scheduler to allow another scheduled task to run. A task that calls this function is still runnable, and will be restarted immediately if there are no other runnable tasks.

task_local_storage (*key*)

Look up the value of a key in the current task's task-local storage.

task_local_storage (*key*, *value*)

Assign a value to a key in the current task's task-local storage.

task_local_storage (*body*, *key*, *value*)

Call the function `body` with a modified task-local storage, in which `value` is assigned to `key`; the previous value of `key`, or lack thereof, is restored afterwards. Useful for emulating dynamic scoping.

Condition ()

Create an edge-triggered event source that tasks can wait for. Tasks that call `wait()` (page 490) on a `Condition` are suspended and queued. Tasks are woken up when `notify()` (page 486) is later called on the `Condition`. Edge triggering means that only tasks waiting at the time `notify()` (page 486) is called can be woken up. For level-triggered notifications, you must keep extra state to keep track of whether a notification has happened. The `Channel` type does this, and so can be used for level-triggered events.

notify (*condition*, *val=nothing*; *all=true*, *error=false*)

Wake up tasks waiting for a condition, passing them `val`. If `all` is `true` (the default), all waiting tasks are woken, otherwise only one is. If `error` is `true`, the passed value is raised as an exception in the woken tasks.

schedule (*t::Task*, [*val*]; *error=false*)

Add a `Task` (page 485) to the scheduler's queue. This causes the task to run constantly when the system is otherwise idle, unless the task performs a blocking operation such as `wait()` (page 490).

If a second argument `val` is provided, it will be passed to the task (via the return value of `yieldto()` (page 485)) when it runs again. If `error` is `true`, the value is raised as an exception in the woken task.

```
julia> a5() = det(rand(1000, 1000));

julia> b = Task(a5);

julia> istaskstarted(b)
false

julia> schedule(b);

julia> yield();

julia> istaskstarted(b)
true

julia> istaskdone(b)
true
```

@schedule()

Wrap an expression in a [Task](#) (page 485) and add it to the local machine’s scheduler queue.

@task()

Wrap an expression in a [Task](#) (page 485) without executing it, and return the [Task](#) (page 485). This only creates a task, and does not run it.

```
julia> a1() = det(rand(1000, 1000));

julia> b = @task a1();

julia> istaskstarted(b)
false

julia> schedule(b);

julia> yield();

julia> istaskdone(b)
true
```

sleep(seconds)

Block the current task for a specified number of seconds. The minimum sleep time is 1 millisecond or input of 0.001.

Channel{T} (sz::Int)

Constructs a `Channel` with an internal buffer that can hold a maximum of `sz` objects of type `T`. [put!\(\)](#) (page 487) calls on a full channel block until an object is removed with [take!\(\)](#) (page 487).

`Channel{0}` constructs an unbuffered channel. `put!` blocks until a matching `take!` is called. And vice-versa.

Other constructors:

- `Channel{Inf}`: equivalent to `Channel{Any}(typemax{Int})`
- `Channel{sz}`: equivalent to `Channel{Any}(sz)`

put!(c::Channel, v)

Appends an item `v` to the channel `c`. Blocks if the channel is full.

For unbuffered channels, blocks until a [take!\(\)](#) (page 487) is performed by a different task.

take!(c::Channel)

Removes and returns a value from a `Channel`. Blocks until data is available.

For unbuffered channels, blocks until a [put!\(\)](#) (page 487) is performed by a different task.

isready(c::Channel)

Determine whether a `Channel` has a value stored to it. Returns immediately, does not block.

For unbuffered channels returns `true` if there are tasks waiting on a [put!\(\)](#) (page 487).

fetch(c::Channel)

Waits for and gets the first available item from the channel. Does not remove the item. `fetch` is unsupported on an unbuffered (0-size) channel.

close(c::Channel)

Closes a channel. An exception is thrown by:

- [put!\(\)](#) (page 487) on a closed channel.
- [take!\(\)](#) (page 487) and [fetch\(\)](#) (page 487) on an empty, closed channel.

46.2 General Parallel Computing Support

addprocs (*np::Integer; restrict=true, kwargs...*) → List of process identifiers

Launches workers using the in-built LocalManager which only launches workers on the local host. This can be used to take advantage of multiple cores. `addprocs(4)` will add 4 processes on the local machine. If `restrict` is `true`, binding is restricted to `127.0.0.1`. Keyword args `dir`, `exename`, `exeflags`, `topology`, and `enable_threaded_blas` have the same effect as documented for `addprocs(machines)`.

addprocs (; *kwargs...*) → List of process identifiers

Equivalent to `addprocs(Sys.CPU_CORES; kwargs...)`

Note that workers do not run a `.juliarc.jl` startup script, nor do they synchronize their global state (such as global variables, new method definitions, and loaded modules) with any of the other running processes.

addprocs (*machines; tunnel=false, sshflags="", max_parallel=10, kwargs...*) → List of process identifiers

Add processes on remote machines via SSH. Requires `julia` to be installed in the same location on each node, or to be available via a shared file system.

`machines` is a vector of machine specifications. Workers are started for each specification.

A machine specification is either a string `machine_spec` or a tuple - (`machine_spec`, `count`).

`machine_spec` is a string of the form `[user@]host[:port] [bind_addr[:port]]`. `user` defaults to current user, `port` to the standard ssh port. If `[bind_addr[:port]]` is specified, other workers will connect to this worker at the specified `bind_addr` and `port`.

`count` is the number of workers to be launched on the specified host. If specified as `:auto` it will launch as many workers as the number of cores on the specific host.

Keyword arguments:

- `tunnel`: if `true` then SSH tunneling will be used to connect to the worker from the master process. Default is `false`.
- `sshflags`: specifies additional ssh options, e.g. `sshflags="-i /home/foo/bar.pem"`
- `max_parallel`: specifies the maximum number of workers connected to in parallel at a host. Defaults to 10.
- `dir`: specifies the working directory on the workers. Defaults to the host's current directory (as found by `pwd()`)
- `enable_threaded_blas`: if `true` then BLAS will run on multiple threads in added processes. Default is `false`.
- `exename`: name of the `julia` executable. Defaults to `"$JULIA_HOME/julia"` or `"$JULIA_HOME/julia-debug"` as the case may be.
- `exeflags`: additional flags passed to the worker processes.
- `topology`: Specifies how the workers connect to each other. Sending a message between unconnected workers results in an error.

– `topology=:all_to_all`: All processes are connected to each other. This is the default.

– `topology=:master_slave`: Only the driver process, i.e. `pid 1` connects to the workers. The workers do not connect to each other.

– `topology=:custom`: The launch method of the cluster manager specifies the connection topology via fields `ident` and `connect_idents` in `WorkerConfig`. A worker with a cluster manager identity `ident` will connect to all workers specified in `connect_idents`.

Environment variables :

If the master process fails to establish a connection with a newly launched worker within 60.0 seconds, the worker treats it as a fatal situation and terminates. This timeout can be controlled via environment variable `JULIA_WORKER_TIMEOUT`. The value of `JULIA_WORKER_TIMEOUT` on the master process specifies the number of seconds a newly launched worker waits for connection establishment.

addprocs (*manager::ClusterManager; kwargs...*) → List of process identifiers

Launches worker processes via the specified cluster manager.

For example, Beowulf clusters are supported via a custom cluster manager implemented in the package `ClusterManagers.jl`.

The number of seconds a newly launched worker waits for connection establishment from the master can be specified via variable `JULIA_WORKER_TIMEOUT` in the worker process's environment. Relevant only when using TCP/IP as transport.

nprocs ()

Get the number of available processes.

nworkers ()

Get the number of available worker processes. This is one less than `nprocs()`. Equal to `nprocs()` if `nprocs() == 1`.

procs ()

Returns a list of all process identifiers.

procs (*pid::Integer*)

Returns a list of all process identifiers on the same physical node. Specifically all workers bound to the same ip-address as `pid` are returned.

workers ()

Returns a list of all worker process identifiers.

rmprocs (*pids...; waitfor=0.0*)

Removes the specified workers. Note that only process 1 can add or remove workers - if another worker tries to call `rmprocs`, an error will be thrown. The optional argument `waitfor` determines how long the first process will wait for the workers to shut down.

interrupt (*pids::AbstractVector=workers()*)

Interrupt the current executing task on the specified workers. This is equivalent to pressing Ctrl-C on the local machine. If no arguments are given, all workers are interrupted.

interrupt (*pids::Integer...*)

Interrupt the current executing task on the specified workers. This is equivalent to pressing Ctrl-C on the local machine. If no arguments are given, all workers are interrupted.

myid ()

Get the id of the current process.

asynccmap (*f, c...*) → collection

Transform collection `c` by applying `@async f` to each element.

For multiple collection arguments, apply `f` elementwise.

pmap (*[::AbstractWorkerPool], f, c...; distributed=true, batch_size=1, on_error=nothing, retry_n=0, retry_max_delay=DEFAULT_RETRY_MAX_DELAY, retry_on=DEFAULT_RETRY_ON*) → collection
Transform collection `c` by applying `f` to each element using available workers and tasks.

For multiple collection arguments, apply `f` elementwise.

Note that `f` must be made available to all worker processes; see [Code Availability and Loading Packages](#) (page 200) for details.

If a worker pool is not specified, all available workers, i.e., the default worker pool is used.

By default, `pmap` distributes the computation over all specified workers. To use only the local process and distribute over tasks, specify `distributed=false`. This is equivalent to `asynccmap()` (page 489).

`pmap` can also use a mix of processes and tasks via the `batch_size` argument. For batch sizes greater than 1, the collection is split into multiple batches, which are distributed across workers. Each such batch is processed in parallel via tasks in each worker. The specified `batch_size` is an upper limit, the actual size of batches may be smaller and is calculated depending on the number of workers available and length of the collection.

Any error stops `pmap` from processing the remainder of the collection. To override this behavior you can specify an error handling function via argument `on_error` which takes in a single argument, i.e., the exception. The function can stop the processing by rethrowing the error, or, to continue, return any value which is then returned inline with the results to the caller.

Failed computation can also be retried via `retry_on`, `retry_n`, `retry_max_delay`, which are passed through to `retry` as arguments `retry_on`, `n` and `max_delay` respectively. If batching is specified, and an entire batch fails, all items in the batch are retried.

The following are equivalent:

- `pmap(f, c; distributed=false)` and `asynccmap(f, c)`
- `pmap(f, c; retry_n=1)` and `asynccmap(retry(remote(f)), c)`
- `pmap(f, c; retry_n=1, on_error=e->e)` and `asynccmap(x->try`
`retry(remote(f))(x) catch e; e end, c)`

RemoteException (*captured*)

Exceptions on remote computations are captured and rethrown locally. A `RemoteException` wraps the `pid` of the worker and a captured exception. A `CapturedException` captures the remote exception and a serializable form of the call stack when the exception was raised.

Future (*pid::Integer=myid()*)

Create a `Future` on process `pid`. The default `pid` is the current process.

RemoteChannel (*pid::Integer=myid()*)

Make a reference to a `Channel{Any}(1)` on process `pid`. The default `pid` is the current process.

RemoteChannel (*f::Function, pid::Integer=myid()*)

Create references to remote channels of a specific size and type. `f()` is a function that when executed on `pid` must return an implementation of an `AbstractChannel`.

For example, `RemoteChannel(()->Channel{Int}(10), pid)`, will return a reference to a channel of type `Int` and size 10 on `pid`.

The default `pid` is the current process.

wait (*[x]*)

Block the current task until some event occurs, depending on the type of the argument:

- `RemoteChannel` (page 490) : Wait for a value to become available on the specified remote channel.
- `Future` (page 490) : Wait for a value to become available for the specified future.
- `Channel`: Wait for a value to be appended to the channel.
- `Condition` (page 486): Wait for `notify()` (page 486) on a condition.
- `Process`: Wait for a process or process chain to exit. The `exitcode` field of a process can be used to determine success or failure.
- `Task` (page 485): Wait for a `Task` to finish, returning its result value. If the task fails with an exception, the exception is propagated (re-thrown in the task that called `wait`).

- RawFD**: Wait for changes on a file descriptor (see `poll_fd()` (page 573) for keyword arguments and return code)

If no argument is passed, the task blocks for an undefined period. A task can only be restarted by an explicit call to `schedule()` (page 486) or `yieldto()` (page 485).

Often `wait` is called within a `while` loop to ensure a waited-for condition is met before proceeding.

fetch(*x*)

Waits and fetches a value from *x* depending on the type of *x*:

- Future** (page 490): Wait for and get the value of a `Future`. The fetched value is cached locally. Further calls to `fetch` on the same reference return the cached value. If the remote value is an exception, throws a `RemoteException` (page 490) which captures the remote exception and backtrace.
- RemoteChannel** (page 490): Wait for and get the value of a remote reference. Exceptions raised are same as for a `Future`.

Does not remove the item fetched.

remotecall(*f*, *id::Integer*, *args...*; *kwargs...*) → `Future`

Call a function *f* asynchronously on the given arguments on the specified process. Returns a `Future` (page 490). Keyword arguments, if any, are passed through to *f*.

remotecall_wait(*f*, *id::Integer*, *args...*; *kwargs...*)

Perform a faster wait (`remotecall(...)`) in one message on the `Worker` specified by worker id *id*. Keyword arguments, if any, are passed through to *f*.

See also `wait()` (page 490) and `remotecall()` (page 491).

remotecall_fetch(*f*, *id::Integer*, *args...*; *kwargs...*)

Perform `fetch(remotecall(...))` in one message. Keyword arguments, if any, are passed through to *f*. Any remote exceptions are captured in a `RemoteException` (page 490) and thrown.

See also `fetch()` (page 487) and `remotecall()` (page 491).

remote_do(*f*, *id::Integer*, *args...*; *kwargs...*) → nothing

Executes *f* on worker *id* asynchronously. Unlike `remotecall()` (page 491), it does not store the result of computation, nor is there a way to wait for its completion.

A successful invocation indicates that the request has been accepted for execution on the remote node.

While consecutive `remotecalls` to the same worker are serialized in the order they are invoked, the order of executions on the remote worker is undetermined. For example, `remote_do(f1, 2); remotecall(f2, 2); remote_do(f3, 2)` will serialize the call to *f1*, followed by *f2* and *f3* in that order. However, it is not guaranteed that *f1* is executed before *f3* on worker 2.

Any exceptions thrown by *f* are printed to `STDERR` (page 561) on the remote worker.

Keyword arguments, if any, are passed through to *f*.

put!(*rr::RemoteChannel*, *args...*)

Store a set of values to the `RemoteChannel` (page 490). If the channel is full, blocks until space is available. Returns its first argument.

put!(*rr::Future*, *v*)

Store a value to a `Future` (page 490) *rr*. Futures are write-once remote references. A `put!` on an already set `Future` throws an `Exception`. All asynchronous remote calls return `Futures` and set the value to the return value of the call upon completion.

take!(*rr::RemoteChannel*, *args...*)

Fetch value(s) from a `RemoteChannel` (page 490) *rr*, removing the value(s) in the process.

isready (*rr::RemoteChannel*, *args...*)

Determine whether a [RemoteChannel](#) (page 490) has a value stored to it. Note that this function can cause race conditions, since by the time you receive its result it may no longer be true. However, it can be safely used on a [Future](#) (page 490) since they are assigned only once.

isready (*rr::Future*)

Determine whether a [Future](#) (page 490) has a value stored to it.

If the argument `Future` is owned by a different node, this call will block to wait for the answer. It is recommended to wait for `rr` in a separate task instead or to use a local `Channel` as a proxy:

```
c = Channel{1}
@async put!(c, remotecall_fetch(long_computation, p))
isready(c) # will not block
```

WorkerPool (*workers::Vector{Int}*)

Create a `WorkerPool` from a vector of worker ids.

CachingPool (*workers::Vector{Int}*)

An implementation of an `AbstractWorkerPool`. [remote\(\)](#) (page 492), [remotecall_fetch\(\)](#) (page 491), [pmap\(\)](#) (page 489) (and other remote calls which execute functions remotely) benefit from caching the serialized/deserialized functions on the worker nodes, especially closures (which may capture large amounts of data).

The remote cache is maintained for the lifetime of the returned `CachingPool` object. To clear the cache earlier, use `clear!(pool)`.

For global variables, only the bindings are captured in a closure, not the data. `let` blocks can be used to capture global data.

For example:

```
const foo=rand(10^8);
wp=CachingPool(workers())
let foo=foo
    pmap(wp, i->sum(foo)+i, 1:100);
end
```

The above would transfer `foo` only once to each worker.

default_worker_pool()

`WorkerPool` containing `idle_workers()` - used by `remote(f)` and `pmap()` (page 489) (by default).

remote (*[::AbstractWorkerPool]*, *f*) → `Function`

Returns a lambda that executes function `f` on an available worker using [remotecall_fetch\(\)](#) (page 491).

remotecall (*f*, *pool::AbstractWorkerPool*, *args...*; *kwargs...*) → `Future`

`WorkerPool` variant of `remotecall(f, pid, ...)`. Waits for and takes a free worker from `pool` and performs a `remotecall` on it.

remotecall_wait (*f*, *pool::AbstractWorkerPool*, *args...*; *kwargs...*) → `Future`

`WorkerPool` variant of `remotecall_wait(f, pid, ...)`. Waits for and takes a free worker from `pool` and performs a `remotecall_wait` on it.

remotecall_fetch (*f*, *pool::AbstractWorkerPool*, *args...*; *kwargs...*) → `result`

`WorkerPool` variant of `remotecall_fetch(f, pid, ...)`. Waits for and takes a free worker from `pool` and performs a `remotecall_fetch` on it.

remote_do (*f*, *pool::AbstractWorkerPool*, *args...*; *kwargs...*) → `nothing`

`WorkerPool` variant of `remote_do(f, pid, ...)`. Waits for and takes a free worker from `pool` and performs a `remote_do` on it.

timedwait (*testcb::Function, secs::Float64; pollint::Float64=0.1*)

Waits until `testcb` returns true or for `secs` seconds, whichever is earlier. `testcb` is polled every `pollint` seconds.

@spawn ()

Creates a closure around an expression and runs it on an automatically-chosen process, returning a *Future* (page 490) to the result.

@spawnat ()

Accepts two arguments, `p` and an expression. A closure is created around the expression and run asynchronously on process `p`. Returns a *Future* (page 490) to the result.

@fetch ()

Equivalent to `fetch(@spawn expr)`. See *fetch()* (page 487) and *@spawn()* (page 493).

@fetchfrom ()

Equivalent to `fetch(@spawnat p expr)`. See *fetch()* (page 487) and *@spawnat()* (page 493).

@async ()

Like *@schedule*, *@async* wraps an expression in a *Task* and adds it to the local machine's scheduler queue. Additionally it adds the task to the set of items that the nearest enclosing *@sync* waits for. *@async* also wraps the expression in a `let x=x, y=y, ...` block to create a new scope with copies of all variables referenced in the expression.

@sync ()

Wait until all dynamically-enclosed uses of *@async*, *@spawn*, *@spawnat* and *@parallel* are complete. All exceptions thrown by enclosed async operations are collected and thrown as a *CompositeException*.

@parallel ()

A parallel for loop of the form :

```
@parallel [reducer] for var = range
    body
end
```

The specified range is partitioned and locally executed across all workers. In case an optional reducer function is specified, *@parallel* performs local reductions on each worker with a final reduction on the calling process.

Note that without a reducer function, *@parallel* executes asynchronously, i.e. it spawns independent tasks on all available workers and returns immediately without waiting for completion. To wait for completion, prefix the call with *@sync()* (page 493), like :

```
@sync @parallel for var = range
    body
end
```

@everywhere expr

Execute an expression under *Main* everywhere. Equivalent to calling `eval(Main, expr)` on all processes. Errors on any of the processes are collected into a *CompositeException* and thrown. For example :

```
@everywhere bar=1
```

will define `Main.bar` on all processes.

Unlike *@spawn()* (page 493) and *@spawnat()* (page 493), *@everywhere* does not capture any local variables. Prefixing *@everywhere* with *@eval()* (page 368) allows us to broadcast local variables using interpolation :

```
foo = 1
@eval @everywhere bar=$foo
```

The expression is evaluated under `Main` irrespective of where `@everywhere` is called from. For example :

```
module FooBar
    foo() = @everywhere bar()=myid()
end
FooBar.foo()
```

will result in `Main.bar` being defined on all processes and not `FooBar.bar`.

clear! (*pool::CachingPool*) → *pool*

Removes all cached functions from all participating workers.

Base.remoteref_id (*r::AbstractRemoteRef*) → *RRID*

Futures and RemoteChannels are identified by fields:

- *where* - refers to the node where the underlying object/storage referred to by the reference actually exists.
- *whence* - refers to the node the remote reference was created from. Note that this is different from the node where the underlying object referred to actually exists. For example calling `RemoteChannel(2)` from the master process would result in a *where* value of 2 and a *whence* value of 1.
- *id* is unique across all references created from the worker specified by *whence*.

Taken together, *whence* and *id* uniquely identify a reference across all workers.

`Base.remoteref_id` is a low-level API which returns a `Base.RRID` object that wraps *whence* and *id* values of a remote reference.

Base.channel_from_id (*id*) → *c*

A low-level API which returns the backing `AbstractChannel` for an *id* returned by `Base.remoteref_id()` (page 494). The call is valid only on the node where the backing channel exists.

Base.worker_id_from_socket (*s*) → *pid*

A low-level API which given a `IO` connection or a `Worker`, returns the *pid* of the worker it is connected to. This is useful when writing custom `serialize()` (page 564) methods for a type, which optimizes the data written out depending on the receiving process id.

Base.cluster_cookie () → *cookie*

Returns the cluster cookie.

Base.cluster_cookie (*cookie*) → *cookie*

Sets the passed cookie as the cluster cookie, then returns it.

46.3 Shared Arrays

SharedArray (*T::Type, dims::NTuple; init=false, pids=Int[]*)

Construct a `SharedArray` of a `bitstype` *T* and size *dims* across the processes specified by *pids* - all of which have to be on the same host.

If *pids* is left unspecified, the shared array will be mapped across all processes on the current host, including the master. But, `localindexes` and `indexpids` will only refer to worker processes. This facilitates work distribution code to use workers for actual computation with the master process acting as a driver.

If an *init* function of the type `initfn(S::SharedArray)` is specified, it is called on all the participating workers.

SharedArray (*filename::AbstractString, T::Type, dims::NTuple, [offset=0]; mode=nothing, init=false, pids=Int[]*)

Construct a `SharedArray` backed by the file *filename*, with element type *T* (must be a `bitstype`) and

size `dims`, across the processes specified by `pids` - all of which have to be on the same host. This file is mmapmed into the host memory, with the following consequences:

- The array data must be represented in binary format (e.g., an ASCII format like CSV cannot be supported)
- Any changes you make to the array values (e.g., `A[3] = 0`) will also change the values on disk

If `pids` is left unspecified, the shared array will be mapped across all processes on the current host, including the master. But, `localindexes` and `indexpids` will only refer to worker processes. This facilitates work distribution code to use workers for actual computation with the master process acting as a driver.

`mode` must be one of "r", "r+", "w+", or "a+", and defaults to "r+" if the file specified by `filename` already exists, or "w+" if not. If an `init` function of the type `initfn(S::SharedArray)` is specified, it is called on all the participating workers. You cannot specify an `init` function if the file is not writable.

`offset` allows you to skip the specified number of bytes at the beginning of the file.

procs (*S::SharedArray*)

Get the vector of processes mapping the shared array.

sdata (*S::SharedArray*)

Returns the actual `Array` object backing `S`.

indexpids (*S::SharedArray*)

Returns the current worker's index in the list of workers mapping the `SharedArray` (i.e. in the same list returned by `procs(S)`), or 0 if the `SharedArray` is not mapped locally.

localindexes (*S::SharedArray*)

Returns a range describing the "default" indexes to be handled by the current process. This range should be interpreted in the sense of linear indexing, i.e., as a sub-range of `1:length(S)`. In multi-process contexts, returns an empty range in the parent process (or any process for which `indexpids()` (page 495) returns 0).

It's worth emphasizing that `localindexes` exists purely as a convenience, and you can partition work on the array among workers any way you wish. For a `SharedArray`, all indexes should be equally fast for each worker process.

46.4 Multi-Threading

This experimental interface supports Julia's multi-threading capabilities. Types and function described here might (and likely will) change in the future.

Threads.threadid()

Get the ID number of the current thread of execution. The master thread has ID 1.

Threads.nthreads()

Get the number of threads available to the Julia process. This is the inclusive upper bound on `threadid()`.

Threads.@threads()

A macro to parallelize a for-loop to run with multiple threads. This spawns `nthreads()` number of threads, splits the iteration space amongst them, and iterates in parallel. A barrier is placed at the end of the loop which waits for all the threads to finish execution, and the loop returns.

Threads.Atomic{T}()

Holds a reference to an object of type `T`, ensuring that it is only accessed atomically, i.e. in a thread-safe manner.

Only certain "simple" types can be used atomically, namely the bitstype integer and float-point types. These are `Int8...Int128`, `UInt8...UInt128`, and `Float16...Float64`.

New atomic objects can be created from a non-atomic values; if none is specified, the atomic object is initialized with zero.

Atomic objects can be accessed using the `[]` notation:

```
x::Atomic{Int}
x[] = 1
val = x[]
```

Atomic operations use an `atomic_` prefix, such as `atomic_add!`, `atomic_xchg!`, etc.

`Threads.atomic_cas!{T}(x::Atomic{T}, cmp::T, newval::T)`

Atomically compare-and-set `x`

Atomically compares the value in `x` with `cmp`. If equal, write `newval` to `x`. Otherwise, leaves `x` unmodified. Returns the old value in `x`. By comparing the returned value to `cmp` (via `===`) one knows whether `x` was modified and now holds the new value `newval`.

For further details, see LLVM's `cmpxchg` instruction.

This function can be used to implement transactional semantics. Before the transaction, one records the value in `x`. After the transaction, the new value is stored only if `x` has not been modified in the mean time.

`Threads.atomic_xchg!{T}(x::Atomic{T}, newval::T)`

Atomically exchange the value in `x`

Atomically exchanges the value in `x` with `newval`. Returns the old value.

For further details, see LLVM's `atomicrmw_xchg` instruction.

`Threads.atomic_add!{T}(x::Atomic{T}, val::T)`

Atomically add `val` to `x`

Performs `x[] += val` atomically. Returns the old (!) value.

For further details, see LLVM's `atomicrmw_add` instruction.

`Threads.atomic_sub!{T}(x::Atomic{T}, val::T)`

Atomically subtract `val` from `x`

Performs `x[] -= val` atomically. Returns the old (!) value.

For further details, see LLVM's `atomicrmw_sub` instruction.

`Threads.atomic_and!{T}(x::Atomic{T}, val::T)`

Atomically bitwise-and `x` with `val`

Performs `x[] &= val` atomically. Returns the old (!) value.

For further details, see LLVM's `atomicrmw_and` instruction.

`Threads.atomic_nand!{T}(x::Atomic{T}, val::T)`

Atomically bitwise-nand (not-and) `x` with `val`

Performs `x[] = ~(x[] & val)` atomically. Returns the old (!) value.

For further details, see LLVM's `atomicrmw_nand` instruction.

`Threads.atomic_or!{T}(x::Atomic{T}, val::T)`

Atomically bitwise-or `x` with `val`

Performs `x[] |= val` atomically. Returns the old (!) value.

For further details, see LLVM's `atomicrmw_or` instruction.

`Threads.atomic_xor!{T}(x::Atomic{T}, val::T)`

Atomically bitwise-xor (exclusive-or) `x` with `val`

Performs `x[] ^= val` atomically. Returns the old (!) value.

For further details, see LLVM’s `atomicrmw xor` instruction.

`Threads.atomic_max!{T}(x::Atomic{T}, val::T)`

Atomically store the maximum of `x` and `val` in `x`

Performs `x[] = max(x[], val)` atomically. Returns the old (!) value.

For further details, see LLVM’s `atomicrmw min` instruction.

`Threads.atomic_min!{T}(x::Atomic{T}, val::T)`

Atomically store the minimum of `x` and `val` in `x`

Performs `x[] = min(x[], val)` atomically. Returns the old (!) value.

For further details, see LLVM’s `atomicrmw max` instruction.

`Threads.atomic_fence()`

Insert a sequential-consistency memory fence

Inserts a memory fence with sequentially-consistent ordering semantics. There are algorithms where this is needed, i.e. where an acquire/release ordering is insufficient.

This is likely a very expensive operation. Given that all other atomic operations in Julia already have acquire/release semantics, explicit fences should not be necessary in most cases.

For further details, see LLVM’s `fence` instruction.

46.5 ccall using a threadpool (Experimental)

`@threadcall((cfunc, clib), rettype, (argtypes...), argvals...)`

The `@threadcall` macro is called in the same way as `ccall()` (page 607) but does the work in a different thread. This is useful when you want to call a blocking C function without causing the main `julia` thread to become blocked. Concurrency is limited by size of the libuv thread pool, which defaults to 4 threads but can be increased by setting the `UV_THREADPOOL_SIZE` environment variable and restarting the `julia` process.

Note that the called function should never call back into Julia.

46.6 Synchronization Primitives

AbstractLock

Abstract supertype describing types that implement the thread-safe synchronization primitives: `lock`, `trylock`, `unlock`, and `islocked`

`lock(the_lock)`

Acquires the lock when it becomes available. If the lock is already locked by a different task/thread, it waits for it to become available.

Each `lock` must be matched by an `unlock`.

`unlock(the_lock)`

Releases ownership of the lock.

If this is a recursive lock which has been acquired before, it just decrements an internal counter and returns immediately.

`trylock(the_lock) → Success (Boolean)`

Acquires the lock if it is available, returning `true` if successful. If the lock is already locked by a different task/thread, returns `false`.

Each successful `trylock` must be matched by an `unlock`.

islocked (*the_lock*) → Status (Boolean)

Check whether the lock is held by any task/thread. This should not be used for synchronization (see instead `trylock`).

ReentrantLock ()

Creates a reentrant lock for synchronizing Tasks. The same task can acquire the lock as many times as required. Each `lock` must be matched with an `unlock`.

This lock is NOT threadsafe. See `Threads.Mutex` for a threadsafe lock.

Mutex ()

These are standard system mutexes for locking critical sections of logic.

On Windows, this is a critical section object, on pthreads, this is a `pthread_mutex_t`.

See also `SpinLock` for a lighter-weight lock.

SpinLock ()

Creates a non-reentrant lock. Recursive use will result in a deadlock. Each `lock` must be matched with an `unlock`.

Test-and-test-and-set spin locks are quickest up to about 30ish contending threads. If you have more contention than that, perhaps a lock is the wrong way to synchronize.

See also `RecursiveSpinLock` for a version that permits recursion.

See also `Mutex` for a more efficient version on one core or if the lock may be held for a considerable length of time.

RecursiveSpinLock ()

Creates a reentrant lock. The same thread can acquire the lock as many times as required. Each `lock` must be matched with an `unlock`.

See also `SpinLock` for a slightly faster version.

See also `Mutex` for a more efficient version on one core or if the lock may be held for a considerable length of time.

Semaphore (*sem_size*)

Creates a counting semaphore that allows at most `sem_size` acquires to be in use at any time. Each acquire must be matched with a release.

This construct is NOT threadsafe.

acquire (*s::Semaphore*)

Wait for one of the `sem_size` permits to be available, blocking until one can be acquired.

release (*s::Semaphore*)

Return one permit to the pool, possibly allowing another task to acquire it and resume execution.

46.7 Cluster Manager Interface

This interface provides a mechanism to launch and manage Julia workers on different cluster environments. Local-Manager, for launching additional workers on the same host and SSHManager, for launching on remote hosts via ssh are present in Base. TCP/IP sockets are used to connect and transport messages between processes. It is possible for Cluster Managers to provide a different transport.

launch (*manager::ClusterManager, params::Dict, launched::Array, launch_ntfy::Condition*)

Implemented by cluster managers. For every Julia worker launched by this function, it should append a

`WorkerConfig` entry to `launched` and notify `launch_ntfy`. The function MUST exit once all workers, requested by `manager` have been launched. `params` is a dictionary of all keyword arguments `addprocs()` (page 488) was called with.

manage (*manager::ClusterManager, id::Integer, config::WorkerConfig, op::Symbol*)

Implemented by cluster managers. It is called on the master process, during a worker's lifetime, with appropriate `op` values:

- with `:register/`:deregister when a worker is added / removed from the Julia worker pool.
- with `:interrupt` when `interrupt(workers)` is called. The `ClusterManager` should signal the appropriate worker with an interrupt signal.
- with `:finalize` for cleanup purposes.

kill (*manager::ClusterManager, pid::Int, config::WorkerConfig*)

Implemented by cluster managers. It is called on the master process, by `rmprocs()` (page 489). It should cause the remote worker specified by `pid` to exit. `Base.kill(manager::ClusterManager....)` executes a remote `exit()` on `pid`.

init_worker (*cookie::AbstractString, manager::ClusterManager=DefaultClusterManager()*)

Called by cluster managers implementing custom transports. It initializes a newly launched process as a worker. Command line argument `--worker` has the effect of initializing a process as a worker using TCP/IP sockets for transport. `cookie` is a `cluster_cookie()`.

connect (*manager::ClusterManager, pid::Int, config::WorkerConfig*) \rightarrow (*instrm::IO, outstrm::IO*)

Implemented by cluster managers using custom transports. It should establish a logical connection to worker with id `pid`, specified by `config` and return a pair of `IO` objects. Messages from `pid` to current process will be read off `instrm`, while messages to be sent to `pid` will be written to `outstrm`. The custom transport implementation must ensure that messages are delivered and received completely and in order. `Base.connect(manager::ClusterManager....)` sets up TCP/IP socket connections in-between workers.

`Base`.**process_messages** (*r_stream::IO, w_stream::IO, incoming::Bool=true*)

Called by cluster managers using custom transports. It should be called when the custom transport implementation receives the first message from a remote worker. The custom transport must manage a logical connection to the remote worker and provide two `IO` objects, one for incoming messages and the other for messages addressed to the remote worker. If `incoming` is `true`, the remote peer initiated the connection. Whichever of the pair initiates the connection sends the cluster cookie and its Julia version number to perform the authentication handshake.

See also `cluster_cookie()`.

Linear Algebra

47.1 Standard Functions

Linear algebra functions in Julia are largely implemented by calling functions from [LAPACK](#). Sparse factorizations call functions from [SuiteSparse](#).

***** (*A::AbstractMatrix, B::AbstractMatrix*)
Matrix multiplication.

Example

```
julia> [1 1; 0 1] * [1 0; 1 1]
2×2 Array{Int64,2}:
 2  1
 1  1
```

**** (*A, B*)

Matrix division using a polyalgorithm. For input matrices *A* and *B*, the result *X* is such that $A * X == B$ when *A* is square. The solver that is used depends upon the structure of *A*. If *A* is upper or lower triangular (or diagonal), no factorization of *A* is required and the system is solved with either forward or backward substitution. For non-triangular square matrices, an LU factorization is used.

For rectangular *A* the result is the minimum-norm least squares solution computed by a pivoted QR factorization of *A* and a rank estimate of *A* based on the *R* factor.

When *A* is sparse, a similar polyalgorithm is used. For indefinite matrices, the `LDLt` factorization does not use pivoting during the numerical factorization and therefore the procedure can fail even for invertible matrices.

Example

```
julia> A = [1 0; 1 -2]; B = [32; -4];

julia> X = A \ B
2-element Array{Float64,1}:
 32.0
 18.0

julia> A * X == B
true
```

dot (*x, y*)

\cdot (*x, y*)

Compute the dot product. For complex vectors, the first vector is conjugated.

Example

```
julia> dot([1; 1], [2; 3])
5

julia> dot([im; im], [1; 1])
0 - 2im
```

vecdot (*x*, *y*)

For any iterable containers *x* and *y* (including arrays of any dimension) of numbers (or any element type for which `dot` is defined), compute the Euclidean dot product (the sum of `dot(x[i], y[i])`) as if they were vectors.

cross (*x*, *y*)**×** (*x*, *y*)

Compute the cross product of two 3-vectors.

Example

```
julia> a = [0;1;0]
3-element Array{Int64,1}:
 0
 1
 0

julia> b = [0;0;1]
3-element Array{Int64,1}:
 0
 0
 1

julia> cross(a,b)
3-element Array{Int64,1}:
 1
 0
 0
```

factorize (*A*)

Compute a convenient factorization of *A*, based upon the type of the input matrix. `factorize` checks *A* to see if it is symmetric/triangular/etc. if *A* is passed as a generic matrix. `factorize` checks every element of *A* to verify/rule out each property. It will short-circuit as soon as it can rule out symmetry/triangular structure. The return value can be reused for efficient solving of multiple systems. For example: `A=factorize(A)`; `x=A\b`; `y=A\C`.

| Properties of <i>A</i> | type of factorization |
|----------------------------|---|
| Positive-definite | Cholesky (see cholfact() (page 509)) |
| Dense Symmetric/Hermitian | Bunch-Kaufman (see bkfact() (page 514)) |
| Sparse Symmetric/Hermitian | LDLt (see ldltfact() (page 511)) |
| Triangular | Triangular |
| Diagonal | Diagonal |
| Bidiagonal | Bidiagonal |
| Tridiagonal | LU (see lufact() (page 507)) |
| Symmetric real tridiagonal | LDLt (see ldltfact() (page 511)) |
| General square | LU (see lufact() (page 507)) |
| General non-square | QR (see qrfact() (page 513)) |

If `factorize` is called on a Hermitian positive-definite matrix, for instance, then `factorize` will return a Cholesky factorization.

Example

```
julia> A = Array(Bidiagonal(ones(5, 5), true))
5×5 Array{Float64,2}:
 1.0  1.0  0.0  0.0  0.0
 0.0  1.0  1.0  0.0  0.0
 0.0  0.0  1.0  1.0  0.0
 0.0  0.0  0.0  1.0  1.0
 0.0  0.0  0.0  0.0  1.0

julia> factorize(A) # factorize will check to see that A is already factorized
5×5 Bidiagonal{Float64}:
 1.0  1.0  .  .  .
 .  1.0  1.0  .  .
 .  .  1.0  1.0  .
 .  .  .  1.0  1.0
 .  .  .  .  1.0
```

This returns a `5×5 Bidiagonal{Float64}`, which can now be passed to other linear algebra functions (e.g. eigensolvers) which will use specialized methods for `Bidiagonal` types.

full (*F*)

Reconstruct the matrix *A* from the factorization `F=factorize(A)`.

Diagonal (*A::AbstractMatrix*)

Constructs a matrix from the diagonal of *A*.

Example

```
julia> A = [1 2 3; 4 5 6; 7 8 9]
3×3 Array{Int64,2}:
 1  2  3
 4  5  6
 7  8  9

julia> Diagonal(A)
3×3 Diagonal{Int64}:
 1  .  .
 .  5  .
 .  .  9
```

Diagonal (*V::AbstractVector*)

Constructs a matrix with *V* as its diagonal.

Example

```
julia> V = [1; 2]
2-element Array{Int64,1}:
 1
 2

julia> Diagonal(V)
2×2 Diagonal{Int64}:
 1  .
 .  2
```

Bidiagonal (*dv, ev, isupper::Bool*)

Constructs an upper (`isupper=true`) or lower (`isupper=false`) bidiagonal matrix using the given diagonal (*dv*) and off-diagonal (*ev*) vectors. The result is of type `Bidiagonal` and provides efficient specialized linear solvers, but may be converted into a regular matrix with `convert()` (page 364) (or `Array()` for short). *ev*'s length must be one less than the length of *dv*.

Example

```

julia> dv = [1; 2; 3; 4]
4-element Array{Int64,1}:
 1
 2
 3
 4

julia> ev = [7; 8; 9]
3-element Array{Int64,1}:
 7
 8
 9

julia> Bu = Bidiagonal(dv, ev, true) # ev is on the first superdiagonal
4×4 Bidiagonal{Int64}:
 1  7  .  .
 .  2  8  .
 .  .  3  9
 .  .  .  4

julia> Bl = Bidiagonal(dv, ev, false) # ev is on the first subdiagonal
4×4 Bidiagonal{Int64}:
 1  .  .  .
 7  2  .  .
 .  8  3  .
 .  .  9  4

```

Bidiagonal (*dv*, *ev*, *uplo::Char*)

Constructs an upper (*uplo*='U') or lower (*uplo*='L') bidiagonal matrix using the given diagonal (*dv*) and off-diagonal (*ev*) vectors. The result is of type `Bidiagonal` and provides efficient specialized linear solvers, but may be converted into a regular matrix with `convert()` (page 364) (or `Array()` for short). *ev*'s length must be one less than the length of *dv*.

Example

```

julia> dv = [1; 2; 3; 4]
4-element Array{Int64,1}:
 1
 2
 3
 4

julia> ev = [7; 8; 9]
3-element Array{Int64,1}:
 7
 8
 9

julia> Bu = Bidiagonal(dv, ev, 'U') #e is on the first superdiagonal
4×4 Bidiagonal{Int64}:
 1  7  .  .
 .  2  8  .
 .  .  3  9
 .  .  .  4

julia> Bl = Bidiagonal(dv, ev, 'L') #e is on the first subdiagonal
4×4 Bidiagonal{Int64}:
 1  .  .  .
 7  2  .  .
 .  8  3  .
 .  .  9  4

```

```

. 8 3 .
. . 9 4

```

Bidiagonal (*A*, *isupper::Bool*)

Construct a Bidiagonal matrix from the main diagonal of *A* and its first super- (if *isupper*=true) or sub-diagonal (if *isupper*=false).

Example

```

julia> A = [1 1 1 1; 2 2 2 2; 3 3 3 3; 4 4 4 4]
4×4 Array{Int64,2}:
 1  1  1  1
 2  2  2  2
 3  3  3  3
 4  4  4  4

julia> Bidiagonal(A, true) #contains the main diagonal and first superdiagonal of A
4×4 Bidiagonal{Int64}:
 1  1  .  .
 .  2  2  .
 .  .  3  3
 .  .  .  4

julia> Bidiagonal(A, false) #contains the main diagonal and first subdiagonal of A
4×4 Bidiagonal{Int64}:
 1  .  .  .
 2  2  .  .
 .  3  3  .
 .  .  4  4

```

SymTridiagonal (*dv*, *ev*)

Construct a symmetric tridiagonal matrix from the diagonal and first sub/super-diagonal, respectively. The result is of type `SymTridiagonal` and provides efficient specialized eigensolvers, but may be converted into a regular matrix with `convert()` (page 364) (or `Array()` for short).

Example

```

julia> dv = [1; 2; 3; 4]
4-element Array{Int64,1}:
 1
 2
 3
 4

julia> ev = [7; 8; 9]
3-element Array{Int64,1}:
 7
 8
 9

julia> SymTridiagonal(dv, ev)
4×4 SymTridiagonal{Int64}:
 1  7  .  .
 7  2  8  .
 .  8  3  9
 .  .  9  4

```

Tridiagonal (*dl*, *d*, *du*)

Construct a tridiagonal matrix from the first subdiagonal, diagonal, and first superdiagonal, respectively. The result is of type `Tridiagonal` and provides efficient specialized linear solvers, but may be converted into a

regular matrix with `convert()` (page 364) (or `Array{<T>}` for short). The lengths of `dl` and `du` must be one less than the length of `d`.

Example

```
julia> dl = [1; 2; 3]
3-element Array{Int64,1}:
 1
 2
 3

julia> du = [4; 5; 6]
3-element Array{Int64,1}:
 4
 5
 6

julia> d = [7; 8; 9; 0]
4-element Array{Int64,1}:
 7
 8
 9
 0

julia> Tridiagonal(dl, d, du)
4×4 Tridiagonal{Int64}:
 7  4  .  .
 1  8  5  .
 .  2  9  6
 .  .  3  0
```

Symmetric(A, uplo=:U)

Construct a `Symmetric` matrix from the upper (if `uplo = :U`) or lower (if `uplo = :L`) triangle of `A`.

Example

```
julia> A = [1 0 2 0 3; 0 4 0 5 0; 6 0 7 0 8; 0 9 0 1 0; 2 0 3 0 4]
5×5 Array{Int64,2}:
 1  0  2  0  3
 0  4  0  5  0
 6  0  7  0  8
 0  9  0  1  0
 2  0  3  0  4

julia> Supper = Symmetric(A)
5×5 Symmetric{Int64,Array{Int64,2}}:
 1  0  2  0  3
 0  4  0  5  0
 2  0  7  0  8
 0  5  0  1  0
 3  0  8  0  4

julia> Slower = Symmetric(A, :L)
5×5 Symmetric{Int64,Array{Int64,2}}:
 1  0  6  0  2
 0  4  0  9  0
 6  0  7  0  3
 0  9  0  1  0
 2  0  3  0  4
```



```
julia> eigfact(Supper)
Base.LinAlg.Eigen{Float64,Float64,Array{Float64,2},Array{Float64,1}}([-2.96684,-2.72015,0.440875
```

`eigfact` will use a method specialized for matrices known to be symmetric. Note that `Supper` will not be equal to `Slower` unless `A` is itself symmetric (e.g. if `A == A'`).

Hermitian (*A*, *uplo*=:*U*)

Construct a Hermitian matrix from the upper (if `uplo == :U`) or lower (if `uplo == :L`) triangle of *A*.

Example

```
julia> A = [1 0 2+2im 0 3-3im; 0 4 0 5 0; 6-6im 0 7 0 8+8im; 0 9 0 1 0; 2+2im 0 3-3im 0 4];

julia> Hupper = Hermitian(A)
5×5 Hermitian{Complex{Int64},Array{Complex{Int64},2}}:
 1+0im 0+0im 2+2im 0+0im 3-3im
 0+0im 4+0im 0+0im 5+0im 0+0im
 2-2im 0+0im 7+0im 0+0im 8+8im
 0+0im 5+0im 0+0im 1+0im 0+0im
 3+3im 0+0im 8-8im 0+0im 4+0im

julia> Hlower = Hermitian(A, :L)
5×5 Hermitian{Complex{Int64},Array{Complex{Int64},2}}:
 1+0im 0+0im 6+6im 0+0im 2-2im
 0+0im 4+0im 0+0im 9+0im 0+0im
 6-6im 0+0im 7+0im 0+0im 3+3im
 0+0im 9+0im 0+0im 1+0im 0+0im
 2+2im 0+0im 3-3im 0+0im 4+0im

julia> eigfact(Hupper)
Base.LinAlg.Eigen{Complex{Float64},Float64,Array{Complex{Float64},2},Array{Float64,1}}([-8.32069
```

`eigfact` will use a method specialized for matrices known to be Hermitian. Note that `Hupper` will not be equal to `Hlower` unless *A* is itself Hermitian (e.g. if `A == A'`).

lu (*A*, *pivot*=*Val{true}*) → *L*, *U*, *p*

Compute the LU factorization of *A*, such that $A[p, :] = L * U$. By default, pivoting is used. This can be overridden by passing `Val{false}` for the second argument.

See also `lu`*fact* () (page 507).

Example

```
julia> A = [4. 3.; 6. 3.]
2×2 Array{Float64,2}:
 4.0  3.0
 6.0  3.0

julia> L, U, p = lu(A)
(
 [1.0 0.0; 0.666667 1.0],

 [6.0 3.0; 0.0 1.0],

 [2,1])

julia> A[p, :] == L * U
true
```

lu*fact* (*A*[, *pivot*=*Val{true}*]) → *F*::*LU*

Compute the LU factorization of A .

In most cases, if A is a subtype S of `AbstractMatrix{T}` with an element type T supporting $+$, $-$, $*$ and $/$, the return type is `LU{T, S{T}}`. If pivoting is chosen (default) the element type should also support `abs` and `<`.

The individual components of the factorization F can be accessed by indexing:

| Component | Description |
|--------------------|---|
| <code>F[:L]</code> | L (lower triangular) part of LU |
| <code>F[:U]</code> | U (upper triangular) part of LU |
| <code>F[:p]</code> | (right) permutation <code>Vector</code> |
| <code>F[:P]</code> | (right) permutation <code>Matrix</code> |

The relationship between F and A is

`F[:L]*F[:U] == A[F[:p], :]`

F further supports the following functions:

| Supported function | LU | <code>LU{T, Tridiagonal{T}}</code> |
|-------------------------------------|----|------------------------------------|
| <code>/()</code> (page 405) | x | |
| <code>\()</code> (page 501) | x | x |
| <code>cond()</code> (page 527) | x | |
| <code>det()</code> (page 528) | x | x |
| <code>logdet()</code> (page 528) | x | x |
| <code>logabsdet()</code> (page 528) | x | x |
| <code>size()</code> (page 451) | x | x |

Example

```
julia> A = [4 3; 6 3]
2×2 Array{Int64,2}:
 4  3
 6  3

julia> F = lufact(A)
Base.LinAlg.LU{Float64,Array{Float64,2}}([4.0 3.0; 1.5 -1.5], [1,2], 0)

julia> F[:L] * F[:U] == A[F[:p], :]
true
```

lufact ($A::\text{SparseMatrixCSC}$) $\rightarrow F::\text{UmfpackLU}$

Compute the LU factorization of a sparse matrix A .

For sparse A with real or complex element type, the return type of F is `UmfpackLU{Tv, Ti}`, with $Tv = \text{Float64}$ or `Complex128` respectively and Ti is an integer type (`Int32` or `Int64`).

The individual components of the factorization F can be accessed by indexing:

| Component | Description |
|----------------------|--|
| <code>F[:L]</code> | L (lower triangular) part of LU |
| <code>F[:U]</code> | U (upper triangular) part of LU |
| <code>F[:p]</code> | right permutation <code>Vector</code> |
| <code>F[:q]</code> | left permutation <code>Vector</code> |
| <code>F[:Rs]</code> | <code>Vector</code> of scaling factors |
| <code>F[:, :]</code> | (L, U, p, q, Rs) components |

The relation between F and A is

`F[:L]*F[:U] == (F[:Rs] .* A)[F[:p], F[:q]]`

F further supports the following functions:

- `\ ()` (page 501)
- `cond ()` (page 527)
- `det ()` (page 528)

**** Implementation note ****

`lufact (A::SparseMatrixCSC)` uses the UMFPACK library that is part of SuiteSparse. As this library only supports sparse matrices with `Float64` or `Complex128` elements, `lufact` converts `A` into a copy that is of type `SparseMatrixCSC{Float64}` or `SparseMatrixCSC{Complex128}` as appropriate.

lufact! (`A`, `pivot=Val{true}`) \rightarrow LU

`lufact!` is the same as `lufact ()` (page 507), but saves space by overwriting the input `A`, instead of creating a copy. An `InexactError` (page 375) exception is thrown if the factorisation produces a number not representable by the element type of `A`, e.g. for integer types.

chol (`A`) \rightarrow U

Compute the Cholesky factorization of a positive definite matrix `A` and return the UpperTriangular matrix `U` such that $A = U'U$.

Example

```
julia> A = [1. 2.; 2. 50.]
2×2 Array{Float64,2}:
 1.0  2.0
 2.0 50.0

julia> U = chol(A)
2×2 UpperTriangular{Float64,Array{Float64,2}}:
 1.0  2.0
 .   6.78233

julia> U'U
2×2 Array{Float64,2}:
 1.0  2.0
 2.0 50.0
```

chol (`x::Number`) \rightarrow y

Compute the square root of a non-negative number `x`.

Example

```
julia> chol(16)
4.0
```

chofact (`A`, `uplo::Symbol`, `Val{false}`) \rightarrow Cholesky

Compute the Cholesky factorization of a dense symmetric positive definite matrix `A` and return a Cholesky factorization. The matrix `A` can either be a `Symmetric` or `Hermitian StridedMatrix` or a *perfectly* symmetric or `Hermitian StridedMatrix`. In the latter case, the optional argument `uplo` may be `:L` for using the lower part or `:U` for the upper part of `A`. The default is to use `:U`. The triangular Cholesky factor can be obtained from the factorization `F` with: `F[:L]` and `F[:U]`. The following functions are available for Cholesky objects: `size`, `\`, `inv`, `det`. A `PosDefException` exception is thrown in case the matrix is not positive definite.

Example

```
julia> A = [4. 12. -16.; 12. 37. -43.; -16. -43. 98.]
3×3 Array{Float64,2}:
 4.0  12.0 -16.0
```

```

12.0  37.0 -43.0
-16.0 -43.0  98.0

julia> C = cholfact(A)
Base.LinAlg.Cholesky{Float64,Array{Float64,2}} with factor:
[2.0 6.0 -8.0; 0.0 1.0 5.0; 0.0 0.0 3.0]

julia> C[:U]
3×3 UpperTriangular{Float64,Array{Float64,2}}:
 2.0  6.0 -8.0
  .   1.0  5.0
  .   .   3.0

julia> C[:L]
3×3 LowerTriangular{Float64,Array{Float64,2}}:
 2.0  .   .
 6.0  1.0  .
-8.0  5.0  3.0

julia> C[:L] * C[:U] == A
true

```

cholfact (A , $uplo::Symbol$, $Val\{true\}$; $tol = 0.0$) \rightarrow CholeskyPivoted

Compute the pivoted Cholesky factorization of a dense symmetric positive semi-definite matrix A and return a CholeskyPivoted factorization. The matrix A can either be a `Symmetric` or `Hermitian` `StridedMatrix` or a *perfectly* symmetric or `Hermitian` `StridedMatrix`. In the latter case, the optional argument `uplo` may be `:L` for using the lower part or `:U` for the upper part of A . The default is to use `:U`. The triangular Cholesky factor can be obtained from the factorization F with: $F[:L]$ and $F[:U]$. The following functions are available for `PivotedCholesky` objects: `size`, `\`, `inv`, `det`, and `rank`. The argument `tol` determines the tolerance for determining the rank. For negative values, the tolerance is the machine precision.

cholfact (A ; $shift = 0.0$, $perm = Int[]$) \rightarrow CHOLMOD.Factor

Compute the Cholesky factorization of a sparse positive definite matrix A . A must be a `SparseMatrixCSC`, `Symmetric{SparseMatrixCSC}`, or `Hermitian{SparseMatrixCSC}`. Note that even if A doesn't have the type tag, it must still be symmetric or Hermitian. A fill-reducing permutation is used. $F = \text{cholfact}(A)$ is most frequently used to solve systems of equations with $F \setminus b$, but also the methods `diag()` (page 524), `det()` (page 528), and `logdet()` (page 528) are defined for F . You can also extract individual factors from F , using $F[:L]$. However, since pivoting is on by default, the factorization is internally represented as $A == P' * L * L' * P$ with a permutation matrix P ; using just L without accounting for P will give incorrect answers. To include the effects of permutation, it's typically preferable to extract "combined" factors like $PtL = F[:PtL]$ (the equivalent of $P' * L$) and $LtP = F[:UP]$ (the equivalent of $L' * P$).

Setting the optional `shift` keyword argument computes the factorization of $A + \text{shift} * I$ instead of A . If the `perm` argument is nonempty, it should be a permutation of $1:\text{size}(A, 1)$ giving the ordering to use (instead of CHOLMOD's default AMD ordering).

Note: This method uses the CHOLMOD library from SuiteSparse, which only supports doubles or complex doubles. Input matrices not of those element types will be converted to `SparseMatrixCSC{Float64}` or `SparseMatrixCSC{Complex128}` as appropriate.

Many other functions from CHOLMOD are wrapped but not exported from the `Base.SparseArrays.CHOLMOD` module.

cholfact! ($F::Factor, A$; $shift = 0.0$) \rightarrow CHOLMOD.Factor

Compute the Cholesky (LL') factorization of A , reusing the symbolic factorization F . A must be a `SparseMatrixCSC`, `Symmetric{SparseMatrixCSC}`, or `Hermitian{SparseMatrixCSC}`.

Note that even if `A` doesn't have the type tag, it must still be symmetric or Hermitian.

See also `cholfact()` (page 509).

Note: This method uses the CHOLMOD library from SuiteSparse, which only supports doubles or complex doubles. Input matrices not of those element types will be converted to `SparseMatrixCSC{Float64}` or `SparseMatrixCSC{Complex128}` as appropriate.

cholfact! (`A`, `uplo::Symbol`], `Val{false}`) \rightarrow Cholesky

The same as `cholfact`, but saves space by overwriting the input `A`, instead of creating a copy. An `InexactError` exception is thrown if the factorisation produces a number not representable by the element type of `A`, e.g. for integer types.

Example

```
julia> A = [1 2; 2 50]
2×2 Array{Int64,2}:
 1  2
 2 50

julia> cholfact!(A)
ERROR: InexactError()
...
```

cholfact! (`A`, `uplo::Symbol`], `Val{true}; tol = 0.0`) \rightarrow CholeskyPivoted

The same as `cholfact`, but saves space by overwriting the input `A`, instead of creating a copy. An `InexactError` exception is thrown if the factorisation produces a number not representable by the element type of `A`, e.g. for integer types.

lowrankupdate (`C::Cholesky`, `v::StridedVector`) \rightarrow `CC::Cholesky`

Update a Cholesky factorization `C` with the vector `v`. If $A = C[:,U]'C[:,U]$ then $CC = \text{cholfact}(C[:,U]'C[:,U] + v*v')$ but the computation of `CC` only uses $O(n^2)$ operations.

lowrankdowndate (`C::Cholesky`, `v::StridedVector`) \rightarrow `CC::Cholesky`

Downdate a Cholesky factorization `C` with the vector `v`. If $A = C[:,U]'C[:,U]$ then $CC = \text{cholfact}(C[:,U]'C[:,U] - v*v')$ but the computation of `CC` only uses $O(n^2)$ operations.

lowrankupdate! (`C::Cholesky`, `v::StridedVector`) \rightarrow `CC::Cholesky`

Update a Cholesky factorization `C` with the vector `v`. If $A = C[:,U]'C[:,U]$ then $CC = \text{cholfact}(C[:,U]'C[:,U] + v*v')$ but the computation of `CC` only uses $O(n^2)$ operations. The input factorization `C` is updated in place such that on exit `C == CC`. The vector `v` is destroyed during the computation.

lowrankdowndate! (`C::Cholesky`, `v::StridedVector`) \rightarrow `CC::Cholesky`

Downdate a Cholesky factorization `C` with the vector `v`. If $A = C[:,U]'C[:,U]$ then $CC = \text{cholfact}(C[:,U]'C[:,U] - v*v')$ but the computation of `CC` only uses $O(n^2)$ operations. The input factorization `C` is updated in place such that on exit `C == CC`. The vector `v` is destroyed during the computation.

ldltfact (`S::SymTridiagonal`) \rightarrow `LDLt`

Compute an `LDLt` factorization of a real symmetric tridiagonal matrix such that $A = L * \text{Diagonal}(d) * L'$ where `L` is a unit lower triangular matrix and `d` is a vector. The main use of an `LDLt` factorization `F = ldltfact(A)` is to solve the linear system of equations $Ax = b$ with `F\b`.

ldltfact (`A`; `shift = 0.0`, `perm=Int[]`) \rightarrow `CHOLMOD.Factor`

Compute the LDL' factorization of a sparse matrix `A`. `A` must be a `SparseMatrixCSC`, `Symmetric{SparseMatrixCSC}`, or `Hermitian{SparseMatrixCSC}`. Note that even if `A` doesn't have the type tag, it must still be symmetric or Hermitian. A fill-reducing permutation is used. `F = ldltfact(A)` is most frequently used to solve systems of equations $A*x = b$ with `F\b`. The returned factorization object `F` also supports the methods `diag()` (page 524), `det()` (page 528), and `logdet()`

(page 528). You can extract individual factors from F using $F[:,L]$. However, since pivoting is on by default, the factorization is internally represented as $A == P' * L * D * L' * P$ with a permutation matrix P ; using just L without accounting for P will give incorrect answers. To include the effects of permutation, it is typically preferable to extract “combined” factors like $PtL = F[:,PtL]$ (the equivalent of $P' * L$) and $LtP = F[:,UP]$ (the equivalent of $L' * P$). The complete list of supported factors is `:L`, `:PtL`, `:D`, `:UP`, `:U`, `:LD`, `:DU`, `:PtLD`, `:DUP`.

Setting the optional `shift` keyword argument computes the factorization of $A + \text{shift} * I$ instead of A . If the `perm` argument is nonempty, it should be a permutation of $1:\text{size}(A,1)$ giving the ordering to use (instead of CHOLMOD’s default AMD ordering).

Note: This method uses the CHOLMOD library from SuiteSparse, which only supports doubles or complex doubles. Input matrices not of those element types will be converted to `SparseMatrixCSC{Float64}` or `SparseMatrixCSC{Complex128}` as appropriate.

Many other functions from CHOLMOD are wrapped but not exported from the `Base.SparseArrays.CHOLMOD` module.

ldltfact! ($F::Factor, A; \text{shift} = 0.0$) \rightarrow CHOLMOD.Factor

Compute the LDL' factorization of A , reusing the symbolic factorization F . A must be a `SparseMatrixCSC`, `Symmetric{SparseMatrixCSC}`, or `Hermitian{SparseMatrixCSC}`. Note that even if A doesn’t have the type tag, it must still be symmetric or Hermitian.

See also `ldltfact()` (page 511).

Note: This method uses the CHOLMOD library from SuiteSparse, which only supports doubles or complex doubles. Input matrices not of those element types will be converted to `SparseMatrixCSC{Float64}` or `SparseMatrixCSC{Complex128}` as appropriate.

ldltfact! ($S::\text{SymTridiagonal}$) \rightarrow LDLt

Same as `ldltfact()` (page 511), but saves space by overwriting the input A , instead of creating a copy.

qr ($v::\text{AbstractVector}$) $\rightarrow w, r$

Computes the polar decomposition of a vector. Returns w , a unit vector in the direction of v , and r , the norm of v .

See also `normalize()` (page 527), `normalize!()` (page 527), and `LinAlg.qr!()` (page 512).

Example

```
julia> v = [1; 2]
2-element Array{Int64,1}:
 1
 2

julia> w, r = qr(v)
([0.447214, 0.894427], 2.23606797749979)

julia> w*r == v
true
```

`LinAlg.qr!` ($v::\text{AbstractVector}$) $\rightarrow w, r$

Computes the polar decomposition of a vector. Instead of returning a new vector as `qr(v::AbstractVector)`, this function mutates the input vector v in place. Returns w , a unit vector in the direction of v (this is a mutation of v), and r , the norm of v .

See also `normalize()` (page 527), `normalize!()` (page 527), and `qr()` (page 512).

qr (*A*, *pivot*=*Val{false}*; *thin*::*Bool*=*true*) → *Q*, *R*, [*p*]

Compute the (pivoted) QR factorization of *A* such that either $A = Q \cdot R$ or $A[:,p] = Q \cdot R$. Also see [qrfact\(\)](#) (page 513). The default is to compute a thin factorization. Note that *R* is not extended with zeros when the full *Q* is requested.

qrfact (*A*, *pivot*=*Val{false}*) → *F*

Computes the QR factorization of *A*. The return type of *F* depends on the element type of *A* and whether pivoting is specified (with *pivot*=*Val{true}*).

| Return type | eltype(<i>A</i>) | pivot | Relationship between <i>F</i> and <i>A</i> |
|-------------|--------------------|-------------------|--|
| QR | not BlasFloat | either | $A == F[:,Q] * F[:,R]$ |
| QRCompactWY | BlasFloat | <i>Val{false}</i> | $A == F[:,Q] * F[:,R]$ |
| QRPivoted | BlasFloat | <i>Val{true}</i> | $A[:,F[:,p]] == F[:,Q] * F[:,R]$ |

BlasFloat refers to any of: Float32, Float64, Complex64 or Complex128.

The individual components of the factorization *F* can be accessed by indexing:

| Component | Description | QR | QRCompactWY | QRPivoted |
|---------------|--|-------------------------------|----------------------------------|-------------------------------|
| <i>F[:,Q]</i> | <i>Q</i> (orthogonal/unitary) part of QR | x
(QR _{PackedQ}) | x
(QR _{CompactWYQ}) | x
(QR _{PackedQ}) |
| <i>F[:,R]</i> | <i>R</i> (upper right triangular) part of QR | x | x | x |
| <i>F[:,p]</i> | pivot Vector | | | x |
| <i>F[:,P]</i> | (pivot) permutation Matrix | | | x |

The following functions are available for the QR objects: [size\(\)](#) (page 451) and [\\(\)](#) (page 501). When *A* is rectangular, [\\(\)](#) will return a least squares solution and if the solution is not unique, the one with smallest norm is returned.

Multiplication with respect to either thin or full *Q* is allowed, i.e. both $F[:,Q] * F[:,R]$ and $F[:,Q] * A$ are supported. A *Q* matrix can be converted into a regular matrix with [full\(\)](#) (page 503) which has a named argument *thin*.

Note: `qrfact` returns multiple types because LAPACK uses several representations that minimize the memory storage requirements of products of Householder elementary reflectors, so that the *Q* and *R* matrices can be stored compactly rather as two separate dense matrices.

The data contained in QR or QR_{Pivoted} can be used to construct the QR_{PackedQ} type, which is a compact representation of the rotation matrix:

$$Q = \prod_{i=1}^{\min(m,n)} (I - \tau_i v_i v_i^T)$$

where τ_i is the scale factor and v_i is the projection vector associated with the i^{th} Householder elementary reflector.

The data contained in QR_{CompactWY} can be used to construct the QR_{CompactWYQ} type, which is a compact representation of the rotation matrix

$$Q = I + YTY^T$$

where *Y* is $m \times r$ lower trapezoidal and *T* is $r \times r$ upper triangular. The *compact WY* representation [\[Schreiber1989\]](#) (page 621) is not to be confused with the older, *WY* representation [\[Bischof1987\]](#) (page 621). (The LAPACK documentation uses *V* in lieu of *Y*.)

qrfact (*A*) → SPQR.Factorization

Compute the QR factorization of a sparse matrix *A*. A fill-reducing permutation is used. The main application of this type is to solve least squares problems with `\`. The function calls the C library SPQR and a few additional functions from the library are wrapped but not exported.

qrfact! (*A*, *pivot*=*Val{false}*)

`qrfact!` is the same as `qrfact()` (page 513) when *A* is a subtype of `StridedMatrix`, but saves space by overwriting the input *A*, instead of creating a copy. An `InexactError` (page 375) exception is thrown if the factorisation produces a number not representable by the element type of *A*, e.g. for integer types.

full (*A*::*Union{QRPackedQ, QRCompactWYQ}*; *thin*::*Bool*=*true*) → *Matrix*

Converts an orthogonal or unitary matrix stored as a `QRCompactWYQ` object, i.e. in the compact WY format [Bischof1987] (page 621), or in the `QRPackedQ` format, to a dense matrix.

Optionally takes a `thin` Boolean argument, which if `true` omits the columns that span the rows of *R* in the QR factorization that are zero. The resulting matrix is the *Q* in a thin QR factorization (sometimes called the reduced QR factorization). If `false`, returns a *Q* that spans all rows of *R* in its corresponding QR factorization.

lqfact! (*A*) → *LQ*

Compute the LQ factorization of *A*, using the input matrix as a workspace. See also `lq()` (page 514).

lqfact (*A*) → *LQ*

Compute the LQ factorization of *A*. See also `lq()` (page 514).

lq (*A*; [*thin*=*true*]) → *L*, *Q*

Perform an LQ factorization of *A* such that $A = L \cdot Q$. The default is to compute a thin factorization. The LQ factorization is the QR factorization of A' . *L* is not extended with zeros if the full *Q* is requested.

bkfact (*A*, *uplo*::*Symbol*=*U*, *symmetric*::*Bool*=*issymmetric(A)*, *rook*::*Bool*=*false*) → *BunchKaufman*

Compute the Bunch-Kaufman [Bunch1977] (page 621) factorization of a symmetric or Hermitian matrix *A* and return a *BunchKaufman* object. *uplo* indicates which triangle of matrix *A* to reference. If *symmetric* is `true`, *A* is assumed to be symmetric. If *symmetric* is `false`, *A* is assumed to be Hermitian. If *rook* is `true`, rook pivoting is used. If *rook* is `false`, rook pivoting is not used. The following functions are available for *BunchKaufman* objects: `size()` (page 451), `\`, `inv()` (page 528), `issymmetric()` (page 532), `ishermitian()` (page 534).

bkfact! (*A*, *uplo*::*Symbol*=*U*, *symmetric*::*Bool*=*issymmetric(A)*, *rook*::*Bool*=*false*) → *BunchKaufman*

`bkfact!` is the same as `bkfact()` (page 514), but saves space by overwriting the input *A*, instead of creating a copy.

eig (*A*, [*irange*], [*vl*], [*vu*], [*permute*=*true*], [*scale*=*true*]) → *D*, *V*

Computes eigenvalues (*D*) and eigenvectors (*V*) of *A*. See `eigfact()` (page 516) for details on the *irange*, *vl*, and *vu* arguments and the *permute* and *scale* keyword arguments. The eigenvectors are returned columnwise.

Example

```
julia> eig([1.0 0.0 0.0; 0.0 3.0 0.0; 0.0 0.0 18.0])
([1.0, 3.0, 18.0],
 [1.0 0.0 0.0; 0.0 1.0 0.0; 0.0 0.0 1.0])
```

`eig` is a wrapper around `eigfact()` (page 516), extracting all parts of the factorization to a tuple; where possible, using `eigfact()` (page 516) is recommended.

eig (*A*, *B*) → *D*, *V*

Computes generalized eigenvalues (*D*) and vectors (*V*) of *A* with respect to *B*.

`eig` is a wrapper around `eigfact()` (page 516), extracting all parts of the factorization to a tuple; where possible, using `eigfact()` (page 516) is recommended.

Example


```
julia> A = [1 0; 0 -1]
2×2 Array{Int64,2}:
 1  0
 0 -1

julia> B = [0 1; 1 0]
2×2 Array{Int64,2}:
 0  1
 1  0

julia> eig(A, B)
(Complex{Float64}[0.0+1.0im,0.0-1.0im],
 Complex{Float64}[0.0-1.0im 0.0+1.0im; -1.0-0.0im -1.0+0.0im])
```

eigvals (A,[irange],[vl],[vu]) → values

Returns the eigenvalues of A. If A is Symmetric, Hermitian or SymTridiagonal, it is possible to calculate only a subset of the eigenvalues by specifying either a UnitRange *irange* covering indices of the sorted eigenvalues, or a pair *vl* and *vu* for the lower and upper boundaries of the eigenvalues.

For general non-symmetric matrices it is possible to specify how the matrix is balanced before the eigenvector calculation. The option `permute=true` permutes the matrix to become closer to upper triangular, and `scale=true` scales the matrix by its diagonal elements to make rows and columns more equal in norm. The default is `true` for both options.

eigvals! (A,[irange],[vl],[vu]) → values

Same as `eigvals()` (page 515), but saves space by overwriting the input A, instead of creating a copy.

eigmax (A; permute::Bool=true, scale::Bool=true)

Returns the largest eigenvalue of A. The option `permute=true` permutes the matrix to become closer to upper triangular, and `scale=true` scales the matrix by its diagonal elements to make rows and columns more equal in norm. Note that if the eigenvalues of A are complex, this method will fail, since complex numbers cannot be sorted.

Example

```
julia> A = [0 im; -im 0]
2×2 Array{Complex{Int64},2}:
 0+0im  0+1im
 0-1im  0+0im

julia> eigmax(A)
1.0

julia> A = [0 im; -1 0]
2×2 Array{Complex{Int64},2}:
 0+0im  0+1im
 -1+0im  0+0im

julia> eigmax(A)
ERROR: DomainError:
in #eigmax#30(::Bool, ::Bool, ::Function, ::Array{Complex{Int64},2}) at ./linalg/eigen.jl:219
in eigmax(::Array{Complex{Int64},2}) at ./linalg/eigen.jl:217
...
```

eigmin (A; permute::Bool=true, scale::Bool=true)

Returns the smallest eigenvalue of A. The option `permute=true` permutes the matrix to become closer to upper triangular, and `scale=true` scales the matrix by its diagonal elements to make rows and columns more equal in norm. Note that if the eigenvalues of A are complex, this method will fail, since complex numbers cannot be sorted.

Example

```
julia> A = [0 im; -im 0]
2×2 Array{Complex{Int64},2}:
 0+0im  0+1im
 0-1im  0+0im

julia> eigmin(A)
-1.0

julia> A = [0 im; -1 0]
2×2 Array{Complex{Int64},2}:
 0+0im  0+1im
-1+0im  0+0im

julia> eigmin(A)
ERROR: DomainError:
 in #eigmin#31(::Bool, ::Bool, ::Function, ::Array{Complex{Int64},2}) at ./linalg/eigen.jl:261
 in eigmin(::Array{Complex{Int64},2}) at ./linalg/eigen.jl:259
...
```

eigvecs (*A*, [*eigvals*], [*permute=true*], [*scale=true*]) → Matrix

Returns a matrix *M* whose columns are the eigenvectors of *A*. (The *k*th eigenvector can be obtained from the slice *M*[:, *k*].) The *permute* and *scale* keywords are the same as for *eigfact*() (page 516).

For *SymTridiagonal* (page 505) matrices, if the optional vector of eigenvalues *eigvals* is specified, returns the specific corresponding eigenvectors.

Example

```
julia> eigvecs([1.0 0.0 0.0; 0.0 3.0 0.0; 0.0 0.0 18.0])
3×3 Array{Float64,2}:
 1.0  0.0  0.0
 0.0  1.0  0.0
 0.0  0.0  1.0
```

eigfact (*A*, [*irange*], [*vl*], [*vu*], [*permute=true*], [*scale=true*]) → Eigen

Computes the eigenvalue decomposition of *A*, returning an Eigen factorization object *F* which contains the eigenvalues in *F*[:values] and the eigenvectors in the columns of the matrix *F*[:vectors]. (The *k*th eigenvector can be obtained from the slice *F*[:vectors][:, *k*].)

The following functions are available for Eigen objects: *inv*() (page 528), *det*() (page 528), and *isposdef*() (page 532).

If *A* is *Symmetric* (page 506), *Hermitian* (page 507) or *SymTridiagonal* (page 505), it is possible to calculate only a subset of the eigenvalues by specifying either a UnitRange *irange* covering indices of the sorted eigenvalues or a pair *vl* and *vu* for the lower and upper boundaries of the eigenvalues.

For general nonsymmetric matrices it is possible to specify how the matrix is balanced before the eigenvector calculation. The option *permute=true* permutes the matrix to become closer to upper triangular, and *scale=true* scales the matrix by its diagonal elements to make rows and columns more equal in norm. The default is true for both options.

Example

```
julia> F = eigfact([1.0 0.0 0.0; 0.0 3.0 0.0; 0.0 0.0 18.0])
Base.LinAlg.Eigen{Float64,Float64,Array{Float64,2},Array{Float64,1}}([1.0, 3.0, 18.0], [1.0 0.0 0.0
julia> F[:values]
3-element Array{Float64,1}:
 1.0
```

```

3.0
18.0

julia> F[:vectors]
3×3 Array{Float64,2}:
 1.0  0.0  0.0
 0.0  1.0  0.0
 0.0  0.0  1.0

```

eigfact (*A*, *B*) → GeneralizedEigen

Computes the generalized eigenvalue decomposition of *A* and *B*, returning a GeneralizedEigen factorization object *F* which contains the generalized eigenvalues in *F*[:values] and the generalized eigenvectors in the columns of the matrix *F*[:vectors]. (The *k*th generalized eigenvector can be obtained from the slice *F*[:vectors][:, *k*].)

eigfact! (*A*[, *B*])

Same as *eigfact*() (page 516), but saves space by overwriting the input *A* (and *B*), instead of creating a copy.

hessfact (*A*) → Hessenberg

Compute the Hessenberg decomposition of *A* and return a Hessenberg object. If *F* is the factorization object, the unitary matrix can be accessed with *F*[:Q] and the Hessenberg matrix with *F*[:H]. When *Q* is extracted, the resulting type is the HessenbergQ object, and may be converted to a regular matrix with *convert*() (page 364) (or *Array*(_) for short).

Example

```

julia> A = [4. 9. 7.; 4. 4. 1.; 4. 3. 2.]
3×3 Array{Float64,2}:
 4.0  9.0  7.0
 4.0  4.0  1.0
 4.0  3.0  2.0

julia> F = hessfact(A);

julia> F[:Q] * F[:H] * F[:Q]'
3×3 Array{Float64,2}:
 4.0  9.0  7.0
 4.0  4.0  1.0
 4.0  3.0  2.0

```

hessfact! (*A*) → Hessenberg

hessfact! is the same as *hessfact*() (page 517), but saves space by overwriting the input *A*, instead of creating a copy.

schurfact (*A*::StridedMatrix) → *F*::Schur

Computes the Schur factorization of the matrix *A*. The (quasi) triangular Schur factor can be obtained from the Schur object *F* with either *F*[:Schur] or *F*[:T] and the orthogonal/unitary Schur vectors can be obtained with *F*[:vectors] or *F*[:Z] such that $A = F[:vectors] * F[:Schur] * F[:vectors]'$. The eigenvalues of *A* can be obtained with *F*[:values].

Example

```

julia> A = [-2. 1. 3.; 2. 1. -1.; -7. 2. 7.]
3×3 Array{Float64,2}:
 -2.0  1.0  3.0
  2.0  1.0 -1.0
 -7.0  2.0  7.0

julia> F = schurfact(A)

```

```
Base.LinAlg.Schur{Float64,Array{Float64,2}}([2.0 0.801792 6.63509; -8.55988e-11 2.0 8.08286; 0.0
julia> F[:vectors] * F[:Schur] * F[:vectors]'
3×3 Array{Float64,2}:
-2.0  1.0  3.0
 2.0  1.0 -1.0
-7.0  2.0  7.0
```

schurfact! (*A::StridedMatrix*) → *F::Schur*

Same as `schurfact` but uses the input argument as workspace.

schur (*A::StridedMatrix*) → *T::Matrix*, *Z::Matrix*, *λ::Vector*

Computes the Schur factorization of the matrix *A*. The methods return the (quasi) triangular Schur factor *T* and the orthogonal/unitary Schur vectors *Z* such that $A = Z * T * Z'$. The eigenvalues of *A* are returned in the vector *λ*.

See `schurfact`.

Example

```
julia> A = [-2. 1. 3.; 2. 1. -1.; -7. 2. 7.]
3×3 Array{Float64,2}:
-2.0  1.0  3.0
 2.0  1.0 -1.0
-7.0  2.0  7.0

julia> T, Z, lambda = schur(A)
(
 [2.0 0.801792 6.63509; -8.55988e-11 2.0 8.08286; 0.0 0.0 1.99999],
 [0.577351 0.154299 -0.801784; 0.577346 -0.77152 0.267262; 0.577354 0.617211 0.534522],
 Complex{Float64}[2.0+8.28447e-6im,2.0-8.28447e-6im,1.99999+0.0im])

julia> Z * T * Z'
3×3 Array{Float64,2}:
-2.0  1.0  3.0
 2.0  1.0 -1.0
-7.0  2.0  7.0
```

ordschur (*F::Schur*, *select::Union{Vector{Bool}, BitVector}*) → *F::Schur*

Reorders the Schur factorization *F* of a matrix $A = Z * T * Z'$ according to the logical array *select* returning the reordered factorization *F* object. The selected eigenvalues appear in the leading diagonal of *F[:Schur]* and the corresponding leading columns of *F[:vectors]* form an orthogonal/unitary basis of the corresponding right invariant subspace. In the real case, a complex conjugate pair of eigenvalues must be either both included or both excluded via *select*.

ordschur! (*F::Schur*, *select::Union{Vector{Bool}, BitVector}*) → *F::Schur*

Same as `ordschur` but overwrites the factorization *F*.

ordschur (*T::StridedMatrix*, *Z::StridedMatrix*, *select::Union{Vector{Bool}, BitVector}*) →

T::StridedMatrix, *Z::StridedMatrix*, *λ::Vector*

Reorders the Schur factorization of a real matrix $A = Z * T * Z'$ according to the logical array *select* returning the reordered matrices *T* and *Z* as well as the vector of eigenvalues *λ*. The selected eigenvalues appear in the leading diagonal of *T* and the corresponding leading columns of *Z* form an orthogonal/unitary basis of the corresponding right invariant subspace. In the real case, a complex conjugate pair of eigenvalues must be either both included or both excluded via *select*.

ordschur! (*T::StridedMatrix*, *Z::StridedMatrix*, *select::Union{Vector{Bool}, BitVector}*) →

T::StridedMatrix, *Z::StridedMatrix*, *λ::Vector*

Same as `ordschur` but overwrites the input arguments.

schurfact (*A::StridedMatrix*, *B::StridedMatrix*) → *F::GeneralizedSchur*

Computes the Generalized Schur (or QZ) factorization of the matrices *A* and *B*. The (quasi) triangular Schur factors can be obtained from the *Schur* object *F* with *F[:S]* and *F[:T]*, the left unitary/orthogonal Schur vectors can be obtained with *F[:left]* or *F[:Q]* and the right unitary/orthogonal Schur vectors can be obtained with *F[:right]* or *F[:Z]* such that $A = F[:left] * F[:S] * F[:right]'$ and $B = F[:left] * F[:T] * F[:right]'$. The generalized eigenvalues of *A* and *B* can be obtained with *F[:alpha] ./ F[:beta]*.

schurfact! (*A::StridedMatrix*, *B::StridedMatrix*) → *F::GeneralizedSchur*

Same as `schurfact` but uses the input matrices *A* and *B* as workspace.

ordschur (*F::GeneralizedSchur*, *select::Union{Vector{Bool}, BitVector}*) → *F::GeneralizedSchur*

Reorders the Generalized Schur factorization *F* of a matrix pair $(A, B) = (Q * S * Z', Q * T * Z')$ according to the logical array *select* and returns a *GeneralizedSchur* object *F*. The selected eigenvalues appear in the leading diagonal of both *F[:S]* and *F[:T]*, and the left and right orthogonal/unitary Schur vectors are also reordered such that $(A, B) = F[:Q] * (F[:S], F[:T]) * F[:Z]'$ still holds and the generalized eigenvalues of *A* and *B* can still be obtained with *F[:alpha] ./ F[:beta]*.

ordschur! (*F::GeneralizedSchur*, *select::Union{Vector{Bool}, BitVector}*) → *F::GeneralizedSchur*

Same as `ordschur` but overwrites the factorization *F*.

ordschur (*S::StridedMatrix*, *T::StridedMatrix*, *Q::StridedMatrix*, *Z::StridedMatrix*, *select*) →

S::StridedMatrix, *T::StridedMatrix*, *Q::StridedMatrix*, *Z::StridedMatrix*, $\alpha::Vector$, $\beta::Vector$
Reorders the Generalized Schur factorization of a matrix pair $(A, B) = (Q * S * Z', Q * T * Z')$ according to the logical array *select* and returns the matrices *S*, *T*, *Q*, *Z* and vectors α and β . The selected eigenvalues appear in the leading diagonal of both *S* and *T*, and the left and right unitary/orthogonal Schur vectors are also reordered such that $(A, B) = Q * (S, T) * Z'$ still holds and the generalized eigenvalues of *A* and *B* can still be obtained with $\alpha ./ \beta$.

ordschur! (*S::StridedMatrix*, *T::StridedMatrix*, *Q::StridedMatrix*, *Z::StridedMatrix*, *select*) →

S::StridedMatrix, *T::StridedMatrix*, *Q::StridedMatrix*, *Z::StridedMatrix*, $\alpha::Vector$, $\beta::Vector$
Same as `ordschur` but overwrites the factorization the input arguments.

schur (*A::StridedMatrix*, *B::StridedMatrix*) → *S::StridedMatrix*, *T::StridedMatrix*, *Q::StridedMatrix*,

Z::StridedMatrix, $\alpha::Vector$, $\beta::Vector$
See `schurfact`.

svdfact (*A*[, *thin=true*]) → SVD

Compute the singular value decomposition (SVD) of *A* and return an SVD object.

U, *S*, *V* and *Vt* can be obtained from the factorization *F* with *F[:U]*, *F[:S]*, *F[:V]* and *F[:Vt]*, such that $A = U * \text{diagm}(S) * Vt$. The algorithm produces *Vt* and hence *Vt* is more efficient to extract than *V*.

If *thin=true* (default), a thin SVD is returned. For a $M \times N$ matrix *A*, *U* is $M \times M$ for a full SVD (*thin=false*) and $M \times \min(M, N)$ for a thin SVD.

Example

```
julia> A = [1. 0. 0. 0. 2.; 0. 0. 3. 0. 0.; 0. 0. 0. 0. 0.; 0. 2. 0. 0. 0.]
```

```
4×5 Array{Float64,2}:
```

```
 1.0  0.0  0.0  0.0  2.0
 0.0  0.0  3.0  0.0  0.0
 0.0  0.0  0.0  0.0  0.0
 0.0  2.0  0.0  0.0  0.0
```

```
julia> F = svdfact(A)
```

```
Base.LinAlg.SVD{Float64,Float64,Array{Float64,2}}([0.0 1.0 0.0 0.0; 1.0 0.0 0.0 0.0; 0.0 0.0 0.0 0.0; 0.0 0.0 0.0 0.0])
```

```
julia> F[:U] * diagm(F[:S]) * F[:Vt]
```

```
4×5 Array{Float64,2}:
 1.0  0.0  0.0  0.0  2.0
 0.0  0.0  3.0  0.0  0.0
 0.0  0.0  0.0  0.0  0.0
 0.0  2.0  0.0  0.0  0.0
```

svdfact! (A , $thin=true$) \rightarrow SVD

svdfact! is the same as `svdfact()` (page 519), but saves space by overwriting the input A , instead of creating a copy.

If $thin=true$ (default), a thin SVD is returned. For a $M \times N$ matrix A , U is $M \times M$ for a full SVD ($thin=false$) and $M \times \min(M, N)$ for a thin SVD.

svd (A , $thin=true$) $\rightarrow U, S, V$

Computes the SVD of A , returning U , vector S , and V such that $A == U \cdot \text{diag}(S) \cdot V'$.

If $thin=true$ (default), a thin SVD is returned. For a $M \times N$ matrix A , U is $M \times M$ for a full SVD ($thin=false$) and $M \times \min(M, N)$ for a thin SVD.

svd is a wrapper around `svdfact(A)`, extracting all parts of the SVD factorization to a tuple. Direct use of `svdfact` is therefore more efficient.

Example

```
julia> A = [1. 0. 0. 0. 2.; 0. 0. 3. 0. 0.; 0. 0. 0. 0. 0.; 0. 2. 0. 0. 0.]
4×5 Array{Float64,2}:
 1.0  0.0  0.0  0.0  2.0
 0.0  0.0  3.0  0.0  0.0
 0.0  0.0  0.0  0.0  0.0
 0.0  2.0  0.0  0.0  0.0

julia> U, S, V = svd(A)
(
 [0.0 1.0 0.0 0.0; 1.0 0.0 0.0 0.0; 0.0 0.0 0.0 -1.0; 0.0 0.0 1.0 0.0],

 [3.0, 2.23607, 2.0, 0.0],

 [-0.0 0.447214 -0.0 0.0; 0.0 0.0 1.0 0.0; ... ; -0.0 0.0 -0.0 1.0; 0.0 0.894427 0.0 0.0])

julia> U*diag(S)*V'
4×5 Array{Float64,2}:
 1.0  0.0  0.0  0.0  2.0
 0.0  0.0  3.0  0.0  0.0
 0.0  0.0  0.0  0.0  0.0
 0.0  2.0  0.0  0.0  0.0
```

svdvals (A)

Returns the singular values of A .

Example

```
julia> A = [1. 0. 0. 0. 2.; 0. 0. 3. 0. 0.; 0. 0. 0. 0. 0.; 0. 2. 0. 0. 0.]
4×5 Array{Float64,2}:
 1.0  0.0  0.0  0.0  2.0
 0.0  0.0  3.0  0.0  0.0
 0.0  0.0  0.0  0.0  0.0
 0.0  2.0  0.0  0.0  0.0

julia> svdvals(A)
4-element Array{Float64,1}:
 3.0
 2.23607
```

```
2.0
0.0
```

svdvals! (*A*)

Returns the singular values of *A*, saving space by overwriting the input.

svdfact (*A*, *B*) → GeneralizedSVD

Compute the generalized SVD of *A* and *B*, returning a GeneralizedSVD factorization object *F*, such that $A = F[:,U] * F[:,D1] * F[:,R0] * F[:,Q]'$ and $B = F[:,V] * F[:,D2] * F[:,R0] * F[:,Q]'$.

For an *M*-by-*N* matrix *A* and *P*-by-*N* matrix *B*,

- $F[:,U]$ is a *M*-by-*M* orthogonal matrix,
- $F[:,V]$ is a *P*-by-*P* orthogonal matrix,
- $F[:,Q]$ is a *N*-by-*N* orthogonal matrix,
- $F[:,R0]$ is a (*K*+*L*)-by-*N* matrix whose rightmost (*K*+*L*)-by-(*K*+*L*) block is nonsingular upper block triangular,
- $F[:,D1]$ is a *M*-by-(*K*+*L*) diagonal matrix with 1s in the first *K* entries,
- $F[:,D2]$ is a *P*-by-(*K*+*L*) matrix whose top right *L*-by-*L* block is diagonal,

K+*L* is the effective numerical rank of the matrix $[A; B]$.

The entries of $F[:,D1]$ and $F[:,D2]$ are related, as explained in the LAPACK documentation for the [generalized SVD](#) and the [xGGSVD3](#) routine which is called underneath (in LAPACK 3.6.0 and newer).

svd (*A*, *B*) → *U*, *V*, *Q*, *D1*, *D2*, *R0*

Wrapper around `svdfact(A, B)()` extracting all parts of the factorization to a tuple. Direct use of `svdfact` is therefore generally more efficient. The function returns the generalized SVD of *A* and *B*, returning *U*, *V*, *Q*, *D1*, *D2*, and *R0* such that $A = U * D1 * R0 * Q'$ and $B = V * D2 * R0 * Q'$.

svdvals (*A*, *B*)

Return the generalized singular values from the generalized singular value decomposition of *A* and *B*.

LinAlg.Givens (*i1*, *i2*, *c*, *s*) → *G*

A Givens rotation linear operator. The fields *c* and *s* represent the cosine and sine of the rotation angle, respectively. The `Givens` type supports left multiplication $G * A$ and conjugated transpose right multiplication $A * G'$. The type doesn't have a `size` and can therefore be multiplied with matrices of arbitrary size as long as $i2 \leq \text{size}(A, 2)$ for $G * A$ or $i2 \leq \text{size}(A, 1)$ for $A * G'$.

See also: [givens\(\)](#) (page 521)

givens{T} (*f::T*, *g::T*, *i1::Integer*, *i2::Integer*) → (*G::Givens*, *r::T*)

Computes the Givens rotation *G* and scalar *r* such that for any vector *x* where

```
x[i1] = f
x[i2] = g
```

the result of the multiplication

```
y = G * x
```

has the property that

```
y[i1] = r
y[i2] = 0
```

See also: [LinAlg.Givens](#) (page 521)

givens ($x::AbstractVector$, $i1::Integer$, $i2::Integer$) \rightarrow ($G::Givens$, r)

Computes the Givens rotation G and scalar r such that the result of the multiplication

$$B = G * x$$

has the property that

$$\begin{aligned} B[i1] &= r \\ B[i2] &= 0 \end{aligned}$$

See also: [LinAlg.Givens](#) (page 521)

givens ($A::AbstractArray$, $i1::Integer$, $i2::Integer$, $j::Integer$) \rightarrow ($G::Givens$, r)

Computes the Givens rotation G and scalar r such that the result of the multiplication

$$B = G * A$$

has the property that

$$\begin{aligned} B[i1, j] &= r \\ B[i2, j] &= 0 \end{aligned}$$

See also: [LinAlg.Givens](#) (page 521)

triu (M)

Upper triangle of a matrix.

Example

```
julia> a = ones(4,4)
4×4 Array{Float64,2}:
 1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0

julia> triu(a)
4×4 Array{Float64,2}:
 1.0  1.0  1.0  1.0
 0.0  1.0  1.0  1.0
 0.0  0.0  1.0  1.0
 0.0  0.0  0.0  1.0
```

triu (M , $k::Integer$)

Returns the upper triangle of M starting from the k th superdiagonal.

Example

```
julia> a = ones(4,4)
4×4 Array{Float64,2}:
 1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0

julia> triu(a,3)
4×4 Array{Float64,2}:
 0.0  0.0  0.0  1.0
 0.0  0.0  0.0  0.0
 0.0  0.0  0.0  0.0
 0.0  0.0  0.0  0.0
```



```
julia> triu(a,-3)
4×4 Array{Float64,2}:
 1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0
```

triu!(*M*)

Upper triangle of a matrix, overwriting *M* in the process. See also *triu()* (page 522).

triu!(*M*, *k::Integer*)

Returns the upper triangle of *M* starting from the *k*th superdiagonal, overwriting *M* in the process.

Example

```
julia> M = [1 2 3 4 5; 1 2 3 4 5; 1 2 3 4 5; 1 2 3 4 5; 1 2 3 4 5]
5×5 Array{Int64,2}:
 1  2  3  4  5
 1  2  3  4  5
 1  2  3  4  5
 1  2  3  4  5
 1  2  3  4  5

julia> triu!(M, 1)
5×5 Array{Int64,2}:
 0  2  3  4  5
 0  0  3  4  5
 0  0  0  4  5
 0  0  0  0  5
 0  0  0  0  0
```

tril(*M*)

Lower triangle of a matrix.

Example

```
julia> a = ones(4,4)
4×4 Array{Float64,2}:
 1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0

julia> tril(a)
4×4 Array{Float64,2}:
 1.0  0.0  0.0  0.0
 1.0  1.0  0.0  0.0
 1.0  1.0  1.0  0.0
 1.0  1.0  1.0  1.0
```

tril(*M*, *k::Integer*)

Returns the lower triangle of *M* starting from the *k*th superdiagonal.

Example

```
julia> a = ones(4,4)
4×4 Array{Float64,2}:
 1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0
```

```
julia> tril(a,3)
4×4 Array{Float64,2}:
 1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0

julia> tril(a,-3)
4×4 Array{Float64,2}:
 0.0  0.0  0.0  0.0
 0.0  0.0  0.0  0.0
 0.0  0.0  0.0  0.0
 1.0  0.0  0.0  0.0
```

tril!(M)

Lower triangle of a matrix, overwriting M in the process. See also `tril()` (page 523).

tril!(M, k::Integer)

Returns the lower triangle of M starting from the kth superdiagonal, overwriting M in the process.

Example

```
julia> M = [1 2 3 4 5; 1 2 3 4 5; 1 2 3 4 5; 1 2 3 4 5; 1 2 3 4 5]
5×5 Array{Int64,2}:
 1  2  3  4  5
 1  2  3  4  5
 1  2  3  4  5
 1  2  3  4  5
 1  2  3  4  5

julia> tril!(M, 2)
5×5 Array{Int64,2}:
 1  2  3  0  0
 1  2  3  4  0
 1  2  3  4  5
 1  2  3  4  5
 1  2  3  4  5
```

diagind(M, k::Integer=0)

A Range giving the indices of the kth diagonal of the matrix M.

Example

```
julia> A = [1 2 3; 4 5 6; 7 8 9]
3×3 Array{Int64,2}:
 1  2  3
 4  5  6
 7  8  9

julia> diagind(A, -1)
2:4:6
```

diag(M, k::Integer=0)

The kth diagonal of a matrix, as a vector. Use `diagm()` (page 525) to construct a diagonal matrix.

Example

```
julia> A = [1 2 3; 4 5 6; 7 8 9]
3×3 Array{Int64,2}:
 1  2  3
```

```

4 5 6
7 8 9

julia> diag(A,1)
2-element Array{Int64,1}:
 2
 6

```

diag(*v*, *k::Integer*=0)

Construct a matrix by placing *v* on the *k*th diagonal.

Example

```

julia> diagm([1,2,3],1)
4×4 Array{Int64,2}:
 0  1  0  0
 0  0  2  0
 0  0  0  3
 0  0  0  0

```

scale!(*A*, *b*)

scale!(*b*, *A*)

Scale an array *A* by a scalar *b* overwriting *A* in-place.

If *A* is a matrix and *b* is a vector, then `scale!(A,b)` scales each column *i* of *A* by *b*[*i*] (similar to `A*Diagonal(b)`), while `scale!(b,A)` scales each row *i* of *A* by *b*[*i*] (similar to `Diagonal(b)*A`), again operating in-place on *A*. An `InexactError` exception is thrown if the scaling produces a number not representable by the element type of *A*, e.g. for integer types.

Example

```

julia> a = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> b = [1; 2]
2-element Array{Int64,1}:
 1
 2

julia> scale!(a,b)
2×2 Array{Int64,2}:
 1  4
 3  8

julia> a = [1 2; 3 4];

julia> b = [1; 2];

julia> scale!(b,a)
2×2 Array{Int64,2}:
 1  2
 6  8

```

Tridiagonal(*dl*, *d*, *du*)

Construct a tridiagonal matrix from the first subdiagonal, diagonal, and first superdiagonal, respectively. The result is of type `Tridiagonal` and provides efficient specialized linear solvers, but may be converted into a regular matrix with `convert()` (page 364) (or `Array()` for short). The lengths of *dl* and *du* must be one less than the length of *d*.

Example

```

julia> dl = [1; 2; 3]
3-element Array{Int64,1}:
 1
 2
 3

julia> du = [4; 5; 6]
3-element Array{Int64,1}:
 4
 5
 6

julia> d = [7; 8; 9; 0]
4-element Array{Int64,1}:
 7
 8
 9
 0

julia> Tridiagonal(dl, d, du)
4×4 Tridiagonal{Int64}:
 7  4  .  .
 1  8  5  .
 .  2  9  6
 .  .  3  0

```

rank(M , $tol::Real$)

Compute the rank of a matrix by counting how many singular values of M have magnitude greater than tol . By default, the value of tol is the largest dimension of M multiplied by the `eps()` (page 366) of the `eltype()` (page 384) of M .

norm(A , $p::Real=2$)

Compute the p -norm of a vector or the operator norm of a matrix A , defaulting to the $p=2$ -norm.

For vectors, p can assume any numeric value (even though not all values produce a mathematically valid vector norm). In particular, `norm(A, Inf)` returns the largest value in `abs(A)`, whereas `norm(A, -Inf)` returns the smallest.

Example

```

julia> v = [3;-2;6]
3-element Array{Int64,1}:
 3
-2
 6

julia> norm(v)
7.0

julia> norm(v, Inf)
6.0

```

For matrices, the matrix norm induced by the vector p -norm is used, where valid values of p are 1, 2, or `Inf`. (Note that for sparse matrices, $p=2$ is currently not implemented.) Use `vecnorm()` (page 527) to compute the Frobenius norm.

Example

```
julia> A = [1 -2 -3; 2 3 -1]
2×3 Array{Int64,2}:
 1  -2  -3
 2   3  -1

julia> norm(A, Inf)
6.0
```

vecnorm(A[, p::Real=2])

For any iterable container A (including arrays of any dimension) of numbers (or any element type for which `norm` is defined), compute the p-norm (defaulting to `p=2`) as if A were a vector of the corresponding length.

For example, if A is a matrix and `p=2`, then this is equivalent to the Frobenius norm.

Example

```
julia> vecnorm([1 2 3; 4 5 6; 7 8 9])
16.881943016134134

julia> vecnorm([1 2 3 4 5 6 7 8 9])
16.881943016134134
```

normalize!(v[, p::Real=2])

Normalize the vector `v` in-place with respect to the p-norm. See also `vecnorm()` (page 527) and `normalize()` (page 527).

normalize(v[, p::Real=2])

Normalize the vector `v` with respect to the p-norm. See also `normalize!()` (page 527) and `vecnorm()` (page 527).

Example

```
julia> a = [1,2,4];

julia> normalize(a)
3-element Array{Float64,1}:
 0.218218
 0.436436
 0.872872

julia> normalize(a,1)
3-element Array{Float64,1}:
 0.142857
 0.285714
 0.571429
```

cond(M, p::Real=2)

Condition number of the matrix M, computed using the operator p-norm. Valid values for `p` are 1, 2 (default), or `Inf`.

condskeel(M[, x, p::Real=Inf])

$$\kappa_S(M, p) = \| |M| |M^{-1}| \|_p$$

$$\kappa_S(M, x, p) = \| |M| |M^{-1}| |x| \|_p$$

Skeel condition number κ_S of the matrix M, optionally with respect to the vector `x`, as computed using the operator p-norm. `p` is `Inf` by default, if not provided. Valid values for `p` are 1, 2, or `Inf`.

This quantity is also known in the literature as the Bauer condition number, relative condition number, or componentwise relative condition number.

trace(*M*)

Matrix trace. Sums the diagonal elements of *M*.

Example

```
julia> A = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> trace(A)
5
```

det(*M*)

Matrix determinant.

Example

```
julia> M = [1 0; 2 2]
2×2 Array{Int64,2}:
 1  0
 2  2

julia> det(M)
2.0
```

logdet(*M*)

Log of matrix determinant. Equivalent to $\log(\det(M))$, but may provide increased accuracy and/or speed.

Example

```
julia> M = [1 0; 2 2]
2×2 Array{Int64,2}:
 1  0
 2  2

julia> logdet(M)
0.6931471805599453

julia> logdet(eye(3))
0.0
```

logabsdet(*M*)

Log of absolute value of determinant of real matrix. Equivalent to $(\log(\text{abs}(\det(M))), \text{sign}(\det(M)))$, but may provide increased accuracy and/or speed.

inv(*M*)

Matrix inverse. Computes matrix *N* such that $M * N = I$, where *I* is the identity matrix. Computed by solving the left-division $N = M \setminus I$.

Example

```
julia> M = [2 5; 1 3]
2×2 Array{Int64,2}:
 2  5
 1  3

julia> N = inv(M)
2×2 Array{Float64,2}:
 3.0 -5.0
-1.0  2.0
```

```
julia> M*N == N*M == eye(2)
true
```

pinv (M , $tol::Real$)

Computes the Moore-Penrose pseudoinverse.

For matrices M with floating point elements, it is convenient to compute the pseudoinverse by inverting only singular values above a given threshold, tol .

The optimal choice of tol varies both with the value of M and the intended application of the pseudoinverse. The default value of tol is $\text{eps}(\text{real}(\text{float}(\text{one}(\text{eltype}(M)))) * \text{maximum}(\text{size}(A)))$, which is essentially machine epsilon for the real part of a matrix element multiplied by the larger matrix dimension. For inverting dense ill-conditioned matrices in a least-squares sense, $tol = \text{sqrt}(\text{eps}(\text{real}(\text{float}(\text{one}(\text{eltype}(M))))))$ is recommended.

For more information, see [\[issue8859\]](#) (page 621), [\[B96\]](#) (page 621), [\[S84\]](#) (page 621), [\[KY88\]](#) (page 621).

Example

```
julia> M = [1.5 1.3; 1.2 1.9]
2×2 Array{Float64,2}:
 1.5  1.3
 1.2  1.9

julia> N = pinv(M)
2×2 Array{Float64,2}:
 1.47287  -1.00775
-0.930233  1.16279

julia> M * N
2×2 Array{Float64,2}:
 1.0      -2.22045e-16
 4.44089e-16  1.0
```

nullspace (M)

Basis for nullspace of M .

Example

```
julia> M = [1 0 0; 0 1 0; 0 0 0]
3×3 Array{Int64,2}:
 1  0  0
 0  1  0
 0  0  0

julia> nullspace(M)
3×1 Array{Float64,2}:
 0.0
 0.0
 1.0
```

repmat (A , $m::Int$, $n::Int=1$)

Construct a matrix by repeating the given matrix m times in dimension 1 and n times in dimension 2.

```
julia> repmat([1, 2, 3], 2)
6-element Array{Int64,1}:
 1
 2
 3
 1
 2
 3
```

```

2
3

julia> repmat([1, 2, 3], 2, 3)
6×3 Array{Int64,2}:
 1  1  1
 2  2  2
 3  3  3
 1  1  1
 2  2  2
 3  3  3

```

repeat (*A::AbstractArray*; *inner=ntuple(x->1, ndims(A))*, *outer=ntuple(x->1, ndims(A))*)

Construct an array by repeating the entries of *A*. The *i*-th element of *inner* specifies the number of times that the individual entries of the *i*-th dimension of *A* should be repeated. The *i*-th element of *outer* specifies the number of times that a slice along the *i*-th dimension of *A* should be repeated. If *inner* or *outer* are omitted, no repetition is performed.

```

julia> repeat(1:2, inner=2)
4-element Array{Int64,1}:
 1
 1
 2
 2

julia> repeat(1:2, outer=2)
4-element Array{Int64,1}:
 1
 2
 1
 2

julia> repeat([1 2; 3 4], inner=(2, 1), outer=(1, 3))
4×6 Array{Int64,2}:
 1  2  1  2  1  2
 1  2  1  2  1  2
 3  4  3  4  3  4
 3  4  3  4  3  4

```

kron (*A*, *B*)

Kronecker tensor product of two vectors or two matrices.

Example

```

julia> A = [1 2; 3 4]
2×2 Array{Int64,2}:
 1  2
 3  4

julia> B = [im 1; 1 -im]
2×2 Array{Complex{Int64},2}:
 0+1im  1+0im
 1+0im  0-1im

julia> kron(A, B)
4×4 Array{Complex{Int64},2}:
 0+1im  1+0im  0+2im  2+0im
 1+0im  0-1im  2+0im  0-2im
 0+3im  3+0im  0+4im  4+0im

```



```
3+0im 0-3im 4+0im 0-4im
```

blkdiag(A...)

Concatenate matrices block-diagonally. Currently only implemented for sparse matrices.

linreg(x, y)

Perform simple linear regression using Ordinary Least Squares. Returns a and b such that $a + b \cdot x$ is the closest straight line to the given points (x, y) , i.e., such that the squared error between y and $a + b \cdot x$ is minimized.

Examples:

```
using PyPlot
x = 1.0:12.0
y = [5.5, 6.3, 7.6, 8.8, 10.9, 11.79, 13.48, 15.02, 17.77, 20.81, 22.0, 22.99]
a, b = linreg(x, y)           # Linear regression
plot(x, y, "o")               # Plot (x, y) points
plot(x, a + b*x)              # Plot line determined by linear regression
```

See also:

`\`, `cov()` (page 427), `std()` (page 425), `mean()` (page 425).

expm(A)

Compute the matrix exponential of A , defined by

$$e^A = \sum_{n=0}^{\infty} \frac{A^n}{n!}.$$

For symmetric or Hermitian A , an eigendecomposition (`eigfact()` (page 516)) is used, otherwise the scaling and squaring algorithm (see [H05] (page 621)) is chosen.

Example

```
julia> A = eye(2, 2)
2×2 Array{Float64,2}:
 1.0  0.0
 0.0  1.0

julia> expm(A)
2×2 Array{Float64,2}:
 2.71828  0.0
 0.0      2.71828
```

logm(A::StridedMatrix)

If A has no negative real eigenvalue, compute the principal matrix logarithm of A , i.e. the unique matrix X such that $e^X = A$ and $-\pi < \text{Im}(\lambda) < \pi$ for all the eigenvalues λ of X . If A has nonpositive eigenvalues, a nonprincipal matrix function is returned whenever possible.

If A is symmetric or Hermitian, its eigendecomposition (`eigfact()` (page 516)) is used, if A is triangular an improved version of the inverse scaling and squaring method is employed (see [AH12] (page 621) and [AHR13] (page 621)). For general matrices, the complex Schur form (`schur()` (page 518)) is computed and the triangular algorithm is used on the triangular factor.

Example

```
julia> A = 2.7182818 * eye(2)
2×2 Array{Float64,2}:
 2.71828  0.0
 0.0      2.71828
```

```
julia> logm(A)
2×2 Array{Float64,2}:
 1.0  0.0
 0.0  1.0
```

sqrtm(A)

If A has no negative real eigenvalues, compute the principal matrix square root of A, that is the unique matrix X with eigenvalues having positive real part such that $X^2 = A$. Otherwise, a nonprincipal square root is returned.

If A is symmetric or Hermitian, its eigendecomposition (*eigfact()* (page 516)) is used to compute the square root. Otherwise, the square root is determined by means of the Björck-Hammarling method, which computes the complex Schur form (*schur()* (page 518)) and then the complex square root of the triangular factor.

Example

```
julia> A = [4 0; 0 4]
2×2 Array{Int64,2}:
 4  0
 0  4

julia> sqrtm(A)
2×2 Array{Float64,2}:
 2.0  0.0
 0.0  2.0
```

lyap(A, C)

Computes the solution X to the continuous Lyapunov equation $AX + XA' + C = 0$, where no eigenvalue of A has a zero real part and no two eigenvalues are negative complex conjugates of each other.

sylvester(A, B, C)

Computes the solution X to the Sylvester equation $AX + XB + C = 0$, where A, B and C have compatible dimensions and A and $-B$ have no eigenvalues with equal real part.

issymmetric(A) → Bool

Test whether a matrix is symmetric.

Example

```
julia> a = [1 2; 2 -1]
2×2 Array{Int64,2}:
 1  2
 2 -1

julia> issymmetric(a)
true

julia> b = [1 im; -im 1]
2×2 Array{Complex{Int64},2}:
 1+0im  0+1im
 0-1im  1+0im

julia> issymmetric(b)
false
```

isposdef(A) → Bool

Test whether a matrix is positive definite.

Example

```
julia> A = [1 2; 2 50]
2×2 Array{Int64,2}:
 1  2
 2 50

julia> isposdef(A)
true
```

isposdef! (*A*) → Bool

Test whether a matrix is positive definite, overwriting *A* in the process.

istril (*A*) → Bool

Test whether a matrix is lower triangular.

Example

```
julia> a = [1 2; 2 -1]
2×2 Array{Int64,2}:
 1  2
 2 -1

julia> istril(a)
false

julia> b = [1 0; -im -1]
2×2 Array{Complex{Int64},2}:
 1+0im  0+0im
 0-1im -1+0im

julia> istril(b)
true
```

istriu (*A*) → Bool

Test whether a matrix is upper triangular.

Example

```
julia> a = [1 2; 2 -1]
2×2 Array{Int64,2}:
 1  2
 2 -1

julia> istriu(a)
false

julia> b = [1 im; 0 -1]
2×2 Array{Complex{Int64},2}:
 1+0im  0+1im
 0+0im -1+0im

julia> istriu(b)
true
```

isdiag (*A*) → Bool

Test whether a matrix is diagonal.

Example

```
julia> a = [1 2; 2 -1]
2×2 Array{Int64,2}:
 1  2
```

```
2 -1

julia> isdiag(a)
false

julia> b = [im 0; 0 -im]
2×2 Array{Complex{Int64},2}:
 0+1im  0+0im
 0+0im  0-1im

julia> isdiag(b)
true
```

ishermitian(*A*) → Bool

Test whether a matrix is Hermitian.

Example

```
julia> a = [1 2; 2 -1]
2×2 Array{Int64,2}:
 1  2
 2 -1

julia> ishermitian(a)
true

julia> b = [1 im; -im 1]
2×2 Array{Complex{Int64},2}:
 1+0im  0+1im
 0-1im  1+0im

julia> ishermitian(b)
true
```

transpose(*A*)

The transposition operator (*.*[']).

Example

```
julia> A = [1 2 3; 4 5 6; 7 8 9]
3×3 Array{Int64,2}:
 1  2  3
 4  5  6
 7  8  9

julia> transpose(A)
3×3 Array{Int64,2}:
 1  4  7
 2  5  8
 3  6  9
```

transpose!(*dest*, *src*)

Transpose array *src* and store the result in the preallocated array *dest*, which should have a size corresponding to `(size(src, 2), size(src, 1))`. No in-place transposition is supported and unexpected results will happen if *src* and *dest* have overlapping memory regions.

ctranspose(*A*)

The conjugate transposition operator (*'*).

Example

```
julia> A = [3+2im 9+2im; 8+7im 4+6im]
2×2 Array{Complex{Int64},2}:
 3+2im  9+2im
 8+7im  4+6im

julia> ctranspose(A)
2×2 Array{Complex{Int64},2}:
 3-2im  8-7im
 9-2im  4-6im
```

ctranspose! (*dest*, *src*)

Conjugate transpose array *src* and store the result in the preallocated array *dest*, which should have a size corresponding to `(size(src, 2), size(src, 1))`. No in-place transposition is supported and unexpected results will happen if *src* and *dest* have overlapping memory regions.

eigs (*A*; *nev*=6, *ncv*=max(20, 2**nev*+1), *which*="LM", *tol*=0.0, *maxiter*=300, *sigma*=nothing, *ritzvec*=true, *v0*=zeros((0,))) → (*d*, [*v*,], *nconv*, *niter*, *nmult*, *resid*)

Computes eigenvalues *d* of *A* using implicitly restarted Lanczos or Arnoldi iterations for real symmetric or general nonsymmetric matrices respectively.

The following keyword arguments are supported:

- nev*: Number of eigenvalues
- ncv*: Number of Krylov vectors used in the computation; should satisfy $nev+1 \leq ncv \leq n$ for real symmetric problems and $nev+2 \leq ncv \leq n$ for other problems, where *n* is the size of the input matrix *A*. The default is $ncv = \max(20, 2*nev+1)$. Note that these restrictions limit the input matrix *A* to be of dimension at least 2.
- which*: type of eigenvalues to compute. See the note below.

| which | type of eigenvalues |
|-------|--|
| :LM | eigenvalues of largest magnitude (default) |
| :SM | eigenvalues of smallest magnitude |
| :LR | eigenvalues of largest real part |
| :SR | eigenvalues of smallest real part |
| :LI | eigenvalues of largest imaginary part (nonsymmetric or complex <i>A</i> only) |
| :SI | eigenvalues of smallest imaginary part (nonsymmetric or complex <i>A</i> only) |
| :BE | compute half of the eigenvalues from each end of the spectrum, biased in favor of the high end. (real symmetric <i>A</i> only) |

- tol*: parameter defining the relative tolerance for convergence of Ritz values (eigenvalue estimates). A Ritz value θ is considered converged when its associated residual is less than or equal to the product of *tol* and $\max(\varepsilon^{2/3}, |\theta|)$, where $\varepsilon = \text{eps}(\text{real}(\text{eltype}(A))) / 2$ is LAPACK's machine epsilon. The residual associated with θ and its corresponding Ritz vector *v* is defined as the norm $\|Av - v\theta\|$. The specified value of *tol* should be positive; otherwise, it is ignored and ε is used instead. Default: ε .

- maxiter*: Maximum number of iterations (default = 300)
- sigma*: Specifies the level shift used in inverse iteration. If nothing (default), defaults to ordinary (forward) iterations. Otherwise, find eigenvalues close to *sigma* using shift and invert iterations.
- ritzvec*: Returns the Ritz vectors *v* (eigenvectors) if true
- v0*: starting vector from which to start the iterations

eigs returns the *nev* requested eigenvalues in *d*, the corresponding Ritz vectors *v* (only if *ritzvec*=true), the number of converged eigenvalues *nconv*, the number of iterations *niter* and the number of matrix vector multiplications *nmult*, as well as the final residual vector *resid*.

Note: The `sigma` and `which` keywords interact: the description of eigenvalues searched for by `which` do *not* necessarily refer to the eigenvalues of A , but rather the linear operator constructed by the specification of the iteration mode implied by `sigma`.

| <code>sigma</code> | iteration mode | which refers to eigenvalues of |
|----------------------|---|--------------------------------|
| <code>nothing</code> | ordinary (forward) | A |
| real or complex | inverse with level shift <code>sigma</code> | $(A - \sigma I)^{-1}$ |

Note: Although `tol` has a default value, the best choice depends strongly on the matrix A . We recommend that users `_always_` specify a value for `tol` which suits their specific needs.

For details of how the errors in the computed eigenvalues are estimated, see:

- 2. (a)Parlett, “The Symmetric Eigenvalue Problem”, SIAM: Philadelphia, 2/e (1998), Ch. 13.2, “Accessing Accuracy in Lanczos Problems”, pp. 290-292 ff.
- 18. (a)Lehoucq and D. C. Sorensen, “Deflation Techniques for an Implicitly Restarted Arnoldi Iteration”, SIAM Journal on Matrix Analysis and Applications (1996), 17(4), 789–821. doi:10.1137/S0895479895281484

eigs (A , B ; `nev=6`, `ncv=max(20, 2*nev+1)`, `which="LM"`, `tol=0.0`, `maxiter=300`, `sigma=nothing`, `ritzvec=true`, `v0=zeros((0,))`) \rightarrow (`d`,`[v,]`,`nconv`,`niter`,`nmult`,`resid`)
 Computes generalized eigenvalues d of A and B using implicitly restarted Lanczos or Arnoldi iterations for real symmetric or general nonsymmetric matrices respectively.

The following keyword arguments are supported:

- `nev`: Number of eigenvalues
- `ncv`: Number of Krylov vectors used in the computation; should satisfy `nev+1 <= ncv <= n` for real symmetric problems and `nev+2 <= ncv <= n` for other problems, where n is the size of the input matrices A and B . The default is `ncv = max(20, 2*nev+1)`. Note that these restrictions limit the input matrix A to be of dimension at least 2.
- `which`: type of eigenvalues to compute. See the note below.

| <code>which</code> | type of eigenvalues |
|--------------------|---|
| <code>:LM</code> | eigenvalues of largest magnitude (default) |
| <code>:SM</code> | eigenvalues of smallest magnitude |
| <code>:LR</code> | eigenvalues of largest real part |
| <code>:SR</code> | eigenvalues of smallest real part |
| <code>:LI</code> | eigenvalues of largest imaginary part (nonsymmetric or complex A only) |
| <code>:SI</code> | eigenvalues of smallest imaginary part (nonsymmetric or complex A only) |
| <code>:BE</code> | compute half of the eigenvalues from each end of the spectrum, biased in favor of the high end. (real symmetric A only) |

- `tol`: relative tolerance used in the convergence criterion for eigenvalues, similar to `tol` in the `eigs()` (page 535) method for the ordinary eigenvalue problem, but effectively for the eigenvalues of $B^{-1}A$ instead of A . See the documentation for the ordinary eigenvalue problem in `eigs()` (page 535) and the accompanying note about `tol`.
- `maxiter`: Maximum number of iterations (default = 300)
- `sigma`: Specifies the level shift used in inverse iteration. If `nothing` (default), defaults to ordinary (forward) iterations. Otherwise, find eigenvalues close to `sigma` using shift and invert iterations.
- `ritzvec`: Returns the Ritz vectors v (eigenvectors) if `true`

- `v0`: starting vector from which to start the iterations

`eigs` returns the `nev` requested eigenvalues in `d`, the corresponding Ritz vectors `v` (only if `ritzvec=true`), the number of converged eigenvalues `nconv`, the number of iterations `niter` and the number of matrix vector multiplications `nmult`, as well as the final residual vector `resid`.

Example

```
X = sprand(10, 5, 0.2)
eigs(X, nsv = 2, tol = 1e-3)
```

Note: The `sigma` and `which` keywords interact: the description of eigenvalues searched for by `which` do *not* necessarily refer to the eigenvalue problem $Av = Bv\lambda$, but rather the linear operator constructed by the specification of the iteration mode implied by `sigma`.

| <code>sigma</code> | iteration mode | which refers to the problem |
|------------------------------|---|-----------------------------------|
| <code>nothing</code> | ordinary (forward) | $Av = Bv\lambda$ |
| <code>real or complex</code> | inverse with level shift <code>sigma</code> | $(A - \sigma B)^{-1}B = v\lambda$ |

svds (`A`; `nsv=6`, `ritzvec=true`, `tol=0.0`, `maxiter=1000`, `ncv=2*nsv`, `u0=zeros((0,))`, `v0=zeros((0,))`) \rightarrow (`SVD([left_sv,], s, [right_sv,])`, `nconv`, `niter`, `nmult`, `resid`)
Computes the largest singular values `s` of `A` using implicitly restarted Lanczos iterations derived from `eigs()` (page 535).

Inputs

- `A`: Linear operator whose singular values are desired. `A` may be represented as a subtype of `AbstractArray`, e.g., a sparse matrix, or any other type supporting the four methods `size(A)`, `eltype(A)`, `A * vector`, and `A' * vector`.
- `nsv`: Number of singular values. Default: 6.
- `ritzvec`: If `true`, return the left and right singular vectors `left_sv` and `right_sv`. If `false`, omit the singular vectors. Default: `true`.
- `tol`: tolerance, see `eigs()` (page 535).
- `maxiter`: Maximum number of iterations, see `eigs()` (page 535). Default: 1000.
- `ncv`: Maximum size of the Krylov subspace, see `eigs()` (page 535) (there called `nev`). Default: `2*nsv`.
- `u0`: Initial guess for the first left Krylov vector. It may have length `m` (the first dimension of `A`), or 0.
- `v0`: Initial guess for the first right Krylov vector. It may have length `n` (the second dimension of `A`), or 0.

Outputs

- `svd`: An SVD object containing the left singular vectors, the requested values, and the right singular vectors. If `ritzvec = false`, the left and right singular vectors will be empty.
- `nconv`: Number of converged singular values.
- `niter`: Number of iterations.
- `nmult`: Number of matrix–vector products used.
- `resid`: Final residual vector.

Example

```
X = sprand(10, 5, 0.2)
svds(X, nsv = 2)
```

Implementation note

`svds(A)` is formally equivalent to calling `eigs` to perform implicitly restarted Lanczos tridiagonalization on the Hermitian matrix $\begin{pmatrix} 0 & A' \\ A & 0 \end{pmatrix}$, whose eigenvalues are plus and minus the singular values of A .

peakflops ($n::Integer=2000$; $parallel::Bool=false$)

`peakflops` computes the peak flop rate of the computer by using double precision `Base.LinAlg.BLAS.gemm!()` (page 541). By default, if no arguments are specified, it multiplies a matrix of size $n \times n$, where $n = 2000$. If the underlying BLAS is using multiple threads, higher flop rates are realized. The number of BLAS threads can be set with `Base.LinAlg.BLAS.set_num_threads()` (page 542).

If the keyword argument `parallel` is set to `true`, `peakflops` is run in parallel on all the worker processors. The flop rate of the entire parallel computer is returned. When running in parallel, only 1 BLAS thread is used. The argument n still refers to the size of the problem that is solved on each processor.

47.2 Low-level matrix operations

Matrix operations involving transpositions operations like $A' \setminus B$ are converted by the Julia parser into calls to specially named functions like `A_c_ldiv_B`. If you want to overload these operations for your own types, then it is useful to know the names of these functions.

Also, in many cases there are in-place versions of matrix operations that allow you to supply a pre-allocated output vector or matrix. This is useful when optimizing critical code in order to avoid the overhead of repeated allocations. These in-place operations are suffixed with `!` below (e.g. `A_mul_B!`) according to the usual Julia convention.

A_ldiv_B! ($[Y], A, B$) $\rightarrow Y$

Compute $A \setminus B$ in-place and store the result in Y , returning the result. If only two arguments are passed, then `A_ldiv_B!(A, B)` overwrites B with the result.

The argument A should *not* be a matrix. Rather, instead of matrices it should be a factorization object (e.g. produced by `factorize()` (page 502) or `cholfact()` (page 509)). The reason for this is that factorization itself is both expensive and typically allocates memory (although it can also be done in-place via, e.g., `luifact!()` (page 509)), and performance-critical situations requiring `A_ldiv_B!` usually also require fine-grained control over the factorization of A .

A_ldiv_Bc (A, B)

For matrices or vectors A and B , calculates $A \setminus B^H$.

A_ldiv_Bt (A, B)

For matrices or vectors A and B , calculates $A \setminus B^T$.

A_mul_B! (Y, A, B) $\rightarrow Y$

Calculates the matrix-matrix or matrix-vector product $A \cdot B$ and stores the result in Y , overwriting the existing value of Y . Note that Y must not be aliased with either A or B .

Example

```
julia> A=[1.0 2.0; 3.0 4.0]; B=[1.0 1.0; 1.0 1.0]; Y = similar(B); A_mul_B!(Y, A, B);

julia> Y
2×2 Array{Float64,2}:
 3.0  3.0
 7.0  7.0
```

A_mul_Bc (A, B)

For matrices or vectors A and B , calculates $A \cdot B^H$.

A_mul_Bt (A, B)

For matrices or vectors A and B , calculates $A \cdot B^T$.

A_rdiv_Bc (A, B)

For matrices or vectors A and B , calculates A/B^H .

A_rdiv_Bt (A, B)

For matrices or vectors A and B , calculates A/B^T .

Ac_ldiv_B (A, B)

For matrices or vectors A and B , calculates $A^H \setminus B$.

Ac_ldiv_B! ($[Y], A, B$) $\rightarrow Y$

Similar to [A_ldiv_B!\(\)](#) (page 538), but return $A^H \setminus B$, computing the result in-place in Y (or overwriting B if Y is not supplied).

Ac_ldiv_Bc (A, B)

For matrices or vectors A and B , calculates $A^H \setminus B^H$.

Ac_mul_B (A, B)

For matrices or vectors A and B , calculates $A^H \cdot B$.

Ac_mul_Bc (A, B)

For matrices or vectors A and B , calculates $A^H B^H$.

Ac_rdiv_B (A, B)

For matrices or vectors A and B , calculates A^H/B .

Ac_rdiv_Bc (A, B)

For matrices or vectors A and B , calculates A^H/B^H .

At_ldiv_B (A, B)

For matrices or vectors A and B , calculates $A^T \setminus B$.

At_ldiv_B! ($[Y], A, B$) $\rightarrow Y$

Similar to [A_ldiv_B!\(\)](#) (page 538), but return $A^T \setminus B$, computing the result in-place in Y (or overwriting B if Y is not supplied).

At_ldiv_Bt (A, B)

For matrices or vectors A and B , calculates $A^T \setminus B^T$.

At_mul_B (A, B)

For matrices or vectors A and B , calculates $A^T \cdot B$.

At_mul_Bt (A, B)

For matrices or vectors A and B , calculates $A^T \cdot B^T$.

At_rdiv_B (A, B)

For matrices or vectors A and B , calculates A^T/B .

At_rdiv_Bt (A, B)

For matrices or vectors A and B , calculates A^T/B^T .

47.3 BLAS Functions

In Julia (as in much of scientific computation), dense linear-algebra operations are based on the [LAPACK library](#), which in turn is built on top of basic linear-algebra building-blocks known as the [BLAS](#). There are highly optimized implementations of BLAS available for every computer architecture, and sometimes in high-performance linear algebra routines it is useful to call the BLAS functions directly.

`Base.LinAlg.BLAS` (page 539) provides wrappers for some of the BLAS functions. Those BLAS functions that overwrite one of the input arrays have names ending in `!`. Usually, a BLAS function has four methods defined, for `Float64`, `Float32`, `Complex128`, and `Complex64` arrays.

dot (*n*, *X*, *incx*, *Y*, *incy*)

Dot product of two vectors consisting of *n* elements of array *X* with stride *incx* and *n* elements of array *Y* with stride *incy*.

dotu (*n*, *X*, *incx*, *Y*, *incy*)

Dot function for two complex vectors.

dotc (*n*, *X*, *incx*, *U*, *incy*)

Dot function for two complex vectors conjugating the first vector.

blascopy! (*n*, *X*, *incx*, *Y*, *incy*)

Copy *n* elements of array *X* with stride *incx* to array *Y* with stride *incy*. Returns *Y*.

nrm2 (*n*, *X*, *incx*)

2-norm of a vector consisting of *n* elements of array *X* with stride *incx*.

asum (*n*, *X*, *incx*)

Sum of the absolute values of the first *n* elements of array *X* with stride *incx*.

axpy! (*a*, *X*, *Y*)

Overwrite *Y* with $a \cdot X + Y$. Returns *Y*.

scal! (*n*, *a*, *X*, *incx*)

Overwrite *X* with $a \cdot X$ for the first *n* elements of array *X* with stride *incx*. Returns *X*.

scal (*n*, *a*, *X*, *incx*)

Returns *X* scaled by *a* for the first *n* elements of array *X* with stride *incx*.

ger! (*alpha*, *x*, *y*, *A*)

Rank-1 update of the matrix *A* with vectors *x* and *y* as $\alpha \cdot x \cdot y' + A$.

syr! (*uplo*, *alpha*, *x*, *A*)

Rank-1 update of the symmetric matrix *A* with vector *x* as $\alpha \cdot x \cdot x'$ + *A*. When *uplo* is 'U' the upper triangle of *A* is updated ('L' for lower triangle). Returns *A*.

syrk! (*uplo*, *trans*, *alpha*, *A*, *beta*, *C*)

Rank-k update of the symmetric matrix *C* as $\alpha \cdot A \cdot A' + \beta \cdot C$ or $\alpha \cdot A' \cdot A + \beta \cdot C$ according to whether *trans* is 'N' or 'T'. When *uplo* is 'U' the upper triangle of *C* is updated ('L' for lower triangle). Returns *C*.

syrk (*uplo*, *trans*, *alpha*, *A*)

Returns either the upper triangle or the lower triangle, according to *uplo* ('U' or 'L'), of $\alpha \cdot A \cdot A'$ or $\alpha \cdot A' \cdot A$, according to *trans* ('N' or 'T').

her! (*uplo*, *alpha*, *x*, *A*)

Methods for complex arrays only. Rank-1 update of the Hermitian matrix *A* with vector *x* as $\alpha \cdot x \cdot x'$ + *A*. When *uplo* is 'U' the upper triangle of *A* is updated ('L' for lower triangle). Returns *A*.

herk! (*uplo*, *trans*, *alpha*, *A*, *beta*, *C*)

Methods for complex arrays only. Rank-k update of the Hermitian matrix *C* as $\alpha \cdot A \cdot A' + \beta \cdot C$ or $\alpha \cdot A' \cdot A + \beta \cdot C$ according to whether *trans* is 'N' or 'T'. When *uplo* is 'U' the upper triangle of *C* is updated ('L' for lower triangle). Returns *C*.

herk (*uplo*, *trans*, *alpha*, *A*)

Methods for complex arrays only. Returns either the upper triangle or the lower triangle, according to *uplo* ('U' or 'L'), of $\alpha \cdot A \cdot A'$ or $\alpha \cdot A' \cdot A$, according to *trans* ('N' or 'T').

gbmv! (*trans, m, kl, ku, alpha, A, x, beta, y*)

Update vector y as $\alpha A x + \beta y$ or $\alpha A' x + \beta y$ according to *trans* ('N' or 'T'). The matrix A is a general band matrix of dimension m by $\text{size}(A, 2)$ with kl sub-diagonals and ku super-diagonals. Returns the updated y .

gbmv (*trans, m, kl, ku, alpha, A, x, beta, y*)

Returns $\alpha A x$ or $\alpha A' x$ according to *trans* ('N' or 'T'). The matrix A is a general band matrix of dimension m by $\text{size}(A, 2)$ with kl sub-diagonals and ku super-diagonals.

sbmv! (*uplo, k, alpha, A, x, beta, y*)

Update vector y as $\alpha A x + \beta y$ where A is a symmetric band matrix of order $\text{size}(A, 2)$ with k super-diagonals stored in the argument A . The storage layout for A is described the reference BLAS module, level-2 BLAS at <http://www.netlib.org/lapack/explore-html/>.

Returns the updated y .

sbmv (*uplo, k, alpha, A, x*)

Returns $\alpha A x$ where A is a symmetric band matrix of order $\text{size}(A, 2)$ with k super-diagonals stored in the argument A .

sbmv (*uplo, k, A, x*)

Returns $A x$ where A is a symmetric band matrix of order $\text{size}(A, 2)$ with k super-diagonals stored in the argument A .

gemm! (*tA, tB, alpha, A, B, beta, C*)

Update C as $\alpha A B + \beta C$ or the other three variants according to tA (transpose A) and tB . Returns the updated C .

gemm (*tA, tB, alpha, A, B*)

Returns $\alpha A B$ or the other three variants according to tA (transpose A) and tB .

gemm (*tA, tB, A, B*)

Returns $A B$ or the other three variants according to tA (transpose A) and tB .

gemv! (*tA, alpha, A, x, beta, y*)

Update the vector y as $\alpha A x + \beta y$ or $\alpha A' x + \beta y$ according to tA (transpose A). Returns the updated y .

gemv (*tA, alpha, A, x*)

Returns $\alpha A x$ or $\alpha A' x$ according to tA (transpose A).

gemv (*tA, A, x*)

Returns $A x$ or $A' x$ according to tA (transpose A).

symm! (*side, ul, alpha, A, B, beta, C*)

Update C as $\alpha A B + \beta C$ or $\alpha B A + \beta C$ according to *side*. A is assumed to be symmetric. Only the *ul* triangle of A is used. Returns the updated C .

symm (*side, ul, alpha, A, B*)

Returns $\alpha A B$ or $\alpha B A$ according to *side*. A is assumed to be symmetric. Only the *ul* triangle of A is used.

symm (*side, ul, A, B*)

Returns $A B$ or $B A$ according to *side*. A is assumed to be symmetric. Only the *ul* triangle of A is used.

symm (*tA, tB, alpha, A, B*)

Returns $\alpha A B$ or the other three variants according to tA (transpose A) and tB .

symv! (*ul, alpha, A, x, beta, y*)

Update the vector y as $\alpha A x + \beta y$. A is assumed to be symmetric. Only the *ul* triangle of A is used. Returns the updated y .

symv (*ul*, *alpha*, *A*, *x*)

Returns $\alpha A x$. *A* is assumed to be symmetric. Only the *ul* triangle of *A* is used.

symv (*ul*, *A*, *x*)

Returns $A x$. *A* is assumed to be symmetric. Only the *ul* triangle of *A* is used.

trmm! (*side*, *ul*, *tA*, *dA*, *alpha*, *A*, *B*)

Update *B* as $\alpha A B$ or one of the other three variants determined by *side* (*A* on left or right) and *tA* (transpose *A*). Only the *ul* triangle of *A* is used. *dA* indicates if *A* is unit-triangular (the diagonal is assumed to be all ones). Returns the updated *B*.

trmm (*side*, *ul*, *tA*, *dA*, *alpha*, *A*, *B*)

Returns $\alpha A B$ or one of the other three variants determined by *side* (*A* on left or right) and *tA* (transpose *A*). Only the *ul* triangle of *A* is used. *dA* indicates if *A* is unit-triangular (the diagonal is assumed to be all ones).

trsm! (*side*, *ul*, *tA*, *dA*, *alpha*, *A*, *B*)

Overwrite *B* with the solution to $A X = \alpha B$ or one of the other three variants determined by *side* (*A* on left or right of *X*) and *tA* (transpose *A*). Only the *ul* triangle of *A* is used. *dA* indicates if *A* is unit-triangular (the diagonal is assumed to be all ones). Returns the updated *B*.

trsm (*side*, *ul*, *tA*, *dA*, *alpha*, *A*, *B*)

Returns the solution to $A X = \alpha B$ or one of the other three variants determined by *side* (*A* on left or right of *X*) and *tA* (transpose *A*). Only the *ul* triangle of *A* is used. *dA* indicates if *A* is unit-triangular (the diagonal is assumed to be all ones).

trmv! (*ul*, *tA*, *dA*, *A*, *b*)

Returns $\text{op}(A) * b$, where *op* is determined by *tA* (N for identity, T for transpose *A*, and C for conjugate transpose *A*). Only the *ul* triangle (U for upper, L for lower) of *A* is used. *dA* indicates if *A* is unit-triangular (the diagonal is assumed to be all ones if U, or non-unit if N). The multiplication occurs in-place on *b*.

trmv (*ul*, *tA*, *dA*, *A*, *b*)

Returns $\text{op}(A) * b$, where *op* is determined by *tA* (N for identity, T for transpose *A*, and C for conjugate transpose *A*). Only the *ul* triangle (U for upper, L for lower) of *A* is used. *dA* indicates if *A* is unit-triangular (the diagonal is assumed to be all ones if U, or non-unit if N).

trsv! (*ul*, *tA*, *dA*, *A*, *b*)

Overwrite *b* with the solution to $A x = b$ or one of the other two variants determined by *tA* (transpose *A*) and *ul* (triangle of *A* used). *dA* indicates if *A* is unit-triangular (the diagonal is assumed to be all ones). Returns the updated *b*.

trsv (*ul*, *tA*, *dA*, *A*, *b*)

Returns the solution to $A x = b$ or one of the other two variants determined by *tA* (transpose *A*) and *ul* (triangle of *A* is used.) *dA* indicates if *A* is unit-triangular (the diagonal is assumed to be all ones).

set_num_threads (*n*)

Set the number of threads the BLAS library should use.

I

An object of type `UniformScaling`, representing an identity matrix of any size.

Example

```
julia> ones(5, 6) * I == ones(5, 6)
true

julia> [1 2im 3; 1im 2 3] * I
2×3 Array{Complex{Int64},2}:
 1+0im  0+2im  3+0im
 0+1im  2+0im  3+0im
```

47.4 LAPACK Functions

`Base.LinAlg.LAPACK` (page 543) provides wrappers for some of the LAPACK functions for linear algebra. Those functions that overwrite one of the input arrays have names ending in `!`.

Usually a function has 4 methods defined, one each for `Float64`, `Float32`, `Complex128` and `Complex64` arrays.

Note that the LAPACK API provided by Julia can and will change in the future. Since this API is not user-facing, there is no commitment to support/deprecate this specific set of functions in future releases.

gbtrf! (*kl, ku, m, AB*) \rightarrow (*AB, ipiv*)

Compute the LU factorization of a banded matrix *AB*. *kl* is the first subdiagonal containing a nonzero band, *ku* is the last superdiagonal containing one, and *m* is the first dimension of the matrix *AB*. Returns the LU factorization in-place and *ipiv*, the vector of pivots used.

gbtrs! (*trans, kl, ku, m, AB, ipiv, B*)

Solve the equation $AB * X = B$. *trans* determines the orientation of *AB*. It may be `N` (no transpose), `T` (transpose), or `C` (conjugate transpose). *kl* is the first subdiagonal containing a nonzero band, *ku* is the last superdiagonal containing one, and *m* is the first dimension of the matrix *AB*. *ipiv* is the vector of pivots returned from `gbtrf!`. Returns the vector or matrix *X*, overwriting *B* in-place.

gebal! (*job, A*) \rightarrow (*ilo, ihi, scale*)

Balance the matrix *A* before computing its eigensystem or Schur factorization. *job* can be one of `N` (*A* will not be permuted or scaled), `P` (*A* will only be permuted), `S` (*A* will only be scaled), or `B` (*A* will be both permuted and scaled). Modifies *A* in-place and returns *ilo*, *ihi*, and *scale*. If permuting was turned on, $A[i, j] = 0$ if $j > i$ and $1 < j < ilo$ or $j > ihi$. *scale* contains information about the scaling/permutations performed.

gebak! (*job, side, ilo, ihi, scale, V*)

Transform the eigenvectors *V* of a matrix balanced using `gebal!` to the unscaled/unpermuted eigenvectors of the original matrix. Modifies *V* in-place. *side* can be `L` (left eigenvectors are transformed) or `R` (right eigenvectors are transformed).

gebrd! (*A*) \rightarrow (*A, d, e, tauq, taup*)

Reduce *A* in-place to bidiagonal form $A = QBP'$. Returns *A*, containing the bidiagonal matrix *B*; *d*, containing the diagonal elements of *B*; *e*, containing the off-diagonal elements of *B*; *tauq*, containing the elementary reflectors representing *Q*; and *taup*, containing the elementary reflectors representing *P*.

gelqf! (*A, tau*)

Compute the LQ factorization of *A*, $A = LQ$. *tau* contains scalars which parameterize the elementary reflectors of the factorization. *tau* must have length greater than or equal to the smallest dimension of *A*.

Returns *A* and *tau* modified in-place.

gelqf! (*A*) \rightarrow (*A, tau*)

Compute the LQ factorization of *A*, $A = LQ$.

Returns *A*, modified in-place, and *tau*, which contains scalars which parameterize the elementary reflectors of the factorization.

geqlf! (*A, tau*)

Compute the QL factorization of *A*, $A = QL$. *tau* contains scalars which parameterize the elementary reflectors of the factorization. *tau* must have length greater than or equal to the smallest dimension of *A*.

Returns *A* and *tau* modified in-place.

geqlf! (*A*) \rightarrow (*A, tau*)

Compute the QL factorization of *A*, $A = QL$.

Returns A , modified in-place, and τ , which contains scalars which parameterize the elementary reflectors of the factorization.

geqrf! (A, τ)

Compute the QR factorization of A , $A = QR$. τ contains scalars which parameterize the elementary reflectors of the factorization. τ must have length greater than or equal to the smallest dimension of A .

Returns A and τ modified in-place.

geqrf! (A) \rightarrow (A, τ)

Compute the QR factorization of A , $A = QR$.

Returns A , modified in-place, and τ , which contains scalars which parameterize the elementary reflectors of the factorization.

geqp3! ($A, jpvt, \tau$)

Compute the pivoted QR factorization of A , $AP = QR$ using BLAS level 3. P is a pivoting matrix, represented by $jpvt$. τ stores the elementary reflectors. $jpvt$ must have length greater than or equal to n if A is an $(m \times n)$ matrix. τ must have length greater than or equal to the smallest dimension of A .

A , $jpvt$, and τ are modified in-place.

geqp3! ($A, jpvt$) \rightarrow ($A, jpvt, \tau$)

Compute the pivoted QR factorization of A , $AP = QR$ using BLAS level 3. P is a pivoting matrix, represented by $jpvt$. $jpvt$ must have length greater than or equal to n if A is an $(m \times n)$ matrix.

Returns A and $jpvt$, modified in-place, and τ , which stores the elementary reflectors.

geqp3! (A) \rightarrow ($A, jpvt, \tau$)

Compute the pivoted QR factorization of A , $AP = QR$ using BLAS level 3.

Returns A , modified in-place, $jpvt$, which represents the pivoting matrix P , and τ , which stores the elementary reflectors.

gerqf! (A, τ)

Compute the RQ factorization of A , $A = RQ$. τ contains scalars which parameterize the elementary reflectors of the factorization. τ must have length greater than or equal to the smallest dimension of A .

Returns A and τ modified in-place.

gerqf! (A) \rightarrow (A, τ)

Compute the RQ factorization of A , $A = RQ$.

Returns A , modified in-place, and τ , which contains scalars which parameterize the elementary reflectors of the factorization.

geqrt! (A, T)

Compute the blocked QR factorization of A , $A = QR$. T contains upper triangular block reflectors which parameterize the elementary reflectors of the factorization. The first dimension of T sets the block size and it must be between 1 and n . The second dimension of T must equal the smallest dimension of A .

Returns A and T modified in-place.

geqrt! (A, nb) \rightarrow (A, T)

Compute the blocked QR factorization of A , $A = QR$. nb sets the block size and it must be between 1 and n , the second dimension of A .

Returns A , modified in-place, and T , which contains upper triangular block reflectors which parameterize the elementary reflectors of the factorization.

geqrt3! (A, T)

Recursively computes the blocked QR factorization of A , $A = QR$. T contains upper triangular block reflectors

which parameterize the elementary reflectors of the factorization. The first dimension of `T` sets the block size and it must be between 1 and `n`. The second dimension of `T` must equal the smallest dimension of `A`.

Returns `A` and `T` modified in-place.

geqrt3! (`A`) \rightarrow (`A`, `T`)

Recursively computes the blocked QR factorization of `A`, $A = QR$.

Returns `A`, modified in-place, and `T`, which contains upper triangular block reflectors which parameterize the elementary reflectors of the factorization.

getrf! (`A`) \rightarrow (`A`, `ipiv`, `info`)

Compute the pivoted LU factorization of `A`, $A = LU$.

Returns `A`, modified in-place, `ipiv`, the pivoting information, and an `info` code which indicates success (`info` = 0), a singular value in `U` (`info` = `i`, in which case $U[i, i]$ is singular), or an error code (`info` < 0).

tzrzf! (`A`) \rightarrow (`A`, `tau`)

Transforms the upper trapezoidal matrix `A` to upper triangular form in-place. Returns `A` and `tau`, the scalar parameters for the elementary reflectors of the transformation.

ormrz! (`side`, `trans`, `A`, `tau`, `C`)

Multiplies the matrix `C` by Q from the transformation supplied by `tzrzf!`. Depending on `side` or `trans` the multiplication can be left-sided (`side` = `L`, $Q \cdot C$) or right-sided (`side` = `R`, $C \cdot Q$) and Q can be unmodified (`trans` = `N`), transposed (`trans` = `T`), or conjugate transposed (`trans` = `C`). Returns matrix `C` which is modified in-place with the result of the multiplication.

gels! (`trans`, `A`, `B`) \rightarrow (`F`, `B`, `ssr`)

Solves the linear equation $A \cdot X = B$, $A' \cdot X = B$, or $A' \cdot X = B$ using a QR or LQ factorization. Modifies the matrix/vector `B` in place with the solution. `A` is overwritten with its QR or LQ factorization. `trans` may be one of `N` (no modification), `T` (transpose), or `C` (conjugate transpose). `gels!` searches for the minimum norm/least squares solution. `A` may be under or over determined. The solution is returned in `B`.

gesv! (`A`, `B`) \rightarrow (`B`, `A`, `ipiv`)

Solves the linear equation $A \cdot X = B$ where `A` is a square matrix using the LU factorization of `A`. `A` is overwritten with its LU factorization and `B` is overwritten with the solution `X`. `ipiv` contains the pivoting information for the LU factorization of `A`.

getrs! (`trans`, `A`, `ipiv`, `B`)

Solves the linear equation $A \cdot X = B$, $A' \cdot X = B$, or $A' \cdot X = B$ for square `A`. Modifies the matrix/vector `B` in place with the solution. `A` is the LU factorization from `getrf!`, with `ipiv` the pivoting information. `trans` may be one of `N` (no modification), `T` (transpose), or `C` (conjugate transpose).

getri! (`A`, `ipiv`)

Computes the inverse of `A`, using its LU factorization found by `getrf!`. `ipiv` is the pivot information output and `A` contains the LU factorization of `getrf!`. `A` is overwritten with its inverse.

gesvx! (`fact`, `trans`, `A`, `AF`, `ipiv`, `equed`, `R`, `C`, `B`) \rightarrow (`X`, `equed`, `R`, `C`, `B`, `rcond`, `ferr`, `berr`, `work`)

Solves the linear equation $A \cdot X = B$ (`trans` = `N`), $A' \cdot X = B$ (`trans` = `T`), or $A' \cdot X = B$ (`trans` = `C`) using the LU factorization of `A`. `fact` may be `E`, in which case `A` will be equilibrated and copied to `AF`; `F`, in which case `AF` and `ipiv` from a previous LU factorization are inputs; or `N`, in which case `A` will be copied to `AF` and then factored. If `fact` = `F`, `equed` may be `N`, meaning `A` has not been equilibrated; `R`, meaning `A` was multiplied by $\text{diagm}(R)$ from the left; `C`, meaning `A` was multiplied by $\text{diagm}(C)$ from the right; or `B`, meaning `A` was multiplied by $\text{diagm}(R)$ from the left and $\text{diagm}(C)$ from the right. If `fact` = `F` and `equed` = `R` or `B` the elements of `R` must all be positive. If `fact` = `F` and `equed` = `C` or `B` the elements of `C` must all be positive.

Returns the solution `X`; `equed`, which is an output if `fact` is not `N`, and describes the equilibration that was performed; `R`, the row equilibration diagonal; `C`, the column equilibration diagonal; `B`, which may be overwritten with its equilibrated form $\text{diagm}(R) \cdot B$ (if `trans` = `N` and `equed` = `R`, `B`) or $\text{diagm}(C) \cdot B$ (if `trans` =

T, C and $\text{equed} = C, B$; rcond , the reciprocal condition number of A after equilibrating; ferr , the forward error bound for each solution vector in X ; berr , the forward error bound for each solution vector in X ; and work , the reciprocal pivot growth factor.

gesvx! (A, B)

The no-equilibration, no-transpose simplification of `gesvx!`.

gelsd! (A, B, rcond) $\rightarrow (B, \text{rnk})$

Computes the least norm solution of $A * X = B$ by finding the SVD factorization of A , then dividing-and-conquering the problem. B is overwritten with the solution X . Singular values below rcond will be treated as zero. Returns the solution in B and the effective rank of A in rnk .

gelsy! (A, B, rcond) $\rightarrow (B, \text{rnk})$

Computes the least norm solution of $A * X = B$ by finding the full QR factorization of A , then dividing-and-conquering the problem. B is overwritten with the solution X . Singular values below rcond will be treated as zero. Returns the solution in B and the effective rank of A in rnk .

gglse! (A, c, B, d) $\rightarrow (X, \text{res})$

Solves the equation $A * x = c$ where x is subject to the equality constraint $B * x = d$. Uses the formula $\|c - A*x\|^2 = 0$ to solve. Returns X and the residual sum-of-squares.

geev! ($\text{jobvl}, \text{jobvr}, A$) $\rightarrow (W, VL, VR)$

Finds the eigensystem of A . If $\text{jobvl} = N$, the left eigenvectors of A aren't computed. If $\text{jobvr} = N$, the right eigenvectors of A aren't computed. If $\text{jobvl} = V$ or $\text{jobvr} = V$, the corresponding eigenvectors are computed. Returns the eigenvalues in W , the right eigenvectors in VR , and the left eigenvectors in VL .

gesdd! (job, A) $\rightarrow (U, S, VT)$

Finds the singular value decomposition of A , $A = U * S * V'$, using a divide and conquer approach. If $\text{job} = A$, all the columns of U and the rows of V' are computed. If $\text{job} = N$, no columns of U or rows of V' are computed. If $\text{job} = O$, A is overwritten with the columns of (thin) U and the rows of (thin) V' . If $\text{job} = S$, the columns of (thin) U and the rows of (thin) V' are computed and returned separately.

gesvd! ($\text{jobu}, \text{jobvt}, A$) $\rightarrow (U, S, VT)$

Finds the singular value decomposition of A , $A = U * S * V'$. If $\text{jobu} = A$, all the columns of U are computed. If $\text{jobvt} = A$ all the rows of V' are computed. If $\text{jobu} = N$, no columns of U are computed. If $\text{jobvt} = N$ no rows of V' are computed. If $\text{jobu} = O$, A is overwritten with the columns of (thin) U . If $\text{jobvt} = O$, A is overwritten with the rows of (thin) V' . If $\text{jobu} = S$, the columns of (thin) U are computed and returned separately. If $\text{jobvt} = S$ the rows of (thin) V' are computed and returned separately. jobu and jobvt can't both be O .

Returns U, S , and Vt , where S are the singular values of A .

ggsvd! ($\text{jobu}, \text{jobv}, \text{jobq}, A, B$) $\rightarrow (U, V, Q, \text{alpha}, \text{beta}, k, l, R)$

Finds the generalized singular value decomposition of A and B , $U' * A * Q = D1 * R$ and $V' * B * Q = D2 * R$. $D1$ has alpha on its diagonal and $D2$ has beta on its diagonal. If $\text{jobu} = U$, the orthogonal/unitary matrix U is computed. If $\text{jobv} = V$ the orthogonal/unitary matrix V is computed. If $\text{jobq} = Q$, the orthogonal/unitary matrix Q is computed. If jobu, jobv or jobq is N , that matrix is not computed. This function is only available in LAPACK versions prior to 3.6.0.

ggsvd3! ($\text{jobu}, \text{jobv}, \text{jobq}, A, B$) $\rightarrow (U, V, Q, \text{alpha}, \text{beta}, k, l, R)$

Finds the generalized singular value decomposition of A and B , $U' * A * Q = D1 * R$ and $V' * B * Q = D2 * R$. $D1$ has alpha on its diagonal and $D2$ has beta on its diagonal. If $\text{jobu} = U$, the orthogonal/unitary matrix U is computed. If $\text{jobv} = V$ the orthogonal/unitary matrix V is computed. If $\text{jobq} = Q$, the orthogonal/unitary matrix Q is computed. If jobu, jobv , or jobq is N , that matrix is not computed. This function requires LAPACK 3.6.0.

geevx! ($\text{balanc}, \text{jobvl}, \text{jobvr}, \text{sense}, A$) $\rightarrow (A, w, VL, VR, \text{ilo}, \text{ihi}, \text{scale}, \text{abnrm}, \text{rconde}, \text{rcondv})$

Finds the eigensystem of A with matrix balancing. If $\text{jobvl} = N$, the left eigenvectors of A aren't computed. If $\text{jobvr} = N$, the right eigenvectors of A aren't computed. If $\text{jobvl} = V$ or $\text{jobvr} = V$, the corresponding

eigenvectors are computed. If `balanc = N`, no balancing is performed. If `balanc = P`, A is permuted but not scaled. If `balanc = S`, A is scaled but not permuted. If `balanc = B`, A is permuted and scaled. If `sense = N`, no reciprocal condition numbers are computed. If `sense = E`, reciprocal condition numbers are computed for the eigenvalues only. If `sense = V`, reciprocal condition numbers are computed for the right eigenvectors only. If `sense = B`, reciprocal condition numbers are computed for the right eigenvectors and the eigenvectors. If `sense = E, B`, the right and left eigenvectors must be computed.

ggeev! (*jobvl, jobvr, A, B*) \rightarrow (*alpha, beta, vl, vr*)

Finds the generalized eigendecomposition of A and B . If `jobvl = N`, the left eigenvectors aren't computed. If `jobvr = N`, the right eigenvectors aren't computed. If `jobvl = V` or `jobvr = V`, the corresponding eigenvectors are computed.

gtsv! (*dl, d, du, B*)

Solves the equation $A * X = B$ where A is a tridiagonal matrix with `dl` on the subdiagonal, `d` on the diagonal, and `du` on the superdiagonal.

Overwrites B with the solution X and returns it.

gttrf! (*dl, d, du*) \rightarrow (*dl, d, du, du2, ipiv*)

Finds the LU factorization of a tridiagonal matrix with `dl` on the subdiagonal, `d` on the diagonal, and `du` on the superdiagonal.

Modifies `dl`, `d`, and `du` in-place and returns them and the second superdiagonal `du2` and the pivoting vector `ipiv`.

gttrs! (*trans, dl, d, du, du2, ipiv, B*)

Solves the equation $A * X = B$ (`trans = N`), $A.' * X = B$ (`trans = T`), or $A' * X = B$ (`trans = C`) using the LU factorization computed by `gttrf!`. B is overwritten with the solution X .

orglq! (*A, tau, k = length(tau)*)

Explicitly finds the matrix Q of a LQ factorization after calling `gelqf!` on A . Uses the output of `gelqf!`. A is overwritten by Q .

orgqr! (*A, tau, k = length(tau)*)

Explicitly finds the matrix Q of a QR factorization after calling `geqrf!` on A . Uses the output of `geqrf!`. A is overwritten by Q .

orgql! (*A, tau, k = length(tau)*)

Explicitly finds the matrix Q of a QL factorization after calling `geqlf!` on A . Uses the output of `geqlf!`. A is overwritten by Q .

orgrq! (*A, tau, k = length(tau)*)

Explicitly finds the matrix Q of a RQ factorization after calling `gerqf!` on A . Uses the output of `gerqf!`. A is overwritten by Q .

ormlq! (*side, trans, A, tau, C*)

Computes $Q * C$ (`trans = N`), $Q.' * C$ (`trans = T`), $Q' * C$ (`trans = C`) for `side = L` or the equivalent right-sided multiplication for `side = R` using Q from a LQ factorization of A computed using `gelqf!`. C is overwritten.

ormqr! (*side, trans, A, tau, C*)

Computes $Q * C$ (`trans = N`), $Q.' * C$ (`trans = T`), $Q' * C$ (`trans = C`) for `side = L` or the equivalent right-sided multiplication for `side = R` using Q from a QR factorization of A computed using `geqrf!`. C is overwritten.

ormql! (*side, trans, A, tau, C*)

Computes $Q * C$ (`trans = N`), $Q.' * C$ (`trans = T`), $Q' * C$ (`trans = C`) for `side = L` or the equivalent right-sided multiplication for `side = R` using Q from a QL factorization of A computed using `geqlf!`. C is overwritten.

ormrq! (*side, trans, A, tau, C*)

Computes $Q * C$ ($trans = N$), $Q.' * C$ ($trans = T$), $Q' * C$ ($trans = C$) for $side = L$ or the equivalent right-sided multiplication for $side = R$ using Q from a RQ factorization of A computed using `gerqf!`. C is overwritten.

gemqrt! (*side, trans, V, T, C*)

Computes $Q * C$ ($trans = N$), $Q.' * C$ ($trans = T$), $Q' * C$ ($trans = C$) for $side = L$ or the equivalent right-sided multiplication for $side = R$ using Q from a QR factorization of A computed using `geqrt!`. C is overwritten.

posv! (*uplo, A, B*) $\rightarrow (A, B)$

Finds the solution to $A * X = B$ where A is a symmetric or Hermitian positive definite matrix. If $uplo = U$ the upper Cholesky decomposition of A is computed. If $uplo = L$ the lower Cholesky decomposition of A is computed. A is overwritten by its Cholesky decomposition. B is overwritten with the solution X .

potrf! (*uplo, A*)

Computes the Cholesky (upper if $uplo = U$, lower if $uplo = L$) decomposition of positive-definite matrix A . A is overwritten and returned with an info code.

potri! (*uplo, A*)

Computes the inverse of positive-definite matrix A after calling `potrf!` to find its (upper if $uplo = U$, lower if $uplo = L$) Cholesky decomposition.

A is overwritten by its inverse and returned.

potrs! (*uplo, A, B*)

Finds the solution to $A * X = B$ where A is a symmetric or Hermitian positive definite matrix whose Cholesky decomposition was computed by `potrf!`. If $uplo = U$ the upper Cholesky decomposition of A was computed. If $uplo = L$ the lower Cholesky decomposition of A was computed. B is overwritten with the solution X .

pstrf! (*uplo, A, tol*) $\rightarrow (A, piv, rank, info)$

Computes the (upper if $uplo = U$, lower if $uplo = L$) pivoted Cholesky decomposition of positive-definite matrix A with a user-set tolerance tol . A is overwritten by its Cholesky decomposition.

Returns A , the pivots piv , the rank of A , and an $info$ code. If $info = 0$, the factorization succeeded. If $info = i > 0$, then A is indefinite or rank-deficient.

ptsv! (*D, E, B*)

Solves $A * X = B$ for positive-definite tridiagonal A . D is the diagonal of A and E is the off-diagonal. B is overwritten with the solution X and returned.

pttrf! (*D, E*)

Computes the LDLt factorization of a positive-definite tridiagonal matrix with D as diagonal and E as off-diagonal. D and E are overwritten and returned.

pttrs! (*D, E, B*)

Solves $A * X = B$ for positive-definite tridiagonal A with diagonal D and off-diagonal E after computing A 's LDLt factorization using `pttrf!`. B is overwritten with the solution X .

trtri! (*uplo, diag, A*)

Finds the inverse of (upper if $uplo = U$, lower if $uplo = L$) triangular matrix A . If $diag = N$, A has non-unit diagonal elements. If $diag = U$, all diagonal elements of A are one. A is overwritten with its inverse.

trtrs! (*uplo, trans, diag, A, B*)

Solves $A * X = B$ ($trans = N$), $A.' * X = B$ ($trans = T$), or $A' * X = B$ ($trans = C$) for (upper if $uplo = U$, lower if $uplo = L$) triangular matrix A . If $diag = N$, A has non-unit diagonal elements. If $diag = U$, all diagonal elements of A are one. B is overwritten with the solution X .

trcon! (*norm, uplo, diag, A*)

Finds the reciprocal condition number of (upper if $uplo = U$, lower if $uplo = L$) triangular matrix A . If

`diag = N`, A has non-unit diagonal elements. If `diag = U`, all diagonal elements of A are one. If `norm = I`, the condition number is found in the infinity norm. If `norm = O` or `1`, the condition number is found in the one norm.

trrevc! (*side, howmny, select, T, VL = similar(T), VR = similar(T)*)

Finds the eigensystem of an upper triangular matrix T . If `side = R`, the right eigenvectors are computed. If `side = L`, the left eigenvectors are computed. If `side = B`, both sets are computed. If `howmny = A`, all eigenvectors are found. If `howmny = B`, all eigenvectors are found and backtransformed using VL and VR . If `howmny = S`, only the eigenvectors corresponding to the values in `select` are computed.

trrfs! (*uplo, trans, diag, A, B, X, Ferr, Berr*) \rightarrow (*Ferr, Berr*)

Estimates the error in the solution to $A * X = B$ (`trans = N`), $A.' * X = B$ (`trans = T`), $A' * X = B$ (`trans = C`) for `side = L`, or the equivalent equations a right-handed `side = R` $X * A$ after computing X using `trtrs!`. If `uplo = U`, A is upper triangular. If `uplo = L`, A is lower triangular. If `diag = N`, A has non-unit diagonal elements. If `diag = U`, all diagonal elements of A are one. *Ferr* and *Berr* are optional inputs. *Ferr* is the forward error and *Berr* is the backward error, each component-wise.

stev! (*job, dv, ev*) \rightarrow (*dv, Zmat*)

Computes the eigensystem for a symmetric tridiagonal matrix with *dv* as diagonal and *ev* as off-diagonal. If `job = N` only the eigenvalues are found and returned in *dv*. If `job = V` then the eigenvectors are also found and returned in *Zmat*.

stebz! (*range, order, vl, vu, il, iu, abstol, dv, ev*) \rightarrow (*dv, iblock, isplit*)

Computes the eigenvalues for a symmetric tridiagonal matrix with *dv* as diagonal and *ev* as off-diagonal. If `range = A`, all the eigenvalues are found. If `range = V`, the eigenvalues in the half-open interval $(vl, vu]$ are found. If `range = I`, the eigenvalues with indices between *il* and *iu* are found. If `order = B`, eigenvalues are ordered within a block. If `order = E`, they are ordered across all the blocks. *abstol* can be set as a tolerance for convergence.

stegr! (*jobz, range, dv, ev, vl, vu, il, iu*) \rightarrow (*w, Z*)

Computes the eigenvalues (`jobz = N`) or eigenvalues and eigenvectors (`jobz = V`) for a symmetric tridiagonal matrix with *dv* as diagonal and *ev* as off-diagonal. If `range = A`, all the eigenvalues are found. If `range = V`, the eigenvalues in the half-open interval $(vl, vu]$ are found. If `range = I`, the eigenvalues with indices between *il* and *iu* are found. The eigenvalues are returned in *w* and the eigenvectors in *Z*.

stein! (*dv, ev_in, w_in, iblock_in, isplit_in*)

Computes the eigenvectors for a symmetric tridiagonal matrix with *dv* as diagonal and *ev_in* as off-diagonal. *w_in* specifies the input eigenvalues for which to find corresponding eigenvectors. *iblock_in* specifies the submatrices corresponding to the eigenvalues in *w_in*. *isplit_in* specifies the splitting points between the submatrix blocks.

syconv! (*uplo, A, ipiv*) \rightarrow (*A, work*)

Converts a symmetric matrix A (which has been factorized into a triangular matrix) into two matrices L and D . If `uplo = U`, A is upper triangular. If `uplo = L`, it is lower triangular. *ipiv* is the pivot vector from the triangular factorization. A is overwritten by L and D .

sysv! (*uplo, A, B*) \rightarrow (*B, A, ipiv*)

Finds the solution to $A * X = B$ for symmetric matrix A . If `uplo = U`, the upper half of A is stored. If `uplo = L`, the lower half is stored. B is overwritten by the solution X . A is overwritten by its Bunch-Kaufman factorization. *ipiv* contains pivoting information about the factorization.

sytrf! (*uplo, A*) \rightarrow (*A, ipiv, info*)

Computes the Bunch-Kaufman factorization of a symmetric matrix A . If `uplo = U`, the upper half of A is stored. If `uplo = L`, the lower half is stored.

Returns A , overwritten by the factorization, a pivot vector *ipiv*, and the error code *info* which is a non-negative integer. If *info* is positive the matrix is singular and the diagonal part of the factorization is exactly zero at position *info*.

sytri! (*uplo*, *A*, *ipiv*)

Computes the inverse of a symmetric matrix *A* using the results of `sytrf!`. If *uplo* = `U`, the upper half of *A* is stored. If *uplo* = `L`, the lower half is stored. *A* is overwritten by its inverse.

sytrs! (*uplo*, *A*, *ipiv*, *B*)

Solves the equation $A * X = B$ for a symmetric matrix *A* using the results of `sytrf!`. If *uplo* = `U`, the upper half of *A* is stored. If *uplo* = `L`, the lower half is stored. *B* is overwritten by the solution *X*.

hesv! (*uplo*, *A*, *B*) → (*B*, *A*, *ipiv*)

Finds the solution to $A * X = B$ for Hermitian matrix *A*. If *uplo* = `U`, the upper half of *A* is stored. If *uplo* = `L`, the lower half is stored. *B* is overwritten by the solution *X*. *A* is overwritten by its Bunch-Kaufman factorization. *ipiv* contains pivoting information about the factorization.

hetrf! (*uplo*, *A*) → (*A*, *ipiv*, *info*)

Computes the Bunch-Kaufman factorization of a Hermitian matrix *A*. If *uplo* = `U`, the upper half of *A* is stored. If *uplo* = `L`, the lower half is stored.

Returns *A*, overwritten by the factorization, a pivot vector *ipiv*, and the error code *info* which is a non-negative integer. If *info* is positive the matrix is singular and the diagonal part of the factorization is exactly zero at position *info*.

hetri! (*uplo*, *A*, *ipiv*)

Computes the inverse of a Hermitian matrix *A* using the results of `sytrf!`. If *uplo* = `U`, the upper half of *A* is stored. If *uplo* = `L`, the lower half is stored. *A* is overwritten by its inverse.

hetrs! (*uplo*, *A*, *ipiv*, *B*)

Solves the equation $A * X = B$ for a Hermitian matrix *A* using the results of `sytrf!`. If *uplo* = `U`, the upper half of *A* is stored. If *uplo* = `L`, the lower half is stored. *B* is overwritten by the solution *X*.

syev! (*jobz*, *uplo*, *A*)

Finds the eigenvalues (*jobz* = `N`) or eigenvalues and eigenvectors (*jobz* = `V`) of a symmetric matrix *A*. If *uplo* = `U`, the upper triangle of *A* is used. If *uplo* = `L`, the lower triangle of *A* is used.

syevr! (*jobz*, *range*, *uplo*, *A*, *vl*, *vu*, *il*, *iu*, *abstol*) → (*W*, *Z*)

Finds the eigenvalues (*jobz* = `N`) or eigenvalues and eigenvectors (*jobz* = `V`) of a symmetric matrix *A*. If *uplo* = `U`, the upper triangle of *A* is used. If *uplo* = `L`, the lower triangle of *A* is used. If *range* = `A`, all the eigenvalues are found. If *range* = `V`, the eigenvalues in the half-open interval (*vl*, *vu*] are found. If *range* = `I`, the eigenvalues with indices between *il* and *iu* are found. *abstol* can be set as a tolerance for convergence.

The eigenvalues are returned in *W* and the eigenvectors in *Z*.

sygvd! (*jobz*, *range*, *uplo*, *A*, *vl*, *vu*, *il*, *iu*, *abstol*) → (*w*, *A*, *B*)

Finds the generalized eigenvalues (*jobz* = `N`) or eigenvalues and eigenvectors (*jobz* = `V`) of a symmetric matrix *A* and symmetric positive-definite matrix *B*. If *uplo* = `U`, the upper triangles of *A* and *B* are used. If *uplo* = `L`, the lower triangles of *A* and *B* are used. If *itype* = 1, the problem to solve is $A * x = \lambda * B * x$. If *itype* = 2, the problem to solve is $A * B * x = \lambda * x$. If *itype* = 3, the problem to solve is $B * A * x = \lambda * x$.

bdsqr! (*uplo*, *d*, *e_*, *Vt*, *U*, *C*) → (*d*, *Vt*, *U*, *C*)

Computes the singular value decomposition of a bidiagonal matrix with *d* on the diagonal and *e_* on the off-diagonal. If *uplo* = `U`, *e_* is the superdiagonal. If *uplo* = `L`, *e_* is the subdiagonal. Can optionally also compute the product $Q' * C$.

Returns the singular values in *d*, and the matrix *C* overwritten with $Q' * C$.

bdsdc! (*uplo*, *compq*, *d*, *e_*) → (*d*, *e*, *u*, *vt*, *q*, *iq*)

Computes the singular value decomposition of a bidiagonal matrix with *d* on the diagonal and *e_* on the off-diagonal using a divide and conquer method. If *uplo* = `U`, *e_* is the superdiagonal. If *uplo* = `L`, *e_* is the subdiagonal. If *compq* = `N`, only the singular values are found. If *compq* = `I`, the singular values and

vectors are found. If `compq = P`, the singular values and vectors are found in compact form. Only works for real types.

Returns the singular values in `d`, and if `compq = P`, the compact singular vectors in `iq`.

gecon! (*normtype*, *A*, *anorm*)

Finds the reciprocal condition number of matrix *A*. If `normtype = I`, the condition number is found in the infinity norm. If `normtype = 0` or `1`, the condition number is found in the one norm. *A* must be the result of `getrf!` and *anorm* is the norm of *A* in the relevant norm.

gehrd! (*ilo*, *ihi*, *A*) → (*A*, *tau*)

Converts a matrix *A* to Hessenberg form. If *A* is balanced with `gebal!` then *ilo* and *ihi* are the outputs of `gebal!`. Otherwise they should be `ilo = 1` and `ihi = size(A, 2)`. *tau* contains the elementary reflectors of the factorization.

orghr! (*ilo*, *ihi*, *A*, *tau*)

Explicitly finds *Q*, the orthogonal/unitary matrix from `gehrd!`. *ilo*, *ihi*, *A*, and *tau* must correspond to the input/output to `gehrd!`.

gees! (*jobvs*, *A*) → (*A*, *vs*, *w*)

Computes the eigenvalues (`jobvs = N`) or the eigenvalues and Schur vectors (`jobvs = V`) of matrix *A*. *A* is overwritten by its Schur form.

Returns *A*, *vs* containing the Schur vectors, and *w*, containing the eigenvalues.

gges! (*jobvsl*, *jobvsr*, *A*, *B*) → (*A*, *B*, *alpha*, *beta*, *vsl*, *vsr*)

Computes the generalized eigenvalues, generalized Schur form, left Schur vectors (`jobvsl = V`), or right Schur vectors (`jobvsr = V`) of *A* and *B*.

The generalized eigenvalues are returned in *alpha* and *beta*. The left Schur vectors are returned in *vsl* and the right Schur vectors are returned in *vsr*.

trexc! (*compq*, *ifst*, *ilst*, *T*, *Q*) → (*T*, *Q*)

Reorder the Schur factorization of a matrix. If `compq = V`, the Schur vectors *Q* are reordered. If `compq = N` they are not modified. *ifst* and *ilst* specify the reordering of the vectors.

trsen! (*compq*, *job*, *select*, *T*, *Q*) → (*T*, *Q*, *w*)

Reorder the Schur factorization of a matrix and optionally finds reciprocal condition numbers. If `job = N`, no condition numbers are found. If `job = E`, only the condition number for this cluster of eigenvalues is found. If `job = V`, only the condition number for the invariant subspace is found. If `job = B` then the condition numbers for the cluster and subspace are found. If `compq = V` the Schur vectors *Q* are updated. If `compq = N` the Schur vectors are not modified. *select* determines which eigenvalues are in the cluster.

Returns *T*, *Q*, and reordered eigenvalues in *w*.

tgse! (*select*, *S*, *T*, *Q*, *Z*) → (*S*, *T*, *alpha*, *beta*, *Q*, *Z*)

Reorders the vectors of a generalized Schur decomposition. *select* specifies the eigenvalues in each cluster.

trsyl! (*transa*, *transb*, *A*, *B*, *C*, *isgn=1*) → (*C*, *scale*)

Solves the Sylvester matrix equation $A * X +/- X * B = scale * C$ where *A* and *B* are both quasi-upper triangular. If `transa = N`, *A* is not modified. If `transa = T`, *A* is transposed. If `transa = C`, *A* is conjugate transposed. Similarly for `transb` and *B*. If `isgn = 1`, the equation $A * X + X * B = scale * C$ is solved. If `isgn = -1`, the equation $A * X - X * B = scale * C$ is solved.

Returns *X* (overwriting *C*) and *scale*.

Constants

nothing

The singleton instance of type `Void`, used by convention when there is no value to return (as in a C `void` function). Can be converted to an empty `Nullable` value.

PROGRAM_FILE

A string containing the script name passed to Julia from the command line. Note that the script name remains unchanged from within included files. Alternatively see `@__FILE__` (page 558).

ARGS

An array of the command line arguments passed to Julia, as strings.

C_NULL

The C null pointer constant, sometimes used when calling external code.

VERSION

A `VersionNumber` object describing which version of Julia is in use. For details see *Version Number Literals* (page 53).

LOAD_PATH

An array of paths (as strings) where the `require` function looks for code.

JULIA_HOME

A string containing the full path to the directory containing the `julia` executable.

ANY

Equivalent to `Any` for dispatch purposes, but signals the compiler to skip code generation specialization for that field.

Sys.CPU_CORES

The number of logical CPU cores available in the system.

See the `Hwloc.jl` package for extended information, including number of physical cores.

Sys.WORD_SIZE

Standard word size on the current machine, in bits.

Sys.KERNEL

A symbol representing the name of the operating system, as returned by `uname` of the build configuration.

Sys.ARCH

A symbol representing the architecture of the build configuration.

Sys.MACHINE

A string containing the build triple.

See also:

[STDIN](#) (page 561) [STDOUT](#) (page 561) [STDERR](#) (page 561) [ENV](#) (page 373) [ENDIAN_BOM](#) (page 574)
[Libc.MS_ASYNC](#) [Libc.MS_INVALIDATE](#) [Libc.MS_SYNC](#) [Libdl.DL_LOAD_PATH](#) (page 616)
[Libdl.RTLD_DEEPBIND](#) (page 615) [Libdl.RTLD_LOCAL](#) (page 615) [Libdl.RTLD_NOLOAD](#) (page 615)
[Libdl.RTLD_LAZY](#) (page 615) [Libdl.RTLD_NOW](#) (page 615) [Libdl.RTLD_GLOBAL](#) (page 615)
[Libdl.RTLD_NODELETE](#) (page 615) [Libdl.RTLD_FIRST](#) (page 615)

Filesystem

pwd () → AbstractString

Get the current working directory.

cd (dir::AbstractString=homedir())

Set the current working directory.

cd (f::Function, dir::AbstractString=homedir())

Temporarily changes the current working directory and applies function *f* before returning.

readdir (dir::AbstractString=".") → Vector{String}

Returns the files and directories in the directory *dir* (or the current working directory if not given).

walkdir (dir; topdown=true, follow_symlinks=false, onerror=throw)

The `walkdir` method returns an iterator that walks the directory tree of a directory. The iterator returns a tuple containing (rootpath, dirs, files). The directory tree can be traversed top-down or bottom-up. If `walkdir` encounters a *SystemError* (page 375) it will rethrow the error by default. A custom error handling function can be provided through `onerror` keyword argument. `onerror` is called with a *SystemError* as argument.

```
for (root, dirs, files) in walkdir(".")
    println("Directories in $root")
    for dir in dirs
        println(joinpath(root, dir)) # path to directories
    end
    println("Files in $root")
    for file in files
        println(joinpath(root, file)) # path to files
    end
end
```

mkdir (path::AbstractString, mode::Unsigned=0o777)

Make a new directory with name *path* and permissions *mode*. *mode* defaults to 0o777, modified by the current file creation mask.

mkpath (path::AbstractString, mode::Unsigned=0o777)

Create all directories in the given *path*, with permissions *mode*. *mode* defaults to 0o777, modified by the current file creation mask.

symlink (target::AbstractString, link::AbstractString)

Creates a symbolic link to *target* with the name *link*.

Note: This function raises an error under operating systems that do not support soft symbolic links, such as

Windows XP.

readlink (*path::AbstractString*) → *AbstractString*

Returns the target location a symbolic link *path* points to.

chmod (*path::AbstractString, mode::Integer; recursive::Bool=false*)

Change the permissions mode of *path* to *mode*. Only integer modes (e.g. 0o777) are currently supported. If *recursive=true* and the path is a directory all permissions in that directory will be recursively changed.

chown (*path::AbstractString, owner::Integer, group::Integer=-1*)

Change the owner and/or group of *path* to *owner* and/or *group*. If the value entered for *owner* or *group* is *-1* the corresponding ID will not change. Only integer owners and groups are currently supported.

stat (*file*)

Returns a structure whose fields contain information about the file. The fields of the structure are:

| Name | Description |
|---------|--|
| size | The size (in bytes) of the file |
| device | ID of the device that contains the file |
| inode | The inode number of the file |
| mode | The protection mode of the file |
| nlink | The number of hard links to the file |
| uid | The user id of the owner of the file |
| gid | The group id of the file owner |
| rdev | If this file refers to a device, the ID of the device it refers to |
| blksize | The file-system preferred block size for the file |
| blocks | The number of such blocks allocated |
| mtime | Unix timestamp of when the file was last modified |
| ctime | Unix timestamp of when the file was created |

lstat (*file*)

Like *stat()* (page 556), but for symbolic links gets the info for the link itself rather than the file it refers to. This function must be called on a file path rather than a file object or a file descriptor.

ctime (*file*)

Equivalent to *stat(file).ctime*

mtime (*file*)

Equivalent to *stat(file).mtime*.

filemode (*file*)

Equivalent to *stat(file).mode*

filesize (*path...*)

Equivalent to *stat(file).size*.

uperm (*file*)

Gets the permissions of the owner of the file as a bitfield of

| Value | Description |
|-------|--------------------|
| 01 | Execute Permission |
| 02 | Write Permission |
| 04 | Read Permission |

For allowed arguments, see *stat()* (page 556).

gperm (*file*)

Like *uperm()* (page 556) but gets the permissions of the group owning the file.

operm (*file*)

Like *uperm* () (page 556) but gets the permissions for people who neither own the file nor are a member of the group owning the file

cp (*src::AbstractString*, *dst::AbstractString*; *remove_destination::Bool=false*, *follow_symlinks::Bool=false*)

Copy the file, link, or directory from *src* to *dst*. *remove_destination=true* will first remove an existing *dst*.

If *follow_symlinks=false*, and *src* is a symbolic link, *dst* will be created as a symbolic link. If *follow_symlinks=true* and *src* is a symbolic link, *dst* will be a copy of the file or directory *src* refers to.

download (*url::AbstractString* [, *localfile::AbstractString*])

Download a file from the given *url*, optionally renaming it to the given local file name. Note that this function relies on the availability of external tools such as *curl*, *wget* or *fetch* to download the file and is provided for convenience. For production use or situations in which more options are needed, please use a package that provides the desired functionality instead.

mv (*src::AbstractString*, *dst::AbstractString*; *remove_destination::Bool=false*)

Move the file, link, or directory from *src* to *dst*. *remove_destination=true* will first remove an existing *dst*.

rm (*path::AbstractString*; *force::Bool=false*, *recursive::Bool=false*)

Delete the file, link, or empty directory at the given path. If *force=true* is passed, a non-existing path is not treated as error. If *recursive=true* is passed and the path is a directory, then all contents are removed recursively.

touch (*path::AbstractString*)

Update the last-modified timestamp on a file to the current time.

tempname ()

Generate a unique temporary file path.

tempdir ()

Obtain the path of a temporary directory (possibly shared with other processes).

mktemp (*parent=tempdir*())

Returns (*path*, *io*), where *path* is the path of a new temporary file in *parent* and *io* is an open file object for this path.

mktemp (*f::Function*, *parent=tempdir*())

Apply the function *f* to the result of *mktemp* (*parent*) and remove the temporary file upon completion.

mktempdir (*parent=tempdir*())

Create a temporary directory in the *parent* directory and return its path.

mktempdir (*f::Function*, *parent=tempdir*())

Apply the function *f* to the result of *mktempdir* (*parent*) and remove the temporary directory upon completion.

isblockdev (*path*) → Bool

Returns true if *path* is a block device, false otherwise.

ischardev (*path*) → Bool

Returns true if *path* is a character device, false otherwise.

isdir (*path*) → Bool

Returns true if *path* is a directory, false otherwise.

isfifo (*path*) → Bool

Returns true if *path* is a FIFO, false otherwise.

isfile (*path*) → Bool

Returns true if *path* is a regular file, false otherwise.

islink (*path*) → Bool

Returns true if *path* is a symbolic link, false otherwise.

ismount (*path*) → Bool

Returns true if *path* is a mount point, false otherwise.

ispath (*path*) → Bool

Returns true if *path* is a valid filesystem path, false otherwise.

issetgid (*path*) → Bool

Returns true if *path* has the setgid flag set, false otherwise.

issetuid (*path*) → Bool

Returns true if *path* has the setuid flag set, false otherwise.

issocket (*path*) → Bool

Returns true if *path* is a socket, false otherwise.

issticky (*path*) → Bool

Returns true if *path* has the sticky bit set, false otherwise.

homedir () → AbstractString

Return the current user's home directory.

dirname (*path::AbstractString*) → AbstractString

Get the directory part of a path.

```
julia> dirname("/home/myuser")
"/home"
```

basename (*path::AbstractString*) → AbstractString

Get the file name part of a path.

```
julia> basename("/home/myuser/example.jl")
"example.jl"
```

@__FILE__ () → AbstractString

@__FILE__ expands to a string with the absolute file path of the file containing the macro. Returns nothing if run from a REPL or an empty string if evaluated by `julia -e <expr>`. Alternatively see [PROGRAM_FILE](#) (page 553).

@__DIR__ () → AbstractString

@__DIR__ expands to a string with the directory part of the absolute path of the file containing the macro. Returns nothing if run from a REPL or an empty string if evaluated by `julia -e <expr>`.

@__LINE__ () → Int

@__LINE__ expands to the line number of the call-site.

isabspath (*path::AbstractString*) → Bool

Determines whether a path is absolute (begins at the root directory).

```
julia> isabspath("/home")
true

julia> isabspath("home")
false
```

isdirpath (*path::AbstractString*) → Bool

Determines whether a path refers to a directory (for example, ends with a path separator).

```
julia> isdirpath("/home")
false

julia> isdirpath("/home/")
true
```

joinpath (*parts...*) → AbstractString

Join path components into a full path. If some argument is an absolute path, then prior components are dropped.

```
julia> joinpath("/home/myuser", "example.jl")
"/home/myuser/example.jl"
```

abspath (*path::AbstractString*) → AbstractString

Convert a path to an absolute path by adding the current directory if necessary.

normpath (*path::AbstractString*) → AbstractString

Normalize a path, removing "." and ".." entries.

```
julia> normpath("/home/myuser/./example.jl")
"/home/example.jl"
```

realpath (*path::AbstractString*) → AbstractString

Canonicalize a path by expanding symbolic links and removing "." and ".." entries.

relpath (*path::AbstractString, startpath::AbstractString = "."*) → AbstractString

Return a relative filepath to *path* either from the current directory or from an optional start directory. This is a path computation: the filesystem is not accessed to confirm the existence or nature of *path* or *startpath*.

expanduser (*path::AbstractString*) → AbstractString

On Unix systems, replace a tilde character at the start of a path with the current user's home directory.

splitdir (*path::AbstractString*) → (AbstractString, AbstractString)

Split a path into a tuple of the directory name and file name.

```
julia> splitdir("/home/myuser")
("/home", "myuser")
```

splitdrive (*path::AbstractString*) → (AbstractString, AbstractString)

On Windows, split a path into the drive letter part and the path part. On Unix systems, the first component is always the empty string.

splitext (*path::AbstractString*) → (AbstractString, AbstractString)

If the last component of a path contains a dot, split the path into everything before the dot and everything including and after the dot. Otherwise, return a tuple of the argument unmodified and the empty string.

```
julia> splitext("/home/myuser/example.jl")
("/home/myuser/example", ".jl")

julia> splitext("/home/myuser/example")
("/home/myuser/example", "")
```

I/O and Network

50.1 General I/O

STDOUT

Global variable referring to the standard out stream.

STDERR

Global variable referring to the standard error stream.

STDIN

Global variable referring to the standard input stream.

open (*filename*::AbstractString[, *read*::Bool, *write*::Bool, *create*::Bool, *truncate*::Bool, *append*::Bool]) → IOStream

Open a file in a mode specified by five boolean arguments. The default is to open files for reading only. Returns a stream for accessing the file.

open (*filename*::AbstractString[, *mode*::AbstractString]) → IOStream

Alternate syntax for open, where a string-based mode specifier is used instead of the five booleans. The values of mode correspond to those from `fopen(3)` or Perl `open`, and are equivalent to setting the following boolean groups:

| Mode | Description |
|------|-------------------------------|
| r | read |
| r+ | read, write |
| w | write, create, truncate |
| w+ | read, write, create, truncate |
| a | write, create, append |
| a+ | read, write, create, append |

open (*command*, *mode*::AbstractString="r", *stdio*=DevNull)

Start running *command* asynchronously, and return a tuple (*stream*, *process*). If mode is "r", then *stream* reads from the process's standard output and *stdio* optionally specifies the process's standard input stream. If mode is "w", then *stream* writes to the process's standard input and *stdio* optionally specifies the process's standard output stream.

open (*f*::Function, *command*, *mode*::AbstractString="r", *stdio*=DevNull)

Similar to `open(command, mode, stdio)`, but calls `f(stream)` on the resulting read or write stream, then closes the stream and waits for the process to complete. Returns the value returned by *f*.

open (*f*::Function, *args*...)

Apply the function *f* to the result of `open(args...)` and close the resulting file descriptor upon completion.

Example: `open(readstring, "file.txt")`

IOBuffer () → IOBuffer

Create an in-memory I/O stream.

IOBuffer (size::Int)

Create a fixed size IOBuffer. The buffer will not grow dynamically.

IOBuffer (string::String)

Create a read-only IOBuffer on the data underlying the given string.

IOBuffer ([data][, readable::Bool=true, writable::Bool=true[, maxsize::Int=typemax(Int)]])

Create an IOBuffer, which may optionally operate on a pre-existing array. If the readable/writable arguments are given, they restrict whether or not the buffer may be read from or written to respectively. By default the buffer is readable but not writable. The last argument optionally specifies a size beyond which the buffer may not be grown.

take! (b::IOBuffer)

Obtain the contents of an IOBuffer as an array, without copying. Afterwards, the IOBuffer is reset to its initial state.

fdio ([name::AbstractString], fd::Integer[, own::Bool=false]) → IOStream

Create an IOStream object from an integer file descriptor. If own is true, closing this object will close the underlying descriptor. By default, an IOStream is closed when it is garbage collected. name allows you to associate the descriptor with a named file.

flush (stream)

Commit all currently buffered writes to the given stream.

close (stream)

Close an I/O stream. Performs a `flush()` (page 562) first.

write (stream::IO, x)

write (filename::AbstractString, x)

Write the canonical binary representation of a value to the given I/O stream or file. Returns the number of bytes written into the stream.

You can write multiple values with the same write call. i.e. the following are equivalent:

```
write(stream, x, y...)
write(stream, x) + write(stream, y...)
```

read (stream::IO, T)

Read a single value of type T from stream, in canonical binary representation.

read (stream::IO, T, dims)

Read a series of values of type T from stream, in canonical binary representation. dims is either a tuple or a series of integer arguments specifying the size of the Array{T} to return.

read! (stream::IO, array::Union{Array, BitArray})

read! (filename::AbstractString, array::Union{Array, BitArray})

Read binary data from an I/O stream or file, filling in array.

readbytes! (stream::IOStream, b::AbstractVector{UInt8}, nb=length(b); all::Bool=true)

Read at most nb bytes from stream into b, returning the number of bytes read. The size of b will be increased if needed (i.e. if nb is greater than length(b) and enough bytes could be read), but it will never be decreased.

See `read()` (page 562) for a description of the all option.

readbytes! (stream::IO, b::AbstractVector{UInt8}, nb=length(b))

Read at most nb bytes from stream into b, returning the number of bytes read. The size of b will be increased if needed (i.e. if nb is greater than length(b) and enough bytes could be read), but it will never be decreased.

read (*s::IO*, *nb=typemax{Int}*)

Read at most *nb* bytes from *s*, returning a `Vector{UInt8}` of the bytes read.

read (*s::IOStream*, *nb::Integer*; *all=true*)

Read at most *nb* bytes from *s*, returning a `Vector{UInt8}` of the bytes read.

If *all* is `true` (the default), this function will block repeatedly trying to read all requested bytes, until an error or end-of-file occurs. If *all* is `false`, at most one `read` call is performed, and the amount of data returned is device-dependent. Note that not all stream types support the *all* option.

read (*filename::AbstractString*, *args...*)

Open a file and read its contents. *args* is passed to `read`: this is equivalent to `open(io->read(io, args...), filename)`.

unsafe_read (*io::IO*, *ref*, *nbytes::UInt*)

Copy *nbytes* from the IO stream object into *ref* (converted to a pointer).

It is recommended that subtypes `T<:IO` override the following method signature to provide more efficient implementations: `unsafe_read(s::T, p::Ptr{UInt8}, n::UInt)`

unsafe_write (*io::IO*, *ref*, *nbytes::UInt*)

Copy *nbytes* from *ref* (converted to a pointer) into the IO object.

It is recommended that subtypes `T<:IO` override the following method signature to provide more efficient implementations: `unsafe_write(s::T, p::Ptr{UInt8}, n::UInt)`

position (*s*)

Get the current position of a stream.

seek (*s*, *pos*)

Seek a stream to the given position.

seekstart (*s*)

Seek a stream to its beginning.

seekend (*s*)

Seek a stream to its end.

skip (*s*, *offset*)

Seek a stream relative to the current position.

mark (*s*)

Add a mark at the current position of stream *s*. Returns the marked position.

See also `unmark()` (page 563), `reset()` (page 563), `ismarked()` (page 563).

unmark (*s*)

Remove a mark from stream *s*. Returns `true` if the stream was marked, `false` otherwise.

See also `mark()` (page 563), `reset()` (page 563), `ismarked()` (page 563).

reset (*s*)

Reset a stream *s* to a previously marked position, and remove the mark. Returns the previously marked position. Throws an error if the stream is not marked.

See also `mark()` (page 563), `unmark()` (page 563), `ismarked()` (page 563).

ismarked (*s*)

Returns `true` if stream *s* is marked.

See also `mark()` (page 563), `unmark()` (page 563), `reset()` (page 563).

eof (*stream*) \rightarrow Bool

Tests whether an I/O stream is at end-of-file. If the stream is not yet exhausted, this function will block to wait

for more data if necessary, and then return `false`. Therefore it is always safe to read one byte after seeing `eof` return `false`. `eof` will return `false` as long as buffered data is still available, even if the remote end of a connection is closed.

isreadonly (*stream*) → Bool

Determine whether a stream is read-only.

iswritable (*io*) → Bool

Returns `true` if the specified IO object is writable (if that can be determined).

isreadable (*io*) → Bool

Returns `true` if the specified IO object is readable (if that can be determined).

isopen (*object*) → Bool

Determine whether an object - such as a stream, timer, or mmap - is not yet closed. Once an object is closed, it will never produce a new event. However, a closed stream may still have data to read in its buffer, use `eof()` (page 563) to check for the ability to read data. Use `poll_fd()` (page 573) to be notified when a stream might be writable or readable.

serialize (*stream*, *value*)

Write an arbitrary value to a stream in an opaque format, such that it can be read back by `deserialize()` (page 564). The read-back value will be as identical as possible to the original. In general, this process will not work if the reading and writing are done by different versions of Julia, or an instance of Julia with a different system image. Ptr values are serialized as all-zero bit patterns (NULL).

deserialize (*stream*)

Read a value written by `serialize()` (page 564). `deserialize` assumes the binary data read from *stream* is correct and has been serialized by a compatible implementation of `serialize()` (page 564). It has been designed with simplicity and performance as a goal and does not validate the data read. Malformed data can result in process termination. The caller has to ensure the integrity and correctness of data read from *stream*.

print_shortest (*io*, *x*)

Print the shortest possible representation, with the minimum number of consecutive non-zero digits, of number *x*, ensuring that it would parse to the exact same number.

fd (*stream*)

Returns the file descriptor backing the stream or file. Note that this function only applies to synchronous `File`'s and `IOStream`'s not to any of the asynchronous streams.

redirect_stdout ([*stream*]) → (rd, wr)

Create a pipe to which all C and Julia level `STDOUT` (page 561) output will be redirected. Returns a tuple (rd, wr) representing the pipe ends. Data written to `STDOUT` (page 561) may now be read from the rd end of the pipe. The wr end is given for convenience in case the old `STDOUT` (page 561) object was cached by the user and needs to be replaced elsewhere.

Note: *stream* must be a TTY, a Pipe, or a TCPSocket.

redirect_stdout (*f::Function*, *stream*)

Run the function *f* while redirecting `STDOUT` (page 561) to *stream*. Upon completion, `STDOUT` (page 561) is restored to its prior setting.

Note: *stream* must be a TTY, a Pipe, or a TCPSocket.

redirect_stderr ([*stream*]) → (rd, wr)

Like `redirect_stdout()` (page 564), but for `STDERR` (page 561).

Note: `stream` must be a TTY, a Pipe, or a TCPSocket.

redirect_stderr (*f::Function, stream*)

Run the function `f` while redirecting `STDERR` (page 561) to `stream`. Upon completion, `STDERR` (page 561) is restored to its prior setting.

Note: `stream` must be a TTY, a Pipe, or a TCPSocket.

redirect_stdin (*[stream]*) → (rd, wr)

Like `redirect_stdout()` (page 564), but for `STDIN` (page 561). Note that the order of the return tuple is still (rd, wr), i.e. data to be read from `STDIN` (page 561) may be written to `wr`.

Note: `stream` must be a TTY, a Pipe, or a TCPSocket.

redirect_stdin (*f::Function, stream*)

Run the function `f` while redirecting `STDIN` (page 561) to `stream`. Upon completion, `STDIN` (page 561) is restored to its prior setting.

Note: `stream` must be a TTY, a Pipe, or a TCPSocket.

readchomp (*x*)

Read the entirety of `x` as a string and remove a single trailing newline. Equivalent to `chomp!(readstring(x))`.

truncate (*file, n*)

Resize the file or buffer given by the first argument to exactly `n` bytes, filling previously unallocated space with ‘\0’ if the file or buffer is grown.

skipchars (*stream, predicate; linecomment::Char*)

Advance the stream until before the first character for which `predicate` returns false. For example `skipchars(stream, isspace)` will skip all whitespace. If keyword argument `linecomment` is specified, characters from that character through the end of a line will also be skipped.

countlines (*io::IO, eol::Char=‘\n’*)

Read `io` until the end of the stream/file and count the number of lines. To specify a file pass the filename as the first argument. EOL markers other than ‘\n’ are supported by passing them as the second argument.

PipeBuffer (*data::Vector{UInt8}=UInt8[], maxsize::Int=typemax(Int)*)

An `IOBuffer` (page 561) that allows reading and performs writes by appending. Seeking and truncating are not supported. See `IOBuffer` (page 561) for the available constructors. If `data` is given, creates a `PipeBuffer` to operate on a data vector, optionally specifying a size beyond which the underlying `Array` may not be grown.

readavailable (*stream*)

Read all available data on the stream, blocking the task only if no data is available. The result is a `Vector{UInt8, 1}`.

IOContext

`IOContext` provides a mechanism for passing output configuration settings among `show()` (page 566) methods.

In short, it is an immutable dictionary that is a subclass of `IO`. It supports standard dictionary operations such as `getindex()` (page 454), and can also be used as an I/O stream.

IOContext (*io::IO, KV::Pair*)

Create an `IOContext` that wraps a given stream, adding the specified `key=>value` pair to the properties of that stream (note that `io` can itself be an `IOContext`).

- `use (key => value) in dict` to see if this particular combination is in the properties set
- `use get(dict, key, default)` to retrieve the most recent value for a particular key

The following properties are in common use:

- `compact`: Boolean specifying that small values should be printed more compactly, e.g. that numbers should be printed with fewer digits. This is set when printing array elements.
- `limit`: Boolean specifying that containers should be truncated, e.g. showing `...` in place of most elements.
- `displaysize`: A `Tuple{Int, Int}` giving the size in rows and columns to use for text output. This can be used to override the display size for called functions, but to get the size of the screen use the `displaysize` function.

IOContext (*io::IO, context::IOContext*)

Create an `IOContext` that wraps an alternate `IO()` but inherits the properties of `context`.

IOContext (*io::IO; properties...*)

The same as `IOContext(io::IO, KV::Pair)`, but accepting properties as keyword arguments.

50.2 Text I/O

show (*x*)

Write an informative text representation of a value to the current output stream. New types should overload `show(io, x)` where the first argument is a stream. The representation used by `show` generally includes Julia-specific formatting and type information.

showcompact (*x*)

Show a compact representation of a value.

This is used for printing array elements without repeating type information (which would be redundant with that printed once for the whole array), and without line breaks inside the representation of an element.

To offer a compact representation different from its standard one, a custom type should test `get(io, :compact, false)` in its normal `show` method.

showall (*x*)

Similar to `show()` (page 566), except shows all elements of arrays.

summary (*x*)

Return a string giving a brief description of a value. By default returns `string(typeof(x))`, e.g. `Int64`.

For arrays, returns a string of size and type info, e.g. `10-element Array{Int64,1}`.

```
julia> summary(1)
"Int64"

julia> summary(zeros(2))
"2-element Array{Float64,1}"
```

print (*io::IO, x*)

Write to `io` (or to the default output stream `STDOUT` (page 561) if `io` is not given) a canonical (un-decorated) text representation of a value if there is one, otherwise call `show()` (page 566). The representation used by `print` includes minimal formatting and tries to avoid Julia-specific details.

println (*io::IO*, *xs...*)

Print (using `print()` (page 566)) *xs* followed by a newline. If *io* is not supplied, prints to `STDOUT` (page 561).

print_with_color (*color::Union{Symbol, Int}* [, *io*], *strings...*)

Print strings in a color specified as a symbol.

color may take any of the values `:normal`, `:bold`, `:black`, `:blue`, `:cyan`, `:green`, `:magenta`, `:red`, `:white`, or `:yellow` or an integer between 0 and 255 inclusive. Note that not all terminals support 256 colors.

info (*msg...*; *prefix*="INFO: ")

Display an informational message. Argument *msg* is a string describing the information to be displayed. The *prefix* keyword argument can be used to override the default prepending of *msg*.

```
julia> info("hello world")
INFO: hello world
```

```
julia> info("hello world"; prefix="MY INFO: ")
MY INFO: hello world
```

warn (*msg*)

Display a warning. Argument *msg* is a string describing the warning to be displayed.

@printf ([*io::IOStream*], "*%Fmt*", *args...*)

Print *args* using `C printf()` style format specification string, with some caveats: `Inf` and `NaN` are printed consistently as `Inf` and `NaN` for flags `%a`, `%A`, `%e`, `%E`, `%f`, `%F`, `%g`, and `%G`.

Optionally, an `IOStream` may be passed as the first argument to redirect output.

Examples

```
julia> @printf("%f %F %f %F\n", Inf, Inf, NaN, NaN)
Inf Inf NaN NaN
```

@sprintf ("*%Fmt*", *args...*)

Return `@printf` formatted output as string.

Examples

```
julia> s = @sprintf "this is a %s %15.1f" "test" 34.567;

julia> println(s)
this is a test          34.6
```

sprint (*f::Function*, *args...*)

Call the given function with an I/O stream and the supplied extra arguments. Everything written to this I/O stream is returned as a string.

```
julia> sprint(showcompact, 66.66666)
"66.6667"
```

showerror (*io*, *e*)

Show a descriptive representation of an exception object.

dump (*x*)

Show every part of the representation of a value.

readstring (*stream::IO*)

readstring (*filename::AbstractString*)

Read the entire contents of an I/O stream or a file as a string. The text is assumed to be encoded in UTF-8.

readline (*stream::IO=STDIN*)

readline (*filename::AbstractString*)

Read a single line of text, including a trailing newline character (if one is reached before the end of the input), from the given I/O stream or file (defaults to `STDIN`). When reading from a file, the text is assumed to be encoded in UTF-8.

readuntil (*stream::IO, delim*)

readuntil (*filename::AbstractString, delim*)

Read a string from an I/O stream or a file, up to and including the given delimiter byte. The text is assumed to be encoded in UTF-8.

readlines (*stream::IO*)

readlines (*filename::AbstractString*)

Read all lines of an I/O stream or a file as a vector of strings. The text is assumed to be encoded in UTF-8.

eachline (*stream::IO*)

eachline (*filename::AbstractString*)

Create an iterable object that will yield each line from an I/O stream or a file. The text is assumed to be encoded in UTF-8.

readdlm (*source, delim::Char, T::Type, eol::Char; header=false, skipstart=0, skipblanks=true, use_mmap, quotes=true, dims, comments=true, comment_char='#'*)

Read a matrix from the source where each line (separated by `eol`) gives one row, with elements separated by the given delimiter. The source can be a text file, stream or byte array. Memory mapped files can be used by passing the byte array representation of the mapped segment as source.

If `T` is a numeric type, the result is an array of that type, with any non-numeric elements as `NaN` for floating-point types, or zero. Other useful values of `T` include `String`, `AbstractString`, and `Any`.

If `header` is `true`, the first row of data will be read as header and the tuple `(data_cells, header_cells)` is returned instead of only `data_cells`.

Specifying `skipstart` will ignore the corresponding number of initial lines from the input.

If `skipblanks` is `true`, blank lines in the input will be ignored.

If `use_mmap` is `true`, the file specified by `source` is memory mapped for potential speedups. Default is `true` except on Windows. On Windows, you may want to specify `true` if the file is large, and is only read once and not written to.

If `quotes` is `true`, columns enclosed within double-quote (") characters are allowed to contain new lines and column delimiters. Double-quote characters within a quoted field must be escaped with another double-quote. Specifying `dims` as a tuple of the expected rows and columns (including header, if any) may speed up reading of large files. If `comments` is `true`, lines beginning with `comment_char` and text following `comment_char` in any line are ignored.

readdlm (*source, delim::Char, eol::Char; options...*)

If all data is numeric, the result will be a numeric array. If some elements cannot be parsed as numbers, a heterogeneous array of numbers and strings is returned.

readdlm (*source, delim::Char, T::Type; options...*)

The end of line delimiter is taken as `\n`.

readdlm (*source, delim::Char; options...*)

The end of line delimiter is taken as `\n`. If all data is numeric, the result will be a numeric array. If some elements cannot be parsed as numbers, a heterogeneous array of numbers and strings is returned.

readdlm (*source, T::Type; options...*)

The columns are assumed to be separated by one or more whitespaces. The end of line delimiter is taken as `\n`.

readdlm (*source; options...*)

The columns are assumed to be separated by one or more whitespaces. The end of line delimiter is taken as

`\n`. If all data is numeric, the result will be a numeric array. If some elements cannot be parsed as numbers, a heterogeneous array of numbers and strings is returned.

writedlm (*f*, *A*, *delim*="\\t"; *opts*)

Write *A* (a vector, matrix, or an iterable collection of iterable rows) as text to *f* (either a filename string or an IO stream) using the given delimiter *delim* (which defaults to tab, but can be any printable Julia object, typically a Char or AbstractString).

For example, two vectors *x* and *y* of the same length can be written as two columns of tab-delimited text to *f* by either `writedlm(f, [x y])` or by `writedlm(f, zip(x, y))`.

readcsv (*source*, [*T::Type*]; *options...*)

Equivalent to `readdlm()` (page 568) with *delim* set to comma, and type optionally defined by *T*.

writcsv (*filename*, *A*; *opts*)

Equivalent to `writedlm()` (page 569) with *delim* set to comma.

Base64EncodePipe (*ostream*)

Returns a new write-only I/O stream, which converts any bytes written to it into base64-encoded ASCII bytes written to *ostream*. Calling `close()` (page 487) on the Base64EncodePipe stream is necessary to complete the encoding (but does not close *ostream*).

Base64DecodePipe (*istream*)

Returns a new read-only I/O stream, which decodes base64-encoded data read from *istream*.

base64encode (*writefunc*, *args...*)

base64encode (*args...*)

Given a `write()` (page 562)-like function *writefunc*, which takes an I/O stream as its first argument, `base64encode(writefunc, args...)` calls *writefunc* to write *args...* to a base64-encoded string, and returns the string. `base64encode(args...)` is equivalent to `base64encode(write, args...)`: it converts its arguments into bytes using the standard `write()` (page 562) functions and returns the base64-encoded string.

base64decode (*string*)

Decodes the base64-encoded *string* and returns a `Vector{UInt8}` of the decoded bytes.

displaysize (*io*) → (lines, columns)

Return the nominal size of the screen that may be used for rendering output to this *io* object

50.3 Multimedia I/O

Just as text output is performed by `print` and user-defined types can indicate their textual representation by overloading `show`, Julia provides a standardized mechanism for rich multimedia output (such as images, formatted text, or even audio and video), consisting of three parts:

- A function `display(x)` to request the richest available multimedia display of a Julia object *x* (with a plain-text fallback).
- Overloading `show` allows one to indicate arbitrary multimedia representations (keyed by standard MIME types) of user-defined types.
- Multimedia-capable display backends may be registered by subclassing a generic `Display` type and pushing them onto a stack of display backends via `pushdisplay`.

The base Julia runtime provides only plain-text display, but richer displays may be enabled by loading external modules or by using graphical Julia environments (such as the IPython-based IJulia notebook).

display (*x*)

display (*d::Display*, *x*)

display (*mime*, *x*)

display (*d::Display*, *mime*, *x*)

Display *x* using the topmost applicable display in the display stack, typically using the richest supported multimedia output for *x*, with plain-text `STDOUT` (page 561) output as a fallback. The `display(d, x)` variant attempts to display *x* on the given display *d* only, throwing a `MethodError` if *d* cannot display objects of this type.

There are also two variants with a *mime* argument (a MIME type string, such as `"image/png"`), which attempt to display *x* using the requested MIME type *only*, throwing a `MethodError` if this type is not supported by either the display(s) or by *x*. With these variants, one can also supply the “raw” data in the requested MIME type by passing *x::AbstractString* (for MIME types with text-based storage, such as `text/html` or `application/postscript`) or *x::Vector{UInt8}* (for binary MIME types).

redisplay (*x*)

redisplay (*d::Display*, *x*)

redisplay (*mime*, *x*)

redisplay (*d::Display*, *mime*, *x*)

By default, the `redisplay` functions simply call `display()` (page 569). However, some display backends may override `redisplay` to modify an existing display of *x* (if any). Using `redisplay` is also a hint to the backend that *x* may be redisplayed several times, and the backend may choose to defer the display until (for example) the next interactive prompt.

displayable (*mime*) → Bool

displayable (*d::Display*, *mime*) → Bool

Returns a boolean value indicating whether the given *mime* type (string) is displayable by any of the displays in the current display stack, or specifically by the display *d* in the second variant.

show (*stream*, *mime*, *x*)

The `display` functions ultimately call `show` in order to write an object *x* as a given *mime* type to a given I/O stream (usually a memory buffer), if possible. In order to provide a rich multimedia representation of a user-defined type *T*, it is only necessary to define a new `show` method for *T*, via: `show(stream, ::MIME"mime", x::T) = ...`, where *mime* is a MIME-type string and the function body calls `write` (or similar) to write that representation of *x* to *stream*. (Note that the `MIME""` notation only supports literal strings; to construct MIME types in a more flexible manner use `MIME{Symbol}("")`.)

For example, if you define a `MyImage` type and know how to write it to a PNG file, you could define a function `show(stream, ::MIME"image/png", x::MyImage) = ...` to allow your images to be displayed on any PNG-capable `Display` (such as `IJulia`). As usual, be sure to `import Base.show` in order to add new methods to the built-in Julia function `show`.

The default MIME type is `MIME"text/plain"`. There is a fallback definition for `text/plain` output that calls `show` with 2 arguments. Therefore, this case should be handled by defining a 2-argument `show(stream::IO, x::MyType)` method.

Technically, the `MIME"mime"` macro defines a singleton type for the given *mime* string, which allows us to exploit Julia’s dispatch mechanisms in determining how to display objects of any given type.

The first argument to `show` can be an `IOContext` (page 565) specifying output format properties. See `IOContext` (page 565) for details.

mimewritable (*mime*, *x*)

Returns a boolean value indicating whether or not the object *x* can be written as the given *mime* type. (By default, this is determined automatically by the existence of the corresponding `show()` (page 566) method for `typeof(x)`.)

reprmime (*mime*, *x*)

Returns an `AbstractString` or `Vector{UInt8}` containing the representation of *x* in the requested *mime* type, as written by `show` (throwing a `MethodError` if no appropriate `show` is available). An `AbstractString` is returned for MIME types with textual representations (such as `"text/html"` or

"application/postscript"), whereas binary data is returned as `Vector{UInt8}`. (The function `istextmime(mime)` returns whether or not Julia treats a given mime type as text.)

As a special case, if `x` is an `AbstractString` (for textual MIME types) or a `Vector{UInt8}` (for binary MIME types), the `reprmime` function assumes that `x` is already in the requested mime format and simply returns `x`. This special case does not apply to the "text/plain" MIME type. This is useful so that raw data can be passed to `display(m::MIME, x)`.

stringmime (*mime*, *x*)

Returns an `AbstractString` containing the representation of `x` in the requested mime type. This is similar to `reprmime()` (page 570) except that binary data is base64-encoded as an ASCII string.

As mentioned above, one can also define new display backends. For example, a module that can display PNG images in a window can register this capability with Julia, so that calling `display(x)` on types with PNG representations will automatically display the image using the module's window.

In order to define a new display backend, one should first create a subtype `D` of the abstract class `Display`. Then, for each MIME type (mime string) that can be displayed on `D`, one should define a function `display(d::D, ::MIME"mime", x) = ...` that displays `x` as that MIME type, usually by calling `reprmime(mime, x)`. A `MethodError` should be thrown if `x` cannot be displayed as that MIME type; this is automatic if one calls `reprmime`. Finally, one should define a function `display(d::D, x)` that queries `mimewritable(mime, x)` for the mime types supported by `D` and displays the "best" one; a `MethodError` should be thrown if no supported MIME types are found for `x`. Similarly, some subtypes may wish to override `redisplay(d::D, ...)`. (Again, one should import `Base.display` to add new methods to `display`.) The return values of these functions are up to the implementation (since in some cases it may be useful to return a display "handle" of some type). The display functions for `D` can then be called directly, but they can also be invoked automatically from `display(x)` simply by pushing a new display onto the display-backend stack with:

pushdisplay (*d::Display*)

Pushes a new display `d` on top of the global display-backend stack. Calling `display(x)` or `display(mime, x)` will display `x` on the topmost compatible backend in the stack (i.e., the topmost backend that does not throw a `MethodError`).

popdisplay ()

popdisplay (*d::Display*)

Pop the topmost backend off of the display-backend stack, or the topmost copy of `d` in the second variant.

TextDisplay (*io::IO*)

Returns a `TextDisplay <: Display`, which can display any object as the text/plain MIME type (only), writing the text representation to the given I/O stream. (The text representation is the same as the way an object is printed in the Julia REPL.)

istextmime (*m::MIME*)

Determine whether a MIME type is text data. MIME types are assumed to be binary data except for a set of types known to be text data (possibly Unicode).

50.4 Memory-mapped I/O

Mmap.Anonymous (*name*, *readonly*, *create*)

Create an IO-like object for creating zeroed-out mmapmed-memory that is not tied to a file for use in `Mmap.mmap`. Used by `SharedArray` for creating shared memory arrays.

Mmap.mmap (*io::Union{IOStream,AbstractString,Mmap.AnonymousMmap}*[], *type::Type{Array{T,N}}*, *dims*, *offset*; *grow::Bool=true*, *shared::Bool=true*)

Mmap.mmap (*type::Type{Array{T,N}}*, *dims*)

Create an `Array` whose values are linked to a file, using memory-mapping. This provides a convenient way of working with data too large to fit in the computer's memory.

The type is an `Array{T,N}` with a bits-type element of `T` and dimension `N` that determines how the bytes of the array are interpreted. Note that the file must be stored in binary format, and no format conversions are possible (this is a limitation of operating systems, not Julia).

`dims` is a tuple or single `Integer` specifying the size or length of the array.

The file is passed via the stream argument, either as an open `IOStream` or filename string. When you initialize the stream, use `"r"` for a “read-only” array, and `"w+"` to create a new array used to write values to disk.

If no type argument is specified, the default is `Vector{UInt8}`.

Optionally, you can specify an offset (in bytes) if, for example, you want to skip over a header in the file. The default value for the offset is the current stream position for an `IOStream`.

The `grow` keyword argument specifies whether the disk file should be grown to accommodate the requested size of array (if the total file size is $<$ requested array size). Write privileges are required to grow the file.

The `shared` keyword argument specifies whether the resulting `Array` and changes made to it will be visible to other processes mapping the same file.

For example, the following code

```
# Create a file for mmaping
# (you could alternatively use mmap to do this step, too)
A = rand(1:20, 5, 30)
s = open("/tmp/mmap.bin", "w+")
# We'll write the dimensions of the array as the first two Ints in the file
write(s, size(A,1))
write(s, size(A,2))
# Now write the data
write(s, A)
close(s)

# Test by reading it back in
s = open("/tmp/mmap.bin") # default is read-only
m = read(s, Int)
n = read(s, Int)
A2 = Mmap.mmap(s, Matrix{Int}, (m,n))
```

creates a `m-by-n Matrix{Int}`, linked to the file associated with stream `s`.

A more portable file would need to encode the word size – 32 bit or 64 bit – and endianness information in the header. In practice, consider encoding binary data using standard formats like HDF5 (which can be used with memory-mapping).

`Mmap.mmap(io, BitArray[, dims, offset])`

Create a `BitArray` whose values are linked to a file, using memory-mapping; it has the same purpose, works in the same way, and has the same arguments, as `mmap()`, but the byte representation is different.

Example: `B = Mmap.mmap(s, BitArray, (25,30000))`

This would create a 25-by-30000 `BitArray`, linked to the file associated with stream `s`.

`Mmap.sync!(array)`

Forces synchronization between the in-memory version of a memory-mapped `Array` or `BitArray` and the on-disk version.

50.5 Network I/O

connect (*[host]*, *port::Integer*) → *TCPSocket*
 Connect to the host *host* on port *port*.

connect (*path::AbstractString*) → *PipeEndpoint*
 Connect to the named pipe / UNIX domain socket at *path*.

listen (*[addr]*, *port::Integer*; *backlog::Integer=BACKLOG_DEFAULT*) → *TCPServer*
 Listen on port *port* on the address specified by *addr*. By default this listens on *localhost* only. To listen on all interfaces pass *IPv4(0)* or *IPv6(0)* as appropriate. *backlog* determines how many connections can be pending (not having called *accept()* (page 573)) before the server will begin to reject them. The default value of *backlog* is 511.

listen (*path::AbstractString*) → *PipeServer*
 Create and listen on a named pipe / UNIX domain socket.

getaddrinfo (*host::AbstractString*) → *IPAddr*
 Gets the IP address of the host (may have to do a DNS lookup)

getsockname (*sock::Union{TCPServer, TCPSocket}*) → (*IPAddr*, *UInt16*)
 Get the IP address and the port that the given *TCPSocket* is connected to (or bound to, in the case of *TCPServer*).

IPv4 (*host::Integer*) → *IPv4*
 Returns an *IPv4* object from ip address *host* formatted as an *Integer*.

IPv6 (*host::Integer*) → *IPv6*
 Returns an *IPv6* object from ip address *host* formatted as an *Integer*.

nb_available (*stream*)
 Returns the number of bytes available for reading before a read from this stream or buffer will block.

accept (*server* [, *client*])
 Accepts a connection on the given server and returns a connection to the client. An uninitialized client stream may be provided, in which case it will be used instead of creating a new stream.

listenany (*[host::IPAddr]*, *port_hint*) → (*UInt16*, *TCPServer*)
 Create a *TCPServer* on any port, using *hint* as a starting point. Returns a tuple of the actual port that the server was created on and the server itself.

poll_fd (*fd*, *timeout_s::Real=-1*; *readable=false*, *writable=false*)
 Monitor a file descriptor *fd* for changes in the read or write availability, and with a timeout given by *timeout_s* seconds.

The keyword arguments determine which of read and/or write status should be monitored; at least one of them must be set to *true*.

The returned value is an object with boolean fields *readable*, *writable*, and *timedout*, giving the result of the polling.

poll_file (*path::AbstractString*, *interval_s::Real=5.007*, *timeout_s::Real=-1*) → (*previous::StatStruct*, *current::StatStruct*)
 Monitor a file for changes by polling every *interval_s* seconds until a change occurs or *timeout_s* seconds have elapsed. The *interval_s* should be a long period; the default is 5.007 seconds.

Returns a pair of *StatStruct* objects (*previous*, *current*) when a change is detected.

To determine when a file was modified, compare *mtime(prev) != mtime(current)* to detect notification of changes. However, using *watch_file()* (page 573) for this operation is preferred, since it is more reliable and efficient, although in some situations it may not be available.

watch_file (*path::AbstractString, timeout_s::Real=-1*)

Watch file or directory *path* for changes until a change occurs or *timeout_s* seconds have elapsed.

The returned value is an object with boolean fields *changed*, *renamed*, and *timedout*, giving the result of watching the file.

This behavior of this function varies slightly across platforms. See <https://nodejs.org/api/fs.html#fs_caveats> for more detailed information.

bind (*socket::Union{UDPSocket, TCPSocket}, host::IPAddr, port::Integer; ipv6only=false, reuseaddr=false, kws...*)

Bind *socket* to the given *host:port*. Note that 0.0.0.0 will listen on all devices.

- The *ipv6only* parameter disables dual stack mode. If *ipv6only=true*, only an IPv6 stack is created.
- If *reuseaddr=true*, multiple threads or processes can bind to the same address without error if they all set *reuseaddr=true*, but only the last to bind will receive any traffic.

send (*socket::UDPSocket, host, port::Integer, msg*)

Send *msg* over *socket* to *host:port*.

recv (*socket::UDPSocket*)

Read a UDP packet from the specified socket, and return the bytes received. This call blocks.

recvfrom (*socket::UDPSocket*) → (*address, data*)

Read a UDP packet from the specified socket, returning a tuple of (*address, data*), where *address* will be either IPv4 or IPv6 as appropriate.

setopt (*sock::UDPSocket; multicast_loop = nothing, multicast_ttl=nothing, enable_broadcast=nothing, ttl=nothing*)

Set UDP socket options.

- multicast_loop*: loopback for multicast packets (default: *true*).
- multicast_ttl*: TTL for multicast packets (default: *nothing*).
- enable_broadcast*: flag must be set to *true* if socket will be used for broadcast messages, or else the UDP system will return an access error (default: *false*).
- ttl*: Time-to-live of packets sent on the socket (default: *nothing*).

ntoh (*x*)

Converts the endianness of a value from Network byte order (big-endian) to that used by the Host.

hton (*x*)

Converts the endianness of a value from that used by the Host to Network byte order (big-endian).

ltoh (*x*)

Converts the endianness of a value from Little-endian to that used by the Host.

htol (*x*)

Converts the endianness of a value from that used by the Host to Little-endian.

ENDIAN_BOM

The 32-bit byte-order-mark indicates the native byte order of the host machine. Little-endian machines will contain the value 0x04030201. Big-endian machines will contain the value 0x01020304.

Punctuation

Extended documentation for mathematical symbols & functions is [here](#) (page 405).

| symbol | meaning |
|--------|---|
| @m | invoke macro m; followed by space-separated expressions |
| ! | prefix “not” operator |
| a! () | at the end of a function name, ! indicates that a function modifies its argument(s) |
| # | begin single line comment |
| #= | begin multi-line comment (these are nestable) |
| =# | end multi-line comment |
| \$ | string and expression interpolation |
| % | remainder operator |
| ^ | exponent operator |
| & | bitwise and |
| && | short-circuiting boolean and |
| | bitwise or |
| | short-circuiting boolean or |
| | bitwise xor operator |
| * | multiply, or matrix multiply |
| () | the empty tuple |
| ~ | bitwise not operator |
| \ | backslash operator |
| ' | complex transpose operator A^H |
| a[] | array indexing |
| [,] | vertical concatenation |
| [;] | also vertical concatenation |
| [] | with space-separated expressions, horizontal concatenation |
| T{ } | parametric type instantiation |
| ; | statement separator |
| , | separate function arguments or tuple components |
| ? | 3-argument conditional operator (conditional ? if_true : if_false) |
| " " | delimit string literals |
| ' ' | delimit character literals |
| ` ` | delimit external process (command) specifications |
| ... | splice arguments into a function call or declare a varargs function or type |
| . | access named fields in objects or names inside modules, also prefixes elementwise operators |
| a:b | range a, a+1, a+2, ..., b |
| a:s:b | range a, a+s, a+2s, ..., b |

Continued on next page

Table 51.1 – continued from previous page

| symbol | meaning |
|--------|---------------------------------------|
| : | index an entire dimension (1:end) |
| :: | type annotation, depending on context |
| : () | quoted expression |
| : a | symbol a |

Sorting and Related Functions

Julia has an extensive, flexible API for sorting and interacting with already-sorted arrays of values. By default, Julia picks reasonable algorithms and sorts in standard ascending order:

```
julia> sort([2,3,1])
3-element Array{Int64,1}:
 1
 2
 3
```

You can easily sort in reverse order as well:

```
julia> sort([2,3,1], rev=true)
3-element Array{Int64,1}:
 3
 2
 1
```

To sort an array in-place, use the “bang” version of the sort function:

```
julia> a = [2,3,1];

julia> sort!(a);

julia> a
3-element Array{Int64,1}:
 1
 2
 3
```

Instead of directly sorting an array, you can compute a permutation of the array’s indices that puts the array into sorted order:

```
julia> v = randn(5)
5-element Array{Float64,1}:
 0.297288
 0.382396
-0.597634
-0.0104452
-0.839027

julia> p = sortperm(v)
5-element Array{Int64,1}:
 5
 3
```

```
4
1
2

julia> v[p]
5-element Array{Float64,1}:
-0.839027
-0.597634
-0.0104452
 0.297288
 0.382396
```

Arrays can easily be sorted according to an arbitrary transformation of their values:

```
julia> sort(v, by=abs)
5-element Array{Float64,1}:
-0.0104452
 0.297288
 0.382396
-0.597634
-0.839027
```

Or in reverse order by a transformation:

```
julia> sort(v, by=abs, rev=true)
5-element Array{Float64,1}:
-0.839027
-0.597634
 0.382396
 0.297288
-0.0104452
```

If needed, the sorting algorithm can be chosen:

```
julia> sort(v, alg=InsertionSort)
5-element Array{Float64,1}:
-0.839027
-0.597634
-0.0104452
 0.297288
 0.382396
```

All the sorting and order related functions rely on a “less than” relation defining a total order on the values to be manipulated. The `isless` function is invoked by default, but the relation can be specified via the `lt` keyword.

52.1 Sorting Functions

sort! (*v*; *alg*::Algorithm=*defalg*(*v*), *lt*=*isless*, *by*=*identity*, *rev*::Bool=*false*, *order*::Ordering=*Forward*)

Sort the vector *v* in place. `QuickSort` is used by default for numeric arrays while `MergeSort` is used for other arrays. You can specify an algorithm to use via the `alg` keyword (see [Sorting Algorithms](#) for available algorithms). The `by` keyword lets you provide a function that will be applied to each element before comparison; the `lt` keyword allows providing a custom “less than” function; use `rev=true` to reverse the sorting order. These options are independent and can be used together in all possible combinations: if both `by` and `lt` are specified, the `lt` function is applied to the result of the `by` function; `rev=true` reverses whatever ordering specified via the `by` and `lt` keywords.

sort (*v*; *alg*::Algorithm=defalg(*v*), *lt*=isless, *by*=identity, *rev*::Bool=false, *order*::Ordering=Forward)

Variant of `sort!()` (page 578) that returns a sorted copy of *v* leaving *v* itself unmodified.

sort (*A*, *dim*::Integer; *alg*::Algorithm=DEFAULT_UNSTABLE, *lt*=isless, *by*=identity, *rev*::Bool=false, *order*::Ordering=Forward, *initialized*::Bool=false)

Sort a multidimensional array *A* along the given dimension. See `sort!()` (page 578) for a description of possible keyword arguments.

sortperm (*v*; *alg*::Algorithm=DEFAULT_UNSTABLE, *lt*=isless, *by*=identity, *rev*::Bool=false, *order*::Ordering=Forward)

Return a permutation vector of indices of *v* that puts it in sorted order. Specify *alg* to choose a particular sorting algorithm (see Sorting Algorithms). MergeSort is used by default, and since it is stable, the resulting permutation will be the lexicographically first one that puts the input array into sorted order – i.e. indices of equal elements appear in ascending order. If you choose a non-stable sorting algorithm such as QuickSort, a different permutation that puts the array into order may be returned. The order is specified using the same keywords as `sort!`.

See also `sortperm!()` (page 579).

sortperm! (*ix*, *v*; *alg*::Algorithm=DEFAULT_UNSTABLE, *lt*=isless, *by*=identity, *rev*::Bool=false, *order*::Ordering=Forward, *initialized*::Bool=false)

Like `sortperm()` (page 579), but accepts a preallocated index vector *ix*. If *initialized* is false (the default), *ix* is initialized to contain the values `1:length(v)`.

sortrows (*A*; *alg*::Algorithm=DEFAULT_UNSTABLE, *lt*=isless, *by*=identity, *rev*::Bool=false, *order*::Ordering=Forward)

Sort the rows of matrix *A* lexicographically. See `sort!()` (page 578) for a description of possible keyword arguments.

sortcols (*A*; *alg*::Algorithm=DEFAULT_UNSTABLE, *lt*=isless, *by*=identity, *rev*::Bool=false, *order*::Ordering=Forward)

Sort the columns of matrix *A* lexicographically. See `sort!()` (page 578) for a description of possible keyword arguments.

52.2 Order-Related Functions

isorted (*v*, *by*=identity, *rev*::Bool=false, *order*::Ordering=Forward)

Test whether a vector is in sorted order. The *by*, *lt* and *rev* keywords modify what order is considered to be sorted just as they do for `sort()` (page 578).

searchsorted (*a*, *x*, [*by*=<transform>], [*lt*=<comparison>], [*rev*=false])

Returns the range of indices of *a* which compare as equal to *x* according to the order specified by the *by*, *lt* and *rev* keywords, assuming that *a* is already sorted in that order. Returns an empty range located at the insertion point if *a* does not contain values equal to *x*.

searchsortedfirst (*a*, *x*, [*by*=<transform>], [*lt*=<comparison>], [*rev*=false])

Returns the index of the first value in *a* greater than or equal to *x*, according to the specified order. Returns `length(a)+1` if *x* is greater than all values in *a*.

searchsortedlast (*a*, *x*, [*by*=<transform>], [*lt*=<comparison>], [*rev*=false])

Returns the index of the last value in *a* less than or equal to *x*, according to the specified order. Returns 0 if *x* is less than all values in *a*.

select! (*v*, *k*, [*by*=<transform>], [*lt*=<comparison>], [*rev*=false])

Partially sort the vector *v* in place, according to the order specified by *by*, *lt* and *rev* so that the value at index *k* (or range of adjacent values if *k* is a range) occurs at the position where it would appear if the array were fully sorted via a non-stable algorithm. If *k* is a single index, that value is returned; if *k* is a range, an array of values at those indices is returned. Note that `select!` does not fully sort the input array.

select (*v*, *k*, [*by*=<transform>], [*lt*=<comparison>], [*rev*=false])

Variant of `select!` which copies *v* before partially sorting it, thereby returning the same thing as `select!` but leaving *v* unmodified.

selectperm (*v*, *k*, [*alg*=<algorithm>], [*by*=<transform>], [*lt*=<comparison>], [*rev*=false])

Return a partial permutation of the vector *v*, according to the order specified by *by*, *lt* and *rev*, so that `v[output]` returns the first *k* (or range of adjacent values if *k* is a range) values of a fully sorted version of *v*. If *k* is a single index (Integer), an array of the first *k* indices is returned; if *k* is a range, an array of those indices is returned. Note that the handling of integer values for *k* is different from `select` in that it returns a vector of *k* elements instead of just the *k*th element. Also note that this is equivalent to, but more efficient than, calling `sortperm(...)[k]`

selectperm! (*ix*, *v*, *k*, [*alg*=<algorithm>], [*by*=<transform>], [*lt*=<comparison>], [*rev*=false], [*initialized*=false])

Like `selectperm`, but accepts a preallocated index vector *ix*. If *initialized* is false (the default), *ix* is initialized to contain the values `1:length(ix)`.

52.3 Sorting Algorithms

There are currently four sorting algorithms available in base Julia:

- `InsertionSort`
- `QuickSort`
- `PartialQuickSort(k)`
- `MergeSort`

`InsertionSort` is an $O(n^2)$ stable sorting algorithm. It is efficient for very small *n*, and is used internally by `QuickSort`.

`QuickSort` is an $O(n \log n)$ sorting algorithm which is in-place, very fast, but not stable – i.e. elements which are considered equal will not remain in the same order in which they originally appeared in the array to be sorted. `QuickSort` is the default algorithm for numeric values, including integers and floats.

`PartialQuickSort(k)` is similar to `QuickSort`, but the output array is only sorted up to index *k* if *k* is an integer, or in the range of *k* if *k* is an `OrdinalRange`. For example:

```
x = rand(1:500, 100)
k = 50
k2 = 50:100
s = sort(x; alg=QuickSort)
ps = sort(x; alg=PartialQuickSort(k))
qs = sort(x; alg=PartialQuickSort(k2))
map(issorted, (s, ps, qs))           # => (true, false, false)
map(x->issorted(x[1:k]), (s, ps, qs)) # => (true, true, false)
map(x->issorted(x[k2]), (s, ps, qs))  # => (true, false, true)
s[1:k] == ps[1:k]                     # => true
s[k2] == qs[k2]                       # => true
```

`MergeSort` is an $O(n \log n)$ stable sorting algorithm but is not in-place – it requires a temporary array of half the size of the input array – and is typically not quite as fast as `QuickSort`. It is the default algorithm for non-numeric data.

The default sorting algorithms are chosen on the basis that they are fast and stable, or *appear* to be so. For numeric types indeed, `QuickSort` is selected as it is faster and indistinguishable in this case from a stable sort (unless the array records its mutations in some way). The stability property comes at a non-negligible cost, so if you don't need it, you may want to explicitly specify your preferred algorithm, e.g. `sort!(v, alg=QuickSort)`.

The mechanism by which Julia picks default sorting algorithms is implemented via the `Base.Sort.defalg` function. It allows a particular algorithm to be registered as the default in all sorting functions for specific arrays. For example, here are the two default methods from `sort.jl`:

```
defalg(v::AbstractArray) = MergeSort
defalg{T<:Number}(v::AbstractArray{T}) = QuickSort
```

As for numeric arrays, choosing a non-stable default algorithm for array types for which the notion of a stable sort is meaningless (i.e. when two values comparing equal can not be distinguished) may make sense.

Package Manager Functions

All package manager functions are defined in the `Pkg` module. None of the `Pkg` module's functions are exported; to use them, you'll need to prefix each function call with an explicit `Pkg.`, e.g. `Pkg.status()` or `Pkg.dir()`.

Functions for package development (e.g. `tag`, `publish`, etc.) have been moved to the [PkgDev](#) package. See [PkgDev README](#) for the documentation of those functions.

dir() → `AbstractString`

Returns the absolute path of the package directory. This defaults to `joinpath(homedir(), ".julia", "v$(VERSION.major).$(VERSION.minor)")` on all platforms (i.e. `~/.julia/v0.6` in UNIX shell syntax). If the `JULIA_PKGDIR` environment variable is set, then that path is used in the returned value as `joinpath(ENV["JULIA_PKGDIR"], "v$(VERSION.major).$(VERSION.minor)")`. If `JULIA_PKGDIR` is a relative path, it is interpreted relative to whatever the current working directory is.

dir(names...) → `AbstractString`

Equivalent to `normpath(Pkg.dir(), names...)` – i.e. it appends path components to the package directory and normalizes the resulting path. In particular, `Pkg.dir(pkg)` returns the path to the package `pkg`.

init(meta::`AbstractString`=`DEFAULT_META`, branch::`AbstractString`=`META_BRANCH`)

Initialize `Pkg.dir()` as a package directory. This will be done automatically when the `JULIA_PKGDIR` is not set and `Pkg.dir()` uses its default value. As part of this process, clones a local METADATA git repository from the site and branch specified by its arguments, which are typically not provided. Explicit (non-default) arguments can be used to support a custom METADATA setup.

resolve()

Determines an optimal, consistent set of package versions to install or upgrade to. The optimal set of package versions is based on the contents of `Pkg.dir("REQUIRE")` and the state of installed packages in `Pkg.dir()`. Packages that are no longer required are moved into `Pkg.dir(".trash")`.

edit()

Opens `Pkg.dir("REQUIRE")` in the editor specified by the `VISUAL` or `EDITOR` environment variables; when the editor command returns, it runs `Pkg.resolve()` to determine and install a new optimal set of installed package versions.

add(pkg, vers...)

Add a requirement entry for `pkg` to `Pkg.dir("REQUIRE")` and call `Pkg.resolve()`. If `vers` are given, they must be `VersionNumber` objects and they specify acceptable version intervals for `pkg`.

rm(pkg)

Remove all requirement entries for `pkg` from `Pkg.dir("REQUIRE")` and call `Pkg.resolve()`.

clone(url[, pkg])

Clone a package directly from the git URL `url`. The package does not need to be registered in

`Pkg.dir("METADATA")`. The package repo is cloned by the name `pkg` if provided; if not provided, `pkg` is determined automatically from `url`.

clone (*pkg*)

If `pkg` has a URL registered in `Pkg.dir("METADATA")`, clone it from that URL on the default branch. The package does not need to have any registered versions.

setprotocol! (*proto*)

Set the protocol used to access GitHub-hosted packages. Defaults to ‘https’, with a blank `proto` delegating the choice to the package developer.

available () → Vector{String}

Returns the names of available packages.

available (*pkg*) → Vector{VersionNumber}

Returns the version numbers available for package `pkg`.

installed () → Dict{String,VersionNumber}

Returns a dictionary mapping installed package names to the installed version number of each package.

installed (*pkg*) → Void | VersionNumber

If `pkg` is installed, return the installed version number. If `pkg` is registered, but not installed, return `nothing`.

status ()

Prints out a summary of what packages are installed and what version and state they’re in.

status (*pkg*)

Prints out a summary of what version and state `pkg`, specifically, is in.

update (*pkgs...*)

Update the metadata repo – kept in `Pkg.dir("METADATA")` – then update any fixed packages that can safely be pulled from their origin; then call `Pkg.resolve()` to determine a new optimal set of packages versions.

Without arguments, updates all installed packages. When one or more package names are provided as arguments, only those packages and their dependencies are updated.

checkout (*pkg*, [*branch*="master"]; *merge*=true, *pull*=true)

Checkout the `Pkg.dir(pkg)` repo to the branch `branch`. Defaults to checking out the “master” branch. To go back to using the newest compatible released version, use `Pkg.free(pkg)`. Changes are merged (fast-forward only) if the keyword argument `merge == true`, and the latest version is pulled from the upstream repo if `pull == true`.

pin (*pkg*)

Pin `pkg` at the current version. To go back to using the newest compatible released version, use `Pkg.free(pkg)`

pin (*pkg*, *version*)

Pin `pkg` at registered version `version`.

free (*pkg*)

Free the package `pkg` to be managed by the package manager again. It calls `Pkg.resolve()` to determine optimal package versions after. This is an inverse for both `Pkg.checkout` and `Pkg.pin`.

You can also supply an iterable collection of package names, e.g., `Pkg.free(("Pkg1", "Pkg2"))` to free multiple packages at once.

build ()

Run the build scripts for all installed packages in depth-first recursive order.

build (*pkgs...*)

Run the build script in `deps/build.jl` for each package in `pkgs` and all of their dependencies in depth-first recursive order. This is called automatically by `Pkg.resolve()` on all installed or updated packages.

test (*; coverage=false*)

Run the tests for all installed packages ensuring that each package's test dependencies are installed for the duration of the test. A package is tested by running its `test/runtests.jl` file and test dependencies are specified in `test/REQUIRE`. Coverage statistics for the packages may be generated by passing `coverage=true`. The default behavior is not to run coverage.

test (*pkgs...; coverage=false*)

Run the tests for each package in `pkgs` ensuring that each package's test dependencies are installed for the duration of the test. A package is tested by running its `test/runtests.jl` file and test dependencies are specified in `test/REQUIRE`. Coverage statistics for the packages may be generated by passing `coverage=true`. The default behavior is not to run coverage.

dependents (*packagename*)

List the packages that have `packagename` as a dependency.

Dates and Time

54.1 Dates and Time Types

Period

Year

Month

Week

Day

Hour

Minute

Second

Millisecond

Period types represent discrete, human representations of time.

CompoundPeriod

A `CompoundPeriod` is useful for expressing time periods that are not a fixed multiple of smaller periods. For example, “a year and a day” is not a fixed number of days, but can be expressed using a `CompoundPeriod`. In fact, a `CompoundPeriod` is automatically generated by addition of different period types, e.g. `Year(1) + Day(1)` produces a `CompoundPeriod` result.

Instant

`Instant` types represent integer-based, machine representations of time as continuous timelines starting from an epoch.

UTInstant{T}

The `UTInstant` represents a machine timeline based on UT time (1 day = one revolution of the earth). The `T` is a `Period` parameter that indicates the resolution or precision of the instant.

TimeType

`TimeType` types wrap `Instant` machine instances to provide human representations of the machine instant. Both `DateTime` and `Date` are subtypes of `TimeType`.

DateTime

`DateTime` wraps a `UTInstant{Millisecond}` and interprets it according to the proleptic Gregorian calendar.

Date

`Date` wraps a `UTInstant{Day}` and interprets it according to the proleptic Gregorian calendar.

54.2 Dates Functions

All Dates functions are defined in the `Dates` module; note that only the `Date`, `DateTime`, and `now` functions are exported; to use all other `Dates` functions, you'll need to prefix each function call with an explicit `Dates.`, e.g. `Dates.dayofweek(dt)`. Alternatively, you can write using `Base.Dates` to bring all exported functions into `Main` to be used without the `Dates.` prefix.

DateTime ($y[, m, d, h, mi, s, ms]$) \rightarrow `DateTime`

Construct a `DateTime` type by parts. Arguments must be convertible to `Int64`.

DateTime ($periods::Period\dots$) \rightarrow `DateTime`

Construct a `DateTime` type by `Period` type parts. Arguments may be in any order. `DateTime` parts not provided will default to the value of `Dates.default(period)`.

DateTime ($f::Function, y[, m, d, h, mi, s]; step=Day(1), negate=false, limit=10000$) \rightarrow `DateTime`

Create a `DateTime` through the adjuster API. The starting point will be constructed from the provided y, m, d, \dots arguments, and will be adjusted until $f::Function$ returns `true`. The step size in adjusting can be provided manually through the `step` keyword. If `negate=true`, then the adjusting will stop when $f::Function$ returns `false` instead of `true`. `limit` provides a limit to the max number of iterations the adjustment API will pursue before throwing an error (in the case that $f::Function$ is never satisfied).

DateTime ($dt::Date$) \rightarrow `DateTime`

Converts a `Date` to a `DateTime`. The hour, minute, second, and millisecond parts of the new `DateTime` are assumed to be zero.

DateTime ($dt::AbstractString, format::AbstractString; locale="english"$) \rightarrow `DateTime`

Construct a `DateTime` by parsing the `dt` date string following the pattern given in the `format` string. The following character codes can be used to construct the `format` string:

| Code | Matches | Comment |
|----------------------|-----------|---|
| <code>y</code> | 1996, 96 | Returns year of 1996, 0096 |
| <code>Y</code> | 1996, 96 | Returns year of 1996, 0096. Equivalent to <code>y</code> |
| <code>m</code> | 1, 01 | Matches 1 or 2-digit months |
| <code>u</code> | Jan | Matches abbreviated months according to the <code>locale</code> keyword |
| <code>U</code> | January | Matches full month names according to the <code>locale</code> keyword |
| <code>d</code> | 1, 01 | Matches 1 or 2-digit days |
| <code>H</code> | 00 | Matches hours |
| <code>M</code> | 00 | Matches minutes |
| <code>S</code> | 00 | Matches seconds |
| <code>s</code> | .500 | Matches milliseconds |
| <code>e</code> | Mon, Tues | Matches abbreviated days of the week |
| <code>E</code> | Monday | Matches full name days of the week |
| <code>yyymmdd</code> | 19960101 | Matches fixed-width year, month, and day |

Characters not listed above are normally treated as delimiters between date and time slots. For example a `dt` string of "1996-01-15T00:00:00.0" would have a `format` string like "y-m-dTH:M:S.s". If you need to use a code character as a delimiter you can escape it using backslash. The date "1995y01m" would have the format "y\ym\m".

format ($dt::TimeType, format::AbstractString; locale="english"$) \rightarrow `AbstractString`

Construct a string by using a `TimeType` object and applying the provided `format`. The following character codes can be used to construct the `format` string:

| Code | Examples | Comment |
|----------------|----------|--|
| <code>y</code> | 6 | Numeric year with a fixed width |
| <code>Y</code> | 1996 | Numeric year with a minimum width |
| <code>m</code> | 1, 12 | Numeric month with a minimum width |
| <code>u</code> | Jan | Month name shortened to 3-chars according to the <code>locale</code> |
| <code>U</code> | January | Full month name according to the <code>locale</code> keyword |
| <code>d</code> | 1, 31 | Day of the month with a minimum width |
| <code>H</code> | 0, 23 | Hour (24-hour clock) with a minimum width |
| <code>M</code> | 0, 59 | Minute with a minimum width |
| <code>S</code> | 0, 59 | Second with a minimum width |
| <code>s</code> | 000, 500 | Millisecond with a minimum width of 3 |
| <code>e</code> | Mon, Tue | Abbreviated days of the week |
| <code>E</code> | Monday | Full day of week name |

The number of sequential code characters indicate the width of the code. A format of `yyyy-mm` specifies that the code `y` should have a width of four while `m` a width of two. Codes that yield numeric digits have an associated mode: fixed-width or minimum-width. The fixed-width mode left-pads the value with zeros when it is shorter than the specified width and truncates the value when longer. Minimum-width mode works the same as fixed-width except that it does not truncate values longer than the width.

When creating a `format` you can use any non-code characters as a separator. For example to generate the string “1996-01-15T00:00:00” you could use `format`: “`yyyy-mm-ddTHH:MM:SS`”. Note that if you need to use a code character as a literal you can use the escape character backslash. The string “1996y01m” can be produced with the format “`yyyy\ymm\m`”.

DateFormat (*format::AbstractString, locale::AbstractString="english"*) → `DateFormat`

Construct a date formatting object that can be used for parsing date strings or formatting a date object as a string. For details on the syntax for `format` see [parsing](#) (page 588) and [formatting](#) (page 588).

DateTime (*dt::AbstractString, df::DateFormat*) → `DateTime`

Construct a `DateTime` by parsing the `dt` date string following the pattern given in the `Dates.DateFormat()` (page 589) object. Similar to `DateTime(::AbstractString, ::AbstractString)` but more efficient when repeatedly parsing similarly formatted date strings with a pre-created `DateFormat` object.

Date (*y[, m, d]*) → `Date`

Construct a `Date` type by parts. Arguments must be convertible to `Int64`.

Date (*period::Period...*) → `Date`

Construct a `Date` type by `Period` type parts. Arguments may be in any order. Date parts not provided will default to the value of `Dates.default(period)`.

Date (*f::Function, y[, m, d]; step=Day(1), negate=false, limit=10000*) → `Date`

Create a `Date` through the adjuster API. The starting point will be constructed from the provided `y`, `m`, `d` arguments, and will be adjusted until `f::Function` returns `true`. The step size in adjusting can be provided manually through the `step` keyword. If `negate=true`, then the adjusting will stop when `f::Function` returns `false` instead of `true`. `limit` provides a limit to the max number of iterations the adjustment API will pursue before throwing an error (given that `f::Function` is never satisfied).

Date (*dt::DateTime*) → `Date`

Converts a `DateTime` to a `Date`. The hour, minute, second, and millisecond parts of the `DateTime` are truncated, so only the year, month and day parts are used in construction.

Date (*dt::AbstractString, format::AbstractString; locale="english"*) → `Date`

Construct a `Date` object by parsing a `dt` date string following the pattern given in the `format` string. Follows the same conventions as `DateTime(::AbstractString, ::AbstractString)`.

Date (*dt::AbstractString, df::DateFormat*) → `Date`

Parse a date from a date string `dt` using a `DateFormat` object `df`.

now () → DateTime

Returns a DateTime corresponding to the user's system time including the system timezone locale.

now (::Type{UTC}) → DateTime

Returns a DateTime corresponding to the user's system time as UTC/GMT.

eps (::DateTime) → Millisecond

eps (::Date) → Day

Returns Millisecond(1) for DateTime values and Day(1) for Date values.

54.2.1 Accessor Functions

year (dt::TimeType) → Int64

The year of a Date or DateTime as an Int64.

month (dt::TimeType) → Int64

The month of a Date or DateTime as an Int64.

week (dt::TimeType) → Int64

Return the [ISO week date](#) of a Date or DateTime as an Int64. Note that the first week of a year is the week that contains the first Thursday of the year which can result in dates prior to January 4th being in the last week of the previous year. For example `week(Date(2005, 1, 1))` is the 53rd week of 2004.

day (dt::TimeType) → Int64

The day of month of a Date or DateTime as an Int64.

hour (dt::DateTime) → Int64

The hour of day of a DateTime as an Int64.

minute (dt::DateTime) → Int64

The minute of a DateTime as an Int64.

second (dt::DateTime) → Int64

The second of a DateTime as an Int64.

millisecond (dt::DateTime) → Int64

The millisecond of a DateTime as an Int64.

Year (dt::TimeType) → Year

The year part of a Date or DateTime as a Year.

Month (dt::TimeType) → Month

The month part of a Date or DateTime as a Month.

Week (dt::TimeType) → Week

The week part of a Date or DateTime as a Week. For details see [week\(\)](#) (page 590).

Day (dt::TimeType) → Day

The day part of a Date or DateTime as a Day.

Hour (dt::DateTime) → Hour

The hour part of a DateTime as a Hour.

Minute (dt::DateTime) → Minute

The minute part of a DateTime as a Minute.

Second (dt::DateTime) → Second

The second part of a DateTime as a Second.

Millisecond (dt::DateTime) → Millisecond

The millisecond part of a DateTime as a Millisecond.

yearmonth (*dt::TimeType*) → (Int64, Int64)

Simultaneously return the year and month parts of a `Date` or `DateTime`.

monthday (*dt::TimeType*) → (Int64, Int64)

Simultaneously return the month and day parts of a `Date` or `DateTime`.

yearmonthday (*dt::TimeType*) → (Int64, Int64, Int64)

Simultaneously return the year, month and day parts of a `Date` or `DateTime`.

54.2.2 Query Functions

dayname (*dt::TimeType; locale="english"*) → AbstractString

Return the full day name corresponding to the day of the week of the `Date` or `DateTime` in the given `locale`.

dayabbr (*dt::TimeType; locale="english"*) → AbstractString

Return the abbreviated name corresponding to the day of the week of the `Date` or `DateTime` in the given `locale`.

dayofweek (*dt::TimeType*) → Int64

Returns the day of the week as an Int64 with 1 = Monday, 2 = Tuesday, etc..

dayofmonth (*dt::TimeType*) → Int64

The day of month of a `Date` or `DateTime` as an Int64.

dayofweekofmonth (*dt::TimeType*) → Int

For the day of week of `dt`, returns which number it is in `dt`'s month. So if the day of the week of `dt` is Monday, then 1 = First Monday of the month, 2 = Second Monday of the month, etc. In the range 1:5.

daysofweekinmonth (*dt::TimeType*) → Int

For the day of week of `dt`, returns the total number of that day of the week in `dt`'s month. Returns 4 or 5. Useful in temporal expressions for specifying the last day of a week in a month by including `dayofweekofmonth(dt) == daysofweekinmonth(dt)` in the `adjuster` function.

monthname (*dt::TimeType; locale="english"*) → AbstractString

Return the full name of the month of the `Date` or `DateTime` in the given `locale`.

monthabbr (*dt::TimeType; locale="english"*) → AbstractString

Return the abbreviated month name of the `Date` or `DateTime` in the given `locale`.

daysinmonth (*dt::TimeType*) → Int

Returns the number of days in the month of `dt`. Value will be 28, 29, 30, or 31.

isleapyear (*dt::TimeType*) → Bool

Returns `true` if the year of `dt` is a leap year.

dayofyear (*dt::TimeType*) → Int

Returns the day of the year for `dt` with January 1st being day 1.

daysinyear (*dt::TimeType*) → Int

Returns 366 if the year of `dt` is a leap year, otherwise returns 365.

quarterofyear (*dt::TimeType*) → Int

Returns the quarter that `dt` resides in. Range of value is 1:4.

dayofquarter (*dt::TimeType*) → Int

Returns the day of the current quarter of `dt`. Range of value is 1:92.

54.2.3 Adjuster Functions

trunc (*dt::TimeType*, *::Type{Period}*) → *TimeType*

Truncates the value of *dt* according to the provided *Period* type. E.g. if *dt* is 1996-01-01T12:30:00, then `trunc(dt, Day) == 1996-01-01T00:00:00`.

firstdayofweek (*dt::TimeType*) → *TimeType*

Adjusts *dt* to the Monday of its week.

lastdayofweek (*dt::TimeType*) → *TimeType*

Adjusts *dt* to the Sunday of its week.

firstdayofmonth (*dt::TimeType*) → *TimeType*

Adjusts *dt* to the first day of its month.

lastdayofmonth (*dt::TimeType*) → *TimeType*

Adjusts *dt* to the last day of its month.

firstdayofyear (*dt::TimeType*) → *TimeType*

Adjusts *dt* to the first day of its year.

lastdayofyear (*dt::TimeType*) → *TimeType*

Adjusts *dt* to the last day of its year.

firstdayofquarter (*dt::TimeType*) → *TimeType*

Adjusts *dt* to the first day of its quarter.

lastdayofquarter (*dt::TimeType*) → *TimeType*

Adjusts *dt* to the last day of its quarter.

tonext (*dt::TimeType*, *dow::Int*; *same::Bool=false*) → *TimeType*

Adjusts *dt* to the next day of week corresponding to *dow* with 1 = Monday, 2 = Tuesday, etc. Setting *same=true* allows the current *dt* to be considered as the next *dow*, allowing for no adjustment to occur.

toprev (*dt::TimeType*, *dow::Int*; *same::Bool=false*) → *TimeType*

Adjusts *dt* to the previous day of week corresponding to *dow* with 1 = Monday, 2 = Tuesday, etc. Setting *same=true* allows the current *dt* to be considered as the previous *dow*, allowing for no adjustment to occur.

tofirst (*dt::TimeType*, *dow::Int*; *of=Month*) → *TimeType*

Adjusts *dt* to the first *dow* of its month. Alternatively, *of=Year* will adjust to the first *dow* of the year.

tolast (*dt::TimeType*, *dow::Int*; *of=Month*) → *TimeType*

Adjusts *dt* to the last *dow* of its month. Alternatively, *of=Year* will adjust to the last *dow* of the year.

tonext (*func::Function*, *dt::TimeType*; *step=Day(1)*, *negate=false*, *limit=10000*, *same=false*) → *TimeType*

Adjusts *dt* by iterating at most *limit* iterations by *step* increments until *func* returns *true*. *func* must take a single *TimeType* argument and return a *Bool*. *same* allows *dt* to be considered in satisfying *func*. *negate* will make the adjustment process terminate when *func* returns *false* instead of *true*.

toprev (*func::Function*, *dt::TimeType*; *step=Day(-1)*, *negate=false*, *limit=10000*, *same=false*) → *TimeType*

Adjusts *dt* by iterating at most *limit* iterations by *step* increments until *func* returns *true*. *func* must take a single *TimeType* argument and return a *Bool*. *same* allows *dt* to be considered in satisfying *func*. *negate* will make the adjustment process terminate when *func* returns *false* instead of *true*.

54.2.4 Periods

Year (*v*)

Month (*v*)

Week (*v*)

Day(*v*)**Hour**(*v*)**Minute**(*v*)**Second**(*v*)**Millisecond**(*v*)

Construct a `Period` type with the given *v* value. Input must be losslessly convertible to an `Int64`.

CompoundPeriod(*periods*) → `CompoundPeriod`

Construct a `CompoundPeriod` from a `Vector` of `Period`s. The constructor will automatically simplify the periods into a canonical form according to the following rules:

- All `Period`s of the same type will be added together
- Any `Period` large enough to be partially representable by a coarser `Period` will be broken into multiple `Period`s (eg. `Hour(30)` becomes `Day(1) + Hour(6)`)
- `Period`s with opposite signs will be combined when possible (eg. `Hour(1) - Day(1)` becomes `-Hour(23)`)

Due to the canonicalization, `CompoundPeriod` is also useful for converting time periods into more human-comprehensible forms.

Examples

```
julia> Dates.CompoundPeriod([Dates.Hour(12), Dates.Hour(13)])
1 day, 1 hour

julia> Dates.CompoundPeriod([Dates.Hour(-1), Dates.Minute(1)])
-59 minutes

julia> Dates.CompoundPeriod([Dates.Month(1), Dates.Week(-2)])
1 month, -2 weeks

julia> Dates.CompoundPeriod(Dates.Minute(50000))
4 weeks, 6 days, 17 hours, 20 minutes
```

default(*p::Period*) → `Period`

Returns a sensible “default” value for the input `Period` by returning one (*p*) for `Year`, `Month`, and `Day`, and zero (*p*) for `Hour`, `Minute`, `Second`, and `Millisecond`.

54.2.5 Rounding Functions

`Date` and `DateTime` values can be rounded to a specified resolution (e.g., 1 month or 15 minutes) with `floor`, `ceil`, or `round`.

floor(*dt::TimeType*, *p::Period*) → `TimeType`

Returns the nearest `Date` or `DateTime` less than or equal to *dt* at resolution *p*.

For convenience, *p* may be a type instead of a value: `floor(dt, Dates.Hour)` is a shortcut for `floor(dt, Dates.Hour(1))`.

```
julia> floor(Date(1985, 8, 16), Dates.Month)
1985-08-01

julia> floor(DateTime(2013, 2, 13, 0, 31, 20), Dates.Minute(15))
2013-02-13T00:30:00

julia> floor(DateTime(2016, 8, 6, 12, 0, 0), Dates.Day)
2016-08-06T00:00:00
```

ceil (*dt::TimeType, p::Period*) → *TimeType*

Returns the nearest *Date* or *DateTime* greater than or equal to *dt* at resolution *p*.

For convenience, *p* may be a type instead of a value: `ceil(dt, Dates.Hour)` is a shortcut for `ceil(dt, Dates.Hour(1))`.

```
julia> ceil(Date(1985, 8, 16), Dates.Month)
1985-09-01

julia> ceil(DateTime(2013, 2, 13, 0, 31, 20), Dates.Minute(15))
2013-02-13T00:45:00

julia> ceil(DateTime(2016, 8, 6, 12, 0, 0), Dates.Day)
2016-08-07T00:00:00
```

round (*dt::TimeType, p::Period* [, *r::RoundingMode*]) → *TimeType*

Returns the *Date* or *DateTime* nearest to *dt* at resolution *p*. By default (`RoundNearestTiesUp`), ties (e.g., rounding 9:30 to the nearest hour) will be rounded up.

For convenience, *p* may be a type instead of a value: `round(dt, Dates.Hour)` is a shortcut for `round(dt, Dates.Hour(1))`.

```
julia> round(Date(1985, 8, 16), Dates.Month)
1985-08-01

julia> round(DateTime(2013, 2, 13, 0, 31, 20), Dates.Minute(15))
2013-02-13T00:30:00

julia> round(DateTime(2016, 8, 6, 12, 0, 0), Dates.Day)
2016-08-07T00:00:00
```

Valid rounding modes for `round(::TimeType, ::Period, ::RoundingMode)` are `RoundNearestTiesUp` (default), `RoundDown` (floor), and `RoundUp` (ceil).

The following functions are not exported:

floorceil (*dt::TimeType, p::Period*) → (*TimeType, TimeType*)

Simultaneously return the floor and ceil of a *Date* or *DateTime* at resolution *p*. More efficient than calling both floor and ceil individually.

epochdays2date (*days*) → *Date*

Takes the number of days since the rounding epoch (0000-01-01T00:00:00) and returns the corresponding *Date*.

epochms2datetime (*milliseconds*) → *DateTime*

Takes the number of milliseconds since the rounding epoch (0000-01-01T00:00:00) and returns the corresponding *DateTime*.

date2epochdays (*dt::Date*) → *Int64*

Takes the given *Date* and returns the number of days since the rounding epoch (0000-01-01T00:00:00) as an *Int64*.

datetime2epochms (*dt::DateTime*) → *Int64*

Takes the given *DateTime* and returns the number of milliseconds since the rounding epoch (0000-01-01T00:00:00) as an *Int64*.

54.2.6 Conversion Functions

today () → *Date*

Returns the date portion of `now()`.

unix2datetime (*x*) → DateTime

Takes the number of seconds since unix epoch 1970-01-01T00:00:00 and converts to the corresponding DateTime.

datetime2unix (*dt::DateTime*) → Float64

Takes the given DateTime and returns the number of seconds since the unix epoch 1970-01-01T00:00:00 as a Float64.

julian2datetime (*julian_days*) → DateTime

Takes the number of Julian calendar days since epoch -4713-11-24T12:00:00 and returns the corresponding DateTime.

datetime2julian (*dt::DateTime*) → Float64

Takes the given DateTime and returns the number of Julian calendar days since the julian epoch -4713-11-24T12:00:00 as a Float64.

rata2datetime (*days*) → DateTime

Takes the number of Rata Die days since epoch 0000-12-31T00:00:00 and returns the corresponding DateTime.

datetime2rata (*dt::TimeType*) → Int64

Returns the number of Rata Die days since epoch from the given Date or DateTime.

54.2.7 Constants

Days of the Week:

| Variable | Abbr. | Value (Int) |
|-----------|-------|-------------|
| Monday | Mon | 1 |
| Tuesday | Tue | 2 |
| Wednesday | Wed | 3 |
| Thursday | Thu | 4 |
| Friday | Fri | 5 |
| Saturday | Sat | 6 |
| Sunday | Sun | 7 |

Months of the Year:

| Variable | Abbr. | Value (Int) |
|-----------|-------|-------------|
| January | Jan | 1 |
| February | Feb | 2 |
| March | Mar | 3 |
| April | Apr | 4 |
| May | May | 5 |
| June | Jun | 6 |
| July | Jul | 7 |
| August | Aug | 8 |
| September | Sep | 9 |
| October | Oct | 10 |
| November | Nov | 11 |
| December | Dec | 12 |

Iteration utilities

zip (*iters...*)

For a set of iterable objects, returns an iterable of tuples, where the *i*th tuple contains the *i*th component of each input iterable.

Note that `zip()` (page 597) is its own inverse: `collect(zip(zip(a...)...)) == collect(a)`.

```
julia> a = 1:5
1:5

julia> b = ["e", "d", "b", "c", "a"]
5-element Array{String,1}:
"e"
"d"
"b"
"c"
"a"

julia> c = zip(a,b)
Base.Iterators.Zip2{UnitRange{Int64},Array{String,1}}(1:5,String["e","d","b","c","a"])

julia> length(c)
5

julia> first(c)
(1, "e")
```

enumerate (*iter*)

An iterator that yields (*i*, *x*) where *i* is a counter starting at 1, and *x* is the *i*th value from the given iterator. It's useful when you need not only the values *x* over which you are iterating, but also the number of iterations so far. Note that *i* may not be valid for indexing *iter*; it's also possible that *x* \neq *iter*[*i*], if *iter* has indices that do not start at 1.

```
julia> a = ["a", "b", "c"];

julia> for (index, value) in enumerate(a)
    println("$index $value")
end

1 a
2 b
3 c
```

rest (*iter, state*)

An iterator that yields the same elements as *iter*, but starting at the given *state*.

countfrom (*start=1, step=1*)

An iterator that counts forever, starting at *start* and incrementing by *step*.

take (*iter, n*)

An iterator that generates at most the first *n* elements of *iter*.

```
julia> a = 1:2:11
1:2:11

julia> collect(a)
6-element Array{Int64,1}:
 1
 3
 5
 7
 9
11

julia> collect(Iterators.take(a, 3))
3-element Array{Int64,1}:
 1
 3
 5
```

drop (*iter, n*)

An iterator that generates all but the first *n* elements of *iter*.

```
julia> a = 1:2:11
1:2:11

julia> collect(a)
6-element Array{Int64,1}:
 1
 3
 5
 7
 9
11

julia> collect(Iterators.drop(a, 4))
2-element Array{Int64,1}:
 9
11
```

cycle (*iter*)

An iterator that cycles through *iter* forever.

repeated (*x* [, *n::Int*])

An iterator that generates the value *x* forever. If *n* is specified, generates *x* that many times (equivalent to `take(repeated(x), n)`).

```
julia> a = Iterators.repeated([1 2], 4);

julia> collect(a)
4-element Array{Array{Int64,2},1}:
 [1 2]
 [1 2]
 [1 2]
 [1 2]
```

product (*iters...*)

Returns an iterator over the product of several iterators. Each generated element is a tuple whose *i*th element comes from the *i*th argument iterator. The first iterator changes the fastest. Example:

```
julia> collect(Iterators.product(1:2, 3:5))
2×3 Array{Tuple{Int64, Int64}, 2}:
 (1, 3) (1, 4) (1, 5)
 (2, 3) (2, 4) (2, 5)
```

flatten (*iter*)

Given an iterator that yields iterators, return an iterator that yields the elements of those iterators. Put differently, the elements of the argument iterator are concatenated. Example:

```
julia> collect(Iterators.flatten((1:2, 8:9)))
4-element Array{Int64, 1}:
 1
 2
 8
 9
```

partition (*collection, n*)

Iterate over a collection *n* elements at a time.

```
julia> collect(Iterators.partition([1, 2, 3, 4, 5], 2))
3-element Array{Array{Int64, 1}, 1}:
 [1, 2]
 [3, 4]
 [5]
```

Unit Testing

56.1 Testing Base Julia

Julia is under rapid development and has an extensive test suite to verify functionality across multiple platforms. If you build Julia from source, you can run this test suite with `make test`. In a binary install, you can run the test suite using `Base.runtests()`.

runtests (`[tests=["all"]`, `numcores=ceil(Int, Sys.CPU_CORES / 2)`])

Run the Julia unit tests listed in `tests`, which can be either a string or an array of strings, using `numcores` processors. (not exported)

56.2 Basic Unit Tests

The `Base.Test` module provides simple *unit testing* functionality. Unit testing is a way to see if your code is correct by checking that the results are what you expect. It can be helpful to ensure your code still works after you make changes, and can be used when developing as a way of specifying the behaviors your code should have when complete.

Simple unit testing can be performed with the `@test()` and `@test_throws()` macros:

@test ex

Tests that the expression `ex` evaluates to `true`. Returns a `Pass Result` if it does, a `Fail Result` if it is `false`, and an `Error Result` if it could not be evaluated.

@test_throws extype ex

Tests that the expression `ex` throws an exception of type `exctype`.

For example, suppose we want to check our new function `foo(x)` works as expected:

```
julia> using Base.Test

julia> foo(x) = length(x)^2
foo (generic function with 1 method)
```

If the condition is true, a `Pass` is returned:

```
julia> @test foo("bar") == 9
Test Passed
  Expression: foo("bar") == 9
  Evaluated: 9 == 9

julia> @test foo("fizz") >= 10
```

```
Test Passed
Expression: foo("fizz") >= 10
Evaluated: 16 >= 10
```

If the condition is false, then a `Fail` is returned and an exception is thrown:

```
julia> @test foo("f") == 20
Test Failed
Expression: foo("f") == 20
Evaluated: 1 == 20
ERROR: There was an error during testing
in record at test.jl:268
in do_test at test.jl:191
```

If the condition could not be evaluated because an exception was thrown, which occurs in this case because `length()` is not defined for symbols, an `Error` object is returned and an exception is thrown:

```
julia> @test foo(:cat) == 1
Error During Test
Test threw an exception of type MethodError
Expression: foo(:cat) == 1
MethodError: `length` has no method matching length(::Symbol)
in foo at none:1
in anonymous at test.jl:159
in do_test at test.jl:180
ERROR: There was an error during testing
in record at test.jl:268
in do_test at test.jl:191
```

If we expect that evaluating an expression *should* throw an exception, then we can use `@test_throws()` to check that this occurs:

```
julia> @test_throws MethodError foo(:cat)
Test Passed
Expression: foo(:cat)
Evaluated: MethodError
```

56.3 Working with Test Sets

Typically a large number of tests are used to make sure functions work correctly over a range of inputs. In the event a test fails, the default behavior is to throw an exception immediately. However, it is normally preferable to run the rest of the tests first to get a better picture of how many errors there are in the code being tested.

The `@testset()` macro can be used to group tests into *sets*. All the tests in a test set will be run, and at the end of the test set a summary will be printed. If any of the tests failed, or could not be evaluated due to an error, the test set will then throw a `TestSetException`.

```
@testset [CustomTestSet] [option=val ...] ["description"] begin ... end
@testset [CustomTestSet] [option=val ...] ["description $v"] for v in (...) ... end
@testset [CustomTestSet] [option=val ...] ["description $v, $w"] for v in (...), w in (...)
```

Starts a new test set, or multiple test sets if a `for` loop is provided.

If no custom testset type is given it defaults to creating a `DefaultTestSet`. `DefaultTestSet` records all the results and, and if there are any `Fails` or `Errors`, throws an exception at the end of the top-level (non-nested) test set, along with a summary of the test results.

Any custom testset type (subtype of `AbstractTestSet`) can be given and it will also be used for any nested `@testset` invocations. The given options are only applied to the test set where they are given. The default test set type does not take any options.

The description string accepts interpolation from the loop indices. If no description is provided, one is constructed based on the variables.

By default the `@testset` macro will return the testset object itself, though this behavior can be customized in other testset types. If a `for` loop is used then the macro collects and returns a list of the return values of the `finish` method, which by default will return a list of the testset objects used in each iteration.

We can put our tests for the `foo(x)` function in a test set:

```
julia> @testset "Foo Tests" begin
    @test foo("a") == 1
    @test foo("ab") == 4
    @test foo("abc") == 9
end
Test Summary: | Pass  Total
Foo Tests    |     3     3
```

Test sets can also be nested:

```
julia> @testset "Foo Tests" begin
    @testset "Animals" begin
        @test foo("cat") == 9
        @test foo("dog") == foo("cat")
    end
    @testset "Arrays $i" for i in 1:3
        @test foo(zeros(i)) == i^2
        @test foo(ones(i)) == i^2
    end
end
Test Summary: | Pass  Total
Foo Tests    |     8     8
```

In the event that a nested test set has no failures, as happened here, it will be hidden in the summary. If we do have a test failure, only the details for the failed test sets will be shown:

```
julia> @testset "Foo Tests" begin
    @testset "Animals" begin
        @testset "Felines" begin
            @test foo("cat") == 9
        end
        @testset "Canines" begin
            @test foo("dog") == 9
        end
    end
    @testset "Arrays" begin
        @test foo(zeros(2)) == 4
        @test foo(ones(4)) == 15
    end
end

Arrays: Test Failed
  Expression: foo(ones(4)) == 15
  Evaluated: 16 == 15
in record at test.jl:297
in do_test at test.jl:191
Test Summary: | Pass  Fail  Total
```

```

Foo Tests      |      3      1      4
Animals       |      2      2
Arrays        |      1      1      2
ERROR: Some tests did not pass: 3 passed, 1 failed, 0 errored, 0 broken.
in finish at test.jl:362

```

56.4 Other Test Macros

As calculations on floating-point values can be imprecise, you can perform approximate equality checks using either `@test a ≈ b` (where `≈`, typed via tab completion of `\approx`, is the `isapprox()` function) or use `isapprox()` directly.

An alternative is the `@test_approx_eq` macro (which differs from `isapprox` in that it treats NaN values as equal and has a smaller default tolerance) or `@test_approx_eq_eps` (which takes an extra argument indicating the relative tolerance):

```

julia> @test 1 ≈ 0.999999999
julia> @test 1 ≈ 0.999999
ERROR: test failed: 1 isapprox 0.999999
in expression: 1 ≈ 0.999999
in error at error.jl:21
in default_handler at test.jl:30
in do_test at test.jl:53

julia> @test_approx_eq 1. 0.999999999
ERROR: assertion failed: |1.0 - 0.999999999| < 2.220446049250313e-12
1.0 = 1.0
0.999999999 = 0.999999999
in test_approx_eq at test.jl:75
in test_approx_eq at test.jl:80

julia> @test_approx_eq 1. 0.999999999999999
julia> @test_approx_eq_eps 1. 0.999 1e-2

julia> @test_approx_eq_eps 1. 0.999 1e-3
ERROR: assertion failed: |1.0 - 0.999| <= 0.001
1.0 = 1.0
0.999 = 0.999
difference = 0.00100000000000000009 > 0.001
in error at error.jl:22
in test_approx_eq at test.jl:68

```

Note that these macros will fail immediately, and are not compatible with `@testset()`, so using `@test isapprox` is encouraged when writing new tests.

`@test_approx_eq(a, b)`

Test two floating point numbers `a` and `b` for equality taking into account small numerical errors.

`@test_approx_eq_eps(a, b, tol)`

Test two floating point numbers `a` and `b` for equality taking into account a margin of tolerance given by `tol`.

`@inferred f(x)`

Tests that the call expression `f(x)` returns a value of the same type inferred by the compiler. It is useful to check for type stability.

`f(x)` can be any call expression. Returns the result of `f(x)` if the types match, and an `Error Result` if it finds different types.

```
julia> using Base.Test

julia> f(a,b,c) = b > 1 ? 1 : 1.0
f (generic function with 1 method)

julia> typeof(f(1,2,3))
Int64

julia> @code_warntype f(1,2,3)
...
Body:
  begin
    unless (Base.slt_int)(1,b::Int64)::Bool goto 3
    return 1
  3:
    return 1.0
  end::UNION{Float64, Int64}

julia> @inferred f(1,2,3)
ERROR: return type Int64 does not match inferred return type Union{Float64, Int64}
in error(::String) at ./error.jl:21
...

julia> @inferred max(1,2)
2
```

56.5 Broken Tests

If a test fails consistently it can be changed to use the `@test_broken()` macro. This will denote the test as `Broken` if the test continues to fail and alerts the user via an `Error` if the test succeeds.

`@test_broken ex`

Indicates a test that should pass but currently consistently fails. Tests that the expression `ex` evaluates to `false` or causes an exception. Returns a `Broken Result` if it does, or an `Error Result` if the expression evaluates to `true`.

`@test_skip()` is also available to skip a test without evaluation, but counting the skipped test in the test set reporting. The test will not run but gives a `Broken Result`.

`@test_skip ex`

Marks a test that should not be executed but should be included in test summary reporting as `Broken`. This can be useful for tests that intermittently fail, or tests of not-yet-implemented functionality.

56.6 Creating Custom `AbstractTestSet` Types

Packages can create their own `AbstractTestSet` subtypes by implementing the `record` and `finish` methods. The subtype should have a one-argument constructor taking a description string, with any options passed in as keyword arguments.

`record(ts::AbstractTestSet, res::Result)`

Record a result to a testset. This function is called by the `@testset` infrastructure each time a contained

`@test` macro completes, and is given the test result (which could be an `Error`). This will also be called with an `Error` if an exception is thrown inside the test block but outside of a `@test` context.

finish (*ts::AbstractTestSet*)

Do any final processing necessary for the given testset. This is called by the `@testset` infrastructure after a test block executes. One common use for this function is to record the testset to the parent's results list, using `get_testset`.

`Base.Test` takes responsibility for maintaining a stack of nested testsets as they are executed, but any result accumulation is the responsibility of the `AbstractTestSet` subtype. You can access this stack with the `get_testset` and `get_testset_depth` methods. Note that these functions are not exported.

get_testset ()

Retrieve the active test set from the task's local storage. If no test set is active, use the fallback default test set.

get_testset_depth ()

Returns the number of active test sets, not including the default test set

`Base.Test` also makes sure that nested `@testset` invocations use the same `AbstractTestSet` subtype as their parent unless it is set explicitly. It does not propagate any properties of the testset. Option inheritance behavior can be implemented by packages using the stack infrastructure that `Base.Test` provides.

Defining a basic `AbstractTestSet` subtype might look like:

```
import Base.Test: record, finish
using Base.Test: AbstractTestSet, Result, Pass, Fail, Error
using Base.Test: get_testset_depth, get_testset
immutable CustomTestSet <: Base.Test.AbstractTestSet
    description::AbstractString
    foo::Int
    results::Vector
    # constructor takes a description string and options keyword arguments
    CustomTestSet(desc; foo=1) = new(desc, foo, [])
end

record(ts::CustomTestSet, child::AbstractTestSet) = push!(ts.results, child)
record(ts::CustomTestSet, res::Result) = push!(ts.results, res)
function finish(ts::CustomTestSet)
    # just record if we're not the top-level parent
    if get_testset_depth() > 0
        record(get_testset(), ts)
    end
    ts
end
```

And using that testset looks like:

```
@testset CustomTestSet foo=4 "custom testset inner 2" begin
    # this testset should inherit the type, but not the argument.
    @testset "custom testset inner" begin
        @test true
    end
end
```

C Interface

ccall *((symbol, library) or function_pointer, ReturnType, (ArgumentType1, ...), ArgumentValue1, ...)*

Call function in C-exported shared library, specified by *(function name, library)* tuple, where each component is a string or symbol.

Note that the argument type tuple must be a literal tuple, and not a tuple-valued variable or expression. Alternatively, `ccall` may also be used to call a function pointer, such as one returned by `dlsym`.

Each `ArgumentValue` to the `ccall` will be converted to the corresponding `ArgumentType`, by automatic insertion of calls to `unsafe_convert(ArgumentType, cconvert(ArgumentType, ArgumentValue))`. (See also the documentation for each of these functions for further details.) In most cases, this simply results in a call to `convert(ArgumentType, ArgumentValue)`.

cglobal *((symbol, library)[, type=Void])*

Obtain a pointer to a global variable in a C-exported shared library, specified exactly as in `ccall()` (page 607). Returns a `Ptr{Type}`, defaulting to `Ptr{Void}` if no `Type` argument is supplied. The values can be read or written by `unsafe_load()` (page 608) or `unsafe_store!()` (page 608), respectively.

cfunction *(function::Function, ReturnType::Type, (ArgumentTypes...))*

Generate C-callable function pointer from Julia function. Type annotation of the return value in the callback function is a must for situations where Julia cannot infer the return type automatically.

For example:

```
function foo()
    # body

    retval::Float64
end

bar = cfunction(foo, Float64, ())
```

unsafe_convert *(T, x)*

Convert `x` to a value of type `T`

In cases where `convert()` (page 364) would need to take a Julia object and turn it into a `Ptr`, this function should be used to define and perform that conversion.

Be careful to ensure that a Julia reference to `x` exists as long as the result of this function will be used. Accordingly, the argument `x` to this function should never be an expression, only a variable name or field reference. For example, `x=a.b.c` is acceptable, but `x=[a,b,c]` is not.

The `unsafe` prefix on this function indicates that using the result of this function after the `x` argument to this function is no longer accessible to the program may cause undefined behavior, including program corruption or segfaults, at any later time.

cconvert (*T*, *x*)

Convert *x* to a value of type *T*, typically by calling `convert` (*T*, *x*)

In cases where *x* cannot be safely converted to *T*, unlike `convert` () (page 364), `cconvert` may return an object of a type different from *T*, which however is suitable for `unsafe_convert` () (page 607) to handle.

Neither `convert` nor `cconvert` should take a Julia object and turn it into a `Ptr`.

unsafe_load (*p::Ptr{T}*, *i::Integer=1*)

Load a value of type *T* from the address of the *i*th element (1-indexed) starting at *p*. This is equivalent to the C expression `p[i-1]`.

The `unsafe` prefix on this function indicates that no validation is performed on the pointer *p* to ensure that it is valid. Incorrect usage may segfault your program or return garbage answers, in the same manner as C.

unsafe_store! (*p::Ptr{T}*, *x*, *i::Integer=1*)

Store a value of type *T* to the address of the *i*th element (1-indexed) starting at *p*. This is equivalent to the C expression `p[i-1] = x`.

The `unsafe` prefix on this function indicates that no validation is performed on the pointer *p* to ensure that it is valid. Incorrect usage may corrupt or segfault your program, in the same manner as C.

unsafe_copy! (*dest::Ptr{T}*, *src::Ptr{T}*, *N*)

Copy *N* elements from a source pointer to a destination, with no checking. The size of an element is determined by the type of the pointers.

The `unsafe` prefix on this function indicates that no validation is performed on the pointers *dest* and *src* to ensure that they are valid. Incorrect usage may corrupt or segfault your program, in the same manner as C.

unsafe_copy! (*dest::Array*, *do*, *src::Array*, *so*, *N*)

Copy *N* elements from a source array to a destination, starting at offset *so* in the source and *do* in the destination (1-indexed).

The `unsafe` prefix on this function indicates that no validation is performed to ensure that *N* is in bounds on either array. Incorrect usage may corrupt or segfault your program, in the same manner as C.

copy! (*dest*, *src*)

Copy all elements from collection *src* to array *dest*. Returns *dest*.

copy! (*dest*, *do*, *src*, *so*, *N*)

Copy *N* elements from collection *src* starting at offset *so*, to array *dest* starting at offset *do*. Returns *dest*.

pointer (*array*[, *index*])

Get the native address of an array or string element. Be careful to ensure that a Julia reference to *a* exists as long as this pointer will be used. This function is “unsafe” like `unsafe_convert`.

Calling `Ref(array[, index])` is generally preferable to this function.

unsafe_wrap (*Array*, *pointer::Ptr{T}*, *dims*, *own=false*)

Wrap a Julia `Array` object around the data at the address given by *pointer*, without making a copy. The pointer element type *T* determines the array element type. *dims* is either an integer (for a 1d array) or a tuple of the array dimensions. *own* optionally specifies whether Julia should take ownership of the memory, calling `free` on the pointer when the array is no longer referenced.

This function is labelled “unsafe” because it will crash if *pointer* is not a valid memory address to data of the requested length.

pointer_from_objref (*x*)

Get the memory address of a Julia object as a `Ptr`. The existence of the resulting `Ptr` will not protect the object from garbage collection, so you must ensure that the object remains referenced for the whole time that the `Ptr` will be used.

unsafe_pointer_to_objref (*p::Ptr*)

Convert a `Ptr` to an object reference. Assumes the pointer refers to a valid heap-allocated Julia object. If this is not the case, undefined behavior results, hence this function is considered “unsafe” and should be used with care.

disable_sigint (*f::Function*)

Disable Ctrl-C handler during execution of a function on the current task, for calling external code that may call julia code that is not interrupt safe. Intended to be called using `do` block syntax as follows:

```
disable_sigint() do
    # interrupt-unsafe code
    ...
end
```

This is not needed on worker threads (`Threads.threadid() != 1`) since the `InterruptException` will only be delivered to the master thread. External functions that do not call julia code or julia runtime automatically disable sigint during their execution.

reenable_sigint (*f::Function*)

Re-enable Ctrl-C handler during execution of a function. Temporarily reverses the effect of `disable_sigint`.

systemerror (*sysfunc, iftrue*)

Raises a `SystemError` for `errno` with the descriptive string `sysfunc` if `iftrue` is `true`

Ptr{T}

A memory address referring to data of type `T`. However, there is no guarantee that the memory is actually valid, or that it actually represents data of the specified type.

Ref{T}

An object that safely references data of type `T`. This type is guaranteed to point to valid, Julia-allocated memory of the correct type. The underlying data is protected from freeing by the garbage collector as long as the `Ref` itself is referenced.

When passed as a `ccall` argument (either as a `Ptr` or `Ref` type), a `Ref` object will be converted to a native pointer to the data it references.

There is no invalid (NULL) `Ref`.

Cchar

Equivalent to the native `char` c-type.

Cuchar

Equivalent to the native unsigned `char` c-type (`UInt8`).

Cshort

Equivalent to the native signed `short` c-type (`Int16`).

Cushort

Equivalent to the native unsigned `short` c-type (`UInt16`).

Cint

Equivalent to the native signed `int` c-type (`Int32`).

Cuint

Equivalent to the native unsigned `int` c-type (`UInt32`).

Clong

Equivalent to the native signed `long` c-type.

Culong

Equivalent to the native unsigned `long` c-type.

Clonglong

Equivalent to the native `signed long long` c-type (Int64).

Culonglong

Equivalent to the native `unsigned long long` c-type (UInt64).

Cintmax_t

Equivalent to the native `intmax_t` c-type (Int64).

Cuintmax_t

Equivalent to the native `uintmax_t` c-type (UInt64).

Csize_t

Equivalent to the native `size_t` c-type (UInt).

Cssize_t

Equivalent to the native `ssize_t` c-type.

Cptrdiff_t

Equivalent to the native `ptrdiff_t` c-type (Int).

Cwchar_t

Equivalent to the native `wchar_t` c-type (Int32).

Cfloat

Equivalent to the native `float` c-type (Float32).

Cdouble

Equivalent to the native `double` c-type (Float64).

LLVM Interface

llvmcall (*IR::String*, *ReturnType*, (*ArgumentType1*, ...), *ArgumentValue1*, ...)

llvmcall ((*declarations::String*, *IR::String*), *ReturnType*, (*ArgumentType1*, ...), *ArgumentValue1*, ...)

Call LLVM IR string in the first argument. Similar to an LLVM function `define` block, arguments are available as consecutive unnamed SSA variables (`%0`, `%1`, etc.).

The optional `declarations` string contains external functions declarations that are necessary for `llvm` to compile the IR string. Multiple declarations can be passed in by separating them with line breaks.

Note that the argument type tuple must be a literal tuple, and not a tuple-valued variable or expression.

Each `ArgumentValue` to `llvmcall` will be converted to the corresponding `ArgumentType`, by automatic insertion of calls to `unsafe_convert(ArgumentType, cconvert(ArgumentType, ArgumentValue))`. (see also the documentation for each of these functions for further details). In most cases, this simply results in a call to `convert(ArgumentType, ArgumentValue)`.

See `test/llvmcall.jl` for usage examples.

C Standard Library

malloc (*size::Integer*) → Ptr{Void}

Call `malloc` from the C standard library.

calloc (*num::Integer, size::Integer*) → Ptr{Void}

Call `calloc` from the C standard library.

realloc (*addr::Ptr, size::Integer*) → Ptr{Void}

Call `realloc` from the C standard library.

See warning in the documentation for `free` regarding only using this on memory originally obtained from `malloc`.

free (*addr::Ptr*)

Call `free` from the C standard library. Only use this on memory obtained from `malloc`, not on pointers retrieved from other C libraries. `Ptr` objects obtained from C libraries should be freed by the `free` functions defined in that library, to avoid assertion failures if multiple `libc` libraries exist on the system.

errno ([*code*])

Get the value of the C library's `errno`. If an argument is specified, it is used to set the value of `errno`.

The value of `errno` is only valid immediately after a `ccall` to a C library routine that sets it. Specifically, you cannot call `errno` at the next prompt in a REPL, because lots of code is executed between prompts.

strerror (*n=errno()*)

Convert a system call error code to a descriptive string

GetLastError ()

Call the Win32 `GetLastError` function [only available on Windows].

FormatMessage (*n=GetLastError()*)

Convert a Win32 system call error code to a descriptive string [only available on Windows].

time (*t::TmStruct*)

Converts a `TmStruct` struct to a number of seconds since the epoch.

strftime ([*format*], *time*)

Convert `time`, given as a number of seconds since the epoch or a `TmStruct`, to a formatted string using the given format. Supported formats are the same as those in the standard C library.

strptime ([*format*], *timestr*)

Parse a formatted time string into a `TmStruct` giving the seconds, minute, hour, date, etc. Supported formats are the same as those in the standard C library. On some platforms, timezones will not be parsed correctly. If the result of this function will be passed to `time` to convert it to seconds since the epoch, the `isdst` field should be filled in manually. Setting it to `-1` will tell the C library to use the current system settings to determine the timezone.

TmStruct (*[seconds]*)

Convert a number of seconds since the epoch to broken-down format, with fields *sec*, *min*, *hour*, *mday*, *month*, *year*, *wday*, *yday*, and *isdst*.

flush_cstdio()

Flushes the C *stdout* and *stderr* streams (which may have been written to by external C code).

Dynamic Linker

The names in `Base.Libdl` are not exported and need to be called e.g. as `Libdl.dlopen()`.

dlopen (*libfile::AbstractString* [, *flags::Integer*])

Load a shared library, returning an opaque handle.

The extension given by the constant `dlext` (`.so`, `.dll`, or `.dylib`) can be omitted from the `libfile` string, as it is automatically appended if needed. If `libfile` is not an absolute path name, then the paths in the array `DL_LOAD_PATH` are searched for `libfile`, followed by the system load path.

The optional `flags` argument is a bitwise-or of zero or more of `RTLD_LOCAL`, `RTLD_GLOBAL`, `RTLD_LAZY`, `RTLD_NOW`, `RTLD_NODELETE`, `RTLD_NOLOAD`, `RTLD_DEEPBIND`, and `RTLD_FIRST`. These are converted to the corresponding flags of the POSIX (and/or GNU libc and/or MacOS) `dlopen` command, if possible, or are ignored if the specified functionality is not available on the current platform. The default flags are platform specific. On MacOS the default `dlopen` flags are `RTLD_LAZY|RTLD_DEEPBIND|RTLD_GLOBAL` while on other platforms the defaults are `RTLD_LAZY|RTLD_DEEPBIND|RTLD_LOCAL`. An important usage of these flags is to specify non default behavior for when the dynamic library loader binds library references to exported symbols and if the bound references are put into process local or global scope. For instance `RTLD_LAZY|RTLD_DEEPBIND|RTLD_GLOBAL` allows the library's symbols to be available for usage in other shared libraries, addressing situations where there are dependencies between shared libraries.

dlopen_e (*libfile::AbstractString* [, *flags::Integer*])

Similar to `dlopen()` (page 615), except returns a `NULL` pointer instead of raising errors.

RTLD_DEEPBIND

RTLD_FIRST

RTLD_GLOBAL

RTLD_LAZY

RTLD_LOCAL

RTLD_NODELETE

RTLD_NOLOAD

RTLD_NOW

Enum constant for `dlopen()` (page 615). See your platform man page for details, if applicable.

dlsym (*handle*, *sym*)

Look up a symbol from a shared library handle, return callable function pointer on success.

dlsym_e (*handle*, *sym*)

Look up a symbol from a shared library handle, silently return `NULL` pointer on lookup failure.

dlclose (*handle*)

Close shared library referenced by handle.

dlext

File extension for dynamic libraries (e.g. `dll`, `dylib`, `so`) on the current platform.

find_library (*names*, *locations*)

Searches for the first library in *names* in the paths in the *locations* list, `DL_LOAD_PATH`, or system library paths (in that order) which can successfully be `dlopen`'d. On success, the return value will be one of the names (potentially prefixed by one of the paths in *locations*). This string can be assigned to a `global const` and used as the library name in future `ccall`'s. On failure, it returns the empty string.

DL_LOAD_PATH

When calling `dlopen()` (page 615), the paths in this list will be searched first, in order, before searching the system locations for a valid library handle.

Profiling

@profile()

@profile <expression> runs your expression while taking periodic backtraces. These are appended to an internal buffer of backtraces.

The methods in `Base.Profile` are not exported and need to be called e.g. as `Profile.print()`.

clear()

Clear any existing backtraces from the internal buffer.

print ([*io::IO = STDOUT*], [*data::Vector*]; *kwargs...*)

Prints profiling results to *io* (by default, `STDOUT`). If you do not supply a data vector, the internal buffer of accumulated backtraces will be used.

The keyword arguments can be any combination of:

- `format` can be `:tree` (default) or `:flat`.
- If `C` is `true`, backtraces from C and Fortran code are shown (normally they are excluded).
- If `combine` is `true` (default), instruction pointers are merged that correspond to the same line of code.
- `maxdepth` can be used to limit the depth of printing in `:tree` format, while `sortedby` can be used to control the order in `:flat` format: `:filefuncline` (default) sorts by the source line, whereas `:count` sorts in order of number of collected samples.
- `noisefloor` only shows frames that exceed the heuristic noise floor of the sample (only applies to format `:tree`). A suggested value to try for this is 2.0 (the default is 0). This parameter hides samples for which $n \leq \text{noisefloor} * \sqrt{N}$, where *n* is the number of samples on this line, and *N* is the number of samples for the callee.
- `mincount` can also be used to limit the printout to only those lines with at least `mincount` occurrences.

print ([*io::IO = STDOUT*], *data::Vector*, *lidict::LineInfoDict*; *kwargs...*)

Prints profiling results to *io*. This variant is used to examine results exported by a previous call to `retrieve()` (page 618). Supply the vector *data* of backtraces and a dictionary *lidict* of line information.

See `Profile.print([io], data)` for an explanation of the valid keyword arguments.

init (; *n::Integer*, *delay::Float64*)

Configure the `delay` between backtraces (measured in seconds), and the number *n* of instruction pointers that may be stored. Each instruction pointer corresponds to a single line of code; backtraces generally consist of a long list of instruction pointers. Default settings can be obtained by calling this function with no arguments, and each can be set independently using keywords or in the order `(n, delay)`.

fetch() → *data*

Returns a reference to the internal buffer of backtraces. Note that subsequent operations, like `clear()`

(page 617), can affect `data` unless you first make a copy. Note that the values in `data` have meaning only on this machine in the current session, because it depends on the exact memory addresses used in JIT-compiling. This function is primarily for internal use; `retrieve()` (page 618) may be a better choice for most users.

retrieve() → `data`, `lidict`

“Exports” profiling results in a portable format, returning the set of all backtraces (`data`) and a dictionary that maps the (session-specific) instruction pointers in `data` to `LineInfo` values that store the file name, function name, and line number. This function allows you to save profiling results for future analysis.

callers(`funcname`[, `data`, `lidict`][, `filename`=<`filename`>][, `linerange`=<`start:stop`>]) → `Vector{Tuple{count, lineinfo}}`

Given a previous profiling run, determine who called a particular function. Supplying the `filename` (and optionally, range of line numbers over which the function is defined) allows you to disambiguate an overloaded method. The returned value is a vector containing a count of the number of calls and line information about the caller. One can optionally supply backtrace `data` obtained from `retrieve()` (page 618); otherwise, the current internal profile buffer is used.

clear_malloc_data()

Clears any stored memory allocation data when running `julia` with `--track-allocation`. Execute the command(s) you want to test (to force JIT-compilation), then call `clear_malloc_data()` (page 618). Then execute your command(s) again, quit Julia, and examine the resulting `*.mem` files.

SIMD Support

Type `VecElement{T}` is intended for building libraries of SIMD operations. Practical use of it requires using `llvmcall`. The type is defined as:

```
immutable VecElement{T}
    value::T
end
```

It has a special compilation rule: a homogeneous tuple of `VecElement{T}` maps to an LLVM vector type when `T` is a bitstype and the tuple length is in the set `{2-6,8-10,16}`.

At `-O3`, the compiler *might* automatically vectorize operations on such tuples. For example, the following program, when compiled with `julia -O3` generates two SIMD addition instructions (`addps`) on x86 systems:

```
typealias m128 NTuple{4,VecElement{Float32}}

function add(a::m128, b::m128)
    (VecElement(a[1].value+b[1].value),
     VecElement(a[2].value+b[2].value),
     VecElement(a[3].value+b[3].value),
     VecElement(a[4].value+b[4].value))
end

triple(c::m128) = add(add(c,c),c)

code_native(triple, (m128,))
```

However, since the automatic vectorization cannot be relied upon, future use will mostly be via libraries that use `llvmcall`.

- [Clarke61] Arthur C. Clarke, *Profiles of the Future* (1961): Clarke's Third Law.
- [Bischof1987] C Bischof and C Van Loan, "The WY representation for products of Householder matrices", *SIAM J Sci Stat Comput* 8 (1987), s2-s13. doi:10.1137/0908009
- [Schreiber1989] R Schreiber and C Van Loan, "A storage-efficient WY representation for products of Householder transformations", *SIAM J Sci Stat Comput* 10 (1989), 53-57. doi:10.1137/0910005
- [Bunch1977] J R Bunch and L Kaufman, Some stable methods for calculating inertia and solving symmetric linear systems, *Mathematics of Computation* 31:137 (1977), 163-179. url.
- [issue8859] Issue 8859, "Fix least squares", <https://github.com/JuliaLang/julia/pull/8859>
- [B96] Åke Björck, "Numerical Methods for Least Squares Problems", SIAM Press, Philadelphia, 1996, "Other Titles in Applied Mathematics", Vol. 51. doi:10.1137/1.9781611971484
- [S84] 7. (a) Stewart, "Rank Degeneracy", *SIAM Journal on Scientific and Statistical Computing*, 5(2), 1984, 403-413. doi:10.1137/0905030
- [KY88] Konstantinos Konstantinides and Kung Yao, "Statistical analysis of effective singular values in matrix rank determination", *IEEE Transactions on Acoustics, Speech and Signal Processing*, 36(5), 1988, 757-763. doi:10.1109/29.1585
- [H05] Nicholas J. Higham, "The squaring and scaling method for the matrix exponential revisited", *SIAM Journal on Matrix Analysis and Applications*, 26(4), 2005, 1179-1193. doi:10.1137/090768539
- [AH12] Awad H. Al-Mohy and Nicholas J. Higham, "Improved inverse scaling and squaring algorithms for the matrix logarithm", *SIAM Journal on Scientific Computing*, 34(4), 2012, C153-C169. doi:10.1137/110852553
- [AHR13] Awad H. Al-Mohy, Nicholas J. Higham and Samuel D. Relton, "Computing the Fréchet derivative of the matrix logarithm and estimating the condition number", *SIAM Journal on Scientific Computing*, 35(4), 2013, C394-C410. doi:10.1137/120885991
- [BH83] Åke Björck and Sven Hammarling, "A Schur method for the square root of a matrix", *Linear Algebra and its Applications*, 52-53, 1983, 127-140. doi:10.1016/0024-3795(83)80010-X

Symbols

*() (in module Base), 405, 443, 501

+() (in module Base), 405

-() (in module Base), 405

.

=() (in module Base), 410

.*() (in module Base), 406

.+() (in module Base), 405

.-() (in module Base), 405

./() (in module Base), 406

==() (in module Base), 410

.≥() (in module Base), 410

.≤() (in module Base), 410

.≠() (in module Base), 410

.^() (in module Base), 406

.\() (in module Base), 406

.>() (in module Base), 410

.>=() (in module Base), 410

.<() (in module Base), 410

.<=() (in module Base), 410

/() (in module Base), 405

//() (in module Base), 408

:() (in module Base), 409

==() (in module Base), 409

===() (in module Base), 362

%() (in module Base), 407

&() (in module Base), 411

γ (in module Base), 435

φ (in module Base), 435

π (in module Base), 435

() (in module Base), 411

∩() (in module Base), 397

·() (in module Base), 501

∪() (in module Base), 397

≡() (in module Base), 362

≥() (in module Base), 410

∈() (in module Base), 383

≤() (in module Base), 410

∃() (in module Base), 383

≠() (in module Base), 410

≠() (in module Base), 410

≠() (in module Base), 383

∉() (in module Base), 383

⊈() (in module Base), 392

⊆() (in module Base), 392, 397

⊂() (in module Base), 392

÷() (in module Base), 407

×() (in module Base), 502

__precompile__() (in module Base), 361

^() (in module Base), 405, 443

~() (in module Base), 411

\() (in module Base), 405, 501

l() (in module Base), 411

l>() (in module Base), 368

>() (in module Base), 410

>=() (in module Base), 410

>>() (in module Base), 408

>>>() (in module Base), 409

<() (in module Base), 410

<:() (in module Base), 365

<=() (in module Base), 410

<<() (in module Base), 408

A

A_ldiv_B

() (in module Base), 538

A_ldiv_Bc() (in module Base), 538

A_ldiv_Bt() (in module Base), 538

A_mul_B

() (in module Base), 538

A_mul_Bc() (in module Base), 538

A_mul_Bt() (in module Base), 538

A_rdiv_Bc() (in module Base), 539

A_rdiv_Bt() (in module Base), 539

abs() (in module Base), 418

abs2() (in module Base), 419

abspath() (in module Base), 559

AbstractLock (in module Base), 497

Ac_ldiv_B

() (in module Base), 539

Ac_ldiv_B() (in module Base), 539

Ac_ldiv_Bc() (in module Base), 539
 Ac_mul_B() (in module Base), 539
 Ac_mul_Bc() (in module Base), 539
 Ac_rdiv_B() (in module Base), 539
 Ac_rdiv_Bc() (in module Base), 539
 accept() (in module Base), 573
 accumulate
 () (in module Base), 471
 accumulate() (in module Base), 471
 acos() (in module Base), 412
 acosd() (in module Base), 413
 acosh() (in module Base), 413
 acot() (in module Base), 413
 acotd() (in module Base), 413
 acoth() (in module Base), 414
 acquire() (in module Base), 498
 acsc() (in module Base), 413
 acscd() (in module Base), 413
 acsch() (in module Base), 414
 add() (in module Base.Pkg), 583
 addprocs() (in module Base), 488, 489
 airy() (in module Base), 422
 airyai() (in module Base), 422
 airyaiprime() (in module Base), 423
 airybi() (in module Base), 423
 airybiprime() (in module Base), 423
 airyprime() (in module Base), 423
 airyx() (in module Base), 423
 all
 () (in module Base), 390
 all() (in module Base), 389, 390
 allunique() (in module Base), 385
 angle() (in module Base), 420
 ans (in module Base), 362
 any
 () (in module Base), 389
 ANY (in module Base), 553
 any() (in module Base), 389, 390
 append
 () (in module Base), 401
 applicable() (in module Base), 368
 apropos() (in module Base), 361
 ARGS (in module Base), 553
 ArgumentError() (in module Base), 374
 Array() (in module Base), 454
 ascii() (in module Base), 444
 asec() (in module Base), 413
 asecd() (in module Base), 413
 asech() (in module Base), 414
 asin() (in module Base), 412
 asind() (in module Base), 413
 asinh() (in module Base), 413
 assert() (in module Base), 374
 AssertionError() (in module Base), 374

asum() (in module Base.LinAlg.BLAS), 540
 AsyncCondition() (in module Base), 376
 asyncmap() (in module Base), 489
 At_ldiv_B
 () (in module Base), 539
 At_ldiv_B() (in module Base), 539
 At_ldiv_Bt() (in module Base), 539
 At_mul_B() (in module Base), 539
 At_mul_Bt() (in module Base), 539
 At_rdiv_B() (in module Base), 539
 At_rdiv_Bt() (in module Base), 539
 atan() (in module Base), 412
 atan2() (in module Base), 413
 atand() (in module Base), 413
 atanh() (in module Base), 414
 atexit() (in module Base), 359
 atreplinit() (in module Base), 359
 available() (in module Base.Pkg), 584
 axpy
 () (in module Base.LinAlg.BLAS), 540

B

backtrace() (in module Base), 374
 baremodule, 139
 base() (in module Base), 433
 Base.add_with_overflow() (in module Base), 419
 Base.channel_from_id() (in module Base), 494
 Base.checked_abs() (in module Base), 418
 Base.checked_add() (in module Base), 418
 Base.checked_cld() (in module Base), 419
 Base.checked_div() (in module Base), 418
 Base.checked_fld() (in module Base), 419
 Base.checked_mod() (in module Base), 419
 Base.checked_mul() (in module Base), 418
 Base.checked_neg() (in module Base), 418
 Base.checked_rem() (in module Base), 419
 Base.checked_sub() (in module Base), 418
 Base.cluster_cookie() (in module Base), 494
 Base.Collections (module), 401
 Base.compilecache() (in module Base), 360
 Base.datatype_module() (in module Base), 377
 Base.datatype_name() (in module Base), 377
 Base.function_module() (in module Base), 377
 Base.function_name() (in module Base), 377
 Base.Iterators (module), 595
 Base.LinAlg (module), 501
 Base.LinAlg.BLAS (module), 539
 Base.LinAlg.LAPACK (module), 543
 Base.linearindexing() (in module Base), 452
 Base.mul_with_overflow() (in module Base), 419
 Base.OneTo() (in module Base), 409
 Base.Pkg (module), 581
 Base.process_messages() (in module Base), 499
 Base.remoteref_id() (in module Base), 494

Base.sub_with_overflow() (in module Base), 419
 Base.summarysize() (in module Base), 360
 Base.Test (module), 601
 Base.worker_id_from_socket() (in module Base), 494
 base64decode() (in module Base), 569
 Base64DecodePipe() (in module Base), 569
 base64encode() (in module Base), 569
 Base64EncodePipe() (in module Base), 569
 basename() (in module Base), 558
 bdsdc
 () (in module Base.LinAlg.LAPACK), 550
 bdsqr
 () (in module Base.LinAlg.LAPACK), 550
 besselh() (in module Base), 423
 besselhx() (in module Base), 423
 besseli() (in module Base), 424
 besselix() (in module Base), 424
 besselj() (in module Base), 423
 besselj0() (in module Base), 423
 besselj1() (in module Base), 423
 besseljx() (in module Base), 423
 besserk() (in module Base), 424
 besserkx() (in module Base), 424
 bessely() (in module Base), 423
 bessely0() (in module Base), 423
 bessely1() (in module Base), 423
 besselyx() (in module Base), 423
 beta() (in module Base), 424
 bfft
 () (in module Base), 428
 bfft() (in module Base), 428
 Bidiagonal() (in module Base), 503–505
 big() (in module Base), 434
 BigFloat() (in module Base), 437
 BigInt() (in module Base), 437
 bin() (in module Base), 433
 bind() (in module Base), 574
 binomial() (in module Base), 420
 BitArray() (in module Base), 455
 BitArray{N}() (in module Base), 455
 bitbroadcast() (in module Base), 460
 bitrand() (in module Base), 441
 bits() (in module Base), 434
 bkfact
 () (in module Base), 514
 bkfact() (in module Base), 514
 blascopy
 () (in module Base.LinAlg.BLAS), 540
 blkdiag() (in module Base), 531
 BoundsError() (in module Base), 374
 brfft() (in module Base), 429
 broadcast
 () (in module Base), 460
 broadcast() (in module Base), 459

broadcast_getindex() (in module Base), 461
 broadcast_setindex
 () (in module Base), 462
 bswap() (in module Base), 434
 build() (in module Base.Pkg), 584
 bytes2hex() (in module Base), 435

C

C_NULL (in module Base), 553
 CachingPool() (in module Base), 492
 callers() (in module Base.Profile), 618
 calloc() (in module Base.Libc), 613
 cat() (in module Base), 462
 catalan (in module Base), 435
 catch_backtrace() (in module Base), 374
 cbrt() (in module Base), 420
 ccall() (in module Base), 607
 Cchar (in module Base), 609
 cconvert() (in module Base), 607
 cd() (in module Base), 555
 Cdouble (in module Base), 610
 ceil() (in module Base), 417
 ceil() (in module Dates), 593
 Cfloat (in module Base), 610
 cfunction() (in module Base), 607
 cgloball() (in module Base), 607
 Channel{T}() (in module Base), 487
 charwidth() (in module Base), 449
 checkbounds() (in module Base), 470
 checkindex() (in module Base), 470
 checkout() (in module Base.Pkg), 584
 chmod() (in module Base), 556
 chol() (in module Base), 509
 cholfact
 () (in module Base), 510, 511
 cholfact() (in module Base), 509, 510
 chomp() (in module Base), 448
 chop() (in module Base), 448
 chown() (in module Base), 556
 chr2ind() (in module Base), 449
 Cint (in module Base), 609
 Cintmax_t (in module Base), 610
 circcopy
 () (in module Base), 464
 circshift
 () (in module Base), 464
 circshift() (in module Base), 464
 cis() (in module Base), 420
 clamp
 () (in module Base), 418
 clamp() (in module Base), 418
 cld() (in module Base), 407
 clear
 () (in module Base), 494

`clear()` (in module `Base.Profile`), 617
`clear_malloc_data()` (in module `Base.Profile`), 618
`clipboard()` (in module `Base`), 360
`clone()` (in module `Base.Pkg`), 583, 584
`Clong` (in module `Base`), 609
`Clonglong` (in module `Base`), 609
`close()` (in module `Base`), 487, 562
`cmp()` (in module `Base`), 411
`code_llvm()` (in module `Base`), 379
`code_lowered()` (in module `Base`), 378
`code_native()` (in module `Base`), 379
`code_typed()` (in module `Base`), 379
`code_warntype()` (in module `Base`), 379
`collect()` (in module `Base`), 392
`colon()` (in module `Base`), 409
`complex()` (in module `Base`), 434
`CompoundPeriod()` (in module `Dates`), 593
`cond()` (in module `Base`), 527
`Condition()` (in module `Base`), 486
`condskeel()` (in module `Base`), 527
`conj`
 `()` (in module `Base`), 453
`conj()` (in module `Base`), 420
`connect()` (in module `Base`), 499, 573
`consume()` (in module `Base`), 486
`contains()` (in module `Base`), 446
`conv()` (in module `Base`), 431
`conv2()` (in module `Base`), 431
`convert()` (in module `Base`), 364
`copy`
 `()` (in module `Base`), 608
`copy()` (in module `Base`), 363
`copysign()` (in module `Base`), 419
`cor()` (in module `Base`), 427
`cos()` (in module `Base`), 412
`cosc()` (in module `Base`), 414
`cosd()` (in module `Base`), 412
`cosh()` (in module `Base`), 412
`cospi()` (in module `Base`), 412
`cot()` (in module `Base`), 413
`cotd()` (in module `Base`), 413
`coth()` (in module `Base`), 413
`count()` (in module `Base`), 390
`count_ones()` (in module `Base`), 438
`count_zeros()` (in module `Base`), 439
`countfrom()` (in module `Base.Iterators`), 597
`countlines()` (in module `Base`), 565
`countnz()` (in module `Base`), 453
`cov()` (in module `Base`), 427
`cp()` (in module `Base`), 557
`Cptrdiff_t` (in module `Base`), 610
`cross()` (in module `Base`), 502
`csc()` (in module `Base`), 413
`cscd()` (in module `Base`), 413

`csch()` (in module `Base`), 413
`Cshort` (in module `Base`), 609
`Csize_t` (in module `Base`), 610
`Cssize_t` (in module `Base`), 610
`ctime()` (in module `Base`), 556
`ctranspose`
 `()` (in module `Base`), 535
`ctranspose()` (in module `Base`), 534
`Cuchar` (in module `Base`), 609
`Cuint` (in module `Base`), 609
`Cuintmax_t` (in module `Base`), 610
`Culong` (in module `Base`), 609
`Culonglong` (in module `Base`), 610
`cummax()` (in module `Base`), 472
`cummin()` (in module `Base`), 472
`cumprod`
 `()` (in module `Base`), 471
`cumprod()` (in module `Base`), 471
`cumsum`
 `()` (in module `Base`), 472
`cumsum()` (in module `Base`), 471
`cumsum_kbn()` (in module `Base`), 472
`current_module()` (in module `Base`), 376
`current_task()` (in module `Base`), 485
`Cushort` (in module `Base`), 609
`Cwchar_t` (in module `Base`), 610
`cycle()` (in module `Base.Iterators`), 598

D

`Date` (in module `Dates`), 587
`Date()` (in module `Dates`), 589
`date2epochdays()` (in module `Dates`), 594
`DateFormat()` (in module `Dates`), 589
`Dates` (module), 585
`DateTime` (in module `Dates`), 587
`DateTime()` (in module `Dates`), 588, 589
`datetime2epochms()` (in module `Dates`), 594
`datetime2julian()` (in module `Dates`), 595
`datetime2rata()` (in module `Dates`), 595
`datetime2unix()` (in module `Dates`), 595
`dawson()` (in module `Base`), 420
`Day` (in module `Dates`), 587, 592
`Day()` (in module `Dates`), 590
`day()` (in module `Dates`), 590
`dayabbr()` (in module `Dates`), 591
`dayname()` (in module `Dates`), 591
`dayofmonth()` (in module `Dates`), 591
`dayofquarter()` (in module `Dates`), 591
`dayofweek()` (in module `Dates`), 591
`dayofweekofmonth()` (in module `Dates`), 591
`dayofyear()` (in module `Dates`), 591
`daysinmonth()` (in module `Dates`), 591
`daysinyear()` (in module `Dates`), 591
`daysofweekinmonth()` (in module `Dates`), 591

dct
 () (in module Base), 430
 dct() (in module Base), 430
 dec() (in module Base), 433
 deconv() (in module Base), 431
 deepcopy() (in module Base), 363
 default() (in module Dates), 593
 default_worker_pool() (in module Base), 492
 deg2rad() (in module Base), 414
 delete
 () (in module Base), 395
 deleteat
 () (in module Base), 399
 denominator() (in module Base), 408
 dependents() (in module Base.Pkg), 585
 dequeue
 () (in module Base.Collections), 402
 deserialize() (in module Base), 564
 det() (in module Base), 528
 detach() (in module Base), 371
 DevNull (in module Base), 370
 diag() (in module Base), 524
 diagind() (in module Base), 524
 diagm() (in module Base), 525
 Diagonal() (in module Base), 503
 Dict() (in module Base), 394
 diff() (in module Base), 472
 digamma() (in module Base), 422
 digits
 () (in module Base), 433
 digits() (in module Base), 433
 DimensionMismatch() (in module Base), 374
 dir() (in module Base.Pkg), 583
 dirname() (in module Base), 558
 disable_sigint() (in module Base), 609
 display() (in module Base), 569
 displayable() (in module Base), 570
 displaysize() (in module Base), 569
 div() (in module Base), 407
 DivideError() (in module Base), 374
 divrem() (in module Base), 407
 DL_LOAD_PATH (in module Base.Libdl), 616
 dlclose() (in module Base.Libdl), 615
 dlex() (in module Base.Libdl), 615
 dlopen() (in module Base.Libdl), 615
 dlopen_e() (in module Base.Libdl), 615
 dlsym() (in module Base.Libdl), 615
 dlsym_e() (in module Base.Libdl), 615
 DomainError() (in module Base), 375
 done() (in module Base), 381
 dot() (in module Base), 501
 dot() (in module Base.LinAlg.BLAS), 540
 dotc() (in module Base.LinAlg.BLAS), 540
 dotu() (in module Base.LinAlg.BLAS), 540

download() (in module Base), 557
 drop() (in module Base.Iterators), 598
 dropzeros
 () (in module Base), 483
 dropzeros() (in module Base), 483
 dump() (in module Base), 567

E

e (in module Base), 435
 eachindex() (in module Base), 452
 eachline() (in module Base), 568
 eachmatch() (in module Base), 445
 edit() (in module Base), 360
 edit() (in module Base.Pkg), 583
 eig() (in module Base), 514
 eigfact
 () (in module Base), 517
 eigfact() (in module Base), 516, 517
 eigmax() (in module Base), 515
 eigmin() (in module Base), 515
 eigs() (in module Base), 535, 536
 eigvals
 () (in module Base), 515
 eigvals() (in module Base), 515
 eigvecs() (in module Base), 516
 eltype() (in module Base), 384
 empty
 () (in module Base), 382
 ENDIAN_BOM (in module Base), 574
 endof() (in module Base), 383
 endswith() (in module Base), 448
 enqueue
 () (in module Base.Collections), 402
 enumerate() (in module Base.Iterators), 597
 ENV (in module Base), 373
 EnvHash() (in module Base), 373
 eof() (in module Base), 563
 EOFError() (in module Base), 375
 epochdays2date() (in module Dates), 594
 epochms2datetime() (in module Dates), 594
 eps() (in module Base), 366
 eps() (in module Dates), 590
 erf() (in module Base), 420
 erfc() (in module Base), 420
 erfcinv() (in module Base), 420
 erfcx() (in module Base), 420
 erfi() (in module Base), 420
 erfinv() (in module Base), 420
 errno() (in module Base.Libc), 613
 error() (in module Base), 374
 ErrorException() (in module Base), 375
 esc() (in module Base), 368
 escape_string() (in module Base), 450
 eta() (in module Base), 424

eu (in module Base), 435
eulergamma (in module Base), 435
eval() (in module Base), 368
evalfile() (in module Base), 368
exit() (in module Base), 359
exp() (in module Base), 415
exp10() (in module Base), 415
exp2() (in module Base), 415
expand() (in module Base), 378
expanduser() (in module Base), 559
expm() (in module Base), 531
expm1() (in module Base), 415
exponent() (in module Base), 434
export, 139
extrema() (in module Base), 386
eye() (in module Base), 458

F

factorial() (in module Base), 420
factorize() (in module Base), 502
falses() (in module Base), 456
fd() (in module Base), 564
fdio() (in module Base), 562
fetch() (in module Base), 487, 491
fetch() (in module Base.Profile), 617
fft
 () (in module Base), 428
fft() (in module Base), 428
fftshift() (in module Base), 430
fieldname() (in module Base), 377
fieldnames() (in module Base), 377
fieldoffset() (in module Base), 366
fieldtype() (in module Base), 367
filemode() (in module Base), 556
filesize() (in module Base), 556
fill
 () (in module Base), 456
fill() (in module Base), 456
filt
 () (in module Base), 431
filt() (in module Base), 431
filter
 () (in module Base), 392
filter() (in module Base), 392
finalize() (in module Base), 363
finalizer() (in module Base), 363
find() (in module Base), 465
find_library() (in module Base.Libdl), 615
findfirst() (in module Base), 466
findin() (in module Base), 384
findlast() (in module Base), 467
findmax
 () (in module Base), 387
findmax() (in module Base), 387

findmin
 () (in module Base), 387
findmin() (in module Base), 387
findn() (in module Base), 465
findnext() (in module Base), 467, 468
findnz() (in module Base), 466
findprev() (in module Base), 468, 469
finish() (in module Base.Test), 606
first() (in module Base), 391
firstdayofmonth() (in module Dates), 592
firstdayofquarter() (in module Dates), 592
firstdayofweek() (in module Dates), 592
firstdayofyear() (in module Dates), 592
flatten() (in module Base.Iterators), 599
fld() (in module Base), 407
fld1() (in module Base), 408
fldmod() (in module Base), 408
fldmod1() (in module Base), 408
flipbits
 () (in module Base), 477
flipdim() (in module Base), 464
flipsign() (in module Base), 419
float() (in module Base), 434
Float32() (in module Base), 437
Float64() (in module Base), 437
floor() (in module Base), 417
floor() (in module Dates), 593
floorceil() (in module Dates), 594
flush() (in module Base), 562
flush_cstdio() (in module Base.Libc), 614
fma() (in module Base), 407
foldl() (in module Base), 385
foldr() (in module Base), 385
foreach() (in module Base), 390
format() (in module Dates), 588
FormatMessage() (in module Base.Libc), 613
free() (in module Base.Libc), 613
free() (in module Base.Pkg), 584
frexp() (in module Base), 415
full() (in module Base), 480, 503, 514
fullname() (in module Base), 377
functionloc() (in module Base), 377
Future() (in module Base), 490

G

gamma() (in module Base), 422
gbmv
 () (in module Base.LinAlg.BLAS), 540
gbmv() (in module Base.LinAlg.BLAS), 541
gbtrf
 () (in module Base.LinAlg.LAPACK), 543
gbtrs
 () (in module Base.LinAlg.LAPACK), 543
gc() (in module Base), 378

- gc_enable() (in module Base), 378
 - gcd() (in module Base), 421
 - gcdx() (in module Base), 421
 - gebak
 - () (in module Base.LinAlg.LAPACK), 543
 - gebal
 - () (in module Base.LinAlg.LAPACK), 543
 - gebrd
 - () (in module Base.LinAlg.LAPACK), 543
 - gecon
 - () (in module Base.LinAlg.LAPACK), 551
 - gees
 - () (in module Base.LinAlg.LAPACK), 551
 - geev
 - () (in module Base.LinAlg.LAPACK), 546
 - geevx
 - () (in module Base.LinAlg.LAPACK), 546
 - gehrd
 - () (in module Base.LinAlg.LAPACK), 551
 - gelqf
 - () (in module Base.LinAlg.LAPACK), 543
 - gels
 - () (in module Base.LinAlg.LAPACK), 545
 - gelsd
 - () (in module Base.LinAlg.LAPACK), 546
 - gelsy
 - () (in module Base.LinAlg.LAPACK), 546
 - gemm
 - () (in module Base.LinAlg.BLAS), 541
 - gemm() (in module Base.LinAlg.BLAS), 541
 - gemqrt
 - () (in module Base.LinAlg.LAPACK), 548
 - gemv
 - () (in module Base.LinAlg.BLAS), 541
 - gemv() (in module Base.LinAlg.BLAS), 541
 - gensym() (in module Base), 368
 - geqlf
 - () (in module Base.LinAlg.LAPACK), 543
 - geqp3
 - () (in module Base.LinAlg.LAPACK), 544
 - geqrf
 - () (in module Base.LinAlg.LAPACK), 544
 - geqrt
 - () (in module Base.LinAlg.LAPACK), 544
 - geqrt3
 - () (in module Base.LinAlg.LAPACK), 544, 545
 - ger
 - () (in module Base.LinAlg.BLAS), 540
 - gerqf
 - () (in module Base.LinAlg.LAPACK), 544
 - gesdd
 - () (in module Base.LinAlg.LAPACK), 546
 - gesv
 - () (in module Base.LinAlg.LAPACK), 545
 - gesvd
 - () (in module Base.LinAlg.LAPACK), 546
 - gesvx
 - () (in module Base.LinAlg.LAPACK), 545, 546
 - get
 - () (in module Base), 394
 - get() (in module Base), 369, 394
 - get_testset() (in module Base.Test), 606
 - get_testset_depth() (in module Base.Test), 606
 - get_zero_subnormals() (in module Base), 438
 - getaddrinfo() (in module Base), 573
 - getfield() (in module Base), 366
 - gethostname() (in module Base), 372
 - getindex() (in module Base), 393, 454, 460
 - getipaddr() (in module Base), 372
 - getkey() (in module Base), 395
 - GetLastError() (in module Base.Libc), 613
 - getpid() (in module Base), 372
 - getrf
 - () (in module Base.LinAlg.LAPACK), 545
 - getri
 - () (in module Base.LinAlg.LAPACK), 545
 - getrs
 - () (in module Base.LinAlg.LAPACK), 545
 - getsockname() (in module Base), 573
 - gges
 - () (in module Base.LinAlg.LAPACK), 551
 - ggevd
 - () (in module Base.LinAlg.LAPACK), 547
 - ggls
 - () (in module Base.LinAlg.LAPACK), 546
 - ggsvd
 - () (in module Base.LinAlg.LAPACK), 546
 - ggsvd3
 - () (in module Base.LinAlg.LAPACK), 546
 - givens() (in module Base), 521, 522
 - givens{T}() (in module Base), 521
 - golden (in module Base), 435
 - gperm() (in module Base), 556
 - gradient() (in module Base), 472
 - graphemes() (in module Base), 445
 - gtsv
 - () (in module Base.LinAlg.LAPACK), 547
 - gttrf
 - () (in module Base.LinAlg.LAPACK), 547
 - gttrs
 - () (in module Base.LinAlg.LAPACK), 547
- ## H
- hankelh1() (in module Base), 423
 - hankelh1x() (in module Base), 423
 - hankelh2() (in module Base), 423
 - hankelh2x() (in module Base), 423
 - hash() (in module Base), 363

haskey() (in module Base), 394
hcat() (in module Base), 462
heapify
 () (in module Base.Collections), 403
heapify() (in module Base.Collections), 403
heappop
 () (in module Base.Collections), 404
heappush
 () (in module Base.Collections), 404
her
 () (in module Base.LinAlg.BLAS), 540
herk
 () (in module Base.LinAlg.BLAS), 540
herk() (in module Base.LinAlg.BLAS), 540
Hermitian() (in module Base), 507
hessfact
 () (in module Base), 517
hessfact() (in module Base), 517
hesv
 () (in module Base.LinAlg.LAPACK), 550
hetrf
 () (in module Base.LinAlg.LAPACK), 550
hetri
 () (in module Base.LinAlg.LAPACK), 550
hetrs
 () (in module Base.LinAlg.LAPACK), 550
hex() (in module Base), 433
hex2bytes() (in module Base), 435
hex2num() (in module Base), 435
homedir() (in module Base), 558
Hour (in module Dates), 587, 592
Hour() (in module Dates), 590
hour() (in module Dates), 590
htol() (in module Base), 574
hton() (in module Base), 574
hvcat() (in module Base), 463
hypot() (in module Base), 414

|
I (in module Base.LinAlg.BLAS), 542
idct
 () (in module Base), 430
idct() (in module Base), 430
identity() (in module Base), 365
ifelse() (in module Base), 362
ifft
 () (in module Base), 428
ifft() (in module Base), 428
ifftshift() (in module Base), 431
ignorestatus() (in module Base), 371
im (in module Base), 435
imag() (in module Base), 420
import, 139
importall, 139

in() (in module Base), 383
include() (in module Base), 361
include_dependency() (in module Base), 361
include_string() (in module Base), 361
ind2chr() (in module Base), 449
ind2sub() (in module Base), 453, 454
indexin() (in module Base), 384
indexpids() (in module Base), 495
indices() (in module Base), 451
indmax() (in module Base), 386
indmin() (in module Base), 386
InexactError() (in module Base), 375
Inf (in module Base), 436
Inf16 (in module Base), 436
Inf32 (in module Base), 436
info() (in module Base), 567
init() (in module Base.Pkg), 583
init() (in module Base.Profile), 617
init_worker() (in module Base), 499
InitError() (in module Base), 375
insert
 () (in module Base), 399
installed() (in module Base.Pkg), 584
instances() (in module Base), 367
Instant (in module Dates), 587
interrupt() (in module Base), 489
InterruptedException() (in module Base), 375
intersect
 () (in module Base), 397
intersect() (in module Base), 397
IntSet() (in module Base), 396
inv() (in module Base), 528
invdigamma() (in module Base), 422
invmod() (in module Base), 422
invoke() (in module Base), 368
invperm() (in module Base), 476
IOBuffer() (in module Base), 561, 562
IOContext() (in module Base), 565, 566
ipermute
 () (in module Base), 477
IPv4() (in module Base), 573
IPv6() (in module Base), 573
irfft() (in module Base), 429
is_apple() (in module Base), 373
is_assigned_char() (in module Base), 445
is_bsd() (in module Base), 374
is_linux() (in module Base), 373
is_unix() (in module Base), 373
is_windows() (in module Base), 374
isa() (in module Base), 362
isabspath() (in module Base), 558
isalnum() (in module Base), 449
isalpha() (in module Base), 449
isapprox() (in module Base), 412

isascii() (in module Base), 449
 isassigned() (in module Base), 462
 isbits() (in module Base), 367
 isblockdev() (in module Base), 557
 ischardev() (in module Base), 557
 iscntrl() (in module Base), 449
 isconst() (in module Base), 377
 isdefined() (in module Base), 363
 isdiag() (in module Base), 533
 isdigit() (in module Base), 449
 isdir() (in module Base), 557
 isdirpath() (in module Base), 558
 isempty() (in module Base), 382
 isequal() (in module Base), 362
 iseven() (in module Base), 439
 isfifo() (in module Base), 557
 isfile() (in module Base), 557
 isfinite() (in module Base), 436
 isgraph() (in module Base), 449
 isheap() (in module Base.Collections), 403
 ishermitian() (in module Base), 534
 isimag() (in module Base), 437
 isimmutable() (in module Base), 367
 isinf() (in module Base), 436
 isinteger() (in module Base), 436
 isinteractive() (in module Base), 359
 isleatype() (in module Base), 367
 isleapyear() (in module Dates), 591
 isless() (in module Base), 362
 islink() (in module Base), 558
 islocked() (in module Base), 498
 islower() (in module Base), 449
 ismarked() (in module Base), 563
 ismatch() (in module Base), 445
 ismount() (in module Base), 558
 isnan() (in module Base), 436
 isnull() (in module Base), 369
 isnumber() (in module Base), 449
 isodd() (in module Base), 439
 isopen() (in module Base), 564
 ispath() (in module Base), 558
 isperm() (in module Base), 476
 isposdef
 () (in module Base), 533
 isposdef() (in module Base), 532
 ispow2() (in module Base), 421
 isprint() (in module Base), 449
 ispunct() (in module Base), 450
 isqrt() (in module Base), 420
 isreadable() (in module Base), 564
 isreadonly() (in module Base), 564
 isready() (in module Base), 487, 491, 492
 isreal() (in module Base), 436
 issetgid() (in module Base), 558

issetuid() (in module Base), 558
 issocket() (in module Base), 558
 issorted() (in module Base), 579
 isspace() (in module Base), 450
 issparse() (in module Base), 480
 issticky() (in module Base), 558
 issubnormal() (in module Base), 436
 issubset() (in module Base), 392, 397
 issubtype() (in module Base), 365
 issymmetric() (in module Base), 532
 istaskdone() (in module Base), 485
 istaskstarted() (in module Base), 485
 istextmime() (in module Base), 571
 istril() (in module Base), 533
 istriu() (in module Base), 533
 isupper() (in module Base), 450
 isvalid() (in module Base), 445
 iswritable() (in module Base), 564
 isxdigit() (in module Base), 450
 iteratorretype() (in module Base), 382
 iteratorsize() (in module Base), 381

J

join() (in module Base), 448
 joinpath() (in module Base), 559
 JULIA_HOME (in module Base), 553
 julian2datetime() (in module Dates), 595

K

KeyError() (in module Base), 375
 keys() (in module Base), 395
 keytype() (in module Base), 396
 kill() (in module Base), 371, 499
 kron() (in module Base), 530

L

last() (in module Base), 391
 lastdayofmonth() (in module Dates), 592
 lastdayofquarter() (in module Dates), 592
 lastdayofweek() (in module Dates), 592
 lastdayofyear() (in module Dates), 592
 launch() (in module Base), 498
 lbeta() (in module Base), 424
 lcfirst() (in module Base), 448
 lcm() (in module Base), 421
 ldexp() (in module Base), 415
 ldlfact
 () (in module Base), 512
 ldlfact() (in module Base), 511
 leading_ones() (in module Base), 439
 leading_zeros() (in module Base), 439
 length() (in module Base), 383, 443, 451
 less() (in module Base), 360
 lexcmp() (in module Base), 363

lexless() (in module Base), 363
lfact() (in module Base), 422
lgamma() (in module Base), 422
Libc (module), 611
Libdl (module), 614
LinAlg.checksquare() (in module Base), 454
LinAlg.Givens() (in module Base), 521
LinAlg.qr
 () (in module Base), 512
linearindices() (in module Base), 452
linreg() (in module Base), 531
linspace() (in module Base), 458
listen() (in module Base), 573
listenany() (in module Base), 573
llvmcall() (in module Base), 611
LOAD_PATH (in module Base), 553
LoadError() (in module Base), 375
localindexes() (in module Base), 495
lock() (in module Base), 497
log() (in module Base), 414
log10() (in module Base), 415
log1p() (in module Base), 415
log2() (in module Base), 415
logabsdet() (in module Base), 528
logdet() (in module Base), 528
logm() (in module Base), 531
logspace() (in module Base), 458
lowercase() (in module Base), 448
lowrankdownate
 () (in module Base.LinAlg), 511
lowrankdownate() (in module Base.LinAlg), 511
lowrankupdate
 () (in module Base.LinAlg), 511
lowrankupdate() (in module Base.LinAlg), 511
lpad() (in module Base), 445
lq() (in module Base), 514
lqfact
 () (in module Base), 514
lqfact() (in module Base), 514
lstat() (in module Base), 556
lstrip() (in module Base), 447
ltoh() (in module Base), 574
lu() (in module Base), 507
lufact
 () (in module Base), 509
lufact() (in module Base), 507, 508
lyap() (in module Base), 532

M

macroexpand() (in module Base), 378
malloc() (in module Base.Libc), 613
manage() (in module Base), 499
map
 () (in module Base), 390

map() (in module Base), 390
mapfoldl() (in module Base), 391
mapfoldr() (in module Base), 391
mapreduce() (in module Base), 390, 391
mapreducedim() (in module Base), 475
mapslices() (in module Base), 475
mark() (in module Base), 563
match() (in module Base), 445
matchall() (in module Base), 445
max() (in module Base), 418
maxabs
 () (in module Base), 387
maxabs() (in module Base), 387
maximum
 () (in module Base), 385
maximum() (in module Base), 385
maxintfloat() (in module Base), 366
mean
 () (in module Base), 425
mean() (in module Base), 425
median
 () (in module Base), 426
median() (in module Base), 426
merge
 () (in module Base), 396
merge() (in module Base), 395
MersenneTwister() (in module Base), 440
method_exists() (in module Base), 368
MethodError() (in module Base), 375
methods() (in module Base), 361
methodswith() (in module Base), 361
middle() (in module Base), 426
midpoints() (in module Base), 426
Millisecond (in module Dates), 587, 592
Millisecond() (in module Dates), 590
millisecond() (in module Dates), 590
mimewritable() (in module Base), 570
min() (in module Base), 417
minabs
 () (in module Base), 388
minabs() (in module Base), 387, 388
minimum
 () (in module Base), 386
minimum() (in module Base), 385, 386
minmax() (in module Base), 418
Minute (in module Dates), 587, 592
Minute() (in module Dates), 590
minute() (in module Dates), 590
mkdir() (in module Base), 555
mkpath() (in module Base), 555
mktemp() (in module Base), 557
mktempdir() (in module Base), 557
Mmap.Anonymous() (in module Base), 571
Mmap.mmap() (in module Base), 571, 572

Mmap.sync

 () (in module Base), 572
 mod() (in module Base), 407
 mod1() (in module Base), 408
 mod2pi() (in module Base), 407
 modf() (in module Base), 415
 module, 139
 module_name() (in module Base), 376
 module_parent() (in module Base), 376
 Month (in module Dates), 587, 592
 Month() (in module Dates), 590
 month() (in module Dates), 590
 monthabbr() (in module Dates), 591
 monthday() (in module Dates), 591
 monthname() (in module Dates), 591
 mtime() (in module Base), 556
 muladd() (in module Base), 407
 Mutex() (in module Base), 498
 mv() (in module Base), 557
 myid() (in module Base), 489

N

names() (in module Base), 377
 NaN (in module Base), 436
 NaN16 (in module Base), 436
 NaN32 (in module Base), 436
 nb_available() (in module Base), 573
 ndigits() (in module Base), 424
 ndims() (in module Base), 451
 next() (in module Base), 381
 nextfloat() (in module Base), 436
 nextind() (in module Base), 449
 nextpow() (in module Base), 422
 nextpow2() (in module Base), 421
 nextprod() (in module Base), 422
 nfields() (in module Base), 377
 nnz() (in module Base), 480
 nonzeros() (in module Base), 482
 norm() (in module Base), 526
 normalize
 () (in module Base), 527
 normalize() (in module Base), 527
 normalize_string() (in module Base), 444
 normpath() (in module Base), 559
 nothing (in module Base), 553
 notify() (in module Base), 486
 now() (in module Dates), 590
 nprocs() (in module Base), 489
 nrm2() (in module Base.LinAlg.BLAS), 540
 ntoh() (in module Base), 574
 ntuple() (in module Base), 363
 Nullable() (in module Base), 369
 NullException() (in module Base), 375
 nullspace() (in module Base), 529

num2hex() (in module Base), 434
 numerator() (in module Base), 408
 nworkers() (in module Base), 489
 nzrange() (in module Base), 482

O

object_id() (in module Base), 363
 oct() (in module Base), 433
 oftype() (in module Base), 365
 one() (in module Base), 435
 ones() (in module Base), 455
 open() (in module Base), 561
 operm() (in module Base), 556
 ordschur
 () (in module Base), 518, 519
 ordschur() (in module Base), 518, 519
 orghr
 () (in module Base.LinAlg.LAPACK), 551
 orglq
 () (in module Base.LinAlg.LAPACK), 547
 orgql
 () (in module Base.LinAlg.LAPACK), 547
 orgqr
 () (in module Base.LinAlg.LAPACK), 547
 orgrq
 () (in module Base.LinAlg.LAPACK), 547
 ormlq
 () (in module Base.LinAlg.LAPACK), 547
 ormql
 () (in module Base.LinAlg.LAPACK), 547
 ormqr
 () (in module Base.LinAlg.LAPACK), 547
 ormrq
 () (in module Base.LinAlg.LAPACK), 547
 ormrz
 () (in module Base.LinAlg.LAPACK), 545
 OutOfMemoryError() (in module Base), 375
 OverflowError() (in module Base), 375

P

parent() (in module Base), 461
 parentindexes() (in module Base), 461
 parse() (in module Base), 369, 434
 ParseError() (in module Base), 375
 partition() (in module Base.Iterators), 599
 peakflops() (in module Base), 538
 peek() (in module Base.Collections), 402
 Period (in module Dates), 587
 permute
 () (in module Base), 477
 permutedims
 () (in module Base), 469
 permutedims() (in module Base), 469
 pi (in module Base), 435

`pin()` (in module `Base.Pkg`), 584
`pinv()` (in module `Base`), 529
`PipeBuffer()` (in module `Base`), 565
`pipeline()` (in module `Base`), 372
`plan_bfft`
 `()` (in module `Base`), 429
`plan_bfft()` (in module `Base`), 429
`plan_brfft()` (in module `Base`), 430
`plan_dct`
 `()` (in module `Base`), 430
`plan_dct()` (in module `Base`), 430
`plan_fft`
 `()` (in module `Base`), 429
`plan_fft()` (in module `Base`), 428
`plan_idct`
 `()` (in module `Base`), 430
`plan_idct()` (in module `Base`), 430
`plan_ift`
 `()` (in module `Base`), 429
`plan_ift()` (in module `Base`), 429
`plan_irfft()` (in module `Base`), 430
`plan_r2r`
 `()` (in module `Base.FFTW`), 431
`plan_r2r()` (in module `Base.FFTW`), 431
`plan_rfft()` (in module `Base`), 430
`pmap()` (in module `Base`), 489
`pointer()` (in module `Base`), 608
`pointer_from_objref()` (in module `Base`), 608
`poll_fd()` (in module `Base`), 573
`poll_file()` (in module `Base`), 573
`polygamma()` (in module `Base`), 422
`pop`
 `()` (in module `Base`), 395, 398
`popdisplay()` (in module `Base`), 571
`position()` (in module `Base`), 563
`posv`
 `()` (in module `Base.LinAlg.LAPACK`), 548
`potrf`
 `()` (in module `Base.LinAlg.LAPACK`), 548
`potri`
 `()` (in module `Base.LinAlg.LAPACK`), 548
`potrs`
 `()` (in module `Base.LinAlg.LAPACK`), 548
`powermod()` (in module `Base`), 422
`precision()` (in module `Base`), 439
`precompile()` (in module `Base`), 379
`prepend`
 `()` (in module `Base`), 401
`prevfloat()` (in module `Base`), 436
`prevind()` (in module `Base`), 449
`prevpow()` (in module `Base`), 422
`prevpow2()` (in module `Base`), 421
`print()` (in module `Base`), 566
`print()` (in module `Base.Profile`), 617

`print_shortest()` (in module `Base`), 564
`print_with_color()` (in module `Base`), 567
`println()` (in module `Base`), 566
`PriorityQueue()` (in module `Base.Collections`), 402
`process_exited()` (in module `Base`), 371
`process_running()` (in module `Base`), 371
`ProcessExitedException()` (in module `Base`), 375
`procs()` (in module `Base`), 489, 495
`prod`
 `()` (in module `Base`), 389
`prod()` (in module `Base`), 388
`produce()` (in module `Base`), 486
`product()` (in module `Base.Iterators`), 598
`Profile` (module), 616
`PROGRAM_FILE` (in module `Base`), 553
`promote()` (in module `Base`), 365
`promote_rule()` (in module `Base`), 366
`promote_shape()` (in module `Base`), 470
`promote_type()` (in module `Base`), 366
`pstrf`
 `()` (in module `Base.LinAlg.LAPACK`), 548
`Ptr{T}` (in module `Base`), 609
`ptsv`
 `()` (in module `Base.LinAlg.LAPACK`), 548
`ptrtf`
 `()` (in module `Base.LinAlg.LAPACK`), 548
`ptrtrs`
 `()` (in module `Base.LinAlg.LAPACK`), 548
`push`
 `()` (in module `Base`), 398
`pushdisplay()` (in module `Base`), 571
`put`
 `()` (in module `Base`), 487, 491
`pwd()` (in module `Base`), 555

Q

`q::AbstractVector{Tq}()` (in module `Base`), 483
`qr()` (in module `Base`), 512
`qrfact`
 `()` (in module `Base`), 514
`qrfact()` (in module `Base`), 513
`quadgk()` (in module `Base`), 432
`quantile`
 `()` (in module `Base`), 427
`quantile()` (in module `Base`), 426
`quarterofyear()` (in module `Dates`), 591
`quit()` (in module `Base`), 359

R

`r2r`
 `()` (in module `Base.FFTW`), 431
`r2r()` (in module `Base.FFTW`), 431
`rad2deg()` (in module `Base`), 414
`rand`

- `()` (in module Base), 440
- `rand()` (in module Base), 440
- `randcycle()` (in module Base), 477
- `randexp`
 - `()` (in module Base), 441
- `randexp()` (in module Base), 441
- `randjump()` (in module Base), 441
- `randn`
 - `()` (in module Base), 441
- `randn()` (in module Base), 441
- `RandomDevice()` (in module Base), 440
- `randperm()` (in module Base), 476
- `randstring()` (in module Base), 449
- `randsubseq`
 - `()` (in module Base), 471
- `randsubseq()` (in module Base), 470
- `range()` (in module Base), 409
- `rank()` (in module Base), 526
- `rata2datetime()` (in module Dates), 595
- `rationalize()` (in module Base), 408
- `read`
 - `()` (in module Base), 562
- `read()` (in module Base), 562, 563
- `readandwrite()` (in module Base), 371
- `readavailable()` (in module Base), 565
- `readbytes`
 - `()` (in module Base), 562
- `readchomp()` (in module Base), 565
- `readcsv()` (in module Base), 569
- `readdir()` (in module Base), 555
- `readdlm()` (in module Base), 568
- `readline()` (in module Base), 567
- `readlines()` (in module Base), 568
- `readlink()` (in module Base), 556
- `ReadOnlyMemoryError()` (in module Base), 375
- `readstring()` (in module Base), 567
- `readuntil()` (in module Base), 568
- `real()` (in module Base), 420
- `realloc()` (in module Base.Libc), 613
- `realmax()` (in module Base), 366
- `realmin()` (in module Base), 366
- `realpath()` (in module Base), 559
- `record()` (in module Base.Test), 605
- `RecursiveSpinLock()` (in module Base), 498
- `recv()` (in module Base), 574
- `recvfrom()` (in module Base), 574
- `redirect_stderr()` (in module Base), 564, 565
- `redirect_stdin()` (in module Base), 565
- `redirect_stdout()` (in module Base), 564
- `redisplay()` (in module Base), 570
- `reduce()` (in module Base), 385
- `reducedim()` (in module Base), 474
- `reenable_sigint()` (in module Base), 609
- `ReentrantLock()` (in module Base), 498
- `Ref{T}` (in module Base), 609
- `reim()` (in module Base), 420
- `reinterpret()` (in module Base), 458
- `release()` (in module Base), 498
- `reload()` (in module Base), 360
- `relpath()` (in module Base), 559
- `rem()` (in module Base), 407
- `remote()` (in module Base), 492
- `remote_do()` (in module Base), 491, 492
- `remotecall()` (in module Base), 491, 492
- `remotecall_fetch()` (in module Base), 491, 492
- `remotecall_wait()` (in module Base), 491, 492
- `RemoteChannel()` (in module Base), 490
- `RemoteException()` (in module Base), 490
- `repeat()` (in module Base), 530
- `repeated()` (in module Base.Iterators), 598
- `replace()` (in module Base), 446
- `repmat()` (in module Base), 529
- `repr()` (in module Base), 443
- `reprmime()` (in module Base), 570
- `require()` (in module Base), 360
- `reset()` (in module Base), 563
- `reshape()` (in module Base), 456
- `resize`
 - `()` (in module Base), 401
- `resolve()` (in module Base.Pkg), 583
- `rest()` (in module Base.Iterators), 597
- `rethrow()` (in module Base), 374
- `retrieve()` (in module Base.Profile), 618
- `retry()` (in module Base), 376
- `reverse`
 - `()` (in module Base), 477
- `reverse()` (in module Base), 446, 477
- `reverseind()` (in module Base), 477
- `rfft()` (in module Base), 429
- `rm()` (in module Base), 557
- `rm()` (in module Base.Pkg), 583
- `rmprocs()` (in module Base), 489
- `rol`
 - `()` (in module Base), 477, 478
- `rol()` (in module Base), 478
- `ror`
 - `()` (in module Base), 478
- `ror()` (in module Base), 478
- `rot180()` (in module Base), 472, 473
- `rotl90()` (in module Base), 473
- `rotr90()` (in module Base), 474
- `round()` (in module Base), 416, 417
- `round()` (in module Dates), 594
- `RoundDown` (in module Base), 417
- `rounding()` (in module Base), 437
- `RoundNearest` (in module Base), 417
- `RoundNearestTiesAway` (in module Base), 417
- `RoundNearestTiesUp` (in module Base), 417

RoundToZero (in module Base), 417
RoundUp (in module Base), 417
rowvals() (in module Base), 482
rpad() (in module Base), 446
rsearch() (in module Base), 446
rsearchindex() (in module Base), 446
rsplit() (in module Base), 447
rstrip() (in module Base), 447
RTLD_DEEPBIND (in module Base.Libdl), 615
RTLD_FIRST (in module Base.Libdl), 615
RTLD_GLOBAL (in module Base.Libdl), 615
RTLD_LAZY (in module Base.Libdl), 615
RTLD_LOCAL (in module Base.Libdl), 615
RTLD_NODELETE (in module Base.Libdl), 615
RTLD_NOLOAD (in module Base.Libdl), 615
RTLD_NOW (in module Base.Libdl), 615
run() (in module Base), 370
runtests() (in module Base), 601

S

sbmv
 () (in module Base.LinAlg.BLAS), 541
sbmv() (in module Base.LinAlg.BLAS), 541
scal
 () (in module Base.LinAlg.BLAS), 540
scal() (in module Base.LinAlg.BLAS), 540
scale
 () (in module Base), 525
schedule() (in module Base), 486
schur() (in module Base), 518, 519
schurfact
 () (in module Base), 518, 519
schurfact() (in module Base), 517, 519
sdata() (in module Base), 495
search() (in module Base), 446
searchindex() (in module Base), 446
searchsorted() (in module Base), 579
searchsortedfirst() (in module Base), 579
searchsortedlast() (in module Base), 579
sec() (in module Base), 413
secd() (in module Base), 413
sech() (in module Base), 413
Second (in module Dates), 587, 592
Second() (in module Dates), 590
second() (in module Dates), 590
seek() (in module Base), 563
seekend() (in module Base), 563
seekstart() (in module Base), 563
select
 () (in module Base), 579
select() (in module Base), 579
selectperm
 () (in module Base), 580
selectperm() (in module Base), 580

Semaphore() (in module Base), 498
send() (in module Base), 574
serialize() (in module Base), 564
Set() (in module Base), 396
set_num_threads() (in module Base.LinAlg.BLAS), 542
set_zero_subnormals() (in module Base), 438
setdiff
 () (in module Base), 397
setdiff() (in module Base), 397
setenv() (in module Base), 372
setfield
 () (in module Base), 366
setindex
 () (in module Base), 393, 461
setopt() (in module Base), 574
setprecision() (in module Base), 440
setprotocol
 () (in module Base.Pkg), 584
setrounding() (in module Base), 437, 438
SharedArray() (in module Base), 494
shift
 () (in module Base), 398
show() (in module Base), 566, 570
showall() (in module Base), 566
showcompact() (in module Base), 566
showerror() (in module Base), 567
shuffle
 () (in module Base), 477
shuffle() (in module Base), 477
sign() (in module Base), 419
signbit() (in module Base), 419
signed() (in module Base), 434
signif() (in module Base), 417
significand() (in module Base), 434
similar() (in module Base), 457
sin() (in module Base), 412
sinc() (in module Base), 414
sind() (in module Base), 412
sinh() (in module Base), 412
sinpi() (in module Base), 412
size() (in module Base), 451
sizehint
 () (in module Base), 396
sizeof() (in module Base), 366, 443
skip() (in module Base), 563
skipchars() (in module Base), 565
sleep() (in module Base), 487
slicedim() (in module Base), 461
sort
 () (in module Base), 578
sort() (in module Base), 578, 579
sortcols() (in module Base), 579
sortperm
 () (in module Base), 579

- sortperm() (in module Base), 579
- sortrows() (in module Base), 579
- sparse() (in module Base), 479, 480
- sparsevec() (in module Base), 479, 480
- spawn() (in module Base), 370
- spdiagm() (in module Base), 481
- speye() (in module Base), 481
- SpinLock() (in module Base), 498
- splice
 - () (in module Base), 399, 400
- split() (in module Base), 446
- splitdir() (in module Base), 559
- splitdrive() (in module Base), 559
- splitext() (in module Base), 559
- spones() (in module Base), 480
- sprand() (in module Base), 481
- sprandn() (in module Base), 482
- sprint() (in module Base), 567
- spzeros() (in module Base), 480
- sqrt() (in module Base), 420
- sqrtn() (in module Base), 532
- squeeze() (in module Base), 469
- srand() (in module Base), 440
- StackOverflowError() (in module Base), 375
- start() (in module Base), 381
- startswith() (in module Base), 447
- stat() (in module Base), 556
- status() (in module Base.Pkg), 584
- std() (in module Base), 425
- STDERR (in module Base), 561
- STDIN (in module Base), 561
- stdm() (in module Base), 425
- STDOUT (in module Base), 561
- stebz
 - () (in module Base.LinAlg.LAPACK), 549
- stegr
 - () (in module Base.LinAlg.LAPACK), 549
- stein
 - () (in module Base.LinAlg.LAPACK), 549
- step() (in module Base), 391
- stev
 - () (in module Base.LinAlg.LAPACK), 549
- strerror() (in module Base.Libc), 613
- strftime() (in module Base.Libc), 613
- stride() (in module Base), 453
- strides() (in module Base), 453
- String() (in module Base), 443
- string() (in module Base), 443
- stringmime() (in module Base), 571
- strip() (in module Base), 447
- strptime() (in module Base.Libc), 613
- strwidth() (in module Base), 449
- sub2ind() (in module Base), 454
- subtypes() (in module Base), 365
- success() (in module Base), 371
- sum
 - () (in module Base), 388
- sum() (in module Base), 388
- sum_kbn() (in module Base), 476
- sumabs
 - () (in module Base), 388
- sumabs() (in module Base), 388
- sumabs2
 - () (in module Base), 388
- sumabs2() (in module Base), 388
- summary() (in module Base), 566
- supertype() (in module Base), 365
- svd() (in module Base), 520, 521
- svdfact
 - () (in module Base), 520
- svdfact() (in module Base), 519, 521
- svds() (in module Base), 537
- svdvals
 - () (in module Base), 521
- svdvals() (in module Base), 520, 521
- syconv
 - () (in module Base.LinAlg.LAPACK), 549
- syev
 - () (in module Base.LinAlg.LAPACK), 550
- syevr
 - () (in module Base.LinAlg.LAPACK), 550
- sygvd
 - () (in module Base.LinAlg.LAPACK), 550
- sylvester() (in module Base), 532
- Symbol() (in module Base), 450
- syndiff
 - () (in module Base), 397
- syndiff() (in module Base), 397
- symlink() (in module Base), 555
- symm
 - () (in module Base.LinAlg.BLAS), 541
- symm() (in module Base.LinAlg.BLAS), 541
- Symmetric() (in module Base), 506
- SymTridiagonal() (in module Base), 505
- symv
 - () (in module Base.LinAlg.BLAS), 541
- symv() (in module Base.LinAlg.BLAS), 541, 542
- syr
 - () (in module Base.LinAlg.BLAS), 540
- syrk
 - () (in module Base.LinAlg.BLAS), 540
- syrk() (in module Base.LinAlg.BLAS), 540
- Sys.ARCH (in module Base), 553
- Sys.CPU_CORES (in module Base), 553
- Sys.get_process_title() (in module Base), 371
- Sys.KERNEL (in module Base), 553
- Sys.MACHINE (in module Base), 553
- Sys.set_process_title() (in module Base), 371

`Sys.WORD_SIZE` (in module `Base`), 553

`SystemError()` (in module `Base`), 375

`systemerror()` (in module `Base`), 609

`sysv`

`()` (in module `Base.LinAlg.LAPACK`), 549

`sytrf`

`()` (in module `Base.LinAlg.LAPACK`), 549

`sytri`

`()` (in module `Base.LinAlg.LAPACK`), 549

`sytrs`

`()` (in module `Base.LinAlg.LAPACK`), 550

T

`take`

`()` (in module `Base`), 487, 491, 562

`take()` (in module `Base.Iterators`), 598

`tan()` (in module `Base`), 412

`tand()` (in module `Base`), 412

`tanh()` (in module `Base`), 412

`Task()` (in module `Base`), 485

`task_local_storage()` (in module `Base`), 486

`tempdir()` (in module `Base`), 557

`tempname()` (in module `Base`), 557

`test()` (in module `Base.Pkg`), 584, 585

`TextDisplay()` (in module `Base`), 571

`tgsgen`

`()` (in module `Base.LinAlg.LAPACK`), 551

`threads()` (in module `Base`), 495

`Threads.atomic_add`

`{T}()` (in module `Base`), 496

`Threads.atomic_and`

`{T}()` (in module `Base`), 496

`Threads.atomic_cas`

`{T}()` (in module `Base`), 496

`Threads.atomic_fence()` (in module `Base`), 497

`Threads.atomic_max`

`{T}()` (in module `Base`), 497

`Threads.atomic_min`

`{T}()` (in module `Base`), 497

`Threads.atomic_nand`

`{T}()` (in module `Base`), 496

`Threads.atomic_or`

`{T}()` (in module `Base`), 496

`Threads.atomic_sub`

`{T}()` (in module `Base`), 496

`Threads.atomic_xchg`

`{T}()` (in module `Base`), 496

`Threads.atomic_xor`

`{T}()` (in module `Base`), 496

`Threads.Atomic{T}()` (in module `Base`), 495

`Threads.nthreads()` (in module `Base`), 495

`Threads.threadid()` (in module `Base`), 495

`throw()` (in module `Base`), 374

`tic()` (in module `Base`), 372

`time()` (in module `Base`), 372

`time()` (in module `Base.Libc`), 613

`time_ns()` (in module `Base`), 372

`timedwait()` (in module `Base`), 492

`Timer()` (in module `Base`), 376

`TimeType` (in module `Dates`), 587

`titlecase()` (in module `Base`), 448

`TmStruct()` (in module `Base.Libc`), 613

`toc()` (in module `Base`), 373

`today()` (in module `Dates`), 594

`tofirst()` (in module `Dates`), 592

`tolast()` (in module `Dates`), 592

`tonext()` (in module `Dates`), 592

`toprev()` (in module `Dates`), 592

`toq()` (in module `Base`), 373

`touch()` (in module `Base`), 557

`trace()` (in module `Base`), 527

`trailing_ones()` (in module `Base`), 439

`trailing_zeros()` (in module `Base`), 439

`transcode()` (in module `Base`), 443

`transpose`

`()` (in module `Base`), 534

`transpose()` (in module `Base`), 534

`trcon`

`()` (in module `Base.LinAlg.LAPACK`), 548

`trevc`

`()` (in module `Base.LinAlg.LAPACK`), 549

`trexc`

`()` (in module `Base.LinAlg.LAPACK`), 551

`Tridiagonal()` (in module `Base`), 505, 525

`trigamma()` (in module `Base`), 422

`tril`

`()` (in module `Base`), 524

`tril()` (in module `Base`), 523

`triu`

`()` (in module `Base`), 523

`triu()` (in module `Base`), 522

`trmm`

`()` (in module `Base.LinAlg.BLAS`), 542

`trmm()` (in module `Base.LinAlg.BLAS`), 542

`trmv`

`()` (in module `Base.LinAlg.BLAS`), 542

`trmv()` (in module `Base.LinAlg.BLAS`), 542

`trrfs`

`()` (in module `Base.LinAlg.LAPACK`), 549

`trsen`

`()` (in module `Base.LinAlg.LAPACK`), 551

`trsm`

`()` (in module `Base.LinAlg.BLAS`), 542

`trsm()` (in module `Base.LinAlg.BLAS`), 542

`trsv`

`()` (in module `Base.LinAlg.BLAS`), 542

`trsv()` (in module `Base.LinAlg.BLAS`), 542

`trsyl`

() (in module Base.LinAlg.LAPACK), 551
 trtri
 () (in module Base.LinAlg.LAPACK), 548
 trtrs
 () (in module Base.LinAlg.LAPACK), 548
 trues() (in module Base), 455, 456
 trunc() (in module Base), 417
 trunc() (in module Dates), 592
 truncate() (in module Base), 565
 trylock() (in module Base), 497
 tryparse() (in module Base), 434
 tuple() (in module Base), 363
 TypeError() (in module Base), 375
 typeintersect() (in module Base), 367
 typejoin() (in module Base), 367
 typemax() (in module Base), 365
 typemin() (in module Base), 365
 typeof() (in module Base), 363
 tzrf
 () (in module Base.LinAlg.LAPACK), 545

U

ucfirst() (in module Base), 448
 UndefRefError() (in module Base), 375
 UndefVarError() (in module Base), 375
 unescape_string() (in module Base), 450
 union
 () (in module Base), 397
 union() (in module Base), 397
 unique() (in module Base), 384, 385
 unix2datetime() (in module Dates), 595
 unlock() (in module Base), 497
 unmark() (in module Base), 563
 unsafe_convert() (in module Base), 607
 unsafe_copy
 () (in module Base), 608
 unsafe_get() (in module Base), 370
 unsafe_load() (in module Base), 608
 unsafe_pointer_to_objref() (in module Base), 608
 unsafe_read() (in module Base), 563
 unsafe_store
 () (in module Base), 608
 unsafe_string() (in module Base), 443
 unsafe_trunc() (in module Base), 417
 unsafe_wrap() (in module Base), 444, 608
 unsafe_write() (in module Base), 563
 unshift
 () (in module Base), 398
 unsigned() (in module Base), 434
 update() (in module Base.Pkg), 584
 uperm() (in module Base), 556
 uppercase() (in module Base), 448
 using, 139
 UTInstant{T} (in module Dates), 587

V

Val{c}() (in module Base), 367
 valtype() (in module Base), 396
 values() (in module Base), 395
 var() (in module Base), 425
 varm() (in module Base), 425
 vcat() (in module Base), 462
 vec() (in module Base), 470
 vecdot() (in module Base), 502
 vecnorm() (in module Base), 527
 VERSION (in module Base), 553
 versioninfo() (in module Base), 362
 view() (in module Base), 461

W

wait() (in module Base), 490
 walkdir() (in module Base), 555
 warn() (in module Base), 567
 watch_file() (in module Base), 573
 Week (in module Dates), 587, 592
 Week() (in module Dates), 590
 week() (in module Dates), 590
 which() (in module Base), 361
 whos() (in module Base), 359
 widemul() (in module Base), 424
 widen() (in module Base), 365
 windows_version() (in module Base), 374
 withenv() (in module Base), 372
 WorkerPool() (in module Base), 492
 workers() (in module Base), 489
 workspace() (in module Base), 362
 write() (in module Base), 562
 writcsv() (in module Base), 569
 writedlm() (in module Base), 569

X

xcorr() (in module Base), 431
 xor() (in module Base), 411

Y

Year (in module Dates), 587, 592
 Year() (in module Dates), 590
 year() (in module Dates), 590
 yearmonth() (in module Dates), 590
 yearmonthday() (in module Dates), 591
 yield() (in module Base), 486
 yieldto() (in module Base), 485

Z

zero() (in module Base), 435
 zeros() (in module Base), 454
 zeta() (in module Base), 424
 zip() (in module Base.Iterators), 597